lkan 1.0.27 - A Specification (with	all registered Vulkan extensions)
Vulkan 1.0.27 -	- A Specification (with all registered Vulkan extensions)

Contents

1	Intr	roduction			1
	1.1	What is the Vulkan Graphics System?		1
		1.1.1 The Programmer's View of Vulkan	 		1
		1.1.2 The Implementor's View of Vulkan	 		2
		1.1.3 Our View of Vulkan	 		2
	1.2	Filing Bug Reports	 		2
	1.3	Terminology	 	. .	2
	1.4	Normative References		3
2	Fun	ndamentals			5
	2.1	Architecture Model	 ,	5
	2.2	Execution Model	 ,	6
		2.2.1 Queue Operation	 	. .	6
	2.3	Object Model	 		8
		2.3.1 Object Lifetime	 		8
	2.4	Command Syntax and Duration	 		10
		2.4.1 Lifetime of Retrieved Results		11
	2.5	Threading Behavior	 		11
	2.6	Errors	 		18
		2.6.1 Valid Usage	 		18
		2.6.1.1 Valid Usage for Object Handles	 		19
		2.6.1.2 Valid Usage for Pointers	 		19
		2.6.1.3 Valid Usage for Enumerated Types	 		19
		2.6.1.4 Valid Usage for Flags	 		20
		2.6.1.5 Valid Usage for Structure Types	 	. .	20
		2.6.1.6 Valid Usage for Structure Pointer Chains	 	. .	20
		2.6.1.7 Valid Usage for Nested Structures	 		20

		2.6.2	Return Codes	21
	2.7	Numer	ic Representation and Computation	23
		2.7.1	Floating-Point Computation	23
		2.7.2	16-Bit Floating-Point Numbers	23
		2.7.3	Unsigned 11-Bit Floating-Point Numbers	23
		2.7.4	Unsigned 10-Bit Floating-Point Numbers	24
		2.7.5	General Requirements	24
	2.8	Fixed-	Point Data Conversions	24
		2.8.1	Conversion from Normalized Fixed-Point to Floating-Point	25
		2.8.2	Conversion from Floating-Point to Normalized Fixed-Point	25
	2.9	API Ve	ersion Numbers and Semantics	25
	2.10	Comm	on Object Types	26
		2.10.1	Offsets	26
		2.10.2	Extents	26
		2.10.3	Rectangles	27
3	Initia	alizatio	n	29
	3.1	Comm	and Function Pointers	29
	3.2	Instanc	res	31
4	Devi	ces and	Queues	37
	4.1	Physic	al Devices	37
	4.2	Device	S	42
		4.2.1	Device Creation	43
		4.2.2	Device Use	45
		4.2.3	Lost Device	45
		4.2.4	Device Destruction	46
	4.3	Queue	8	47
		4.3.1	Queue Family Properties	47
		4.3.2	Queue Creation	48
		4.3.3	Queue Family Index	50
		4.3.4	Queue Priority	50
		4.3.5	Queue Submission	50
			4.3.5.1 Sparse Memory Binding	51
		4.3.6	Queue Destruction	51

5	nmand Buffers	53							
	5.1	Command Pools	54						
	5.2	Command Buffer Allocation and Management	58						
	5.3	Command Buffer Recording	62						
	5.4	Command Buffer Submission	67						
	5.5	Queue Forward Progress	72						
	5.6	Secondary Command Buffer Execution	72						
6	Syno	ynchronization and Cache Control 7							
	6.1	Fences	75						
	6.2	Semaphores	81						
	6.3	Events	85						
	6.4	Execution And Memory Dependencies	96						
	6.5	Pipeline Barriers	98						
		6.5.1 Subpass Self-dependency	101						
		6.5.2 Pipeline Stage Flags	102						
		6.5.3 Memory Barriers	104						
		6.5.4 Global Memory Barriers	105						
		6.5.5 Buffer Memory Barriers	108						
		6.5.6 Image Memory Barriers	109						
		6.5.7 Wait Idle Operations	112						
	6.6	Host Write Ordering Guarantees	113						
7	Ren	der Pass	115						
	7.1	Render Pass Creation	116						
	7.2	Render Pass Compatibility	129						
	7.3	Framebuffers	129						
	7.4	Render Pass Commands	133						
8	Shac	ders	141						
	8.1	Shader Modules	141						
	8.2	Shader Execution	144						
	8.3	Shader Memory Access Ordering	144						
	8.4	Shader Inputs and Outputs	145						
	8.5	Vertex Shaders	146						
		8.5.1 Vertex Shader Execution	146						
	8.6	Tessellation Control Shaders	146						

		8.6.1 Tessellation Control Shader Execution	. 146
	8.7	Tessellation Evaluation Shaders	. 147
		8.7.1 Tessellation Evaluation Shader Execution	. 147
	8.8	Geometry Shaders	. 147
		8.8.1 Geometry Shader Execution	. 147
	8.9	Fragment Shaders	. 147
		8.9.1 Fragment Shader Execution	. 147
		8.9.2 Early Fragment Tests	. 148
	8.10	Compute Shaders	. 148
	8.11	Interpolation Decorations	. 148
	8.12	Static Use	. 149
	8.13	Invocation and Derivative Groups	. 149
Λ	D: al		151
9	Pipel	Compute Pipelines	
	9.1	• •	
	9.2	Graphics Pipelines	
	0.2	9.2.1 Valid Combinations of Stages for Graphics Pipelines	
	9.3	Pipeline destruction	
	9.4	Multiple Pipeline Creation	
	9.5	Pipeline Derivatives	
	9.6	Pipeline Cache	
	9.7	Specialization Constants	
	9.8	Pipeline Binding	. 178
10	Mem	ory Allocation	181
	10.1	Host Memory	. 181
	10.2	Device Memory	. 187
		10.2.1 Host Access to Device Memory Objects	. 198
		10.2.2 Lazily Allocated Memory	. 203
11	ъ		205
11		urce Creation	205
		Buffers	
		Buffer Views	
		Images	
		Image Layouts	
		Image Views	
		Resource Memory Association	
		Resource Sharing Mode	
	11.8	Memory Aliasing	. 245

12	Sam	plers		247
13	Reso	urce De	escriptors	253
	13.1	Descrip	ptor Types	254
		13.1.1	Storage Image	254
		13.1.2	Sampler	255
		13.1.3	Sampled Image	256
		13.1.4	Combined Image Sampler	256
		13.1.5	Uniform Texel Buffer	257
		13.1.6	Storage Texel Buffer	257
		13.1.7	Uniform Buffer	258
		13.1.8	Storage Buffer	259
		13.1.9	Dynamic Uniform Buffer	259
		13.1.10	Dynamic Storage Buffer	260
		13.1.11	Input Attachment	260
	13.2	Descrip	ptor Sets	260
		13.2.1	Descriptor Set Layout	260
		13.2.2	Pipeline Layouts	267
			13.2.2.1 Pipeline Layout Compatibility	272
		13.2.3	Allocation of Descriptor Sets	273
		13.2.4	Descriptor Set Updates	282
		13.2.5	Descriptor Set Binding	288
		13.2.6	Push Constant Updates	290
14	Shad	ler Inte	rfaces	293
	14.1	Shader	Input and Output Interfaces	293
		14.1.1	Built-in Interface Block	293
		14.1.2	User-defined Variable Interface	294
		14.1.3	Interface Matching	294
		14.1.4	Location Assignment	295
		14.1.5	Component Assignment	296
	14.2	Vertex	Input Interface	296
	14.3	Fragme	ent Output Interface	296
	14.4	Fragme	ent Input Attachment Interface	297
	14.5	Shader	Resource Interface	298
		14.5.1	Push Constant Interface	298
		14.5.2	Descriptor Set Interface	298
		14.5.3	DescriptorSet and Binding Assignment	300
		14.5.4	Offset and Stride Assignment	301
	14.6	Built-I	n Variables	302

15	Imag	ge Oper	rations	313
	15.1	Image	Operations Overview	313
		15.1.1	Texel Coordinate Systems	314
	15.2	Conver	rsion Formulas	316
		15.2.1	RGB to Shared Exponent Conversion	316
		15.2.2	Shared Exponent to RGB	317
	15.3	Texel I	nput Operations	318
		15.3.1	Texel Input Validation Operations	318
			15.3.1.1 Instruction/Sampler/Image Validation	318
			15.3.1.2 Integer Texel Coordinate Validation	319
			15.3.1.3 Cube Map Edge Handling	320
			15.3.1.4 Sparse Validation	320
		15.3.2	Format Conversion	321
		15.3.3	Texel Replacement	321
		15.3.4	Depth Compare Operation	322
		15.3.5	Conversion to RGBA	323
		15.3.6	Component Swizzle	323
		15.3.7	Sparse Residency	324
	15.4	Texel (Output Operations	324
		15.4.1	Texel Output Validation Operations	324
			15.4.1.1 Texel Format Validation	324
		15.4.2	Integer Texel Coordinate Validation	324
		15.4.3	Sparse Texel Operation	324
		15.4.4	Texel Output Format Conversion	325
	15.5	Deriva	tive Operations	325
	15.6	Norma	lized Texel Coordinate Operations	326
		15.6.1	Projection Operation	326
		15.6.2	Derivative Image Operations	326
		15.6.3	Cube Map Face Selection and Transformations	327
		15.6.4	Cube Map Face Selection	327
		15.6.5	Cube Map Coordinate Transformation	328
		15.6.6	Cube Map Derivative Transformation	328
		15.6.7	Scale Factor Operation, Level-of-Detail Operation and Image Level(s) Selection $\ \ldots \ \ldots \ \ldots$	328
			15.6.7.1 Scale Factor Operation	328
			15.6.7.2 Level-of-Detail Operation	330

		15.6.7.3 Image Level(s) Selection	330
		15.6.8 (s,t,r,q,a) to (u,v,w,a) Transformation	331
	15.7	Unnormalized Texel Coordinate Operations	331
		$15.7.1 \ \ (u,v,w,a) \ to \ (i,j,k,l,n) \ Transformation \ And \ Array \ Layer \ Selection \ . \ . \ . \ . \ . \ . \ . \ . \ . \ $	331
	15.8	Image Sample Operations	332
		15.8.1 Wrapping Operation	332
		15.8.2 Texel Gathering	333
		15.8.3 Texel Filtering	333
		15.8.4 Texel Anisotropic Filtering	335
	15.9	Image Operation Steps	335
16	Quer	ries	337
	_	Query Pools	337
		Query Operation	
		Occlusion Queries	
		Pipeline Statistics Queries	
	16.5	Timestamp Queries	352
17	Class	r Commands	355
1/			
		Clearing Images Outside A Render Pass Instance	
		Clearing Images Inside A Render Pass Instance	
		Clear Values	
		Filling Buffers	
	17.5	Updating Buffers	364
18	Copy	y Commands	367
	18.1	Common Operation	367
	18.2	Copying Data Between Buffers	368
	18.3	Copying Data Between Images	370
	18.4	Copying Data Between Buffers and Images	375
	18.5	Image Copies with Scaling	381
	18.6	Resolving Multisample Images	387

19	Drav	wing Commands	391
	19.1	Primitive Topologies	392
		19.1.1 Points	392
		19.1.2 Separate Lines	393
		19.1.3 Line Strips	393
		19.1.4 Triangle Strips	393
		19.1.5 Triangle Fans	393
		19.1.6 Separate Triangles	393
		19.1.7 Lines With Adjacency	394
		19.1.8 Line Strips With Adjacency	394
		19.1.9 Triangle List With Adjacency	394
		19.1.10 Triangle Strips With Adjacency	395
		19.1.11 Separate Patches	396
		19.1.12 General Considerations For Polygon Primitives	396
	19.2	Programmable Primitive Shading	397
20	Fixe	d-Function Vertex Processing	417
		Vertex Attributes	
		20.1.1 Attribute Location and Component Assignment	
	20.2	Vertex Input Description	
		Example	
21		ellation	427
		Tessellator	
		Tessellator Patch Discard	
		Tessellator Spacing	
		Quad Tessellation	432
		Isoline Tessellation	434
	21.7	Tessellation Pipeline State	434
22	Geor	metry Shading	437
	22.1	Geometry Shader Input Primitives	437
	22.2	Geometry Shader Output Primitives	438
	22.3	Multiple Invocations of Geometry Shaders	438
	22.4	Geometry Shader Primitive Ordering	439

23	3 Fixed-Function Vertex Post-Processing	441
	23.1 Flat Shading	441
	23.2 Primitive Clipping	442
	23.3 Clipping Shader Outputs	443
	23.4 Coordinate Transformations	444
	23.5 Controlling the Viewport	444
24	4 Rasterization	449
	24.1 Discarding Primitives Before Rasterization	452
	24.2 Rasterization Order	452
	24.3 Multisampling	453
	24.4 Sample Shading	455
	24.5 Points	455
	24.5.1 Basic Point Rasterization	456
	24.6 Line Segments	456
	24.6.1 Basic Line Segment Rasterization	457
	24.7 Polygons	459
	24.7.1 Basic Polygon Rasterization	459
	24.7.2 Polygon Mode	461
	24.7.3 Depth Bias	462
25	5 Fragment Operations	465
	25.1 Early Per-Fragment Tests	
	25.2 Scissor Test	
	25.3 Sample Mask	
	25.4 Early Fragment Test Mode	
	25.5 Late Per-Fragment Tests	
	25.6 Multisample Coverage	
	25.7 Depth and Stencil Operations	
	25.8 Depth Bounds Test	
	25.9 Stencil Test	
	25.10Depth Test	
	25.11Sample Counting	
	25/11/5ample Counting	
26	6 The Framebuffer	479
	26.1 Blending	
	26.1.1 Blend Factors	
	26.1.2 Dual-Source Blending	
	26.1.3 Blend Operations	
	26.2 Logical Operations	486

27	Disp	atching	g Commands	489
28	Spar	se Reso	ources	495
	28.1	Sparse	Resource Features	495
	28.2	Sparse	Buffers and Fully-Resident Images	496
		28.2.1	Sparse Buffer and Fully-Resident Image Block Size	497
	28.3	Sparse	Partially-Resident Buffers	497
	28.4	Sparse	Partially-Resident Images	497
		28.4.1	Accessing Unbound Regions	497
		28.4.2	Mip Tail Regions	498
		28.4.3	Standard Sparse Image Block Shapes	502
		28.4.4	Custom Sparse Image Block Shapes	504
		28.4.5	Multiple Aspects	505
			28.4.5.1 Metadata	505
	28.5	Sparse	Memory Aliasing	506
	28.6	Sparse	Resource Implementation Guidelines	506
	28.7	Sparse	Resource API	508
		28.7.1	Physical Device Features	508
			28.7.1.1 Sparse Physical Device Features	508
		28.7.2	Physical Device Sparse Properties	509
		28.7.3	Sparse Image Format Properties	510
			28.7.3.1 Sparse Image Format Properties API	510
		28.7.4	Sparse Resource Creation	512
		28.7.5	Sparse Resource Memory Requirements	512
			28.7.5.1 Buffer and Fully-Resident Images	512
			28.7.5.2 Partially Resident Images	512
			28.7.5.3 Sparse Image Memory Requirements	512
		28.7.6	Binding Resource Memory	514
			28.7.6.1 Sparse Memory Binding Functions	515
	28.8	Examp	oles	523
		28.8.1	Basic Sparse Resources	523
		28.8.2	Advanced Sparse Resources	524

29	Wind	dow System Integration (WSI)	529
	29.1	WSI Platform	529
	29.2	WSI Surface	529
		29.2.1 Android Platform	530
		29.2.2 Mir Platform	531
		29.2.3 Wayland Platform	533
		29.2.4 Win32 Platform	535
		29.2.5 XCB Platform	537
		29.2.6 Xlib Platform	538
		29.2.7 Platform-Independent Information	540
	29.3	Presenting Directly to Display Devices	541
		29.3.1 Display Enumeration	541
		29.3.1.1 Display Planes	544
		29.3.1.2 Display Modes	546
		29.3.2 Display Surfaces	552
	29.4	Querying for WSI Support	555
		29.4.1 Android Platform	556
		29.4.2 Mir Platform	556
		29.4.3 Wayland Platform	557
		29.4.4 Win32 Platform	557
		29.4.5 XCB Platform	558
		29.4.6 Xlib Platform	559
	29.5	Surface Queries	559
	29.6	WSI Swapchain	566
30	Exte	nded Functionality	585
			585
		30.1.1 Device Layer Deprecation	587
	30.2	Extensions	
		30.2.1 Instance Extensions and Device Extensions	
21	E 4		502
31		ures, Limits, and Formats	593
	31.1	Features	
	21.2	31.1.1 Feature Requirements	
	31.2	Limits	
		31.2.1 Limit Requirements	013

	31.3	Formats	619
		31.3.1 Format Definition	619
		31.3.1.1 Packed Formats	635
		31.3.1.2 Identification of Formats	636
		31.3.1.3 Representation	637
		31.3.1.4 Depth/Stencil Formats	639
		31.3.1.5 Format Compatibility Classes	639
		31.3.2 Format Properties	644
		31.3.3 Required Format Support	646
	31.4	Additional Image Capabilities	658
		31.4.1 Supported Sample Counts	663
		31.4.2 Allowed Extent Values Based On Image Type	664
32	Debu	ugging	665
	32.1	Object Annotation	665
	32.2	Command Buffer Markers	671
33	Glos	ssary	675
34	Com	nmon Abbreviations	687
35	Prefi	ixes	689
A	Vulk	can Environment for SPIR-V	691
	A.1	Required Versions and Formats	691
	A.2	Capabilities	691
	A.3	Validation Rules within a Module	692
	A.4	Precision and Operation of SPIR-V Instructions	694
В	Com	npressed Image Formats	697
	B.1	Block-Compressed Image Formats	698
	B.2	ETC Compressed Image Formats	699
	B.3	ASTC Compressed Image Formats	700

C	Lay	ers & E	xtensions 7	01
	C.1	VK_K	HR_sampler_mirror_clamp_to_edge	01
		C.1.1	New Enum Constants	02
		C.1.2	Example	02
		C.1.3	Version History	02
	C.2	Windo	w System Integration (WSI) Extensions	03
		C.2.1	Editors	03
		C.2.2	VK_KHR_surface	03
			C.2.2.1 New Object Types	04
			C.2.2.2 New Enum Constants	04
			C.2.2.3 New Enums	04
			C.2.2.4 New Structures	04
			C.2.2.5 New Functions	04
			C.2.2.6 Examples	05
			C.2.2.7 Issues	07
			C.2.2.8 Version History	08
		C.2.3	VK_KHR_swapchain	11
			C.2.3.1 New Object Types	12
			C.2.3.2 New Enum Constants	12
			C.2.3.3 New Enums	12
			C.2.3.4 New Structures	12
			C.2.3.5 New Functions	12
			C.2.3.6 Issues	13
			C.2.3.7 Examples	19
			C.2.3.8 Version History	30
		C.2.4	VK_KHR_display	37
			C.2.4.1 New Object Types	38
			C.2.4.2 New Enum Constants	38
			C.2.4.3 New Enums	38
			C.2.4.4 New Structures	38
			C.2.4.5 New Functions	38
			C.2.4.6 Issues	39
			C.2.4.7 Examples	41
			C.2.4.8 Version History	44
		C.2.5	VK_KHR_display_swapchain	47

	C.2.5.1	New Object Types
	C.2.5.2	New Enum Constants
	C.2.5.3	New Enums
	C.2.5.4	New Structures
	C.2.5.5	New Functions
	C.2.5.6	Issues
	C.2.5.7	Examples
	C.2.5.8	Version History
C.2.6	VK_KHI	R_android_surface
	C.2.6.1	New Object Types
	C.2.6.2	New Enum Constants
	C.2.6.3	New Enums
	C.2.6.4	New Structures
	C.2.6.5	New Functions
	C.2.6.6	Issues
	C.2.6.7	Version History
C.2.7	VK_KHI	R_mir_surface
	C.2.7.1	New Object Types
	C.2.7.2	New Enum Constants
	C.2.7.3	New Enums
	C.2.7.4	New Structures
	C.2.7.5	New Functions
	C.2.7.6	Issues
	C.2.7.7	Version History
C.2.8	VK_KHI	R_wayland_surface
	C.2.8.1	New Object Types
	C.2.8.2	New Enum Constants
	C.2.8.3	New Enums
	C.2.8.4	New Structures
	C.2.8.5	New Functions
	C.2.8.6	Issues
	C.2.8.7	Version History
C.2.9	VK_KHI	R_win32_surface
	C.2.9.1	New Object Types
	C.2.9.2	New Enum Constants

		C.2.9.3	New Enums		759		
		C.2.9.4	New Structures		759		
		C.2.9.5	New Functions		759		
		C.2.9.6	Issues		759		
		C.2.9.7	Version History		759		
	C.2.10	VK_KHR	R_xcb_surface		760		
		C.2.10.1	New Object Types		761		
		C.2.10.2	New Enum Constants		761		
		C.2.10.3	New Enums		761		
		C.2.10.4	New Structures		761		
		C.2.10.5	New Functions		761		
		C.2.10.6	Issues		761		
		C.2.10.7	Version History		761		
	C.2.11	VK_KHR	R_xlib_surface		762		
		C.2.11.1	New Object Types		763		
		C.2.11.2	New Enum Constants		763		
		C.2.11.3	New Enums		763		
		C.2.11.4	New Structures		763		
		C.2.11.5	New Functions		763		
		C.2.11.6	Issues		763		
		C.2.11.7	Version History		764		
C.3	VK_EXT_debug_marker						
	C.3.1	New Obje	ect Types		765		
	C.3.2	New Enu	ım Constants		765		
	C.3.3	New Enu	ıms		765		
	C.3.4	New Stru	actures		765		
	C.3.5	New Fund	actions		765		
	C.3.6	Examples	s		766		
	C.3.7	Issues .			767		
	C.3.8	Version H	History		768		
C.4	VK_EX	XT_debug_	_report		768		
C.5	VK_A	MD_draw_	_indirect_count		774		
	C.5.1	New Fund	actions		775		
	C.5.2	Version H	History		775		
C.6	VK_Al	MD_gcn_s	shader		775		

	C.6.1 V	Yersion History	776
C.7	VK_AMI	D_rasterization_order	776
	C.7.1 N	Tew Object Types	777
	C.7.2 N	Tew Enum Constants	777
	C.7.3 N	Tew Enums	778
	C.7.4 N	Tew Structures	778
	C.7.5 N	Tew Functions	778
	C.7.6 Is	ssues	778
	C.7.7 E	xamples	778
	C.7.8 Is	ssues	778
	C.7.9 V	Yersion History	779
C.8	VK_AMI	D_shader_explicit_vertex_parameter	779
	C.8.1 V	Yersion History	779
C.9	VK_AMI	D_shader_trinary_minmax	780
	C.9.1 V	Yersion History	780
C.10	VK_AMI	D_negative_viewport_height	780
C.11	VK_IMG	f_filter_cubic	781
	C.11.1 N	Tew Enum Constants	782
	C.11.2 E	xample	782
	C.11.3 V	Yersion History	782
C.12	VK_NV_	dedicated_allocation	782
	C.12.1 N	Tew Object Types	783
	C.12.2 N	Tew Enum Constants	783
	C.12.3 N	lew Enums	783
	C.12.4 N	Tew Structures	784
	C.12.5 N	lew Functions	784
	C.12.6 Is	ssues	784
	C.12.7 E	xamples	784
	C.12.8 V	Yersion History	785
C.13	VK_NV_	glsl_shader	785
	C.13.1 N	Iew Object Types	786
	C.13.2 N	Iew Enum Constants	786
	C.13.3 N	Iew Enums	786
	C.13.4 N	lew Structures	786
	C.13.5 N	lew Functions	786

	C.13.6 Issues	786
	C.13.7 Examples	786
	C.13.8 Version History	787
C.14	VK_NV_external_memory_capabilities	787
	C.14.1 New Object Types	788
	C.14.2 New Enum Constants	788
	C.14.3 New Enums	788
	C.14.4 New Structs	788
	C.14.5 New Functions	788
	C.14.6 Issues	788
C.15	VK_NV_external_memory	789
	C.15.1 New Object Types	790
	C.15.2 New Enum Constants	790
	C.15.3 New Enums	790
	C.15.4 New Structures	790
	C.15.5 New Functions	790
	C.15.6 Issues	790
	C.15.7 Examples	791
	C.15.8 Version History	791
C.16	VK_NV_external_memory_win32	791
	C.16.1 New Object Types	792
	C.16.2 New Enum Constants	792
	C.16.3 New Enums	792
	C.16.4 New Structures	792
	C.16.5 New Functions	792
	C.16.6 Issues	793
	C.16.7 Examples	793
	C.16.8 Version History	796
C.17	VK_NV_win32_keyed_mutex	796
	C.17.1 New Object Types	797
	C.17.2 New Enum Constants	797
	C.17.3 New Enums	797
	C.17.4 New Structures	797
	C.17.5 New Functions	797
	C.17.6 Issues	797
	C.17.7 Examples	797
	C.17.8 Version History	800

D	API Boilerplate				
	D.1	Structu	re Types	801	
	D.2	Flag Ty	pes	802	
	D.3	Macro	Definitions in vulkan.h	805	
		D.3.1	Vulkan Version Number Macros	805	
		D.3.2	Vulkan Header File Version Number	806	
		D.3.3	Vulkan Handle macros	806	
	D.4	Platfori	n-Specific Macro Definitions in vk_platform.h	807	
		D.4.1	Platform-Specific Calling Conventions	807	
		D.4.2	Platform-Specific Header Control	807	
		D.4.3	Window System-Specific Header Control	808	
E	Inva	riance		809	
	E.1	Repeata	ability	809	
	E.2	Multi-p	ass Algorithms	809	
	E.3	Invaria	nce Rules	810	
	E.4	Tessellation Invariance			
F	Cred	lits		813	

Copyright © 2014-2016 The Khronos Group Inc. All Rights Reserved.

This specification is protected by copyright laws and contains material proprietary to the Khronos Group, Inc. It or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast or otherwise exploited in any manner without the express prior written permission of Khronos Group. You may use this specification for implementing the functionality therein, without altering or removing any trademark, copyright or other notice from the specification, but the receipt or possession of this specification does not convey any rights to reproduce, disclose, or distribute its contents, or to manufacture, use, or sell anything that it may describe, in whole or in part.

Khronos Group grants express permission to any current Promoter, Contributor or Adopter member of Khronos to copy and redistribute UNMODIFIED versions of this specification in any fashion, provided that NO CHARGE is made for the specification and the latest available update of the specification for any version of the API is used whenever possible. Such distributed specification may be reformatted AS LONG AS the contents of the specification are not changed in any way. The specification may be incorporated into a product that is sold as long as such product includes significant independent work developed by the seller. A link to the current version of this specification on the Khronos Group web-site should be included whenever possible with specification distributions.

This specification has been created under the Khronos Intellectual Property Rights Policy, which is Attachment A of the Khronos Group Membership Agreement available at www.khronos.org/files/member_agreement.pdf. This specification contains substantially unmodified functionality from, and is a successor to, Khronos specifications including OpenGL, OpenGL ES and OpenCL.

Some parts of this Specification are purely informative and do not define requirements necessary for compliance and so are outside the Scope of this Specification. These parts of the Specification are marked by the "Note" icon or designated "Informative".

Where this Specification uses terms, defined in the Glossary or otherwise, that refer to enabling technologies that are not expressly set forth as being required for compliance, those enabling technologies are outside the Scope of this Specification.

Where this Specification uses the terms "may", or "optional", such features or behaviors do not define requirements necessary for compliance and so are outside the Scope of this Specification.

Where this Specification uses the terms "not required", such features or behaviors may be omitted from certain implementations, but when they are included, they define requirements necessary for compliance and so are INCLUDED in the Scope of this Specification.

Where this Specification includes normative references to external documents, the specifically identified sections and functionality of those external documents are in Scope. Requirements defined by external documents not created by Khronos may contain contributions from non-members of Khronos not covered by the Khronos Intellectual Property Rights Policy.

Khronos Group makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation, any implied warranties of merchantability or fitness for a particular purpose or non-infringement of any intellectual property. Khronos Group makes no, and expressly disclaims any, warranties, express or implied, regarding the correctness, accuracy, completeness, timeliness, and reliability of the specification. Under no circumstances will the Khronos Group, or any of its Promoters, Contributors or Members or their respective partners, officers, directors, employees, agents or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Khronos and Vulkan are trademarks of The Khronos Group Inc. OpenCL is a trademark of Apple Inc. and OpenGL is a registered trademark of Silicon Graphics International, both used under license by Khronos.

Chapter 1

Introduction

This chapter is Informative except for the sections on Terminology and Normative References.

This document, referred to as the "Vulkan Specification" or just the "Specification" hereafter, describes the Vulkan graphics system: what it is, how it acts, and what is required to implement it. We assume that the reader has at least a rudimentary understanding of computer graphics. This means familiarity with the essentials of computer graphics algorithms and terminology as well as with modern GPUs (Graphic Processing Units).

The canonical version of the Specification is available in the official Vulkan Registry, located at URL

http://www.khronos.org/registry/vulkan/

1.1 What is the Vulkan Graphics System?

Vulkan is an API (Application Programming Interface) for graphics and compute hardware. The API consists of many commands that allow a programmer to specify shader programs, compute kernels, objects, and operations involved in producing high-quality graphical images, specifically color images of three-dimensional objects.

1.1.1 The Programmer's View of Vulkan

To the programmer, Vulkan is a set of commands that allow the specification of *shader programs* or *shaders*, *kernels*, data used by kernels or shaders, and state controlling aspects of Vulkan outside the scope of shaders. Typically, the data represents geometry in two or three dimensions and texture images, while the shaders and kernels control the processing of the data, rasterization of the geometry, and the lighting and shading of *fragments* generated by rasterization, resulting in the rendering of geometry into the framebuffer.

A typical Vulkan program begins with platform-specific calls to open a window or otherwise prepare a display device onto which the program will draw. Then, calls are made to open *queues* to which *command buffers* are submitted. The command buffers contain lists of commands which will be executed by the underlying hardware. The application can also allocate device memory, associate *resources* with memory and refer to these resources from within command buffers. Drawing commands cause application-defined shader programs to be invoked, which can then consume the data in the resources and use them to produce graphical images. To display the resulting images, further platform-specific commands are made to transfer the resulting image to a display device or window.

1.1.2 The Implementor's View of Vulkan

To the implementor, Vulkan is a set of commands that allow the construction and submission of command buffers to a device. Modern devices accelerate virtually all Vulkan operations, storing data and framebuffer images in high-speed memory and executing shaders in dedicated GPU processing resources.

The implementor's task is to provide a software library on the host which implements the Vulkan API, while mapping the work for each Vulkan command to the graphics hardware as appropriate for the capabilities of the device.

1.1.3 Our View of Vulkan

We view Vulkan as a pipeline having some programmable stages and some state-driven fixed-function stages that are invoked by a set of specific drawing operations. We expect this model to result in a specification that satisfies the needs of both programmers and implementors. It does not, however, necessarily provide a model for implementation. An implementation must produce results conforming to those produced by the specified methods, but may carry out particular computations in ways that are more efficient than the one specified.

1.2 Filing Bug Reports

Issues with and bug reports on the Vulkan Specification and the API Registry can be filed in the Khronos Vulkan GitHub repository, located at URL

http://github.com/KhronosGroup/Vulkan-Docs

Please tag issues with appropriate labels, such as "Specification", "Ref Pages" or "Registry", to help us triage and assign them appropriately. Unfortunately, GitHub does not currently let users who do not have write access to the repository set GitHub labels on issues. In the meantime, they can be added to the title line of the issue set in brackets, e.g. *'[Specification]*".

1.3 Terminology

The key words **must**, **required**, **shall should**, **recommend**, **may**, and **optional** in this document are to be interpreted as described in RFC 2119:

http://www.ietf.org/rfc/rfc2119.txt

must

When used alone, this word, or the term **required**, means that the definition is an absolute requirement of the specification. When followed by **not** ("must not"), the phrase means that the definition is an absolute prohibition of the specification.

should

When used alone, this word, or the adjective **recommended**, means that there may exist valid reasons in particular circumstances to ignore a particular item, but the full implications must be understood and carefully weighed before choosing a different course. When followed by **not** ("should not"), the phrase means that there may exist valid reasons in particular circumstances when the particular behavior is acceptable or even useful, but the full implications should be understood and the case carefully weighed before implementing any behavior described with this label.

may

This word, or the adjective **optional**, means that an item is truly optional. One vendor may choose to include the item because a particular marketplace requires it or because the vendor feels that it enhances the product while another vendor may omit the same item. An implementation which does not include a particular option must be prepared to interoperate with another implementation which does include the option, though perhaps with reduced functionality. In the same vein an implementation which does include a particular option must be prepared to interoperate with another implementation which does not include the option (except, of course, for the feature the option provides).

The additional terms **can** and **cannot** are to be interpreted as follows:

can

This word means that the particular behavior described is a valid choice for an application, and is never used to refer to implementation behavior.

cannot

This word means that the particular behavior described is not achievable by an application. For example, an entry point does not exist, or shader code is not capable of expressing an operation.



August, 2008.

Note

There is an important distinction between **cannot** and **must not**, as used in this Specification. **Cannot** means something the application literally is unable to express or accomplish through the API, while **must not** means something that the application is capable of expressing through the API, but that the consequences of doing so are undefined and potentially unrecoverable for the implementation.

1.4 Normative References

Normative references are references to external documents or resources to which implementers of Vulkan must comply. *IEEE Standard for Floating-Point Arithmetic*, IEEE Std 754-2008, http://dx.doi.org/10.1109/IEEESTD.2008.4610935,

A. Garrard, *Khronos Data Format Specification*, version 1.1, https://www.khronos.org/registry/dataformat/specs/1.1/dataformat.1.1.html, June, 2016.

- J. Kessenich, SPIR-V Extended Instructions for GLSL, Version 1.00, https://www.khronos.org/registry/spir-v/, February 10, 2016.
- J. Kessenich and B. Ouriel, *The Khronos SPIR-V Specification*, *Version 1.00*, https://www.khronos.org/registry/spir-v/, February 10, 2016.
- J. Leech and T. Hector, *Vulkan Documentation and Extensions: Procedures and Conventions*, https://www.khronos.org/registry/vulkan/, July 11, 2016

Vulkan Loader Specification and Architecture Overview, https://github.com/KhronosGroup/Vulkan-LoaderAndValidationLayers/blob/master/loader/LoaderAndLayerInterface.md, August, 2016.

Chapter 2

Fundamentals

This chapter introduces fundamental concepts including the Vulkan architecture and execution model, API syntax, queues, pipeline configurations, numeric representation, state and state queries, and the different types of objects and shaders. It provides a framework for interpreting more specific descriptions of commands and behavior in the remainder of the Specification.

2.1 Architecture Model

Vulkan is designed for, and the API is written for, CPU, GPU, and other hardware accelerator architectures with the following properties:

- Runtime support for 8, 16, 32 and 64-bit signed and unsigned twos-complement integers, all addressable at the granularity of their size in bytes.
- Runtime support for 32- and 64-bit floating-point types satisfying the range and precision constraints in the Floating Point Computation section.
- The representation and endianness of these types must be identical for the host and the physical devices.



Note

Since a variety of data types and structures in Vulkan may be mapped back and forth between host and physical device memory, host and device architectures must both be able to access such data efficiently in order to write portable and performant applications.

Where the Specification leaves choices open that would affect Application Binary Interface compatibility on a given platform supporting Vulkan, those choices are usually made to be compliant to the preferred ABI defined by the platform vendor. Some choices, such as function calling conventions, may be made in platform-specific portions of the vk_platform.h header file.



Note

For example, the Android ABI is defined by Google, and the Linux ABI is defined by a combination of gcc defaults, distribution vendor choices, and external standards such as the Linux Standard Base.

2.2 Execution Model

This section outlines the execution model of a Vulkan system.

Vulkan exposes one or more *devices*, each of which exposes one or more *queues* which may process work asynchronously to one another. The set of queues supported by a device is partitioned into *families*. Each family supports one or more types of functionality and may contain multiple queues with similar characteristics. Queues within a single family are considered *compatible* with one another, and work produced for a family of queues can be executed on any queue within that family. This specification defines four types of functionality that queues may support: graphics, compute, transfer, and sparse memory management.



Note

A single device may report multiple similar queue families rather than, or as well as, reporting multiple members of one or more of those families. This indicates that while members of those families have similar capabilities, they are *not* directly compatible with one another.

Device memory is explicitly managed by the application. Each device may advertise one or more heaps, representing different areas of memory. Memory heaps are either device local or host local, but are always visible to the device. Further detail about memory heaps is exposed via memory types available on that heap. Examples of memory areas that may be available on an implementation include:

- device local is memory that is physically connected to the device.
- device local, host visible is device local memory that is visible to the host.
- host local, host visible is memory that is local to the host and visible to the device and host.

On other architectures, there may only be a single heap that can be used for any purpose.

A Vulkan application controls a set of devices through the submission of command buffers which have recorded device commands issued via Vulkan library calls. The content of command buffers is specific to the underlying hardware and is opaque to the application. Once constructed, a command buffer can be submitted once or many times to a queue for execution. Multiple command buffers can be built in parallel by employing multiple threads within the application.

Command buffers submitted to different queues may execute in parallel or even out of order with respect to one another. Command buffers submitted to a single queue respect the submission order, as described further in Queue Operation. Command buffer execution by the device is also asynchronous to host execution. Once a command buffer is submitted to a queue, control may return to the application immediately. Synchronization between the device and host, and between different queues is the responsibility of the application.

2.2.1 Queue Operation

Vulkan queues provide an interface to the execution engines of a device. Commands for these execution engines are recorded into command buffers ahead of execution time. These command buffers are then submitted to queues with a *queue submission* command for execution in a number of *batches*. Once submitted to a queue, these commands will begin and complete execution without further application intervention, though the order of this execution is dependent on a number of implicit and explicit ordering constraints.

Work is submitted to queues using queue submission commands that typically take the form **vkQueue*** (e.g. vkQueueSubmit, vkQueueBindSparse), and optionally take a list of semaphores upon which to wait before work begins and a list of semaphores to signal once work has completed. The work itself, as well as signaling and waiting on the semaphores are all *queue operations*.

Queue operations on different queues have no implicit ordering constraints, and may execute in any order. Explicit ordering constraints between queues can be expressed with semaphores and fences.

Command buffer submissions to a single queue must always adhere to command order and API order, but otherwise may overlap or execute out of order. Other types of batches and queue submissions against a single queue (e.g. sparse memory binding) have no implicit ordering constraints with any other queue submission or batch. Additional explicit ordering constraints between queue submissions and individual batches can be expressed with semaphores and fences.

Before a fence or semaphore is signaled, it is guaranteed that any previously submitted queue operations have completed execution, and that memory writes from those queue operations are available to future queue operations. Waiting on a signaled semaphore or fence guarantees that previous writes that are available are also visible to subsequent commands.

Command buffer boundaries, both between primary command buffers of the same or different batches or submissions as well as between primary and secondary command buffers, do not introduce any implicit ordering constraints. In other words, submitting the set of command buffers (which can include executing secondary command buffers) between any semaphore or fence operations execute the recorded commands as if they had all been recorded into a single primary command buffer, except that the current state is reset on each boundary. Explicit ordering constraints can be expressed with events and pipeline barriers.

There are a few implicit ordering constraints between commands within a command buffer, but only covering a subset of execution. Additional explicit ordering constraints can be expressed with events, pipeline barriers and subpass dependencies.



Note

Implementations have significant freedom to overlap execution of work submitted to a queue, and this is common due to deep pipelining and parallelism in Vulkan devices.

Commands recorded in command buffers either perform actions (draw, dispatch, clear, copy, query/timestamp operations, begin/end subpass operations), set state (bind pipelines, descriptor sets, and buffers, set dynamic state, push constants, set render pass/subpass state), or perform synchronization (set/wait events, pipeline barrier, render pass/subpass dependencies). Some commands perform more than one of these tasks. State setting commands update the *current state* of the command buffer. Some commands that perform actions (e.g. draw/dispatch) do so based on the current state set cumulatively since the start of the command buffer. The work involved in performing action commands is often allowed to overlap or to be reordered, but doing so must not alter the state to be used by each action command. In general, action commands are those commands that alter framebuffer attachments, read/write buffer or image memory, or write to query pools.

Synchronization commands introduce explicit execution and memory dependencies between two sets of action commands, where the second set of commands depends on the first set of commands. These dependencies enforce that both the execution of certain pipeline stages in the later set occur after the execution of certain stages in the source set, and that the effects of memory accesses performed by certain pipeline stages occur in order and are visible to each other. When not enforced by an explicit dependency or otherwise forbidden by the specification, action commands may overlap execution or execute out of order, and may not see the side effects of each other's memory accesses.

The execution order of an action command with respect to any synchronization commands that affect that action command must match the recording and submission order, within submissions to a single queue.

API order sorts primitives:

- First, by the action command that generates them.
- Second, by the order they are processed by primitive assembly.

Within this order, implementations also sort primitives:

- Third, by an implementation-dependent ordering of new primitives generated by tessellation, if a tessellation shader is
 active.
- Fourth, by the order new primitives are generated by geometry shading, if geometry shading is active.
- Fifth, by an implementation-dependent ordering of primitives generated due to the polygon mode.

The device executes queue operations asynchronously with respect to the host. Control is returned to an application immediately following command buffer submission to a queue. The application must synchronize work between the host and device as needed.

2.3 Object Model

The devices, queues, and other entities in Vulkan are represented by Vulkan objects. At the API level, all objects are referred to by handles. There are two classes of handles, dispatchable and non-dispatchable. *Dispatchable* handle types are a pointer to an opaque type. This pointer may be used by layers as part of intercepting API commands, and thus each API command takes a dispatchable type as its first parameter. Each object of a dispatchable type must have a unique handle value during its lifetime.

Non-dispatchable handle types are a 64-bit integer type whose meaning is implementation-dependent, and may encode object information directly in the handle rather than pointing to a software structure. Objects of a non-dispatchable type may not have unique handle values within a type or across types. If handle values are not unique, then destroying one such handle must not cause identical handles of other types to become invalid, and must not cause identical handles of the same type to become invalid if that handle value has been created more times than it has been destroyed.

All objects created or allocated from a VkDevice (i.e. with a VkDevice as the first parameter) are private to that device, and must not be used on other devices.

2.3.1 Object Lifetime

Objects are created or allocated by **vkCreate*** and **vkAllocate*** commands, respectively. Once an object is created or allocated, its "structure" is considered to be immutable, though the contents of certain object types is still free to change. Objects are destroyed or freed by **vkDestroy*** and **vkFree*** commands, respectively.

Objects that are allocated (rather than created) take resources from an existing pool object or memory heap, and when freed return resources to that pool or heap. While object creation and destruction are generally expected to be low-frequency occurrences during runtime, allocating and freeing objects can occur at high frequency. Pool objects help accommodate improved performance of the allocations and frees.

It is an application's responsibility to track the lifetime of Vulkan objects, and not to destroy them while they are still in

Application-owned memory is immediately consumed by any Vulkan command it is passed into. The application can alter or free this memory as soon as the commands that consume it have returned.

The following object types are consumed when they are passed into a Vulkan command and not further accessed by the objects they are used to create. They must not be destroyed in the duration of any API command they are passed into:

- VkShaderModule
- VkPipelineCache

A VkPipelineLayout object must not be destroyed while any command buffer that uses it is in the recording state.

VkDescriptorSetLayout objects may be accessed by commands that operate on descriptor sets allocated using that layout, and those descriptor sets must not be updated with vkUpdateDescriptorSets after the descriptor set layout has been destroyed. Otherwise, descriptor set layouts can be destroyed any time they are not in use by an API command.

The application must not destroy any other type of Vulkan object until all uses of that object by the device (such as via command buffer execution) have completed.

The following Vulkan objects must not be destroyed while any command buffers using the object are recording or pending execution:

- VkEvent
- VkQueryPool
- VkBuffer
- VkBufferView
- VkImage
- VkImageView
- VkPipeline
- VkSampler
- VkDescriptorPool
- VkFramebuffer
- VkRenderPass
- VkCommandPool
- VkDeviceMemory
- VkDescriptorSet

The following Vulkan objects must not be destroyed while any queue is executing commands that use the object:

- VkFence
- VkSemaphore
- VkCommandBuffer
- VkCommandPool

In general, objects can be destroyed or freed in any order, even if the object being freed is involved in the use of another object (e.g. use of a resource in a view, use of a view in a descriptor set, use of an object in a command buffer, binding of a memory allocation to a resource), as long as any object that uses the freed object is not further used in any way except to be destroyed or to be reset in such a way that it no longer uses the other object (such as resetting a command buffer). If the object has been reset, then it can be used as if it never used the freed object. An exception to this is when there is a parent/child relationship between objects. In this case, the application must not destroy a parent object before its children, except when the parent is explicitly defined to free its children when it is destroyed (e.g. for pool objects, as defined below).

VkCommandPool objects are parents of VkCommandBuffer objects. VkDescriptorPool objects are parents of VkDescriptorSet objects. VkDevice objects are parents of many object types (all that take a VkDevice as a parameter to their creation).

The following Vulkan objects have specific restrictions for when they can be destroyed:

- VkQueue objects cannot be explicitly destroyed. Instead, they are implicitly destroyed when the VkDevice object they are retrieved from is destroyed.
- Destroying a pool object implicitly frees all objects allocated from that pool. Specifically, destroying VkCommandPool frees all VkCommandBuffer objects that were allocated from it, and destroying VkDescriptorPool frees all VkDescriptorSet objects that were allocated from it.
- VkDevice objects can be destroyed when all VkQueue objects retrieved from them are idle, and all objects created from them have been destroyed. This includes the following objects:
 - VkFence
 - VkSemaphore
 - VkEvent
 - VkQueryPool
 - VkBuffer
 - VkBufferView
 - VkImage
 - VkImageView
 - VkShaderModule
 - VkPipelineCache
 - VkPipeline
 - VkPipelineLayout
 - VkSampler
 - VkDescriptorSetLayout
 - VkDescriptorPool
 - VkFramebuffer
 - VkRenderPass
 - VkCommandPool
 - VkCommandBuffer
 - VkDeviceMemory
- VkPhysicalDevice objects cannot be explicitly destroyed. Instead, they are implicitly destroyed when the VkInstance object they are retrieved from is destroyed.
- VkInstance objects can be destroyed once all VkDevice objects created from any of its VkPhysicalDevice objects have been destroyed.

2.4 Command Syntax and Duration

The Specification describes Vulkan commands as functions or procedures using C99 syntax. Language bindings for other languages such as C++ and JavaScript may allow for stricter parameter passing, or object-oriented interfaces.

Vulkan uses the standard C types for the base type of scalar parameters (e.g. types from stdint.h), with exceptions described below, or elsewhere in the text when appropriate:

VkBool32 represents boolean **True** and **False** values, since C does not have a sufficiently portable built-in boolean type:

```
typedef uint32_t VkBool32;
```

VkDeviceSize represents device memory size and offset values:

```
typedef uint64_t VkDeviceSize;
```

Commands that create Vulkan objects are of the form **vkCreate*** and take Vk*CreateInfo structures with the parameters needed to create the object. These Vulkan objects are destroyed with commands of the form **vkDestroy***.

The last in-parameter to each command that creates or destroys a Vulkan object is *pAllocator*. The *pAllocator* parameter can be set to a non-NULL value such that allocations for the given object are delegated to an application provided callback; refer to the Memory Allocation chapter for further details.

Commands that allocate Vulkan objects owned by pool objects are of the form **vkAllocate***, and take Vk*AllocateInfo structures. These Vulkan objects are freed with commands of the form **vkFree***. These objects do not take allocators; if host memory is needed, they will use the allocator that was specified when their parent pool was created.

Commands are recorded into a command buffer by calling API commands of the form **vkCmd***. Each such command may have different restrictions on where it can be used: in a primary and/or secondary command buffer, inside and/or outside a render pass, and in one or more of the supported queue types. These restrictions are documented together with the definition of each such command.

The duration of a Vulkan command refers to the interval between calling the command and its return to the caller.

2.4.1 Lifetime of Retrieved Results

Information is retrieved from the implementation with commands of the form vkGet* and vkEnumerate*.

Unless otherwise specified for an individual command, the results are *invariant*; that is, they will remain unchanged when retrieved again by calling the same command with the same parameters, so long as those parameters themselves all remain valid.

2.5 Threading Behavior

Vulkan is intended to provide scalable performance when used on multiple host threads. All commands support being called concurrently from multiple threads, but certain parameters, or components of parameters are defined to be *externally synchronized*. This means that the caller must guarantee that no more than one thread is using such a parameter at a given time.

More precisely, Vulkan commands use simple stores to update software structures representing Vulkan objects. A parameter declared as externally synchronized may have its software structures updated at any time during the host execution of the command. If two commands operate on the same object and at least one of the commands declares the object to be externally synchronized, then the caller must guarantee not only that the commands do not execute simultaneously, but also that the two commands are separated by an appropriate memory barrier (if needed).



Note

Memory barriers are particularly relevant on the ARM CPU architecture which is more weakly ordered than many developers are accustomed to from x86/x64 programming. Fortunately, most higher-level synchronization primitives (like the pthread library) perform memory barriers as a part of mutual exclusion, so mutexing Vulkan objects via these primitives will have the desired effect.

Many object types are *immutable*, meaning the objects cannot change once they have been created. These types of objects never need external synchronization, except that they must not be destroyed while they are in use on another thread. In certain special cases, mutable object parameters are internally synchronized such that they do not require external synchronization. One example of this is the use of a VkPipelineCache in vkCreateGraphicsPipelines and vkCreateComputePipelines, where external synchronization around such a heavyweight command would be impractical. The implementation must internally synchronize the cache in this example, and may be able to do so in the form of a much finer-grained mutex around the command. Any command parameters that are not labeled as externally synchronized are either not mutated by the command or are internally synchronized. Additionally, certain objects related to a command's parameters (e.g. command pools and descriptor pools) may be affected by a command, and must also be externally synchronized. These implicit parameters are documented as described below.

Parameters of commands that are externally synchronized are listed below.

Externally Synchronized Parameters

- The instance parameter in vkDestroyInstance
- The device parameter in vkDestroyDevice
- The queue parameter in vkQueueSubmit
- The fence parameter in vkQueueSubmit
- The memory parameter in vkFreeMemory
- The memory parameter in vkMapMemory
- $\bullet \ \ The \ \textit{memory} \ parameter \ in \ \texttt{vkUnmapMemory}$
- The buffer parameter in vkBindBufferMemory
- The image parameter in vkBindImageMemory
- The queue parameter in vkQueueBindSparse
- The fence parameter in vkQueueBindSparse
- The fence parameter in vkDestroyFence
- The semaphore parameter in vkDestroySemaphore
- The event parameter in vkDestroyEvent
- The event parameter in vkSetEvent
- The event parameter in vkResetEvent
- The queryPool parameter in vkDestroyQueryPool
- The buffer parameter in vkDestroyBuffer
- The bufferView parameter in vkDestroyBufferView
- The image parameter in vkDestroyImage
- The imageView parameter in vkDestroyImageView

- The shaderModule parameter in vkDestroyShaderModule
- The pipelineCache parameter in vkDestroyPipelineCache
- The dstCache parameter in vkMergePipelineCaches
- The pipeline parameter in vkDestroyPipeline
- The pipelineLayout parameter in vkDestroyPipelineLayout
- The sampler parameter in vkDestroySampler
- The descriptorSetLayout parameter in vkDestroyDescriptorSetLayout
- The descriptorPool parameter in vkDestroyDescriptorPool
- The descriptorPool parameter in vkResetDescriptorPool
- The descriptorPool member of the pAllocateInfo parameter in vkAllocateDescriptorSets
- $\bullet \ \ \textbf{The} \ \textit{descriptorPool} \ \textbf{parameter} \ \textbf{in} \ \textbf{vkFreeDescriptorSets}$
- \bullet The framebuffer parameter in vkDestroyFramebuffer
- The renderPass parameter in vkDestroyRenderPass
- The commandPool parameter in vkDestroyCommandPool
- The commandPool parameter in vkResetCommandPool
- The commandPool member of the pAllocateInfo parameter in vkAllocateCommandBuffers
- $\bullet \ \ The \ \textit{commandPool} \ parameter \ in \ \texttt{vkFreeCommandBuffers}$
- The commandBuffer parameter in vkBeginCommandBuffer
- The commandBuffer parameter in vkEndCommandBuffer
- The commandBuffer parameter in vkResetCommandBuffer
- The commandBuffer parameter in vkCmdBindPipeline
- $\bullet \ \ The \ \textit{commandBuffer} \ parameter \ in \ \texttt{vkCmdSetViewport}$
- The commandBuffer parameter in vkCmdSetScissor
- The commandBuffer parameter in vkCmdSetLineWidth
- $\bullet \ \ The \ \textit{commandBuffer} \ parameter \ in \ \texttt{vkCmdSetDepthBias}$
- The commandBuffer parameter in vkCmdSetBlendConstants
- The commandBuffer parameter in vkCmdSetDepthBounds
- The commandBuffer parameter in vkCmdSetStencilCompareMask
- The commandBuffer parameter in vkCmdSetStencilWriteMask
- The commandBuffer parameter in vkCmdSetStencilReference
- The commandBuffer parameter in vkCmdBindDescriptorSets

- The commandBuffer parameter in vkCmdBindIndexBuffer
- The commandBuffer parameter in vkCmdBindVertexBuffers
- The commandBuffer parameter in vkCmdDraw
- $\bullet \ \ The \ \textit{commandBuffer} \ parameter \ in \ \texttt{vkCmdDrawIndexed}$
- The commandBuffer parameter in vkCmdDrawIndirect
- The commandBuffer parameter in vkCmdDrawIndexedIndirect
- The commandBuffer parameter in vkCmdDispatch
- The commandBuffer parameter in vkCmdDispatchIndirect
- The commandBuffer parameter in vkCmdCopyBuffer
- The commandBuffer parameter in vkCmdCopyImage
- The commandBuffer parameter in vkCmdBlitImage
- The commandBuffer parameter in vkCmdCopyBufferToImage
- The commandBuffer parameter in vkCmdCopyImageToBuffer
- The commandBuffer parameter in vkCmdUpdateBuffer
- The commandBuffer parameter in vkCmdFillBuffer
- The commandBuffer parameter in vkCmdClearColorImage
- $\bullet \ \ The \ \textit{commandBuffer} \ parameter \ in \ \texttt{vkCmdClearDepthStencilImage}$
- The commandBuffer parameter in vkCmdClearAttachments
- The commandBuffer parameter in vkCmdResolveImage
- The commandBuffer parameter in vkCmdSetEvent
- The commandBuffer parameter in vkCmdResetEvent
- $\bullet \ \ The \ \textit{commandBuffer} \ parameter \ in \ \texttt{vkCmdWaitEvents}$
- The commandBuffer parameter in vkCmdPipelineBarrier
- The commandBuffer parameter in vkCmdBeginQuery
- The commandBuffer parameter in vkCmdEndQuery
- The commandBuffer parameter in vkCmdResetQueryPool
- The commandBuffer parameter in vkCmdWriteTimestamp
- The commandBuffer parameter in vkCmdCopyQueryPoolResults
- The commandBuffer parameter in vkCmdPushConstants
- The commandBuffer parameter in vkCmdBeginRenderPass
- The commandBuffer parameter in vkCmdNextSubpass

- The commandBuffer parameter in vkCmdEndRenderPass
- The commandBuffer parameter in vkCmdExecuteCommands
- The surface parameter in vkDestroySurfaceKHR
- The surface member of the pCreateInfo parameter in vkCreateSwapchainKHR
- The oldSwapchain member of the pCreateInfo parameter in vkCreateSwapchainKHR
- The swapchain parameter in vkDestroySwapchainKHR
- The swapchain parameter in vkAcquireNextImageKHR
- The semaphore parameter in vkAcquireNextImageKHR
- The fence parameter in vkAcquireNextImageKHR
- The queue parameter in vkQueuePresentKHR
- The display parameter in vkCreateDisplayModeKHR
- The mode parameter in vkGetDisplayPlaneCapabilitiesKHR
- The callback parameter in vkDestroyDebugReportCallbackEXT
- The object member of the pTaqInfo parameter in vkDebuqMarkerSetObjectTaqEXT
- The object member of the pNameInfo parameter in vkDebuqMarkerSetObjectNameEXT
- The commandBuffer parameter in vkCmdDrawIndirectCountAMD
- The commandBuffer parameter in vkCmdDrawIndexedIndirectCountAMD

There are also a few instances where a command can take in a user allocated list whose contents are externally synchronized parameters. In these cases, the caller must guarantee that at most one thread is using a given element within the list at a given time. These parameters are listed below.

Externally Synchronized Parameter Lists

- Each element of the pWaitSemaphores member of each element of the pSubmits parameter in vkQueueSubmit
- Each element of the pSignalSemaphores member of each element of the pSubmits parameter in vkQueueSubmit
- Each element of the pWaitSemaphores member of each element of the pBindInfo parameter in vkQueueBindSparse
- Each element of the <code>pSignalSemaphores</code> member of each element of the <code>pBindInfo</code> parameter in <code>vkQueueBindSparse</code>

- The buffer member of each element of the pBufferBinds member of each element of the pBindInfo parameter in vkQueueBindSparse
- The *image* member of each element of the *pImageOpaqueBinds* member of each element of the *pBindInfo* parameter in vkQueueBindSparse
- The *image* member of each element of the *pImageBinds* member of each element of the *pBindInfo* parameter in vkQueueBindSparse
- Each element of the pFences parameter in vkResetFences
- Each element of the pDescriptorSets parameter in vkFreeDescriptorSets
- The dstSet member of each element of the pDescriptorWrites parameter in vkUpdateDescriptorSets
- The dstSet member of each element of the pDescriptorCopies parameter in vkUpdateDescriptorSets
- Each element of the pCommandBuffers parameter in vkFreeCommandBuffers
- Each element of the pWaitSemaphores member of the pPresentInfo parameter in vkQueuePresentKHR
- Each element of the pSwapchains member of the pPresentInfo parameter in vkQueuePresentKHR
- The *surface* member of each element of the *pCreateInfos* parameter in vkCreateSharedSwapchainsKHR
- The oldSwapchain member of each element of the pCreateInfos parameter in vkCreateSharedSwapchainsKHR

In addition, there are some implicit parameters that need to be externally synchronized. For example, all <code>commandBuffer</code> parameters that need to be externally synchronized imply that the <code>commandPool</code> that was passed in when creating that command buffer also needs to be externally synchronized. The implicit parameters and their associated object are listed below.

Implicit Externally Synchronized Parameters

- All VkQueue objects created from device in vkDeviceWaitIdle
- Any VkDescriptorSet objects allocated from descriptorPool in vkResetDescriptorPool
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBindPipeline
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetViewport
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetScissor
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetLineWidth
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetDepthBias
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetBlendConstants

- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetDepthBounds
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetStencilCompareMask
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetStencilWriteMask
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetStencilReference
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBindDescriptorSets
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBindIndexBuffer
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBindVertexBuffers
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDraw
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDrawIndexed
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDrawIndirect
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDrawIndexedIndirect
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDispatch
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDispatchIndirect
- The VkCommandPool that commandBuffer was allocated from, in vkCmdCopyBuffer
- The VkCommandPool that commandBuffer was allocated from, in vkCmdCopyImage
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBlitImage
- The VkCommandPool that commandBuffer was allocated from, in vkCmdCopyBufferToImage
- The VkCommandPool that commandBuffer was allocated from, in vkCmdCopyImageToBuffer
- The VkCommandPool that commandBuffer was allocated from, in vkCmdUpdateBuffer
- The VkCommandPool that commandBuffer was allocated from, in vkCmdFillBuffer
- The VkCommandPool that commandBuffer was allocated from, in vkCmdClearColorImage
- The VkCommandPool that commandBuffer was allocated from, in vkCmdClearDepthStencilImage
- The VkCommandPool that commandBuffer was allocated from, in vkCmdClearAttachments
- $\bullet \ \ The \ {\tt VkCommandPool} \ \ that \ \ {\tt commandBuffer} \ \ was \ allocated \ from, in \ \ {\tt vkCmdResolveImage}$
- The VkCommandPool that commandBuffer was allocated from, in vkCmdSetEvent
- The VkCommandPool that commandBuffer was allocated from, in vkCmdResetEvent
- The VkCommandPool that commandBuffer was allocated from, in vkCmdWaitEvents
- The VkCommandPool that commandBuffer was allocated from, in vkCmdPipelineBarrier
- $\bullet \ \ The \ {\tt VkCommandPool} \ \ that \ \ {\tt commandBuffer} \ \ was \ allocated \ from, \ in \ {\tt vkCmdBeginQuery}$
- The VkCommandPool that commandBuffer was allocated from, in vkCmdEndQuery
- The VkCommandPool that commandBuffer was allocated from, in vkCmdResetQueryPool

- The VkCommandPool that commandBuffer was allocated from, in vkCmdWriteTimestamp
- The VkCommandPool that commandBuffer was allocated from, in vkCmdCopyQueryPoolResults
- The VkCommandPool that commandBuffer was allocated from, in vkCmdPushConstants
- The VkCommandPool that commandBuffer was allocated from, in vkCmdBeginRenderPass
- The VkCommandPool that commandBuffer was allocated from, in vkCmdNextSubpass
- The VkCommandPool that commandBuffer was allocated from, in vkCmdEndRenderPass
- The VkCommandPool that commandBuffer was allocated from, in vkCmdExecuteCommands
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDebugMarkerBeginEXT
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDebuqMarkerEndEXT
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDebugMarkerInsertEXT
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDrawIndirectCountAMD
- The VkCommandPool that commandBuffer was allocated from, in vkCmdDrawIndexedIndirectCountAMD

2.6 Errors

Vulkan is a layered API. The lowest layer is the core Vulkan layer, as defined by this Specification. The application can use additional layers above the core for debugging, validation, and other purposes.

One of the core principles of Vulkan is that building and submitting command buffers should be highly efficient. Thus error checking and validation of state in the core layer is minimal, although more rigorous validation can be enabled through the use of layers.

The core layer assumes applications are using the API correctly. Except as documented elsewhere in the Specification, the behavior of the core layer to an application using the API incorrectly is undefined, and may include program termination. However, implementations must ensure that incorrect usage by an application does not affect the integrity of the operating system, the Vulkan implementation, or other Vulkan client applications in the system, and does not allow one application to access data belonging to another application. Applications can request stronger robustness guarantees by enabling the <code>robustBufferAccess</code> feature as described in Chapter 31.

Validation of correct API usage is left to validation layers. Applications should be developed with validation layers enabled, to help catch and eliminate errors. Once validated, released applications should not enable validation layers by default.

2.6.1 Valid Usage

Valid usage defines a set of conditions which must be met in order to achieve well-defined run-time behavior in an application. These conditions depend only on Vulkan state, and the parameters or objects whose usage is constrained by the condition.

Some valid usage conditions have dependencies on run-time limits or feature availability. It is possible to validate these conditions against Vulkan's minimum supported values for these limits and features, or some subset of other known values.

Valid usage conditions do not cover conditions where well-defined behavior (including returning an error code) exists.

Valid usage conditions should apply to the command or structure where complete information about the condition would be known during execution of an application. This is such that a validation layer or linter can be written directly against these statements at the point they are specified.

Note



This does lead to some non-obvious places for valid usage statements. For instance, the valid values for a structure might depend on a separate value in the calling command. In this case, the structure itself will not reference this valid usage as it is impossible to determine validity from the structure that it is invalid - instead this valid usage would be attached to the calling command.

Another example is draw state - the state setters are independent, and can cause a legitimately invalid state configuration between draw calls; so the valid usage statements are attached to the place where all state needs to be valid - at the draw command.

Certain usage rules apply to all commands in the API unless explicitly denoted differently for a command. These rules are as follows.

2.6.1.1 Valid Usage for Object Handles

Any input parameter to a command that is an object handle must be a valid object handle, unless otherwise specified. An object handle is valid if:

- It has been created or allocated by a previous, successful call to the API. Such calls are noted in the specification.
- It has not been deleted or freed by a previous call to the API. Such calls are noted in the specification.
- Any objects used by that object, either as part of creation or execution, must also be valid.

The reserved handle VK_NULL_HANDLE can be passed in place of valid object handles when *explicitly called out in the specification*. Any command that creates an object successfully must not return VK_NULL_HANDLE. It is valid to pass VK_NULL_HANDLE to any **vkDestroy*** or **vkFree*** command, which will silently ignore these values.

2.6.1.2 Valid Usage for Pointers

Any parameter that is a pointer must be a valid pointer. A pointer is valid if it points at memory containing values of the number and type(s) expected by the command, and all fundamental types accessed through the pointer (e.g. as elements of an array or as members of a structure) satisfy the alignment requirements of the host processor.

2.6.1.3 Valid Usage for Enumerated Types

Any parameter of an enumerated type must be a valid enumerant for that type. A enumerant is valid if:

- The enumerant is defined as part of the enumerated type.
- The enumerant is not one of the special values defined for the enumerated type, which are suffixed with _BEGIN_ RANGE, END RANGE, RANGE SIZE or MAX ENUM.

2.6.1.4 Valid Usage for Flags

A collection of flags is represented by a bitmask using the type VkFlags:

```
typedef uint32_t VkFlags;
```

Bitmasks are passed to many commands and structures to compactly represent options, but VkFlags is not used directly in the API. Instead, a Vk*Flags type which is an alias of VkFlags, and whose name matches the corresponding Vk*FlagBits that are valid for that type, is used. These aliases are described in the Flag Types appendix of the Specification.

Any Vk*Flags member or parameter used in the API must be a valid combination of bit flags. A valid combination is either zero or the bitwise OR of valid bit flags. A bit flag is valid if:

- The bit flag is defined as part of the Vk*FlagBits type, where the bits type is obtained by taking the flag type and replacing the trailing Flags with FlagBits. For example, a flag value of type VkColorComponentFlags must contain only bit flags defined by VkColorComponentFlagBits.
- The flag is allowed in the context in which it is being used. For example, in some cases, certain bit flags or combinations of bit flags are mutually exclusive.

2.6.1.5 Valid Usage for Structure Types

Any parameter that is a structure containing a sType member must have a value of sType which is a valid VkStructureType value matching the type of the structure. As a general rule, the name of this value is obtained by taking the structure name, stripping the leading Vk, prefixing each capital letter with _, converting the entire resulting string to upper case, and prefixing it with $VK_STRUCTURE_TYPE_$. For example, structures of type VkImageCreateInfo must have a sType value of $VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO$.

The values VK_STRUCTURE_TYPE_LOADER_INSTANCE_CREATE_INFO and VK_STRUCTURE_TYPE_LOADER_DEVICE_CREATE_INFO are reserved for internal use by the loader, and do not have corresponding Vulkan structures in this specification.

The list of supported structure types is defined in an appendix.

2.6.1.6 Valid Usage for Structure Pointer Chains

Any parameter that is a structure containing a void*pNext member must have a value of pNext that is either NULL, or points to a valid structure defined by an extension, containing sType and pNext members as described in the Vulkan Documentation and Extensions document in the section "Extension Interactions". If that extension is supported by the implementation, then it must be enabled. Any component of the implementation (the loader, any enabled layers, and drivers) must skip over, without processing (other than reading the sType and pNext members) any chained structures with sType values not defined by extensions supported by that component.

Extension structures are not described in the base Vulkan specification, but either in layered specifications incorporating those extensions, or in separate vendor-provided documents.

2.6.1.7 Valid Usage for Nested Structures

The above rules also apply recursively to members of structures provided as input to a command, either as a direct argument to the command, or themselves a member of another structure.

Specifics on valid usage of each command are covered in their individual sections.

2.6.2 Return Codes

While the core Vulkan API is not designed to capture incorrect usage, some circumstances still require return codes. Commands in Vulkan return their status via return codes that are in one of two categories:

- Successful completion codes are returned when a command needs to communicate success or status information. All successful completion codes are non-negative values.
- Run time error codes are returned when a command needs to communicate a failure that could only be detected at run time. All run time error codes are negative values.

All return codes in Vulkan are reported via VkResult return values. The possible codes are:

```
typedef enum VkResult {
   VK\_SUCCESS = 0,
   VK_NOT_READY = 1,
   VK\_TIMEOUT = 2,
   VK\_EVENT\_SET = 3,
   VK\_EVENT\_RESET = 4,
   VK_INCOMPLETE = 5,
    VK\_ERROR\_OUT\_OF\_HOST\_MEMORY = -1,
    VK\_ERROR\_OUT\_OF\_DEVICE\_MEMORY = -2,
   VK ERROR INITIALIZATION FAILED = -3,
   VK_ERROR_DEVICE_LOST = -4,
   VK\_ERROR\_MEMORY\_MAP\_FAILED = -5,
   VK\_ERROR\_LAYER\_NOT\_PRESENT = -6,
   VK ERROR EXTENSION NOT PRESENT = -7,
   VK\_ERROR\_FEATURE\_NOT\_PRESENT = -8,
   VK\_ERROR\_INCOMPATIBLE\_DRIVER = -9,
   VK_ERROR_TOO_MANY_OBJECTS = -10,
   VK\_ERROR\_FORMAT\_NOT\_SUPPORTED = -11,
   VK\_ERROR\_FRAGMENTED\_POOL = -12,
    VK\_ERROR\_SURFACE\_LOST\_KHR = -1000000000,
   VK_ERROR_NATIVE_WINDOW_IN_USE_KHR = -1000000001,
    VK_SUBOPTIMAL_KHR = 1000001003,
    VK\_ERROR\_OUT\_OF\_DATE\_KHR = -1000001004,
    VK_ERROR_INCOMPATIBLE_DISPLAY_KHR = -1000003001,
    VK\_ERROR\_VALIDATION\_FAILED\_EXT = -1000011001,
    VK\_ERROR\_INVALID\_SHADER\_NV = -1000012000,
} VkResult;
```

SUCCESS CODES

- VK_SUCCESS Command successfully completed
- VK_NOT_READY A fence or query has not yet completed
- VK_TIMEOUT A wait operation has not completed in the specified time
- VK EVENT SET An event is signaled
- VK_EVENT_RESET An event is unsignaled
- VK INCOMPLETE A return array was too small for the result
- VK_SUBOPTIMAL_KHR A swapchain no longer matches the surface properties exactly, but can still be used to present to the surface successfully.

ERROR CODES

- VK_ERROR_OUT_OF_HOST_MEMORY A host memory allocation has failed.
- VK ERROR OUT OF DEVICE MEMORY A device memory allocation has failed.
- VK_ERROR_INITIALIZATION_FAILED Initialization of an object could not be completed for implementation-specific reasons.
- VK_ERROR_DEVICE_LOST The logical or physical device has been lost. See Lost Device
- VK_ERROR_MEMORY_MAP_FAILED Mapping of a memory object has failed.
- VK_ERROR_LAYER_NOT_PRESENT A requested layer is not present or could not be loaded.
- VK_ERROR_EXTENSION_NOT_PRESENT A requested extension is not supported.
- VK_ERROR_FEATURE_NOT_PRESENT A requested feature is not supported.
- VK_ERROR_INCOMPATIBLE_DRIVER The requested version of Vulkan is not supported by the driver or is otherwise incompatible for implementation-specific reasons.
- VK_ERROR_TOO_MANY_OBJECTS Too many objects of the type have already been created.
- VK_ERROR_FORMAT_NOT_SUPPORTED A requested format is not supported on this device.
- VK_ERROR_FRAGMENTED_POOL A requested pool allocation has failed due to fragmentation of the pool's memory.
- VK_ERROR_SURFACE_LOST_KHR A surface is no longer available.
- VK_ERROR_NATIVE_WINDOW_IN_USE_KHR The requested window is already connected to a VkSurfaceKHR, or to some other non-Vulkan API.
- VK_ERROR_OUT_OF_DATE_KHR A surface has changed in such a way that it is no longer compatible with the swapchain, and further presentation requests using the swapchain will fail. Applications must query the new surface properties and recreate their swapchain if they wish to continue presenting to the surface.
- VK_ERROR_INCOMPATIBLE_DISPLAY_KHR The display used by a swapchain does not use the same presentable image layout, or is incompatible in a way that prevents sharing an image.
- VK_ERROR_INVALID_SHADER_NV One or more shaders failed to compile or link. More details are reported back to the application via VK_EXT_debug_report if enabled.

If a command returns a run time error, it will leave any result pointers unmodified, unless other behavior is explicitly defined in the specification.

Out of memory errors do not damage any currently existing Vulkan objects. Objects that have already been successfully created can still be used by the application.

Performance-critical commands generally do not have return codes. If a run time error occurs in such commands, the implementation will defer reporting the error until a specified point. For commands that record into command buffers (**vkCmd***) run time errors are reported by **vkEndCommandBuffer**.

2.7 Numeric Representation and Computation

Implementations normally perform computations in floating-point, and must meet the range and precision requirements defined under "Floating-Point Computation" below.

These requirements only apply to computations performed in Vulkan operations outside of shader execution, such as texture image specification and sampling, and per-fragment operations. Range and precision requirements during shader execution differ and are specified by the Precision and Operation of SPIR-V Instructions section.

In some cases, the representation and/or precision of operations is implicitly limited by the specified format of vertex or texel data consumed by Vulkan. Specific floating-point formats are described later in this section.

2.7.1 Floating-Point Computation

Most floating-point computation is performed in SPIR-V shader modules. The properties of computation within shaders are constrained as defined by the Precision and Operation of SPIR-V Instructions section.

Some floating-point computation is performed outside of shaders, such as viewport and depth range calculations. For these computations, we do not specify how floating-point numbers are to be represented, or the details of how operations on them are performed, but only place minimal requirements on representation and precision as described in the remainder of this section.

We require simply that numbers' floating-point parts contain enough bits and that their exponent fields are large enough so that individual results of floating-point operations are accurate to about 1 part in 10^5 . The maximum representable magnitude for all floating-point values must be at least 2^{32} . $x \cdot 0 = 0 \cdot x = 0$ for any non-infinite and non-NaN x. $1 \cdot x = x \cdot 1 = x$. x + 0 = 0 + x = x. $0^0 = 1$.

Occasionally, further requirements will be specified. Most single-precision floating-point formats meet these requirements.

The special values Inf and -Inf encode values with magnitudes too large to be represented; the special value NaN encodes "Not A Number" values resulting from undefined arithmetic operations such as 0/0. Implementations may support Infs and NaNs in their floating-point computations.

Any representable floating-point value is legal as input to a Vulkan command that requires floating-point data. The result of providing a value that is not a floating-point number to such a command is unspecified, but must not lead to Vulkan interruption or termination. In IEEE 754 arithmetic, for example, providing a negative zero or a denormalized number to an Vulkan command must yield deterministic results, while providing a *NaN* or *Inf* yields unspecified results.

2.7.2 16-Bit Floating-Point Numbers

16-bit floating point numbers are defined in the "16-bit floating point numbers" section of the Khronos Data Format Specification.

Any representable 16-bit floating-point value is legal as input to a Vulkan command that accepts 16-bit floating-point data. The result of providing a value that is not a floating-point number (such as Inf or NaN) to such a command is unspecified, but must not lead to Vulkan interruption or termination. Providing a denormalized number or negative zero to Vulkan must yield deterministic results.

2.7.3 Unsigned 11-Bit Floating-Point Numbers

Unsigned 11-bit floating point numbers are defined in the "Unsigned 11-bit floating point numbers" section of the Khronos Data Format Specification.

When a floating-point value is converted to an unsigned 11-bit floating-point representation, finite values are rounded to the closest representable finite value.

While less accurate, implementations are allowed to always round in the direction of zero. This means negative values are converted to zero. Likewise, finite positive values greater than 65024 (the maximum finite representable unsigned 11-bit floating-point value) are converted to 65024. Additionally: negative infinity is converted to zero; positive infinity is converted to positive infinity; and both positive and negative *NaN* are converted to positive *NaN*.

Any representable unsigned 11-bit floating-point value is legal as input to a Vulkan command that accepts 11-bit floating-point data. The result of providing a value that is not a floating-point number (such as *Inf* or *NaN*) to such a command is unspecified, but must not lead to Vulkan interruption or termination. Providing a denormalized number to Vulkan must yield deterministic results.

2.7.4 Unsigned 10-Bit Floating-Point Numbers

Unsigned 10-bit floating point numbers are defined in the "Unsigned 10-bit floating point numbers" section of the Khronos Data Format Specification.

When a floating-point value is converted to an unsigned 10-bit floating-point representation, finite values are rounded to the closest representable finite value.

While less accurate, implementations are allowed to always round in the direction of zero. This means negative values are converted to zero. Likewise, finite positive values greater than 64512 (the maximum finite representable unsigned 10-bit floating-point value) are converted to 64512. Additionally: negative infinity is converted to zero; positive infinity is converted to positive infinity; and both positive and negative *NaN* are converted to positive *NaN*.

Any representable unsigned 10-bit floating-point value is legal as input to a Vulkan command that accepts 10-bit floating-point data. The result of providing a value that is not a floating-point number (such as *Inf* or *NaN*) to such a command is unspecified, but must not lead to Vulkan interruption or termination. Providing a denormalized number to Vulkan must yield deterministic results.

2.7.5 General Requirements

Some calculations require division. In such cases (including implied divisions performed by vector normalization), division by zero produces an unspecified result but must not lead to Vulkan interruption or termination.

2.8 Fixed-Point Data Conversions

When generic vertex attributes and pixel color or depth *components* are represented as integers, they are often (but not always) considered to be *normalized*. Normalized integer values are treated specially when being converted to and from floating-point values, and are usually referred to as *normalized fixed-point*.

In the remainder of this section, b denotes the bit width of the fixed-point integer representation. When the integer is one of the types defined by the API, b is the bit width of that type. When the integer comes from an image containing color or depth component texels, b is the number of bits allocated to that component in its specified image format.

The signed and unsigned fixed-point representations are assumed to be b-bit binary two's-complement integers and binary unsigned integers, respectively.

2.8.1 Conversion from Normalized Fixed-Point to Floating-Point

Unsigned normalized fixed-point integers represent numbers in the range [0,1]. The conversion from an unsigned normalized fixed-point value c to the corresponding floating-point value f is defined as

$$f = \frac{c}{2^b - 1}$$

Signed normalized fixed-point integers represent numbers in the range [-1,1]. The conversion from a signed normalized fixed-point value c to the corresponding floating-point value f is performed using

$$f = \max(\frac{c}{2^{b-1} - 1}, -1.0)$$

Only the range $[-2^{b-1}+1,2^{b-1}-1]$ is used to represent signed fixed-point values in the range [-1,1]. For example, if b=8, then the integer value -127 corresponds to -1.0 and the value 127 corresponds to 1.0. Note that while zero is exactly expressible in this representation, one value (-128 in the example) is outside the representable range, and must be clamped before use. This equation is used everywhere that signed normalized fixed-point values are converted to floating-point.

2.8.2 Conversion from Floating-Point to Normalized Fixed-Point

The conversion from a floating-point value f to the corresponding unsigned normalized fixed-point value c is defined by first clamping f to the range [0,1], then computing

$$c = \text{convertFloatToUint}(f \times (2^b - 1), b)$$

where convertFloatToUint(r,b) returns one of the two unsigned binary integer values with exactly b bits which are closest to the floating-point value r. Implementations should round to nearest. If r is equal to an integer, then that integer value is returned. In particular, if f is equal to 0.0 or 1.0, then c must be assigned 0 or $2^b - 1$, respectively.

The conversion from a floating-point value f to the corresponding signed normalized fixed-point value c is performed by clamping f to the range [-1,1], then computing

$$c = \text{convertFloatToInt}(f \times (2^{b-1} - 1), b)$$

where convertFloatToInt(r,b) returns one of the two signed two's-complement binary integer values with exactly b bits which are closest to the floating-point value r. Implementations should round to nearest. If r is equal to an integer, then that integer value must be returned. In particular, if f is equal to -1.0, 0.0, or 1.0, then c must be assigned $-(2^{b-1}-1)$, 0, or $2^{b-1}-1$, respectively.

This equation is used everywhere that floating-point values are converted to signed normalized fixed-point.

2.9 API Version Numbers and Semantics

The Vulkan version number is used in several places in the API. In each such use, the API *major version number*, *minor version number*, and *patch version number* are packed into a 32-bit integer as follows:

- The major version number is a 10-bit integer packed into bits 31-22.
- The minor version number is a 10-bit integer packed into bits 21-12.
- The patch version number is a 12-bit integer packed into bits 11-0.

Differences in any of the Vulkan version numbers indicates a change to the API in some way, with each part of the version number indicating a different scope of changes.

A difference in patch version numbers indicates that some usually small part of the specification or header has been modified, typically to fix a bug, and may have an impact on the behavior of existing functionality. Differences in this version number should not affect either *full compatibility* or *backwards compatibility* between two versions, or add additional interfaces to the API.

A difference in minor version numbers indicates that some amount of new functionality has been added. This will usually include new interfaces in the header, and may also include behavior changes and bug fixes. Functionality may be deprecated in a minor revision, but will not be removed. When a new minor version is introduced, the patch version is reset to 0, and each minor revision maintains its own set of patch versions. Differences in this version should not affect backwards compatibility, but will affect full compatibility.

A difference in major version numbers indicates a large set of changes to the API, potentially including new functionality and header interfaces, behavioral changes, removal of deprecated features, modification or outright replacement of any feature, and is thus very likely to break any and all compatibility. Differences in this version will typically require significant modification to an application in order for it to function.

C language macros for manipulating version numbers are defined in the Version Number Macros appendix.

2.10 Common Object Types

Some types of Vulkan objects are used in many different structures and command parameters, and are described here. These types include *offsets*, *extents*, and *rectangles*.

2.10.1 Offsets

Offsets are used to describe a pixel location within an image or framebuffer, as an (x,y) location for two-dimensional images, or an (x,y,z) location for three-dimensional images.

A two-dimensional offsets is defined by the structure:

```
typedef struct VkOffset2D {
  int32_t x;
  int32_t y;
} VkOffset2D;
```

A three-dimensional offset is defined by the structure:

```
typedef struct VkOffset3D {
   int32_t x;
   int32_t y;
   int32_t z;
} VkOffset3D;
```

2.10.2 Extents

Extents are used to describe the size of a rectangular region of pixels within an image or framebuffer, as (width,height) for two-dimensional images, or as (width,height,depth) for three-dimensional images.

A two-dimensional extent is defined by the structure:

```
typedef struct VkExtent2D {
   uint32_t width;
   uint32_t height;
} VkExtent2D;
```

A three-dimensional extent is defined by the structure:

```
typedef struct VkExtent3D {
   uint32_t width;
   uint32_t height;
   uint32_t depth;
} VkExtent3D;
```

2.10.3 Rectangles

Rectangles are used to describe a specified rectangular region of pixels within an image or framebuffer. Rectangles include both an offset and an extent of the same dimensionality, as described above. Two-dimensional rectangles are defined by the structure

```
typedef struct VkRect2D {
    VkOffset2D offset;
    VkExtent2D extent;
} VkRect2D;
```

Chapter 3

Initialization

Before using Vulkan, an application must initialize it by loading the Vulkan commands, and creating a VkInstance object.

3.1 Command Function Pointers

Vulkan commands are not necessarily exposed statically on a platform. Function pointers for all Vulkan commands can be obtained with the command:

- *instance* is the instance that the function pointer will be compatible with, or NULL for commands not dependent on any instance.
- pName is the name of the command to obtain.

vkGetInstanceProcAddr itself is obtained in a platform- and loader- specific manner. Typically, the loader library will export this command as a function symbol, so applications can link against the loader library, or load it dynamically and look up the symbol using platform-specific APIs. Loaders are encouraged to export function symbols for all other core Vulkan commands as well; if this is done, then applications that use only the core Vulkan commands have no need to use **vkGetInstanceProcAddr**.

The table below defines the various use cases for **vkGetInstanceProcAddr** and expected return value ("fp" is function pointer) for each case.

The returned function pointer is of type PFN_vkVoidFunction, and must be cast to the type of the command being queried.

1able 3.1:	vkGeunstanceProcAddr	benavior

instance	pName	return value
*	NULL	undefined
invalid instance	*	undefined

Table 3.1: (continued)

instance	pName	return value	
NULL	vkEnumerateInstanc	fp	
	eExtensionPropert		
	ies		
NULL	vkEnumerateInstanc	fp	
	eLayerProperties		
NULL	vkCreateInstance	fp	
NULL	* (any pName not covered	NULL	
	above)		
instance	core Vulkan command	fp ¹	
instance	enabled instance extension	fp ¹	
	commands for instance		
instance	available device extension ²	fp ¹	
	commands for instance		
instance	* (any pName not covered	NULL	
	above)		

The returned function pointer must only be called with a dispatchable object (the first parameter) that is <code>instance</code> or a child of <code>instance</code>. e.g. VkInstance, VkPhysicalDevice, VkDevice, VkQueue, or VkCommandBuffer.

2

An "available extension" is an extension function supported by any of the loader, driver or layer.

Valid Usage

- If instance is not NULL, instance must be a valid VkInstance handle
- pName must be a null-terminated string

In order to support systems with multiple Vulkan implementations comprising heterogeneous collections of hardware and software, the function pointers returned by **vkGetInstanceProcAddr** may point to dispatch code, which calls a different real implementation for different VkDevice objects (and objects created from them). The overhead of this internal dispatch can be avoided by obtaining device-specific function pointers for any commands that use a device or device-child object as their dispatchable object. Such function pointers can be obtained with the command:

The table below defines the various use cases for **vkGetDeviceProcAddr** and expected return value for each case.

The returned function pointer is of type PFN_vkVoidFunction, and must be cast to the type of the command being queried.

device	pName	return value	
NULL	*	undefined	
invalid device	*	undefined	
device	NULL	undefined	
device	core Vulkan command	fp ¹	
device	enabled extension	fp ¹	
	commands		
device	* (any pName not covered	NULL	
	above)		

Table 3.2: vkGetDeviceProcAddr behavior

1 The returned function pointer must only be called with a dispatchable object (the first parameter) that is <code>device</code> or a child of <code>device</code>. e.g. VkDevice, VkQueue, or VkCommandBuffer.

Valid Usage

- device must be a valid VkDevice handle
- pName must be a null-terminated string

The definition of PFN_vkVoidFunction is:

```
typedef void (VKAPI_PTR *PFN_vkVoidFunction) (void);
```

3.2 Instances

There is no global state in Vulkan and all per-application state is stored in a VkInstance object. Creating a VkInstance object initializes the Vulkan library and allows the application to pass information about itself to the implementation.

Instances are represented by VkInstance handles:

```
VK_DEFINE_HANDLE (VkInstance)
```

To create an instance object, call:

- pCreateInfo points to an instance of VkInstanceCreateInfo controlling creation of the instance.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pInstance points a VkInstance handle in which the resulting instance is returned.

vkCreateInstance creates the instance, then enables and initializes global layers and extensions requested by the application. If an extension is provided by a layer, both the layer and extension must be specified at **vkCreateInstance** time. If a specified layer cannot be found, no VkInstance will be created and the function will return VK_ERROR_LAYER_NOT_PRESENT. Likewise, if a specified extension cannot be found the call will return VK_ERROR_EXTENSION_NOT_PRESENT.

Valid Usage

- pCreateInfo must be a pointer to a valid VkInstanceCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pInstance must be a pointer to a VkInstance handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_INITIALIZATION_FAILED
- VK_ERROR_LAYER_NOT_PRESENT
- VK_ERROR_EXTENSION_NOT_PRESENT
- VK_ERROR_INCOMPATIBLE_DRIVER

The VkInstanceCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- papplicationInfo is NULL or a pointer to an instance of VkapplicationInfo. If not NULL, this information helps implementations recognize behavior inherent to classes of applications. VkapplicationInfo is defined in detail below.
- enabledLayerCount is the number of global layers to enable.
- ppEnabledLayerNames is a pointer to an array of enabledLayerCount null-terminated UTF-8 strings containing the names of layers to enable for the created instance. See the Layers section for further details.
- enabledExtensionCount is the number of global extensions to enable.
- ppEnabledExtensionNames is a pointer to an array of enabledExtensionCount null-terminated UTF-8 strings containing the names of extensions to enable.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- If papplicationInfo is not NULL, papplicationInfo must be a pointer to a valid VkapplicationInfo structure
- If enabledLayerCount is not 0, ppEnabledLayerNames must be a pointer to an array of enabledLayerCount null-terminated strings
- If <code>enabledExtensionCount</code> is not 0, <code>ppEnabledExtensionNames</code> must be a pointer to an array of <code>enabledExtensionCount</code> null-terminated strings

The VkApplicationInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- papplicationName is a pointer to a null-terminated UTF-8 string containing the name of the application.
- applicationVersion is an unsigned integer variable containing the developer-supplied version number of the application.
- pEngineName is a pointer to a null-terminated UTF-8 string containing the name of the engine (if any) used to create the application.
- engineVersion is an unsigned integer variable containing the developer-supplied version number of the engine used to create the application.
- apiVersion is the version of the Vulkan API against which the application expects to run, encoded as described in the API Version Numbers and Semantics section. If apiVersion is 0 the implementation must ignore it, otherwise if the implementation does not support the requested apiVersion it must return VK_ERROR_INCOMPATIBLE_DRIVER. The patch version number specified in apiVersion is ignored when creating an instance object. Only the major and minor versions of the instance must match those requested in apiVersion.

Valid Usage

- sType must be VK STRUCTURE TYPE APPLICATION INFO
- pNext must be NULL
- If pApplicationName is not NULL, pApplicationName must be a null-terminated string
- If pEngineName is not NULL, pEngineName must be a null-terminated string
- apiVersion must be zero, or otherwise it must be a version that the implementation supports, or supports an effective substitute for

To destroy an instance, call:

- instance is the handle of the instance to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- If instance is not NULL, instance must be a valid VkInstance handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- All child objects created using instance must have been destroyed prior to destroying instance
- If VkAllocationCallbacks were provided when <code>instance</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when instance was created, pAllocator must be NULL

Host Synchronization

• Host access to instance must be externally synchronized

Chapter 4

Devices and Queues

Once Vulkan is initialized, devices and queues are the primary objects used to interact with a Vulkan implementation.

Vulkan separates the concept of *physical* and *logical* devices. A physical device usually represents a single device in a system (perhaps made up of several individual hardware devices working together), of which there are a finite number. A logical device represents an application's view of the device.

Physical devices are represented by VkPhysicalDevice handles:

VK_DEFINE_HANDLE (VkPhysicalDevice)

4.1 Physical Devices

To retrieve a list of physical device objects representing the physical devices installed in the system, call:

- instance is a handle to a Vulkan instance previously created with vkCreateInstance.
- pPhysicalDeviceCount is a pointer to an integer related to the number of physical devices available or queried, as
 described below.
- $\bullet \ \textit{pPhysicalDevices} \ \textbf{is} \ \textbf{either} \ \texttt{NULL} \ \textbf{or} \ \textbf{a} \ \textbf{pointer} \ \textbf{to} \ \textbf{an} \ \textbf{array} \ \textbf{of} \ \texttt{VkPhysicalDevice} \ \textbf{handles}.$

If <code>pPhysicalDevices</code> is <code>NULL</code>, then the number of physical devices available is returned in <code>pPhysicalDeviceCount</code>. Otherwise, <code>pPhysicalDeviceCount</code> must point to a variable set by the user to the number of elements in the <code>pPhysicalDevices</code> array, and on return the variable is overwritten with the number of structures actually written to <code>pPhysicalDevices</code>. If <code>pPhysicalDeviceCount</code> is less than the number of physical devices available, at most <code>pPhysicalDeviceCount</code> structures will be written. If <code>pPhysicalDeviceCount</code> is smaller than the number of physical devices available, <code>VK_INCOMPLETE</code> will be returned instead of <code>VK_SUCCESS</code>, to indicate that not all the available physical devices were returned.

Valid Usage

- instance must be a valid VkInstance handle
- pPhysicalDeviceCount must be a pointer to a uint32_t value
- If the value referenced by pPhysicalDeviceCount is not 0, and pPhysicalDevices is not NULL, pPhysicalDevices must be a pointer to an array of pPhysicalDeviceCount VkPhysicalDevice handles

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_INITIALIZATION_FAILED

To query general properties of physical devices once enumerated, call:

- physicalDevice is the handle to the physical device whose properties will be queried.
- pProperties points to an instance of the VkPhysicalDeviceProperties structure, that will be filled with returned information.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pProperties must be a pointer to a VkPhysicalDeviceProperties structure

The VkPhysicalDeviceProperties structure is defined as:

```
typedef struct VkPhysicalDeviceProperties {
   uint32_t
                                        apiVersion;
   uint32_t
                                        driverVersion;
   uint32_t
                                        vendorID;
   uint32_t
                                        deviceID;
   VkPhysicalDeviceType
                                        deviceType;
                                        deviceName[VK_MAX_PHYSICAL_DEVICE_NAME_SIZE];
   char
   uint8_t
                                        pipelineCacheUUID[VK_UUID_SIZE];
   VkPhysicalDeviceLimits
                                        limits;
   VkPhysicalDeviceSparseProperties
                                        sparseProperties;
} VkPhysicalDeviceProperties;
```

- apiVersion is the version of Vulkan supported by the device, encoded as described in the API Version Numbers and Semantics section.
- driverVersion is the vendor-specified version of the driver.
- vendorID is a unique identifier for the vendor (see below) of the physical device.
- deviceID is a unique identifier for the physical device among devices available from the vendor.
- deviceType is a VkPhysicalDeviceType specifying the type of device.
- deviceName is a null-terminated UTF-8 string containing the name of the device.
- pipelineCacheUUID is an array of size VK_UUID_SIZE, containing 8-bit values that represent a universally unique identifier for the device.
- limits is the VkPhysicalDeviceLimits structure which specifies device-specific limits of the physical device. See Limits for details.
- *sparseProperties* is the VkPhysicalDeviceSparseProperties structure which specifies various sparse related properties of the physical device. See Sparse Properties for details.

The <code>vendorID</code> and <code>deviceID</code> fields are provided to allow applications to adapt to device characteristics that are not adequately exposed by other Vulkan queries. These may include performance profiles, hardware errata, or other characteristics. In PCI-based implementations, the low sixteen bits of <code>vendorID</code> and <code>deviceID</code> must contain (respectively) the PCI vendor and device IDs associated with the hardware device, and the remaining bits must be set to zero. In non-PCI implementations, the choice of what values to return may be dictated by operating system or platform policies. It is otherwise at the discretion of the implementer, subject to the following constraints and guidelines:

- For purposes of physical device identification, the *vendor* of a physical device is the entity responsible for the most salient characteristics of the hardware represented by the physical device handle. In the case of a discrete GPU, this should be the GPU chipset vendor. In the case of a GPU or other accelerator integrated into a system-on-chip (SoC), this should be the supplier of the silicon IP used to create the GPU or other accelerator.
- If the vendor of the physical device has a valid PCI vendor ID issued by PCI-SIG, that ID should be used to construct <code>vendorID</code> as described above for PCI-based implementations. Implementations that do not return a PCI vendor ID in <code>vendorID</code> must return a valid Khronos vendor ID, obtained as described in the Vulkan Documentation and Extensions document in the section "Registering a Vendor ID with Khronos". Khronos vendor IDs are allocated starting at 0x10000, to distinguish them from the PCI vendor ID namespace.

• The vendor of the physical device is responsible for selecting <code>deviceID</code>. The value selected should uniquely identify both the device version and any major configuration options (for example, core count in the case of multicore devices). The same device ID should be used for all physical implementations of that device version and configuration. For example, all uses of a specific silicon IP GPU version and configuration should use the same device ID, even if those uses occur in different SoCs.

The physical devices types are:

```
typedef enum VkPhysicalDeviceType {
    VK_PHYSICAL_DEVICE_TYPE_OTHER = 0,
    VK_PHYSICAL_DEVICE_TYPE_INTEGRATED_GPU = 1,
    VK_PHYSICAL_DEVICE_TYPE_DISCRETE_GPU = 2,
    VK_PHYSICAL_DEVICE_TYPE_VIRTUAL_GPU = 3,
    VK_PHYSICAL_DEVICE_TYPE_CPU = 4,
} VkPhysicalDeviceType;
```

- VK_PHYSICAL_DEVICE_TYPE_OTHER The device does not match any other available types.
- VK_PHYSICAL_DEVICE_TYPE_INTEGRATED_GPU The device is typically one embedded in or tightly coupled with the host.
- VK_PHYSICAL_DEVICE_TYPE_DISCRETE_GPU The device is typically a separate processor connected to the host via an interlink.
- VK_PHYSICAL_DEVICE_TYPE_VIRTUAL_GPU The device is typically a virtual node in a virtualization environment.
- VK_PHYSICAL_DEVICE_TYPE_CPU The device is typically running on the same processors as the host.

The physical device type is advertised for informational purposes only, and does not directly affect the operation of the system. However, the device type may correlate with other advertised properties or capabilities of the system, such as how many memory heaps there are.

To query properties of queues available on a physical device, call:

- physical Device is the handle to the physical device whose properties will be queried.
- pQueueFamilyPropertyCount is a pointer to an integer related to the number of queue families available or queried, as described below.
- pQueueFamilyProperties is either NULL or a pointer to an array of VkQueueFamilyProperties structures.

If pQueueFamilyProperties is NULL, then the number of queue families available is returned in pQueueFamilyPropertyCount. Otherwise, pQueueFamilyPropertyCount must point to a variable set by the user to the number of elements in the pQueueFamilyProperties array, and on return the variable is overwritten with the number of structures actually written to pQueueFamilyProperties. If pQueueFamilyPropertyCount is less than the number of queue families available, at most pQueueFamilyPropertyCount structures will be written.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pQueueFamilyPropertyCount must be a pointer to a uint32_t value
- If the value referenced by <code>pQueueFamilyPropertyCount</code> is not 0, and <code>pQueueFamilyProperties</code> is not <code>NULL</code>, <code>pQueueFamilyProperties</code> must be a pointer to an array of <code>pQueueFamilyPropertyCount</code> <code>VkQueueFamilyProperties</code> structures

The VkQueueFamilyProperties structure is defined as:

- queueFlags contains flags indicating the capabilities of the queues in this queue family.
- queueCount is the unsigned integer count of queues in this queue family.
- timestampValidBits is the unsigned integer count of meaningful bits in the timestamps written via **vkCmdWriteTimestamp**. The valid range for the count is 36..64 bits, or a value of 0, indicating no support for timestamps. Bits outside the valid range are guaranteed to be zeros.
- minImageTransferGranularity is the minimum granularity supported for image transfer operations on the queues in this queue family.

The bits specified in queueFlags are:

```
typedef enum VkQueueFlagBits {
    VK_QUEUE_GRAPHICS_BIT = 0x00000001,
    VK_QUEUE_COMPUTE_BIT = 0x00000002,
    VK_QUEUE_TRANSFER_BIT = 0x00000004,
    VK_QUEUE_SPARSE_BINDING_BIT = 0x00000008,
} VkQueueFlagBits;
```

- if VK_QUEUE_GRAPHICS_BIT is set, then the queues in this queue family support graphics operations.
- if VK_QUEUE_COMPUTE_BIT is set, then the queues in this queue family support compute operations.
- if VK_QUEUE_TRANSFER_BIT is set, then the queues in this queue family support transfer operations.
- if VK_QUEUE_SPARSE_BINDING_BIT is set, then the queues in this queue family support sparse memory management operations (see Sparse Resources). If any of the sparse resource features are enabled, then at least one queue family must support this bit.

If an implementation exposes any queue family that supports graphics operations, at least one queue family of at least one physical device exposed by the implementation must support both graphics and compute operations.



Note

All commands that are allowed on a queue that supports transfer operations are also allowed on a queue that supports either graphics or compute operations thus if the capabilities of a queue family include VK_QUEUE_GRAPHICS_BIT or VK_QUEUE_COMPUTE_BIT then reporting the VK_QUEUE_TRANSFER_BIT capability separately for that queue family is optional.

For further details see Queues.

The value returned in minImageTransferGranularity has a unit of compressed texel blocks for images having a block-compressed format, and a unit of texels otherwise.

Possible values of minImageTransferGranularity are:

- (0,0,0) which indicates that only whole mip levels must be transferred using the image transfer operations on the corresponding queues. In this case, the following restrictions apply to all offset and extent parameters of image transfer operations:
 - The x, y, and z members of a VkOffset 3D parameter must always be zero.
- The width, height, and depth members of a VkExtent3D parameter must always match the width, height, and depth of the image subresource corresponding to the parameter, respectively.
- (Ax, Ay, Az) where Ax, Ay, and Az are all integer powers of two. In this case the following restrictions apply to all image transfer operations:
 - -x, y, and z of a VkOffset 3D parameter must be integer multiples of Ax, Ay, and Az, respectively.
 - width of a VkExtent3D parameter must be an integer multiple of Ax, or else (x + width) must equal the width of the image subresource corresponding to the parameter.
 - height of a VkExtent3D parameter must be an integer multiple of Ay, or else (y + height) must equal the height of the image subresource corresponding to the parameter.
 - depth of a VkExtent3D parameter must be an integer multiple of Az, or else (z + depth) must equal the depth of the image subresource corresponding to the parameter.
 - If the format of the image corresponding to the parameters is one of the block-compressed formats then for the
 purposes of the above calculations the granularity must be scaled up by the compressed texel block dimensions.

Queues supporting graphics and/or compute operations must report (1,1,1) in minImageTransferGranularity, meaning that there are no additional restrictions on the granularity of image transfer operations for these queues. Other queues supporting image transfer operations are only required to support whole mip level transfers, thus minImageTransferGranularity for queues belonging to such queue families may be (0,0,0).

The Device Memory section describes memory properties queried from the physical device.

For physical device feature queries see the Features chapter.

4.2 Devices

Device objects represent logical connections to physical devices. Each device exposes a number of *queue families* each having one or more *queues*. All queues in a queue family support the same operations.

As described in Physical Devices, a Vulkan application will first query for all physical devices in a system. Each physical device can then be queried for its capabilities, including its queue and queue family properties. Once an acceptable physical device is identified, an application will create a corresponding logical device. An application must create a

separate logical device for each physical device it will use. The created logical device is then the primary interface to the physical device.

How to enumerate the physical devices in a system and query those physical devices for their queue family properties is described in the Physical Device Enumeration section above.

4.2.1 Device Creation

Logical devices are represented by VkDevice handles:

```
VK_DEFINE_HANDLE(VkDevice)
```

A logical device is created as a connection to a physical device. To create a logical device, call:

- physicalDevice must be one of the device handles returned from a call to **vkEnumeratePhysicalDevices** (see Physical Device Enumeration).
- pCreateInfo is a pointer to a VkDeviceCreateInfo structure containing information about how to create the device.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pDevice points to a handle in which the created VkDevice is returned.

Multiple logical devices can be created from the same physical device. Logical device creation may fail due to lack of device-specific resources (in addition to the other errors). If that occurs, **vkCreateDevice** will return VK_ERROR_TOO_MANY_OBJECTS.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pCreateInfo must be a pointer to a valid VkDeviceCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pDevice must be a pointer to a VkDevice handle

Return Codes

Success

• VK SUCCESS

Failure

- VK ERROR OUT OF HOST MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_INITIALIZATION_FAILED
- VK_ERROR_EXTENSION_NOT_PRESENT
- VK_ERROR_FEATURE_NOT_PRESENT
- VK_ERROR_TOO_MANY_OBJECTS
- VK_ERROR_DEVICE_LOST

The VkDeviceCreateInfo structure is defined as:

```
typedef struct VkDeviceCreateInfo {
   VkStructureType
                                      sType;
   const void*
                                     pNext;
   VkDeviceCreateFlags
                                     flags;
   uint32 t
                                     queueCreateInfoCount;
   const VkDeviceQueueCreateInfo* pQueueCreateInfos;
   uint32_t
                                      enabledLayerCount;
                                    ppEnabledLayerNames;
   const char* const*
   uint32_t
                                      enabledExtensionCount;
                                    ppEnabledExtensionNames;
   const char* const*
   const VkPhysicalDeviceFeatures*     pEnabledFeatures;
} VkDeviceCreateInfo;
```

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- queueCreateInfoCount is the unsigned integer size of the pQueueCreateInfos array. Refer to the Queue Creation section below for further details.
- pQueueCreateInfos is a pointer to an array of VkDeviceQueueCreateInfo structures describing the queues that are requested to be created along with the logical device. Refer to the Queue Creation section below for further details.
- enabledLayerCount is deprecated and ignored.
- ppEnabledLayerNames is deprecated and ignored. See Device Layer Deprecation.
- enabledExtensionCount is the number of device extensions to enable.
- ppEnabledExtensionNames is a pointer to an array of enabledExtensionCount null-terminated UTF-8 strings containing the names of extensions to enable for the created device. See the Extensions section for further details.

• pEnabledFeatures is NULL or a pointer to a VkPhysicalDeviceFeatures structure that contains boolean indicators of all the features to be enabled. Refer to the Features section for further details.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- pQueueCreateInfos must be a pointer to an array of queueCreateInfoCount valid VkDeviceQueueCreateInfo structures
- If enabledLayerCount is not 0, ppEnabledLayerNames must be a pointer to an array of enabledLayerCount null-terminated strings
- If enabledExtensionCount is not 0, ppEnabledExtensionNames must be a pointer to an array of enabledExtensionCount null-terminated strings
- If pEnabledFeatures is not NULL, pEnabledFeatures must be a pointer to a valid VkPhysicalDeviceFeatures structure
- queueCreateInfoCount must be greater than 0
- The queueFamilyIndex member of any given element of pQueueCreateInfos must be unique within pQueueCreateInfos

4.2.2 Device Use

The following is a high-level list of VkDevice uses along with references on where to find more information:

- Creation of queues. See the Queues section below for further details.
- Creation and tracking of various synchronization constructs. See Synchronization and Cache Control for further details.
- Allocating, freeing, and managing memory. See Memory Allocation and Resource Creation for further details.
- Creation and destruction of command buffers and command buffer pools. See Command Buffers for further details.
- Creation, destruction, and management of graphics state. See Pipelines and Resource Descriptors, among others, for further details.

4.2.3 Lost Device

A logical device may become *lost* because of hardware errors, execution timeouts, power management events and/or platform-specific events. This may cause pending and future command execution to fail and cause hardware resources to be corrupted. When this happens, certain commands will return VK_ERROR_DEVICE_LOST (see Error Codes for a list

of such commands). After any such event, the logical device is considered *lost*. It is not possible to reset the logical device to a non-lost state, however the lost state is specific to a logical device (VkDevice), and the corresponding physical device (VkPhysicalDevice) may be otherwise unaffected. In some cases, the physical device may also be lost, and attempting to create a new logical device will fail, returning VK_ERROR_DEVICE_LOST. This is usually indicative of a problem with the underlying hardware, or its connection to the host. If the physical device has not been lost, and a new logical device is successfully created from that physical device, it must be in the non-lost state.

Note



Whilst logical device loss may be recoverable, in the case of physical device loss, it is unlikely that an application will be able to recover unless additional, unaffected physical devices exist on the system. The error is largely informational and intended only to inform the user that their hardware has probably developed a fault or become physically disconnected, and should be investigated further. In many cases, physical device loss may cause other more serious issues such as the operating system crashing; in which case it may not be reported via the Vulkan API.



Note

Undefined behavior caused by an application error may cause a device to become lost. However, such undefined behavior may also cause unrecoverable damage to the process, and it is then not guaranteed that the API objects, including the VkPhysicalDevice or the VkInstance are still valid or that the error is recoverable.

When a device is lost, its child objects are not implicitly destroyed and their handles are still valid. Those objects must still be destroyed before their parents or the device can be destroyed (see the Object Lifetime section). The host address space corresponding to device memory mapped using vkMapMemory is still valid, and host memory accesses to these mapped regions are still valid, but the contents are undefined. It is still legal to call any API command on the device and child objects.

Once a device is lost, command execution may fail, and commands that return a VkResult may return VK_ERROR_ DEVICE_LOST. Commands that do not allow run-time errors must still operate correctly for valid usage and, if applicable, return valid data.

Commands that wait indefinitely for device execution (namely vkDeviceWaitIdle, vkQueueWaitIdle, vkWaitForFences or vkAcquireNextImageKHR with a maximum timeout, and vkGetQueryPoolResults with the VK_QUERY_RESULT_WAIT_BIT bit set in flags) must return in finite time even in the case of a lost device, and return either VK_SUCCESS or VK_ERROR_DEVICE_LOST. For any command that may return VK_ERROR_DEVICE_LOST, for the purpose of determining whether a command buffer is pending execution, or whether resources are considered in-use by the device, a return value of VK_ERROR_DEVICE_LOST is equivalent to VK_SUCCESS.

4.2.4 Device Destruction

To destroy a device, call:

- device is the logical device to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

To ensure that no work is active on the device, vkDeviceWaitIdle can be used to gate the destruction of the device. Prior to destroying a device, an application is responsible for destroying/freeing any Vulkan objects that were created using that device as the first parameter of the corresponding vkCreate* or vkAllocate* command.



Note

The lifetime of each of these objects is bound by the lifetime of the VkDevice object. Therefore, to avoid resource leaks, it is critical that an application explicitly free all of these resources prior to calling **vkDestroy Device**.

Valid Usage

- If device is not NULL, device must be a valid VkDevice handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- All child objects created on device must have been destroyed prior to destroying device
- If VkAllocationCallbacks were provided when <code>device</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when device was created, pAllocator must be NULL

Host Synchronization

Host access to device must be externally synchronized

4.3 Queues

4.3.1 Queue Family Properties

As discussed in the Physical Device Enumeration section above, the

 ${\tt vkGetPhysicalDeviceQueueFamilyProperties} \ \ {\tt command} \ \ {\tt is} \ \ {\tt used} \ \ {\tt to} \ \ {\tt retrieve} \ \ {\tt details} \ \ {\tt about} \ \ {\tt the} \ \ {\tt queue} \ \ {\tt families} \ \ {\tt and} \ \ {\tt queues} \ \ {\tt supported} \ \ {\tt by} \ \ {\tt a} \ \ {\tt device}.$

Each index in the pQueueFamilyProperties array returned by

vkGetPhysicalDeviceQueueFamilyProperties describes a unique queue family on that physical device. These indices are used when creating queues, and they correspond directly with the <code>queueFamilyIndex</code> that is passed to the vkCreateDevice command via the VkDeviceQueueCreateInfo structure as described in the Queue Creation section below.

Grouping of queue families within a physical device is implementation-dependent.



Note

The general expectation is that a physical device groups all queues of matching capabilities into a single family. However, this is a recommendation to implementations and it is possible that a physical device may return two separate queue families with the same capabilities.

Once an application has identified a physical device with the queue(s) that it desires to use, it will create those queues in conjunction with a logical device. This is described in the following section.

4.3.2 Queue Creation

Creating a logical device also creates the queues associated with that device. The queues to create are described by a set of VkDeviceQueueCreateInfo structures that are passed to vkCreateDevice in pQueueCreateInfos.

Queues are represented by VkQueue handles:

```
VK_DEFINE_HANDLE (VkQueue)
```

The VkDeviceQueueCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- queueFamilyIndex is an unsigned integer indicating the index of the queue family to create on this device. This index corresponds to the index of an element of the pQueueFamilyProperties array that was returned by vkGetPhysicalDeviceQueueFamilyProperties.
- queueCount is an unsigned integer specifying the number of queues to create in the queue family indicated by queueFamilyIndex.
- pQueuePriorities is an array of queueCount normalized floating point values, specifying priorities of work that will be submitted to each created queue. See Queue Priority for more information.

Valid Usage

• sType must be VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO

- pNext must be NULL
- flags must be 0
- pQueuePriorities must be a pointer to an array of queueCount float values
- queueCount must be greater than 0
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties
- queueCount must be less than or equal to the queueCount member of the VkQueueFamilyProperties
 structure, as returned by vkGetPhysicalDeviceQueueFamilyProperties in the
 pQueueFamilyProperties[queueFamilyIndex]
- Each element of pQueuePriorities must be between 0.0 and 1.0 inclusive

To retrieve a handle to a VkQueue object, call:

- device is the logical device that owns the queue.
- queueFamilyIndex is the index of the queue family to which the queue belongs.
- queueIndex is the index within this queue family of the queue to retrieve.
- pQueue is a pointer to a VkQueue object that will be filled with the handle for the requested queue.

Valid Usage

- device must be a valid VkDevice handle
- pQueue must be a pointer to a VkQueue handle
- queueFamilyIndex must be one of the queue family indices specified when device was created, via the VkDeviceQueueCreateInfo structure
- queueIndex must be less than the number of queues created for the specified queue family index when device was created, via the queueCount member of the VkDeviceQueueCreateInfo structure

4.3.3 Queue Family Index

The queue family index is used in multiple places in Vulkan in order to tie operations to a specific family of queues.

When retrieving a handle to the queue via **vkGetDeviceQueue**, the queue family index is used to select which queue family to retrieve the VkQueue handle from as described in the previous section.

When creating a VkCommandPool object (see Command Pools), a queue family index is specified in the VkCommandPoolCreateInfo structure. Command buffers from this pool can only be submitted on queues corresponding to this queue family.

When creating VkImage (see Images) and VkBuffer (see Buffers) resources, a set of queue families is included in the VkImageCreateInfo and VkBufferCreateInfo structures to specify the queue families that can access the resource.

When inserting a VkBufferMemoryBarrier or VkImageMemoryBarrier (see Section 6.3) a source and destination queue family index is specified to allow the ownership of a buffer or image to be transferred from one queue family to another. See the Resource Sharing section for details.

4.3.4 Queue Priority

Each queue is assigned a priority, as set in the VkDeviceQueueCreateInfo structures when creating the device. The priority of each queue is a normalized floating point value between 0.0 and 1.0, which is then translated to a discrete priority level by the implementation. Higher values indicate a higher priority, with 0.0 being the lowest priority and 1.0 being the highest.

Within the same device, queues with higher priority may be allotted more processing time than queues with lower priority. The implementation makes no guarantees with regards to ordering or scheduling among queues with the same priority, other than the constraints defined by explicit scheduling primitives. The implementation make no guarantees with regards to queues across different devices.

An implementation may allow a higher-priority queue to starve a lower-priority queue on the same VkDevice until the higher-priority queue has no further commands to execute. The relationship of queue priorities must not cause queues on one VkDevice to starve queues on another VkDevice.

No specific guarantees are made about higher priority queues receiving more processing time or better quality of service than lower priority queues.

4.3.5 Queue Submission

Work is submitted to a queue via *queue submission* commands such as vkQueueSubmit. Queue submission commands define a set of *queue operations* to be executed by the underlying physical device, including synchronization with semaphores and fences.

Submission commands take as parameters a target queue, zero or more *batches* of work, and an optional fence to signal upon completion. Each batch consists of three distinct parts:

- 1. Zero or more semaphores to wait on before execution of the rest of the batch.
 - If present, these describe a semaphore wait operation.
- 2. Zero or more work items to execute.
 - If present, these describe a *queue operation* matching the work described.
- 3. Zero or more semaphores to signal upon completion of the work items.

• If present, these describe a semaphore signal operation.

If a fence is present in a queue submission, it describes a fence signal operation.

All work described by a queue submission command must be submitted to the queue before the command returns.

4.3.5.1 Sparse Memory Binding

In Vulkan it is possible to sparsely bind memory to buffers and images as described in the Sparse Resource chapter. Sparse memory binding is a queue operation. A queue whose flags include the VK_QUEUE_SPARSE_BINDING_BIT must be able to support the mapping of a virtual address to a physical address on the device. This causes an update to the page table mappings on the device. This update must be synchronized on a queue to avoid corrupting page table mappings during execution of graphics commands. By binding the sparse memory resources on queues, all commands that are dependent on the updated bindings are synchronized to only execute after the binding is updated. See the Synchronization and Cache Control chapter for how this synchronization is accomplished.

4.3.6 Queue Destruction

Queues are created along with a logical device during **vkCreateDevice**. All queues associated with a logical device are destroyed when **vkDestroyDevice** is called on that device.

Chapter 5

Command Buffers

Command buffers are objects used to record commands which can be subsequently submitted to a device queue for execution. There are two levels of command buffers - *primary command buffers*, which can execute secondary command buffers, and which are submitted to queues, and *secondary command buffers*, which can be executed by primary command buffers, and which are not directly submitted to queues.

Command buffers are represented by VkCommandBuffer handles:

VK_DEFINE_HANDLE (VkCommandBuffer)

Recorded commands include commands to bind pipelines and descriptor sets to the command buffer, commands to modify dynamic state, commands to draw (for graphics rendering), commands to dispatch (for compute), commands to execute secondary command buffers (for primary command buffers only), commands to copy buffers and images, and other commands.

Each command buffer manages state independently of other command buffers. There is no inheritance of state across primary and secondary command buffers, or between secondary command buffers. When a command buffer begins recording, all state in that command buffer is undefined. When secondary command buffer(s) are recorded to execute on a primary command buffer, the secondary command buffer inherits no state from the primary command buffer, and all state of the primary command buffer is undefined after an execute secondary command buffer command is recorded. There is one exception to this rule - if the primary command buffer is inside a render pass instance, then the render pass and subpass state is not disturbed by executing secondary command buffers. Whenever the state of a command buffer is undefined, the application must set all relevant state on the command buffer before any state dependent commands such as draws and dispatches are recorded, otherwise the behavior of executing that command buffer is undefined.

Unless otherwise specified, and without explicit synchronization, the various commands submitted to a queue via command buffers may execute in arbitrary order relative to each other, and/or concurrently. Also, the memory side-effects of those commands may not be directly visible to other commands without memory barriers. This is true within a command buffer, and across command buffers submitted to a given queue. See Section 6.3, Section 6.5 and Section 6.5.3 about synchronization primitives suitable to guarantee execution order and side-effect visibility between commands on a given queue.

Each command buffer is always in one of three states:

- *Initial state*: Before vkBeginCommandBuffer. Either vkBeginCommandBuffer has never been called, or the command buffer has been reset since it last recorded commands.
- Recording state: Between vkBeginCommandBuffer and vkEndCommandBuffer. The command buffer is in a state where it can record commands.

• Executable state: After vkEndCommandBuffer. The command buffer is in a state where it has finished recording commands and can be executed.

Resetting a command buffer is an operation that discards any previously recorded commands and puts a command buffer in the initial state. Resetting occurs as a result of vkResetCommandBuffer or vkResetCommandPool, or as part of vkBeginCommandBuffer (which additionally puts the command buffer in the recording state).

5.1 Command Pools

Command pools are opaque objects that command buffer memory is allocated from, and which allow the implementation to amortize the cost of resource creation across multiple command buffers. Command pools are application-synchronized, meaning that a command pool must not be used concurrently in multiple threads. That includes use via recording commands on any command buffers allocated from the pool, as well as operations that allocate, free, and reset command buffers or the pool itself.

Command pools are represented by VkCommandPool handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkCommandPool)
```

To create a command pool, call:

- device is the logical device that creates the command pool.
- pCreateInfo contains information used to create the command pool.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pCommandPool points to a VkCommandPool handle in which the created pool is returned.

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkCommandPoolCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pCommandPool must be a pointer to a VkCommandPool handle

Return Codes

Success

• VK SUCCESS

Failure

- VK ERROR OUT OF HOST MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkCommandPoolCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is a bitmask indicating usage behavior for the pool and command buffers allocated from it. Bits which can be set include:

```
typedef enum VkCommandPoolCreateFlagBits {
    VK_COMMAND_POOL_CREATE_TRANSIENT_BIT = 0x00000001,
    VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT = 0x00000002,
} VkCommandPoolCreateFlagBits;
```

- VK_COMMAND_POOL_CREATE_TRANSIENT_BIT indicates that command buffers allocated from the pool will be short-lived, meaning that they will be reset or freed in a relatively short timeframe. This flag may be used by the implementation to control memory allocation behavior within the pool.
- VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT controls whether command buffers allocated from the pool can be individually reset. If this flag is set, individual command buffers allocated from the pool can be reset either explicitly, by calling vkResetCommandBuffer, or implicitly, by calling vkBeginCommandBuffer on an executable command buffer. If this flag is not set, then vkResetCommandBuffer and vkBeginCommandBuffer (on an executable command buffer) must not be called on the command buffers allocated from the pool, and they can only be reset in bulk by calling vkResetCommandPool.
- queueFamilyIndex designates a queue family as described in section Queue Family Properties. All command buffers allocated from this command pool must be submitted on queues from the same queue family.

- sType must be VK_STRUCTURE_TYPE_COMMAND_POOL_CREATE_INFO
- pNext must be NULL
- flags must be a valid combination of VkCommandPoolCreateFlagBits values
- queueFamilyIndex must be the index of a queue family available in the calling command's device parameter

To reset a command pool, call:

- device is the logical device that owns the command pool.
- commandPool is the command pool to reset.
- flags contains additional flags controlling the behavior of the reset. Bits which can be set include:

```
typedef enum VkCommandPoolResetFlagBits {
    VK_COMMAND_POOL_RESET_RELEASE_RESOURCES_BIT = 0x00000001,
} VkCommandPoolResetFlagBits;
```

If flags includes VK_COMMAND_POOL_RESET_RELEASE_RESOURCES_BIT, resetting a command pool recycles all of the resources from the command pool back to the system.

Resetting a command pool recycles all of the resources from all of the command buffers allocated from the command pool back to the command pool. All command buffers that have been allocated from the command pool are put in the initial state.

- device must be a valid VkDevice handle
- commandPool must be a valid VkCommandPool handle
- flags must be a valid combination of VkCommandPoolResetFlagBits values
- commandPool must have been created, allocated, or retrieved from device
- All VkCommandBuffer objects allocated from commandPool must not currently be pending execution

• Host access to commandPool must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To destroy a command pool, call:

- device is the logical device that destroys the command pool.
- *commandPool* is the handle of the command pool to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

When a pool is destroyed, all command buffers allocated from the pool are implicitly freed and become invalid. Command buffers allocated from a given pool do not need to be freed before destroying that command pool.

- device must be a valid VkDevice handle
- If commandPool is not VK_NULL_HANDLE, commandPool must be a valid VkCommandPool handle
- $\bullet \ \ If \ \textit{pAllocator} \ is \ not \ \texttt{NULL}, \ \textit{pAllocator} \ must \ be \ a \ pointer \ to \ a \ valid \ \texttt{VkAllocationCallbacks} \ structure$
- If commandPool is a valid handle, it must have been created, allocated, or retrieved from device
- All VkCommandBuffer objects allocated from commandPool must not be pending execution

- If VkAllocationCallbacks were provided when *commandPool* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when commandPool was created, pAllocator must be NULL

• Host access to commandPool must be externally synchronized

5.2 Command Buffer Allocation and Management

To allocate command buffers, call:

- device is the logical device that owns the command pool.
- pAllocateInfo is a pointer to an instance of the VkCommandBufferAllocateInfo structure describing parameters of the allocation.
- pCommandBuffers is a pointer to an array of VkCommandBuffer handles in which the resulting command buffer objects are returned. The array must be at least the length specified by the commandBufferCount member of pAllocateInfo. Each allocated command buffer begins in the initial state.

- device must be a valid VkDevice handle
- $\bullet \ \textit{pAllocateInfo} \ \textbf{must} \ \textbf{be} \ \textbf{a} \ \textbf{pointer} \ \textbf{to} \ \textbf{a} \ \textbf{valid} \ \texttt{VkCommandBufferAllocateInfo} \ \textbf{structure}$
- pCommandBuffers must be a pointer to an array of pAllocateInfo→commandBufferCount VkCommandBuffer handles

• Host access to pAllocateInfo→commandPool must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkCommandBufferAllocateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- commandPool is the name of the command pool that the command buffers allocate their memory from.
- level determines whether the command buffers are primary or secondary command buffers. Possible values include:

```
typedef enum VkCommandBufferLevel {
   VK_COMMAND_BUFFER_LEVEL_PRIMARY = 0,
   VK_COMMAND_BUFFER_LEVEL_SECONDARY = 1,
} VkCommandBufferLevel;
```

• commandBufferCount is the number of command buffers to allocate from the pool.

- sType must be $VK_STRUCTURE_TYPE_COMMAND_BUFFER_ALLOCATE_INFO$
- pNext must be NULL
- commandPool must be a valid VkCommandPool handle
- level must be a valid VkCommandBufferLevel value
- commandBufferCount must be greater than 0

To reset command buffers, call:

- commandBuffer is the command buffer to reset. The command buffer can be in any state, and is put in the initial state.
- flags is a bitmask controlling the reset operation. Bits which can be set include:

```
typedef enum VkCommandBufferResetFlagBits {
    VK_COMMAND_BUFFER_RESET_RELEASE_RESOURCES_BIT = 0x00000001,
} VkCommandBufferResetFlagBits;
```

If <code>flags</code> includes <code>VK_COMMAND_BUFFER_RESET_RELEASE_RESOURCES_BIT</code>, then most or all memory resources currently owned by the command buffer should be returned to the parent command pool. If this flag is not set, then the command buffer may hold onto memory resources and reuse them when recording commands.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- flags must be a valid combination of VkCommandBufferResetFlagBits values
- commandBuffer must not currently be pending execution
- commandBuffer must have been allocated from a pool that was created with the VK_COMMAND_POOL_ CREATE_RESET_COMMAND_BUFFER_BIT

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To free command buffers, call:

- device is the logical device that owns the command pool.
- commandPool is the handle of the command pool that the command buffers were allocated from.
- commandBufferCount is the length of the pCommandBuffers array.
- pCommandBuffers is an array of handles of command buffers to free.

- device must be a valid VkDevice handle
- commandPool must be a valid VkCommandPool handle
- commandBufferCount must be greater than 0
- commandPool must have been created, allocated, or retrieved from device
- Each element of pCommandBuffers that is a valid handle must have been created, allocated, or retrieved from commandPool
- All elements of pCommandBuffers must not be pending execution
- pCommandBuffers must be a pointer to an array of commandBufferCount VkCommandBuffer handles, each element of which must either be a valid handle or **NULL**

- Host access to commandPool must be externally synchronized
- Host access to each member of pCommandBuffers must be externally synchronized

5.3 Command Buffer Recording

To begin recording a command buffer, call:

```
VkResult vkBeginCommandBuffer(
     VkCommandBuffer commandBuffer,
     const VkCommandBufferBeginInfo* pBeginInfo);
```

- commandBuffer is the handle of the command buffer which is to be put in the recording state.
- pBeginInfo is an instance of the VkCommandBufferBeginInfo structure, which defines additional information about how the command buffer begins recording.

- commandBuffer must be a valid VkCommandBuffer handle
- pBeginInfo must be a pointer to a valid VkCommandBufferBeginInfo structure
- commandBuffer must not be in the recording state
- commandBuffer must not currently be pending execution
- If commandBuffer was allocated from a VkCommandPool which did not have the VK_COMMAND_POOL_ CREATE_RESET_COMMAND_BUFFER_BIT flag set, commandBuffer must be in the initial state
- If commandBuffer is a secondary command buffer, the pInheritanceInfo member of pBeginInfo must be a valid VkCommandBufferInheritanceInfo structure
- If commandBuffer is a secondary command buffer and either the occlusionQueryEnable member of the pInheritanceInfo member of pBeginInfo is VK_FALSE, or the precise occlusion queries feature is not enabled, the queryFlags member of the pInheritanceInfo member pBeginInfo must not contain VK_QUERY CONTROL PRECISE BIT

• Host access to commandBuffer must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkCommandBufferBeginInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is a bitmask indicating usage behavior for the command buffer. Bits which can be set include:

```
typedef enum VkCommandBufferUsageFlagBits {
   VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT = 0x00000001,
   VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT = 0x00000002,
   VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT = 0x00000004,
} VkCommandBufferUsageFlagBits;
```

- VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT indicates that each recording of the command buffer will only be submitted once, and the command buffer will be reset and recorded again between each submission.
- VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT indicates that a secondary command buffer is considered to be entirely inside a render pass. If this is a primary command buffer, then this bit is ignored.
- Setting VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT allows the command buffer to be resubmitted to a queue or recorded into a primary command buffer while it is pending execution.
- pInheritanceInfo is a pointer to a VkCommandBufferInheritanceInfo structure, which is used if commandBuffer is a secondary command buffer. If this is a primary command buffer, then this value is ignored.

- sType must be VK_STRUCTURE_TYPE_COMMAND_BUFFER_BEGIN_INFO
- pNext must be NULL
- flags must be a valid combination of VkCommandBufferUsageFlagBits values
- If flags contains VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT, the renderPass member of pInheritanceInfo must be a valid VkRenderPass
- If flags contains VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT, the subpass member of pInheritanceInfo must be a valid subpass index within the renderPass member of pInheritanceInfo
- If flags contains VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT, the framebuffer member of pInheritanceInfo must be either VK_NULL_HANDLE, or a valid VkFramebuffer that is compatible with the renderPass member of pInheritanceInfo

If the command buffer is a secondary command buffer, then the VkCommandBufferInheritanceInfo structure defines any state that will be inherited from the primary command buffer:

```
typedef struct VkCommandBufferInheritanceInfo {
   VkStructureType
                                     sType;
   const void*
                                    pNext;
   VkRenderPass
                                    renderPass;
   uint32_t
                                    subpass;
   VkFramebuffer
                                    framebuffer;
                                    occlusionQueryEnable;
   VkBool32
   VkQueryControlFlags
                                    queryFlags;
   VkQueryPipelineStatisticFlags
                                    pipelineStatistics;
 VkCommandBufferInheritanceInfo;
```

- renderPass is a VkRenderPass object defining which render passes the VkCommandBuffer will be compatible with and can be executed within. If the VkCommandBuffer will not be executed within a render pass instance, renderPass is ignored.
- subpass is the index of the subpass within renderPass that the VkCommandBuffer will be executed within. If the VkCommandBuffer will not be executed within a render pass instance, subpass is ignored.
- framebuffer optionally refers to the VkFramebuffer object that the VkCommandBuffer will be rendering to if it is executed within a render pass instance. It can be VK_NULL_HANDLE if the framebuffer is not known, or if the VkCommandBuffer will not be executed within a render pass instance.



Note

Specifying the exact framebuffer that the secondary command buffer will be executed with may result in better performance at command buffer execution time.

- occlusionQueryEnable indicates whether the command buffer can be executed while an occlusion query is active in the primary command buffer. If this is VK_TRUE, then this command buffer can be executed whether the primary command buffer has an occlusion query active or not. If this is VK_FALSE, then the primary command buffer must not have an occlusion query active.
- queryFlags indicates the query flags that can be used by an active occlusion query in the primary command buffer
 when this secondary command buffer is executed. If this value includes the VK_QUERY_CONTROL_PRECISE_BIT
 bit, then the active query can return boolean results or actual sample counts. If this bit is not set, then the active query
 must not use the VK_QUERY_CONTROL_PRECISE_BIT bit. If this is a primary command buffer, then this value is
 ignored.
- pipelineStatistics indicates the set of pipeline statistics that can be counted by an active query in the primary command buffer when this secondary command buffer is executed. If this value includes a given bit, then this command buffer can be executed whether the primary command buffer has a pipeline statistics query active that includes this bit or not. If this value excludes a given bit, then the active pipeline statistics query must not be from a query pool that counts that statistic.

- sType must be VK_STRUCTURE_TYPE_COMMAND_BUFFER_INHERITANCE_INFO
- pNext must be NULL
- Both of framebuffer, and renderPass that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- If the inherited queries feature is not enabled, occlusionQueryEnable must be VK_FALSE
- If the inherited queries feature is enabled, *queryFlags* must be a valid combination of VkQueryControlFlagBits values
- If the pipeline statistics queries feature is not enabled, pipelineStatistics must be 0

A primary command buffer is considered to be pending execution from the time it is submitted via **vkQueueSubmit** until that submission completes.

A secondary command buffer is considered to be pending execution from the time its execution is recorded into a primary buffer (via **vkCmdExecuteCommands**) until the final time that primary buffer's submission to a queue completes. If, after the primary buffer completes, the secondary command buffer is recorded to execute on a different primary buffer, the first primary buffer must not be resubmitted until after it is reset with vkResetCommandBuffer unless the secondary command buffer was recorded with VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT.

If VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT is not set on a secondary command buffer, that command buffer must not be used more than once in a given primary command buffer. Furthermore, if a secondary command buffer without VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT set is recorded to execute in a primary command buffer with VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT set, the primary command buffer must not be pending execution more than once at a time.



Note

On some implementations, not using the VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT bit enables command buffers to be patched in-place if needed, rather than creating a copy of the command buffer.

If a command buffer is in the executable state and the command buffer was allocated from a command pool with the VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT flag set, then **vkBeginCommandBuffer** implicitly resets the command buffer, behaving as if **vkResetCommandBuffer** had been called with VK_COMMAND_BUFFER_RESET_RELEASE_RESOURCES_BIT not set. It then puts the command buffer in the recording state.

Once recording starts, an application records a sequence of commands (**vkCmd***) to set state in the command buffer, draw, dispatch, and other commands.

To complete recording of a command buffer, call:

• commandBuffer is the command buffer to complete recording. The command buffer must have been in the recording state, and is moved to the executable state.

If there was an error during recording, the application will be notified by an unsuccessful return code returned by **vkEndCommandBuffer**. If the application wishes to further use the command buffer, the command buffer must be reset.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- If commandBuffer is a primary command buffer, there must not be an active render pass instance
- All queries made active during the recording of commandBuffer must have been made inactive

Host Synchronization

Host access to commandBuffer must be externally synchronized

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK ERROR OUT OF DEVICE MEMORY

When a command buffer is in the executable state, it can be submitted to a queue for execution.

5.4 Command Buffer Submission

To submit command buffers to a queue, call:

- queue is the queue that the command buffers will be submitted to.
- submitCount is the number of elements in the pSubmits array.
- pSubmits is a pointer to an array of VkSubmitInfo structures, each specifying a command buffer submission batch.
- fence is an optional handle to a fence to be signaled. If fence is not VK_NULL_HANDLE, it defines a fence signal operation.



Note

Submission can be a high overhead operation, and applications should attempt to batch work together into as few calls to **vkQueueSubmit** as possible.

vkQueueSubmit is a queue submission command, with each batch defined by an element of pSubmits as an instance of the VkSubmitInfo structure.

Fence and semaphore operations submitted with vkQueueSubmit have additional ordering constraints compared to other submission commands, with dependencies involving previous and subsequent queue operations. Information about these additional constraints can be found in the semaphore and fence sections of the synchronization chapter.

Details on the interaction of pWaitDstStageMask with synchronization are described in the semaphore wait operation section of the synchronization chapter.

- queue must be a valid VkQueue handle
- If submitCount is not 0, pSubmits must be a pointer to an array of submitCount valid VkSubmitInfo structures
- If fence is not VK_NULL_HANDLE, fence must be a valid VkFence handle
- Both of fence, and queue that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- If fence is not VK_NULL_HANDLE, fence must be unsignaled
- If fence is not VK_NULL_HANDLE, fence must not be associated with any other queue command that has not yet completed execution on that queue

Host Synchronization

- Host access to queue must be externally synchronized
- Host access to pSubmits[].pWaitSemaphores[] must be externally synchronized
- Host access to pSubmits[].pSignalSemaphores[] must be externally synchronized
- Host access to fence must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
-	-	Any

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

The VkSubmitInfo structure is defined as:

```
typedef struct VkSubmitInfo {
   VkStructureType
                                sType;
   const void*
                                pNext;
   uint32_t
                                waitSemaphoreCount;
   uint32_t
const VkSemaphore*
const VkPipelineStageFlags*
pWaitDstStageMask;
   uint32_t
                                commandBufferCount;
   uint32_t
                                signalSemaphoreCount;
   const VkSemaphore*
                               pSignalSemaphores;
} VkSubmitInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- waitSemaphoreCount is the number of semaphores upon which to wait before executing the command buffers for the batch.
- pWaitSemaphores is a pointer to an array of semaphores upon which to wait before the command buffers for this batch begin execution. If semaphores to wait on are provided, they define a semaphore wait operation.
- pWaitDstStageMask is a pointer to an array of pipeline stages at which each corresponding semaphore wait will
 occur.
- commandBufferCount is the number of command buffers to execute in the batch.
- pCommandBuffers is a pointer to an array of command buffers to execute in the batch. The command buffers submitted in a batch begin execution in the order they appear in pCommandBuffers, but may complete out of order.
- signalSemaphoreCount is the number of semaphores to be signaled once the commands specified in pCommandBuffers have completed execution.
- pSignalSemaphores is a pointer to an array of semaphores which will be signaled when the command buffers for this batch have completed execution. If semaphores to be signaled are provided, they define a semaphore signal operation.

Valid Usage

• sType must be VK_STRUCTURE_TYPE_SUBMIT_INFO

- pNext must be NULL
- If waitSemaphoreCount is not 0, pWaitSemaphores must be a pointer to an array of waitSemaphoreCount valid VkSemaphore handles
- If waitSemaphoreCount is not 0, pWaitDstStageMask must be a pointer to an array of waitSemaphoreCount valid combinations of VkPipelineStageFlagBits values
- Each element of pWaitDstStageMask must not be 0
- If commandBufferCount is not 0, pCommandBuffers must be a pointer to an array of commandBufferCount valid VkCommandBuffer handles
- If signalSemaphoreCount is not 0, pSignalSemaphores must be a pointer to an array of signalSemaphoreCount valid VkSemaphore handles
- Each of the elements of pCommandBuffers, the elements of pSignalSemaphores, and the elements of pWaitSemaphores that are valid handles must have been created, allocated, or retrieved from the same VkDevice.
- Any given element of pSignalSemaphores must currently be unsignaled
- Any given element of pCommandBuffers must either have been recorded with the VK_COMMAND_BUFFER_ USAGE_SIMULTANEOUS_USE_BIT, or not currently be executing on the device
- Any given element of pCommandBuffers must be in the executable state
- If any given element of pCommandBuffers contains commands that execute secondary command buffers, those secondary command buffers must have been recorded with the VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS USE BIT, or not currently be executing on the device
- If any given element of pCommandBuffers was recorded with VK_COMMAND_BUFFER_USAGE_ONE_TIME_ SUBMIT_BIT, it must not have been previously submitted without re-recording that command buffer
- If any given element of pCommandBuffers contains commands that execute secondary command buffers recorded with VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT, each such secondary command buffer must not have been previously submitted without re-recording that command buffer
- Any given element of pCommandBuffers must not contain commands that execute a secondary command buffer, if that secondary command buffer has been recorded in another primary command buffer after it was recorded into this VkCommandBuffer
- Any given element of pCommandBuffers must have been allocated from a VkCommandPool that was created for the same queue family that the calling command's queue belongs to
- Any given element of pCommandBuffers must not have been allocated with VK_COMMAND_BUFFER_LEVEL_ SECONDARY
- Any given element of VkSemaphore in pWaitSemaphores must refer to a prior signal of that VkSemaphore that will not be consumed by any other wait on that semaphore
- If the geometry shaders feature is not enabled, any given element of pWaitDstStageMask must not contain VK_ PIPELINE STAGE GEOMETRY SHADER BIT
- If the tessellation shaders feature is not enabled, any given element of pWaitDstStageMask must not contain VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT or VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT

When submitting work that operates on memory imported from a Direct3D 11 resource to a queue, the keyed mutex mechanism may be used in addition to Vulkan semaphores to synchronize the work. Keyed mutexes are a property of a properly created shareable Direct3D 11 resource. They can only be used if the imported resource was created with the D3D11_RESOURCE_MISC_SHARED_KEYEDMUTEX flag.

To acquire keyed mutexes before submitted work and/or release them after, add a

VkWin32KeyedMutexAcquireReleaseInfoNV structure to the pNext chain of the VkSubmitInfo structure.

The $\mbox{VkWin32KeyedMutexAcquireReleaseInfoNV}$ structure is defined as:

- acquireCount is the number of entries in the pAcquireSyncs, pAcquireKeys, and pAcquireTimeoutMilliseconds arrays.
- pAcquireSyncs is a pointer to an array of VkDeviceMemory objects which were imported from Direct3D 11 resources.
- pAcquireKeys is a pointer to an array of mutex key values to wait for prior to beginning the submitted work. Entries refer to the keyed mutex associated with the corresponding entries in pAcquireSyncs.
- pAcquireTimeoutMilliseconds is an array of timeout values, in millisecond units, for each acquire specified in pAcquireKeys.
- releaseCount is the number of entries in the pReleaseSyncs and pReleaseKeys arrays.
- pReleaseSyncs is a pointer to an array of VkDeviceMemory objects which were imported from Direct3D 11 resources.
- pReleaseKeys is a pointer to an array of mutex key values to set when the submitted work has completed. Entries refer to the keyed mutex associated with the corresponding entries in pReleaseSyncs.

- sType must be VK_STRUCTURE_TYPE_WIN32_KEYED_MUTEX_ACQUIRE_RELEASE_INFO_NV
- pNext must be NULL
- pAcquireSyncs must be a pointer to a valid VkDeviceMemory handle
- pAcquireKeys must be a pointer to a valid uint 64_t value

- pAcquireTimeoutMilliseconds must be a pointer to a valid uint 32_t value
- pReleaseSyncs must be a pointer to a valid VkDeviceMemory handle
- pReleaseKeys must be a pointer to a valid uint 64_t value
- Both of pAcquireSyncs, and pReleaseSyncs must have been created, allocated, or retrieved from the same VkDevice

5.5 Queue Forward Progress

The application must ensure that command buffer submissions will be able to complete without any subsequent operations by the application on any queue. After any call to **vkQueueSubmit**, for every queued wait on a semaphore there must be a prior signal of that semaphore that will not be consumed by a different wait on the semaphore.

Command buffers in the submission can include vkCmdWaitEvents commands that wait on events that will not be signaled by earlier commands in the queue. Such events must be signaled by the application using vkSetEvent, and the vkCmdWaitEvents commands that wait upon them must not be inside a render pass instance. Implementations may have limits on how long the command buffer will wait, in order to avoid interfering with progress of other clients of the device. If the event is not signaled within these limits, results are undefined and may include device loss.

5.6 Secondary Command Buffer Execution

A secondary command buffer must not be directly submitted to a queue. Instead, secondary command buffers are recorded to execute as part of a primary command buffer with the command:

- commandBuffer is a handle to a primary command buffer that the secondary command buffers are executed in.
- commandBufferCount is the length of the pCommandBuffers array.
- pCommandBuffers is an array of secondary command buffer handles, which are recorded to execute in the primary command buffer in the order they are listed in the array.

Once **vkCmdExecuteCommands** has been called, any prior executions of the secondary command buffers specified by *pCommandBuffers* in any other primary command buffer become invalidated, unless those secondary command buffers were recorded with VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT.

Valid Usage

• commandBuffer must be a valid VkCommandBuffer handle

- pCommandBuffers must be a pointer to an array of commandBufferCount valid VkCommandBuffer handles
- commandBuffer must be in the recording state
- The VkCommandPool that *commandBuffer* was allocated from must support transfer, graphics, or compute operations
- commandBuffer must be a primary VkCommandBuffer
- commandBufferCount must be greater than 0
- Both of commandBuffer, and the elements of pCommandBuffers must have been created, allocated, or retrieved from the same VkDevice
- commandBuffer must have been allocated with a level of VK_COMMAND_BUFFER_LEVEL_PRIMARY
- Any given element of pCommandBuffers must have been allocated with a level of VK_COMMAND_BUFFER_ LEVEL_SECONDARY
- Any given element of pCommandBuffers must not be already pending execution in commandBuffer, or appear twice in pCommandBuffers, unless it was recorded with the VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT flag
- Any given element of pCommandBuffers must not be already pending execution in any other VkCommandBuffer, unless it was recorded with the VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT flag
- ullet Any given element of pCommandBuffers must be in the executable state
- Any given element of pCommandBuffers must have been allocated from a VkCommandPool that was created for the same queue family as the VkCommandPool from which commandBuffer was allocated
- If **vkCmdExecuteCommands** is being called within a render pass instance, that render pass instance must have been begun with the *contents* parameter of **vkCmdBeginRenderPass** set to VK_SUBPASS_CONTENTS_SECONDARY_COMMAND_BUFFERS
- If **vkCmdExecuteCommands** is being called within a render pass instance, any given element of pCommandBuffers must have been recorded with the VK_COMMAND_BUFFER_USAGE_RENDER_PASS_ CONTINUE_BIT
- If **vkCmdExecuteCommands** is being called within a render pass instance, any given element of *pCommandBuffers* must have been recorded with VkCommandBufferInheritanceInfo::subpass set to the index of the subpass which the given command buffer will be executed in
- If **vkCmdExecuteCommands** is being called within a render pass instance, any given element of *pCommandBuffers* must have been recorded with a render pass that is compatible with the current render pass see Section 7.2
- If vkCmdExecuteCommands is being called within a render pass instance, and any given element of pCommandBuffers was recorded with VkCommandBufferInheritanceInfo::framebuffer not equal to VK_NULL_HANDLE, that VkFramebuffer must match the VkFramebuffer used in the current render pass instance
- If vkCmdExecuteCommands is not being called within a render pass instance, any given element of pCommandBuffers must not have been recorded with the VK_COMMAND_BUFFER_USAGE_RENDER_PASS_ CONTINUE BIT

- If the inherited queries feature is not enabled, commandBuffer must not have any queries active
- If commandBuffer has a VK_QUERY_TYPE_OCCLUSION query active, then each element of pCommandBuffers must have been recorded with VkCommandBufferInheritanceInfo::occlusionQueryEnable set to VK_TRUE
- If commandBuffer has a VK_QUERY_TYPE_OCCLUSION query active, then each element of pCommandBuffers must have been recorded with VkCommandBufferInheritanceInfo::queryFlags having all bits set that are set for the query
- If commandBuffer has a VK_QUERY_TYPE_PIPELINE_STATISTICS query active, then each element of pCommandBuffers must have been recorded with VkCommandBufferInheritanceInfo::pipelineStatistics having all bits set that are set in the VkQueryPool the query uses
- Any given element of pCommandBuffers must not begin any query types that are active in commandBuffer

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	TRANSFER
		GRAPHICS
		COMPUTE

Chapter 6

Synchronization and Cache Control

Synchronization of access to resources is primarily the responsibility of the application. In Vulkan, there are four forms of concurrency during execution: between the host and device, between the queues, between queue submissions, and between commands within a command buffer. Vulkan provides the application with a set of synchronization primitives for these purposes. Further, memory caches and other optimizations mean that the normal flow of command execution does not guarantee that all memory transactions from a command are immediately visible to other agents with views into a given range of memory. Vulkan also provides barrier operations to ensure this type of synchronization.

Four synchronization primitive types are exposed by Vulkan. These are:

- Fences
- · Semaphores
- Events
- Barriers

Each is covered in detail in its own subsection of this chapter. Fences are used to communicate completion of execution of command buffer submissions to queues back to the application. Fences can therefore be used as a coarse-grained synchronization mechanism. Semaphores are generally associated with resources or groups of resources and can be used to marshal ownership of shared data. Their status is not visible to the host. Events provide a finer-grained synchronization primitive which can be signaled at command level granularity by both device and host, and can be waited upon by either. Barriers provide execution and memory synchronization between sets of commands.

6.1 Fences

Fences can be used by the host to determine completion of execution of queue operations.

A fence's status is always either *signaled* or *unsignaled*. The host can poll the status of a single fence, or wait for any or all of a group of fences to become signaled.

Fences are represented by VkFence handles:

VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkFence)

To create a new fence object, use the command

- device is the logical device that creates the fence.
- pCreateInfo points to a VkFenceCreateInfo structure specifying the state of the fence object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pFence points to a handle in which the resulting fence object is returned.

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkFenceCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pFence must be a pointer to a VkFence handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkFenceCreateInfo structure is defined as:

• flags defines the initial state and behavior of the fence. Bits which can be set include:

```
typedef enum VkFenceCreateFlagBits {
    VK_FENCE_CREATE_SIGNALED_BIT = 0x00000001,
} VkFenceCreateFlagBits;
```

If flags contains VK_FENCE_CREATE_SIGNALED_BIT then the fence object is created in the signaled state. Otherwise it is created in the unsignaled state.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_FENCE_CREATE_INFO
- pNext must be NULL
- flags must be a valid combination of VkFenceCreateFlagBits values

To destroy a fence, call:

- device is the logical device that destroys the fence.
- fence is the handle of the fence to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If fence is not VK_NULL_HANDLE, fence must be a valid VkFence handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If fence is a valid handle, it must have been created, allocated, or retrieved from device
- · fence must not be associated with any queue command that has not yet completed execution on that queue
- If VkAllocationCallbacks were provided when fence was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when fence was created, pAllocator must be NULL

• Host access to fence must be externally synchronized

To query the status of a fence from the host, use the command

- device is the logical device that owns the fence.
- fence is the handle of the fence to query.

Upon success, **vkGetFenceStatus** returns the status of the fence, which is one of:

- VK_SUCCESS indicates that the fence is signaled.
- VK_NOT_READY indicates that the fence is unsignaled.

Valid Usage

- device must be a valid VkDevice handle
- fence must be a valid VkFence handle
- fence must have been created, allocated, or retrieved from device

Return Codes

Success

- VK_SUCCESS
- VK NOT READY

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

To reset the status of one or more fences to the unsignaled state, use the command:

- device is the logical device that owns the fences.
- fenceCount is the number of fences to reset.
- pFences is a pointer to an array of fenceCount fence handles to reset.

If a fence is already in the unsignaled state, then resetting it has no effect.

Valid Usage

- device must be a valid VkDevice handle
- pFences must be a pointer to an array of fenceCount valid VkFence handles
- fenceCount must be greater than 0
- Each element of pFences must have been created, allocated, or retrieved from device
- Any given element of pFences must not currently be associated with any queue command that has not yet completed execution on that queue

Host Synchronization

• Host access to each member of pFences must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

Fences can be signaled by including them in a queue submission command, defining a queue operation to signal that fence. This *fence signal operation* defines the first half of a memory dependency, guaranteeing that all memory accesses defined by the queue submission are made available, and that queue operations described by that submission have completed execution. This half of the memory dependency does not include host availability of memory accesses. The second half of the dependency can be defined by vkWaitForFences.

Fence signal operations for vkQueueSubmit additionally include all queue operations previously submitted via vkQueueSubmit in their half of a memory dependency.

To cause the host to wait until any one or all of a group of fences is signaled, use the command:

- device is the logical device that owns the fences.
- fenceCount is the number of fences to wait on.
- pFences is a pointer to an array of fenceCount fence handles.
- waitAll is the condition that must be satisfied to successfully unblock the wait. If waitAll is VK_TRUE, then the condition is that all fences in pFences are signaled. Otherwise, the condition is that at least one fence in pFences is signaled.
- timeout is the timeout period in units of nanoseconds. timeout is adjusted to the closest value allowed by the implementation-dependent timeout accuracy, which may be substantially longer than one nanosecond, and may be longer than the requested period.

If the condition is satisfied when **vkWaitForFences** is called, then **vkWaitForFences** returns immediately. If the condition is not satisfied at the time **vkWaitForFences** is called, then **vkWaitForFences** will block and wait up to timeout nanoseconds for the condition to become satisfied.

If timeout is zero, then **vkWaitForFences** does not wait, but simply returns the current state of the fences. VK_TIMEOUT will be returned in this case if the condition is not satisfied, even though no actual wait was performed.

If the specified timeout period expires before the condition is satisfied, **vkWaitForFences** returns VK_TIMEOUT. If the condition is satisfied before *timeout* nanoseconds has expired, **vkWaitForFences** returns VK_SUCCESS.

vkWaitForFences defines the second half of a memory dependency with the host, for each fence being waited on. The memory dependency defined by signaling a fence and waiting on the host does not guarantee that the results of memory accesses will be visible to the host, or that the memory is available. To provide that guarantee, the application must insert a memory barrier between the device writes and the end of the submission that will signal the fence, with <code>dstAccessMask</code> having the VK_ACCESS_HOST_READ_BIT bit set, with <code>dstStageMask</code> having the VK_PIPELINE_STAGE_HOST_BIT bit set, and with the appropriate <code>srcStageMask</code> and <code>srcAccessMask</code> members set to guarantee completion of the writes. If the memory was allocated without the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT set, then **vkInvalidateMappedMemoryRanges** must be called after the fence is signaled in order to ensure the writes are visible to the host, as described in Host Access to Device Memory Objects.

- device must be a valid VkDevice handle
- pFences must be a pointer to an array of fenceCount valid VkFence handles
- fenceCount must be greater than 0
- Each element of pFences must have been created, allocated, or retrieved from device

Return Codes

Success

- VK_SUCCESS
- VK_TIMEOUT

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

6.2 Semaphores

Semaphores are used to coordinate queue operations both within a queue and between different queues. A semaphore's status is always either *signaled* or *unsignaled*.

Semaphores are represented by VkSemaphore handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkSemaphore)
```

To create a new semaphore object, use the command

- device is the logical device that creates the semaphore.
- pCreateInfo points to a VkSemaphoreCreateInfo structure specifying the state of the semaphore object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

• pSemaphore points to a handle in which the resulting semaphore object is returned. The semaphore is created in the unsignaled state.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkSemaphoreCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSemaphore must be a pointer to a VkSemaphore handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkSemaphoreCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.

- sType must be VK STRUCTURE TYPE SEMAPHORE CREATE INFO
- pNext must be NULL
- flags must be 0

To destroy a semaphore, call:

- device is the logical device that destroys the semaphore.
- semaphore is the handle of the semaphore to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If semaphore is not VK_NULL_HANDLE, semaphore must be a valid VkSemaphore handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If semaphore is a valid handle, it must have been created, allocated, or retrieved from device
- semaphore must not be associated with any queue command that has not yet completed execution on that queue
- If VkAllocationCallbacks were provided when *semaphore* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when semaphore was created, pAllocator must be NULL

Host Synchronization

• Host access to semaphore must be externally synchronized

Semaphores can be signaled by including them in a batch as part of a queue submission command, defining a queue operation to signal that semaphore. This *semaphore signal operation* defines the first half of a memory dependency, guaranteeing that all memory accesses defined by the submitted queue operations in the batch are made available, and that those queue operations have completed execution.

Semaphore signal operations for vkQueueSubmit additionally include all queue operations previously submitted via vkQueueSubmit in their half of a memory dependency, and all batches that are stored at a lower index in the same pSubmits array.

Signaling of semaphores can be waited on by similarly including them in a batch, defining a queue operation to wait for a signal. A semaphore wait operation defines the second half of a memory dependency for the semaphores being waited on. This half of the memory dependency guarantees that the first half has completed execution, and also guarantees that all available memory accesses are made visible to the queue operations in the batch.

Semaphore wait operations for vkQueueSubmit additionally include all queue operations subsequently submitted via vkQueueSubmit in their half of a memory dependency, and all batches that are stored at a higher index in the same pSubmits array.

When queue execution reaches a semaphore wait operation, the queue will stall execution of queue operations in the batch until each semaphore becomes signaled. Once all semaphores are signaled, the semaphores will be reset to the unsignaled state, and subsequent queue operations will be permitted to execute.

Semaphore wait operations defined by vkQueueSubmit only wait at specific pipeline stages, rather than delaying all of each command buffer's execution, with the pipeline stages determined by the corresponding element of the <code>pWaitDstStageMask</code> member of VkSubmitInfo. Execution of work by those stages in subsequent commands is stalled until the corresponding semaphore reaches the signaled state.

Note

A common scenario for using pWaitDstStageMask with values other than $VK_PIPELINE_STAGE_ALL_COMMANDS_BIT$ is when synchronizing a window system presentation operation against subsequent command buffers which render the next frame. In this case, a presentation image must not be overwritten until the presentation operation completes, but other pipeline stages can execute without waiting. A mask of $VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT$ prevents subsequent color attachment writes from executing until the semaphore signals. Some implementations may be able to execute transfer operations and/or vertex processing work before the semaphore is signaled.

If an image layout transition needs to be performed on a swapchain image before it is used in a framebuffer, that can be performed as the first operation submitted to the queue after acquiring the image, and should not prevent other work from overlapping with the presentation operation. For example, a VkImageMemoryBarrier could use:

srcStageMask = VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT



- srcAccessMask = VK_ACCESS_MEMORY_READ_BIT
- dstStageMask = VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT
- dstaccessMask = VK_ACCESS_COLOR_ATTACHMENT_READ_BIT | VK_ACCESS_COLOR_ATTACH MENT WRITE BIT.
- oldLayout = VK_IMAGE_LAYOUT_PRESENT_SRC_KHR
- newLayout = VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL

Alternatively, oldLayout can be VK_IMAGE_LAYOUT_UNDEFINED, if the image's contents need not be preserved.

This barrier accomplishes a dependency chain between previous presentation operations and subsequent color attachment output operations, with the layout transition performed in between, and does not introduce a dependency between previous work and any vertex processing stages. More precisely, the semaphore signals after the presentation operation completes, then the semaphore wait stalls the VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT stage, then there is a dependency from that same stage to itself with the layout transition performed in between.

(The primary use case for this example is with the presentation extensions, thus the VK_IMAGE_LAYOUT_ PRESENT_SRC_KHR token is used even though it is not defined in the core Vulkan specification.)

6.3 Events

Events represent a fine-grained synchronization primitive that can be used to gauge progress through a sequence of commands executed on a queue by Vulkan. An event is initially in the unsignaled state. It can be signaled by a device, using commands inserted into the command buffer, or by the host. It can also be reset to the unsignaled state by a device or the host. The host can query the state of an event. A device can wait for one or more events to become signaled.

Events are represented by VkEvent handles:

VK_DEFINE NON_DISPATCHABLE_HANDLE(VkEvent)

To create an event, call:

```
const VkAllocationCallbacks* pAllocator,
VkEvent* pEvent);
```

- device is the logical device that creates the event.
- pCreateInfo is a pointer to an instance of the VkEventCreateInfo structure which contains information about how the event is to be created.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pEvent points to a handle in which the resulting event object is returned.

When created, the event object is in the unsignaled state.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkEventCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pEvent must be a pointer to a VkEvent handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The $\mbox{VkEventCreateInfo}$ structure is defined as:

```
typedef struct VkEventCreateInfo {
    VkStructureType     sType;
    const void*     pNext;
    VkEventCreateFlags    flags;
} VkEventCreateInfo;
```

• flags is reserved for future use.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_EVENT_CREATE_INFO
- pNext must be NULL
- flags must be 0

To destroy an event, call:

- device is the logical device that destroys the event.
- event is the handle of the event to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If event is not VK_NULL_HANDLE, event must be a valid VkEvent handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If event is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to event must have completed execution
- If VkAllocationCallbacks were provided when <code>event</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when event was created, pAllocator must be NULL

Host Synchronization

Host access to event must be externally synchronized

To query the state of an event from the host, call:

- device is the logical device that owns the event.
- event is the handle of the event to query.

Upon success, **vkGetEventStatus** returns the state of the event object with the following return codes:

Table 6.1: Event Object Status Codes

Status	Meaning
VK_EVENT_SET	The event specified by event is signaled.
VK_EVENT_RESET	The event specified by event is unsignaled.

If a **vkCmdSetEvent** or **vkCmdResetEvent** command is pending execution, then the value returned by this command may immediately be out of date.

The state of an event can be updated by the host. The state of the event is immediately changed, and subsequent calls to **vkGetEventStatus** will return the new state. If an event is already in the requested state, then updating it to the same state has no effect.

- device must be a valid VkDevice handle
- event must be a valid VkEvent handle
- event must have been created, allocated, or retrieved from device

Success

- VK_EVENT_SET
- VK_EVENT_RESET

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

To set the state of an event to signaled from the host, call:

- device is the logical device that owns the event.
- event is the event to set.

Valid Usage

- device must be a valid VkDevice handle
- event must be a valid VkEvent handle
- event must have been created, allocated, or retrieved from device

Host Synchronization

• Host access to event must be externally synchronized

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To set the state of an event to unsignaled from the host, call:

- device is the logical device that owns the event.
- event is the event to reset.

Valid Usage

- device must be a valid VkDevice handle
- event must be a valid VkEvent handle
- event must have been created, allocated, or retrieved from device
- event must not be waited on by a vkCmdWaitEvents command that is currently executing

Host Synchronization

• Host access to event must be externally synchronized

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The state of an event can also be updated on the device by commands inserted in command buffers. To set the state of an event to signaled from a device, call:

- commandBuffer is the command buffer into which the command is recorded.
- event is the event that will be signaled.
- stageMask specifies the pipeline stage at which the state of event is updated as described below.

- commandBuffer must be a valid VkCommandBuffer handle
- event must be a valid VkEvent handle
- stageMask must be a valid combination of VkPipelineStageFlagBits values
- stageMask must not be 0
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance
- Both of commandBuffer, and event must have been created, allocated, or retrieved from the same VkDevice
- If the geometry shaders feature is not enabled, <code>stageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>stageMask</code> must not contain <code>VK_PIPELINE_STAGE_</code>
 <code>TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_</code>
 <code>SHADER_BIT</code>

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		COMPUTE

To set the state of an event to unsignaled from a device, call:

- commandBuffer is the command buffer into which the command is recorded.
- event is the event that will be reset.
- stageMask specifies the pipeline stage at which the state of event is updated as described below.

- commandBuffer must be a valid VkCommandBuffer handle
- event must be a valid VkEvent handle
- stageMask must be a valid combination of VkPipelineStageFlagBits values
- stageMask must not be 0
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance

- Both of commandBuffer, and event must have been created, allocated, or retrieved from the same VkDevice
- If the geometry shaders feature is not enabled, <code>stageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>stageMask</code> must not contain <code>VK_PIPELINE_STAGE_</code>
 <code>TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_</code>
 <code>SHADER_BIT</code>
- When this command executes, event must not be waited on by a vkCmdWaitEvents command that is currently executing

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope Supported Queue Types	
Primary	Outside	GRAPHICS
Secondary		COMPUTE

For both **vkCmdSetEvent** and **vkCmdResetEvent**, the status of *event* is updated once the pipeline stages specified by *stageMask* (see Section 6.5.2) have completed executing prior commands. The command modifying the event is passed through the pipeline bound to the command buffer at time of execution.

To wait for one or more events to enter the signaled state on a device, call:

```
void vkCmdWaitEvents(
   VkCommandBuffer
                                                 commandBuffer,
   uint32_t
                                                 eventCount,
   const VkEvent*
                                                 pEvents,
   VkPipelineStageFlags
                                                 srcStageMask,
   VkPipelineStageFlags
                                                 dstStageMask,
   uint32_t
                                                 memoryBarrierCount,
   const VkMemoryBarrier*
                                                 pMemoryBarriers,
   uint32_t
                                                 bufferMemoryBarrierCount,
                                                 pBufferMemoryBarriers,
   const VkBufferMemoryBarrier*
                                                 imageMemoryBarrierCount,
   uint32_t
    const VkImageMemoryBarrier*
                                                 pImageMemoryBarriers);
```

- commandBuffer is the command buffer into which the command is recorded.
- eventCount is the length of the pEvents array.
- pEvents is an array of event object handles to wait on.
- srcStageMask (see Section 6.5.2) is the bitwise OR of the pipeline stages used to signal the event object handles in pEvents.
- dstStageMask is the pipeline stages at which the wait will occur.
- pMemoryBarriers is a pointer to an array of memoryBarrierCount VkMemoryBarrier structures.
- pBufferMemoryBarriers is a pointer to an array of bufferMemoryBarrierCount VkBufferMemoryBarrier structures.
- pImageMemoryBarriers is a pointer to an array of imageMemoryBarrierCount VkImageMemoryBarrier structures. See Section 6.5.3 for more details about memory barriers.

vkCmdWaitEvents waits for events set by either **vkSetEvent** or **vkCmdSetEvent** to become signaled. Logically, it has three phases:

- 1. Wait at the pipeline stages specified by <code>dstStageMask</code> (see Section 6.5.2) until the <code>eventCount</code> event objects specified by <code>pEvents</code> become signaled. Implementations may wait for each event object to become signaled in sequence (starting with the first event object in <code>pEvents</code>, and ending with the last), or wait for all of the event objects to become signaled at the same time.
- 2. Execute the memory barriers specified by pMemoryBarriers, pBufferMemoryBarriers and pImageMemoryBarriers (see Section 6.5.3).
- 3. Resume execution of pipeline stages specified by dstStageMask

Implementations may not execute commands in a pipelined manner, so **vkCmdWaitEvents** may not observe the results of a subsequent **vkCmdSetEvent** or **vkCmdResetEvent** command, even if the stages in *dstStageMask* occur after the stages in *srcStageMask*.

Commands that update the state of events in different pipeline stages may execute out of order, unless the ordering is enforced by execution dependencies.



Note

Applications should be careful to avoid race conditions when using events. For example, an event should only be reset if no **vkCmdWaitEvents** command is executing that waits upon that event.

- commandBuffer must be a valid VkCommandBuffer handle
- pEvents must be a pointer to an array of eventCount valid VkEvent handles

- srcStageMask must be a valid combination of VkPipelineStageFlagBits values
- srcStageMask must not be 0
- dstStageMask must be a valid combination of VkPipelineStageFlagBits values
- dstStageMask must not be 0
- If memoryBarrierCount is not 0, pMemoryBarriers must be a pointer to an array of memoryBarrierCount valid VkMemoryBarrier structures
- If bufferMemoryBarrierCount is not 0, pBufferMemoryBarriers must be a pointer to an array of bufferMemoryBarrierCount valid VkBufferMemoryBarrier structures
- If imageMemoryBarrierCount is not 0, pImageMemoryBarriers must be a pointer to an array of imageMemoryBarrierCount valid VkImageMemoryBarrier structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- eventCount must be greater than 0
- Both of commandBuffer, and the elements of pEvents must have been created, allocated, or retrieved from the same VkDevice
- srcStageMask must be the bitwise OR of the stageMask parameter used in previous calls to **vkCmdSetEvent** with any of the members of pEvents and VK_PIPELINE_STAGE_HOST_BIT if any of the members of pEvents was set using **vkSetEvent**
- If the geometry shaders feature is not enabled, <code>srcStageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the geometry shaders feature is not enabled, <code>dstStageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>srcStageMask</code> must not contain <code>VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>dstStageMask</code> must not contain <code>VK_PIPELINE_STAGE_</code> <code>TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT</code>
- If pEvents includes one or more events that will be signaled by **vkSetEvent** after commandBuffer has been submitted to a queue, then **vkCmdWaitEvents** must not be called inside a render pass instance

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope Supported Queue Types	
Primary	Both	GRAPHICS
Secondary		COMPUTE

An act of setting or resetting an event in one queue may not affect or be visible to other queues. For cross-queue synchronization, semaphores can be used.

6.4 Execution And Memory Dependencies

Synchronization commands introduce explicit execution and memory dependencies between two sets of action commands, where the second set of commands depends on the first set of commands. The two sets can be:

- First set: commands before a vkCmdSetEvent command.

 Second set: commands after a vkCmdWaitEvents command in the same queue, using the same event.
- First set: commands in a lower numbered subpass (or before a render pass instance).
 Second set: commands in a higher numbered subpass (or after a render pass instance), where there is a subpass dependency between the two subpasses (or between a subpass and VK_SUBPASS_EXTERNAL).
- First set: commands before a pipeline barrier.

 Second set: commands after that pipeline barrier in the same queue (possibly limited to within the same subpass).

An execution dependency is a single dependency between a set of source and destination pipeline stages, which guarantees that all work performed by the set of pipeline stages included in <code>srcStageMask</code> (see Pipeline Stage Flags) of the first set of commands completes before any work performed by the set of pipeline stages included in <code>dstStageMask</code> of the second set of commands begins.

An *execution dependency chain* from a set of source pipeline stages *A* to a set of destination pipeline stages *B* is a sequence of execution dependencies submitted to a queue in order between a first set of commands and a second set of commands, satisfying the following conditions:

- the first dependency includes A or VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT or VK_PIPELINE_STAGE_ALL COMMANDS BIT in the srcStageMask. And,
- the final dependency includes *B* or VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT or VK_PIPELINE_STAGE_ALL_COMMANDS_BIT in the *dstStageMask*. And,
- for each dependency in the sequence (except the first) at least one of the following conditions is true:
 - *srcStageMask* of the current dependency includes at least one bit *C* that is present in the *dstStageMask* of the previous dependency. Or,
 - srcStageMask of the current dependency includes VK_PIPELINE_STAGE_ALL_COMMANDS_BIT or VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT. Or,

- dstStageMask of the previous dependency includes VK_PIPELINE_STAGE_ALL_COMMANDS_BIT or VK_PIPELINE STAGE TOP OF PIPE BIT. Or,
- srcStageMask of the current dependency includes VK_PIPELINE_STAGE_ALL_GRAPHICS_BIT, and dstStageMask of the previous dependency includes at least one graphics pipeline stage. Or,
- dstStageMask of the previous dependency includes VK_PIPELINE_STAGE_ALL_GRAPHICS_BIT, and srcStageMask of the current dependency includes at least one graphics pipeline stage.
- for each dependency in the sequence (except the first), at least one of the following conditions is true:
 - the current dependency is a vkCmdSetEvent/vkCmdWaitEvents pair (where the vkCmdWaitEvents may be inside or outside a render pass instance), or a vkCmdPipelineBarrier outside of a render pass instance, or a subpass dependency with srcSubpass equal to VK_SUBPASS_EXTERNAL for a render pass instance that begins with a vkCmdBeginRenderPass command, and the previous dependency is any of:
 - * a vkCmdSetEvent/vkCmdWaitEvents pair or a vkCmdPipelineBarrier, either one outside of a render pass instance, that precedes the current dependency in the queue execution order. Or,
 - * a subpass dependency, with dstSubpass equal to VK_SUBPASS_EXTERNAL, for a renderpass instance that was ended with a vkCmdEndRenderPass command that precedes the current dependency in the queue execution order.
 - the current dependency is a subpass dependency for a render pass instance, and the previous dependency is any of:
 - * another dependency for the same render pass instance, with a *dstSubpass* equal to the *srcSubpass* of the current dependency. Or,
 - * a **vkCmdPipelineBarrier** of the same render pass instance, recorded for the subpass indicated by the *srcSubpass* of the current dependency. Or,
 - * a vkCmdSetEvent/vkCmdWaitEvents pair, where the vkCmdWaitEvents is inside the same render pass instance, recorded for the subpass indicated by the srcSubpass of the current dependency.
 - the current dependency is a vkCmdPipelineBarrier inside a subpass of a render pass instance, and the previous dependency is any of:
 - * a subpass dependency for the same render pass instance, with a dstSubpass equal to the subpass of the vkCmdPipelineBarrier. Or,
 - * a **vkCmdPipelineBarrier** of the same render pass instance, recorded for the same subpass, before the current dependency. Or,
 - * a **vkCmdSetEvent/vkCmdWaitEvents** pair, where the **vkCmdWaitEvents** is inside the same render pass instance, recorded for the same subpass, before the current dependency.

A pair of consecutive execution dependencies in an execution dependency chain accomplishes a dependency between the stages A and B via intermediate stages C, even if no work is executed between them that uses the pipeline stages included in C.

An execution dependency chain guarantees that the work performed by the pipeline stages A in the first set of commands completes before the work performed by pipeline stages B in the second set of commands begins.

A command C_1 is said to *happen-before* an execution dependency D_2 for a pipeline stage S if all the following conditions are true:

- C_1 is in the first set of commands for an execution dependency D_1 that includes S in its srcStageMask. And,
- there exists an execution dependency chain that includes D_1 and D_2 , where D_2 follows D_1 in the execution dependency sequence.

Similarly, a command C_2 is said to *happen-after* an execution dependency D_1 for a pipeline stage S if all the following conditions are true:

- C₂ is in the second set of commands for an execution dependency D₂ that includes S in its dstStageMask. And,
- there exists an execution dependency chain that includes D_1 and D_2 , where D_2 follows D_1 in the execution dependency sequence.

An execution dependency is by-region if its dependencyFlags parameter includes VK_DEPENDENCY_BY_REGION_BIT. Such a barrier describes a per-region (x,y,layer) dependency. That is, for each region, the implementation must ensure that the source stages for the first set of commands complete execution before any destination stages begin execution in the second set of commands for the same region. Since fragment shader invocations are not specified to run in any particular groupings, the size of a region is implementation-dependent, not known to the application, and must be assumed to be no larger than a single pixel. If dependencyFlags does not include VK_DEPENDENCY_BY_REGION_BIT, it describes a global dependency, that is for all pixel regions, the source stages must have completed for preceding commands before any destination stages starts for subsequent commands.

Memory dependencies are coupled to execution dependencies, and synchronize accesses to memory between two sets of commands. They operate according to two "halves" of a dependency to synchronize two sets of commands, the commands that happen-before the execution dependency for the <code>srcStageMask</code> vs the commands that happen-after the execution dependency for the <code>dstStageMask</code>, as described above. The first half of the dependency makes memory accesses using the set of access types in <code>srcAccessMask</code> performed in pipeline stages in <code>srcStageMask</code> by the first set of commands complete and writes be <code>available</code> for subsequent commands. The second half of the dependency makes any available writes from previous commands <code>visible</code> to pipeline stages in <code>dstStageMask</code> using the set of access types in <code>dstAccessMask</code> for the second set of commands, if those writes have been made available with the first half of the same or a previous dependency. The two halves of a memory dependency can either be expressed as part of a single command, or can be part of separate barriers as long as there is an execution dependency chain between them. The application must use memory dependencies to make writes visible before subsequent reads can rely on them, and before subsequent writes can overwrite them. Failure to do so causes the result of the reads to be undefined, and the order of writes to be undefined.

Global memory barriers apply to all resources owned by the device. Buffer and image memory barriers apply to the buffer range(s) or image subresource(s) included in the command. For accesses to a byte of a buffer or image subresource of an image to be synchronized between two sets of commands, the byte or image subresource must be included in both the first and second halves of the dependencies described above, but need not be included in each step of the execution dependency chain between them.

An execution dependency chain is *by-region* if all stages in all dependencies in the chain are framebuffer-space pipeline stages, and if the VK_DEPENDENCY_BY_REGION_BIT bit is included in all dependencies in the chain. Otherwise, the execution dependency chain is not by-region. The two halves of a memory dependency form a by-region dependency if **all** execution dependency chains between them are by-region. In other words, if there is any execution dependency between two sets of commands that is not by-region, then the memory dependency is not by-region.

When an image memory barrier includes a layout transition, the barrier first makes writes via <code>srcStageMask</code> and <code>srcAccessMask</code> available, then performs the layout transition, then makes the contents of the image subresource(s) in the new layout visible to memory accesses in <code>dstStageMask</code> and <code>dstAccessMask</code>, as if there is an execution and memory dependency between the source masks and the transition, as well as between the transition and the destination masks. Any writes that have previously been made available are included in the layout transition, but any previous writes that have not been made available may become lost or corrupt the image.

All dependencies must include at least one bit in each of the srcStageMask and dstStageMask.

Memory dependencies are used to solve data hazards, e.g. to ensure that write operations are visible to subsequent read operations (read-after-write hazard), as well as write-after-write hazards. Write-after-read and read-after-read hazards only require execution dependencies to synchronize.

6.5 Pipeline Barriers

A *pipeline barrier* inserts an execution dependency and a set of memory dependencies between a set of commands earlier in the command buffer and a set of commands later in the command buffer.

To record a pipeline barrier, call:

```
void vkCmdPipelineBarrier(
                                                 commandBuffer,
   VkCommandBuffer
   VkPipelineStageFlags
                                                 srcStageMask,
   VkPipelineStageFlags
                                                 dstStageMask,
   VkDependencyFlags
                                                 dependencyFlags,
   uint32_t
                                                 memoryBarrierCount,
   const VkMemoryBarrier*
                                                 pMemoryBarriers,
   uint32_t
                                                 bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier*
                                                 pBufferMemoryBarriers,
    uint32_t
                                                 imageMemoryBarrierCount,
    const VkImageMemoryBarrier*
                                                 pImageMemoryBarriers);
```

- commandBuffer is the command buffer into which the command is recorded.
- srcStageMask is a bitmask of VkPipelineStageFlagBits specifying a set of source pipeline stages (see Section 6.5.2).
- dstStageMask is a bitmask specifying a set of destination pipeline stages.

The pipeline barrier specifies an execution dependency such that all work performed by the set of pipeline stages included in <code>srcStageMask</code> of the first set of commands completes before any work performed by the set of pipeline stages included in <code>dstStageMask</code> of the second set of commands begins.

- dependencyFlags is a bitmask of VkDependencyFlagBits. The execution dependency is by-region if the mask includes VK DEPENDENCY BY REGION BIT.
- memoryBarrierCount is the length of the pMemoryBarriers array.
- pMemoryBarriers is a pointer to an array of VkMemoryBarrier structures.
- bufferMemoryBarrierCount is the length of the pBufferMemoryBarriers array.
- pBufferMemoryBarriers is a pointer to an array of VkBufferMemoryBarrier structures.
- imageMemoryBarrierCount is the length of the pImageMemoryBarriers array.
- pImageMemoryBarriers is a pointer to an array of VkImageMemoryBarrier structures.

Each element of the pMemoryBarriers, pBufferMemoryBarriers and pImageMemoryBarriers arrays specifies two halves of a memory dependency, as defined above. Specifics of each type of memory barrier and the memory access types are defined further in Memory Barriers.

If **vkCmdPipelineBarrier** is called outside a render pass instance, then the first set of commands is all prior commands submitted to the queue and recorded in the command buffer and the second set of commands is all subsequent commands recorded in the command buffer and submitted to the queue. If **vkCmdPipelineBarrier** is called inside a render pass instance, then the first set of commands is all prior commands in the same subpass and the second set of commands is all subsequent commands in the same subpass.

Valid Usage

• commandBuffer must be a valid VkCommandBuffer handle

- $\bullet \ \textit{srcStageMask} \ \textbf{must} \ \textbf{be} \ \textbf{a} \ \textbf{valid} \ \textbf{combination} \ \textbf{of} \ \texttt{VkPipelineStageFlagBits} \ \textbf{values}$
- srcStageMask must not be 0
- dstStageMask must be a valid combination of VkPipelineStageFlagBits values
- dstStageMask must not be 0
- dependencyFlags must be a valid combination of VkDependencyFlagBits values
- If memoryBarrierCount is not 0, pMemoryBarriers must be a pointer to an array of memoryBarrierCount valid VkMemoryBarrier structures
- If bufferMemoryBarrierCount is not 0, pBufferMemoryBarriers must be a pointer to an array of bufferMemoryBarrierCount valid VkBufferMemoryBarrier structures
- If imageMemoryBarrierCount is not 0, pImageMemoryBarriers must be a pointer to an array of imageMemoryBarrierCount valid VkImageMemoryBarrier structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support transfer, graphics, or compute operations
- If the geometry shaders feature is not enabled, <code>srcStageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the geometry shaders feature is not enabled, <code>dstStageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY</code> SHADER <code>BIT</code>
- If the tessellation shaders feature is not enabled, <code>srcStageMask</code> must not contain <code>VK_PIPELINE_STAGE_</code>
 <code>TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_</code>
 <code>SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>dstStageMask</code> must not contain <code>VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT</code>
- If **vkCmdPipelineBarrier** is called within a render pass instance, the render pass must have been created with a VkSubpassDependency instance in *pDependencies* that expresses a dependency from the current subpass to itself. Additionally:
 - srcStageMask must contain a subset of the bit values in the srcStageMask member of that instance of VkSubpassDependency
 - dstStageMask must contain a subset of the bit values in the dstStageMask member of that instance of VkSubpassDependency
 - The srcAccessMask of any element of pMemoryBarriers or pImageMemoryBarriers must contain a subset of the bit values the srcAccessMask member of that instance of VkSubpassDependency
 - The dstAccessMask of any element of pMemoryBarriers or pImageMemoryBarriers must contain a subset of the bit values the dstAccessMask member of that instance of VkSubpassDependency
 - dependencyFlags must be equal to the dependencyFlags member of that instance of VkSubpassDependency
- If vkCmdPipelineBarrier is called within a render pass instance, bufferMemoryBarrierCount must be

- If vkCmdPipelineBarrier is called within a render pass instance, the <code>image</code> member of any element of <code>pImageMemoryBarriers</code> must be equal to one of the elements of <code>pAttachments</code> that the current <code>framebuffer</code> was created with, that is also referred to by one of the elements of the <code>pColorAttachments</code>, <code>pResolveAttachments</code> or <code>pDepthStencilAttachment</code> members of the <code>VkSubpassDescription</code> instance that the current subpass was created with
- If vkCmdPipelineBarrier is called within a render pass instance, the <code>oldLayout</code> and <code>newLayout</code> members of any element of <code>pImageMemoryBarriers</code> must be equal to the <code>layout</code> member of an element of the <code>pColorAttachments</code>, <code>pResolveAttachments</code> or <code>pDepthStencilAttachment</code> members of the <code>VkSubpassDescription</code> instance that the current subpass was created with, that refers to the same <code>image</code>
- If **vkCmdPipelineBarrier** is called within a render pass instance, the *oldLayout* and *newLayout* members of an element of *plmageMemoryBarriers* must be equal
- If vkCmdPipelineBarrier is called within a render pass instance, the srcQueueFamilyIndex and dstQueueFamilyIndex members of any element of pImageMemoryBarriers must be VK_QUEUE_FAMILY_
 IGNORED

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope Supported Queue Types	
Primary	Both	TRANSFER
Secondary		GRAPHICS
		COMPUTE

6.5.1 Subpass Self-dependency

If **vkCmdPipelineBarrier** is called inside a render pass instance, the following restrictions apply. For a given subpass to allow a pipeline barrier, the render pass must declare a *self-dependency* from that subpass to itself. That is, there must exist a VkSubpassDependency in the subpass dependency list for the render pass with *srcSubpass* and *dstSubpass* equal to that subpass index. More than one self-dependency can be declared for each subpass. Self-dependencies must only include pipeline stage bits that are graphics stages. Self-dependencies must not have any earlier pipeline stages depend on any later pipeline stages. More precisely, this means that whatever is the last pipeline stage in *srcStageMask* must be no later than whatever is the first pipeline stage in *dstStageMask* (the latest source

stage can be equal to the earliest destination stage). If the source and destination stage masks both include framebuffer-space stages, then <code>dependencyFlags</code> must include <code>VK_DEPENDENCY_BY_REGION_BIT</code>.

A **vkCmdPipelineBarrier** command inside a render pass instance must be a *subset* of one of the self-dependencies of the subpass it is used in, meaning that the stage masks and access masks must each include only a subset of the bits of the corresponding mask in that self-dependency. If the self-dependency has VK_DEPENDENCY_BY_REGION_BIT set, then so must the pipeline barrier. Pipeline barriers within a render pass instance can only be types VkMemoryBarrier or VkImageMemoryBarrier is used, the image and image subresource range specified in the barrier must be a subset of one of the image views used by the framebuffer in the current subpass. Additionally, <code>oldLayout</code> must be equal to <code>newLayout</code>, and both the <code>srcQueueFamilyIndex</code> and <code>dstQueueFamilyIndex</code> must be VK_QUEUE_FAMILY_IGNORED.

6.5.2 Pipeline Stage Flags

Several of the event commands, **vkCmdPipelineBarrier**, and VkSubpassDependency depend on being able to specify where in the logical pipeline events can be signaled, or the source and destination of an execution dependency. These pipeline stages are specified using a bitmask:

```
typedef enum VkPipelineStageFlagBits {
   VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT = 0x00000001,
   VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT = 0x00000002,
   VK_PIPELINE_STAGE_VERTEX_INPUT_BIT = 0x00000004,
   VK_PIPELINE_STAGE_VERTEX_SHADER_BIT = 0x00000008,
   VK PIPELINE STAGE TESSELLATION CONTROL SHADER BIT = 0x00000010,
   VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT = 0x00000020,
   VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT = 0x00000040,
   VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT = 0x00000080,
   VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT = 0x00000100,
   VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT = 0x00000200,
   VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT = 0x00000400,
   VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT = 0x00000800,
   VK_PIPELINE_STAGE_TRANSFER_BIT = 0x00001000,
   VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT = 0x00002000,
   VK_PIPELINE_STAGE_HOST_BIT = 0x00004000,
   VK PIPELINE STAGE ALL GRAPHICS BIT = 0x00008000,
   VK_PIPELINE_STAGE_ALL_COMMANDS_BIT = 0x00010000,
} VkPipelineStageFlagBits;
```

The meaning of each bit is:

- VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT: Stage of the pipeline where commands are initially received by the queue.
- VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT: Stage of the pipeline where Draw/DispatchIndirect data structures are consumed.
- VK_PIPELINE_STAGE_VERTEX_INPUT_BIT: Stage of the pipeline where vertex and index buffers are consumed.
- VK_PIPELINE_STAGE_VERTEX_SHADER_BIT: Vertex shader stage.
- VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT: Tessellation control shader stage.
- VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT: Tessellation evaluation shader stage.
- VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT: Geometry shader stage.

- VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT: Fragment shader stage.
- VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT: Stage of the pipeline where early fragment tests (depth and stencil tests before fragment shading) are performed.
- VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT: Stage of the pipeline where late fragment tests (depth and stencil tests after fragment shading) are performed.
- VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT: Stage of the pipeline after blending where the final color values are output from the pipeline. This stage also includes resolve operations that occur at the end of a subpass. Note that this does not necessarily indicate that the values have been committed to memory.
- VK_PIPELINE_STAGE_TRANSFER_BIT: Execution of copy commands. This includes the operations resulting
 from all transfer commands. The set of transfer commands comprises vkCmdCopyBuffer, vkCmdCopyImage,
 vkCmdBlitImage, vkCmdCopyBufferToImage, vkCmdCopyImageToBuffer, vkCmdUpdateBuffer,
 vkCmdFillBuffer, vkCmdClearColorImage, vkCmdClearDepthStencilImage,
 vkCmdResolveImage, and vkCmdCopyQueryPoolResults.
- VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT: Execution of a compute shader.
- VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT: Final stage in the pipeline where commands complete execution.
- VK_PIPELINE_STAGE_HOST_BIT: A pseudo-stage indicating execution on the host of reads/writes of device memory.
- VK PIPELINE STAGE ALL GRAPHICS BIT: Execution of all graphics pipeline stages.
- VK_PIPELINE_STAGE_ALL_COMMANDS_BIT: Execution of all stages supported on the queue.

Note



BIT differ from VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT in that they correspond to all (or all graphics) stages, rather than to a specific stage at the end of the pipeline. An execution dependency with only VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT in dstStageMask will not delay subsequent commands, while including either of the other two bits will. Similarly, when defining a memory dependency, if the stage mask(s) refer to all stages, then the indicated access types from all stages will be made available and/or visible, but using only VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT would not make any accesses available and/or visible because this stage does not access memory. The VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT is useful for accomplishing memory barriers and layout transitions when the next accesses will be done in a different queue or by a presentation engine; in these cases subsequent commands in the same queue do not need to wait, but the barrier or transition must complete before semaphores associated with the batch signal.

The VK_PIPELINE_STAGE_ALL_COMMANDS_BIT and VK_PIPELINE_STAGE_ALL_GRAPHICS_

Note



If an implementation is unable to update the state of an event at any specific stage of the pipeline, it may instead update the event at any logically later stage. For example, if an implementation is unable to signal an event immediately after vertex shader execution is complete, it may instead signal the event after color attachment output has completed. In the limit, an event may be signaled after all graphics stages complete. If an implementation is unable to wait on an event at any specific stage of the pipeline, it may instead wait on it at any logically earlier stage.

Similarly, if an implementation is unable to implement an execution dependency at specific stages of the pipeline, it may implement the dependency in a way where additional source pipeline stages complete and/or where additional destination pipeline stages' execution is blocked to satisfy the dependency.

If an implementation makes such a substitution, it must not affect the semantics of execution or memory dependencies or image and buffer memory barriers.

Certain pipeline stages are only available on queues that support a particular set of operations. The following table lists, for each pipeline stage flag, which queue capability flag must be supported by the queue. When multiple flags are enumerated in the second column of the table, it means that the pipeline stage is supported on the queue if it supports any of the listed capability flags. For further details on queue capabilities see Physical Device Enumeration and Queues.

Table 6.2: Supported pipeline stage flags

Pipeline stage flag	Required queue capability flag
VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT	None
VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT	VK_QUEUE_GRAPHICS_BIT or
	VK_QUEUE_COMPUTE_BIT
VK_PIPELINE_STAGE_VERTEX_INPUT_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_VERTEX_SHADER_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_	VK_QUEUE_GRAPHICS_BIT
BIT	
VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT	VK_QUEUE_COMPUTE_BIT
VK_PIPELINE_STAGE_TRANSFER_BIT	VK_QUEUE_GRAPHICS_BIT,
	VK_QUEUE_COMPUTE_BIT or
	VK_QUEUE_TRANSFER_BIT
VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT	None
VK_PIPELINE_STAGE_HOST_BIT	None
VK_PIPELINE_STAGE_ALL_GRAPHICS_BIT	VK_QUEUE_GRAPHICS_BIT
VK_PIPELINE_STAGE_ALL_COMMANDS_BIT	None

6.5.3 Memory Barriers

Memory barriers express the two halves of a memory dependency between an earlier set of memory accesses against a later set of memory accesses. Vulkan provides three types of memory barriers: global memory, buffer memory, and

image memory.

6.5.4 Global Memory Barriers

The global memory barrier type is specified with an instance of the VkMemoryBarrier structure. This type of barrier applies to memory accesses involving all memory objects that exist at the time of its execution.

The VkMemoryBarrier structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- srcAccessMask is a bitmask of the classes of memory accesses performed by the first set of commands that will participate in the dependency.
- dstAccessMask is a bitmask of the classes of memory accesses performed by the second set of commands that will participate in the dependency.

<code>srcAccessMask</code> and <code>dstAccessMask</code>, along with <code>srcStageMask</code> and <code>dstStageMask</code> from <code>vkCmdPipelineBarrier</code>, define the two halves of a memory dependency and an execution dependency. Memory accesses using the set of access types in <code>srcAccessMask</code> performed in pipeline stages in <code>srcStageMask</code> by the first set of commands must complete and be available to later commands. The side effects of the first set of commands will be visible to memory accesses using the set of access types in <code>dstAccessMask</code> performed in pipeline stages in <code>dstStageMask</code> by the second set of commands. If the barrier is by-region, these requirements only apply to invocations within the same framebuffer-space region, for pipeline stages that perform framebuffer-space work. The execution dependency guarantees that execution of work by the destination stages of the second set of commands will not begin until execution of work by the source stages of the first set of commands has completed.

A common type of memory dependency is to avoid a read-after-write hazard. In this case, the source access mask and stages will include writes from a particular stage, and the destination access mask and stages will indicate how those writes will be read in subsequent commands. However, barriers can also express write-after-read dependencies and write-after-write dependencies, and are even useful to express read-after-read dependencies across an image layout change.

Bits which can be set in VkMemoryBarrier::srcAccessMask and VkMemoryBarrier::dstAccessMask include:

```
typedef enum VkAccessFlagBits {
   VK_ACCESS_INDIRECT_COMMAND_READ_BIT = 0x00000001,
   VK_ACCESS_INDEX_READ_BIT = 0x00000002,
   VK_ACCESS_VERTEX_ATTRIBUTE_READ_BIT = 0x00000004,
   VK_ACCESS_UNIFORM_READ_BIT = 0x00000008,
   VK_ACCESS_INPUT_ATTACHMENT_READ_BIT = 0x00000010,
   VK_ACCESS_SHADER_READ_BIT = 0x00000020,
   VK_ACCESS_SHADER_WRITE_BIT = 0x00000040,
   VK_ACCESS_COLOR_ATTACHMENT_READ_BIT = 0x000000080,
   VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT = 0x000000100,
   VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_READ_BIT = 0x000000200,
```

```
VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT = 0x00000400,
VK_ACCESS_TRANSFER_READ_BIT = 0x00000800,
VK_ACCESS_TRANSFER_WRITE_BIT = 0x00001000,
VK_ACCESS_HOST_READ_BIT = 0x00002000,
VK_ACCESS_HOST_WRITE_BIT = 0x00004000,
VK_ACCESS_MEMORY_READ_BIT = 0x00008000,
VK_ACCESS_MEMORY_WRITE_BIT = 0x00010000,
} VkACCESS_MEMORY_WRITE_BIT = 0x00010000,
```

- VK_ACCESS_INDIRECT_COMMAND_READ_BIT indicates that the access is an indirect command structure read as part of an indirect drawing command.
- VK ACCESS INDEX READ BIT indicates that the access is an index buffer read.
- VK_ACCESS_VERTEX_ATTRIBUTE_READ_BIT indicates that the access is a read via the vertex input bindings.
- VK_ACCESS_UNIFORM_READ_BIT indicates that the access is a read via a uniform buffer or dynamic uniform buffer descriptor.
- VK_ACCESS_INPUT_ATTACHMENT_READ_BIT indicates that the access is a read via an input attachment descriptor.
- VK_ACCESS_SHADER_READ_BIT indicates that the access is a read from a shader via any other descriptor type.
- VK_ACCESS_SHADER_WRITE_BIT indicates that the access is a write or atomic from a shader via the same descriptor types as in VK_ACCESS_SHADER_READ_BIT.
- VK_ACCESS_COLOR_ATTACHMENT_READ_BIT indicates that the access is a read via a color attachment.
- VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT indicates that the access is a write via a color or resolve attachment.
- VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_READ_BIT indicates that the access is a read via a depth/stencil attachment.
- VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT indicates that the access is a write via a depth/stencil attachment.
- VK_ACCESS_TRANSFER_READ_BIT indicates that the access is a read from a transfer (copy, blit, resolve, etc.) operation. For the complete set of transfer operations, see VK_PIPELINE_STAGE_TRANSFER_BIT.
- VK_ACCESS_TRANSFER_WRITE_BIT indicates that the access is a write from a transfer (copy, blit, resolve, etc.) operation. For the complete set of transfer operations, see VK_PIPELINE_STAGE_TRANSFER_BIT.
- VK_ACCESS_HOST_READ_BIT indicates that the access is a read via the host.
- VK_ACCESS_HOST_WRITE_BIT indicates that the access is a write via the host.
- VK_ACCESS_MEMORY_READ_BIT indicates that the access is a read via a non-specific unit attached to the memory. This unit may be external to the Vulkan device or otherwise not part of the core Vulkan pipeline. When included in dstAccessMask, all writes using access types in srcAccessMask performed by pipeline stages in srcStageMask must be visible in memory.
- VK_ACCESS_MEMORY_WRITE_BIT indicates that the access is a write via a non-specific unit attached to the memory. This unit may be external to the Vulkan device or otherwise not part of the core Vulkan pipeline. When included in <code>srcAccessMask</code>, all access types in <code>dstAccessMask</code> from pipeline stages in <code>dstStageMask</code> will observe the side effects of commands that executed before the barrier. When included in <code>dstAccessMask</code> all writes using access types in <code>srcAccessMask</code> performed by pipeline stages in <code>srcStageMask</code> must be visible in memory.

Color attachment reads and writes are automatically (without memory or execution dependencies) coherent and ordered against themselves and each other for a given sample within a subpass of a render pass instance, executing in rasterization order. Similarly, depth/stencil attachment reads and writes are automatically coherent and ordered against themselves and each other in the same circumstances.

Shader reads and/or writes through two variables (in the same or different shader invocations) decorated with **Coherent** and which use the same image view or buffer view are automatically coherent with each other, but require execution dependencies if a specific order is desired. Similarly, shader atomic operations are coherent with each other and with **Coherent** variables. Non-**Coherent** shader memory accesses require memory dependencies for writes to be available and reads to be visible.

Certain memory access types are only supported on queues that support a particular set of operations. The following table lists, for each access flag, which queue capability flag must be supported by the queue. When multiple flags are enumerated in the second column of the table it means that the access type is supported on the queue if it supports any of the listed capability flags. For further details on queue capabilities see Physical Device Enumeration and Queues.

Table 6.3: Supported access flags

Access flag	Required queue capability flag
VK_ACCESS_INDIRECT_COMMAND_READ_BIT	VK_QUEUE_GRAPHICS_BIT or
	VK_QUEUE_COMPUTE_BIT
VK_ACCESS_INDEX_READ_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_VERTEX_ATTRIBUTE_READ_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_UNIFORM_READ_BIT	VK_QUEUE_GRAPHICS_BIT or
	VK_QUEUE_COMPUTE_BIT
VK_ACCESS_INPUT_ATTACHMENT_READ_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_SHADER_READ_BIT	VK_QUEUE_GRAPHICS_BIT or
	VK_QUEUE_COMPUTE_BIT
VK_ACCESS_SHADER_WRITE_BIT	VK_QUEUE_GRAPHICS_BIT or
	VK_QUEUE_COMPUTE_BIT
VK_ACCESS_COLOR_ATTACHMENT_READ_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_READ_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT	VK_QUEUE_GRAPHICS_BIT
VK_ACCESS_TRANSFER_READ_BIT	VK_QUEUE_GRAPHICS_BIT,
	VK_QUEUE_COMPUTE_BIT or
	VK_QUEUE_TRANSFER_BIT
VK_ACCESS_TRANSFER_WRITE_BIT	VK_QUEUE_GRAPHICS_BIT,
	VK_QUEUE_COMPUTE_BIT or
	VK_QUEUE_TRANSFER_BIT
VK_ACCESS_HOST_READ_BIT	None
VK_ACCESS_HOST_WRITE_BIT	None
VK_ACCESS_MEMORY_READ_BIT	None
VK_ACCESS_MEMORY_WRITE_BIT	None

Valid Usage

- sType must be VK STRUCTURE TYPE MEMORY BARRIER
- pNext must be NULL
- srcAccessMask must be a valid combination of VkAccessFlagBits values
- dstAccessMask must be a valid combination of VkAccessFlagBits values

6.5.5 Buffer Memory Barriers

The buffer memory barrier type is specified with an instance of the VkBufferMemoryBarrier structure. This type of barrier only applies to memory accesses involving a specific range of the specified buffer object. That is, a memory dependency formed from a buffer memory barrier is scoped to the specified range of the buffer. It is also used to transfer ownership of a buffer range from one queue family to another, as described in the Resource Sharing section.

The VkBufferMemoryBarrier structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- *srcAccessMask* is a bitmask of the classes of memory accesses performed by the first set of commands that will participate in the dependency.
- dstAccessMask is a bitmask of the classes of memory accesses performed by the second set of commands that will participate in the dependency.
- srcQueueFamilyIndex is the queue family that is relinquishing ownership of the range of buffer to another queue, or VK_QUEUE_FAMILY_IGNORED if there is no transfer of ownership.
- dstQueueFamilyIndex is the queue family that is acquiring ownership of the range of buffer from another queue, or VK_QUEUE_FAMILY_IGNORED if there is no transfer of ownership.
- buffer is a handle to the buffer whose backing memory is affected by the barrier.
- offset is an offset in bytes into the backing memory for buffer; this is relative to the base offset as bound to the buffer (see vkBindBufferMemory).
- size is a size in bytes of the affected area of backing memory for buffer, or VK_WHOLE_SIZE to use the range from offset to the end of the buffer.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_BUFFER_MEMORY_BARRIER
- pNext must be NULL
- srcAccessMask must be a valid combination of VkAccessFlagBits values
- dstAccessMask must be a valid combination of VkAccessFlagBits values
- buffer must be a valid VkBuffer handle
- offset must be less than the size of buffer
- If size is not equal to VK_WHOLE_SIZE, size must be greater than 0
- If size is not equal to VK_WHOLE_SIZE, size must be less than or equal to than the size of buffer minus offset
- If buffer was created with a sharing mode of VK_SHARING_MODE_CONCURRENT, srcQueueFamilyIndex and dstQueueFamilyIndex must both be VK_QUEUE_FAMILY_IGNORED
- If buffer was created with a sharing mode of VK_SHARING_MODE_EXCLUSIVE, srcQueueFamilyIndex and dstQueueFamilyIndex must either both be VK_QUEUE_FAMILY_IGNORED, or both be a valid queue family (see Section 4.3.1)
- If buffer was created with a sharing mode of VK_SHARING_MODE_EXCLUSIVE, and srcQueueFamilyIndex and dstQueueFamilyIndex are valid queue families, at least one of them must be the same as the family of the queue that will execute this barrier

6.5.6 Image Memory Barriers

The image memory barrier type is specified with an instance of the VkImageMemoryBarrier structure. This type of barrier only applies to memory accesses involving a specific image subresource range of the specified image object. That is, a memory dependency formed from an image memory barrier is scoped to the specified image subresources of the image. It is also used to perform a layout transition for an image subresource range, or to transfer ownership of an image subresource range from one queue family to another as described in the Resource Sharing section.

The VkImageMemoryBarrier structure is defined as:

```
typedef struct VkImageMemoryBarrier {
   VkStructureType
                             sType;
                             pNext;
   const void*
                             srcAccessMask;
   VkAccessFlags
   VkAccessFlags
                             dstAccessMask;
   VkImageLayout
                             oldLayout;
   VkImageLayout
                             newLayout;
   uint32_t
                             srcQueueFamilyIndex;
   uint32_t
                             dstQueueFamilyIndex;
   VkImage
                             image;
   VkImageSubresourceRange subresourceRange;
```

} VkImageMemoryBarrier;

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- srcAccessMask is a bitmask of the classes of memory accesses performed by the first set of commands that will participate in the dependency.
- dstAccessMask is a bitmask of the classes of memory accesses performed by the second set of commands that will
 participate in the dependency.
- oldLayout describes the current layout of the image subresource(s).
- newLayout describes the new layout of the image subresource(s).
- srcQueueFamilyIndex is the queue family that is relinquishing ownership of the image subresource(s) to another queue, or VK_QUEUE_FAMILY_IGNORED if there is no transfer of ownership).
- dstQueueFamilyIndex is the queue family that is acquiring ownership of the image subresource(s) from another queue, or VK_QUEUE_FAMILY_IGNORED if there is no transfer of ownership).
- *image* is a handle to the image whose backing memory is affected by the barrier.
- subresourceRange describes an area of the backing memory for image (see Section 11.5 for the description of VkImageSubresourceRange), as well as the set of image subresources whose image layouts are modified.

If oldLayout differs from newLayout, a layout transition occurs as part of the image memory barrier, affecting the data contained in the region of the image defined by the subresourceRange. If oldLayout is VK_IMAGE_LAYOUT_UNDEFINED, then the data is undefined after the layout transition. This may allow a more efficient transition, since the data may be discarded. The layout transition must occur after all operations using the old layout are completed and before all operations using the new layout are started. This is achieved by ensuring that there is a memory dependency between previous accesses and the layout transition, as well as between the layout transition and subsequent accesses, where the layout transition occurs between the two halves of a memory dependency in an image memory barrier.

Layout transitions that are performed via image memory barriers are automatically ordered against other layout transitions, including those that occur as part of a render pass instance.



Note

See Section 11.4 for details on available image layouts and their usages.

- sType must be VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER
- pNext must be NULL

- srcAccessMask must be a valid combination of VkAccessFlagBits values
- dstAccessMask must be a valid combination of VkAccessFlagBits values
- oldLayout must be a valid VkImageLayout value
- newLayout must be a valid VkImageLayout value
- image must be a valid VkImage handle
- subresourceRange must be a valid VkImageSubresourceRange structure
- oldLayout must be VK_IMAGE_LAYOUT_UNDEFINED or the current layout of the image subresources affected by the barrier
- newLayout must not be VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED
- If image was created with a sharing mode of VK_SHARING_MODE_CONCURRENT, srcQueueFamilyIndex and dstQueueFamilyIndex must both be VK_QUEUE_FAMILY_IGNORED
- If image was created with a sharing mode of VK_SHARING_MODE_EXCLUSIVE, <code>srcQueueFamilyIndex</code> and <code>dstQueueFamilyIndex</code> must either both be VK_QUEUE_FAMILY_IGNORED, or both be a valid queue family (see Section 4.3.1)
- If image was created with a sharing mode of VK_SHARING_MODE_EXCLUSIVE, and srcQueueFamilyIndex and dstQueueFamilyIndex are valid queue families, at least one of them must be the same as the family of the queue that will execute this barrier
- subresourceRange must be a valid image subresource range for the image (see Section 11.5)
- If image has a depth/stencil format with both depth and stencil components, then aspectMask member of subresourceRange must include both VK_IMAGE_ASPECT_DEPTH_BIT and VK_IMAGE_ASPECT_STENCIL_BIT
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL then image must have been created with VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT set
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL then image must have been created with VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT set
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL then image must have been created with VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT set
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL then image must have been created with VK_IMAGE_USAGE_SAMPLED_BIT or VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT set
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL then image must have been created with VK_IMAGE_USAGE_TRANSFER_SRC_BIT set
- If either oldLayout or newLayout is VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL then image must have been created with VK IMAGE USAGE TRANSFER DST BIT set

6.5.7 Wait Idle Operations

To wait on the host for the completion of outstanding queue operations for a given queue, call:

• queue is the queue on which to wait.

vkQueueWaitIdle is equivalent to submitting a fence to a queue and waiting with an infinite timeout for that fence to signal.

Valid Usage

• queue must be a valid VkQueue handle

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types	
-	-	Any	

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

To wait on the host for the completion of outstanding queue operations for all queues on a given logical device, call:

 • device is the logical device to idle.

vkDeviceWaitIdle is equivalent to calling vkQueueWaitIdle for all queues owned by device.

Valid Usage

• device must be a valid VkDevice handle

Host Synchronization

• Host access to all VkQueue objects created from device must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

6.6 Host Write Ordering Guarantees

When submitting batches of command buffers to a queue via vkQueueSubmit, it is guaranteed that:

Host writes to mappable device memory that occurred before the call to vkQueueSubmit are visible to the queue
operation resulting from that submission, if the device memory is coherent or if the memory range was flushed with
vkFlushMappedMemoryRanges.

Chapter 7

Render Pass

A *render pass* represents a collection of attachments, subpasses, and dependencies between the subpasses, and describes how the attachments are used over the course of the subpasses. The use of a render pass in a command buffer is a *render pass instance*.

Render passes are represented by VkRenderPass handles:

VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkRenderPass)

An *attachment description* describes the properties of an attachment including its format, sample count, and how its contents are treated at the beginning and end of each render pass instance.

A *subpass* represents a phase of rendering that reads and writes a subset of the attachments in a render pass. Rendering commands are recorded into a particular subpass of a render pass instance.

A *subpass description* describes the subset of attachments that is involved in the execution of a subpass. Each subpass can read from some attachments as *input attachments*, write to some as *color attachments* or *depth/stencil attachments*, and do resolve operations to others as *resolve attachments*. A subpass description can also include a set of *preserve attachments*, which are attachments that are not read or written by the subpass but whose contents must be preserved throughout the subpass.

A subpass *uses* an attachment if the attachment is a color, depth/stencil, resolve, or input attachment for that subpass. A subpass does not use an attachment if that attachment is preserved by the subpass. The first use of an attachment is in the lowest numbered subpass that uses that attachment. Similarly, the last use of an attachment is in the highest numbered subpass that uses that attachment.

The subpasses in a render pass all render to the same dimensions, and fragments for pixel (x,y,layer) in one subpass can only read attachment contents written by previous subpasses at that same (x,y,layer) location.

Note



By describing a complete set of subpasses in advance, render passes provide the implementation an opportunity to optimize the storage and transfer of attachment data between subpasses.

In practice, this means that subpasses with a simple framebuffer-space dependency may be merged into a single tiled rendering pass, keeping the attachment data on-chip for the duration of a render pass instance. However, it is also guite common for a render pass to only contain a single subpass.

Subpass dependencies describe ordering restrictions between pairs of subpasses. If no dependencies are specified, implementations may reorder or overlap portions (e.g., certain shader stages) of the execution of subpasses.

Dependencies limit the extent of overlap or reordering, and are defined using masks of pipeline stages and memory access types. Each dependency acts as an execution and memory dependency, similarly to how pipeline barriers are defined. Dependencies are needed if two subpasses operate on attachments with overlapping ranges of the same VkDeviceMemory object and at least one subpass writes to that range.

A *subpass dependency chain* is a sequence of subpass dependencies in a render pass, where the source subpass of each subpass dependency (after the first) equals the destination subpass of the previous dependency.

A render pass describes the structure of subpasses and attachments independent of any specific image views for the attachments. The specific image views that will be used for the attachments, and their dimensions, are specified in VkFramebuffer objects. Framebuffers are created with respect to a specific render pass that the framebuffer is compatible with (see Render Pass Compatibility). Collectively, a render pass and a framebuffer define the complete render target state for one or more subpasses as well as the algorithmic dependencies between the subpasses.

The various pipeline stages of the drawing commands for a given subpass may execute concurrently and/or out of order, both within and across drawing commands. However for a given (x,y,layer,sample) sample location, certain per-sample operations are performed in rasterization order.

7.1 Render Pass Creation

To create a render pass, call:

- device is the logical device that creates the render pass.
- pCreateInfo is a pointer to an instance of the VkRenderPassCreateInfo structure that describes the parameters of the render pass.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pRenderPass points to a VkRenderPass handle in which the resulting render pass object is returned.

- device must be a valid VkDevice handle
- \bullet pCreateInfo must be a pointer to a valid <code>VkRenderPassCreateInfo</code> structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pRenderPass must be a pointer to a VkRenderPass handle

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkRenderPassCreateInfo structure is defined as:

```
typedef struct VkRenderPassCreateInfo {
   VkStructureType
                                    sType;
   const void*
                                   pNext;
                                 flags;
   VkRenderPassCreateFlags
   uint32_t
                                  attachmentCount;
   const VkAttachmentDescription* pAttachments;
   uint32_t
                                   subpassCount;
   const VkSubpassDescription* pSubpasses;
   uint32_t
                                   dependencyCount;
                             pDependencies;
   const VkSubpassDependency*
} VkRenderPassCreateInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- attachmentCount is the number of attachments used by this render pass, or zero indicating no attachments. Attachments are referred to by zero-based indices in the range [0,attachmentCount).
- pAttachments points to an array of attachmentCount number of VkAttachmentDescription structures describing properties of the attachments, or NULL if attachmentCount is zero.
- subpassCount is the number of subpasses to create for this render pass. Subpasses are referred to by zero-based indices in the range [0,subpassCount). A render pass must have at least one subpass.
- pSubpasses points to an array of subpassCount number of VkSubpassDescription structures describing properties of the subpasses.
- dependencyCount is the number of dependencies between pairs of subpasses, or zero indicating no dependencies.
- pDependencies points to an array of dependencyCount number of VkSubpassDependency structures describing dependencies between pairs of subpasses, or NULL if dependencyCount is zero.

Valid Usage

- sType must be VK STRUCTURE TYPE RENDER PASS CREATE INFO
- pNext must be NULL
- flags must be 0
- If attachmentCount is not 0, pAttachments must be a pointer to an array of attachmentCount valid VkAttachmentDescription structures
- pSubpasses must be a pointer to an array of subpassCount valid VkSubpassDescription structures
- If dependencyCount is not 0, pDependencies must be a pointer to an array of dependencyCount valid VkSubpassDependency structures
- subpassCount must be greater than 0
- If any two subpasses operate on attachments with overlapping ranges of the same VkDeviceMemory object, and at least one subpass writes to that area of VkDeviceMemory, a subpass dependency must be included (either directly or via some intermediate subpasses) between them
- If the attachment member of any element of pInputAttachments, pColorAttachments, pResolveAttachments or pDepthStencilAttachment, or the attachment indexed by any element of pPreserveAttachments in any given element of pSubpasses is bound to a range of a VkDeviceMemory object that overlaps with any other attachment in any subpass (including the same subpass), the VkAttachmentDescription structures describing them must include VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT in flags
- If the attachment member of any element of pInputAttachments, pColorAttachments, pResolveAttachments or pDepthStencilAttachment, or any element of pPreserveAttachments in any given element of pSubpasses is not VK_ATTACHMENT_UNUSED, it must be less than attachmentCount
- The value of any element of the pPreserveAttachments member in any given element of pSubpasses must not be VK ATTACHMENT UNUSED

The VkAttachmentDescription structure is defined as:

```
typedef struct VkAttachmentDescription {
   VkAttachmentDescriptionFlags flags;
   VkFormat
                                format;
   VkSampleCountFlagBits
                                samples;
   VkAttachmentLoadOp
                                loadOp;
   VkAttachmentStoreOp
                                storeOp;
   VkAttachmentLoadOp
                                stencilLoadOp;
   VkAttachmentStoreOp
                                 stencilStoreOp;
   VkImageLayout
                                  initialLayout;
   VkImageLayout
                                  finalLayout;
} VkAttachmentDescription;
```

• flags is a bitmask describing additional properties of the attachment. Bits which can be set include:

```
typedef enum VkAttachmentDescriptionFlagBits {
    VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT = 0x00000001,
```

```
} VkAttachmentDescriptionFlagBits;
```

- format is a VkFormat value specifying the format of the image that will be used for the attachment.
- samples is the number of samples of the image as defined in VkSampleCountFlagBits.
- loadOp specifies how the contents of color and depth components of the attachment are treated at the beginning of the subpass where it is first used:

```
typedef enum VkAttachmentLoadOp {
   VK_ATTACHMENT_LOAD_OP_LOAD = 0,
   VK_ATTACHMENT_LOAD_OP_CLEAR = 1,
   VK_ATTACHMENT_LOAD_OP_DONT_CARE = 2,
} VkAttachmentLoadOp;
```

- VK_ATTACHMENT_LOAD_OP_LOAD means the contents within the render area will be preserved.
- VK_ATTACHMENT_LOAD_OP_CLEAR means the contents within the render area will be cleared to a uniform value, which is specified when a render pass instance is begun.
- VK_ATTACHMENT_LOAD_OP_DONT_CARE means the contents within the area need not be preserved; the
 contents of the attachment will be undefined inside the render area.
- storeOp specifies how the contents of color and depth components of the attachment are treated at the end of the subpass where it is last used:

```
typedef enum VkAttachmentStoreOp {
    VK_ATTACHMENT_STORE_OP_STORE = 0,
    VK_ATTACHMENT_STORE_OP_DONT_CARE = 1,
} VkAttachmentStoreOp;
```

- VK_ATTACHMENT_STORE_OP_STORE means the contents within the render area are written to memory and will be available for reading after the render pass instance completes once the writes have been synchronized with VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT (for color attachments) or VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT (for depth/stencil attachments).
- VK_ATTACHMENT_STORE_OP_DONT_CARE means the contents within the render area are not needed after rendering, and may be discarded; the contents of the attachment will be undefined inside the render area.
- stencilLoadOp specifies how the contents of stencil components of the attachment are treated at the beginning of the subpass where it is first used, and must be one of the same values allowed for loadOp above.
- stencilStoreOp specifies how the contents of stencil components of the attachment are treated at the end of the last subpass where it is used, and must be one of the same values allowed for storeOp above.
- initialLayout is the layout the attachment image subresource will be in when a render pass instance begins.
- finalLayout is the layout the attachment image subresource will be transitioned to when a render pass instance ends. During a render pass instance, an attachment can use a different layout in each subpass, if desired.

If the attachment uses a color format, then <code>loadOp</code> and <code>storeOp</code> are used, and <code>stencilLoadOp</code> and <code>stencilStoreOp</code> are ignored. If the format has depth and/or stencil components, <code>loadOp</code> and <code>storeOp</code> apply only to the depth data, while <code>stencilLoadOp</code> and <code>stencilStoreOp</code> define how the stencil data is handled.

During a render pass instance, input/color attachments with color formats that have a component size of 8, 16, or 32 bits must be represented in the attachment's format throughout the instance. Attachments with other floating- or fixed-point color formats, or with depth components may be represented in a format with a precision higher than the attachment

format, but must be represented with the same range. When such a component is loaded via the <code>loadOp</code>, it will be converted into an implementation-dependent format used by the render pass. Such components must be converted from the render pass format, to the format of the attachment, before they are stored or resolved at the end of a render pass instance via <code>storeOp</code>. Conversions occur as described in Numeric Representation and Computation and Fixed-Point Data Conversions.

If flags includes VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT, then the attachment is treated as if it shares physical memory with another attachment in the same render pass. This information limits the ability of the implementation to reorder certain operations (like layout transitions and the loadOp) such that it is not improperly reordered against other uses of the same physical memory via a different attachment. This is described in more detail below.

Valid Usage

- flags must be a valid combination of VkAttachmentDescriptionFlagBits values
- format must be a valid VkFormat value
- samples must be a valid VkSampleCountFlagBits value
- loadOp must be a valid VkAttachmentLoadOp value
- storeOp must be a valid VkAttachmentStoreOp value
- stencilLoadOp must be a valid VkAttachmentLoadOp value
- stencilStoreOp must be a valid VkAttachmentStoreOp value
- initialLayout must be a valid VkImageLayout value
- finalLayout must be a valid VkImageLayout value
- finalLayout must not be VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED

If a render pass uses multiple attachments that alias the same device memory, those attachments must each include the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT bit in their attachment description flags. Attachments aliasing the same memory occurs in multiple ways:

- Multiple attachments being assigned the same image view as part of framebuffer creation.
- Attachments using distinct image views that correspond to the same image subresource of an image.
- Attachments using views of distinct image subresources which are bound to overlapping memory.

Render passes must include subpass dependencies (either directly or via a subpass dependency chain) between any two subpasses that operate on the same attachment or aliasing attachments and those subpass dependencies must include execution and memory dependencies separating uses of the aliases, if at least one of those subpasses writes to one of the aliases. Those dependencies must not include the VK_DEPENDENCY_BY_REGION_BIT if the aliases are views of distinct image subresources which overlap in memory.

Multiple attachments that alias the same memory must not be used in a single subpass. A given attachment index must not be used multiple times in a single subpass, with one exception: two subpass attachments can use the same attachment

index if at least one use is as an input attachment and neither use is as a resolve or preserve attachment. In other words, the same view can be used simultaneously as an input and color or depth/stencil attachment, but must not be used as multiple color or depth/stencil attachments nor as resolve or preserve attachments. This valid scenario is described in more detail below.

If a set of attachments alias each other, then all except the first to be used in the render pass must use an initialLayout of VK_IMAGE_LAYOUT_UNDEFINED, since the earlier uses of the other aliases make their contents undefined. Once an alias has been used and a different alias has been used after it, the first alias must not be used in any later subpasses. However, an application can assign the same image view to multiple aliasing attachment indices, which allows that image view to be used multiple times even if other aliases are used in between. Once an attachment needs the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT bit, there should be no additional cost of introducing additional aliases, and using these additional aliases may allow more efficient clearing of the attachments on multiple uses via VK_ATTACHMENT_LOAD_OP_CLEAR.



Note

The exact set of attachment indices that alias with each other is not known until a framebuffer is created using the render pass, so the above conditions cannot be validated at render pass creation time.

The VkSubpassDescription structure is defined as:

```
typedef struct VkSubpassDescription {
   VkSubpassDescriptionFlags
   VkPipelineBindPoint
                                    pipelineBindPoint;
                                    inputAttachmentCount;
   uint32 t
   const VkAttachmentReference*
                                    pInputAttachments;
   uint32_t
                                    colorAttachmentCount;
   const VkAttachmentReference*
                                    pColorAttachments;
    const VkAttachmentReference*
                                    pResolveAttachments;
   const VkAttachmentReference*
                                    pDepthStencilAttachment;
   uint32_t
                                    preserveAttachmentCount;
   const uint32 t*
                                    pPreserveAttachments;
} VkSubpassDescription;
```

- *flags* is reserved for future use.
- pipelineBindPoint is a VkPipelineBindPoint value specifying whether this is a compute or graphics subpass. Currently, only graphics subpasses are supported.
- inputAttachmentCount is the number of input attachments.
- pInputAttachments is an array of VkAttachmentReference structures (defined below) that lists which of the render pass's attachments can be read in the shader during the subpass, and what layout each attachment will be in during the subpass. Each element of the array corresponds to an input attachment unit number in the shader, i.e. if the shader declares an input variable layout (input_attachment_index=X, set=Y, binding=Z) then it uses the attachment provided in pInputAttachments[X]. Input attachments must also be bound to the pipeline with a descriptor set, with the input attachment descriptor written in the location (set=Y, binding=Z).
- colorAttachmentCount is the number of color attachments.
- pColorAttachments is an array of colorAttachmentCount VkAttachmentReference structures that lists which of the render pass's attachments will be used as color attachments in the subpass, and what layout each attachment will be in during the subpass. Each element of the array corresponds to a fragment shader output location, i.e. if the shader declared an output variable layout (location=X) then it uses the attachment provided in pColorAttachments[X].

- pResolveAttachments is NULL or an array of colorAttachmentCount VkAttachmentReference structures that lists which of the render pass's attachments are resolved to at the end of the subpass, and what layout each attachment will be in during the resolve. If pResolveAttachments is not NULL, each of its elements corresponds to a color attachment (the element in pColorAttachments at the same index). At the end of each subpass, the subpass's color attachments are resolved to corresponding resolve attachments, unless the resolve attachment index is VK_ATTACHMENT_UNUSED or pResolveAttachments is NULL. If the first use of an attachment in a render pass is as a resolve attachment, then the loadOp is effectively ignored as the resolve is guaranteed to overwrite all pixels in the render area.
- pDepthStencilAttachment is a pointer to a VkAttachmentReference specifying which attachment will be used for depth/stencil data and the layout it will be in during the subpass. Setting the attachment index to VK_ ATTACHMENT_UNUSED or leaving this pointer as NULL indicates that no depth/stencil attachment will be used in the subpass.
- preserveAttachmentCount is the number of preserved attachments.
- pPreserveAttachments is an array of preserveAttachmentCount render pass attachment indices describing the attachments that are not used by a subpass, but whose contents must be preserved throughout the subpass.

The contents of an attachment within the render area become undefined at the start of a subpass S if all of the following conditions are true:

- The attachment is used as a color, depth/stencil, or resolve attachment in any subpass in the render pass.
- There is a subpass S1 that uses or preserves the attachment, and a subpass dependency from S1 to S.
- The attachment is not used or preserved in subpass S.

Once the contents of an attachment become undefined in subpass S, they remain undefined for subpasses in subpass dependency chains starting with subpass S until they are written again. However, they remain valid for subpasses in other subpass dependency chains starting with subpass S1 if those subpasses use or preserve the attachment.

- flags must be 0
- pipelineBindPoint must be a valid VkPipelineBindPoint value
- If inputAttachmentCount is not 0, pInputAttachments must be a pointer to an array of inputAttachmentCount valid VkAttachmentReference structures
- If colorAttachmentCount is not 0, pColorAttachments must be a pointer to an array of colorAttachmentCount valid VkAttachmentReference structures
- If colorAttachmentCount is not 0, and pResolveAttachments is not NULL, pResolveAttachments must be a pointer to an array of colorAttachmentCount valid VkAttachmentReference structures
- If pDepthStencilAttachment is not NULL, pDepthStencilAttachment must be a pointer to a valid VkAttachmentReference structure
- If preserveAttachmentCount is not 0, pPreserveAttachments must be a pointer to an array of preserveAttachmentCount uint32_t values

- pipelineBindPoint must be VK_PIPELINE_BIND_POINT_GRAPHICS
- colorCount must be less than or equal to VkPhysicalDeviceLimits::maxColorAttachments
- If the first use of an attachment in this render pass is as an input attachment, and the attachment is not also used as a color or depth/stencil attachment in the same subpass, then <code>loadOp</code> must not be <code>VK_ATTACHMENT_LOAD_OP_CLEAR</code>
- If pResolveAttachments is not NULL, for each resolve attachment that does not have the value VK_ATTACHMENT_UNUSED, the corresponding color attachment must not have the value VK_ATTACHMENT_UNUSED
- If pResolveAttachments is not NULL, the sample count of each element of pColorAttachments must be anything other than VK_SAMPLE_COUNT_1_BIT
- Any given element of pResolveAttachments must have a sample count of VK SAMPLE COUNT 1 BIT
- Any given element of pResolveAttachments must have the same VkFormat as its corresponding color attachment
- All attachments in pColorAttachments and pDepthStencilAttachment that are not VK_ATTACHMENT_UNUSED must have the same sample count
- If any input attachments are VK_ATTACHMENT_UNUSED, then any pipelines bound during the subpass must not access those input attachments from the fragment shader
- The attachment member of any element of pPreserveAttachments must not be VK_ATTACHMENT_ UNUSED
- Any given element of pPreserveAttachments must not also be an element of any other member of the subpass description
- If any attachment is used as both an input attachment and a color or depth/stencil attachment, then each use must use the same <code>layout</code>

The VkAttachmentReference structure is defined as:

```
typedef struct VkAttachmentReference {
    uint32_t          attachment;
    VkImageLayout          layout;
} VkAttachmentReference;
```

- attachment is the index of the attachment of the render pass, and corresponds to the index of the corresponding element in the pAttachments array of the VkRenderPassCreateInfo structure. If any color or depth/stencil attachments are VK_ATTACHMENT_UNUSED, then no writes occur for those attachments.
- layout is a VkImageLayout value specifying the layout the attachment uses during the subpass. The implementation will automatically perform layout transitions as needed between subpasses to make each subpass use the requested layouts.

Valid Usage

- layout must be a valid VkImageLayout value
- layout must not be VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED

The VkSubpassDependency structure is defined as:

- srcSubpass and dstSubpass are the subpass indices of the producer and consumer subpasses, respectively. srcSubpass and dstSubpass can also have the special value VK_SUBPASS_EXTERNAL. The source subpass must always be a lower numbered subpass than the destination subpass (excluding external subpasses and self-dependencies), so that the order of subpass descriptions is a valid execution ordering, avoiding cycles in the dependency graph.
- srcStageMask, dstStageMask, srcAccessMask, dstAccessMask, and dependencyFlags describe an execution and memory dependency between subpasses. The bits that can be included in dependencyFlags are:

```
typedef enum VkDependencyFlagBits {
    VK_DEPENDENCY_BY_REGION_BIT = 0x00000001,
} VkDependencyFlagBits;
```

- If dependencyFlags contains VK_DEPENDENCY_BY_REGION_BIT, then the dependency is by-region as defined in Execution And Memory Dependencies.

Each subpass dependency defines an execution and memory dependency between two sets of commands, with the second set depending on the first set. When srcSubpass does not equal dstSubpass then the first set of commands is:

- All commands in the subpass indicated by srcSubpass, if srcSubpass is not VK_SUBPASS_EXTERNAL.
- All commands before the render pass instance, if srcSubpass is VK_SUBPASS_EXTERNAL.

While the corresponding second set of commands is:

- All commands in the subpass indicated by dstSubpass, if dstSubpass is not VK_SUBPASS_EXTERNAL.
- All commands after the render pass instance, if dstSubpass is VK_SUBPASS_EXTERNAL.

When <code>srcSubpass</code> equals <code>dstSubpass</code> then the first set consists of commands in the subpass before a call to <code>vkCmdPipelineBarrier</code> and the second set consists of commands in the subpass following that same call as described in the Subpass Self-dependency section.

The srcStageMask, dstStageMask, srcAccessMask, dstAccessMask, and dependencyFlags parameters of the dependency are interpreted the same way as for other dependencies, as described in Synchronization and Cache Control.

- srcStageMask must be a valid combination of VkPipelineStageFlagBits values
- srcStageMask must not be 0
- dstStageMask must be a valid combination of VkPipelineStageFlagBits values
- dstStageMask must not be 0
- srcAccessMask must be a valid combination of VkAccessFlagBits values
- dstAccessMask must be a valid combination of VkAccessFlagBits values
- dependencyFlags must be a valid combination of VkDependencyFlagBits values
- If the geometry shaders feature is not enabled, srcStageMask must not contain VK_PIPELINE_STAGE_ GEOMETRY SHADER BIT
- If the geometry shaders feature is not enabled, <code>dstStageMask</code> must not contain <code>VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, <code>srcStageMask</code> must not contain <code>VK_PIPELINE_STAGE_</code> <code>TESSELLATION_CONTROL_SHADER_BIT</code> or <code>VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_</code> <code>SHADER_BIT</code>
- If the tessellation shaders feature is not enabled, dstStageMask must not contain VK_PIPELINE_STAGE_
 TESSELLATION_CONTROL_SHADER_BIT or VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_
 SHADER_BIT
- srcSubpass must be less than or equal to dstSubpass, unless one of them is VK_SUBPASS_EXTERNAL, to avoid cyclic dependencies and ensure a valid execution order
- srcSubpass and dstSubpass must not both be equal to VK_SUBPASS_EXTERNAL
- If srcSubpass is equal to dstSubpass, srcStageMask and dstStageMask must only contain one of VK_ PIPELINE_STAGE_TOP_OF_PIPE_BIT, VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT, VK_ PIPELINE_STAGE_VERTEX_SHADER_BIT, VK_ PIPELINE_STAGE_VERTEX_SHADER_BIT, VK_ PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT, VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT, VK_ PIPELINE_STAGE_GEOMETRY_SHADER_BIT, VK_ PIPELINE_STAGE_FRAGMENT_TESTS_BIT, VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT, VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT, VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT, or VK_PIPELINE_STAGE_ALL_GRAPHICS_BIT
- If srcSubpass is equal to dstSubpass, the highest bit value included in srcStageMask must be less than or equal to the lowest bit value in dstStageMask

Automatic image layout transitions between subpasses also interact with the subpass dependencies. If two subpasses are connected by a dependency and those two subpasses use the same attachment in a different layout, then the layout transition will occur after the memory accesses via <code>srcAccessMask</code> have completed in all pipeline stages included in <code>srcStageMask</code> in the source subpass, and before any memory accesses via <code>dstAccessMask</code> occur in any pipeline stages included in <code>dstStageMask</code> in the destination subpass.

The automatic image layout transitions from <code>initialLayout</code> to the first used layout (if it is different) are performed according to the following rules:

- If the attachment does not include the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT bit and there is no subpass dependency from VK_SUBPASS_EXTERNAL to the first subpass that uses the attachment, then it is as if there were such a dependency with srcStageMask = srcAccessMask = 0 and dstStageMask and dstAccessMask including all relevant bits (all graphics pipeline stages and all access types that use image resources), with the transition executing as part of that dependency. In other words, it may overlap work before the render pass instance and is complete before the subpass begins.
- If the attachment does not include the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT bit and there is a subpass dependency from VK_SUBPASS_EXTERNAL to the first subpass that uses the attachment, then the transition executes as part of that dependency and according to its stage and access masks. It must not overlap work that came before the render pass instance that is included in the source masks, but it may overlap work in previous subpasses.
- If the attachment includes the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT bit, then the transition executes according to all the subpass dependencies with <code>dstSubpass</code> equal to the first subpass index that the attachment is used in. That is, it occurs after all memory accesses in the source stages and masks from all the source subpasses have completed and are available, and before the union of all the destination stages begin, and the new layout is visible to the union of all the destination access types. If there are no incoming subpass dependencies, then this case follows the first rule.

Similar rules apply for the transition to the finalLayout, using dependencies with dstSubpass equal to VK_SUBPASS_EXTERNAL

If an attachment specifies the VK_ATTACHMENT_LOAD_OP_CLEAR load operation, then it will logically be cleared at the start of the first subpass where it is used.

Note



Implementations may move clears earlier as long as it does not affect the operation of a render pass instance. For example, an implementation may choose to clear all attachments at the start of the render pass instance. If an attachment has the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT flag set, then the clear must occur at the start of subpass where the attachment is first used, in order to preserve the operation of the render pass instance.

The first use of an attachment must not specify a layout equal to VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL or VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL if the attachment specifies that the <code>loadOp</code> is VK_ATTACHMENT_LOAD_OP_CLEAR. If a subpass uses the same attachment as both an input attachment and either a color attachment or a depth/stencil attachment, then both uses must observe the result of the clear.

Similarly, if an attachment specifies that the storeOp is VK_ATTACHMENT_STORE_OP_STORE, then it will logically be stored at the end of the last subpass where it is used.



Note

Implementations may move stores later as long as it does not affect the operation of a render pass instance. If an attachment has the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT flag set, then the store must occur at the end of the highest numbered subpass that uses the attachment.

If an attachment is not used by any subpass, then the <code>loadOp</code> and the <code>storeOp</code> are ignored and the attachment's memory contents will not be modified by execution of a render pass instance.

It will be common for a render pass to consist of a simple linear graph of dependencies, where subpass N depends on subpass N-1 for all N, and the operation of the memory barriers and layout transitions is fairly straightforward to reason about for those simple cases. But for more complex graphs, there are some rules that govern when there must be dependencies between subpasses.

As stated earlier, render passes must include subpass dependencies which (either directly or via a subpass dependency chain) separate any two subpasses that operate on the same attachment or aliasing attachments, if at least one of those subpasses writes to the attachment. If an image layout changes between those two subpasses, the implementation uses the stageMasks and accessMasks indicated by the subpass dependency as the masks that control when the layout transition must occur. If there is not a layout change on the attachment, or if an implementation treats the two layouts identically, then it may treat the dependency as a simple execution/memory barrier.

If two subpasses use the same attachment in different layouts but both uses are read-only (i.e. input attachment, or read-only depth/stencil attachment), the application does not need to express a dependency between the two subpasses. Implementations that treat the two layouts differently may deduce and insert a dependency between the subpasses, with the implementation choosing the appropriate stage masks and access masks based on whether the attachment is used as an input or depth/stencil attachment, and may insert the appropriate layout transition along with the execution/memory barrier. Implementations that treat the two layouts identically need not insert a barrier, and the two subpasses may execute simultaneously. The stage masks and access masks are chosen as follows:

- for input attachments, stage mask = VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT, access mask = VK_ACCESS_INPUT_ATTACHMENT_READ_BIT.
- for depth/stencil attachments, stage mask = VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT | VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT, access mask = VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_READ_BIT

where <code>srcStageMask</code> and <code>srcAccessMask</code> are taken based on usage in the source subpass and <code>dstStageMask</code> and <code>dstAccessMask</code> are taken based on usage in the destination subpass.

If a subpass uses the same attachment as both an input attachment and either a color attachment or a depth/stencil attachment, reads from the input attachment are not automatically coherent with writes through the color or depth/stencil attachment. In order to achieve well-defined results, one of two criteria must be satisfied. First, if the color components or depth/stencil components read by the input attachment are mutually exclusive with the components written by the color or depth/stencil attachment then there is no *feedback loop* and the reads and writes both function normally, with the reads observing values from the previous subpass(es) or from memory. This option requires the graphics pipelines used by the subpass to disable writes to color components that are read as inputs via the <code>colorWriteMask</code>, and to disable writes to depth/stencil components that are read as inputs via <code>depthWriteEnable</code> or <code>stencilTestEnable</code>.

Second, if the input attachment reads components that are written by the color or depth/stencil attachment, then there is a feedback loop and a pipeline barrier must be used between when the attachment is written and when it is subsequently read by later fragments. This pipeline barrier must follow the rules of a self-dependency as described in Subpass Self-dependency, where the barrier's flags include:

- dstStageMask = VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,
- dstAccessMask = VK_ACCESS_INPUT_ATTACHMENT_READ_BIT, and
- srcAccessMask = VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT (for color attachments) or srcAccessMask = VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT (for depth/stencil attachments).
- srcStageMask = VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT (for color attachments) or srcStageMask = VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT | VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT (for depth/stencil attachments).

• dependencyFlags = VK_DEPENDENCY_BY_REGION_BIT.

A pipeline barrier is needed each time a fragment will read a particular (x,y,layer,sample) location if that location has been written since the most recent pipeline barrier, or since the start of the subpass if there have been no pipeline barriers since the start of the subpass.

An attachment used as both an input attachment and color attachment must be in the VK_IMAGE_LAYOUT_GENERAL layout. An attachment used as both an input attachment and depth/stencil attachment must be in either the VK_IMAGE_LAYOUT_GENERAL or VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL layout. Since an attachment in the VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL layout is read-only, this situation is not a feedback loop.

To destroy a render pass, call:

- device is the logical device that destroys the render pass.
- renderPass is the handle of the render pass to destroy.
- pallocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If renderPass is not VK_NULL_HANDLE, renderPass must be a valid VkRenderPass handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If renderPass is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to renderPass must have completed execution
- If VkAllocationCallbacks were provided when renderPass was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when renderPass was created, pAllocator must be NULL

Host Synchronization

• Host access to renderPass must be externally synchronized

7.2 Render Pass Compatibility

Framebuffers and graphics pipelines are created based on a specific render pass object. They must only be used with that render pass object, or one compatible with it.

Two attachment references are compatible if they have matching format and sample count, or are both VK_ATTACHMENT_UNUSED or the pointer that would contain the reference is NULL.

Two arrays of attachment references are compatible if all corresponding pairs of attachments are compatible. If the arrays are of different lengths, attachment references not present in the smaller array are treated as VK_ATTACHMENT_UNUSED.

Two render passes that contain only a single subpass are compatible if their corresponding color, input, resolve, and depth/stencil attachment references are compatible.

If two render passes contain more than one subpass, they are compatible if they are identical except for:

- Initial and final image layout in attachment descriptions
- · Load and store operations in attachment descriptions
- Image layout in attachment references

A framebuffer is compatible with a render pass if it was created using the same render pass or a compatible render pass.

7.3 Framebuffers

Render passes operate in conjunction with *framebuffers*. Framebuffers represent a collection of specific memory attachments that a render pass instance uses.

Framebuffers are represented by VkFramebuffer handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkFramebuffer)
```

To create a framebuffer, call:

- device is the logical device that creates the framebuffer.
- pCreateInfo points to a VkFramebufferCreateInfo structure which describes additional information about framebuffer creation.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pFramebuffer points to a VkFramebuffer handle in which the resulting framebuffer object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkFramebufferCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pFramebuffer must be a pointer to a VkFramebuffer handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkFramebufferCreateInfo structure is defined as:

```
typedef struct VkFramebufferCreateInfo {
   VkStructureType
                            sType;
   const void*
                            pNext;
   VkFramebufferCreateFlags flags;
                            renderPass;
   VkRenderPass
                            attachmentCount;
   uint32_t
                          pAttachments;
   const VkImageView*
   uint32_t
                             width;
                            height;
   uint32_t
   uint32_t
                             layers;
} VkFramebufferCreateInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- renderPass is a render pass that defines what render passes the framebuffer will be compatible with. See Render Pass Compatibility for details.
- attachmentCount is the number of attachments.
- pAttachments is an array of VkImageView handles, each of which will be used as the corresponding attachment in a render pass instance.

• width, height and layers define the dimensions of the framebuffer.

Image subresources used as attachments must not be used via any non-attachment usage for the duration of a render pass instance.



Note

This restriction means that the render pass has full knowledge of all uses of all of the attachments, so that the implementation is able to make correct decisions about when and how to perform layout transitions, when to overlap execution of subpasses, etc.

It is legal for a subpass to use no color or depth/stencil attachments, and rather use shader side effects such as image stores and atomics to produce an output. In this case, the subpass continues to use the <code>width</code>, <code>height</code>, and <code>layers</code> of the framebuffer to define the dimensions of the rendering area, and the <code>rasterizationSamples</code> from each pipeline's <code>VkPipelineMultisampleStateCreateInfo</code> to define the number of samples used in rasterization; however, if <code>VkPhysicalDeviceFeatures::variableMultisampleRate</code> is <code>VK_FALSE</code>, then all pipelines to be bound with a given zero-attachment subpass must have the same value for

VkPipelineMultisampleStateCreateInfo::rasterizationSamples.

- stype must be VK STRUCTURE TYPE FRAMEBUFFER CREATE INFO
- pNext must be NULL
- flags must be 0
- renderPass must be a valid VkRenderPass handle
- If attachmentCount is not 0, pAttachments must be a pointer to an array of attachmentCount valid VkImageView handles
- Both of renderPass, and the elements of pAttachments that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- attachmentCount must be equal to the attachment count specified in renderPass
- Any given element of pAttachments that is used as a color attachment or resolve attachment by renderPass must have been created with a usage value including VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT
- Any given element of pAttachments that is used as a depth/stencil attachment by renderPass must have been created with a usage value including VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT
- Any given element of pAttachments that is used as an input attachment by renderPass must have been created with a usage value including VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT
- Any given element of pAttachments must have been created with an VkFormat value that matches the VkFormat specified by the corresponding VkAttachmentDescription in renderPass
- Any given element of pAttachments must have been created with a samples value that matches the samples value specified by the corresponding VkAttachmentDescription in renderPass

- Any given element of pAttachments must have dimensions at least as large as the corresponding framebuffer dimension
- Any given element of pAttachments must only specify a single mip level
- Any given element of pAttachments must have been created with the identity swizzle
- width must be less than or equal to VkPhysicalDeviceLimits::maxFramebufferWidth
- height must be less than or equal to VkPhysicalDeviceLimits::maxFramebufferHeight
- layers must be less than or equal to VkPhysicalDeviceLimits::maxFramebufferLayers

To destroy a framebuffer, call:

- device is the logical device that destroys the framebuffer.
- framebuffer is the handle of the framebuffer to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If framebuffer is not VK_NULL_HANDLE, framebuffer must be a valid VkFramebuffer handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If framebuffer is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to framebuffer must have completed execution
- If VkAllocationCallbacks were provided when framebuffer was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when framebuffer was created, pAllocator must be NULL

Host Synchronization

• Host access to framebuffer must be externally synchronized

7.4 Render Pass Commands

An application records the commands for a render pass instance one subpass at a time, by beginning a render pass instance, iterating over the subpasses to record commands for that subpass, and then ending the render pass instance.

To begin a render pass instance, call:

- commandBuffer is the command buffer in which to record the command.
- pRenderPassBegin is a pointer to a VkRenderPassBeginInfo structure (defined below) which indicates the render pass to begin an instance of, and the framebuffer the instance uses.
- contents specifies how the commands in the first subpass will be provided, and is one of the values:

```
typedef enum VkSubpassContents {
    VK_SUBPASS_CONTENTS_INLINE = 0,
    VK_SUBPASS_CONTENTS_SECONDARY_COMMAND_BUFFERS = 1,
} VkSubpassContents;
```

If contents is VK_SUBPASS_CONTENTS_INLINE, the contents of the subpass will be recorded inline in the primary command buffer, and secondary command buffers must not be executed within the subpass. If contents is VK_SUBPASS_CONTENTS_SECONDARY_COMMAND_BUFFERS, the contents are recorded in secondary command buffers that will be called from the primary command buffer, and vkCmdExecuteCommands is the only valid command on the command buffer until vkCmdNextSubpass or vkCmdEndRenderPass.

After beginning a render pass instance, the command buffer is ready to record the commands for the first subpass of that render pass.

- commandBuffer must be a valid VkCommandBuffer handle
- pRenderPassBegin must be a pointer to a valid VkRenderPassBeginInfo structure
- contents must be a valid VkSubpassContents value
- commandBuffer must be in the recording state

- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called outside of a render pass instance
- commandBuffer must be a primary VkCommandBuffer
- If any of the <code>initialLayout</code> or <code>finalLayout</code> member of the <code>VkAttachmentDescription</code> structures or the <code>layout</code> member of the <code>VkAttachmentReference</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is <code>VK_IMAGE_LAYOUT_COLOR_</code> ATTACHMENT_OPTIMAL then the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code> must have been created with <code>VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT</code> set
- If any of the <code>initialLayout</code> or <code>finalLayout</code> member of the <code>VkAttachmentDescription</code> structures or the <code>layout</code> member of the <code>VkAttachmentReference</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is <code>VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL</code> or <code>VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL</code> then the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code> must have been created with <code>VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT</code> set
- If any of the <code>initialLayout</code> or <code>finalLayout</code> member of the <code>VkAttachmentDescription</code> structures or the <code>layout</code> member of the <code>VkAttachmentReference</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is <code>VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL</code> then the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code> must have been created with <code>VK_IMAGE_USAGE_SAMPLED_BIT</code> or <code>VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT</code> set
- If any of the <code>initialLayout</code> or <code>finalLayout</code> member of the <code>VkAttachmentDescription</code> structures or the <code>layout</code> member of the <code>VkAttachmentReference</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is <code>VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL</code> then the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code> must have been created with <code>VK_IMAGE_USAGE_TRANSFER_SRC_BIT</code> set
- If any of the <code>initialLayout</code> or <code>finalLayout</code> member of the <code>VkAttachmentDescription</code> structures or the <code>layout</code> member of the <code>VkAttachmentReference</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is <code>VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL</code> then the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code> must have been created with <code>VK_IMAGE_USAGE_TRANSFER_DST_DST_BIT</code> set
- If any of the <code>initialLayout</code> members of the <code>VkAttachmentDescription</code> structures specified when creating the render pass specified in the <code>renderPass</code> member of <code>pRenderPassBegin</code> is not <code>VK_IMAGE_LAYOUT_UNDEFINED</code>, then each such <code>initialLayout</code> must be equal to the current layout of the corresponding attachment image subresource of the framebuffer specified in the <code>framebuffer</code> member of <code>pRenderPassBegin</code>

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS

The VkRenderPassBeginInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- renderPass is the render pass to begin an instance of.
- framebuffer is the framebuffer containing the attachments that are used with the render pass.
- renderArea is the render area that is affected by the render pass instance, and is described in more detail below.
- clearValueCount is the number of elements in pClearValues.
- pClearValues is an array of VkClearValue structures that contains clear values for each attachment, if the attachment uses a loadOp value of VK_ATTACHMENT_LOAD_OP_CLEAR or if the attachment has a depth/stencil format and uses a stencilLoadOp value of VK_ATTACHMENT_LOAD_OP_CLEAR. The array is indexed by attachment number. Only elements corresponding to cleared attachments are used. Other elements of pClearValues are ignored.

renderArea is the render area that is affected by the render pass instance. The effects of attachment load, store and resolve operations are restricted to the pixels whose x and y coordinates fall within the render area on all attachments. The render area extends to all layers of framebuffer. The application must ensure (using scissor if necessary) that all rendering is contained within the render area, otherwise the pixels outside of the render area become undefined and shader side effects may occur for fragments outside the render area. The render area must be contained within the framebuffer dimensions.



Note

There may be a performance cost for using a render area smaller than the framebuffer, unless it matches the render area granularity for the render pass.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_RENDER_PASS_BEGIN_INFO
- pNext must be NULL
- renderPass must be a valid VkRenderPass handle
- framebuffer must be a valid VkFramebuffer handle
- If clearValueCount is not 0, pClearValues must be a pointer to an array of clearValueCount valid VkClearValue unions
- Both of framebuffer, and renderPass must have been created, allocated, or retrieved from the same VkDevice
- clearValueCount must be greater than the largest attachment index in renderPass that specifies a loadOp (or stencilLoadOp, if the attachment has a depth/stencil format) of VK_ATTACHMENT_LOAD_OP_CLEAR

To query the render area granularity, call:

- device is the logical device that owns the render pass.
- renderPass is a handle to a render pass.
- pGranularity points to a VkExtent2D structure in which the granularity is returned.

The conditions leading to an optimal renderArea are:

- the offset.x member in renderArea is a multiple of the width member of the returned VkExtent2D (the horizontal granularity).
- the offset.y member in renderArea is a multiple of the height of the returned VkExtent2D (the vertical granularity).
- either the offset.width member in renderArea is a multiple of the horizontal granularity or offset.x+offset. width is equal to the width of the framebuffer in the VkRenderPassBeginInfo.

• either the offset.height member in renderArea is a multiple of the vertical granularity or offset.y+offset. height is equal to the height of the framebuffer in the VkRenderPassBeginInfo.

Subpass dependencies are not affected by the render area, and apply to the entire image subresources attached to the framebuffer. Similarly, pipeline barriers are valid even if their effect extends outside the render area.

Valid Usage

- device must be a valid VkDevice handle
- renderPass must be a valid VkRenderPass handle
- pGranularity must be a pointer to a VkExtent2D structure
- renderPass must have been created, allocated, or retrieved from device

To transition to the next subpass in the render pass instance after recording the commands for a subpass, call:

- commandBuffer is the command buffer in which to record the command.
- contents specifies how the commands in the next subpass will be provided, in the same fashion as the corresponding parameter of vkCmdBeqinRenderPass.

The subpass index for a render pass begins at zero when **vkCmdBeginRenderPass** is recorded, and increments each time **vkCmdNextSubpass** is recorded.

Moving to the next subpass automatically performs any multisample resolve operations in the subpass being ended. End-of-subpass multisample resolves are treated as color attachment writes for the purposes of synchronization. That is, they are considered to execute in the VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT pipeline stage and their writes are synchronized with VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT. Synchronization between rendering within a subpass and any resolve operations at the end of the subpass occurs automatically, without need for explicit dependencies or pipeline barriers. However, if the resolve attachment is also used in a different subpass, an explicit dependency is needed.

After transitioning to the next subpass, the application can record the commands for that subpass.

- commandBuffer must be a valid VkCommandBuffer handle
- contents must be a valid VkSubpassContents value

- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- commandBuffer must be a primary VkCommandBuffer
- The current subpass index must be less than the number of subpasses in the render pass minus one

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS

To record a command to end a render pass instance after recording the commands for the last subpass, call:

```
void vkCmdEndRenderPass(
     VkCommandBuffer commandBuffer);
```

• commandBuffer is the command buffer in which to end the current render pass instance.

Ending a render pass instance performs any multisample resolve operations on the final subpass.

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations

- This command must only be called inside of a render pass instance
- commandBuffer must be a primary VkCommandBuffer
- The current subpass index must be equal to the number of subpasses in the render pass minus one

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types	
Primary	Inside	GRAPHICS	

Chapter 8

Shaders

A shader specifies programmable operations that execute for each vertex, control point, tessellated vertex, primitive, fragment, or workgroup in the corresponding stage(s) of the graphics and compute pipelines.

Graphics pipelines include vertex shader execution as a result of primitive assembly, followed, if enabled, by tessellation control and evaluation shaders operating on patches, geometry shaders, if enabled, operating on primitives, and fragment shaders, if present, operating on fragments generated by Rasterization. In this specification, vertex, tessellation control, tessellation evaluation and geometry shaders are collectively referred to as vertex processing stages and occur in the logical pipeline before rasterization. The fragment shader occurs logically after rasterization.

Only the compute shader stage is included in a compute pipeline. Compute shaders operate on compute invocations in a workgroup.

Shaders can read from input variables, and read from and write to output variables. Input and output variables can be used to transfer data between shader stages, or to allow the shader to interact with values that exist in the execution environment. Similarly, the execution environment provides constants that describe capabilities.

Shader variables are associated with execution environment-provided inputs and outputs using *built-in* decorations in the shader. The available decorations for each stage are documented in the following subsections.

8.1 Shader Modules

Shader modules contain shader code and one or more entry points. Shaders are selected from a shader module by specifying an entry point as part of pipeline creation. The stages of a pipeline can use shaders that come from different modules. The shader code defining a shader module must be in the SPIR-V format, as described by the Vulkan Environment for SPIR-V appendix.

Shader modules are represented by VkShaderModule handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkShaderModule)
```

To create a shader module, call:

- device is the logical device that creates the shader module.
- pCreateInfo parameter is a pointer to an instance of the VkShaderModuleCreateInfo structure.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pShaderModule points to a VkShaderModule handle in which the resulting shader module object is returned.

Once a shader module has been created, any entry points it contains can be used in pipeline shader stages as described in Compute Pipelines and Graphics Pipelines.

If the shader stage fails to compile VK_ERROR_INVALID_SHADER_NV will be generated and the compile log will be reported back to the application by VK_EXT_debug_report if enabled.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkShaderModuleCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pShaderModule must be a pointer to a VkShaderModule handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK ERROR INVALID SHADER NV

The VkShaderModuleCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- codeSize is the size, in bytes, of the code pointed to by pCode.
- pCode points to code that is used to create the shader module. The type and format of the code is determined from the content of the memory addressed by pCode.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_SHADER_MODULE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- pCode must be a pointer to an array of codeSize uint32_t values
- codeSize must be greater than 0
- codeSize must be a multiple of 4. If the VK_NV_glsl_shader extension is enabled and pCode references GLSL code codeSize can be a multiple of 1
- pCode must point to valid SPIR-V code, formatted and packed as described by the Khronos SPIR-V Specification. If the VK_NV_glsl_shader extension is enabled pCode can instead reference valid GLSL code and must be written to the GL_KHR_vulkan_glsl extension specification
- pCode must adhere to the validation rules described by the Validation Rules within a Module section of the SPIR-V Environment appendix. If the VK_NV_glsl_shader extension is enabled pCode can be valid GLSL code with respect to the GL_KHR_vulkan_glsl GLSL extension specification
- pCode must declare the **Shader** capability for SPIR-V code
- pCode must not declare any capability that is not supported by the API, as described by the Capabilities section of the SPIR-V Environment appendix
- If pCode declares any of the capabilities that are listed as not required by the implementation, the relevant feature must be enabled, as listed in the SPIR-V Environment appendix

To destroy a shader module, call:

• device is the logical device that destroys the shader module.

- *shaderModule* is the handle of the shader module to destroy.
- pallocator controls host memory allocation as described in the Memory Allocation chapter.

A shader module can be destroyed while pipelines created using its shaders are still in use.

Valid Usage

- device must be a valid VkDevice handle
- If shaderModule is not VK NULL HANDLE, shaderModule must be a valid VkShaderModule handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If shaderModule is a valid handle, it must have been created, allocated, or retrieved from device
- If VkAllocationCallbacks were provided when *shaderModule* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when shaderModule was created, pAllocator must be NULL.

Host Synchronization

• Host access to shaderModule must be externally synchronized

8.2 Shader Execution

At each stage of the pipeline, multiple invocations of a shader may execute simultaneously. Further, invocations of a single shader produced as the result of different commands may execute simultaneously. The relative execution order of invocations of the same shader type is undefined. Shader invocations may complete in a different order than that in which the primitives they originated from were drawn or dispatched by the application. However, fragment shader outputs are written to attachments in rasterization order.

The relative order of invocations of different shader types is largely undefined. However, when invoking a shader whose inputs are generated from a previous pipeline stage, the shader invocations from the previous stage are guaranteed to have executed far enough to generate input values for all required inputs.

8.3 Shader Memory Access Ordering

The order in which image or buffer memory is read or written by shaders is largely undefined. For some shader types (vertex, tessellation evaluation, and in some cases, fragment), even the number of shader invocations that may perform loads and stores is undefined.

In particular, the following rules apply:

- Vertex and tessellation evaluation shaders will be invoked at least once for each unique vertex, as defined in those sections.
- Fragment shaders will be invoked zero or more times, as defined in that section.
- The relative order of invocations of the same shader type are undefined. A store issued by a shader when working on primitive B might complete prior to a store for primitive A, even if primitive A is specified prior to primitive B. This applies even to fragment shaders; while fragment shader outputs are always written to the framebuffer in primitive order, stores executed by fragment shader invocations are not.
- The relative order of invocations of different shader types is largely undefined.



Note

The above limitations on shader invocation order make some forms of synchronization between shader invocations within a single set of primitives unimplementable. For example, having one invocation poll memory written by another invocation assumes that the other invocation has been launched and will complete its writes in finite time.

Stores issued to different memory locations within a single shader invocation may not be visible to other invocations in the order they were performed. The <code>OpMemoryBarrier</code> instruction can be used to provide stronger ordering of reads and writes performed by a single invocation. <code>OpMemoryBarrier</code> guarantees that any memory transactions issued by the shader invocation prior to the instruction complete prior to the memory transactions issued after the instruction. Memory barriers are needed for algorithms that require multiple invocations to access the same memory and require the operations to be performed in a partially-defined relative order. For example, if one shader invocation does a series of writes, followed by an <code>OpMemoryBarrier</code> instruction, followed by another write, then the results of the series of writes before the barrier become visible to other shader invocations at a time earlier or equal to when the results of the final write become visible to those invocations. In practice it means that another invocation that sees the results of the final write would also see the previous writes. Without the memory barrier, the final write may be visible before the previous writes.

The built-in atomic memory transaction instructions can be used to read and write a given memory address atomically. While built-in atomic functions issued by multiple shader invocations are executed in undefined order relative to each other, these functions perform both a read and a write of a memory address and guarantee that no other memory transaction will write to the underlying memory between the read and write.



Note

Atomics allow shaders to use shared global addresses for mutual exclusion or as counters, among other uses.

8.4 Shader Inputs and Outputs

Data is passed into and out of shaders using variables with input or output storage class, respectively. User-defined inputs and outputs are connected between stages by matching their **Location** decorations. Additionally, data can be provided by or communicated to special functions provided by the execution environment using **BuiltIn** decorations.

In many cases, the same **BuiltIn** decoration can be used in multiple shader stages with similar meaning. The specific behavior of variables decorated as **BuiltIn** is documented in the following sections.

8.5 Vertex Shaders

Each vertex shader invocation operates on one vertex and its associated vertex attribute data, and outputs one vertex and associated data. Graphics pipelines must include a vertex shader, and the vertex shader stage is always the first shader stage in the graphics pipeline.

8.5.1 Vertex Shader Execution

A vertex shader must be executed at least once for each vertex specified by a draw command. During execution, the shader is presented with the index of the vertex and instance for which it has been invoked. Input variables declared in the vertex shader are filled by the implementation with the values of vertex attributes associated with the invocation being executed.

If the same vertex is specified multiple times in a draw command (e.g. by including the same index value multiple times in an index buffer) the implementation may reuse the results of vertex shading if it can statically determine that the vertex shader invocations will produce identical results.



Note

It is implementation-dependent when and if results of vertex shading are reused, and thus how many times the vertex shader will be executed. This is true also if the vertex shader contains stores or atomic operations (see vertexPipelineStoresAndAtomics).

8.6 Tessellation Control Shaders

The tessellation control shader is used to read an input patch provided by the application and to produce an output patch. Each tessellation control shader invocation operates on an input patch (after all control points in the patch are processed by a vertex shader) and its associated data, and outputs a single control point of the output patch and its associated data, and can also output additional per-patch data. The input patch is sized according to the <code>patchControlPoints</code> member of <code>VkPipelineTessellationStateCreateInfo</code>, as part of input assembly. The size of the output patch is controlled by the <code>OpExecutionModeOutputVertices</code> specified in the tessellation control or tessellation evaluation shaders, which must be specified in at least one of the shaders. The size of the input and output patches must each be greater than zero and less than or equal to <code>VkPhysicalDeviceLimits::maxTessellationPatchSize</code>.

8.6.1 Tessellation Control Shader Execution

A tessellation control shader is invoked at least once for each *output* vertex in a patch.

Inputs to the tessellation control shader are generated by the vertex shader. Each invocation of the tessellation control shader can read the attributes of any incoming vertices and their associated data. The invocations corresponding to a given patch execute logically in parallel, with undefined relative execution order. However, the <code>OpControlBarrier</code> instruction can be used to provide limited control of the execution order by synchronizing invocations within a patch, effectively dividing tessellation control shader execution into a set of phases. Tessellation control shaders will read undefined values if one invocation reads a per-vertex or per-patch attribute written by another invocation at any point during the same phase, or if two invocations attempt to write different values to the same per-patch output in a single phase.

8.7 Tessellation Evaluation Shaders

The Tessellation Evaluation Shader operates on an input patch of control points and their associated data, and a single input barycentric coordinate indicating the invocation's relative position within the subdivided patch, and outputs a single vertex and its associated data.

8.7.1 Tessellation Evaluation Shader Execution

A tessellation evaluation shader is invoked at least once for each unique vertex generated by the tessellator.

8.8 Geometry Shaders

The geometry shader operates on a group of vertices and their associated data assembled from a single input primitive, and emits zero or more output primitives and the group of vertices and their associated data required for each output primitive.

8.8.1 Geometry Shader Execution

A geometry shader is invoked at least once for each primitive produced by the tessellation stages, or at least once for each primitive generated by primitive assembly when tessellation is not in use. The number of geometry shader invocations per input primitive is determined from the invocation count of the geometry shader specified by the **OpExecutionMode Invocations** in the geometry shader. If the invocation count is not specified, then a default of one invocation is executed.

8.9 Fragment Shaders

Fragment shaders are invoked as the result of rasterization in a graphics pipeline. Each fragment shader invocation operates on a single fragment and its associated data. With few exceptions, fragment shaders do not have access to any data associated with other fragments and are considered to execute in isolation of fragment shader invocations associated with other fragments.

8.9.1 Fragment Shader Execution

For each fragment generated by rasterization, a fragment shader may be invoked. A fragment shader must not be invoked if the Early Per-Fragment Tests cause it to have no coverage.

Furthermore, if it is determined that a fragment generated as the result of rasterizing a first primitive will have its outputs entirely overwritten by a fragment generated as the result of rasterizing a second primitive in the same subpass, and the fragment shader used for the fragment has no other side effects, then the fragment shader may not be executed for the fragment from the first primitive.

Relative ordering of execution of different fragment shader invocations is not defined.

The number of fragment shader invocations produced per-pixel is determined as follows:

- If per-sample shading is enabled, the fragment shader is invoked once per covered sample.
- Otherwise, the fragment shader is invoked at least once per fragment but no more than once per covered sample.

In addition to the conditions outlined above for the invocation of a fragment shader, a fragment shader invocation may be produced as a *helper invocation*. A helper invocation is a fragment shader invocation that is created solely for the purposes of evaluating derivatives for use in non-helper fragment shader invocations. Stores and atomics performed by helper invocations must not have any effect on memory, and values returned by atomic instructions in helper invocations are undefined.

8.9.2 Early Fragment Tests

An explicit control is provided to allow fragment shaders to enable early fragment tests. If the fragment shader specifies the **EarlyFragmentTests OpExecutionMode**, the per-fragment tests described in Early Fragment Test Mode are performed prior to fragment shader execution. Otherwise, they are performed after fragment shader execution.

8.10 Compute Shaders

Compute shaders are invoked via vkCmdDispatch and vkCmdDispatchIndirect commands. In general, they have access to similar resources as shader stages executing as part of a graphics pipeline.

Compute workloads are formed from groups of work items called workgroups and processed by the compute shader in the current compute pipeline. A workgroup is a collection of shader invocations that execute the same shader, potentially in parallel. Compute shaders execute in *global workgroups* which are divided into a number of *local workgroups* with a size that can be set by assigning a value to the **LocalSize** execution mode or via an object decorated by the **WorkgroupSize** decoration. An invocation within a local workgroup can share data with other members of the local workgroup through shared variables and issue memory and control flow barriers to synchronize with other members of the local workgroup.

8.11 Interpolation Decorations

Interpolation decorations control the behavior of attribute interpolation in the fragment shader stage. Interpolation decorations can be applied to **Input** storage class variables in the fragment shader stage's interface, and control the interpolation behavior of those variables.

Inputs that could be interpolated can be decorated by at most one of the following decorations:

- Flat: no interpolation
- **NoPerspective**: linear interpolation (for lines and polygons).

Fragment input variables decorated with neither **Flat** nor **NoPerspective** use perspective-correct interpolation (for lines and polygons).

The presence of and type of interpolation is controlled by the above interpolation decorations as well as the auxiliary decorations **Centroid** and **Sample**.

A variable decorated with **Flat** will not be interpolated. Instead, it will have the same value for every fragment within a triangle. This value will come from a single provoking vertex. A variable decorated with **Flat** can also be decorated with **Centroid** or **Sample**, which will mean the same thing as decorating it only as **Flat**.

For fragment shader input variables decorated with neither **Centroid** nor **Sample**, the assigned variable may be interpolated anywhere within the pixel and a single value may be assigned to each sample within the pixel.

Centroid and **Sample** can be used to control the location and frequency of the sampling of the decorated fragment shader input. If a fragment shader input is decorated with **Centroid**, a single value may be assigned to that variable for

all samples in the pixel, but that value must be interpolated to a location that lies in both the pixel and in the primitive being rendered, including any of the pixel's samples covered by the primitive. Because the location at which the variable is interpolated may be different in neighboring pixels, and derivatives may be computed by computing differences between neighboring pixels, derivatives of centroid-sampled inputs may be less accurate than those for non-centroid interpolated variables. If a fragment shader input is decorated with Sample, a separate value must be assigned to that variable for each covered sample in the pixel, and that value must be sampled at the location of the individual sample. When rasterizationSamples is VK_SAMPLE_COUNT_1_BIT, the pixel center must be used for Centroid, Sample, and undecorated attribute interpolation.

Fragment shader inputs that are signed or unsigned integers, integer vectors, or any double-precision floating-point type must be decorated with **Flat**.

When the VK_AMD_shader_explicit_vertex_parameter device extension is enabled inputs can be also decorated with the <code>CustomInterpAMD</code> interpolation decoration, including fragment shader inputs that are signed or unsigned integers, integer vectors, or any double-precision floating-point type. Inputs decorated with <code>CustomInterpAMD</code> can only be accessed by the extended instruction <code>InterpolateAtVertexAMD</code> and allows accessing the value of the input for individual vertices of the primitive.

8.12 Static Use

A SPIR-V module declares a global object in memory using the **OpVariable** instruction, which results in a pointer **x** to that object. A specific entry point in a SPIR-V module is said to *statically use* that object if that entry-point's call tree contains a function that contains a memory instruction or image instruction with **x** as an **id** operand. See the "Memory Instructions" and "Image Instructions" subsections of section 3 "Binary Form" of the SPIR-V specification for the complete list of SPIR-V memory instructions.

Static use is not used to control the behavior of variables with **Input** and **Output** storage. The effects of those variables are applied based only on whether they are present in a shader entry point's interface.

8.13 Invocation and Derivative Groups

An *invocation group* (see the subsection "Control Flow" of section 2 of the SPIR-V specification) for a compute shader is the set of invocations in a single local workgroup. For graphics shaders, an invocation group is an implementation-dependent subset of the set of shader invocations of a given shader stage which are produced by a single drawing command. For indirect drawing commands with <code>drawCount</code> greater than one, invocations from separate draws are in distinct invocation groups.



Note

Because the partitioning of invocations into invocation groups is implementation-dependent and not observable, applications generally need to assume the worst case of all invocations in a draw belonging to a single invocation group.

A *derivative group* (see the subsection "Control Flow" of section 2 of the SPIR-V 1.00 Revision 4 specification) for a fragment shader is the set of invocations generated by a single primitive (point, line, or triangle), including any helper invocations generated by that primitive. Derivatives are undefined for a sampled image instruction if the instruction is in flow control that is not uniform across the derivative group.

Chapter 9

Pipelines

The following figure shows a block diagram of the Vulkan pipelines. Some Vulkan commands specify geometric objects to be drawn or computational work to be performed, while others specify state controlling how objects are handled by the various pipeline stages, or control data transfer between memory organized as images and buffers. Commands are effectively sent through a processing pipeline, either a *graphics pipeline* or a *compute pipeline*.

The first stage of the graphics pipeline (Input Assembler) assembles vertices to form geometric primitives such as points, lines, and triangles, based on a requested primitive topology. In the next stage (Vertex Shader) vertices can be transformed, computing positions and attributes for each vertex. If tessellation and/or geometry shaders are supported, they can then generate multiple primitives from a single input primitive, possibly changing the primitive topology or generating additional attribute data in the process.

The final resulting primitives are clipped to a clip volume in preparation for the next stage, Rasterization. The rasterizer produces a series of framebuffer addresses and values using a two-dimensional description of a point, line segment, or triangle. Each *fragment* so produced is fed to the next stage (Fragment Shader) that performs operations on individual fragments before they finally alter the framebuffer. These operations include conditional updates into the framebuffer based on incoming and previously stored depth values (to effect depth buffering), blending of incoming fragment colors with stored colors, as well as masking, stenciling, and other logical operations on fragment values.

Framebuffer operations read and write the color and depth/stencil attachments of the framebuffer for a given subpass of a render pass instance. The attachments can be used as input attachments in the fragment shader in a later subpass of the same render pass.

The compute pipeline is a separate pipeline from the graphics pipeline, which operates on one-, two-, or three-dimensional workgroups which can read from and write to buffer and image memory.

This ordering is meant only as a tool for describing Vulkan, not as a strict rule of how Vulkan is implemented, and we present it only as a means to organize the various operations of the pipelines.



Figure 9.1: Block diagram of the Vulkan pipeline

Each pipeline is controlled by a monolithic object created from a description of all of the shader stages and any relevant fixed-function stages. Linking the whole pipeline together allows the optimization of shaders based on their input/outputs and eliminates expensive draw time state validation.

A pipeline object is bound to the device state in command buffers. Any pipeline object state that is marked as dynamic is not applied to the device state when the pipeline is bound. Dynamic state not set by binding the pipeline object can be modified at any time and persists for the lifetime of the command buffer, or until modified by another dynamic state

command or another pipeline bind. No state, including dynamic state, is inherited from one command buffer to another. Only dynamic state that is required for the operations performed in the command buffer needs to be set. For example, if blending is disabled by the pipeline state then the dynamic color blend constants do not need to be specified in the command buffer, even if this state is marked as dynamic in the pipeline state object. If a new pipeline object is bound with state not marked as dynamic after a previous pipeline object with that same state as dynamic, the new pipeline object state will override the dynamic state. Modifying dynamic state that is not set as dynamic by the pipeline state object will lead to undefined results.

Compute and graphics pipelines are each represented by VkPipeline handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkPipeline)
```

9.1 Compute Pipelines

Compute pipelines consist of a single static compute shader stage and the pipeline layout.

The compute pipeline represents a compute shader and is created by calling **vkCreateComputePipelines** with <code>module</code> and <code>pName</code> selecting an entry point from a shader module, where that entry point defines a valid compute shader, in the <code>VkPipelineShaderStageCreateInfo</code> structure contained within the <code>VkComputePipelineCreateInfo</code> structure.

To create compute pipelines, call:

- device is the logical device that creates the compute pipelines.
- pipelineCache is either VK_NULL_HANDLE, indicating that pipeline caching is disabled; or the handle of a valid pipeline cache object, in which case use of that cache is enabled for the duration of the command.
- createInfoCount is the length of the pCreateInfos and pPipelines arrays.
- pCreateInfos is an array of VkComputePipelineCreateInfo structures.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pPipelines is a pointer to an array in which the resulting compute pipeline objects are returned.

- device must be a valid VkDevice handle
- If pipelineCache is not VK_NULL_HANDLE, pipelineCache must be a valid VkPipelineCache handle

- pCreateInfos must be a pointer to an array of createInfoCount valid VkComputePipelineCreateInfo structures
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pPipelines must be a pointer to an array of createInfoCount VkPipeline handles
- createInfoCount must be greater than 0
- If pipelineCache is a valid handle, it must have been created, allocated, or retrieved from device
- If the flags member of any given element of pCreateInfos contains the VK_PIPELINE_CREATE_ DERIVATIVE_BIT flag, and the basePipelineIndex member of that same element is not -1, basePipelineIndex must be less than the index into pCreateInfos that corresponds to that element

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_INVALID_SHADER_NV

The VkComputePipelineCreateInfo structure is defined as:

```
typedef struct VkComputePipelineCreateInfo {
   VkStructureType
                                      sType;
   const void*
                                      pNext;
   VkPipelineCreateFlags
                                      flags;
   VkPipelineShaderStageCreateInfo
                                      stage;
   VkPipelineLayout
                                       layout;
   VkPipeline
                                       basePipelineHandle;
   int32 t
                                       basePipelineIndex;
} VkComputePipelineCreateInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags provides options for pipeline creation, and is of type VkPipelineCreateFlagBits.
- stage is a VkPipelineShaderStageCreateInfo describing the compute shader.
- layout is the description of binding locations used by both the pipeline and descriptor sets used with the pipeline.

- basePipelineHandle is a pipeline to derive from
- basePipelineIndex is an index into the pCreateInfos parameter to use as a pipeline to derive from

The parameters basePipelineHandle and basePipelineIndex are described in more detail in Pipeline Derivatives. stage points to a structure of type VkPipelineShaderStageCreateInfo.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_COMPUTE_PIPELINE_CREATE_INFO
- pNext must be NULL
- flags must be a valid combination of VkPipelineCreateFlagBits values
- stage must be a valid VkPipelineShaderStageCreateInfo structure
- layout must be a valid VkPipelineLayout handle
- Both of basePipelineHandle, and layout that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineIndex is not 1, basePipelineHandle must be VK_NULL_HANDLE
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineIndex is not 1, it must be a valid index into the calling command's pCreateInfos parameter
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, basePipelineIndex must be -1
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, basePipelineHandle must be a valid VkPipeline handle
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, it must be a valid handle to a compute VkPipeline
- The stage member of stage must be VK_SHADER_STAGE_COMPUTE_BIT
- The shader code for the entry point identified by stage and the rest of the state identified by this structure must adhere to the pipeline linking rules described in the Shader Interfaces chapter
- layout must be consistent with all shaders specified in pStages

The VkPipelineShaderStageCreateInfo structure is defined as:

```
const char*
  const VkSpecializationInfo*
  VkPipelineShaderStageCreateInfo;
pName;
pSpecializationInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- stage names a single pipeline stage. Bits which can be set include:

```
typedef enum VkShaderStageFlagBits {
    VK_SHADER_STAGE_VERTEX_BIT = 0x00000001,
    VK_SHADER_STAGE_TESSELLATION_CONTROL_BIT = 0x00000002,
    VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT = 0x000000004,
    VK_SHADER_STAGE_GEOMETRY_BIT = 0x000000008,
    VK_SHADER_STAGE_FRAGMENT_BIT = 0x00000010,
    VK_SHADER_STAGE_COMPUTE_BIT = 0x00000020,
    VK_SHADER_STAGE_ALL_GRAPHICS = 0x0000001F,
    VK_SHADER_STAGE_ALL = 0x7FFFFFFFF,
} VkShaderStageFlagBits;
```

- module is a VkShaderModule object that contains the shader for this stage.
- pName is a pointer to a null-terminated UTF-8 string specifying the entry point name of the shader for this stage.
- pSpecializationInfo is a pointer to VkSpecializationInfo, as described in Specialization Constants, and can be NULL.

- sType must be VK_STRUCTURE_TYPE_PIPELINE_SHADER_STAGE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- stage must be a valid VkShaderStageFlagBits value
- module must be a valid VkShaderModule handle
- pName must be a null-terminated string
- If pSpecializationInfo is not NULL, pSpecializationInfo must be a pointer to a valid VkSpecializationInfo structure
- If the geometry shaders feature is not enabled, stage must not be VK_SHADER_STAGE_GEOMETRY_BIT
- If the tessellation shaders feature is not enabled, <code>stage</code> must not be <code>VK_SHADER_STAGE_TESSELLATION_CONTROL_BIT</code> or <code>VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT</code>

- stage must not be VK_SHADER_STAGE_ALL_GRAPHICS, or VK_SHADER_STAGE_ALL
- pName must be the name of an OpEntryPoint in module with an execution model that matches stage
- If the identified entry point includes any variable in its interface that is declared with the **ClipDistance BuiltIn** decoration, that variable must not have an array size greater than

 VkPhysicalDeviceLimits::maxClipDistances
- If the identified entry point includes any variable in its interface that is declared with the **CullDistance BuiltIn** decoration, that variable must not have an array size greater than

 VkPhysicalDeviceLimits::maxCullDistances
- If the identified entry point includes any variables in its interface that are declared with the ClipDistance or CullDistance BuiltIn decoration, those variables must not have array sizes which sum to more than VkPhysicalDeviceLimits::maxCombinedClipAndCullDistances
- If the identified entry point includes any variable in its interface that is declared with the **SampleMask BuiltIn** decoration, that variable must not have an array size greater than VkPhysicalDeviceLimits::maxSampleMaskWords
- If stage is VK_SHADER_STAGE_VERTEX_BIT, the identified entry point must not include any input variable in its interface that is decorated with **CullDistance**
- If stage is VK_SHADER_STAGE_TESSELLATION_CONTROL_BIT or VK_SHADER_STAGE_ TESSELLATION_EVALUATION_BIT, and the identified entry point has an **OpExecutionMode** instruction that specifies a patch size with **OutputVertices**, the patch size must be greater than 0 and less than or equal to VkPhysicalDeviceLimits::maxTessellationPatchSize
- If stage is VK_SHADER_STAGE_GEOMETRY_BIT, the identified entry point must have an OpExecutionMode instruction that specifies a maximum output vertex count that is greater than 0 and less than or equal to VkPhysicalDeviceLimits::maxGeometryOutputVertices
- If stage is VK_SHADER_STAGE_GEOMETRY_BIT, the identified entry point must have an **OpExecutionMode** instruction that specifies an invocation count that is greater than 0 and less than or equal to VkPhysicalDeviceLimits::maxGeometryShaderInvocations
- If stage is VK_SHADER_STAGE_GEOMETRY_BIT, and the identified entry point writes to **Layer** for any primitive, it must write the same value to **Layer** for all vertices of a given primitive
- If stage is VK_SHADER_STAGE_GEOMETRY_BIT, and the identified entry point writes to **ViewportIndex** for any primitive, it must write the same value to **ViewportIndex** for all vertices of a given primitive
- If stage is VK_SHADER_STAGE_FRAGMENT_BIT, the identified entry point must not include any output variables in its interface decorated with **CullDistance**
- If stage is VK_SHADER_STAGE_FRAGMENT_BIT, and the identified entry point writes to **FragDepth** in any execution path, it must write to **FragDepth** in all execution paths

9.2 Graphics Pipelines

Graphics pipelines consist of multiple shader stages, multiple fixed-function pipeline stages, and a pipeline layout. To create graphics pipelines, call:

- device is the logical device that creates the graphics pipelines.
- pipelineCache is either VK_NULL_HANDLE, indicating that pipeline caching is disabled; or the handle of a valid pipeline cache object, in which case use of that cache is enabled for the duration of the command.
- createInfoCount is the length of the pCreateInfos and pPipelines arrays.
- pCreateInfos is an array of VkGraphicsPipelineCreateInfo structures.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pPipelines is a pointer to an array in which the resulting graphics pipeline objects are returned.

The VkGraphicsPipelineCreateInfo structure includes an array of shader create info structures containing all the desired active shader stages, as well as creation info to define all relevant fixed-function stages, and a pipeline layout.

- device must be a valid VkDevice handle
- If pipelineCache is not VK_NULL_HANDLE, pipelineCache must be a valid VkPipelineCache handle
- pCreateInfos must be a pointer to an array of createInfoCount valid VkGraphicsPipelineCreateInfo structures
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pPipelines must be a pointer to an array of createInfoCount VkPipeline handles
- createInfoCount must be greater than 0
- If pipelineCache is a valid handle, it must have been created, allocated, or retrieved from device
- If the flags member of any given element of pCreateInfos contains the VK_PIPELINE_CREATE_ DERIVATIVE_BIT flag, and the basePipelineIndex member of that same element is not -1, basePipelineIndex must be less than the index into pCreateInfos that corresponds to that element

Return Codes

Success

• VK SUCCESS

Failure

- VK ERROR OUT OF HOST MEMORY
- VK ERROR OUT OF DEVICE MEMORY
- VK_ERROR_INVALID_SHADER_NV

The VkGraphicsPipelineCreateInfo structure is defined as:

```
typedef struct VkGraphicsPipelineCreateInfo {
   VkStructureType
                                                    sType;
   const void*
                                                    pNext;
   VkPipelineCreateFlags
                                                    flags;
   uint32_t
                                                    stageCount;
   const VkPipelineShaderStageCreateInfo*
                                                    pStages;
   const VkPipelineVertexInputStateCreateInfo*
                                                    pVertexInputState;
   const VkPipelineInputAssemblyStateCreateInfo* pInputAssemblyState;
   const VkPipelineTessellationStateCreateInfo*
                                                   pTessellationState;
                                                   pViewportState;
   const VkPipelineViewportStateCreateInfo*
   const VkPipelineRasterizationStateCreateInfo* pRasterizationState;
   const VkPipelineMultisampleStateCreateInfo*
                                                   pMultisampleState;
   const VkPipelineDepthStencilStateCreateInfo*
                                                   pDepthStencilState;
   const VkPipelineColorBlendStateCreateInfo*
                                                   pColorBlendState;
   const VkPipelineDynamicStateCreateInfo*
                                                    pDynamicState;
   VkPipelineLayout
                                                    layout;
   VkRenderPass
                                                    renderPass;
   uint32_t
                                                    subpass;
   VkPipeline
                                                    basePipelineHandle;
   int32_t
                                                    basePipelineIndex;
} VkGraphicsPipelineCreateInfo;
```

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is a bitmask of VkPipelineCreateFlagBits controlling how the pipeline will be generated, as described below.
- stageCount is the number of entries in the pStages array.
- pStages is an array of size stageCount structures of type VkPipelineShaderStageCreateInfo describing the set of the shader stages to be included in the graphics pipeline.
- pVertexInputState is a pointer to an instance of the VkPipelineVertexInputStateCreateInfo structure.
- pInputAssemblyState is a pointer to an instance of the VkPipelineInputAssemblyStateCreateInfo structure which determines input assembly behavior, as described in Drawing Commands.

- pTessellationState is a pointer to an instance of the VkPipelineTessellationStateCreateInfo structure, or NULL if the pipeline does not include a tessellation control shader stage and tessellation evaluation shader stage.
- pViewportState is a pointer to an instance of the VkPipelineViewportStateCreateInfo structure, or NULL if the pipeline has rasterization disabled.
- pRasterizationState is a pointer to an instance of the VkPipelineRasterizationStateCreateInfo structure.
- pMultisampleState is a pointer to an instance of the VkPipelineMultisampleStateCreateInfo, or NULL if the pipeline has rasterization disabled.
- pDepthStencilState is a pointer to an instance of the VkPipelineDepthStencilStateCreateInfo structure, or NULL if the pipeline has rasterization disabled or if the subpass of the render pass the pipeline is created against does not use a depth/stencil attachment.
- pColorBlendState is a pointer to an instance of the VkPipelineColorBlendStateCreateInfo structure, or NULL if the pipeline has rasterization disabled or if the subpass of the render pass the pipeline is created against does not use any color attachments.
- pDynamicState is a pointer to VkPipelineDynamicStateCreateInfo and is used to indicate which properties of the pipeline state object are dynamic and can be changed independently of the pipeline state. This can be NULL, which means no state in the pipeline is considered dynamic.
- layout is the description of binding locations used by both the pipeline and descriptor sets used with the pipeline.
- renderPass is a handle to a render pass object describing the environment in which the pipeline will be used; the pipeline can be used with an instance of any render pass compatible with the one provided. See Render Pass Compatibility for more information.
- subpass is the index of the subpass in renderPass where this pipeline will be used.
- basePipelineHandle is a pipeline to derive from.
- basePipelineIndex is an index into the pCreateInfos parameter to use as a pipeline to derive from.

The parameters basePipelineHandle and basePipelineIndex are described in more detail in Pipeline Derivatives.

pStages points to an array of VkPipelineShaderStageCreateInfo structures, which were previously described in Compute Pipelines.

Bits which can be set in flags are:

```
typedef enum VkPipelineCreateFlagBits {
    VK_PIPELINE_CREATE_DISABLE_OPTIMIZATION_BIT = 0x00000001,
    VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT = 0x00000002,
    VK_PIPELINE_CREATE_DERIVATIVE_BIT = 0x00000004,
} VkPipelineCreateFlagBits;
```

- VK_PIPELINE_CREATE_DISABLE_OPTIMIZATION_BIT specifies that the created pipeline will not be optimized. Using this flag may reduce the time taken to create the pipeline.
- VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT specifies that the pipeline to be created is allowed to be the parent of a pipeline that will be created in a subsequent call to vkCreateGraphicsPipelines.
- VK_PIPELINE_CREATE_DERIVATIVE_BIT specifies that the pipeline to be created will be a child of a previously created parent pipeline.

It is valid to set both VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT and VK_PIPELINE_CREATE_DERIVATIVE_BIT. This allows a pipeline to be both a parent and possibly a child in a pipeline hierarchy. See Pipeline Derivatives for more information.

pDynamicState points to a structure of type VkPipelineDynamicStateCreateInfo.

If any shader stage fails to compile, the compile log will be reported back to the application, and VK_ERROR_INVALID_SHADER_NV will be generated.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_GRAPHICS_PIPELINE_CREATE_INFO
- pNext must be NULL
- flags must be a valid combination of VkPipelineCreateFlagBits values
- pStages must be a pointer to an array of stageCount valid VkPipelineShaderStageCreateInfo structures
- pVertexInputState must be a pointer to a valid VkPipelineVertexInputStateCreateInfo structure
- pInputAssemblyState must be a pointer to a valid VkPipelineInputAssemblyStateCreateInfo structure
- ullet pRasterizationState must be a pointer to a valid VkPipelineRasterizationStateCreateInfo structure
- If pDynamicState is not NULL, pDynamicState must be a pointer to a valid VkPipelineDynamicStateCreateInfo structure
- layout must be a valid VkPipelineLayout handle
- renderPass must be a valid VkRenderPass handle
- stageCount must be greater than 0
- Each of basePipelineHandle, layout, and renderPass that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineIndex is not 1, basePipelineHandle must be VK_NULL_HANDLE
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineIndex is not 1, it must be a valid index into the calling command's pCreateInfos parameter
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, basePipelineIndex must be -1
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, basePipelineHandle must be a valid VkPipeline handle
- If flags contains the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag, and basePipelineHandle is not VK_NULL_HANDLE, it must be a valid handle to a graphics VkPipeline

- The stage member of each element of pStages must be unique
- The stage member of one element of pStages must be VK_SHADER_STAGE_VERTEX_BIT
- The stage member of any given element of pStages must not be VK SHADER STAGE COMPUTE BIT
- If pStages includes a tessellation control shader stage, it must include a tessellation evaluation shader stage
- If pStages includes a tessellation evaluation shader stage, it must include a tessellation control shader stage
- If pStages includes a tessellation control shader stage and a tessellation evaluation shader stage, pTessellationState must not be NULL
- If pStages includes tessellation shader stages, the shader code of at least one stage must contain an **OpExecutionMode** instruction that specifies the type of subdivision in the pipeline
- If pStages includes tessellation shader stages, and the shader code of both stages contain an **OpExecutionMode** instruction that specifies the type of subdivision in the pipeline, they must both specify the same subdivision mode
- If pStages includes tessellation shader stages, the shader code of at least one stage must contain an **OpExecutionMode** instruction that specifies the output patch size in the pipeline
- If pStages includes tessellation shader stages, and the shader code of both contain an **OpExecutionMode** instruction that specifies the out patch size in the pipeline, they must both specify the same patch size
- If pStages includes tessellation shader stages, the topology member of pInputAssembly must be VK_ PRIMITIVE TOPOLOGY PATCH LIST
- If the topology member of pInputAssembly is VK_PRIMITIVE_TOPOLOGY_PATCH_LIST, pStages must include tessellation shader stages
- If pStages includes a geometry shader stage, and does not include any tessellation shader stages, its shader code must contain an **OpExecutionMode** instruction that specifies an input primitive type that is compatible with the primitive topology specified in pInputAssembly
- If pStages includes a geometry shader stage, and also includes tessellation shader stages, its shader code must contain an OpExecutionMode instruction that specifies an input primitive type that is compatible with the primitive topology that is output by the tessellation stages
- If pStages includes a fragment shader stage and a geometry shader stage, and the fragment shader code reads from an input variable that is decorated with **PrimitiveID**, then the geometry shader code must write to a matching output variable, decorated with **PrimitiveID**, in all execution paths
- If pStages includes a fragment shader stage, its shader code must not read from any input attachment that is defined as VK_ATTACHMENT_UNUSED in subpass
- The shader code for the entry points identified by pStages, and the rest of the state identified by this structure must adhere to the pipeline linking rules described in the Shader Interfaces chapter
- If subpass uses a depth/stencil attachment in renderpass that has a layout of VK_IMAGE_LAYOUT_DEPTH_ STENCIL_READ_ONLY_OPTIMAL in the VkAttachmentReference defined by subpass, and pDepthStencilState is not NULL, the depthWriteEnable member of pDepthStencilState must be VK_FALSE

- If subpass uses a depth/stencil attachment in renderpass that has a layout of VK_IMAGE_LAYOUT_DEPTH_ STENCIL_READ_ONLY_OPTIMAL in the VkAttachmentReference defined by subpass, and pDepthStencilState is not NULL, the failOp, passOp and depthFailOp members of each of the front and back members of pDepthStencilState must be VK_STENCIL_OP_KEEP
- If pColorBlendState is not NULL, the blendEnable member of each element of the pAttachment member of pColorBlendState must be VK_FALSE if the format of the attachment referred to in subpass of renderPass does not support color blend operations, as specified by the VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BLEND_BIT flag in VkFormatProperties::linearTilingFeatures or VkFormatProperties::optimalTilingFeatures returned by

${\tt vkGetPhysicalDeviceFormatProperties}$

- If pColorBlendState is not NULL, The attachmentCount member of pColorBlendState must be equal to the colorAttachmentCount used to create subpass
- If no element of the pDynamicStates member of pDynamicState is VK_DYNAMIC_STATE_VIEWPORT, the pViewports member of pViewportState must be a pointer to an array of pViewportState —viewportCount VkViewport structures
- If no element of the pDynamicStates member of pDynamicState is VK_DYNAMIC_STATE_SCISSOR, the pScissors member of pViewportState must be a pointer to an array of pViewportState —>scissorCount VkRect2D structures
- If the wide lines feature is not enabled, and no element of the pDynamicStates member of pDynamicState is VK_DYNAMIC_STATE_LINE_WIDTH, the lineWidth member of pRasterizationState must be 1.0
- If the rasterizerDiscardEnable member of pRasterizationState is VK_FALSE, pViewportState must be a pointer to a valid VkPipelineViewportStateCreateInfo structure
- If the rasterizerDiscardEnable member of pRasterizationState is VK_FALSE, pMultisampleState must be a pointer to a valid VkPipelineMultisampleStateCreateInfo structure
- If the rasterizerDiscardEnable member of pRasterizationState is VK_FALSE, and subpass uses a depth/stencil attachment, pDepthStencilState must be a pointer to a valid VkPipelineDepthStencilStateCreateInfo structure
- If the rasterizerDiscardEnable member of pRasterizationState is VK_FALSE, and subpass uses color attachments, pColorBlendState must be a pointer to a valid VkPipelineColorBlendStateCreateInfo structure
- If the depth bias clamping feature is not enabled, no element of the pDynamicStates member of pDynamicState is VK_DYNAMIC_STATE_DEPTH_BIAS, and the depthBiasEnable member of pDepthStencil is VK_TRUE, the depthBiasClamp member of pDepthStencil must be 0.0
- If no element of the pDynamicStates member of pDynamicState is VK_DYNAMIC_STATE_DEPTH_ BOUNDS, and the depthBoundsTestEnable member of pDepthStencil is VK_TRUE, the minDepthBounds and maxDepthBounds members of pDepthStencil must be between 0.0 and 1.0, inclusive
- layout must be consistent with all shaders specified in pStages
- If subpass uses color and/or depth/stencil attachments, then the rasterizationSamples member of pMultisampleState must be the same as the sample count for those subpass attachments
- If subpass does not use any color and/or depth/stencil attachments, then the rasterizationSamples member of pMultisampleState must follow the rules for a zero-attachment subpass
- subpass must be a valid subpass within renderpass

The VkPipelineDynamicStateCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- dynamicStateCount is the number of elements in the pDynamicStates array.
- pDynamicStates is an array of VkDynamicState enums which indicate which pieces of pipeline state will use the values from dynamic state commands rather than from the pipeline state creation info.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_PIPELINE_DYNAMIC_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- pDynamicStates must be a pointer to an array of dynamicStateCount valid VkDynamicState values
- dynamicStateCount must be greater than 0

The source of difference pieces of dynamic state is determined by the

VkPipelineDynamicStateCreateInfo::pDynamicStates property of the currently active pipeline, which takes the following values:

```
typedef enum VkDynamicState {
    VK_DYNAMIC_STATE_VIEWPORT = 0,
    VK_DYNAMIC_STATE_SCISSOR = 1,
    VK_DYNAMIC_STATE_LINE_WIDTH = 2,
    VK_DYNAMIC_STATE_DEPTH_BIAS = 3,
    VK_DYNAMIC_STATE_DEPTH_BOUNDS = 5,
    VK_DYNAMIC_STATE_DEPTH_BOUNDS = 5,
    VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK = 6,
    VK_DYNAMIC_STATE_STENCIL_WRITE_MASK = 7,
    VK_DYNAMIC_STATE_STENCIL_REFERENCE = 8,
} VkDynamicState;
```

- VK_DYNAMIC_STATE_VIEWPORT indicates that the *pViewports* state in VkPipelineViewportStateCreateInfo will be ignored and must be set dynamically with vkCmdSetViewport before any draw commands. The number of viewports used by a pipeline is still specified by the *viewportCount* member of VkPipelineViewportStateCreateInfo.
- VK_DYNAMIC_STATE_SCISSOR indicates that the pscissors state in VkPipelineViewportStateCreateInfo will be ignored and must be set dynamically with vkCmdSetScissor before any draw commands. The number of scissor rectangles used by a pipeline is still specified by the scissorCount member of VkPipelineViewportStateCreateInfo.
- VK_DYNAMIC_STATE_LINE_WIDTH indicates that the <code>lineWidth</code> state in VkPipelineRasterizationStateCreateInfo will be ignored and must be set dynamically with vkCmdSetLineWidth before any draw commands that generate line primitives for the rasterizer.
- VK_DYNAMIC_STATE_DEPTH_BIAS indicates that the depthBiasConstantFactor, depthBiasClamp and depthBiasSlopeFactor states in VkPipelineRasterizationStateCreateInfo will be ignored and must be set dynamically with vkCmdSetDepthBias before any draws are performed with depthBiasEnable in VkPipelineRasterizationStateCreateInfo set to VK_TRUE.
- VK_DYNAMIC_STATE_BLEND_CONSTANTS indicates that the <code>blendConstants</code> state in VkPipelineColorBlendStateCreateInfo will be ignored and must be set dynamically with vkCmdSetBlendConstants before any draws are performed with a pipeline state with VkPipelineColorBlendAttachmentState member <code>blendEnable</code> set to VK_TRUE and any of the blend functions using a constant blend color.
- VK_DYNAMIC_STATE_DEPTH_BOUNDS indicates that the minDepthBounds and maxDepthBounds states of VkPipelineDepthStencilStateCreateInfo will be ignored and must be set dynamically with vkCmdSetDepthBounds before any draws are performed with a pipeline state with VkPipelineDepthStencilStateCreateInfo member depthBoundsTestEnable set to VK_TRUE.
- VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK indicates that the <code>compareMask</code> state in VkPipelineDepthStencilStateCreateInfo for both <code>front</code> and <code>back</code> will be ignored and must be set dynamically with <code>vkCmdSetStencilCompareMask</code> before any draws are performed with a pipeline state with VkPipelineDepthStencilStateCreateInfo member <code>stencilTestEnable</code> set to VK_TRUE
- VK_DYNAMIC_STATE_STENCIL_WRITE_MASK indicates that the writeMask state in VkPipelineDepthStencilStateCreateInfo for both front and back will be ignored and must be set dynamically with vkCmdSetStencilWriteMask before any draws are performed with a pipeline state with VkPipelineDepthStencilStateCreateInfo member stencilTestEnable set to VK_TRUE
- VK_DYNAMIC_STATE_STENCIL_REFERENCE indicates that the reference state in VkPipelineDepthStencilStateCreateInfo for both front and back will be ignored and must be set dynamically with vkCmdSetStencilReference before any draws are performed with a pipeline state with VkPipelineDepthStencilStateCreateInfo member stencilTestEnable set to VK_TRUE

9.2.1 Valid Combinations of Stages for Graphics Pipelines

If tessellation shader stages are omitted, the tessellation shading and fixed-function stages of the pipeline are skipped.

If a geometry shader is omitted, the geometry shading stage is skipped.

If a fragment shader is omitted, the results of fragment processing are undefined. Specifically, any fragment color outputs are considered to have undefined values, and the fragment depth is considered to be unmodified. This can be useful for depth-only rendering.

Presence of a shader stage in a pipeline is indicated by including a valid VkPipelineShaderStageCreateInfo with module and pName selecting an entry point from a shader module, where that entry point is valid for the stage specified by stage.

Presence of some of the fixed-function stages in the pipeline is implicitly derived from enabled shaders and provided state. For example, the fixed-function tessellator is always present when the pipeline has valid Tessellation Control and Tessellation Evaluation shaders.

FOR EXAMPLE:

- Depth/stencil-only rendering in a subpass with no color attachments
 - Active Pipeline Shader Stages
 - * Vertex Shader
 - Required: Fixed-Function Pipeline Stages
 - * VkPipelineVertexInputStateCreateInfo
 - * VkPipelineInputAssemblyStateCreateInfo
 - * VkPipelineViewportStateCreateInfo
 - * VkPipelineRasterizationStateCreateInfo
 - * VkPipelineMultisampleStateCreateInfo
 - * VkPipelineDepthStencilStateCreateInfo
- Color-only rendering in a subpass with no depth/stencil attachment
 - Active Pipeline Shader Stages
 - * Vertex Shader
 - * Fragment Shader
 - Required: Fixed-Function Pipeline Stages
 - * VkPipelineVertexInputStateCreateInfo
 - * VkPipelineInputAssemblyStateCreateInfo
 - * VkPipelineViewportStateCreateInfo
 - * VkPipelineRasterizationStateCreateInfo
 - * VkPipelineMultisampleStateCreateInfo
 - * VkPipelineColorBlendStateCreateInfo
- Rendering pipeline with tessellation and geometry shaders
 - Active Pipeline Shader Stages
 - * Vertex Shader
 - * Tessellation Control Shader
 - * Tessellation Evaluation Shader
 - * Geometry Shader
 - * Fragment Shader
 - Required: Fixed-Function Pipeline Stages
 - * VkPipelineVertexInputStateCreateInfo
 - * VkPipelineInputAssemblyStateCreateInfo
 - * VkPipelineTessellationStateCreateInfo
 - * VkPipelineViewportStateCreateInfo
 - * VkPipelineRasterizationStateCreateInfo
 - * VkPipelineMultisampleStateCreateInfo
 - * VkPipelineDepthStencilStateCreateInfo
 - * VkPipelineColorBlendStateCreateInfo

9.3 Pipeline destruction

To destroy a graphics or compute pipeline, call:

- device is the logical device that destroys the pipeline.
- pipeline is the handle of the pipeline to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If pipeline is not VK_NULL_HANDLE, pipeline must be a valid VkPipeline handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If pipeline is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to pipeline must have completed execution
- If VkAllocationCallbacks were provided when pipeline was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when pipeline was created, pAllocator must be NULL

Host Synchronization

• Host access to pipeline must be externally synchronized

9.4 Multiple Pipeline Creation

Multiple pipelines can be created simultaneously by passing an array of VkGraphicsPipelineCreateInfo or VkComputePipelineCreateInfo structures into the vkCreateGraphicsPipelines and

vkCreateComputePipelines commands, respectively. Applications can group together similar pipelines to be created in a single call, and implementations are encouraged to look for reuse opportunities within a group-create.

When an application attempts to create many pipelines in a single command, it is possible that some subset may fail creation. In that case, the corresponding entries in the <code>pPipelines</code> output array will be filled with <code>VK_NULL_HANDLE</code> values. If any pipeline fails creation (for example, due to out of memory errors), the <code>vkCreate*Pipelines</code> commands will return an error code. The implementation will attempt to create all pipelines, and only return <code>VK_NULL_HANDLE</code> values for those that actually failed.

9.5 Pipeline Derivatives

A pipeline derivative is a child pipeline created from a parent pipeline, where the child and parent are expected to have much commonality. The goal of derivative pipelines is that they be cheaper to create using the parent as a starting point, and that it be more efficient (on either host or device) to switch/bind between children of the same parent.

A derivative pipeline is created by setting the VK_PIPELINE_CREATE_DERIVATIVE_BIT flag in the Vk*PipelineCreateInfo structure. If this is set, then exactly one of basePipelineHandle or basePipelineIndex members of the structure must have a valid handle/index, and indicates the parent pipeline. If basePipelineHandle is used, the parent pipeline must have already been created. If basePipelineIndex is used, then the parent is being created in the same command. VK_NULL_HANDLE acts as the invalid handle for basePipelineHandle, and -1 is the invalid index for basePipelineIndex. If basePipelineIndex is used, the base pipeline must appear earlier in the array. The base pipeline must have been created with the VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT flag set.

9.6 Pipeline Cache

Pipeline cache objects allow the result of pipeline construction to be reused between pipelines and between runs of an application. Reuse between pipelines is achieved by passing the same pipeline cache object when creating multiple related pipelines. Reuse across runs of an application is achieved by retrieving pipeline cache contents in one run of an application, saving the contents, and using them to preinitialize a pipeline cache on a subsequent run. The contents of the pipeline cache objects are managed by the implementation. Applications can manage the host memory consumed by a pipeline cache object and control the amount of data retrieved from a pipeline cache object.

Pipeline cache objects are represented by VkPipelineCache handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkPipelineCache)
```

To create pipeline cache objects, call:

- device is the logical device that creates the pipeline cache object.
- pCreateInfo is a pointer to a VkPipelineCacheCreateInfo structure that contains the initial parameters for the pipeline cache object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

• pPipelineCache is a pointer to a VkPipelineCache handle in which the resulting pipeline cache object is returned.



Note

Applications can track and manage the total host memory size of a pipeline cache object using the <code>pAllocator</code>. Applications can limit the amount of data retrieved from a pipeline cache object in <code>vkGetPipelineCacheD</code> ata. Implementations should not internally limit the total number of entries added to a pipeline cache object or the total host memory consumed.

Once created, a pipeline cache can be passed to the **vkCreateGraphicsPipelines** and **vkCreateComputePipelines** commands. If the pipeline cache passed into these commands is not VK_NULL_HANDLE, the implementation will query it for possible reuse opportunities and update it with new content. The use of the pipeline cache object in these commands is internally synchronized, and the same pipeline cache object can be used in multiple threads simultaneously.



Note

Implementations should make every effort to limit any critical sections to the actual accesses to the cache, which is expected to be significantly shorter than the duration of the **vkCreateGraphicsPipelines** and **vkCreateComputePipelines** commands.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkPipelineCacheCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pPipelineCache must be a pointer to a VkPipelineCache handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkPipelineCacheCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- initialDataSize is the number of bytes in pInitialData. If initialDataSize is zero, the pipeline cache will initially be empty.
- pInitialData is a pointer to previously retrieved pipeline cache data. If the pipeline cache data is incompatible (as defined below) with the device, the pipeline cache will be initially empty. If initialDataSize is zero, pInitialData is ignored.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_PIPELINE_CACHE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- If initialDataSize is not 0, pInitialData must be a pointer to an array of initialDataSize bytes
- If initialDataSize is not 0, it must be equal to the size of pInitialData, as returned by vkGetPipelineCacheData when pInitialData was originally retrieved
- If initialDataSize is not 0, pInitialData must have been retrieved from a previous call to vkGetPipelineCacheData

Pipeline cache objects can be merged using the command:

- device is the logical device that owns the pipeline cache objects.
- dstCache is the handle of the pipeline cache to merge results into.

- srcCacheCount is the length of the pSrcCaches array.
- pSrcCaches is an array of pipeline cache handles, which will be merged into dstCache. The previous contents of dstCache are included after the merge.



Note

The details of the merge operation are implementation dependent, but implementations should merge the contents of the specified pipelines and prune duplicate entries.

Valid Usage

- device must be a valid VkDevice handle
- dstCache must be a valid VkPipelineCache handle
- pSrcCaches must be a pointer to an array of srcCacheCount valid VkPipelineCache handles
- srcCacheCount must be greater than 0
- dstCache must have been created, allocated, or retrieved from device
- Each element of pSrcCaches must have been created, allocated, or retrieved from device
- dstCache must not appear in the list of source caches

Host Synchronization

• Host access to dstCache must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

Data can be retrieved from a pipeline cache object using the command:

- device is the logical device that owns the pipeline cache.
- pipelineCache is the pipeline cache to retrieve data from.
- pDataSize is a pointer to a value related to the amount of data in the pipeline cache, as described below.
- pData is either NULL or a pointer to a buffer.

If pData is NULL, then the maximum size of the data that can be retrieved from the pipeline cache, in bytes, is returned in pDataSize. Otherwise, pDataSize must point to a variable set by the user to the size of the buffer, in bytes, pointed to by pData, and on return the variable is overwritten with the amount of data actually written to pData.

If pDataSize is less than the maximum size that can be retrieved by the pipeline cache, at most pDataSize bytes will be written to pData, and vkGetPipelineCacheData will return VK_INCOMPLETE. Any data written to pData is valid and can be provided as the pInitialData member of the VkPipelineCacheCreateInfo structure passed to vkCreatePipelineCache.

Two calls to **vkGetPipelineCacheData** with the same parameters must retrieve the same data unless a command that modifies the contents of the cache is called between them.

Applications can store the data retrieved from the pipeline cache, and use these data, possibly in a future run of the application, to populate new pipeline cache objects. The results of pipeline compiles, however, may depend on the vendor ID, device ID, driver version, and other details of the device. To enable applications to detect when previously retrieved data is incompatible with the device, the initial bytes written to pData must be a header consisting of the following members:

Table 9.1: Layout for pipeline cache header version VK_PIPELINE_CACHE_HEADER_VERSION_ONE

Offset	Size	Meaning			
0	4	length in bytes of the entire pipeline cache header written as a			
		stream of bytes, with the least significant byte first			
4	4	a VkPipelineCacheHeaderVersion value written as a			
		stream of bytes, with the least significant byte first			
8	4	a vendor ID equal to			
		VkPhysicalDeviceProperties::vendorID written as a			
		stream of bytes, with the least significant byte first			
12	4	a device ID equal to			
		VkPhysicalDeviceProperties::deviceID written as a			
		stream of bytes, with the least significant byte first			
16	VK_UUID_SIZE	a pipeline cache ID equal to			
		VkPhysicalDeviceProperties::pipelineCacheUUID			

The first four bytes encode the length of the entire pipeline header, in bytes. This value includes all fields in the header including the pipeline cache version field and the size of the length field.

The next four bytes encode the pipeline cache version. This field is interpreted as a VkPipelineCacheHeaderVersion value, and must have one of the following values:

```
typedef enum VkPipelineCacheHeaderVersion {
    VK_PIPELINE_CACHE_HEADER_VERSION_ONE = 1,
} VkPipelineCacheHeaderVersion;
```

A consumer of the pipeline cache should use the cache version to interpret the remainder of the cache header.

If pDataSize is less than what is necessary to store this header, nothing will be written to pData and zero will be written to pDataSize.

Valid Usage

- device must be a valid VkDevice handle
- pipelineCache must be a valid VkPipelineCache handle
- pDataSize must be a pointer to a size_t value
- If the value referenced by pDataSize is not 0, and pData is not NULL, pData must be a pointer to an array of pDataSize bytes
- pipelineCache must have been created, allocated, or retrieved from device

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To destroy a pipeline cache, call:

- device is the logical device that destroys the pipeline cache object.
- pipelineCache is the handle of the pipeline cache to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If pipelineCache is not VK_NULL_HANDLE, pipelineCache must be a valid VkPipelineCache handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If pipelineCache is a valid handle, it must have been created, allocated, or retrieved from device
- If VkAllocationCallbacks were provided when <code>pipelineCache</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when pipelineCache was created, pAllocator must be NULL

Host Synchronization

• Host access to pipelineCache must be externally synchronized

9.7 Specialization Constants

Specialization constants are a mechanism whereby constants in a SPIR-V module can have their constant value specified at the time the VkPipeline is created. This allows a SPIR-V module to have constants that can be modified while executing an application that uses the Vulkan API.



Note

Specialization constants are useful to allow a compute shader to have its local workgroup size changed at runtime by the user, for example.

Each instance of the VkPipelineShaderStageCreateInfo structure contains a parameter pSpecializationInfo, which can be NULL to indicate no specialization constants, or point to a VkSpecializationInfo structure.

The VkSpecializationInfo structure is defined as:

- mapEntryCount is the number of entries in the pMapEntries array.
- pMapEntries is a pointer to an array of VkSpecializationMapEntry which maps constant IDs to offsets in pData.
- dataSize is the byte size of the pData buffer.
- pData contains the actual constant values to specialize with.

pMapEntries points to a structure of type VkSpecializationMapEntry.

Valid Usage

- If mapEntryCount is not 0, pMapEntries must be a pointer to an array of mapEntryCount valid VkSpecializationMapEntry structures
- If dataSize is not 0, pData must be a pointer to an array of dataSize bytes
- The offset member of any given element of pMapEntries must be less than dataSize
- For any given element of pMapEntries, size must be less than or equal to dataSize minus offset

The VkSpecializationMapEntry structure is defined as:

```
typedef struct VkSpecializationMapEntry {
    uint32_t constantID;
    uint32_t offset;
    size_t size;
} VkSpecializationMapEntry;
```

- constant ID is the ID of the specialization constant in SPIR-V.
- offset is the byte offset of the specialization constant value within the supplied data buffer.
- size is the byte size of the specialization constant value within the supplied data buffer.

If a constant ID value is not a specialization constant ID used in the shader, that map entry does not affect the behavior of the pipeline.

Valid Usage

• For a constant ID specialization constant declared in a shader, size must match the byte size of the constant ID. If the specialization constant is of type **boolean**, size must be the byte size of VkBool32

In human readable SPIR-V:

```
OpDecorate %x SpecId 13; decorate .x component of WorkgroupSize with ID 13
OpDecorate %y SpecId 42; decorate .y component of WorkgroupSize with ID 42
OpDecorate %z SpecId 3; decorate .z component of WorkgroupSize with ID 3
OpDecorate %wgsize BuiltIn WorkgroupSize; decorate WorkgroupSize onto constant
%i32 = OpTypeInt 32 0; declare an unsigned 32-bit type
%uvec3 = OpTypeVector %i32 3; declare a 3 element vector type of unsigned 32-bit
%x = OpSpecConstant %i32 1; declare the .x component of WorkgroupSize
%y = OpSpecConstant %i32 1; declare the .y component of WorkgroupSize
%z = OpSpecConstant %i32 1; declare the .z component of WorkgroupSize
%wgsize = OpSpecConstantComposite %uvec3 %x %y %z; declare WorkgroupSize
```

From the above we have three specialization constants, one for each of the x, y & z elements of the WorkgroupSize vector.

Now to specialize the above via the specialization constants mechanism:

```
const VkSpecializationMapEntry entries[] =
{
                                      // constantID
       13,
       0 * sizeof(uint32_t),
                                      // offset
       sizeof(uint32_t)
                                      // size
   },
                                     // constantID
       42,
       1 * sizeof(uint32_t),
                                      // offset
                                       // size
       sizeof(uint32_t)
   },
                                  // constantID
// offset
       3,
2 * sizeof(uint32_t),
sizeof(uint32_t)
       sizeof(uint32_t)
                                      // size
   }
};
const uint32_t data[] = { 16, 8, 4 }; // our workgroup size is 16x8x4
const VkSpecializationInfo info =
                                      // mapEntryCount
   3,
                                       // pMapEntries
   entries,
                                       // dataSize
   3 * sizeof(uint32_t),
                                       // pData
   data,
};
```

Then when calling vkCreateComputePipelines, and passing the VkSpecializationInfo we defined as the pSpecializationInfo parameter of VkPipelineShaderStageCreateInfo, we will create a compute pipeline with the runtime specified local workgroup size.

Another example would be that an application has a SPIR-V module that has some platform-dependent constants they wish to use.

In human readable SPIR-V:

```
OpDecorate %1 SpecId 0 ; decorate our signed 32-bit integer constant
OpDecorate %2 SpecId 12 ; decorate our 32-bit floating-point constant
%i32 = OpTypeInt 32 1 ; declare a signed 32-bit type
%float = OpTypeFloat 32 ; declare a 32-bit floating-point type
%1 = OpSpecConstant %i32 -1 ; some signed 32-bit integer constant
%2 = OpSpecConstant %float 0.5 ; some 32-bit floating-point constant
```

From the above we have two specialization constants, one is a signed 32-bit integer and the second is a 32-bit floating-point.

Now to specialize the above via the specialization constants mechanism:

```
struct SpecializationData {
   int32_t data0;
   float data1;
};
const VkSpecializationMapEntry entries[] =
        Ο,
                                              // constantID
       offsetof(SpecializationData, data0), // offset
       sizeof(SpecializationData::data0)
                                             // size
    },
                                             // constantID
       offsetof(SpecializationData, data1), // offset
       sizeof(SpecializationData::data1)
                                             // size
   }
};
SpecializationData data;
data.data0 = -42; // set the data for the 32-bit integer
data.data1 = 42.0f; // set the data for the 32-bit floating-point
const VkSpecializationInfo info =
   2,
                                       // mapEntryCount
   entries,
                                        // pMapEntries
                                        // dataSize
   sizeof (data),
                                        // pData
   &data,
};
```

It is legal for a SPIR-V module with specializations to be compiled into a pipeline where no specialization info was provided. SPIR-V specialization constants contain default values such that if a specialization is not provided, the default value will be used. In the examples above, it would be valid for an application to only specialize some of the specialization constants within the SPIR-V module, and let the other constants use their default values encoded within the OpSpecConstant declarations.

9.8 Pipeline Binding

Once a pipeline has been created, it can be bound to the command buffer using the command:

- commandBuffer is the command buffer that the pipeline will be bound to.
- pipelineBindPoint specifies the bind point, and must have one of the values

```
typedef enum VkPipelineBindPoint {
    VK_PIPELINE_BIND_POINT_GRAPHICS = 0,
    VK_PIPELINE_BIND_POINT_COMPUTE = 1,
} VkPipelineBindPoint;
```

specifying whether <code>pipeline</code> will be bound as a compute (VK_PIPELINE_BIND_POINT_COMPUTE) or graphics (VK_PIPELINE_BIND_POINT_GRAPHICS) pipeline. There are separate bind points for each of graphics and compute, so binding one does not disturb the other.

• pipeline is the pipeline to be bound.

Once bound, a pipeline binding affects subsequent graphics or compute commands in the command buffer until a different pipeline is bound to the bind point. The pipeline bound to VK_PIPELINE_BIND_POINT_COMPUTE controls the behavior of vkCmdDispatch and vkCmdDispatchIndirect. The pipeline bound to VK_PIPELINE_BIND_POINT_GRAPHICS controls the behavior of vkCmdDraw, vkCmdDrawIndexed, vkCmdDrawIndirect, and vkCmdDrawIndexedIndirect. No other commands are affected by the pipeline state.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pipelineBindPoint must be a valid VkPipelineBindPoint value
- pipeline must be a valid VkPipeline handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- Both of commandBuffer, and pipeline must have been created, allocated, or retrieved from the same VkDevice
- If pipelineBindPoint is VK_PIPELINE_BIND_POINT_COMPUTE, the VkCommandPool that commandBuffer was allocated from must support compute operations
- If pipelineBindPoint is VK_PIPELINE_BIND_POINT_GRAPHICS, the VkCommandPool that commandBuffer was allocated from must support graphics operations

- If pipelineBindPoint is VK_PIPELINE_BIND_POINT_COMPUTE, pipeline must be a compute pipeline
- If pipelineBindPoint is VK_PIPELINE_BIND_POINT_GRAPHICS, pipeline must be a graphics pipeline
- If the variable multisample rate feature is not supported, pipeline is a graphics pipeline, the current subpass has no attachments, and this is not the first call to this function with a graphics pipeline after transitioning to the current subpass, then the sample count specified by this pipeline must match that set in the previous pipeline

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

Chapter 10

Memory Allocation

Vulkan memory is broken up into two categories, host memory and device memory.

10.1 Host Memory

Host memory is memory needed by the Vulkan implementation for non-device-visible storage. This storage may be used for e.g. internal software structures.

Vulkan provides applications the opportunity to perform host memory allocations on behalf of the Vulkan implementation. If this feature is not used, the implementation will perform its own memory allocations. Since most memory allocations are off the critical path, this is not meant as a performance feature. Rather, this can be useful for certain embedded systems, for debugging purposes (e.g. putting a guard page after all host allocations), or for memory allocation logging.

Allocators are provided by the application as a pointer to a VkAllocationCallbacks structure:

- pUserData is a value to be interpreted by the implementation of the callbacks. When any of the callbacks in VkAllocationCallbacks are called, the Vulkan implementation will pass this value as the first parameter to the callback. This value can vary each time an allocator is passed into a command, even when the same object takes an allocator in multiple commands.
- pfnAllocation is a pointer to an application-defined memory allocation function of type PFN_vkAllocationFunction.
- pfnReallocation is a pointer to an application-defined memory reallocation function of type PFN_vkReallocationFunction.
- pfnFree is a pointer to an application-defined memory free function of type PFN_vkFreeFunction.

- pfnInternalAllocation is a pointer to an application-defined function that is called by the implementation when the implementation makes internal allocations, and it is of type PFN_vkInternalAllocationNotification.
- pfnInternalFree is a pointer to an application-defined function that is called by the implementation when the implementation frees internal allocations, and it is of type PFN_vkInternalFreeNotification.

Valid Usage

- pfnAllocation must be a pointer to a valid user-defined PFN vkAllocationFunction
- pfnReallocation must be a pointer to a valid user-defined PFN_vkReallocationFunction
- pfnFree must be a pointer to a valid user-defined PFN_vkFreeFunction
- If either of pfnInternalAllocation or pfnInternalFree is not NULL, both must be valid callbacks

The type of pfnAllocation is:

- pUserData is the value specified for VkAllocationCallbacks::pUserData in the allocator specified by the application.
- size is the size in bytes of the requested allocation.
- alignment is the requested alignment of the allocation in bytes and must be a power of two.
- allocationScope is a VkSystemAllocationScope value specifying the scope of the lifetime of the allocation, as described here.

If pfnAllocation is unable to allocate the requested memory, it must return NULL. If the allocation was successful, it must return a valid pointer to memory allocation containing at least size bytes, and with the pointer value being a multiple of alignment.

Note



Correct Vulkan operation cannot be assumed if the application does not follow these rules.

For example, pfnAllocation (or pfnReallocation) could cause termination of running Vulkan instance(s) on a failed allocation for debugging purposes, either directly or indirectly. In these circumstances, it cannot be assumed that any part of any affected VkInstance objects are going to operate correctly (even vkDestroyIn stance), and the application must ensure it cleans up properly via other means (e.g. process termination).

If pfnAllocation returns NULL, and if the implementation is unable to continue correct processing of the current command without the requested allocation, it must treat this as a run-time error, and generate VK_ERROR_OUT_OF_HOST_MEMORY at the appropriate time for the command in which the condition was detected, as described in Return Codes.

If the implementation is able to continue correct processing of the current command without the requested allocation, then it may do so, and must not generate VK_ERROR_OUT_OF_HOST_MEMORY as a result of this failed allocation.

The type of pfnReallocation is:

- pUserData is the value specified for VkAllocationCallbacks::pUserData in the allocator specified by the application.
- pOriginal must be either NULL or a pointer previously returned by pfnReallocation or pfnAllocation of the same allocator.
- size is the size in bytes of the requested allocation.
- alignment is the requested alignment of the allocation in bytes and must be a power of two.
- allocationScope is a VkSystemAllocationScope value specifying the scope of the lifetime of the allocation, as described here.

pfnReallocation must return an allocation with enough space for size bytes, and the contents of the original allocation from bytes zero to min(original size, new size) -1 must be preserved in the returned allocation. If size is larger than the old size, the contents of the additional space are undefined. If satisfying these requirements involves creating a new allocation, then the old allocation should be freed.

If poriginal is NULL, then pfnReallocation must behave equivalently to a call to PFN_vkAllocationFunction with the same parameter values (without poriginal).

If size is zero, then pfnReallocation must behave equivalently to a call to PFN_vkFreeFunction with the same pUserData parameter value, and pMemory equal to pOriginal.

If poriginal is non-NULL, the implementation must ensure that alignment is equal to the alignment used to originally allocate poriginal.

If this function fails and poriginal is non-NULL the application must not free the old allocation.

pfnReallocation must follow the same rules for return values as PFN_vkAllocationFunction.

The type of pfnFree is:

- pUserData is the value specified for VkAllocationCallbacks::pUserData in the allocator specified by the application.
- pMemory is the allocation to be freed.

pMemory may be NULL, which the callback must handle safely. If pMemory is non-NULL, it must be a pointer previously allocated by pfnAllocation or pfnReallocation. The application should free this memory.

The type of pfnInternalAllocation is:

- pUserData is the value specified for VkAllocationCallbacks::pUserData in the allocator specified by the application.
- size is the requested size of an allocation.
- allocationType is the requested type of an allocation.
- allocationScope is a VkSystemAllocationScope value specifying the scope of the lifetime of the allocation, as described here.

This is a purely informational callback.

The type of pfnInternalFree is:

- pUserData is the value specified for VkAllocationCallbacks::pUserData in the allocator specified by the application.
- size is the requested size of an allocation.
- allocationType is the requested type of an allocation.
- allocationScope is a VkSystemAllocationScope value specifying the scope of the lifetime of the allocation, as described here.

Each allocation has a *scope* which defines its lifetime and which object it is associated with. The scope is provided in the *allocationScope* parameter passed to callbacks defined in VkAllocationCallbacks. Possible values for this parameter are defined by VkSystemAllocationScope:

```
typedef enum VkSystemAllocationScope {
   VK_SYSTEM_ALLOCATION_SCOPE_COMMAND = 0,
   VK_SYSTEM_ALLOCATION_SCOPE_OBJECT = 1,
   VK_SYSTEM_ALLOCATION_SCOPE_CACHE = 2,
   VK_SYSTEM_ALLOCATION_SCOPE_DEVICE = 3,
   VK_SYSTEM_ALLOCATION_SCOPE_INSTANCE = 4,
} VkSystemAllocationScope;
```

• VK_SYSTEM_ALLOCATION_SCOPE_COMMAND - The allocation is scoped to the duration of the Vulkan command.

- VK_SYSTEM_ALLOCATION_SCOPE_OBJECT The allocation is scoped to the lifetime of the Vulkan object that is being created or used.
- VK_SYSTEM_ALLOCATION_SCOPE_CACHE The allocation is scoped to the lifetime of a VkPipelineCache object.
- VK_SYSTEM_ALLOCATION_SCOPE_DEVICE The allocation is scoped to the lifetime of the Vulkan device.
- VK_SYSTEM_ALLOCATION_SCOPE_INSTANCE The allocation is scoped to the lifetime of the Vulkan instance.

Most Vulkan commands operate on a single object, or there is a sole object that is being created or manipulated. When an allocation uses a scope of VK_SYSTEM_ALLOCATION_SCOPE_OBJECT or VK_SYSTEM_ALLOCATION_SCOPE_CACHE, the allocation is scoped to the object being created or manipulated.

When an implementation requires host memory, it will make callbacks to the application using the most specific allocator and scope available:

- If an allocation is scoped to the duration of a command, the allocator will use the VK_SYSTEM_ALLOCATION_ SCOPE_COMMAND scope. The most specific allocator available is used: if the object being created or manipulated has an allocator, that object's allocator will be used, else if the parent VkDevice has an allocator it will be used, else if the parent VkInstance has an allocator it will be used. Else,
- If an allocation is associated with an object of type VkPipelineCache, the allocator will use the VK_SYSTEM_ ALLOCATION_SCOPE_CACHE scope. The most specific allocator available is used (pipeline cache, else device, else instance). Else,
- If an allocation is scoped to the lifetime of an object, that object is being created or manipulated by the command, and that object's type is not VkDevice or VkInstance, the allocator will use a scope of VK_SYSTEM_ALLOCATION_SCOPE_OBJECT. The most specific allocator available is used (object, else device, else instance). Else,
- If an allocation is scoped to the lifetime of a device, the allocator will use scope of VK_SYSTEM_ALLOCATION_ SCOPE_DEVICE. The most specific allocator available is used (device, else instance). Else,
- If the allocation is scoped to the lifetime of an instance and the instance has an allocator, its allocator will be used with a scope of VK_SYSTEM_ALLOCATION_SCOPE_INSTANCE.
- · Otherwise an implementation will allocate memory through an alternative mechanism that is unspecified.

Objects that are allocated from pools do not specify their own allocator. When an implementation requires host memory for such an object, that memory is sourced from the object's parent pool's allocator.

The application is not expected to handle allocating memory that is intended for execution by the host due to the complexities of differing security implementations across multiple platforms. The implementation will allocate such memory internally and invoke an application provided informational callback when these *internal allocations* are allocated and freed. Upon allocation of executable memory, <code>pfnInternalAllocation</code> will be called. Upon freeing executable memory, <code>pfnInternalFree</code> will be called. An implementation will only call an informational callback for executable memory allocations and frees.

The allocationType parameter to the pfnInternalAllocation and pfnInternalFree functions may be one of the following values:

```
typedef enum VkInternalAllocationType {
    VK_INTERNAL_ALLOCATION_TYPE_EXECUTABLE = 0,
} VkInternalAllocationType;
```

VK_INTERNAL_ALLOCATION_TYPE_EXECUTABLE - The allocation is intended for execution by the host.

An implementation must only make calls into an application-provided allocator from within the scope of an API command. An implementation must only make calls into an application-provided allocator from the same thread that called the provoking API command. The implementation should not synchronize calls to any of the callbacks. If synchronization is needed, the callbacks must provide it themselves. The informational callbacks are subject to the same restrictions as the allocation callbacks.

If an implementation intends to make calls through an VkAllocationCallbacks structure between the time a **vkCreate*** command returns and the time a corresponding **vkDestroy*** command begins, that implementation must save a copy of the allocator before the **vkCreate*** command returns. The callback functions and any data structures they rely upon must remain valid for the lifetime of the object they are associated with.

If an allocator is provided to a **vkCreate*** command, a *compatible* allocator must be provided to the corresponding **vkDestroy*** command. Two VkAllocationCallbacks structures are compatible if memory allocated with *pfnAllocation* or *pfnReallocation* in each can be freed with *pfnReallocation* or *pfnFree* in the other. An allocator must not be provided to a **vkDestroy*** command if an allocator was not provided to the corresponding **vkCreate*** command.

If a non-NULL allocator is used, the <code>pfnAllocation</code>, <code>pfnReallocation</code> and <code>pfnFree</code> members must be non-NULL and point to valid implementations of the callbacks. An application can choose to not provide informational callbacks by setting both <code>pfnInternalAllocation</code> and <code>pfnInternalFree</code> to <code>NULL</code>. <code>pfnInternalAllocation</code> and <code>pfnInternalFree</code> must either both be <code>NULL</code> or both be non-NULL.

If pfnAllocation or pfnReallocation fail, the implementation may fail object creation and/or generate an $VK_ERROR_OUT_OF_HOST_MEMORY$ error, as appropriate.

Allocation callbacks must not call any Vulkan commands.

The following sets of rules define when an implementation is permitted to call the allocator callbacks.

 ${\it pfnAllocation}\ or\ {\it pfnReallocation}\ may\ be\ called\ in\ the\ following\ situations:$

- Host memory scoped to the lifetime of a VkDevice or VkInstance may be allocated from any API command.
- Host memory scoped to the duration of a command may be allocated from any API command.
- Host memory scoped to the lifetime of a VkPipelineCache may only be allocated from:
 - vkCreatePipelineCache
 - vkMergePipelineCaches for dstCache
 - vkCreateGraphicsPipelines for pPipelineCache
 - vkCreateComputePipelines for pPipelineCache
- Host memory scoped to the lifetime of a VkDescriptorPool may only be allocated from:
 - any command that takes the pool as a direct argument
 - vkAllocateDescriptorSets for the descriptorPool member of its pAllocateInfo parameter
 - vkCreateDescriptorPool
- Host memory scoped to the lifetime of a VkCommandPool may only be allocated from:
 - any command that takes the pool as a direct argument
 - vkCreateCommandPool
 - vkAllocateCommandBuffers for the commandPool member of its pAllocateInfo parameter
 - any vkCmd* command whose commandBuffer was allocated from that VkCommandPool

• Host memory scoped to the lifetime of any other object may only be allocated in that object's vkCreate* command.

pfnFree may be called in the following situations:

- Host memory scoped to the lifetime of a VkDevice or VkInstance may be freed from any API command.
- Host memory scoped to the duration of a command must be freed by any API command which allocates such memory.
- Host memory scoped to the lifetime of a VkPipelineCache may be freed from vkDestroyPipelineCache.
- Host memory scoped to the lifetime of a VkDescriptorPool may be freed from
 - any command that takes the pool as a direct argument
- Host memory scoped to the lifetime of a VkCommandPool may be freed from:
 - any command that takes the pool as a direct argument
 - vkResetCommandBuffer whose commandBuffer was allocated from that VkCommandPool
- Host memory scoped to the lifetime of any other object may be freed in that object's vkDestroy* command.
- Any command that allocates host memory may also free host memory of the same scope.

10.2 Device Memory

Device memory is memory that is visible to the device, for example the contents of opaque images that can be natively used by the device, or uniform buffer objects that reside in on-device memory.

Memory properties of a physical device describe the memory heaps and memory types available.

To query memory properties, call:

- physicalDevice is the handle to the device to query.
- pMemoryProperties points to an instance of VkPhysicalDeviceMemoryProperties structure in which the properties are returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pMemoryProperties must be a pointer to a VkPhysicalDeviceMemoryProperties structure

The VkPhysicalDeviceMemoryProperties structure is defined as:

- memoryTypeCount is the number of valid elements in the pMemoryRanges array.
- memoryTypes is an array of VkMemoryType structures describing the memory types that can be used to access memory allocated from the heaps specified by memoryHeaps.
- memoryHeapCount is the number of valid elements in the pMemoryRanges array.
- memoryHeaps is an array of VkMemoryHeap structures describing the memory heaps from which memory can be allocated.

The VkPhysicalDeviceMemoryProperties structure describes a number of *memory heaps* as well as a number of *memory types* that can be used to access memory allocated in those heaps. Each heap describes a memory resource of a particular size, and each memory type describes a set of memory properties (e.g. host cached vs uncached) that can be used with a given memory heap. Allocations using a particular memory type will consume resources from the heap indicated by that memory type's heap index. More than one memory type may share each heap, and the heaps and memory types provide a mechanism to advertise an accurate size of the physical memory resources while allowing the memory to be used with a variety of different properties.

The number of memory heaps is given by <code>memoryHeapCount</code> and is less than or equal to <code>VK_MAX_MEMORY_HEAPS</code>. Each heap is described by an element of the <code>memoryHeaps</code> array, as a <code>VkMemoryHeap</code> structure. The number of memory types available across all memory heaps is given by <code>memoryTypeCount</code> and is less than or equal to <code>VK_MAX_MEMORY_TYPES</code>. Each memory type is described by an element of the <code>memoryTypes</code> array, as a <code>VkMemoryType</code> structure.

At least one heap must include VK_MEMORY_HEAP_DEVICE_LOCAL_BIT in VkMemoryHeap::flags. If there are multiple heaps that all have similar performance characteristics, they may all include VK_MEMORY_HEAP_DEVICE_LOCAL_BIT. In a unified memory architecture (UMA) system, there is often only a single memory heap which is considered to be equally "local" to the host and to the device, and such an implementation must advertise the heap as device-local.

Each memory type returned by vkGetPhysicalDeviceMemoryProperties must have its propertyFlags set to one of the following values:

- ()
- VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_COHERENT_BIT
- VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_CACHED_BIT
- VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_CACHED_BIT|VK_MEMORY_PROPERTY_HOST_COHERENT_BIT
- VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT
- VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT|VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_COHERENT_BIT
- VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT|VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_CACHED_BIT

- VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT|VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT|VK_MEMORY_PROPERTY_HOST_CACHED_BIT|VK_MEMORY_PROPERTY_HOST_COHERENT_BIT
- VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT | VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT

There must be at least one memory type with both the VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT and VK_MEMORY_PROPERTY_HOST_COHERENT_BIT bits set in its propertyFlags. There must be at least one memory type with the VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT bit set in its propertyFlags.

The memory types are sorted according to a preorder which serves to aid in easily selecting an appropriate memory type. Given two memory types X and Y, the preorder defines $X \le Y$ if:

- the memory property bits set for X are a strict subset of the memory property bits set for Y. Or,
- the memory property bits set for X are the same as the memory property bits set for Y, and X uses a memory heap with greater or equal performance (as determined in an implementation-specific manner).

Memory types are ordered in the list such that X is assigned a lesser memoryTypeIndex than Y if $X \le Y \land \neg (Y \le X)$ according to the preorder. Note that the list of all allowed memory property flag combinations above satisfies this preorder, but other orders would as well. The goal of this ordering is to enable applications to use a simple search loop in selecting the proper memory type, along the lines of:

```
// Find a memory type in "memoryTypeBits" that includes all of "properties"
int32_t FindProperties(uint32_t memoryTypeBits, VkMemoryPropertyFlags properties)
{
    for (int32_t i = 0; i < memoryTypeCount; ++i)</pre>
        if ((memoryTypeBits & (1 << i)) &&</pre>
            ((memoryTypes[i].propertyFlags & properties) == properties))
            return i;
    return -1;
// Try to find an optimal memory type, or if it does not exist
// find any compatible memory type
VkMemoryRequirements memoryRequirements;
vkGetImageMemoryRequirements(device, image, &memoryRequirements);
int32_t memoryType = FindProperties(memoryRequirements.memoryTypeBits, \leftrightarrow
   optimalProperties);
if (memoryType == -1)
    memoryType = FindProperties (memoryRequirements.memoryTypeBits, requiredProperties) \leftrightarrow
```

The loop will find the first supported memory type that has all bits requested in **properties** set. If there is no exact match, it will find a closest match (i.e. a memory type with the fewest additional bits set), which has some additional bits set but which are not detrimental to the behaviors requested by **properties**. The application can first search for the optimal properties, e.g. a memory type that is device-local or supports coherent cached accesses, as appropriate for the intended usage, and if such a memory type is not present can fallback to searching for a less optimal but guaranteed set of properties such as "0" or "host-visible and coherent".

The VkMemoryHeap structure is defined as:

- size is the total memory size in bytes in the heap.
- flags is a bitmask of attribute flags for the heap. The bits specified in flags are:

```
typedef enum VkMemoryHeapFlagBits {
    VK_MEMORY_HEAP_DEVICE_LOCAL_BIT = 0x00000001,
} VkMemoryHeapFlagBits;
```

 if flags contains VK_MEMORY_HEAP_DEVICE_LOCAL_BIT, it means the heap corresponds to device local memory. Device local memory may have different performance characteristics than host local memory, and may support different memory property flags.

The VkMemoryType structure is defined as:

- heapIndex describes which memory heap this memory type corresponds to, and must be less than memoryHeapCount from the VkPhysicalDeviceMemoryProperties structure.
- propertyFlags is a bitmask of properties for this memory type. The bits specified in propertyFlags are:

```
typedef enum VkMemoryPropertyFlagBits {
    VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT = 0x00000001,
    VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT = 0x00000002,
    VK_MEMORY_PROPERTY_HOST_COHERENT_BIT = 0x00000004,
    VK_MEMORY_PROPERTY_HOST_CACHED_BIT = 0x00000008,
    VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT = 0x00000010,
} VkMemoryPropertyFlagBits;
```

- if propertyFlags has the VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT bit set, memory allocated with this
 type is the most efficient for device access. This property will only be set for memory types belonging to heaps with
 the VK_MEMORY_HEAP_DEVICE_LOCAL_BIT set.
- if propertyFlags has the VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT bit set, memory allocated with this type can be mapped using vkMapMemory so that it can be accessed on the host.
- if propertyFlags has the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT bit set, host cache management
 commands vkFlushMappedMemoryRanges and vkInvalidateMappedMemoryRanges are not needed to
 make host writes visible to the device or device writes visible to the host, respectively.
- if propertyFlags has the VK_MEMORY_PROPERTY_HOST_CACHED_BIT bit set, memory allocated with this
 type is cached on the host. Host memory accesses to uncached memory are slower than to cached memory, however
 uncached memory is always host coherent.
- if propertyFlags has the VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT bit set, the memory type only allows device access to the memory. Memory types must not have both VK_MEMORY_PROPERTY_LAZILY_ ALLOCATED_BIT and VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT set. Additionally, the object's backing memory may be provided by the implementation lazily as specified in Lazily Allocated Memory.

A Vulkan device operates on data in device memory via memory objects that are represented in the API by a VkDeviceMemory handle.

Memory objects are represented by VkDeviceMemory handles:

VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkDeviceMemory)

To allocate memory objects, call:

- device is the logical device that owns the memory.
- pAllocateInfo is a pointer to an instance of the VkMemoryAllocateInfo structure describing parameters of the allocation. A successful returned allocation must use the requested parameters no substitution is permitted by the implementation.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pMemory is a pointer to a VkDeviceMemory handle in which information about the allocated memory is returned.

Allocations returned by **vkAllocateMemory** are guaranteed to meet any alignment requirement by the implementation. For example, if an implementation requires 128 byte alignment for images and 64 byte alignment for buffers, the device memory returned through this mechanism would be 128-byte aligned. This ensures that applications can correctly suballocate objects of different types (with potentially different alignment requirements) in the same memory object.

When memory is allocated, its contents are undefined.

There is an implementation-dependent maximum number of memory allocations which can be simultaneously created on a device. This is specified by the maxMemoryAllocationCount member of the VkPhysicalDeviceLimits structure. If maxMemoryAllocationCount is exceeded, vkAllocateMemory will return VK_ERROR_TOO_MANY_OBJECTS.



Note

Some platforms may have a limit on the maximum size of a single allocation. For example, certain systems may fail to create allocations with a size greater than or equal to 4GB. Such a limit is implementation-dependent, and if such a failure occurs then the error VK_ERROR_OUT_OF_DEVICE_MEMORY should be returned.

Valid Usage

- device must be a valid VkDevice handle
- pAllocateInfo must be a pointer to a valid VkMemoryAllocateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pMemory must be a pointer to a VkDeviceMemory handle
- The number of currently valid memory objects, allocated from device, must be less than VkPhysicalDeviceLimits::maxMemoryAllocationCount

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_TOO_MANY_OBJECTS

The VkMemoryAllocateInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- allocationSize is the size of the allocation in bytes
- memoryTypeIndex is the memory type index, which selects the properties of the memory to be allocated, as well as the heap the memory will come from.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO
- pNext must be NULL, or a pointer to a valid instance of VkDedicatedAllocationMemoryAllocateInfoNV
- allocationSize must be less than or equal to the amount of memory available to the VkMemoryHeap specified by memoryTypeIndex and the calling command's VkDevice
- allocationSize must be greater than 0

If the pNext list includes a VkDedicatedAllocationMemoryAllocateInfoNV structure, then that structure includes a handle of the sole buffer or image resource that the memory can be bound to.

The VkDedicatedAllocationMemoryAllocateInfoNV structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- image is VK_NULL_HANDLE or a handle of an image which this memory will be bound to.
- buffer is VK_NULL_HANDLE or a handle of a buffer which this memory will be bound to.

Valid Usage

- stype must be VK STRUCTURE TYPE DEDICATED ALLOCATION MEMORY ALLOCATE INFO NV
- pNext must be NULL
- If image is not VK_NULL_HANDLE, image must be a valid VkImage handle
- If buffer is not VK NULL HANDLE, buffer must be a valid VkBuffer handle
- Both of buffer, and image that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- At least one of image and buffer must be VK_NULL_HANDLE
- If image is not VK_NULL_HANDLE, the image must have been created with VkDedicatedAllocationImageCreateInfoNV::dedicatedAllocation equal to VK_TRUE
- If buffer is not VK_NULL_HANDLE, the buffer must have been created with VkDedicatedAllocationBufferCreateInfoNV::dedicatedAllocation equal to VK_TRUE
- If image is not VK_NULL_HANDLE, VkMemoryAllocateInfo::allocationSize must equal the VkMemoryRequirements::size of the image
- If buffer is not VK_NULL_HANDLE, VkMemoryAllocateInfo::allocationSize must equal the VkMemoryRequirements::size of the buffer

When allocating memory that may be exported to another process or Vulkan instance, add a VkExportMemoryAllocateInfoNV structure to the pNext chain of the VkMemoryAllocateInfo structure, specifying the handle types that may be exported.

The VkMemoryAllocateInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- handleTypes is a bitmask of VkExternalMemoryHandleTypeFlagBitsNV specifying one or more memory handle types that may be exported. Multiple handle types may be requested for the same allocation as long as they are compatible, as reported by vkGetPhysicalDeviceExternalImageFormatPropertiesNV.

Valid Usage

- $\bullet \ \, \textit{stype} \ \, \textit{must} \ \, \textit{be} \ \, \textit{VK_STRUCTURE_TYPE_EXPORT_MEMORY_ALLOCATE_INFO_NV} \\$
- pNext must be NULL
- handleTypes must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- handleTypes must not be 0

When VkExportMemoryAllocateInfoNV:: handleTypes includes VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_BIT_NV, add a VkExportMemoryWin32HandleInfoNV to the pNext chain of the VkExportMemoryAllocateInfoNV structure to specify security attributes and access rights for the memory object's external handle.

The VkExportMemoryWin32HandleInfoNV structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- pAttributes is a pointer to a Windows SECURITY_ATTRIBUTES structure specifying security attributes of the handle.
- dwAccess is a DWORD specifying access rights of the handle.

If this structure is not present, or if pAttributes is set to NULL, default security descriptor values will be used, and child processes created by the application will not inherit the handle, as described in the MSDN documentation for "Synchronization Object Security and Access Rights"[1]. Further, if the structure is not present, the access rights will be

```
etext:DXGI_SHARED_RESOURCE_READ | etext:DXGI_SHARED_RESOURCE_WRITE
```

[1] https://msdn.microsoft.com/en-us/library/windows/desktop/ms686670.aspx

Valid Usage

- stype must be VK_STRUCTURE_TYPE_EXPORT_MEMORY_WIN32_HANDLE_INFO_NV
- pNext must be NULL
- pAttributes must be a pointer to a valid SECURITY_ATTRIBUTES value

To import memory created on the same physical device but outside of the current Vulkan instance, add a VkImportMemoryWin32HandleInfoNV structure to the pNext chain of the VkMemoryAllocateInfo structure, specifying a handle to and the type of the memory.

The VkImportMemoryWin32HandleInfoNV structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- handleType is 0 or a flag specifying the type of memory handle in handle. Flags which may be specified are:

```
typedef enum VkExternalMemoryHandleTypeFlagBitsNV {
    VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_BIT_NV = 0x00000001,
    VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_KMT_BIT_NV = 0x00000002,
    VK_EXTERNAL_MEMORY_HANDLE_TYPE_D3D11_IMAGE_BIT_NV = 0x00000004,
    VK_EXTERNAL_MEMORY_HANDLE_TYPE_D3D11_IMAGE_KMT_BIT_NV = 0x00000008,
} VkExternalMemoryHandleTypeFlagBitsNV;
```

- handle is a Windows HANDLE referring to the memory.
 - if handleType is VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_BIT_NV or VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_KMT_BIT_NV, handle must be a valid handle returned by vkGetMemoryWin32HandleNV or, in the case of VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_BIT_NV, a handle duplicated from such a handle using DuplicateHandle().
 - if handleType is VK_EXTERNAL_MEMORY_HANDLE_TYPE_D3D11_IMAGE_BIT_NV, handle must be a valid NT handle returned by IDXGIResource1::ftext:CreateSharedHandle() or a handle duplicated from such a handle using DuplicateHandle().

- if handleType is VK_EXTERNAL_MEMORY_HANDLE_TYPE_D3D11_IMAGE_KMT_BIT_NV, handle must be a valid handle returned by IDXGIResource::GetSharedHandle().

Valid Usage

- stype must be VK_STRUCTURE_TYPE_IMPORT_MEMORY_WIN32_HANDLE_INFO_NV
- pNext must be NULL
- handleType must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- handleType must not be 0

To retrieve the handle corresponding to a device memory object created with

VkExportMemoryAllocateInfoNV::handleTypes set to include VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_BIT_NV or VK_EXTERNAL_MEMORY_HANDLE_TYPE_OPAQUE_WIN32_KMT_BIT_NV, call:

VkResult vkGetMemoryWin32HandleNV(

VkDevice device,
VkDeviceMemory memory,
VkExternalMemoryHandleTypeFlagsNV handleType,
HANDLE* pHandle);

- device is the logical device that owns the memory.
- memory is the VkDeviceMemory object.
- handleType is a bitmask of VkExternalMemoryHandleTypeFlagBitsNV containing a single bit specifying the type of handle requested.
- handle points to a Windows HANDLE in which the handle is returned.

- device must be a valid VkDevice handle
- memory must be a valid VkDeviceMemory handle
- handleType must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- handleType must not be 0
- pHandle must be a pointer to a HANDLE value
- memory must have been created, allocated, or retrieved from device
- handleType must be a flag specified in VkExportMemoryAllocateInfoNV::handleTypes when allocating memory

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_TOO_MANY_OBJECTS
- VK_ERROR_OUT_OF_HOST_MEMORY

To free a memory object, call:

- device is the logical device that owns the memory.
- memory is the VkDeviceMemory object to be freed.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Before freeing a memory object, an application must ensure the memory object is no longer in use by the device—for example by command buffers queued for execution. The memory can remain bound to images or buffers at the time the memory object is freed, but any further use of them (on host or device) for anything other than destroying those objects will result in undefined behavior. If there are still any bound images or buffers, the memory may not be immediately released by the implementation, but must be released by the time all bound images and buffers have been destroyed. Once memory is released, it is returned to the heap from which it was allocated.

How memory objects are bound to Images and Buffers is described in detail in the Resource Memory Association section.

If a memory object is mapped at the time it is freed, it is implicitly unmapped.

- device must be a valid VkDevice handle
- If memory is not VK_NULL_HANDLE, memory must be a valid VkDeviceMemory handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If memory is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to memory (via images or buffers) must have completed execution

Host Synchronization

• Host access to memory must be externally synchronized

10.2.1 Host Access to Device Memory Objects

Memory objects created with **vkAllocateMemory** are not directly host accessible.

Memory objects created with the memory property VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT are considered *mappable*. Memory objects must be mappable in order to be successfully mapped on the host.

To retrieve a host virtual address pointer to a region of a mappable memory object, call:

- device is the logical device that owns the memory.
- memory is the VkDeviceMemory object to be mapped.
- offset is a zero-based byte offset from the beginning of the memory object.
- size is the size of the memory range to map, or VK WHOLE SIZE to map from offset to the end of the allocation.
- flags is reserved for future use.
- ppData points to a pointer in which is returned a host-accessible pointer to the beginning of the mapped range. This pointer minus offset must be aligned to at least VkPhysicalDeviceLimits::minMemoryMapAlignment.

It is an application error to call **vkMapMemory** on a memory object that is already mapped.

vkMapMemory does not check whether the device memory is currently in use before returning the host-accessible pointer. The application must guarantee that any previously submitted command that writes to this range has completed before the host reads from or writes to that range, and that any previously submitted command that reads from that range has completed before the host writes to that region (see here for details on fulfilling such a guarantee). If the device memory was allocated without the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT set, these guarantees must be made for an extended range: the application must round down the start of the range to the nearest multiple of VkPhysicalDeviceLimits::nonCoherentAtomSize, and round the end of the range up to the nearest multiple of VkPhysicalDeviceLimits::nonCoherentAtomSize.

While a range of device memory is mapped for host access, the application is responsible for synchronizing both device and host access to that memory range.



Note

It is important for the application developer to become meticulously familiar with all of the mechanisms described in the chapter on Synchronization and Cache Control as they are crucial to maintaining memory access ordering.

Valid Usage

- device must be a valid VkDevice handle
- memory must be a valid VkDeviceMemory handle
- flags must be 0
- ppData must be a pointer to a pointer
- memory must have been created, allocated, or retrieved from device
- memory must not currently be mapped
- offset must be less than the size of memory
- If size is not equal to VK_WHOLE_SIZE, size must be greater than 0
- If size is not equal to VK_WHOLE_SIZE, size must be less than or equal to the size of the memory minus offset
- memory must have been created with a memory type that reports VK_MEMORY_PROPERTY_HOST_VISIBLE_ BIT

Host Synchronization

• Host access to memory must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_MEMORY_MAP_FAILED

Two commands are provided to enable applications to work with non-coherent memory allocations: **vkFlushMappedMemoryRanges** and **vkInvalidateMappedMemoryRanges**.



Note

If the memory object was created with the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT set, **vkFlus hMappedMemoryRanges** and **vkInvalidateMappedMemoryRanges** are unnecessary and may have performance cost.

To flush ranges of non-coherent memory from the host caches, call:

- device is the logical device that owns the memory ranges.
- memoryRangeCount is the length of the pMemoryRanges array.
- pMemoryRanges is a pointer to an array of VkMappedMemoryRange structures describing the memory ranges to flush.

vkFlushMappedMemoryRanges must be used to guarantee that host writes to non-coherent memory are visible to the device. It must be called after the host writes to non-coherent memory have completed and before command buffers that will read or write any of those memory locations are submitted to a queue.

- device must be a valid VkDevice handle
- pMemoryRanges must be a pointer to an array of memoryRangeCount valid VkMappedMemoryRange structures
- memoryRangeCount must be greater than 0

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To invalidate ranges of non-coherent memory from the host caches, call:

- device is the logical device that owns the memory ranges.
- memoryRangeCount is the length of the pMemoryRanges array.
- pMemoryRanges is a pointer to an array of VkMappedMemoryRange structures describing the memory ranges to invalidate.

vkInvalidateMappedMemoryRanges must be used to guarantee that device writes to non-coherent memory are visible to the host. It must be called after command buffers that execute and flush (via memory barriers) the device writes have completed, and before the host will read or write any of those locations. If a range of non-coherent memory is written by the host and then invalidated without first being flushed, its contents are undefined.

Valid Usage

- device must be a valid VkDevice handle
- pMemoryRanges must be a pointer to an array of memoryRangeCount valid VkMappedMemoryRange structures
- memoryRangeCount must be greater than 0

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkMappedMemoryRange structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- memory is the memory object to which this range belongs.
- offset is the zero-based byte offset from the beginning of the memory object.
- size is either the size of range, or VK_WHOLE_SIZE to affect the range from offset to the end of the current mapping of the allocation.

- sType must be VK_STRUCTURE_TYPE_MAPPED_MEMORY_RANGE
- pNext must be NULL
- memory must be a valid VkDeviceMemory handle
- memory must currently be mapped
- If size is not equal to VK_WHOLE_SIZE, offset and size must specify a range contained within the currently mapped range of memory
- If size is equal to VK_WHOLE_SIZE, offset must be within the currently mapped range of memory
- offset must be a multiple of VkPhysicalDeviceLimits::nonCoherentAtomSize
- If size is not equal to VK_WHOLE_SIZE, size must be a multiple of VkPhysicalDeviceLimits::nonCoherentAtomSize

Host-visible memory types that advertise the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT property still require memory barriers between host and device in order to be coherent, but do not require additional cache management operations to achieve coherency. For host writes to be seen by subsequent command buffer operations, a pipeline barrier from a source of VK_ACCESS_HOST_WRITE_BIT and VK_PIPELINE_STAGE_HOST_BIT to a destination of the relevant device pipeline stages and access types must be performed. Note that such a barrier is performed implicitly upon each command buffer submission, so an explicit barrier is only rarely needed (e.g. if a command buffer waits upon an event signaled by the host, where the host wrote some data after submission). For device writes to be seen by subsequent host reads, a pipeline barrier is required to make the writes visible.

To unmap a memory object once host access to it is no longer needed by the application, call:

- device is the logical device that owns the memory.
- memory is the memory object to be unmapped.

Valid Usage

- device must be a valid VkDevice handle
- memory must be a valid VkDeviceMemory handle
- memory must have been created, allocated, or retrieved from device
- memory must currently be mapped

Host Synchronization

Host access to memory must be externally synchronized

10.2.2 Lazily Allocated Memory

If the memory object is allocated from a heap with the VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT bit set, that object's backing memory may be provided by the implementation lazily. The actual committed size of the memory may initially be as small as zero (or as large as the requested size), and monotonically increases as additional memory is needed.

A memory type with this flag set is only allowed to be bound to a VkImage whose usage flags include VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT.



Note

Using lazily allocated memory objects for framebuffer attachments that are not needed once a render pass instance has completed may allow some implementations to never allocate memory for such attachments.

To determine the amount of lazily-allocated memory that is currently committed for a memory object, call:

- device is the logical device that owns the memory.
- memory is the memory object being queried.
- pCommittedMemoryInBytes is a pointer to a VkDeviceSize value in which the number of bytes currently committed is returned, on success.

The implementation may update the commitment at any time, and the value returned by this query may be out of date.

The implementation guarantees to allocate any committed memory from the heapIndex indicated by the memory type that the memory object was created with.

- device must be a valid VkDevice handle
- memory must be a valid VkDeviceMemory handle
- pCommittedMemoryInBytes must be a pointer to a VkDeviceSize value
- memory must have been created, allocated, or retrieved from device
- memory must have been created with a memory type that reports VK_MEMORY_PROPERTY_LAZILY_ ALLOCATED_BIT

Chapter 11

Resource Creation

Vulkan supports two primary resource types: *buffers* and *images*. Resources are views of memory with associated formatting and dimensionality. Buffers are essentially unformatted arrays of bytes whereas images contain format information, can be multidimensional and may have associated metadata.

11.1 Buffers

Buffers represent linear arrays of data which are used for various purposes by binding them to a graphics or compute pipeline via descriptor sets or via certain commands, or by directly specifying them as parameters to certain commands.

Buffers are represented by VkBuffer handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkBuffer)
```

To create buffers, call:

- device is the logical device that creates the buffer object.
- pCreateInfo is a pointer to an instance of the VkBufferCreateInfo structure containing parameters affecting creation of the buffer.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pBuffer points to a VkBuffer handle in which the resulting buffer object is returned.

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkBufferCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pBuffer must be a pointer to a VkBuffer handle
- If the flags member of pCreateInfo includes VK_BUFFER_CREATE_SPARSE_BINDING_BIT, creating this VkBuffer must not cause the total required sparse memory for all currently valid sparse resources on the device to exceed VkPhysicalDeviceLimits::sparseAddressSpaceSize

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkBufferCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is a bitmask describing additional parameters of the buffer. See VkBufferCreateFlagBits below for a description of the supported bits.
- size is the size in bytes of the buffer to be created.
- *usage* is a bitmask describing the allowed usages of the buffer. See VkBufferUsageFlagBits below for a description of the supported bits.

- *sharingMode* is the sharing mode of the buffer when it will be accessed by multiple queue families, see VkSharingMode in the Resource Sharing section below for supported values.
- queueFamilyIndexCount is the number of entries in the pQueueFamilyIndices array.
- pQueueFamilyIndices is a list of queue families that will access this buffer (ignored if sharingMode is not VK_SHARING_MODE_CONCURRENT).

Bits which can be set in usage are:

```
typedef enum VkBufferUsageFlagBits {
    VK_BUFFER_USAGE_TRANSFER_SRC_BIT = 0x00000001,
    VK_BUFFER_USAGE_TRANSFER_DST_BIT = 0x00000002,
    VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT = 0x000000004,
    VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT = 0x000000008,
    VK_BUFFER_USAGE_UNIFORM_BUFFER_BIT = 0x00000010,
    VK_BUFFER_USAGE_STORAGE_BUFFER_BIT = 0x00000020,
    VK_BUFFER_USAGE_INDEX_BUFFER_BIT = 0x000000040,
    VK_BUFFER_USAGE_VERTEX_BUFFER_BIT = 0x000000080,
    VK_BUFFER_USAGE_INDIRECT_BUFFER_BIT = 0x000000100,
} VkBufferUsageFlagBits;
```

- VK_BUFFER_USAGE_TRANSFER_SRC_BIT indicates that the buffer can be used as the source of a *transfer command* (see the definition of VK_PIPELINE_STAGE_TRANSFER_BIT).
- VK_BUFFER_USAGE_TRANSFER_DST_BIT indicates that the buffer can be used as the destination of a transfer command.
- VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT indicates that the buffer can be used to create a VkBufferView suitable for occupying a VkDescriptorSet slot of type VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER.
- VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT indicates that the buffer can be used to create a VkBufferView suitable for occupying a VkDescriptorSet slot of type VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER.
- VK_BUFFER_USAGE_UNIFORM_BUFFER_BIT indicates that the buffer can be used in a VkDescriptorBufferInfo suitable for occupying a VkDescriptorSet slot either of type VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC.
- VK_BUFFER_USAGE_STORAGE_BUFFER_BIT indicates that the buffer can be used in a VkDescriptorBufferInfo suitable for occupying a VkDescriptorSet slot either of type VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC.
- VK_BUFFER_USAGE_INDEX_BUFFER_BIT indicates that the buffer is suitable for passing as the buffer parameter to vkCmdBindIndexBuffer.
- VK_BUFFER_USAGE_VERTEX_BUFFER_BIT indicates that the buffer is suitable for passing as an element of the pBuffers array to vkCmdBindVertexBuffers.
- VK_BUFFER_USAGE_INDIRECT_BUFFER_BIT indicates that the buffer is suitable for passing as the buffer parameter to vkCmdDrawIndirect, vkCmdDrawIndexedIndirect, or vkCmdDispatchIndirect.

Any combination of bits can be specified for *usage*, but at least one of the bits must be set in order to create a valid buffer.

Bits which can be set in flags are:

```
typedef enum VkBufferCreateFlagBits {
    VK_BUFFER_CREATE_SPARSE_BINDING_BIT = 0x00000001,
    VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT = 0x00000002,
    VK_BUFFER_CREATE_SPARSE_ALIASED_BIT = 0x00000004,
} VkBufferCreateFlagBits;
```

These bits have the following meanings:

- VK_BUFFER_CREATE_SPARSE_BINDING_BIT indicates that the buffer will be backed using sparse memory binding.
- VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT indicates that the buffer can be partially backed using sparse memory binding. Buffers created with this flag must also be created with the VK_BUFFER_CREATE_SPARSE_BINDING_BIT flag.
- VK_BUFFER_CREATE_SPARSE_ALIASED_BIT indicates that the buffer will be backed using sparse memory binding with memory ranges that might also simultaneously be backing another buffer (or another portion of the same buffer). Buffers created with this flag must also be created with the VK_BUFFER_CREATE_SPARSE_BINDING_BIT flag.

See Sparse Resource Features and Physical Device Features for details of the sparse memory features supported on a device.

- stype must be $VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO$
- pNext must be NULL, or a pointer to a valid instance of VkDedicatedAllocationBufferCreateInfoNV
- flags must be a valid combination of VkBufferCreateFlagBits values
- usage must be a valid combination of VkBufferUsageFlagBits values
- usage must not be 0
- sharingMode must be a valid VkSharingMode value
- size must be greater than 0
- If sharingMode is VK_SHARING_MODE_CONCURRENT, pQueueFamilyIndices must be a pointer to an array of queueFamilyIndexCount uint32 t values
- If sharingMode is VK_SHARING_MODE_CONCURRENT, queueFamilyIndexCount must be greater than 1
- If the sparse bindings feature is not enabled, flags must not contain VK_BUFFER_CREATE_SPARSE_BINDING_BIT
- If the sparse buffer residency feature is not enabled, flags must not contain VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT

- If the sparse aliased residency feature is not enabled, flags must not contain VK_BUFFER_CREATE_SPARSE_ALIASED_BIT
- If flags contains VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT or VK_BUFFER_CREATE_SPARSE_ALIASED_BIT, it must also contain VK_BUFFER_CREATE_SPARSE_BINDING_BIT

If the *pNext* list includes a VkDedicatedAllocationBufferCreateInfoNV structure, then that structure includes an enable controlling whether the buffer will have a dedicated memory allocation bound to it.

The VkDedicatedAllocationBufferCreateInfoNV structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- dedicatedAllocation indicates whether the buffer will have a dedicated allocation bound to it.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_BUFFER_CREATE_INFO_NV
- pNext must be NULL
- If dedicatedAllocation is VK_TRUE, VkBufferCreateInfo::flags must not include VK_BUFFER_CREATE_SPARSE_BINDING_BIT, VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT, or VK_BUFFER_CREATE_SPARSE_ALIASED_BIT

To destroy a buffer, call:

- device is the logical device that destroys the buffer.
- buffer is the buffer to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If buffer is not VK_NULL_HANDLE, buffer must be a valid VkBuffer handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If buffer is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to <code>buffer</code>, either directly or via a <code>VkBufferView</code>, must have completed execution
- If VkAllocationCallbacks were provided when buffer was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when buffer was created, pAllocator must be NULL

Host Synchronization

• Host access to buffer must be externally synchronized

11.2 Buffer Views

A *buffer view* represents a contiguous range of a buffer and a specific format to be used to interpret the data. Buffer views are used to enable shaders to access buffer contents interpreted as formatted data. In order to create a valid buffer view, the buffer must have been created with at least one of the following usage flags:

- VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT
- VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT

Buffer views are represented by VkBufferView handles:

VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkBufferView)

To create a buffer view, call:

- device is the logical device that creates the buffer view.
- pCreateInfo is a pointer to an instance of the VkBufferViewCreateInfo structure containing parameters to be used to create the buffer.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pView points to a VkBufferView handle in which the resulting buffer view object is returned.

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkBufferViewCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pView must be a pointer to a VkBufferView handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkBufferViewCreateInfo structure is defined as:

• *sType* is the type of this structure.

- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- buffer is a VkBuffer on which the view will be created.
- format is a VkFormat describing the format of the data elements in the buffer.
- offset is an offset in bytes from the base address of the buffer. Accesses to the buffer view from shaders use addressing that is relative to this starting offset.
- range is a size in bytes of the buffer view. If range is equal to VK_WHOLE_SIZE, the range from offset to the end of the buffer is used. If VK_WHOLE_SIZE is used and the remaining size of the buffer is not a multiple of the element size of format, then the nearest smaller multiple is used.

- sType must be VK_STRUCTURE_TYPE_BUFFER_VIEW_CREATE_INFO
- pNext must be NULL
- flags must be 0
- buffer must be a valid VkBuffer handle
- format must be a valid VkFormat value
- offset must be less than the size of buffer
- $\bullet \ \textit{offset must be a multiple of } VkPhysicalDeviceLimits:: \textit{minTexelBufferOffsetAlignment}\\$
- If range is not equal to VK_WHOLE_SIZE:
 - range must be greater than 0
 - range must be a multiple of the element size of format
 - range divided by the size of an element of format, must be less than or equal to VkPhysicalDeviceLimits::maxTexelBufferElements
 - the sum of offset and range must be less than or equal to the size of buffer
- buffer must have been created with a usage value containing at least one of VK_BUFFER_USAGE_ UNIFORM_TEXEL_BUFFER_BIT or VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT
- If buffer was created with usage containing VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT, format must be supported for uniform texel buffers, as specified by the VK_FORMAT_FEATURE_UNIFORM_TEXEL_BUFFER_BIT flag in VkFormatProperties::bufferFeatures returned by vkGetPhysicalDeviceFormatProperties
- If buffer was created with usage containing VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT, format must be supported for storage texel buffers, as specified by the VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_BIT flag in VkFormatProperties::bufferFeatures returned by vkGetPhysicalDeviceFormatProperties

To destroy a buffer view, call:

- device is the logical device that destroys the buffer view.
- bufferView is the buffer view to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If bufferView is not VK_NULL_HANDLE, bufferView must be a valid VkBufferView handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If bufferView is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to <code>bufferView</code> must have completed execution
- If VkAllocationCallbacks were provided when bufferView was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when bufferView was created, pAllocator must be NULL

Host Synchronization

• Host access to bufferView must be externally synchronized

11.3 Images

Images represent multidimensional - up to 3 - arrays of data which can be used for various purposes (e.g. attachments, textures), by binding them to a graphics or compute pipeline via descriptor sets, or by directly specifying them as parameters to certain commands.

Images are represented by VkImage handles:

VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkImage)

To create images, call:

- device is the logical device that creates the image.
- pCreateInfo is a pointer to an instance of the VkImageCreateInfo structure containing parameters to be used to create the image.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pImage points to a VkImage handle in which the resulting image object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkImageCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pImage must be a pointer to a VkImage handle
- If the flags member of pCreateInfo includes VK_IMAGE_CREATE_SPARSE_BINDING_BIT, creating this VkImage must not cause the total required sparse memory for all currently valid sparse resources on the device to exceed VkPhysicalDeviceLimits::sparseAddressSpaceSize

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkImageCreateInfo structure is defined as:

```
typedef struct VkImageCreateInfo {
   VkStructureType sType;
   VkImageCreateFlags flags;
VkImageType
   VkImageType
                           imageType;
   VkFormat
                           format;
   VkExtent3D
                           extent;
   uint32_t
                          mipLevels;
   uint32_t
                          arrayLayers;
   VkSampleCountFlagBits samples;
   VkImageTiling
                          tiling;
   VkImageUsageFlags
                          usage;
                          sharingMode;
   VkSharingMode
                          queueFamilyIndexCount;
   uint32_t
   const uint32_t* pQueueFamilyIndices;
VkImageLayout initialLayout;
} VkImageCreateInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is a bitmask describing additional parameters of the image. See VkImageCreateFlagBits below for a description of the supported bits.
- *imageType* is a VkImageType specifying the basic dimensionality of the image, as described below. Layers in array textures do not count as a dimension for the purposes of the image type.
- format is a VkFormat describing the format and type of the data elements that will be contained in the image.
- extent is a VkExtent 3D describing the number of data elements in each dimension of the base level.
- mipLevels describes the number of levels of detail available for minified sampling of the image.
- arrayLayers is the number of layers in the image.
- samples is the number of sub-data element samples in the image as defined in VkSampleCountFlagBits. See Multisampling.
- tiling is a VkImageTiling specifying the tiling arrangement of the data elements in memory, as described below.
- *usage* is a bitmask describing the intended usage of the image. See VkImageUsageFlagBits below for a description of the supported bits.
- sharingMode is the sharing mode of the image when it will be accessed by multiple queue families, and must be one of the values described for VkSharingMode in the Resource Sharing section below.
- queueFamilyIndexCount is the number of entries in the pQueueFamilyIndices array.
- pQueueFamilyIndices is a list of queue families that will access this image (ignored if sharingMode is not VK_SHARING_MODE_CONCURRENT).
- initialLayout selects the initial VkImageLayout state of all image subresources of the image. See Image Layouts. initialLayout must be VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_ PREINITIALIZED.

Valid limits for the image <code>extent</code>, <code>mipLevels</code>, <code>arrayLayers</code> and <code>samples</code> members are queried with the <code>vkGetPhysicalDeviceImageFormatProperties</code> command.

Images created with tiling equal to VK_IMAGE_TILING_LINEAR have further restrictions on their limits and capabilities compared to images created with tiling equal to VK_IMAGE_TILING_OPTIMAL. Creation of images with tiling VK_IMAGE_TILING_LINEAR may not be supported unless other parameters meet all of the constraints:

- imageType is VK_IMAGE_TYPE_2D
- format is not a depth/stencil format
- mipLevels is 1
- arrayLayers is 1
- samples is VK_SAMPLE_COUNT_1_BIT
- usage only includes VK_IMAGE_USAGE_TRANSFER_SRC_BIT and/or VK_IMAGE_USAGE_TRANSFER_DST_ BIT

Implementations may support additional limits and capabilities beyond those listed above. To determine the specific capabilities of an implementation, query the valid *usage* bits by calling

vkGetPhysicalDeviceFormatProperties and the valid limits for mipLevels and arrayLayers by calling vkGetPhysicalDeviceImageFormatProperties.

- sType must be VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO
- pNext must be NULL, or a pointer to a valid instance of VkDedicatedAllocationImageCreateInfoNV
- flags must be a valid combination of VkImageCreateFlagBits values
- imageType must be a valid VkImageType value
- format must be a valid VkFormat value
- samples must be a valid VkSampleCountFlagBits value
- usage must be a valid combination of VkImageUsageFlagBits values
- usage must not be 0
- sharingMode must be a valid VkSharingMode value
- initialLayout must be a valid VkImageLayout value
- If sharingMode is VK_SHARING_MODE_CONCURRENT, pQueueFamilyIndices must be a pointer to an array of queueFamilyIndexCount uint32_t values
- If sharingMode is VK_SHARING_MODE_CONCURRENT, queueFamilyIndexCount must be greater than 1
- format must not be VK FORMAT UNDEFINED

- The width, height, and depth members of extent must all be greater than 0
- mipLevels must be greater than 0
- arrayLayers must be greater than 0
- If flags contains VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT, imageType must be VK_IMAGE_TYPE 2D
- If imageType is VK_IMAGE_TYPE_1D, extent.width must be less than or equal to VkPhysicalDeviceLimits::maxImageDimension1D, or VkImageFormatProperties::maxExtent.width (as returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, usage, and flags equal to those in this structure) whichever is higher
- If imageType is VK_IMAGE_TYPE_2D and flags does not contain VK_IMAGE_CREATE_CUBE_
 COMPATIBLE_BIT, extent.width and extent.height must be less than or equal to
 VkPhysicalDeviceLimits::maxImageDimension2D, or VkImageFormatProperties::maxExtent.
 width/height (as returned by vkGetPhysicalDeviceImageFormatProperties with format, type,
 tiling, usage, and flags equal to those in this structure) whichever is higher
- If imageType is VK_IMAGE_TYPE_2D and flags contains VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT, extent.width and extent.height must be less than or equal to VkPhysicalDeviceLimits::maxImageDimensionCube, or VkImageFormatProperties::maxExtent.width/height (as returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, usage, and flags equal to those in this structure) whichever is higher
- If imageType is VK_IMAGE_TYPE_2D and flags contains VK_IMAGE_CREATE_CUBE_COMPATIBLE_ BIT, extent.width and extent.height must be equal and arrayLayers must be greater than or equal to 6
- If imageType is VK_IMAGE_TYPE_3D, extent.width, extent.height and extent.depth must be less than or equal to VkPhysicalDeviceLimits::maxImageDimension3D, or VkImageFormatProperties::maxExtent.width/height/depth (as returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, usage, and flags equal to those in this structure) whichever is higher
- If imageType is VK_IMAGE_TYPE_1D, both extent.height and extent.depth must be 1
- If imageType is VK_IMAGE_TYPE_2D, extent.depth must be 1
- mipLevels must be less than or equal to $\lfloor log_2(max(extent.width, extent.height, extent.depth)) \rfloor + 1$
- If any of extent.width, extent.height, or extent.depth are greater than the equivalently named members of VkPhysicalDeviceLimits::maxImageDimension3D, mipLevels must be less than or equal to VkImageFormatProperties::maxMipLevels (as returned by
 - **vkGetPhysicalDeviceImageFormatProperties** with format, type, tiling, usage, and flags equal to those in this structure)
- arrayLayers must be less than or equal to VkPhysicalDeviceLimits::maxImageArrayLayers, or VkImageFormatProperties::maxArrayLayers (as returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, usage, and flags equal to those in this structure) whichever is higher
- If samples is not VK_SAMPLE_COUNT_1_BIT, imageType must be VK_IMAGE_TYPE_2D, flags must not contain VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT, tiling must be VK_IMAGE_TILING_OPTIMAL, and mipLevels must be equal to 1

- If usage includes VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT, then bits other than VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT, and VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT mustnot: be set
- If usage includes VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, VK_IMAGE_USAGE_DEPTH_ STENCIL_ATTACHMENT_BIT, VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT, or VK_IMAGE_ USAGE_INPUT_ATTACHMENT_BIT, extent.width must be less than or equal to VkPhysicalDeviceLimits::maxFramebufferWidth
- If usage includes VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, VK_IMAGE_USAGE_DEPTH_ STENCIL_ATTACHMENT_BIT, VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT, or VK_IMAGE_ USAGE_INPUT_ATTACHMENT_BIT, extent.height must be less than or equal to VkPhysicalDeviceLimits::maxFramebufferHeight
- samples must be a bit value that is set in VkImageFormatProperties::sampleCounts returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, usage, and flags equal to those in this structure
- If the ETC2 texture compression feature is not enabled, <code>format</code> must not be VK_FORMAT_ETC2_R8G8B8_ UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK, VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK, VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK, VK_FORMAT_EAC_R11_UNORM_BLOCK, VK_FORMAT_EAC_R11_SNORM_BLOCK, VK_FORMAT_EAC_R11G11_UNORM_BLOCK, or VK_FORMAT_EAC_R11G11_SNORM_BLOCK
- If the ASTC LDR texture compression feature is not enabled, format must not be VK_FORMAT_ASTC_4x4_ UNORM_BLOCK, VK_FORMAT_ASTC_4x4_SRGB_BLOCK, VK_FORMAT_ASTC_5x4_UNORM_BLOCK, VK_FORMAT_ASTC_5x4_SRGB_BLOCK, VK_FORMAT_ASTC_5x5_UNORM_BLOCK, VK_FORMAT_ASTC_5x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x6_UNORM_BLOCK, VK_FORMAT_ASTC_6x6_SRGB_BLOCK, VK_FORMAT_ASTC_8x5_UNORM_BLOCK, VK_FORMAT_ASTC_8x5_SRGB_BLOCK, VK_FORMAT_ASTC_8x6_UNORM_BLOCK, VK_FORMAT_ASTC_8x6_SRGB_BLOCK, VK_FORMAT_ASTC_8x8_UNORM_BLOCK, VK_FORMAT_ASTC_8x8_SRGB_BLOCK, VK_FORMAT_ASTC_10x5_UNORM_BLOCK, VK_FORMAT_ASTC_10x5_SRGB_BLOCK, VK_FORMAT_ASTC_10x6_UNORM_BLOCK, VK_FORMAT_ASTC_10x6_SRGB_BLOCK, VK_FORMAT_ASTC_10x10_UNORM_BLOCK, VK_FORMAT_ASTC_10x10_SRGB_BLOCK, VK_FORMAT_ASTC_12x10_UNORM_BLOCK, VK_FORMAT_ASTC_12x10_SRGB_BLOCK, VK_FORMAT_ASTC_12x12_UNORM_BLOCK, OR VK_FORMAT_ASTC_12x12_SRGB_BLOCK, VK_FORMAT_ASTC_12x12_UNORM_BLOCK, OR VK_FORMAT_ASTC_12x12_SRGB_BLOCK
- If the BC texture compression feature is not enabled, <code>format</code> must not be VK_FORMAT_BC1_RGB_UNORM_BLOCK, VK_FORMAT_BC1_RGB_SRGB_BLOCK, VK_FORMAT_BC1_RGBA_UNORM_BLOCK, VK_FORMAT_BC1_RGBA_SRGB_BLOCK, VK_FORMAT_BC2_UNORM_BLOCK, VK_FORMAT_BC2_SRGB_BLOCK, VK_FORMAT_BC3_UNORM_BLOCK, VK_FORMAT_BC3_SRGB_BLOCK, VK_FORMAT_BC4_UNORM_BLOCK, VK_FORMAT_BC5_UNORM_BLOCK, VK_FORMAT_BC5_SNORM_BLOCK, VK_FORMAT_BC6H_UFLOAT_BLOCK, VK_FORMAT_BC6H_SFLOAT_BLOCK, VK_FORMAT_BC7_UNORM_BLOCK, or VK_FORMAT_BC7_SRGB_BLOCK
- If the multisampled storage images feature is not enabled, and usage contains VK_IMAGE_USAGE_STORAGE_BIT, samples must be VK_SAMPLE_COUNT_1_BIT
- If the sparse bindings feature is not enabled, flags must not contain VK_IMAGE_CREATE_SPARSE_BINDING_BIT

- If the sparse residency for 2D images feature is not enabled, and <code>imageType</code> is VK_IMAGE_TYPE_2D, <code>flags</code> must not contain VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT
- If the sparse residency for 3D images feature is not enabled, and <code>imageType</code> is VK_IMAGE_TYPE_3D, <code>flags</code> must not contain VK IMAGE CREATE SPARSE RESIDENCY BIT
- If the sparse residency for images with 2 samples feature is not enabled, <code>imageType</code> is VK_IMAGE_TYPE_2D, and <code>samples</code> is VK_SAMPLE_COUNT_2_BIT, <code>flags</code> must not contain VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT
- If the sparse residency for images with 4 samples feature is not enabled, <code>imageType</code> is VK_IMAGE_TYPE_2D, and <code>samples</code> is VK_SAMPLE_COUNT_4_BIT, <code>flags</code> must not contain VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT
- If the sparse residency for images with 8 samples feature is not enabled, <code>imageType</code> is VK_IMAGE_TYPE_2D, and <code>samples</code> is VK_SAMPLE_COUNT_8_BIT, <code>flags</code> must not contain VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT
- If the sparse residency for images with 16 samples feature is not enabled, <code>imageType</code> is VK_IMAGE_TYPE_2D, and <code>samples</code> is VK_SAMPLE_COUNT_16_BIT, <code>flags</code> must not contain VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT
- If tiling is VK_IMAGE_TILING_LINEAR, format must be a format that has at least one supported feature bit present in the value of VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If tiling is VK_IMAGE_TILING_LINEAR, and VkFormatProperties::linearTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT, usage must not contain VK_IMAGE_USAGE_SAMPLED_BIT
- If tiling is VK_IMAGE_TILING_LINEAR, and VkFormatProperties::linearTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT, usage must not contain VK_IMAGE_USAGE_STORAGE_BIT
- If tiling is VK_IMAGE_TILING_LINEAR, and VkFormatProperties::linearTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT, usage must not contain VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT
- If tiling is VK_IMAGE_TILING_LINEAR, and VkFormatProperties::linearTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT, usage must not contain VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT
- If tiling is VK_IMAGE_TILING_OPTIMAL, format must be a format that has at least one supported feature bit present in the value of VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If tiling is VK_IMAGE_TILING_OPTIMAL, and VkFormatProperties::optimalTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT, usage must not contain VK_IMAGE_USAGE_SAMPLED_BIT

- If tiling is VK_IMAGE_TILING_OPTIMAL, and VkFormatProperties::optimalTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT, usage must not contain VK_IMAGE_USAGE_STORAGE_BIT
- If tiling is VK_IMAGE_TILING_OPTIMAL, and VkFormatProperties::optimalTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT, usage must not contain VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT
- If tiling is VK_IMAGE_TILING_OPTIMAL, and VkFormatProperties::optimalTilingFeatures (as returned by vkGetPhysicalDeviceFormatProperties with the same value of format) does not include VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT, usage must not contain VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT
- If flags contains VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT or VK_IMAGE_CREATE_SPARSE_ALIASED_BIT, it must also contain VK_IMAGE_CREATE_SPARSE_BINDING_BIT

If the *pNext* list includes a VkDedicatedAllocationImageCreateInfoNV structure, then that structure includes an enable controlling whether the image will have a dedicated memory allocation bound to it.

The VkDedicatedAllocationImageCreateInfoNV structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- · dedicatedAllocation indicates whether the image will have a dedicated allocation bound to it.



Note

Using a dedicated allocation for color and depth/stencil attachments or other large images may improve performance on some devices.

- sType must be VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_IMAGE_CREATE_INFO_NV
- pNext must be NULL

• If dedicatedAllocation is VK_TRUE, VkImageCreateInfo::flags must not include VK_IMAGE_CREATE_SPARSE_BINDING_BIT, VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT, or VK_IMAGE_CREATE_SPARSE_ALIASED_BIT

If the *pNext* list includes a VkExternalMemoryImageCreateInfoNV structure, then that structure defines a set of external memory handle types that may be used as backing store for the image.

The VkExternalMemoryImageCreateInfoNV structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- handleTypes is a bitmask of VkExternalMemoryHandleTypeFlagBitsNV specifying one or more external memory handle types. The types must all be compatible with each other and the other image creation parameters, as reported by vkGetPhysicalDeviceExternalImageFormatPropertiesNV.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_EXTERNAL_MEMORY_IMAGE_CREATE_INFO_NV
- pNext must be NULL
- handleTypes must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- handleTypes must not be 0

The intended usage of an image is specified by the bitmask VkImageCreateInfo::usage. Bits which can be set include:

```
typedef enum VkImageUsageFlagBits {
    VK_IMAGE_USAGE_TRANSFER_SRC_BIT = 0x00000001,
    VK_IMAGE_USAGE_TRANSFER_DST_BIT = 0x00000002,
    VK_IMAGE_USAGE_SAMPLED_BIT = 0x00000004,
    VK_IMAGE_USAGE_STORAGE_BIT = 0x00000008,
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT = 0x00000010,
    VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT = 0x000000020,
    VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT = 0x000000040,
    VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT = 0x000000080,
} VkImageUsageFlagBits;
```

These bits have the following meanings:

- VK_IMAGE_USAGE_TRANSFER_SRC_BIT indicates that the image can be used as the source of a transfer command.
- VK_IMAGE_USAGE_TRANSFER_DST_BIT indicates that the image can be used as the destination of a transfer command.
- VK_IMAGE_USAGE_SAMPLED_BIT indicates that the image can be used to create a VkImageView suitable for occupying a VkDescriptorSet slot either of type VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE or VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, and be sampled by a shader.
- VK_IMAGE_USAGE_STORAGE_BIT indicates that the image can be used to create a VkImageView suitable for occupying a VkDescriptorSet slot of type VK_DESCRIPTOR_TYPE_STORAGE_IMAGE.
- VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT indicates that the image can be used to create a VkImageView suitable for use as a color or resolve attachment in a VkFramebuffer.
- VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT indicates that the image can be used to create a VkImageView suitable for use as a depth/stencil attachment in a VkFramebuffer.
- VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT indicates that the memory bound to this image will have been allocated with the VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT (see Chapter 10 for more detail). This bit can be set for any image that can be used to create a VkImageView suitable for use as a color, resolve, depth/stencil, or input attachment.
- VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT indicates that the image can be used to create a VkImageView suitable for occupying VkDescriptorSet slot of type VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT; be read from a shader as an input attachment; and be used as an input attachment in a framebuffer.

Additional parameters of an image are specified by VkImageCreateInfo::flags. Bits which can be set include:

```
typedef enum VkImageCreateFlagBits {
   VK_IMAGE_CREATE_SPARSE_BINDING_BIT = 0x00000001,
   VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT = 0x00000002,
   VK_IMAGE_CREATE_SPARSE_ALIASED_BIT = 0x00000004,
   VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT = 0x00000008,
   VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT = 0x00000010,
} VkImageCreateFlagBits;
```

These bits have the following meanings:

- VK_IMAGE_CREATE_SPARSE_BINDING_BIT indicates that the image will be backed using sparse memory binding.
- VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT indicates that the image can be partially backed using sparse memory binding. Images created with this flag must also be created with the VK_IMAGE_CREATE_SPARSE_BINDING_BIT flag.
- VK_IMAGE_CREATE_SPARSE_ALIASED_BIT indicates that the image will be backed using sparse memory binding with memory ranges that might also simultaneously be backing another image (or another portion of the same image). Images created with this flag must also be created with the VK_IMAGE_CREATE_SPARSE_BINDING_BIT flag
- VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT indicates that the image can be used to create a VkImageView with a different format from the image.

• VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT indicates that the image can be used to create a VkImageView of type VK_IMAGE_VIEW_TYPE_CUBE or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY.

If any of the bits VK_IMAGE_CREATE_SPARSE_BINDING_BIT, VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT, or VK_IMAGE_CREATE_SPARSE_ALIASED_BIT are set, VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT must not also be set.

See Sparse Resource Features and Sparse Physical Device Features for more details.

The basic dimensionality of an image is specified by VkImageCreateInfo::imageType, which must be one of the values

```
typedef enum VkImageType {
    VK_IMAGE_TYPE_1D = 0,
    VK_IMAGE_TYPE_2D = 1,
    VK_IMAGE_TYPE_3D = 2,
} VkImageType;
```

These values specify one-, two-, or three-dimensional images, respectively.

The tiling arrangement of data elements in an image is specified by VkImageCreateInfo::tiling, which must be one of the values

```
typedef enum VkImageTiling {
   VK_IMAGE_TILING_OPTIMAL = 0,
   VK_IMAGE_TILING_LINEAR = 1,
} VkImageTiling;
```

VK_IMAGE_TILING_OPTIMAL specifies optimal tiling (texels are laid out in an implementation-dependent arrangement, for more optimal memory access), and VK_IMAGE_TILING_LINEAR specifies linear tiling (texels are laid out in memory in row-major order, possibly with some padding on each row).

To query the host access layout of an image subresource, for an image created with linear tiling, call:

- device is the logical device that owns the image.
- *image* is the image whose layout is being queried.
- pSubresource is a pointer to a VkImageSubresource structure selecting a specific image for the image subresource.
- pLayout points to a VkSubresourceLayout structure in which the layout is returned.

vkGetImageSubresourceLayout is invariant for the lifetime of a single image.

```
Valid Usage
```

- device must be a valid VkDevice handle
- image must be a valid VkImage handle
- pSubresource must be a pointer to a valid VkImageSubresource structure
- pLayout must be a pointer to a VkSubresourceLayout structure
- image must have been created, allocated, or retrieved from device
- image must have been created with tiling equal to VK_IMAGE_TILING_LINEAR
- The aspectMask member of pSubresource must only have a single bit set

The VkImageSubresource structure is defined as:

- aspectMask is a VkImageAspectFlags selecting the image aspect.
- mipLevel selects the mipmap level.
- arrayLayer selects the array layer.

Valid Usage

- aspectMask must be a valid combination of VkImageAspectFlagBits values
- aspectMask must not be 0
- mipLevel must be less than the mipLevels specified in VkImageCreateInfo when the image was created
- arrayLayer must be less than the arrayLayers specified in VkImageCreateInfo when the image was created

Information about the layout of the image subresource is returned in a VkSubresourceLayout structure:

```
typedef struct VkSubresourceLayout {
   VkDeviceSize offset;
   VkDeviceSize size;
   VkDeviceSize rowPitch;
   VkDeviceSize arrayPitch;
   VkDeviceSize depthPitch;
}
```

- offset is the byte offset from the start of the image where the image subresource begins.
- size is the size in bytes of the image subresource. size includes any extra memory that is required based on rowPitch.
- rowPitch describes the number of bytes between each row of texels in an image.
- arrayPitch describes the number of bytes between each array layer of an image.
- depthPitch describes the number of bytes between each slice of 3D image.

For images created with linear tiling, <code>rowPitch</code>, <code>arrayPitch</code> and <code>depthPitch</code> describe the layout of the image subresource in linear memory. For uncompressed formats, <code>rowPitch</code> is the number of bytes between texels with the same x coordinate in adjacent rows (y coordinates differ by one). <code>arrayPitch</code> is the number of bytes between texels with the same x and y coordinate in adjacent array layers of the image (array layer values differ by one). <code>depthPitch</code> is the number of bytes between texels with the same x and y coordinate in adjacent slices of a 3D image (z coordinates differ by one). Expressed as an addressing formula, the starting byte of a texel in the image subresource has address:

```
// (x,y,z,layer) are in texel coordinates
address(x,y,z,layer) = layer*arrayPitch + z*depthPitch + y*rowPitch + x*texelSize + ←
    offset
```

For compressed formats, the <code>rowPitch</code> is the number of bytes between compressed texel blocks in adjacent rows. <code>arrayPitch</code> is the number of bytes between compressed texel blocks in adjacent array layers. <code>depthPitch</code> is the number of bytes between compressed texel blocks in adjacent slices of a 3D image.

```
// (x,y,z,layer) are in compressed texel block coordinates
address(x,y,z,layer) = layer*arrayPitch + z*depthPitch + y*rowPitch + x* ←
    compressedTexelBlockByteSize + offset;
```

arrayPitch is undefined for images that were not created as arrays. depthPitch is defined only for 3D images.

For color formats, the <code>aspectMask</code> member of <code>VkImageSubresource</code> must be <code>VK_IMAGE_ASPECT_COLOR_BIT</code>. For depth/stencil formats, <code>aspectMask</code> must be either <code>VK_IMAGE_ASPECT_DEPTH_BIT</code> or <code>VK_IMAGE_ASPECT_STENCIL_BIT</code>. On implementations that store depth and stencil aspects separately, querying each of these image subresource layouts will return a different <code>offset</code> and <code>size</code> representing the region of memory used for that aspect. On implementations that store depth and stencil aspects interleaved, the same <code>offset</code> and <code>size</code> are returned and represent the interleaved memory allocation.

To destroy an image, call:

- device is the logical device that destroys the image.
- *image* is the image to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If image is not VK_NULL_HANDLE, image must be a valid VkImage handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If image is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to *image*, either directly or via a VkImageView, must have completed execution
- If VkAllocationCallbacks were provided when *image* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when image was created, pAllocator must be NULL

Host Synchronization

Host access to image must be externally synchronized

11.4 Image Layouts

Images are stored in implementation-dependent opaque layouts in memory. Implementations may support several opaque layouts, and the layout used at any given time is determined by the VkImageLayout state of the image subresource. Each layout has limitations on what kinds of operations are supported for image subresources using the layout. Applications have control over which layout each image subresource uses, and can transition an image subresource from one layout to another. Transitions can happen with an image memory barrier, included as part of a **vkCmdPipelineBarrier** or a **vkCmdWaitEvents** command buffer command (see Section 6.5.6), or as part of a subpass dependency within a render pass (see VkSubpassDependency). The image layout state is per-image subresource, and separate image subresources of the same image can be in different layouts at the same time with one exception - depth and stencil aspects of a given image subresource must always be in the same layout.

Note



Each layout may offer optimal performance for a specific usage of image memory. For example, an image with a layout of VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL may provide optimal performance for use as a color attachment, but be unsupported for use in transfer commands. Applications can transition an image subresource from one layout to another in order to achieve optimal performance when the image subresource is used for multiple kinds of operations. After initialization, applications need not use any layout other than the general layout, though this may produce suboptimal performance on some implementations.

Upon creation, all image subresources of an image are initially in the same layout, where that layout is selected by the VkImageCreateInfo::initialLayout member. The initialLayout must be either VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED. If it is VK_IMAGE_LAYOUT_PREINITIALIZED, then the image data can be preinitialized by the host while using this layout, and the transition away from this layout will preserve that data. If it is VK_IMAGE_LAYOUT_UNDEFINED, then the contents of the data are considered to be undefined, and the transition away from this layout is not guaranteed to preserve that data. For either of these initial layouts, any image subresources must be transitioned to another layout before they are accessed by the device.

Host access to image memory is only well-defined for images created with VK_IMAGE_TILING_LINEAR tiling and for image subresources of those images which are currently in either the VK_IMAGE_LAYOUT_PREINITIALIZED or VK_IMAGE_LAYOUT_GENERAL layout. Calling vkGetImageSubresourceLayout for a linear image returns a subresource layout mapping that is valid for either of those image layouts.

The set of image layouts consists of:

```
typedef enum VkImageLayout {
    VK_IMAGE_LAYOUT_UNDEFINED = 0,
    VK_IMAGE_LAYOUT_GENERAL = 1,
    VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL = 2,
    VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL = 3,
    VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL = 4,
    VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL = 5,
    VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL = 6,
    VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL = 7,
    VK_IMAGE_LAYOUT_PREINITIALIZED = 8,
    VK_IMAGE_LAYOUT_PRESENT_SRC_KHR = 1000001002,
} VkImageLayout;
```

The type(s) of device access supported by each layout are:

- VK_IMAGE_LAYOUT_UNDEFINED: Supports no device access. This layout must only be used as the initialLayout member of VkImageCreateInfo or VkAttachmentDescription, or as the oldLayout in an image transition. When transitioning out of this layout, the contents of the memory are not guaranteed to be preserved.
- VK_IMAGE_LAYOUT_PREINITIALIZED: Supports no device access. This layout must only be used as the initialLayout member of VkImageCreateInfo or VkAttachmentDescription, or as the oldLayout in an image transition. When transitioning out of this layout, the contents of the memory are preserved. This layout is intended to be used as the initial layout for an image whose contents are written by the host, and hence the data can be written to memory immediately, without first executing a layout transition. Currently, VK_IMAGE_LAYOUT_PREINITIALIZED is only useful with VK_IMAGE_TILING_LINEAR images because there is not a standard layout defined for VK_IMAGE_TILING_OPTIMAL images.
- VK IMAGE LAYOUT GENERAL: Supports all types of device access.
- VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL: must only be used as a color or resolve attachment in a VkFramebuffer. This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT usage bit enabled.
- VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL: must only be used as a depth/stencil attachment in a VkFramebuffer. This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT usage bit enabled.
- VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL: must only be used as a read-only depth/stencil attachment in a VkFramebuffer and/or as a read-only image in a shader (which can be read as a sampled image, combined image/sampler and/or input attachment). This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT_usage bit enabled.

- VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL: must only be used as a read-only image in a shader (which can be read as a sampled image, combined image/sampler and/or input attachment). This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_SAMPLED_BIT or VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT usage bit enabled.
- VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL: must only be used as a source image of a transfer command (see the definition of VK_PIPELINE_STAGE_TRANSFER_BIT). This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_TRANSFER_SRC_BIT usage bit enabled.
- VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL: must only be used as a destination image of a transfer command. This layout is valid only for image subresources of images created with the VK_IMAGE_USAGE_ TRANSFER_DST_BIT usage bit enabled.
- VK_IMAGE_LAYOUT_PRESENT_SRC_KHR: must only be used for presenting a swapchain image for display. A swapchain's image must be transitioned to this layout before calling vkQueuePresentKHR, and must be transitioned away from this layout after calling vkAcquireNextImageKHR.

For each mechanism of accessing an image in the API, there is a parameter or structure member that controls the image layout used to access the image. For transfer commands, this is a parameter to the command (see Chapter 17 and Chapter 18). For use as a framebuffer attachment, this is a member in the substructures of the VkRenderPassCreateInfo (see Render Pass). For use in a descriptor set, this is a member in the VkDescriptorImageInfo structure (see Section 13.2.4). At the time that any command buffer command accessing an image executes on any queue, the layouts of the image subresources that are accessed must all match the layout specified via the API controlling those accesses.

The image layout of each image subresource must be well-defined at each point in the image subresource's lifetime. This means that when performing a layout transition on the image subresource, the old layout value must either equal the current layout of the image subresource (at the time the transition executes), or else be VK_IMAGE_LAYOUT_
UNDEFINED (implying that the contents of the image subresource need not be preserved). The new layout used in a transition must not be VK_IMAGE_LAYOUT_UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED.

11.5 Image Views

Image objects are not directly accessed by pipeline shaders for reading or writing image data. Instead, *image views* representing contiguous ranges of the image subresources and containing additional metadata are used for that purpose. Views must be created on images of compatible types, and must represent a valid subset of image subresources.

Image views are represented by VkImageView handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkImageView)
```

The types of image views that can be created are:

```
typedef enum VkImageViewType {
    VK_IMAGE_VIEW_TYPE_1D = 0,
    VK_IMAGE_VIEW_TYPE_2D = 1,
    VK_IMAGE_VIEW_TYPE_3D = 2,
    VK_IMAGE_VIEW_TYPE_CUBE = 3,
    VK_IMAGE_VIEW_TYPE_1D_ARRAY = 4,
    VK_IMAGE_VIEW_TYPE_1D_ARRAY = 5,
    VK_IMAGE_VIEW_TYPE_2D_ARRAY = 6,
} VkImageViewType;
```

The exact image view type is partially implicit, based on the image's type and sample count, as well as the view creation parameters as described in the table below. This table also shows which SPIR-V OpTypeImage Dim and Arrayed parameters correspond to each image view type.

To create an image view, call:

- device is the logical device that creates the image view.
- pCreateInfo is a pointer to an instance of the VkImageViewCreateInfo structure containing parameters to be used to create the image view.
- pallocator controls host memory allocation as described in the Memory Allocation chapter.
- pView points to a VkImageView handle in which the resulting image view object is returned.

Some of the image creation parameters are inherited by the view. The remaining parameters are contained in the pCreateInfo.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkImageViewCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pView must be a pointer to a VkImageView handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkImageViewCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- *image* is a Vk Image on which the view will be created.
- viewType is the type of the image view.
- format is a VkFormat describing the format and type used to interpret data elements in the image.
- *components* specifies a remapping of color components (or of depth or stencil components after they have been converted into color components). See VkComponentMapping.
- subresourceRange is a VkImageSubresourceRange selecting the set of mipmap levels and array layers to be accessible to the view.

If *image* was created with the VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT flag, *format* can be different from the image's format, but if they are not equal they must be *compatible*. Image format compatibility is defined in the Format Compatibility Classes section.

		v requirements

Dim, Arrayed, MS	Image parameters	View parameters
1D, 0, 0	<pre>imageType = VK_IMAGE_TYPE_1D</pre>	<pre>viewType = VK_VIEW_TYPE_1D</pre>
	width >= 1	baseArrayLayer >= 0
	height = 1	layerCount = 1
	depth = 1	
	arrayLayers >= 1	
	samples = 1	
1D, 1, 0	<pre>imageType = VK_IMAGE_TYPE_1D</pre>	<pre>viewType = VK_VIEW_TYPE_1D_ARRAY</pre>
	width $>= 1$	baseArrayLayer >= 0
	height = 1	layerCount >= 1
	depth = 1	
	arrayLayers >= 1	
	samples = 1	
2D, 0, 0	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	<pre>viewType = VK_VIEW_TYPE_2D</pre>
	width $>= 1$	baseArrayLayer >= 0
	height >= 1	layerCount = 1
	depth = 1	
	arrayLayers >= 1	
	samples = 1	
1		· · · · · · · · · · · · · · · · · · ·

Table 11.1: (continued)

Dim, Arrayed, MS	Image parameters	View parameters
2D, 1, 0	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	viewType = VK_VIEW_TYPE_2D_ARRAY
	width >= 1	baseArrayLayer >= 0
	height >= 1	layerCount >= 1
	depth = 1	
	arrayLayers >= 1	
	samples = 1	
2D, 0, 1	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	<pre>viewType = VK_VIEW_TYPE_2D</pre>
	width $>= 1$	baseArrayLayer >= 0
	height >= 1	layerCount = 1
	depth = 1	
	arrayLayers >= 1	
	samples > 1	
2D, 1, 1	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	<pre>viewType = VK_VIEW_TYPE_2D_ARRAY</pre>
	width >= 1	baseArrayLayer >= 0
	height >= 1	layerCount >= 1
	depth = 1	
	arrayLayers >= 1	
	samples > 1	
CUBE, 0, 0	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	<pre>viewType = VK_VIEW_TYPE_CUBE</pre>
	width $>= 1$	baseArrayLayer >= 0
	height = width	layerCount = 6
	depth = 1	
	arrayLayers >= 6	
	samples = 1	
	flags include VK_IMAGE_CREATE_	
	CUBE_COMPATIBLE_BIT	
CUBE, 1, 0	<pre>imageType = VK_IMAGE_TYPE_2D</pre>	<pre>viewType = VK_VIEW_TYPE_CUBE_ARRAY</pre>
	width >= 1	baseArrayLayer >= 0
	height = width	N >= 1
	depth = 1	$layerCount = 6 \times N$
	N >= 1	
	$arrayLayers >= 6 \times N$	
	samples = 1	
	flags include VK_IMAGE_CREATE_	
	CUBE_COMPATIBLE_BIT	
3D, 0, 0	<pre>imageType = VK_IMAGE_TYPE_3D</pre>	<pre>viewType = VK_VIEW_TYPE_3D</pre>
	width >= 1	baseArrayLayer = 0
	height >= 1	layerCount = 1
	depth >= 1	
	arrayLayers = 1	
	samples = 1	

Valid Usage

- sType must be VK_STRUCTURE_TYPE_IMAGE_VIEW_CREATE_INFO
- pNext must be NULL
- flags must be 0
- image must be a valid VkImage handle
- viewType must be a valid VkImageViewType value
- format must be a valid VkFormat value
- components must be a valid VkComponentMapping structure
- subresourceRange must be a valid VkImageSubresourceRange structure
- If image was not created with VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT then viewType must not be VK IMAGE VIEW TYPE CUBE or VK IMAGE VIEW TYPE CUBE ARRAY
- If the image cubemap arrays feature is not enabled, viewType must not be VK_IMAGE_VIEW_TYPE_CUBE_ ARRAY
- If the ETC2 texture compression feature is not enabled, <code>format</code> must not be VK_FORMAT_ETC2_R8G8B8_ UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK, VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK, VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK, VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK, VK_FORMAT_EAC_R11_UNORM_BLOCK, VK_FORMAT_EAC_R11_SNORM_BLOCK, VK_FORMAT_EAC_R11G11_UNORM_BLOCK, or VK_FORMAT_EAC_R11G11_SNORM_BLOCK
- If the ASTC LDR texture compression feature is not enabled, format must not be VK_FORMAT_ASTC_4x4_ UNORM_BLOCK, VK_FORMAT_ASTC_4x4_ SRGB_BLOCK, VK_FORMAT_ASTC_5x4_UNORM_BLOCK, VK_FORMAT_ASTC_5x4_SRGB_BLOCK, VK_FORMAT_ASTC_5x5_UNORM_BLOCK, VK_FORMAT_ASTC_5x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x5_SRGB_BLOCK, VK_FORMAT_ASTC_6x6_UNORM_BLOCK, VK_FORMAT_ASTC_6x6_SRGB_BLOCK, VK_FORMAT_ASTC_8x5_UNORM_BLOCK, VK_FORMAT_ASTC_8x6_UNORM_BLOCK, VK_FORMAT_ASTC_8x6_UNORM_BLOCK, VK_FORMAT_ASTC_8x6_SRGB_BLOCK, VK_FORMAT_ASTC_8x8_UNORM_BLOCK, VK_FORMAT_ASTC_8x8_SRGB_BLOCK, VK_FORMAT_ASTC_10x5_UNORM_BLOCK, VK_FORMAT_ASTC_10x5_SRGB_BLOCK, VK_FORMAT_ASTC_10x6_SRGB_BLOCK, VK_FORMAT_ASTC_10x6_SRGB_BLOCK, VK_FORMAT_ASTC_10x10_UNORM_BLOCK, VK_FORMAT_ASTC_10x10_SRGB_BLOCK, VK_FORMAT_ASTC_12x10_UNORM_BLOCK, VK_FORMAT_ASTC_12x10_SRGB_BLOCK, VK_FORMAT_ASTC_12x12_UNORM_BLOCK, OR VK_FORMAT_ASTC_12x12_SRGB_BLOCK, VK_FORMAT_ASTC_12x12_UNORM_BLOCK, OR VK_FORMAT_ASTC_12x12_SRGB_BLOCK
- If the BC texture compression feature is not enabled, <code>format</code> must not be VK_FORMAT_BC1_RGB_UNORM_BLOCK, VK_FORMAT_BC1_RGB_SRGB_BLOCK, VK_FORMAT_BC1_RGBA_UNORM_BLOCK, VK_FORMAT_BC1_RGBA_SRGB_BLOCK, VK_FORMAT_BC2_UNORM_BLOCK, VK_FORMAT_BC2_SRGB_BLOCK, VK_FORMAT_BC3_UNORM_BLOCK, VK_FORMAT_BC3_SRGB_BLOCK, VK_FORMAT_BC4_UNORM_BLOCK, VK_FORMAT_BC5_UNORM_BLOCK, VK_FORMAT_BC5_SNORM_BLOCK, VK_FORMAT_BC6H_UFLOAT_BLOCK, VK_FORMAT_BC6H_SFLOAT_BLOCK, VK_FORMAT_BC7_UNORM_BLOCK, or VK_FORMAT_BC7_SRGB_BLOCK
- If image was created with VK_IMAGE_TILING_LINEAR, format must be format that has at least one supported feature bit present in the value of VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format

- If image was created with VK_IMAGE_TILING_LINEAR and usage containing VK_IMAGE_USAGE_
 SAMPLED_BIT, format must be supported for sampled images, as specified by the VK_FORMAT_FEATURE_
 SAMPLED_IMAGE_BIT flag in VkFormatProperties::linearTilingFeatures returned by
 vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_LINEAR and usage containing VK_IMAGE_USAGE_ STORAGE_BIT, format must be supported for storage images, as specified by the VK_FORMAT_FEATURE_ STORAGE_IMAGE_BIT flag in VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_LINEAR and usage containing VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, format must be supported for color attachments, as specified by the VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT flag in VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_LINEAR and usage containing VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT, format must be supported for depth/stencil attachments, as specified by the VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT flag in VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_OPTIMAL, format must be format that has at least one supported feature bit present in the value of VkFormatProperties::optimalTilingFeatures returned by

vkGetPhysicalDeviceFormatProperties with the same value of format

- If image was created with VK_IMAGE_TILING_OPTIMAL and usage containing VK_IMAGE_USAGE_ SAMPLED_BIT, format must be supported for sampled images, as specified by the VK_FORMAT_FEATURE_ SAMPLED_IMAGE_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_OPTIMAL and usage containing VK_IMAGE_USAGE_ STORAGE_BIT, format must be supported for storage images, as specified by the VK_FORMAT_FEATURE_ STORAGE_IMAGE_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_OPTIMAL and usage containing VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, format must be supported for color attachments, as specified by the VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- If image was created with VK_IMAGE_TILING_OPTIMAL and usage containing VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT, format must be supported for depth/stencil attachments, as specified by the VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties with the same value of format
- subresourceRange must be a valid image subresource range for image (see Section 11.5)
- If image was created with the VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT flag, format must be compatible with the format used to create image, as defined in Format Compatibility Classes
- If image was not created with the VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT flag, format must be identical to the format used to create image

subResourceRange and viewType must be compatible with the image, as described in the compatibility table

The VkImageSubresourceRange structure is defined as:

- aspectMask is a bitmask indicating which aspect(s) of the image are included in the view. See VkImageAspectFlagBits.
- baseMipLevel is the first mipmap level accessible to the view.
- levelCount is the number of mipmap levels (starting from baseMipLevel) accessible to the view.
- baseArrayLayer is the first array layer accessible to the view.
- layerCount is the number of array layers (starting from baseArrayLayer) accessible to the view.

The number of mipmap levels and array layers must be a subset of the image subresources in the image. If an application wants to use all mip levels or layers in an image after the <code>baseMipLevel</code> or <code>baseArrayLayer</code>, it can set <code>levelCount</code> and <code>layerCount</code> to the special values <code>VK_REMAINING_MIP_LEVELS</code> and <code>VK_REMAINING_ARRAY_LAYERS</code> without knowing the exact number of mip levels or layers.

For cube and cube array image views, the layers of the image view starting at baseArrayLayer correspond to faces in the order +X, -X, +Y, -Y, +Z, -Z. For cube arrays, each set of six sequential layers is a single cube, so the number of cube maps in a cube map array view is layerCount /6, and image array layer baseArrayLayer + i is face index $i \mod 6$ of cube i/6. If the number of layers in the view, whether set explicitly in layerCount or implied by VK_REMAINING_ARRAY LAYERS, is not a multiple of 6, behavior when indexing the last cube is undefined.

aspectMask is a bitmask indicating the format being used. Bits which may be set include:

```
typedef enum VkImageAspectFlagBits {
   VK_IMAGE_ASPECT_COLOR_BIT = 0x00000001,
   VK_IMAGE_ASPECT_DEPTH_BIT = 0x00000002,
   VK_IMAGE_ASPECT_STENCIL_BIT = 0x00000004,
   VK_IMAGE_ASPECT_METADATA_BIT = 0x00000008,
} VkImageAspectFlagBits;
```

The mask must be only VK_IMAGE_ASPECT_COLOR_BIT, VK_IMAGE_ASPECT_DEPTH_BIT or VK_IMAGE_ASPECT_STENCIL_BIT if format is a color, depth-only or stencil-only format, respectively. If using a depth/stencil format with both depth and stencil components, aspectMask must include at least one of VK_IMAGE_ASPECT_DEPTH_BIT and VK_IMAGE_ASPECT_STENCIL_BIT, and can include both.

When using an imageView of a depth/stencil image to populate a descriptor set (e.g. for sampling in the shader, or for use as an input attachment), the <code>aspectMask</code> must only include one bit and selects whether the imageView is used for depth reads (i.e. using a floating-point sampler or input attachment in the shader) or stencil reads (i.e. using an unsigned integer sampler or input attachment in the shader). When an imageView of a depth/stencil image is used as a depth/stencil framebuffer attachment, the <code>aspectMask</code> is ignored and both depth and stencil image subresources are used.

The *components* member is of type VkComponentMapping, and describes a remapping from components of the image to components of the vector returned by shader image instructions. This remapping must be identity for storage image descriptors, input attachment descriptors, and framebuffer attachments.

Valid Usage

- aspectMask must be a valid combination of VkImageAspectFlagBits values
- aspectMask must not be 0
- If levelCount is not VK_REMAINING_MIP_LEVELS, (baseMipLevel + levelCount) must be less than or equal to the mipLevels specified in VkImageCreateInfo when the image was created
- If <code>layerCount</code> is not <code>VK_REMAINING_ARRAY_LAYERS</code>, (<code>baseArrayLayer+layerCount</code>) must be less than or equal to the <code>arrayLayers</code> specified in <code>VkImageCreateInfo</code> when the image was created

The VkComponentMapping structure is defined as:

```
typedef struct VkComponentMapping {
   VkComponentSwizzle     r;
   VkComponentSwizzle     g;
   VkComponentSwizzle     b;
   VkComponentSwizzle     a;
} VkComponentMapping;
```

- r determines the component value placed in the R component of the output vector.
- q determines the component value placed in the G component of the output vector.
- b determines the component value placed in the B component of the output vector.
- a determines the component value placed in the A component of the output vector.

Each of r, g, b, and a is one of the values:

```
typedef enum VkComponentSwizzle {
    VK_COMPONENT_SWIZZLE_IDENTITY = 0,
    VK_COMPONENT_SWIZZLE_ZERO = 1,
    VK_COMPONENT_SWIZZLE_ONE = 2,
    VK_COMPONENT_SWIZZLE_R = 3,
    VK_COMPONENT_SWIZZLE_R = 4,
    VK_COMPONENT_SWIZZLE_B = 5,
    VK_COMPONENT_SWIZZLE_B = 6,
} VkComponentSwizzle;
```

- VK_COMPONENT_SWIZZLE_IDENTITY: the component is set to the identity swizzle.
- VK_COMPONENT_SWIZZLE_ZERO: the component is set to zero.

- VK_COMPONENT_SWIZZLE_ONE: the component is set to either 1 or 1.0 depending on whether the type of the image view format is integer or floating-point respectively, as determined by the Format Definition section for each VkFormat.
- VK_COMPONENT_SWIZZLE_R: the component is set to the value of the R component of the image.
- VK_COMPONENT_SWIZZLE_G: the component is set to the value of the G component of the image.
- VK_COMPONENT_SWIZZLE_B: the component is set to the value of the B component of the image.
- VK COMPONENT SWIZZLE A: the component is set to the value of the A component of the image.

Setting the identity swizzle on a component is equivalent to setting the identity mapping on that component. That is:

Table 11.2: Component Mappings Equivalent To VK_COMPONENT_SW IZZLE_IDENTITY

Component	Identity Mapping
components.r	VK_COMPONENT_SWIZZLE_R
components.g	VK_COMPONENT_SWIZZLE_G
components.b	VK_COMPONENT_SWIZZLE_B
components.a	VK_COMPONENT_SWIZZLE_A

Valid Usage

- r must be a valid VkComponentSwizzle value
- g must be a valid VkComponentSwizzle value
- \bullet \emph{b} must be a valid <code>VkComponentSwizzle</code> value
- a must be a valid VkComponentSwizzle value

To destroy an image view, call:

- device is the logical device that destroys the image view.
- *imageView* is the image view to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If imageView is not VK_NULL_HANDLE, imageView must be a valid VkImageView handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If imageView is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to imageView must have completed execution
- If VkAllocationCallbacks were provided when <code>imageView</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when imageView was created, pAllocator must be NULL

Host Synchronization

• Host access to imageView must be externally synchronized

11.6 Resource Memory Association

Resources are initially created as *virtual allocations* with no backing memory. Device memory is allocated separately (see Section 10.2) and then associated with the resource. This association is done differently for sparse and non-sparse resources.

Resources created with any of the sparse creation flags are considered sparse resources. Resources created without these flags are non-sparse. The details on resource memory association for sparse resources is described in Chapter 28.

Non-sparse resources must be bound completely and contiguously to a single VkDeviceMemory object before the resource is passed as a parameter to any of the following operations:

- · creating image or buffer views
- updating descriptor sets
- recording commands in a command buffer

Once bound, the memory binding is immutable for the lifetime of the resource.

To determine the memory requirements for a buffer resource, call:

- device is the logical device that owns the buffer.
- buffer is the buffer to query.
- pMemoryRequirements points to an instance of the VkMemoryRequirements structure in which the memory requirements of the buffer object are returned.

Valid Usage

- device must be a valid VkDevice handle
- buffer must be a valid VkBuffer handle
- pMemoryRequirements must be a pointer to a VkMemoryRequirements structure
- buffer must have been created, allocated, or retrieved from device

To determine the memory requirements for an image resource, call:

- device is the logical device that owns the image.
- image is the image to query.
- pMemoryRequirements points to an instance of the VkMemoryRequirements structure in which the memory requirements of the image object are returned.

Valid Usage

- device must be a valid VkDevice handle
- image must be a valid VkImage handle
- pMemoryRequirements must be a pointer to a VkMemoryRequirements structure
- image must have been created, allocated, or retrieved from device

The VkMemoryRequirements structure is defined as:

```
typedef struct VkMemoryRequirements {
    VkDeviceSize    size;
    VkDeviceSize    alignment;
    uint32_t    memoryTypeBits;
} VkMemoryRequirements;
```

- size is the size, in bytes, of the memory allocation required for the resource.
- alignment is the alignment, in bytes, of the offset within the allocation required for the resource.
- memoryTypeBits is a bitmask and contains one bit set for every supported memory type for the resource. Bit i is set if and only if the memory type i in the VkPhysicalDeviceMemoryProperties structure for the physical device is supported for the resource.

The implementation guarantees certain properties about the memory requirements returned by vkGetBufferMemoryRequirements and vkGetImageMemoryRequirements:

- The memoryTypeBits member always contains at least one bit set.
- If buffer is a VkBuffer, or if image is a VkImage that was created with a VK_IMAGE_TILING_LINEAR value in the tiling member of the VkImageCreateInfo structure passed to **vkCreateImage**, then the memoryTypeBits member always contains at least one bit set corresponding to a VkMemoryType with a propertyFlags that has both the VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT bit and the VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT bit and the VK_MEMORY_PROPERTY_HOST_COHERENT_BIT bit set. In other words, mappable coherent memory can always be attached to these objects.
- The memoryTypeBits member always contains at least one bit set corresponding to a VkMemoryType with a propertyFlags that has the VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT bit set.
- The memoryTypeBits member is identical for all VkBuffer objects created with the same value for the flags and usage members in the VkBufferCreateInfo structure passed to vkCreateBuffer. Further, if usage1 and usage2 of type VkBufferUsageFlags are such that the bits set in usage2 are a subset of the bits set in usage1, and they have the same flags, then the bits set in memoryTypeBits returned for usage1 must be a subset of the bits set in memoryTypeBits returned for usage2, for all values of flags.
- The alignment member is identical for all VkBuffer objects created with the same combination of values for the usage and flags members in the VkBufferCreateInfo structure passed to **vkCreateBuffer**.
- The memoryTypeBits member is identical for all VkImage objects created with the same combination of values for the tiling member and the VK_IMAGE_CREATE_SPARSE_BINDING_BIT bit of the flags member and the VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT of the usage member in the VkImageCreateInfo structure passed to vkCreateImage.
- If the memory requirements are for a VkImage, the memoryTypeBits member must not refer to a VkMemoryType with a propertyFlags that has the VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT bit set if the vkGetImageMemoryRequirements::image did not have VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT bit set in the usage member of the VkImageCreateInfo structure passed to vkCreateImage.
- If the memory requirements are for a VkBuffer, the memoryTypeBits member must not refer to a VkMemoryType with a propertyFlags that has the VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT bit set.



Note

The implication of this requirement is that lazily allocated memory is disallowed for buffers in all cases.

To attach memory to a buffer object, call:

```
VkResult vkBindBufferMemory(

VkDevice

VkBuffer

VkDeviceMemory

VkDeviceSize

device,

buffer,

memory,

memory,

memoryOffset);
```

- device is the logical device that owns the buffer and memory.
- buffer is the buffer.
- memory is a VkDeviceMemory object describing the device memory to attach.
- memoryOffset is the start offset of the region of memory which is to be bound to the buffer. The number of bytes returned in the VkMemoryRequirements::size member in memory, starting from memoryOffset bytes, will be bound to the specified buffer.

Valid Usage

- device must be a valid VkDevice handle
- buffer must be a valid VkBuffer handle
- memory must be a valid VkDeviceMemory handle
- buffer must have been created, allocated, or retrieved from device
- memory must have been created, allocated, or retrieved from device
- buffer must not already be backed by a memory object
- buffer must not have been created with any sparse memory binding flags
- memoryOffset must be less than the size of memory
- If buffer was created with the VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT or VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT, memoryOffset must be a multiple of VkPhysicalDeviceLimits::minTexelBufferOffsetAlignment
- If buffer was created with the VK_BUFFER_USAGE_UNIFORM_BUFFER_BIT, memoryOffset must be a multiple of VkPhysicalDeviceLimits::minUniformBufferOffsetAlignment
- If buffer was created with the VK_BUFFER_USAGE_STORAGE_BUFFER_BIT, memoryOffset must be a multiple of VkPhysicalDeviceLimits::minStorageBufferOffsetAlignment

- memory must have been allocated using one of the memory types allowed in the memoryTypeBits member of the VkMemoryRequirements structure returned from a call to vkGetBufferMemoryRequirements with buffer
- memoryOffset must be an integer multiple of the alignment member of the VkMemoryRequirements structure returned from a call to vkGetBufferMemoryRequirements with buffer
- The size member of the VkMemoryRequirements structure returned from a call to **vkGetBufferMemoryRequirements** with buffer must be less than or equal to the size of memory minus memoryOffset
- If buffer was created with VkDedicatedAllocationBufferCreateInfoNV::dedicatedAllocation equal to VK_TRUE, memory must have been created with VkDedicatedAllocationMemoryAllocateInfoNV::buffer equal to buffer and memoryOffset must be zero
- If buffer was not created with VkDedicatedAllocationBufferCreateInfoNV::dedicatedAllocation equal to VK_TRUE, memory must not have been allocated dedicated for a specific buffer or image

Host Synchronization

• Host access to buffer must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To attach memory to an image object, call:

- device is the logical device that owns the image and memory.
- image is the image.
- memory is the a VkDeviceMemory object describing the device memory to attach.
- memoryOffset is the start offset of the region of memory which is to be bound to the image. The number of bytes returned in the VkMemoryRequirements::size member in memory, starting from memoryOffset bytes, will be bound to the specified image.

Valid Usage

- device must be a valid VkDevice handle
- image must be a valid Vk Image handle
- memory must be a valid VkDeviceMemory handle
- image must have been created, allocated, or retrieved from device
- memory must have been created, allocated, or retrieved from device
- image must not already be backed by a memory object
- image must not have been created with any sparse memory binding flags
- memoryOffset must be less than the size of memory
- memory must have been allocated using one of the memory types allowed in the memoryTypeBits member of the VkMemoryRequirements structure returned from a call to **vkGetImageMemoryRequirements** with image
- memoryOffset must be an integer multiple of the alignment member of the VkMemoryRequirements structure returned from a call to vkGetImageMemoryRequirements with image
- The size member of the VkMemoryRequirements structure returned from a call to **vkGetImageMemoryRequirements** with image must be less than or equal to the size of memory minus memoryOffset
- If image was created with VkDedicatedAllocationImageCreateInfoNV::dedicatedAllocation equal to VK_TRUE, memory must have been created with VkDedicatedAllocationMemoryAllocateInfoNV::image equal to image and memoryOffset must be zero
- If image was not created with VkDedicatedAllocationImageCreateInfoNV::dedicatedAllocation equal to VK_TRUE, memory must not have been allocated dedicated for a specific buffer or image

Host Synchronization

Host access to image must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

Buffer-Image Granularity

There is an implementation-dependent limit, <code>bufferImageGranularity</code>, which specifies a page-like granularity at which buffer, linear image and optimal image resources must be placed in adjacent memory locations to avoid aliasing. Two resources which do not satisfy this granularity requirement are said to alias. Linear image resource are images created with <code>VK_IMAGE_TILING_LINEAR</code> and optimal image resources are those created with <code>VK_IMAGE_TILING_LINEAR</code> and optimal image resources are those created with <code>VK_IMAGE_TILING_OPTIMAL</code>. <code>bufferImageGranularity</code> is specified in bytes, and must be a power of two. Implementations which do not require such an additional granularity may report a value of one.



Note

bufferImageGranularity is really a granularity between "linear" resources, including buffers and images with linear tiling, vs. "optimal" resources, i.e. images with optimal tiling. It would have been better named "linearOptimalGranularity".

Given resourceA at the lower memory offset and resourceB at the higher memory offset in the same VkDeviceMemory object, where one of the resources is a buffer or a linear image and the other is an optimal image, and the following:

```
resourceA.end = resourceA.memoryOffset + resourceA.size - 1
resourceA.endPage = resourceA.end & ~(bufferImageGranularity-1)
resourceB.start = resourceB.memoryOffset
resourceB.startPage = resourceB.start & ~(bufferImageGranularity-1)
```

The following property must hold:

```
resourceA.endPage < resourceB.startPage
```

That is, the end of the first resource (A) and the beginning of the second resource (B) must be on separate "pages" of size <code>bufferImageGranularity</code>. <code>bufferImageGranularity</code> may be different than the physical page size of the memory heap. This restriction is only needed when a buffer or a linear image is at adjacent memory location with an optimal

image and both will be used simultaneously. Adjacent buffers' or adjacent images' memory ranges can be closer than bufferImageGranularity, provided they meet the alignment requirement for the objects in question.

Sparse block size in bytes and sparse image and buffer memory alignments must all be multiples of the bufferImageGranularity. Therefore, memory bound to sparse resources naturally satisfies the bufferImageGranularity.

11.7 Resource Sharing Mode

Buffer and image objects are created with a *sharing mode* controlling how they can be accessed from queues. The supported sharing modes are:

```
typedef enum VkSharingMode {
    VK_SHARING_MODE_EXCLUSIVE = 0,
    VK_SHARING_MODE_CONCURRENT = 1,
} VkSharingMode;
```

- VK_SHARING_MODE_EXCLUSIVE specifies that access to any range or image subresource of the object will be exclusive to a single queue family at a time.
- VK_SHARING_MODE_CONCURRENT specifies that concurrent access to any range or image subresource of the object from multiple queue families is supported.



Note

VK_SHARING_MODE_CONCURRENT may result in lower performance access to the buffer or image than VK_SHARING_MODE_EXCLUSIVE.

Ranges of buffers and image subresources of image objects created using VK_SHARING_MODE_EXCLUSIVE must only be accessed by queues in the same queue family at any given time. In order for a different queue family to be able to interpret the memory contents of a range or image subresource, the application must transfer exclusive ownership of the range or image subresource between the source and destination queue families with the following sequence of operations:

- 1. Release exclusive ownership from the source queue family to the destination queue family.
- 2. Use semaphores to ensure proper execution control for the ownership transfer.
- 3. Acquire exclusive ownership for the destination queue family from the source queue family.

To release exclusive ownership of a range of a buffer or image subresource of an image object, the application must execute a buffer or image memory barrier, respectively (see VkBufferMemoryBarrier and VkImageMemoryBarrier) on a queue from the source queue family. The <code>srcQueueFamilyIndex</code> parameter of the barrier must be set to the source queue family index, and the <code>dstQueueFamilyIndex</code> parameter to the destination queue family index.

To acquire exclusive ownership, the application must execute the same buffer or image memory barrier on a queue from the destination queue family.

Upon creation, resources using VK_SHARING_MODE_EXCLUSIVE are not owned by any queue family. A buffer or image memory barrier is not required to acquire ownership when no queue family owns the resource - it is implicitly acquired upon first use within a queue. However, images still require a layout transition from VK_IMAGE_LAYOUT_

UNDEFINED or VK_IMAGE_LAYOUT_PREINITIALIZED before being used on the first queue. This layout transition can either be accomplished by an image memory barrier or by use in a render pass instance.

Once a queue family has used a range or image subresource of an VK_SHARING_MODE_EXCLUSIVE resource, its contents are undefined to other queue families unless ownership is transferred. The contents may also become undefined for other reasons, e.g. as a result of writes to an image subresource that aliases the same memory. A queue family can take ownership of a range or image subresource without an ownership transfer in the same way as for a resource that was just created, however doing so means any contents written by other queue families or via incompatible aliases are undefined.

11.8 Memory Aliasing

A range of a VkDeviceMemory allocation is *aliased* if it is bound to multiple resources simultaneously, via vkBindImageMemory, vkBindBufferMemory, or via sparse memory bindings. A memory range aliased between two images or two buffers is defined to be the intersection of the memory ranges bound to the two resources. A memory range aliased between two resources where one is a buffer or a linear image, and the other is an optimal image, is defined to be the intersection of the memory ranges bound to the two resources, where each range is first padded to be aligned to the bufferImageGranularity. Applications can alias memory, but use of multiple aliases is subject to several constraints.



Note

Memory aliasing can be useful to reduce the total device memory footprint of an application, if some large resources are used for disjoint periods of time.

When an opaque, non-VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT image is bound to an aliased range, all image subresources of the image *overlap* the range. When a linear image is bound to an aliased range, the image subresources that (according to the image's advertised layout) include bytes from the aliased range overlap the range. When a VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT image has sparse image blocks bound to an aliased range, only image subresources including those sparse image blocks overlap the range, and when the memory bound to the image's mip tail overlaps an aliased range all image subresources in the mip tail overlap the range.

Buffers, and linear image subresources in either the VK_IMAGE_LAYOUT_PREINITIALIZED or VK_IMAGE_LAYOUT_GENERAL layouts, are *host-accessible subresources*. That is, the host has a well-defined addressing scheme to interpret the contents, and thus the layout of the data in memory can be consistently interpreted across aliases if each of those aliases is a host-accessible subresource. Opaque images and linear image subresources in other layouts are not host-accessible.

If two aliases are both host-accessible, then they interpret the contents of the memory in consistent ways, and data written to one alias can be read by the other alias.

If either of two aliases is not host-accessible, then the aliases interpret the contents of the memory differently, and writes via one alias make the contents of memory partially or completely undefined to the other alias. If the first alias is a host-accessible subresource, then the bytes affected are those written by the memory operations according to its addressing scheme. If the first alias is not host-accessible, then the bytes affected are those overlapped by the image subresources that were written. If the second alias is a host-accessible subresource, the affected bytes become undefined. If the second alias is a not host-accessible, all sparse image blocks (for sparse partially-resident images) or all image subresources (for non-sparse image and fully resident sparse images) that overlap the affected bytes become undefined.

If any image subresources are made undefined due to writes to an alias, then each of those image subresources must have its layout transitioned from VK_IMAGE_LAYOUT_UNDEFINED to a valid layout before it is used, or from VK_

IMAGE_LAYOUT_PREINITIALIZED if the memory has been written by the host. If any sparse blocks of a sparse image have been made undefined, then only the image subresources containing them must be transitioned.

Use of an overlapping range by two aliases must be separated by a memory dependency using the appropriate access types if at least one of those uses performs writes, whether the aliases interpret memory consistently or not. If buffer or image memory barriers are used, the scope of the barrier must contain the entire range and/or set of image subresources that overlap.

If two aliasing image views are used in the same framebuffer, then the render pass must declare the attachments using the VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT, and follow the other rules listed in that section.

Access to resources which alias memory from shaders using variables decorated with **Coherent** are not automatically coherent with each other.



Note

Memory recycled via an application suballocator (i.e. without freeing and reallocating the memory objects) is not substantially different from memory aliasing. However, a suballocator usually waits on a fence before recycling a region of memory, and signaling a fence involves sufficient implicit dependencies to satisfy all the above requirements.

Chapter 12

Samplers

VkSampler objects represent the state of an image sampler which is used by the implementation to read image data and apply filtering and other transformations for the shader.

Samplers are represented by VkSampler handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkSampler)
```

To create a sampler object, call:

- device is the logical device that creates the sampler.
- pCreateInfo is a pointer to an instance of the VkSamplerCreateInfo structure specifying the state of the sampler object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pSampler points to a VkSampler handle in which the resulting sampler object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkSamplerCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSampler must be a pointer to a VkSampler handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_TOO_MANY_OBJECTS

The VkSamplerCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- magFilter is the magnification filter to apply to lookups, and is of type:

```
typedef enum VkFilter {
    VK_FILTER_NEAREST = 0,
    VK_FILTER_LINEAR = 1,
    VK_FILTER_CUBIC_IMG = 1000015000,
} VkFilter;
```

• minFilter is the minification filter to apply to lookups, and is of type VkFilter.

• mipmapMode is the mipmap filter to apply to lookups as described in the Texel Filtering section, and is of type:

```
typedef enum VkSamplerMipmapMode {
   VK_SAMPLER_MIPMAP_MODE_NEAREST = 0,
   VK_SAMPLER_MIPMAP_MODE_LINEAR = 1,
} VkSamplerMipmapMode;
```

- addressModeU is the addressing mode for outside [0..1] range for U coordinate. See VkSamplerAddressMode.
- $\bullet \ \ \textit{addressModeV} \ is \ the \ addressing \ mode \ for \ outside \ [0..1] \ range \ for \ V \ coordinate. \ See \ VkSamplerAddressMode.$
- addressModeW is the addressing mode for outside [0..1] range for W coordinate. See VkSamplerAddressMode.
- mipLodBias is the bias to be added to mipmap LOD calculation and bias provided by image sampling functions in SPIR-V, as described in the Level-of-Detail Operation section.
- anisotropyEnable is VK_TRUE to enable anisotropic filtering, as described in the Texel Anisotropic Filtering section, or VK_FALSE otherwise.
- maxAnisotropy is the anisotropy value clamp.
- compareEnable is VK_TRUE to enable comparison against a reference value during lookups, or VK_FALSE otherwise.
 - Note: Some implementations will default to shader state if this member does not match.
- compareOp is the comparison function to apply to fetched data before filtering as described in the Depth Compare Operation section. See VkCompareOp.
- minLod and maxLod are the values used to clamp the computed level-of-detail value, as described in the Level-of-Detail Operation section. maxLod must be greater than or equal to minLod.
- borderColor is the predefined border color to use, as described in the Texel Replacement section, and is of type:

```
typedef enum VkBorderColor {
    VK_BORDER_COLOR_FLOAT_TRANSPARENT_BLACK = 0,
    VK_BORDER_COLOR_INT_TRANSPARENT_BLACK = 1,
    VK_BORDER_COLOR_FLOAT_OPAQUE_BLACK = 2,
    VK_BORDER_COLOR_INT_OPAQUE_BLACK = 3,
    VK_BORDER_COLOR_FLOAT_OPAQUE_WHITE = 4,
    VK_BORDER_COLOR_INT_OPAQUE_WHITE = 5,
} VkBorderColor;
```

- unnormalizedCoordinates controls whether to use unnormalized or normalized texel coordinates to address texels of the image. When set to VK_TRUE, the range of the image coordinates used to lookup the texel is in the range of zero to the image dimensions for x, y and z. When set to VK_FALSE the range of image coordinates is zero to one. When unnormalizedCoordinates is VK_TRUE, samplers have the following requirements:
- minFilter and magFilter must be equal.
- mipmapMode must be VK_SAMPLER_MIPMAP_MODE_NEAREST.
- minLod and maxLod must be zero.
- addressModeU and addressModeV must each be either VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE or VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER.
- anisotropyEnable must be VK FALSE.
- compareEnable must be VK_FALSE.

- When unnormalizedCoordinates is VK_TRUE, images the sampler is used with in the shader have the following requirements:
 - The viewType must be either VK_IMAGE_VIEW_TYPE_1D or VK_IMAGE_VIEW_TYPE_2D.
 - The image view must have a single layer and a single mip level.
- When unnormalizedCoordinates is VK_TRUE, image built-in functions in the shader that use the sampler have the following requirements:
 - The functions must not use projection.
 - The functions must not use offsets.

Mapping of OpenGL to Vulkan filter modes

magFilter values of VK_FILTER_NEAREST and VK_FILTER_LINEAR directly correspond to **GL_NEA REST** and **GL_LINEAR** magnification filters. minFilter and mipmapMode combine to correspond to the
similarly named OpenGL minification filter of **GL_minFilter_MIPMAP_mipmapMode** (e.g. minFilter of
VK_FILTER_LINEAR and mipmapMode of VK_SAMPLER_MIPMAP_MODE_NEAREST correspond to **GL_LINEAR_MIPMAP_NEAREST**).



There are no Vulkan filter modes that directly correspond to OpenGL minification filters of **GL_LINEAR** or **GL_NEAREST**, but they can be emulated using VK_SAMPLER_MIPMAP_MODE_NEAREST, minLod = 0, and maxLod = 0.25, and using $minFilter = VK_FILTER_LINEAR$ or $minFilter = VK_FILTER_NEAREST$, respectively.

Note that using a maxLod of zero would cause magnification to always be performed, and the magFilter to always be used. This is valid, just not an exact match for OpenGL behavior. Clamping the maximum LOD to 0.25 allows the λ value to be non-zero and minification to be performed, while still always rounding down to the base level. If the minFilter and magFilter are equal, then using a maxLod of zero also works.

 ${\it address} {\it ModeV}, and {\it address} {\it ModeW} \ must \ each \ have \ one \ of \ the \ following \ values:$

```
typedef enum VkSamplerAddressMode {
   VK_SAMPLER_ADDRESS_MODE_REPEAT = 0,
   VK_SAMPLER_ADDRESS_MODE_MIRRORED_REPEAT = 1,
   VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE = 2,
   VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER = 3,
   VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE = 4,
} VkSamplerAddressMode;
```

These values control the behavior of sampling with coordinates outside the range [0,1] for the respective u, v, or w coordinate as defined in the Wrapping Operation section.

- VK_SAMPLER_ADDRESS_MODE_REPEAT indicates that the repeat wrap mode will be used.
- VK_SAMPLER_ADDRESS_MODE_MIRRORED_REPEAT indicates that the mirrored repeat wrap mode will be used.
- VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE indicates that the clamp to edge wrap mode will be used.
- VK SAMPLER ADDRESS MODE CLAMP TO BORDER indicates that the clamp to border wrap mode will be used.
- VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE indicates that the mirror clamp to edge wrap mode will be used. This is only valid if the VK_KHR_mirror_clamp_to_edge extension is enabled.

The maximum number of sampler objects which can be simultaneously created on a device is implementation-dependent and specified by the maxSamplerAllocationCount member of the VkPhysicalDeviceLimits structure. If maxSamplerAllocationCount is exceeded, vkCreateSampler will return VK_ERROR_TOO_MANY_OBJECTS.

Since VkSampler is a non-dispatchable handle type, implementations may return the same handle for sampler state vectors that are identical. In such cases, all such objects would only count once against the <code>maxSamplerAllocationCount limit</code>.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_SAMPLER_CREATE_INFO
- pNext must be NULL
- flags must be 0
- magFilter must be a valid VkFilter value
- minFilter must be a valid VkFilter value
- mipmapMode must be a valid VkSamplerMipmapMode value
- addressModeU must be a valid VkSamplerAddressMode value
- addressModeV must be a valid VkSamplerAddressMode value
- addressModeW must be a valid VkSamplerAddressMode value
- The absolute value of mipLodBias must be less than or equal to
 VkPhysicalDeviceLimits::maxSamplerLodBias
- If the anisotropic sampling feature is not enabled, anisotropyEnable must be VK FALSE
- If anisotropyEnable is VK_TRUE, maxAnisotropy must be between 1.0 and VkPhysicalDeviceLimits::maxSamplerAnisotropy, inclusive
- If unnormalizedCoordinates is VK_TRUE, minFilter and magFilter must be equal
- If unnormalizedCoordinates is VK_TRUE, mipmapMode must be VK_SAMPLER_MIPMAP_MODE_ NEAREST
- If ${\tt unnormalizedCoordinates}$ is ${\tt VK_TRUE}, {\tt minLod}$ and ${\tt maxLod}$ must be zero
- If unnormalizedCoordinates is VK_TRUE, addressModeU and addressModeV must each be either VK_ SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE or VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER
- If unnormalizedCoordinates is VK_TRUE, anisotropyEnable must be VK_FALSE
- $\bullet \ \, \text{If unnormalizedCoordinates is VK_TRUE, compareEnable must be VK_FALSE } \\$
- If any of addressModeU, addressModeV or addressModeW are VK_SAMPLER_ADDRESS_MODE_CLAMP_ TO BORDER, borderColor must be a valid VkBorderColor value
- If the VK_KHR_sampler_mirror_clamp_to_edge extension is not enabled, addressModeU, addressModeV and addressModeW must not be VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE

- If compareEnable is VK_TRUE, compareOp must be a valid VkCompareOp value
- If either magFilter or minFilter is VK_FILTER_CUBIC_IMG, anisotropyEnable must be VK_FALSE

To destroy a sampler, call:

- device is the logical device that destroys the sampler.
- sampler is the sampler to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If sampler is not VK_NULL_HANDLE, sampler must be a valid VkSampler handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If sampler is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to sampler must have completed execution
- If VkAllocationCallbacks were provided when sampler was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when sampler was created, pAllocator must be NULL

Host Synchronization

• Host access to sampler must be externally synchronized

Chapter 13

Resource Descriptors

Shaders access buffer and image resources by using special shader variables which are indirectly bound to buffer and image views via the API. These variables are organized into sets, where each set of bindings is represented by a *descriptor set* object in the API and a descriptor set is bound all at once. A *descriptor* is an opaque data structure representing a shader resource such as a buffer view, image view, sampler, or combined image sampler. The content of each set is determined by its *descriptor set layout* and the sequence of set layouts that can be used by resource variables in shaders within a pipeline is specified in a *pipeline layout*.

Each shader can use up to <code>maxBoundDescriptorSets</code> (see Limits) descriptor sets, and each descriptor set can include bindings for descriptors of all descriptor types. Each shader resource variable is assigned a tuple of (set number, binding number, array element) that defines its location within a descriptor set layout. In GLSL, the set number and binding number are assigned via layout qualifiers, and the array element is implicitly assigned consecutively starting with index equal to zero for the first element of an array (and array element is zero for non-array variables):

GLSL example

```
// Assign set number = M, binding number = N, array element = 0
layout (set=M, binding=N) uniform sampler2D variableName;

// Assign set number = M, binding number = N for all array elements, and
// array element = I for the I'th member of the array.
layout (set=M, binding=N) uniform sampler2D variableNameArray[I];
```

SPIR-V example

```
// Assign set number = M, binding number = N for all array elements, and
// array element = I for the I'th member of the array.
          %1 = OpExtInstImport "GLSL.std.450"
               OpName %13 "variableNameArray"
               OpDecorate %13 DescriptorSet M
               OpDecorate %13 Binding N
          %2 = OpTvpeVoid
          %3 = OpTypeFunction %2
          %6 = OpTypeFloat 32
          %7 = OpTypeImage %6 2D 0 0 1 Unknown
         %8 = OpTypeSampledImage %7
         %9 = OpTypeInt 32 0
         %10 = OpConstant %9 I
         %11 = OpTypeArray %8 %10
         %12 = OpTypePointer UniformConstant %11
         %13 = OpVariable %12 UniformConstant
```

13.1 Descriptor Types

The following sections outline the various descriptor types supported by Vulkan. Each section defines a descriptor type, and each descriptor type has a manifestation in the shading language and SPIR-V as well as in descriptor sets. There is mostly a one-to-one correspondence between descriptor types and classes of opaque types in the shading language, where the opaque types in the shading language must refer to a descriptor in the pipeline layout of the corresponding descriptor type. But there is an exception to this rule as described in Combined Image Sampler.

13.1.1 Storage Image

A *storage image* (VK_DESCRIPTOR_TYPE_STORAGE_IMAGE) is a descriptor type that is used for load, store, and atomic operations on image memory from within shaders bound to pipelines.

Loads from storage images do not use samplers and are unfiltered and do not support coordinate wrapping or clamping. Loads are supported in all shader stages for image formats which report support for the VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT feature bit via vkGetPhysicalDeviceFormatProperties.

Stores to storage images are supported in compute shaders for image formats which report support for the VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT feature.

Storage images also support atomic operations in compute shaders for image formats which report support for the VK_FORMAT_FEATURE_STORAGE_IMAGE_ATOMIC_BIT feature.

Load and store operations on storage images can only be done on images in VK_IMAGE_LAYOUT_GENERAL layout.

When the fragmentStoresAndAtomics feature is enabled, stores and atomic operations are also supported for storage images in fragment shaders with the same set of image formats as supported in compute shaders. When the vertexPipelineStoresAndAtomics feature is enabled, stores and atomic operations are also supported in vertex, tessellation, and geometry shaders with the same set of image formats as supported in compute shaders.

Storage image declarations must specify the image format in the shader if the variable is used for atomic operations.

If the shaderStorageImageReadWithoutFormat feature is not enabled, storage image declarations must specify the image format in the shader if the variable is used for load operations.

If the shaderStorageImageWriteWithoutFormat feature is not enabled, storage image declarations must specify the image format in the shader if the variable is used for store operations.

Storage images are declared in GLSL shader source using uniform **image** variables of the appropriate dimensionality as well as a format layout qualifier (if necessary):

GLSL example

```
layout (set=m, binding=n, r32f) uniform image2D myStorageImage;
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
    ...
    OpName %9 "myStorageImage"
    OpDecorate %9 DescriptorSet m
    OpDecorate %9 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeImage %6 2D 0 0 0 2 R32f
%8 = OpTypePointer UniformConstant %7
%9 = OpVariable %8 UniformConstant
    ...
```

13.1.2 Sampler

A *sampler* (VK_DESCRIPTOR_TYPE_SAMPLER) represents a set of parameters which control address calculations, filtering behavior, and other properties, that can be used to perform filtered loads from *sampled images* (see Sampled Image).

Samplers are declared in GLSL shader source using uniform **sampler** variables, where the sampler type has no associated texture dimensionality:

GLSL Example

```
layout (set=m, binding=n) uniform sampler mySampler;
```

SPIR-V Example

```
%1 = OpExtInstImport "GLSL.std.450"
    ...
    OpName %8 "mySampler"
    OpDecorate %8 DescriptorSet m
    OpDecorate %8 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeSampler
%7 = OpTypePointer UniformConstant %6
%8 = OpVariable %7 UniformConstant
    ...
```

13.1.3 Sampled Image

A *sampled image* (VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE) can be used (usually in conjunction with a sampler) to retrieve sampled image data. Shaders use a sampled image handle and a sampler handle to sample data, where the image handle generally defines the shape and format of the memory and the sampler generally defines how coordinate addressing is performed. The same sampler can be used to sample from multiple images, and it is possible to sample from the same sampled image with multiple samplers, each containing a different set of sampling parameters.

Sampled images are declared in GLSL shader source using uniform **texture** variables of the appropriate dimensionality:

GLSL example

```
layout (set=m, binding=n) uniform texture2D mySampledImage;
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
    ...
    OpName %9 "mySampledImage"
    OpDecorate %9 DescriptorSet m
    OpDecorate %9 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeImage %6 2D 0 0 0 1 Unknown
%8 = OpTypePointer UniformConstant %7
%9 = OpVariable %8 UniformConstant
    ...
```

13.1.4 Combined Image Sampler

A *combined image sampler* (VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER) represents a sampled image along with a set of sampling parameters. It is logically considered a sampled image and a sampler bound together.



Note

On some implementations, it may be more efficient to sample from an image using a combination of sampler and sampled image that are stored together in the descriptor set in a combined descriptor.

Combined image samplers are declared in GLSL shader source using uniform **sampler** variables of the appropriate dimensionality:

GLSL example

```
layout (set=m, binding=n) uniform sampler2D myCombinedImageSampler;
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
    ...
    OpName %10 "myCombinedImageSampler"
```

```
OpDecorate %10 DescriptorSet m
   OpDecorate %10 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeImage %6 2D 0 0 0 1 Unknown
%8 = OpTypeSampledImage %7
%9 = OpTypePointer UniformConstant %8
%10 = OpVariable %9 UniformConstant
....
```

VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER descriptor set entries can also be accessed via separate sampler and sampled image shader variables. Such variables refer exclusively to the corresponding half of the descriptor, and can be combined in the shader with samplers or sampled images that can come from the same descriptor or from other combined or separate descriptor types. There are no additional restrictions on how a separate sampler or sampled image variable is used due to it originating from a combined descriptor.

13.1.5 Uniform Texel Buffer

A uniform texel buffer (VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER) represents a tightly packed array of homogeneous formatted data that is stored in a buffer and is made accessible to shaders. Uniform texel buffers are read-only.

Uniform texel buffers are declared in GLSL shader source using uniform **samplerBuffer** variables:

GLSL example

```
layout (set=m, binding=n) uniform samplerBuffer myUniformTexelBuffer;
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
...
OpName %10 "myUniformTexelBuffer"
OpDecorate %10 DescriptorSet m
OpDecorate %10 Binding n

%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeImage %6 Buffer 0 0 0 1 Unknown
%8 = OpTypeSampledImage %7
%9 = OpTypePointer UniformConstant %8
%10 = OpVariable %9 UniformConstant
...
```

13.1.6 Storage Texel Buffer

A *storage texel buffer* (VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER) represents a tightly packed array of homogeneous formatted data that is stored in a buffer and is made accessible to shaders. Storage texel buffers differ from uniform texel buffers in that they support stores and atomic operations in shaders, may support a different maximum length, and may have different performance characteristics.

Storage texel buffers are declared in GLSL shader source using uniform **imageBuffer** variables:

GLSL example

```
layout (set=m, binding=n, r32f) uniform imageBuffer myStorageTexelBuffer;
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
...
    OpName %9 "myStorageTexelBuffer"
    OpDecorate %9 DescriptorSet m
    OpDecorate %9 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeFloat 32
%7 = OpTypeImage %6 Buffer 0 0 0 2 R32f
%8 = OpTypePointer UniformConstant %7
%9 = OpVariable %8 UniformConstant
...
```

13.1.7 Uniform Buffer

A *uniform buffer* (VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER) is a region of structured storage that is made accessible for read-only access to shaders. It is typically used to store medium sized arrays of constants such as shader parameters, matrices and other related data.

Uniform buffers are declared in GLSL shader source using the uniform storage qualifier and block syntax:

GLSL example

```
layout (set=m, binding=n) uniform myUniformBuffer
{
    vec4 myElement[32];
};
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
     OpName %11 "myUniformBuffer"
     OpMemberName %11 0 "myElement"
     OpName %13 ""
     OpDecorate %10 ArrayStride 16
      OpMemberDecorate %11 0 Offset 0
     OpDecorate %11 Block
     OpDecorate %13 DescriptorSet m
     OpDecorate %13 Binding n
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeFloat 32
%7 = OpTypeVector %6 4
%8 = OpTypeInt 32 0
%9 = OpConstant %8 32
%10 = OpTypeArray %7 %9
%11 = OpTypeStruct %10
%12 = OpTypePointer Uniform %11
```

```
%13 = OpVariable %12 Uniform
...
```

13.1.8 Storage Buffer

A storage buffer (VK_DESCRIPTOR_TYPE_STORAGE_BUFFER) is a region of structured storage that supports both read and write access for shaders. In addition to general read and write operations, some members of storage buffers can be used as the target of atomic operations. In general, atomic operations are only supported on members that have unsigned integer formats.

Storage buffers are declared in GLSL shader source using buffer storage qualifier and block syntax:

GLSL example

```
layout (set=m, binding=n) buffer myStorageBuffer
{
    vec4 myElement[];
};
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
      OpName %9 "myStorageBuffer"
      OpMemberName %9 0 "myElement"
      OpName %11 ""
      OpDecorate %8 ArrayStride 16
      OpMemberDecorate %9 0 Offset 0
      OpDecorate %9 BufferBlock
      OpDecorate %11 DescriptorSet m
      OpDecorate %11 Binding n
 %2 = OpTypeVoid
 %3 = OpTypeFunction %2
 %6 = OpTypeFloat 32
 %7 = OpTypeVector %6 4
 %8 = OpTypeRuntimeArray %7
%9 = OpTypeStruct %8
%10 = OpTypePointer Uniform %9
%11 = OpVariable %10 Uniform
```

13.1.9 Dynamic Uniform Buffer

A dynamic uniform buffer (VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC) differs from a uniform buffer only in how its address and length are specified. Uniform buffers bind a buffer address and length that is specified in the descriptor set update by a buffer handle, offset and range (see Descriptor Set Updates). With dynamic uniform buffers the buffer handle, offset and range specified in the descriptor set define the base address and length. The dynamic offset which is relative to this base address is taken from the <code>pDynamicOffsets</code> parameter to

vkCmdBindDescriptorSets (see Descriptor Set Binding). The address used for a dynamic uniform buffer is the sum of the buffer base address and the relative offset. The length is unmodified and remains the range as specified in the descriptor update. The shader syntax is identical for uniform buffers and dynamic uniform buffers.

13.1.10 Dynamic Storage Buffer

A dynamic storage buffer (VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC) differs from a storage buffer only in how its address and length are specified. The difference is identical to the difference between uniform buffers and dynamic uniform buffers (see Dynamic Uniform Buffer). The shader syntax is identical for storage buffers and dynamic storage buffers.

13.1.11 Input Attachment

An *input attachment* (VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT) is an image view that can be used for pixel local load operations from within fragment shaders bound to pipelines. Loads from input attachments are unfiltered. All image formats that are supported for color attachments (VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT) or depth/stencil attachments (VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT) for a given image tiling mode are also supported for input attachments.

In the shader, input attachments must be decorated with their input attachment index in addition to descriptor set and binding numbers.

GLSL example

```
layout (input_attachment_index=i, set=m, binding=n) uniform subpassInput \leftrightarrow myInputAttachment;
```

SPIR-V example

13.2 Descriptor Sets

Descriptors are grouped together into descriptor set objects. A descriptor set object is an opaque object that contains storage for a set of descriptors, where the types and number of descriptors is defined by a descriptor set layout. The layout object may be used to define the association of each descriptor binding with memory or other hardware resources. The layout is used both for determining the resources that need to be associated with the descriptor set, and determining the interface between shader stages and shader resources.

13.2.1 Descriptor Set Layout

A descriptor set layout object is defined by an array of zero or more descriptor bindings. Each individual descriptor binding is specified by a descriptor type, a count (array size) of the number of descriptors in the binding, a set of shader stages that can access the binding, and (if using immutable samplers) an array of sampler descriptors.

Descriptor set layout objects are represented by VkDescriptorSetLayout handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE (VkDescriptorSetLayout)
```

To create descriptor set layout objects, call:

- device is the logical device that creates the descriptor set layout.
- pCreateInfo is a pointer to an instance of the VkDescriptorSetLayoutCreateInfo structure specifying the state of the descriptor set layout object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pSetLayout points to a VkDescriptorSetLayout handle in which the resulting descriptor set layout object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkDescriptorSetLayoutCreateInfo structure
- $\bullet \ \ If \ \textit{pAllocator} \ \textbf{is not} \ \texttt{NULL}, \ \textit{pAllocator} \ \textbf{must} \ \textbf{be} \ \textbf{a pointer to} \ \textbf{a valid} \ \texttt{VkAllocationCallbacks} \ \textbf{structure}$
- pSetLayout must be a pointer to a VkDescriptorSetLayout handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

Information about the descriptor set layout is passed in an instance of the VkDescriptorSetLayoutCreateInfo structure:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- bindingCount is the number of elements in pBindings.
- pBindings is a pointer to an array of VkDescriptorSetLayoutBinding structures.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO
- pNext must be NULL
- flags must be 0
- If bindingCount is not 0, pBindings must be a pointer to an array of bindingCount valid VkDescriptorSetLayoutBinding structures

The VkDescriptorSetLayoutBinding structure is defined as:

- binding is the binding number of this entry and corresponds to a resource of the same binding number in the shader stages.
- descriptorType is a VkDescriptorType specifying which type of resource descriptors are used for this binding.
- descriptorCount is the number of descriptors contained in the binding, accessed in a shader as an array. If descriptorCount is zero this binding entry is reserved and the resource must not be accessed from any stage via this binding within any pipeline using the set layout.

- stageFlags member is a bitmask of VkShaderStageFlagBits specifying which pipeline shader stages can access a resource for this binding. VK_SHADER_STAGE_ALL is a shorthand specifying that all defined shader stages, including any additional stages defined by extensions, can access the resource.
 - If a shader stage is not included in stageFlags, then a resource must not be accessed from that stage via this binding within any pipeline using the set layout. There are no limitations on what combinations of stages can be used by a descriptor binding, and in particular a binding can be used by both graphics stages and the compute stage.
- pImmutableSamplers affects initialization of samplers. If descriptorType specifies a VK_DESCRIPTOR_TYPE_SAMPLER or VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER type descriptor, then pImmutableSamplers can be used to initialize a set of immutable samplers. Immutable samplers are permanently bound into the set layout; later binding a sampler into an immutable sampler slot in a descriptor set is not allowed. If pImmutableSamplers is not NULL, then it is considered to be a pointer to an array of sampler handles that will be consumed by the set layout and used for the corresponding binding. If pImmutableSamplers is NULL, then the sampler slots are dynamic and sampler handles must be bound into descriptor sets using this layout. If descriptorType is not one of these descriptor types, then pImmutableSamplers is ignored.

The above layout definition allows the descriptor bindings to be specified sparsely such that not all binding numbers between 0 and the maximum binding number need to be specified in the <code>pBindings</code> array. However, all binding numbers between 0 and the maximum binding number may consume memory in the descriptor set layout even if not all descriptor bindings are used, though it should not consume additional memory from the descriptor pool.



Note

The maximum binding number specified should be as compact as possible to avoid wasted memory.

Valid Usage

- descriptorType must be a valid VkDescriptorType value
- If descriptorType is VK_DESCRIPTOR_TYPE_SAMPLER or VK_DESCRIPTOR_TYPE_COMBINED_ IMAGE_SAMPLER, and descriptorCount is not 0 and pImmutableSamplers is not NULL, pImmutableSamplers must be a pointer to an array of descriptorCount valid VkSampler handles
- ullet If descriptorCount is not 0, stageFlags must be a valid combination of VkShaderStageFlagBits values

The following examples show a shader snippet using two descriptor sets, and application code that creates corresponding descriptor set layouts.

GLSL example

```
//
// binding to a single sampled image descriptor in set 0
//
layout (set=0, binding=0) uniform texture2D mySampledImage;
```

```
//
// binding to an array of sampled image descriptors in set 0
//
layout (set=0, binding=1) uniform texture2D myArrayOfSampledImages[12];

//
// binding to a single uniform buffer descriptor in set 1
//
layout (set=1, binding=0) uniform myUniformBuffer
{
    vec4 myElement[32];
};
```

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
      OpName %9 "mySampledImage"
      OpName %14 "myArrayOfSampledImages"
      OpName %18 "myUniformBuffer"
      OpMemberName %18 0 "myElement"
      OpName %20 ""
      OpDecorate %9 DescriptorSet 0
      OpDecorate %9 Binding 0
      OpDecorate %14 DescriptorSet 0
      OpDecorate %14 Binding 1
      OpDecorate %17 ArrayStride 16
      OpMemberDecorate %18 0 Offset 0
      OpDecorate %18 Block
      OpDecorate %20 DescriptorSet 1
      OpDecorate %20 Binding 0
 %2 = OpTypeVoid
 %3 = OpTypeFunction %2
 %6 = OpTypeFloat 32
 %7 = OpTypeImage %6 2D 0 0 1 Unknown
%8 = OpTypePointer UniformConstant %7
%9 = OpVariable %8 UniformConstant
%10 = OpTypeInt 32 0
%11 = OpConstant %10 12
%12 = OpTypeArray %7 %11
%13 = OpTypePointer UniformConstant %12
%14 = OpVariable %13 UniformConstant
%15 = OpTypeVector %6 4
%16 = OpConstant %10 32
%17 = OpTypeArray %15 %16
%18 = OpTypeStruct %17
%19 = OpTypePointer Uniform %18
%20 = OpVariable %19 Uniform
```

API example

```
VkResult myResult;
const VkDescriptorSetLayoutBinding myDescriptorSetLayoutBinding[] =
{
```

```
// binding to a single image descriptor
    {
                                              // binding
       VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE,
                                              // descriptorType
                                               // descriptorCount
       VK_SHADER_STAGE_FRAGMENT_BIT,
                                               // stageFlags
       NULL
                                               // pImmutableSamplers
    },
    // binding to an array of image descriptors
                                              // binding
       1,
       VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE,
                                              // descriptorType
                                              // descriptorCount
                                              // stageFlags
       VK_SHADER_STAGE_FRAGMENT_BIT,
       NULL
                                               // pImmutableSamplers
    },
    // binding to a single uniform buffer descriptor
        Ο,
                                               // binding
       VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER,
                                             // descriptorType
                                               // descriptorCount
       VK_SHADER_STAGE_FRAGMENT_BIT,
                                              // stageFlags
       NULL
                                               // pImmutableSamplers
    }
};
const VkDescriptorSetLayoutCreateInfo myDescriptorSetLayoutCreateInfo[] =
    // Create info for first descriptor set with two descriptor bindings
       VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
                                                                // sType
       NULL,
                                                                // pNext
        Ο,
                                                                // flags
                                                                // bindingCount
        &myDescriptorSetLayoutBinding[0]
                                                                // pBindings
    },
    // Create info for second descriptor set with one descriptor binding
       VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
                                                               // sType
       NULL,
                                                                // pNext
        Ο,
                                                                // flags
                                                               // bindingCount
       1,
        &myDescriptorSetLayoutBinding[2]
                                                                // pBindings
};
VkDescriptorSetLayout myDescriptorSetLayout[2];
11
// Create first descriptor set layout
myResult = vkCreateDescriptorSetLayout(
   myDevice,
  &myDescriptorSetLayoutCreateInfo[0],
```

```
NULL,
    &myDescriptorSetLayout[0]);

//

// Create second descriptor set layout
//

myResult = vkCreateDescriptorSetLayout(
    myDevice,
    &myDescriptorSetLayoutCreateInfo[1],
    NULL,
    &myDescriptorSetLayout[1]);
```

To destroy a descriptor set layout, call:

- device is the logical device that destroys the descriptor set layout.
- descriptorSetLayout is the descriptor set layout to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If descriptorSetLayout is not VK_NULL_HANDLE, descriptorSetLayout must be a valid VkDescriptorSetLayout handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If descriptorSetLayout is a valid handle, it must have been created, allocated, or retrieved from device
- If VkAllocationCallbacks were provided when descriptorSetLayout was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when descriptorSetLayout was created, pAllocator must be NULL

Host Synchronization

• Host access to descriptorSetLayout must be externally synchronized

13.2.2 Pipeline Layouts

Access to descriptor sets from a pipeline is accomplished through a *pipeline layout*. Zero or more descriptor set layouts and zero or more push constant ranges are combined to form a pipeline layout object which describes the complete set of resources that can be accessed by a pipeline. The pipeline layout represents a sequence of descriptor sets with each having a specific layout. This sequence of layouts is used to determine the interface between shader stages and shader resources. Each pipeline is created using a pipeline layout.

Pipeline layout objects are represented by VkPipelineLayout handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkPipelineLayout)
```

To create a pipeline layout, call:

- device is the logical device that creates the pipeline layout.
- pCreateInfo is a pointer to an instance of the VkPipelineLayoutCreateInfo structure specifying the state of the pipeline layout object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pPipelineLayout points to a VkPipelineLayout handle in which the resulting pipeline layout object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkPipelineLayoutCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pPipelineLayout must be a pointer to a VkPipelineLayout handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkPipelineLayoutCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- setLayoutCount is the number of descriptor sets included in the pipeline layout.
- pSetLayouts is a pointer to an array of VkDescriptorSetLayout objects.
- pushConstantRangeCount is the number of push constant ranges included in the pipeline layout.
- pPushConstantRanges is a pointer to an array of VkPushConstantRange structures defining a set of push constant ranges for use in a single pipeline layout. In addition to descriptor set layouts, a pipeline layout also describes how many push constants can be accessed by each stage of the pipeline.



Note

Push constants represent a high speed path to modify constant data in pipelines that is expected to outperform memory-backed resource updates.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO
- pNext must be NULL
- flags must be 0

- If setLayoutCount is not 0, pSetLayouts must be a pointer to an array of setLayoutCount valid VkDescriptorSetLayout handles
- If pushConstantRangeCount is not 0, pPushConstantRanges must be a pointer to an array of pushConstantRangeCount valid VkPushConstantRange structures
- setLayoutCount must be less than or equal to VkPhysicalDeviceLimits::maxBoundDescriptorSets
- The total number of descriptors of the type VK_DESCRIPTOR_TYPE_SAMPLER and VK_DESCRIPTOR_
 TYPE_COMBINED_IMAGE_SAMPLER accessible to any given shader stage across all elements of pSetLayouts
 must be less than or equal to VkPhysicalDeviceLimits::maxPerStageDescriptorSamplers
- The total number of descriptors of the type VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER and VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC accessible to any given shader stage across all elements of pSetLayouts must be less than or equal to VkPhysicalDeviceLimits::maxPerStageDescriptorUniformBuffers
- The total number of descriptors of the type VK_DESCRIPTOR_TYPE_STORAGE_BUFFER and VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC accessible to any given shader stage across all elements of pSetLayouts must be less than or equal to VkPhysicalDeviceLimits::maxPerStageDescriptorStorageBuffers
- The total number of descriptors of the type VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, and VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER accessible to any given shader stage across all elements of pSetLayouts must be less than or equal to VkPhysicalDeviceLimits::maxPerStageDescriptorSampledImages
- The total number of descriptors of the type VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, and VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER accessible to any given shader stage across all elements of pSetLayouts must be less than or equal to VkPhysicalDeviceLimits::maxPerStageDescriptorStorageImages

The VkPushConstantRange structure is defined as:

- stageFlags is a set of stage flags describing the shader stages that will access a range of push constants. If a particular stage is not included in the range, then accessing members of that range of push constants from the corresponding shader stage will result in undefined data being read.
- offset and size are the start offset and size, respectively, consumed by the range. Both offset and size are in units of bytes and must be a multiple of 4. The layout of the push constant variables is specified in the shader.

- stageFlags must be a valid combination of VkShaderStageFlagBits values
- stageFlags must not be 0
- offset must be less than VkPhysicalDeviceLimits::maxPushConstantsSize
- size must be greater than 0
- size must be a multiple of 4
- size must be less than or equal to VkPhysicalDeviceLimits::maxPushConstantsSize minus offset

Once created, pipeline layouts are used as part of pipeline creation (see Pipelines), as part of binding descriptor sets (see Descriptor Set Binding), and as part of setting push constants (see Push Constant Updates). Pipeline creation accepts a pipeline layout as input, and the layout may be used to map (set, binding, arrayElement) tuples to hardware resources or memory locations within a descriptor set. The assignment of hardware resources depends only on the bindings defined in the descriptor sets that comprise the pipeline layout, and not on any shader source.

All resource variables statically used in all shaders in a pipeline must be declared with a (set,binding,arrayElement) that exists in the corresponding descriptor set layout and is of an appropriate descriptor type and includes the set of shader stages it is used by in stageFlags. The pipeline layout can include entries that are not used by a particular pipeline, or that are dead-code eliminated from any of the shaders. The pipeline layout allows the application to provide a consistent set of bindings across multiple pipeline compiles, which enables those pipelines to be compiled in a way that the implementation may cheaply switch pipelines without reprogramming the bindings.

Similarly, the push constant block declared in each shader (if present) must only place variables at offsets that are each included in a push constant range with <code>stageFlags</code> including the bit corresponding to the shader stage that uses it. The pipeline layout can include ranges or portions of ranges that are not used by a particular pipeline, or for which the variables have been dead-code eliminated from any of the shaders.

There is a limit on the total number of resources of each type that can be included in bindings in all descriptor set layouts in a pipeline layout as shown in Pipeline Layout Resource Limits. The "Total Resources Available" column gives the limit on the number of each type of resource that can be included in bindings in all descriptor sets in the pipeline layout. Some resource types count against multiple limits. Additionally, there are limits on the total number of each type of resource that can be used in any pipeline stage as described in Shader Resource Limits.

Table 13.1: Pipeline Layout Resource Limits

Total Resources Available	Resource Types	
maxDescriptorSetSamplers	sampler	
maxDescriptorsetsamplers	combined image sampler	
	sampled image	
maxDescriptorSetSampledImages	combined image sampler	
	uniform texel buffer	
maxDescriptorSetStorageImages	storage image	
maxDescriptorsetstoragemages	storage texel buffer	
max Dagarintar Sat Uniform Buffors	uniform buffer	
maxDescriptorSetUniformBuffers	uniform buffer dynamic	
maxDescriptorSetUniformBuffersDynamic	uniform buffer dynamic	
maxDescriptorSetStorageBuffers	storage buffer	

Table 13.1: (continued)

Total Resources Available	Resource Types
	storage buffer dynamic
maxDescriptorSetStorageBuffersDynamic	storage buffer dynamic
maxDescriptorSetInputAttachments	input attachment

To destroy a pipeline layout, call:

- device is the logical device that destroys the pipeline layout.
- pipelineLayout is the pipeline layout to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

Valid Usage

- device must be a valid VkDevice handle
- If pipelineLayout is not VK_NULL_HANDLE, pipelineLayout must be a valid VkPipelineLayout handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If pipelineLayout is a valid handle, it must have been created, allocated, or retrieved from device
- If VkAllocationCallbacks were provided when <code>pipelineLayout</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when pipelineLayout was created, pAllocator must be NULL

Host Synchronization

• Host access to pipelineLayout must be externally synchronized

13.2.2.1 Pipeline Layout Compatibility

Two pipeline layouts are defined to be "compatible for push constants" if they were created with identical push constant ranges. Two pipeline layouts are defined to be "compatible for set N" if they were created with matching (the same, or identically defined) descriptor set layouts for sets zero through N, and if they were created with identical push constant ranges.

When binding a descriptor set (see Descriptor Set Binding) to set number N, if the previously bound descriptor sets for sets zero through N-1 were all bound using compatible pipeline layouts, then performing this binding does not disturb any of the lower numbered sets. If, additionally, the previous bound descriptor set for set N was bound using a pipeline layout compatible for set N, then the bindings in sets numbered greater than N are also not disturbed.

Similarly, when binding a pipeline, the pipeline can correctly access any previously bound descriptor sets which were bound with compatible pipeline layouts, as long as all lower numbered sets were also bound with compatible layouts.

Layout compatibility means that descriptor sets can be bound to a command buffer for use by any pipeline created with a compatible pipeline layout, and without having bound a particular pipeline first. It also means that descriptor sets can remain valid across a pipeline change, and the same resources will be accessible to the newly bound pipeline.

Implementor's Note

A consequence of layout compatibility is that when the implementation compiles a pipeline layout and assigns hardware units to resources, the mechanism to assign hardware units for set N should only be a function of sets [0..N].



Note

Place the least frequently changing descriptor sets near the start of the pipeline layout, and place the descriptor sets representing the most frequently changing resources near the end. When pipelines are switched, only the descriptor set bindings that have been invalidated will need to be updated and the remainder of the descriptor set bindings will remain in place.

The maximum number of descriptor sets that can be bound to a pipeline layout is queried from physical device properties (see maxBoundDescriptorSets in Limits).

API example

```
} ;
const VkPipelineLayoutCreateInfo createInfo =
    VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO,
                                                     // sType
   NULL,
                                                      // pNext
                                                      // flags
    Ο,
    2,
                                                      // setLayoutCount
    layouts,
                                                      // pSetLayouts
    2,
                                                      // pushConstantRangeCount
                                                      // pPushConstantRanges
    ranges
};
VkPipelineLayout myPipelineLayout;
myResult = vkCreatePipelineLayout(
    myDevice,
    &createInfo,
    NULL,
    &myPipelineLayout);
```

13.2.3 Allocation of Descriptor Sets

A *descriptor pool* maintains a pool of descriptors, from which descriptor sets are allocated. Descriptor pools are externally synchronized, meaning that the application must not allocate and/or free descriptor sets from the same pool in multiple threads simultaneously.

Descriptor pools are represented by VkDescriptorPool handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkDescriptorPool)
```

To create a descriptor pool object, call:

- device is the logical device that creates the descriptor pool.
- pCreateInfo is a pointer to an instance of the VkDescriptorPoolCreateInfo structure specifying the state of the descriptor pool object.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pDescriptorPool points to a VkDescriptorPool handle in which the resulting descriptor pool object is returned.

pAllocator controls host memory allocation as described in the Memory Allocation chapter.

The created descriptor pool is returned in pDescriptorPool.

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkDescriptorPoolCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pDescriptorPool must be a pointer to a VkDescriptorPool handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

 $Additional\ information\ about\ the\ pool\ is\ passed\ in\ an\ instance\ of\ the\ \verb|VkDescriptorPoolCreateInfo| structure:$

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags specifies certain supported operations on the pool. Bits which can be set include:

```
typedef enum VkDescriptorPoolCreateFlagBits {
    VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT = 0x00000001,
} VkDescriptorPoolCreateFlagBits;
```

If flags includes VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT, then descriptor sets can return their individual allocations to the pool, i.e. all of **vkAllocateDescriptorSets**,

vkFreeDescriptorSets, and **vkResetDescriptorPool** are allowed. Otherwise, descriptor sets allocated from the pool must not be individually freed back to the pool, i.e. only **vkAllocateDescriptorSets** and **vkResetDescriptorPool** are allowed.

- maxSets is the maximum number of descriptor sets that can be allocated from the pool.
- poolSizeCount is the number of elements in pPoolSizes.
- pPoolSizes is a pointer to an array of VkDescriptorPoolSize structures, each containing a descriptor type and number of descriptors of that type to be allocated in the pool.

If multiple VkDescriptorPoolSize structures appear in the *pPoolSizes* array then the pool will be created with enough storage for the total number of descriptors of each type.

Fragmentation of a descriptor pool is possible and may lead to descriptor set allocation failures. A failure due to fragmentation is defined as failing a descriptor set allocation despite the sum of all outstanding descriptor set allocations from the pool plus the requested allocation requiring no more than the total number of descriptors requested at pool creation. Implementations provide certain guarantees of when fragmentation must not cause allocation failure, as described below.

If a descriptor pool has not had any descriptor sets freed since it was created or most recently reset then fragmentation must not cause an allocation failure (note that this is always the case for a pool created without the VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT bit set). Additionally, if all sets allocated from the pool since it was created or most recently reset use the same number of descriptors (of each type) and the requested allocation also uses that same number of descriptors (of each type), then fragmentation must not cause an allocation failure.

If an allocation failure occurs due to fragmentation, an application can create an additional descriptor pool to perform further descriptor set allocations.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DESCRIPTOR_POOL_CREATE_INFO
- pNext must be NULL
- flags must be a valid combination of VkDescriptorPoolCreateFlagBits values
- pPoolSizes must be a pointer to an array of poolSizeCount valid VkDescriptorPoolSize structures
- poolSizeCount must be greater than 0
- maxSets must be greater than 0

The VkDescriptorPoolSize structure is defined as:

- type is the type of descriptor.
- descriptorCount is the number of descriptors of that type to allocate.

- type must be a valid VkDescriptorType value
- descriptorCount must be greater than 0

To destroy a descriptor pool, call:

- device is the logical device that destroys the descriptor pool.
- descriptorPool is the descriptor pool to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

When a pool is destroyed, all descriptor sets allocated from the pool are implicitly freed and become invalid. Descriptor sets allocated from a given pool do not need to be freed before destroying that descriptor pool.

Valid Usage

- device must be a valid VkDevice handle
- If descriptorPool is not VK_NULL_HANDLE, descriptorPool must be a valid VkDescriptorPool handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If descriptorPool is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to <code>descriptorPool</code> (via any allocated descriptor sets) must have completed execution
- If VkAllocationCallbacks were provided when <code>descriptorPool</code> was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when descriptorPool was created, pAllocator must be NULL

Host Synchronization

• Host access to descriptorPool must be externally synchronized

Descriptor sets are allocated from descriptor pool objects, and are represented by VkDescriptorSet handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkDescriptorSet)
```

To allocate descriptor sets from a descriptor pool, call:

- device is the logical device that owns the descriptor pool.
- pAllocateInfo is a pointer to an instance of the VkDescriptorSetAllocateInfo structure describing parameters of the allocation.
- pDescriptorSets is a pointer to an array of VkDescriptorSet handles in which the resulting descriptor set objects are returned. The array must be at least the length specified by the descriptorSetCount member of pAllocateInfo.

The allocated descriptor sets are returned in pDescriptorSets.

When a descriptor set is allocated, the initial state is largely uninitialized and all descriptors are undefined. However, the descriptor set can be bound in a command buffer without causing errors or exceptions. All entries that are statically used by a pipeline in a drawing or dispatching command must have been populated before the descriptor set is bound for use by that command. Entries that are not statically used by a pipeline can have uninitialized descriptors or descriptors of resources that have been destroyed, and executing a draw or dispatch with such a descriptor set bound does not cause undefined behavior. This means applications need not populate unused entries with dummy descriptors.

If an allocation fails due to fragmentation, an indeterminate error is returned with an unspecified error code. Any returned error other than VK_ERROR_FRAGMENTED_POOL does not imply its usual meaning: applications should assume that the allocation failed due to fragmentation, and create a new descriptor pool.

Note



Applications should check for a negative return value when allocating new descriptor sets, assume that any error effectively means VK_ERROR_FRAGMENTED_POOL, and try to create a new descriptor pool. If VK_ERROR_FRAGMENTED_POOL is the actual return value, it adds certainty to that decision.

The reason for this is that VK_ERROR_FRAGMENTED_POOL was only added in a later revision of the 1.0 specification, and so drivers may return other errors if they were written against earlier revisions. To ensure full compatibility with earlier patch revisions, these other errors are allowed.

- device must be a valid VkDevice handle
- pAllocateInfo must be a pointer to a valid VkDescriptorSetAllocateInfo structure
- pDescriptorSets must be a pointer to an array of pAllocateInfo→descriptorSetCount VkDescriptorSet handles

Host Synchronization

• Host access to pAllocateInfo→descriptorPool must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_FRAGMENTED_POOL

The VkDescriptorSetAllocateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- descriptorPool is the pool which the sets will be allocated from.

- descriptorSetCount determines the number of descriptor sets to be allocated from the pool.
- pSetLayouts is an array of descriptor set layouts, with each member specifying how the corresponding descriptor set is allocated.

- sType must be VK_STRUCTURE_TYPE_DESCRIPTOR_SET_ALLOCATE_INFO
- pNext must be NULL
- descriptorPool must be a valid VkDescriptorPool handle
- pSetLayouts must be a pointer to an array of descriptorSetCount valid VkDescriptorSetLayout handles
- descriptorSetCount must be greater than 0
- Both of descriptorPool, and the elements of pSetLayouts must have been created, allocated, or retrieved from the same VkDevice
- descriptorSetCount must not be greater than the number of sets that are currently available for allocation in descriptorPool
- descriptorPool must have enough free descriptor capacity remaining to allocate the descriptor sets of the specified layouts

To free allocated descriptor sets, call:

- device is the logical device that owns the descriptor pool.
- descriptorPool is the descriptor pool from which the descriptor sets were allocated.
- descriptorSetCount is the number of elements in the pDescriptorSets array.
- pDescriptorSets is an array of handles to VkDescriptorSet objects.

After a successful call to **vkFreeDescriptorSets**, all descriptor sets in *pDescriptorSets* are invalid.

- device must be a valid VkDevice handle
- descriptorPool must be a valid VkDescriptorPool handle
- descriptorSetCount must be greater than 0
- descriptorPool must have been created, allocated, or retrieved from device
- Each element of pDescriptorSets that is a valid handle must have been created, allocated, or retrieved from descriptorPool
- All submitted commands that refer to any element of pDescriptorSets must have completed execution
- pDescriptorSets must be a pointer to an array of descriptorSetCount VkDescriptorSet handles, each element of which must either be a valid handle or VK_NULL_HANDLE
- Each valid handle in pDescriptorSets must have been allocated from descriptorPool
- descriptorPool must have been created with the VK_DESCRIPTOR_POOL_CREATE_FREE_ DESCRIPTOR_SET_BIT flag

Host Synchronization

- Host access to descriptorPool must be externally synchronized
- Host access to each member of pDescriptorSets must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

To return all descriptor sets allocated from a given pool to the pool, rather than freeing individual descriptor sets, call:

- device is the logical device that owns the descriptor pool.
- descriptorPool is the descriptor pool to be reset.
- flags is reserved for future use.

Resetting a descriptor pool recycles all of the resources from all of the descriptor sets allocated from the descriptor pool back to the descriptor pool, and the descriptor sets are implicitly freed.

Valid Usage

- device must be a valid VkDevice handle
- descriptorPool must be a valid VkDescriptorPool handle
- flags must be 0
- descriptorPool must have been created, allocated, or retrieved from device
- All uses of descriptorPool (via any allocated descriptor sets) must have completed execution

Host Synchronization

- Host access to descriptorPool must be externally synchronized
- Host access to any VkDescriptorSet objects allocated from <code>descriptorPool</code> must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

13.2.4 Descriptor Set Updates

Once allocated, descriptor sets can be updated with a combination of write and copy operations. To update descriptor sets, call:

- device is the logical device that updates the descriptor sets.
- descriptorWriteCount is the number of elements in the pDescriptorWrites array.
- pDescriptorWrites is a pointer to an array of VkWriteDescriptorSet structures describing the descriptor sets to write to.
- descriptorCopyCount is the number of elements in the pDescriptorCopies array.
- pDescriptorCopies is a pointer to an array of VkCopyDescriptorSet structures describing the descriptor sets to copy between.

The operations described by pDescriptorWrites are performed first, followed by the operations described by pDescriptorCopies. Within each array, the operations are performed in the order they appear in the array.

Each element in the pDescriptorWrites array describes an operation updating the descriptor set using descriptors for resources specified in the structure.

Each element in the pDescriptorCopies array is a VkCopyDescriptorSet structure describing an operation copying descriptors between sets.

Valid Usage

- device must be a valid VkDevice handle
- If descriptorWriteCount is not 0, pDescriptorWrites must be a pointer to an array of descriptorWriteCount valid VkWriteDescriptorSet structures
- If descriptorCopyCount is not 0, pDescriptorCopies must be a pointer to an array of descriptorCopyCount valid VkCopyDescriptorSet structures

Host Synchronization

- Host access to pDescriptorWrites[].dstSet must be externally synchronized
- Host access to pDescriptorCopies[].dstSet must be externally synchronized

The VkWriteDescriptorSet structure is defined as:

```
typedef struct VkWriteDescriptorSet {
    VkStructureType
                                          sType;
    const void*
                                         pNext;
    VkDescriptorSet
                                         dstSet;
    uint32_t
                                         dstBinding;
    uint32_t
                                         dstArrayElement;
    uint32_t
                                         descriptorCount;
    VkDescriptorType
                                         descriptorType;
    const VkDescriptorImageInfo* pImageInfo;
const VkDescriptorBufferInfo* pBufferInfo;
                                         pTexelBufferView;
    const VkBufferView*
} VkWriteDescriptorSet;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- dstSet is the destination descriptor set to update.
- dstBinding is the descriptor binding within that set.
- dstArrayElement is the starting element in that array.
- descriptorCount is the number of descriptors to update (the number of elements in pImageInfo, pBufferInfo, or pTexelBufferView).
- descriptorType is a VkDescriptorType specifying the type of each descriptor in pImageInfo, pBufferInfo, or pTexelBufferView, as described below. It must be the same type as that specified in VkDescriptorSetLayoutBinding for dstSet at dstBinding. The type of the descriptor also controls which array the descriptors are taken from.
- pImageInfo points to an array of VkDescriptorImageInfo structures or is ignored, as described below.
- pBufferInfo points to an array of VkDescriptorBufferInfo structures or is ignored, as described below.
- pTexelBufferView points to an array of VkBufferView handles as described in the Buffer Views section or is ignored, as described below.

Only one of pImageInfo, pBufferInfo, or pTexelBufferView members is used according to the descriptor type specified in the descriptorType member of the containing VkWriteDescriptorSet structure, as specified below.

If the dstBinding has fewer than descriptorCount array elements remaining starting from dstArrayElement, then the remainder will be used to update the subsequent binding - dstBinding+1 starting at array element zero. This

behavior applies recursively, with the update affecting consecutive bindings as needed to update all <code>descriptorCount</code> descriptors. All consecutive bindings updated via a single <code>VkWriteDescriptorSet</code> structure must have identical <code>descriptorType</code> and <code>stageFlags</code>, and must all either use immutable samplers or must all not use immutable samplers.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_WRITE_DESCRIPTOR_SET
- pNext must be NULL
- dstSet must be a valid VkDescriptorSet handle
- descriptorType must be a valid VkDescriptorType value
- descriptorCount must be greater than 0
- Both of dstSet, and the elements of pTexelBufferView that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- dstBinding must be a valid binding point within dstSet
- descriptorType must match the type of dstBinding within dstSet
- The sum of dstArrayElement and descriptorCount must be less than or equal to the number of array elements in the descriptor set binding specified by dstBinding, and all applicable consecutive bindings, as described by consecutive binding updates
- If descriptorType is VK_DESCRIPTOR_TYPE_SAMPLER, VK_DESCRIPTOR_TYPE_COMBINED_
 IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, VK_DESCRIPTOR_TYPE_STORAGE_
 IMAGE, or VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT, pImageInfo must be a pointer to an array of descriptorCount valid VkDescriptorImageInfo structures
- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER or VK_DESCRIPTOR_ TYPE_STORAGE_TEXEL_BUFFER, pTexelBufferView must be a pointer to an array of descriptorCount valid VkBufferView handles
- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER, VK_DESCRIPTOR_TYPE_ STORAGE_BUFFER, VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC, or VK_DESCRIPTOR_ TYPE_STORAGE_BUFFER_DYNAMIC, pBufferInfo must be a pointer to an array of descriptorCount valid VkDescriptorBufferInfo structures
- If descriptorType is VK_DESCRIPTOR_TYPE_SAMPLER or VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, and dstSet was not allocated with a layout that included immutable samplers for dstBinding with descriptorType, the sampler member of any given element of pImageInfo must be a valid VkSampler object
- If descriptorType is VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, or VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT, the imageView and imageLayout members of any given element of pImageInfo must be a valid VkImageView and VkImageLayout, respectively

- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_ UNIFORM_BUFFER_DYNAMIC, the offset member of any given element of pBufferInfo must be a multiple of VkPhysicalDeviceLimits::minUniformBufferOffsetAlignment
- If descriptorType is VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_ STORAGE_BUFFER_DYNAMIC, the offset member of any given element of pBufferInfo must be a multiple of VkPhysicalDeviceLimits::minStorageBufferOffsetAlignment
- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_ UNIFORM_BUFFER_DYNAMIC, the buffer member of any given element of pBufferInfo must have been created with VK_BUFFER_USAGE_UNIFORM_BUFFER_BIT set
- If descriptorType is VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC, the buffer member of any given element of pBufferInfo must have been created with VK_BUFFER_USAGE_STORAGE_BUFFER_BIT set
- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_ UNIFORM_BUFFER_DYNAMIC, the range member of any given element of pBufferInfo, or the effective range if range is VK_WHOLE_SIZE, must be less than or equal to VkPhysicalDeviceLimits::maxUniformBufferRange
- If descriptorType is VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC, the range member of any given element of pBufferInfo, or the effective range if range is VK_WHOLE_SIZE, must be less than or equal to VkPhysicalDeviceLimits::maxStorageBufferRange
- If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER, the VkBuffer that any given element of pTexelBufferView was created from must have been created with VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT set
- If descriptorType is VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER, the VkBuffer that any given element of pTexelBufferView was created from must have been created with VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT set
- If descriptorType is VK_DESCRIPTOR_TYPE_STORAGE_IMAGE or VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT, the imageView member of any given element of pImageInfo must have been created with the identity swizzle

The type of descriptors in a descriptor set is specified by VkWriteDescriptorSet::descriptorType, which must be one of the values:

```
typedef enum VkDescriptorType {
   VK_DESCRIPTOR_TYPE_SAMPLER = 0,
   VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER = 1,
   VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE = 2,
   VK_DESCRIPTOR_TYPE_STORAGE_IMAGE = 3,
   VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER = 4,
   VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER = 5,
   VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER = 6,
   VK_DESCRIPTOR_TYPE_STORAGE_BUFFER = 7,
   VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC = 8,
   VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC = 9,
   VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT = 10,
} VkDescriptorType;
```

If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER, VK_DESCRIPTOR_TYPE_STORAGE_BUFFER, VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC, or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC, the elements of the VkWriteDescriptorSet::pBufferInfo array of VkDescriptorBufferInfo structures will be used to update the descriptors, and other arrays will be ignored.

If descriptorType is VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER, the VkWriteDescriptorSet::pTexelBufferView array will be used to update the descriptors, and other arrays will be ignored.

If descriptorType is VK_DESCRIPTOR_TYPE_SAMPLER, VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, or VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT, the elements of the VkWriteDescriptorSet::pImageInfo array of VkDescriptorImageInfo structures will be used to update the descriptors, and other arrays will be ignored.

The VkDescriptorBufferInfo structure is defined as:

- buffer is the buffer resource.
- offset is the offset in bytes from the start of buffer. Access to buffer memory via this descriptor uses addressing that is relative to this starting offset.
- range is the size in bytes that is used for this descriptor update, or VK_WHOLE_SIZE to use the range from offset to the end of the buffer.



Note

When using VK_WHOLE_SIZE , the effective range must not be larger than the maximum range for the descriptor type (maxUniformBufferRange or maxStorageBufferRange). This means that VK_WHOLE_SIZE is not typically useful in the common case where uniform buffer descriptors are suballocated from a buffer that is much larger than maxUniformBufferRange.

For VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC and VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC descriptor types, offset is the base offset from which the dynamic offset is applied and range is the static size used for all dynamic offsets.

Valid Usage

- buffer must be a valid VkBuffer handle
- offset must be less than the size of buffer
- If range is not equal to VK_WHOLE_SIZE, range must be greater than 0
- If range is not equal to VK_WHOLE_SIZE, range must be less than or equal to the size of buffer minus offset

The VkDescriptorImageInfo structure is defined as:

- sampler is a sampler handle, and is used in descriptor updates for types VK_DESCRIPTOR_TYPE_SAMPLER and VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER if the binding being updated does not use immutable samplers.
- *imageView* is an image view handle, and is used in descriptor updates for types VK_DESCRIPTOR_TYPE_ SAMPLED_IMAGE, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, VK_DESCRIPTOR_TYPE_COMBINED_ IMAGE_SAMPLER, and VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT.
- *imageLayout* is the layout that the image will be in at the time this descriptor is accessed. *imageLayout* is used in descriptor updates for types VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, and VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT.

Members of VkDescriptorImageInfo that are not used in an update (as described above) are ignored.

Valid Usage

• Both of imageView, and sampler that are valid handles must have been created, allocated, or retrieved from the same VkDevice

The VkCopyDescriptorSet structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- srcSet, srcBinding, and srcArrayElement are the source set, binding, and array element, respectively.
- dstSet, dstBinding, and dstArrayElement are the destination set, binding, and array element, respectively.

• descriptorCount is the number of descriptors to copy from the source to destination. If descriptorCount is greater than the number of remaining array elements in the source or destination binding, those affect consecutive bindings in a manner similar to VkWriteDescriptorSet above.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_COPY_DESCRIPTOR_SET
- pNext must be NULL
- srcSet must be a valid VkDescriptorSet handle
- dstSet must be a valid VkDescriptorSet handle
- Both of dstSet, and srcSet must have been created, allocated, or retrieved from the same VkDevice
- srcBinding must be a valid binding within srcSet
- The sum of <code>srcArrayElement</code> and <code>descriptorCount</code> must be less than or equal to the number of array elements in the descriptor set binding specified by <code>srcBinding</code>, and all applicable consecutive bindings, as described by consecutive binding updates
- dstBinding must be a valid binding within dstSet
- The sum of dstArrayElement and descriptorCount must be less than or equal to the number of array elements in the descriptor set binding specified by dstBinding, and all applicable consecutive bindings, as described by consecutive binding updates
- If srcSet is equal to dstSet, then the source and destination ranges of descriptors must not overlap, where the ranges may include array elements from consecutive bindings as described by consecutive binding updates

13.2.5 Descriptor Set Binding

To bind one or more descriptor sets to a command buffer, call:

```
void vkCmdBindDescriptorSets(
   VkCommandBuffer
                                                  commandBuffer,
    VkPipelineBindPoint
                                                  pipelineBindPoint,
   VkPipelineLayout
                                                 layout,
   uint32 t
                                                 firstSet,
   uint32_t
                                                 descriptorSetCount,
    const VkDescriptorSet*
                                                 pDescriptorSets,
    uint32_t
                                                 dynamicOffsetCount,
    const uint32_t*
                                                  pDynamicOffsets);
```

- commandBuffer is the command buffer that the descriptor sets will be bound to.
- pipelineBindPoint is a VkPipelineBindPoint indicating whether the descriptors will be used by graphics pipelines or compute pipelines. There is a separate set of bind points for each of graphics and compute, so binding one does not disturb the other.

- layout is a VkPipelineLayout object used to program the bindings.
- firstSet is the set number of the first descriptor set to be bound.
- descriptorSetCount is the number of elements in the pDescriptorSets array.
- pDescriptorSets is an array of handles to VkDescriptorSet objects describing the descriptor sets to write to.
- dynamicOffsetCount is the number of dynamic offsets in the pDynamicOffsets array.
- pDynamicOffsets is a pointer to an array of uint 32_t values specifying dynamic offsets.

vkCmdBindDescriptorSets causes the sets numbered [firstSet. firstSet+descriptorSetCount-1] to use the bindings stored in pDescriptorSets[0..descriptorSetCount-1] for subsequent rendering commands (either compute or graphics, according to the pipelineBindPoint). Any bindings that were previously applied via these sets are no longer valid.

Once bound, a descriptor set affects rendering of subsequent graphics or compute commands in the command buffer until a different set is bound to the same set number, or else until the set is disturbed as described in Pipeline Layout Compatibility.

A compatible descriptor set must be bound for all set numbers that any shaders in a pipeline access, at the time that a draw or dispatch command is recorded to execute using that pipeline. However, if none of the shaders in a pipeline statically use any bindings with a particular set number, then no descriptor set need be bound for that set number, even if the pipeline layout includes a non-trivial descriptor set layout for that set number.

If any of the sets being bound include dynamic uniform or storage buffers, then pDynamicOffsets includes one element for each array element in each dynamic descriptor type binding in each set. Values are taken from pDynamicOffsets in an order such that all entries for set N come before set N+1; within a set, entries are ordered by the binding numbers in the descriptor set layouts; and within a binding array, elements are in order.

dynamicOffsetCount must equal the total number of dynamic descriptors in the sets being bound.

The effective offset used for dynamic uniform and storage buffer bindings is the sum of the relative offset taken from <code>pDynamicOffsets</code>, and the base address of the buffer plus base offset in the descriptor set. The length of the dynamic uniform and storage buffer bindings is the buffer range as specified in the descriptor set.

Each of the pDescriptorSets must be compatible with the pipeline layout specified by layout. The layout used to program the bindings must also be compatible with the pipeline used in subsequent graphics or compute commands, as defined in the Pipeline Layout Compatibility section.

The descriptor set contents bound by a call to **vkCmdBindDescriptorSets** may be consumed during host execution of the command, or during shader execution of the resulting draws, or any time in between. Thus, the contents must not be altered (overwritten by an update command, or freed) between when the command is recorded and when the command completes executing on the queue. The contents of <code>pDynamicOffsets</code> are consumed immediately during execution of **vkCmdBindDescriptorSets**. Once all pending uses have completed, it is legal to update and reuse a descriptor set.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pipelineBindPoint must be a valid VkPipelineBindPoint value
- layout must be a valid VkPipelineLayout handle

- pDescriptorSets must be a pointer to an array of descriptorSetCount valid VkDescriptorSet handles
- If dynamicOffsetCount is not 0, pDynamicOffsets must be a pointer to an array of dynamicOffsetCount uint32_t values
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- descriptorSetCount must be greater than 0
- Each of commandBuffer, layout, and the elements of pDescriptorSets must have been created, allocated, or retrieved from the same VkDevice
- Any given element of pDescriptorSets must have been allocated with a VkDescriptorSetLayout that matches (is the same as, or defined identically to) the VkDescriptorSetLayout at set n in layout, where n is the sum of firstSet and the index into pDescriptorSets
- dynamicOffsetCount must be equal to the total number of dynamic descriptors in pDescriptorSets
- The sum of firstSet and descriptorSetCount must be less than or equal to VkPipelineLayoutCreateInfo::setLayoutCount provided when layout was created
- pipelineBindPoint must be supported by the commandBuffer's parent VkCommandPool's queue family
- Any given element of pDynamicOffsets must satisfy the required alignment for the corresponding descriptor binding's descriptor type

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

13.2.6 Push Constant Updates

As described above in section Pipeline Layouts, the pipeline layout defines shader push constants which are updated via Vulkan commands rather than via writes to memory or copy commands.



Note

Push constants represent a high speed path to modify constant data in pipelines that is expected to outperform memory-backed resource updates.

The values of push constants are undefined at the start of a command buffer.

To update push constants, call:

- commandBuffer is the command buffer in which the push constant update will be recorded.
- layout is the pipeline layout used to program the push constant updates.
- stageFlags is a bitmask of VkShaderStageFlagBits specifying the shader stages that will use the push constants in the updated range.
- offset is the start offset of the push constant range to update, in units of bytes.
- size is the size of the push constant range to update, in units of bytes.
- pValues is an array of size bytes containing the new push constant values.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- layout must be a valid VkPipelineLayout handle
- $\bullet \ \textit{stageFlags} \ \textbf{must} \ \textbf{be} \ \textbf{a} \ \textbf{valid} \ \textbf{combination} \ \textbf{of} \ \textbf{VkShaderStageFlagBits} \ \textbf{values}$
- stageFlags must not be 0
- pValues must be a pointer to an array of size bytes
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- size must be greater than 0
- Both of commandBuffer, and layout must have been created, allocated, or retrieved from the same VkDevice
- stageFlags must match exactly the shader stages used in layout for the range specified by offset and size

- offset must be a multiple of 4
- size must be a multiple of 4
- $\bullet \ \textit{offset must be less than} \ \texttt{VkPhysicalDeviceLimits::} \textit{maxPushConstantsSize}$
- size must be less than or equal to VkPhysicalDeviceLimits::maxPushConstantsSize minus offset

Host Synchronization

 $\bullet \ \ Host\ access\ to\ \textit{commandBuffer}\ must\ be\ externally\ synchronized$

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

Chapter 14

Shader Interfaces

When a pipeline is created, the set of shaders specified in the corresponding Vk * PipelineCreateInfo structure are implicitly linked at a number of different interfaces.

- Shader Input and Output Interface
- Vertex Input Interface
- Fragment Output Interface
- Fragment Input Attachment Interface
- Shader Resource Interface

14.1 Shader Input and Output Interfaces

When multiple stages are present in a pipeline, the outputs of one stage form an interface with the inputs of the next stage. When such an interface involves a shader, shader outputs are matched against the inputs of the next stage, and shader inputs are matched against the outputs of the previous stage.

There are two classes of variables that can be matched between shader stages, built-in variables and user-defined variables. Each class has a different set of matching criteria. Generally, when non-shader stages are between shader stages, the user-defined variables, and most built-in variables, form an interface between the shader stages.

The variables forming the input or output *interfaces* are listed as operands to the **OpEntryPoint** instruction and are declared with the **Input** or **Output** storage classes, respectively, in the SPIR-V module.

Output variables of a shader stage have undefined values until the shader writes to them or uses the **Initializer** operand when declaring the variable.

14.1.1 Built-in Interface Block

Shader built-in variables meeting the following requirements define the built-in interface block. They must

- be explicitly declared (there are no implicit built-ins),
- be identified with a BuiltIn decoration,

- form object types as described in the Built-in Variables section, and
- be declared in a block whose top-level members are the built-ins.

Built-ins only participate in interface matching if they are declared in such a block. They must not have any **Location** or **Component** decorations.

There must be no more than one built-in interface block per shader per interface.

14.1.2 User-defined Variable Interface

The remaining variables listed by **OpEntryPoint** with the **Input** or **Output** storage class form the *user-defined* variable interface. These variables must be identified with a **Location** decoration and can also be identified with a **Component** decoration.

14.1.3 Interface Matching

A user-defined output variable is considered to match an input variable in the subsequent stage if the two variables are declared with the same **Location** and **Component** decoration and match in type and decoration, except that interpolation decorations are not required to match. For the purposes of interface matching, variables declared without a **Component** decoration are considered to have a **Component** decoration of zero.

Variables or block members declared as structures are considered to match in type if and only if the structure members match in type, decoration, number, and declaration order. Variables or block members declared as arrays are considered to match in type only if both declarations specify the same element type and size.

Tessellation control shader per-vertex output variables and blocks, and tessellation control, tessellation evaluation, and geometry shader per-vertex input variables and blocks are required to be declared as arrays, with each element representing input or output values for a single vertex of a multi-vertex primitive. For the purposes of interface matching, the outermost array dimension of such variables and blocks is ignored.

At an interface between two non-fragment shader stages, the built-in interface block must match exactly, as described above. At an interface involving the fragment shader inputs, the presence or absence of any built-in output does not affect the interface matching.

At an interface between two shader stages, the user-defined variable interface must match exactly, as described above.

Any input value to a shader stage is well-defined as long as the preceding stages writes to a matching output, as described above.

Additionally, scalar and vector inputs are well-defined if there is a corresponding output satisfying all of the following conditions:

- the input and output match exactly in decoration,
- the output is a vector with the same basic type and has at least as many components as the input, and
- the common component type of the input and output is 32-bit integer or floating-point (64-bit component types are excluded).

In this case, the components of the input will be taken from the first components of the output, and any extra components of the output will be ignored.

14.1.4 Location Assignment

This section describes how many locations are consumed by a given type. As mentioned above, geometry shader inputs, tessellation control shader inputs and outputs, and tessellation evaluation inputs all have an additional level of arrayness relative to other shader inputs and outputs. This outer array level is removed from the type before considering how many locations the type consumes.

The **Location** value specifies an interface slot comprised of a 32-bit four-component vector conveyed between stages. The **Component** specifies components within these vector locations. Only types with widths of 32 or 64 are supported in shader interfaces.

Inputs and outputs of the following types consume a single interface location:

- 32-bit scalar and vector types, and
- 64-bit scalar and 2-component vector types.

64-bit three- and four-component vectors consume two consecutive locations.

If a declared input or output is an array of size n and each element takes m locations, it will be assigned $m \times n$ consecutive locations starting with the location specified.

If the declared input or output is an $n \times m$ 32- or 64-bit matrix, it will be assigned multiple locations starting with the location specified. The number of locations assigned for each matrix will be the same as for an n-element array of m-component vectors.

The layout of a structure type used as an **Input** or **Output** depends on whether it is also a **Block** (i.e. has a **Block** decoration).

If it is a not a **Block**, then the structure type must have a **Location** decoration. Its members are assigned consecutive locations in their declaration order, with the first member assigned to the location specified for the structure type. The members, and their nested types, must not themselves have **Location** decorations.

If the structure type is a **Block** but without a **Location**, then each of its members must have a **Location** decoration. If it is a **Block** with a **Location** decoration, then its members are assigned consecutive locations in declaration order, starting from the first member which is initially assigned the location specified for the **Block**. Any member with its own **Location** decoration is assigned that location. Each remaining member is assigned the location after the immediately preceding member in declaration order.

The locations consumed by block and structure members are determined by applying the rules above in a depth-first traversal of the instantiated members as though the structure or block member were declared as an input or output variable of the same type.

Any two inputs listed as operands on the same **OpEntryPoint** must not be assigned the same location, either explicitly or implicitly. Any two outputs listed as operands on the same **OpEntryPoint** must not be assigned the same location, either explicitly or implicitly.

The number of input and output locations available for a shader input or output interface are limited, and dependent on the shader stage as described in Table 14.1.

Shader Interface	Locations Available	
vertex input	maxVertexInputAttributes	
vertex output	maxVertexOutputComponents/4	
tessellation control input	maxTessellationControlPerVertexInputComponents/4	
tessellation control output	maxTessellationControlPerVertexOutputComponents/4	

Table 14.1: Shader Input and Output Locations

Table 14.1: (continued)

Shader Interface	Locations Available	
tessellation evaluation input	maxTessellationEvaluationInputComponents/4	
tessellation evaluation output	maxTessellationEvaluationOutputComponents/4	
geometry input	maxGeometryInputComponents/4	
geometry output	maxGeometryOutputComponents/4	
fragment input	maxFragmentInputComponents/4	
fragment output	maxFragmentOutputAttachments	

14.1.5 Component Assignment

The Component decoration allows the Location to be more finely specified for scalars and vectors, down to the individual components within a location that are consumed. The components within a location are 0, 1, 2, and 3. A variable or block member starting at component N will consume components N, N+1, N+2, ... up through its size. For single precision types, it is invalid if this sequence of components gets larger than 3. A scalar 64-bit type will consume two of these components in sequence, and a two-component 64-bit vector type will consume all four components available within a location. A three- or four-component 64-bit vector type must not specify a Component decoration. A three-component 64-bit vector type will consume all four components of the first location and components 0 and 1 of the second location. This leaves components 2 and 3 available for other component-qualified declarations.

A scalar or two-component 64-bit data type must not specify a **Component** decoration of 1 or 3. A **Component** decoration must not be specified for any type that is not a scalar or vector.

14.2 Vertex Input Interface

When the vertex stage is present in a pipeline, the vertex shader input variables form an interface with the vertex input attributes. The vertex shader input variables are matched by the **Location** and **Component** decorations to the vertex input attributes specified in the <code>pVertexInputState</code> member of the <code>VkGraphicsPipelineCreateInfo</code> structure.

The vertex shader input variables listed by **OpEntryPoint** with the **Input** storage class form the *vertex input interface*. These variables must be identified with a **Location** decoration and can also be identified with a **Component** decoration.

For the purposes of interface matching: variables declared without a **Component** decoration are considered to have a **Component** decoration of zero. The number of available vertex input locations is given by the <code>maxVertexInputAttributes</code> member of the VkPhysicalDeviceLimits structure.

See Section 20.1.1 for details.

All vertex shader inputs declared as above must have a corresponding attribute and binding in the pipeline.

14.3 Fragment Output Interface

When the fragment stage is present in a pipeline, the fragment shader outputs form an interface with the output attachments of the current subpass. The fragment shader output variables are matched by the **Location** and **Component** decorations to the color attachments specified in the *pColorAttachments* array of the VkSubpassDescription structure that describes the subpass that the fragment shader is executed in.

The fragment shader output variables listed by **OpEntryPoint** with the **Output** storage class form the *fragment* output interface. These variables must be identified with a **Location** decoration. They can also be identified with a **Component** decoration and/or an **Index** decoration. For the purposes of interface matching: variables declared without a **Component** decoration are considered to have a **Component** decoration of zero, and variables declared without an **Index** decoration are considered to have an **Index** decoration of zero.

A fragment shader output variable identified with a **Location** decoration of i is directed to the color attachment indicated by pColorAttachments[i], after passing through the blending unit as described in Section 26.1, if enabled. Locations are consumed as described in Location Assignment. The number of available fragment output locations is given by the maxFragmentOutputAttachments member of the VkPhysicalDeviceLimits structure.

Components of the output variables are assigned as described in Component Assignment. Output components identified as 0, 1, 2, and 3 will be directed to the R, G, B, and A inputs to the blending unit, respectively, or to the output attachment if blending is disabled. If two variables are placed within the same location, they must have the same underlying type (floating-point or integer). The input to blending or color attachment writes is undefined for components which do not correspond to a fragment shader output.

Fragment outputs identified with an **Index** of zero are directed to the first input of the blending unit associated with the corresponding **Location**. Outputs identified with an **Index** of one are directed to the second input of the corresponding blending unit.

No *component aliasing* of output variables is allowed, that is there must not be two output variables which have the same location, component, and index, either explicitly declared or implied.

Output values written by a fragment shader must be declared with either **OpTypeFloat** or **OpTypeInt**, and a Width of 32. Composites of these types are also permitted. If the color attachment has a signed or unsigned normalized fixed-point format, color values are assumed to be floating-point and are converted to fixed-point as described in Section 2.8.1; otherwise no type conversion is applied. If the type of the values written by the fragment shader do not match the format of the corresponding color attachment, the result is undefined for those components.

14.4 Fragment Input Attachment Interface

When a fragment stage is present in a pipeline, the fragment shader subpass inputs form an interface with the input attachments of the current subpass. The fragment shader subpass input variables are matched by **InputAttachmentIndex** decorations to the input attachments specified in the pInputAttachments array of the VkSubpassDescription structure that describes the subpass that the fragment shader is executed in.

The fragment shader subpass input variables with the **UniformConstant** storage class and a decoration of **InputAttachmentIndex** that are statically used by **OpEntryPoint** form the *fragment input attachment interface*. These variables must be declared with a type of **OpTypeImage**, a **Dim** operand of **SubpassData**, and a **Sampled** operand of 2.

A subpass input variable identified with an **InputAttachmentIndex** decoration of *i* reads from the input attachment indicated by pInputAttachments[*i*] member of VkSubpassDescription. If the subpass input variable is declared as an array of size N, it consumes N consecutive input attachments, starting with the index specified. There must not be more than one input variable with the same **InputAttachmentIndex** whether explicitly declared or implied by an array declaration. The number of available input attachment indices is given by the maxPerStageDescriptorInputAttachments member of the VkPhysicalDeviceLimits structure.

Variables identified with the **InputAttachmentIndex** must only be used by a fragment stage. The basic data type (floating-point, integer, unsigned integer) of the subpass input must match the basic format of the corresponding input attachment, or the values of subpass loads from these variables are undefined.

See Section 13.1.11 for more details.

14.5 Shader Resource Interface

When a shader stage accesses buffer or image resources, as described in the Resource Descriptors section, the shader resource variables must be matched with the pipeline layout that is provided at pipeline creation time.

The set of shader resources that form the *shader resource interface* for a stage are the variables statically used by **OpEntryPoint** with the storage class of **Uniform**, **UniformConstant**, or **PushConstant**. For the fragment shader, this includes the fragment input attachment interface.

The shader resource interface consists of two sub-interfaces: the push constant interface and the descriptor set interface.

14.5.1 Push Constant Interface

The shader variables defined with a storage class of **PushConstant** that are statically used by the shader entry-points for the pipeline define the *push constant interface*. They must be:

- typed as OpTypeStruct,
- · identified with a Block decoration, and
- laid out explicitly using the Offset, ArrayStride, and MatrixStride decorations as specified in Offset and Stride Assignment.

There must be no more than one push constant block statically used per shader entry-point.

Each variable in a push constant block must be placed at an **Offset** such that the entire constant value is entirely contained within the VkPushConstantRange for each **OpEntryPoint** that uses it, and the <code>stageFlags</code> for that range must specify the appropriate VkShaderStageFlagBits for that stage. The **Offset** decoration for any variable in a push constant block must not cause the space required for that variable to extend outside the range [0, maxPushConstantsSize).

Any variable in a push constant block that is declared as an array must only be accessed with dynamically uniform indices.

14.5.2 Descriptor Set Interface

The *descriptor set interface* is comprised of the shader variables with the storage class of **Uniform** or **UniformConstant** (including the variables in the fragment input attachment interface) that are statically used by the shader entry-points for the pipeline.

These variables must have **DescriptorSet** and **Binding** decorations specified, which are assigned and matched with the VkDescriptorSetLayout objects in the pipeline layout as described in DescriptorSet and Binding Assignment.

Variables identified with the **UniformConstant** storage class are used only as handles to refer to opaque resources. Such variables must be typed as **OpTypeImage**, **OpTypeSampler**, **OpTypeSampledImage**, or arrays of only these types. Variables of type **OpTypeImage** must have a **Sampled** operand of 1 (sampled image) or 2 (storage image).

Any array of these types must only be indexed with constant integral expressions, except under the following conditions:

• For arrays of OpTypeImage variables with Sampled operand of 2, if the shaderStorageImageArrayDynamicIndexing feature is enabled and the shader module declares the StorageImageArrayDynamicIndexing capability, the array must only be indexed by dynamically uniform expressions.

• For arrays of OpTypeSampler, OpTypeSampledImage variables, or OpTypeImage variables with Sampled operand of 1, if the <code>shaderSampledImageArrayDynamicIndexing</code> feature is enabled and the shader module declares the SampledImageArrayDynamicIndexing capability, the array must only be indexed by dynamically uniform expressions.

The **Sampled Type** of an **OpTypeImage** declaration must match the same basic data type as the corresponding resource, or the values obtained by reading or sampling from this image are undefined.

The Image Format of an OpTypeImage declaration must not be Unknown, for variables which are used for OpImageRead or OpImageWrite operations, except under the following conditions:

- For OpImageWrite, if the shaderStorageImageWriteWithoutFormat feature is enabled and the shader module declares the StorageImageWriteWithoutFormat capability.
- For OpImageRead, if the shaderStorageImageReadWithoutFormat feature is enabled and the shader module declares the StorageImageReadWithoutFormat capability.

Variables identified with the **Uniform** storage class are used to access transparent buffer backed resources. Such variables must be:

• typed as **OpTypeStruct**, or arrays of only this type,

storage image

- identified with a **Block** or **BufferBlock** decoration, and
- laid out explicitly using the Offset, ArrayStride, and MatrixStride decorations as specified in Offset and Stride Assignment.

Any array of these types must only be indexed with constant integral expressions, except under the following conditions.

- For arrays of Block variables, if the shaderUniformBufferArrayDynamicIndexing feature is enabled and the shader module declares the UniformBufferArrayDynamicIndexing capability, the array must only be indexed by dynamically uniform expressions.
- For arrays of **BufferBlock** variables, if the *shaderStorageBufferArrayDynamicIndexing* feature is enabled and the shader module declares the **StorageBufferArrayDynamicIndexing** capability, the array must only be indexed by dynamically uniform expressions.

The **Offset** decoration for any variable in a **Block** must not cause the space required for that variable to extend outside the range [0, maxUniformBufferRange). The **Offset** decoration for any variable in a **BufferBlock** must not cause the space required for that variable to extend outside the range [0, maxStorageBufferRange).

Variables identified with a storage class of **UniformConstant** and a decoration of **InputAttachmentIndex** must be declared as described in Fragment Input Attachment Interface.

Each shader variable declaration must refer to the same type of resource as is indicated by the <code>descriptorType</code>. See Shader Resource and Descriptor Type Correspondence for the relationship between shader declarations and descriptor types.

Resource type	Descriptor Type
sampler	VK_DESCRIPTOR_TYPE_SAMPLER
sampled image	VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE

VK_DESCRIPTOR_TYPE_STORAGE_IMAGE

Table 14.2: Shader Resource and Descriptor Type Correspondence

Table 14.2: (continued)

Resource type	Descriptor Type
combined image sampler	VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER
uniform texel buffer	VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER
storage texel buffer	VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER
uniform buffer	VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER
	VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC
storage buffer	VK_DESCRIPTOR_TYPE_STORAGE_BUFFER
	VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC
input attachment	VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT

Table 14.3: Shader Resource and Storage Class Correspondence

Resource type	Storage Class	Туре	Decoration(s) ¹
sampler	UniformConstant	OpTypeSampler	
sampled image	UniformConstant	OpTypeImage	
		(Sampled=1)	
storage image	UniformConstant	OpTypeImage	
		(Sampled=2)	
combined image	UniformConstant	OpTypeSampledImage	
sampler			
uniform texel buffer	UniformConstant	OpTypeImage	
		(Dim=Buffer,	
		Sampled=1)	
storage texel buffer	UniformConstant	OpTypeImage	
		(Dim=Buffer,	
		Sampled=2)	
uniform buffer	Uniform	OpTypeStruct	Block, Offset,
			(ArrayStride),
			(MatrixStride)
storage buffer	Uniform	OpTypeStruct	BufferBlock, Offset,
			(ArrayStride),
			(MatrixStride)
input attachment	UniformConstant	OpTypeImage	InputAttachmentIndex
		(Dim=SubpassData,	
		Sampled=2)	

in addition to DescriptorSet and Binding

1

14.5.3 DescriptorSet and Binding Assignment

A variable identified with a **DescriptorSet** decoration of s and a **Binding** decoration of b indicates that this variable is associated with the VkDescriptorSetLayoutBinding that has a binding equal to b in pSetLayouts[s] that was specified in VkPipelineLayoutCreateInfo.

The range of descriptor sets is between zero and <code>maxBoundDescriptorSets</code> minus one. If a descriptor set value is statically used by an entry-point there must be an associated <code>pSetLayout</code> in the corresponding pipeline layout as described in Pipeline Layouts consistency.

If the **Binding** decoration is used with an array, the entire array is identified with that binding value. The size of the array declaration must be no larger than the <code>descriptorCount</code> of that <code>VkDescriptorSetLayoutBinding</code>. The index of each element of the array is referred to as the <code>arrayElement</code>. For the purposes of interface matching and descriptor set operations, if a resource variable is not an array, it is treated as if it has an arrayElement of zero.

The binding can be any 32-bit unsigned integer value, as described in Section 13.2.1. Each descriptor set has its own binding name space.

There is a limit on the number of resources of each type that can be accessed by a pipeline stage as shown in Shader Resource Limits. The "Resources Per Stage" column gives the limit on the number each type of resource that can be statically used for an entry-point in any given stage in a pipeline. The "Resource Types" column lists which resource types are counted against the limit. Some resource types count against multiple limits.

If multiple entry-points in the same pipeline refer to the same set and binding, all variable definitions with that **DescriptorSet** and **Binding** must have the same basic type.

Not all descriptor sets and bindings specified in a pipeline layout need to be used in a particular shader stage or pipeline, but if a **DescriptorSet** and **Binding** decoration is specified for a variable that is statically used in that shader there must be a pipeline layout entry identified with that descriptor set and binding and the corresponding stageFlags must specify the appropriate VkShaderStageFlagBits for that stage.

Resources per Stage	Resource Types	
maxPerStageDescriptorSamplers	sampler	
maxrerstageDescriptorsamplers	combined image sampler	
	sampled image	
maxPerStageDescriptorSampledImages	combined image sampler	
	uniform texel buffer	
maxPerStageDescriptorStorageImages	storage image	
maxrerstageDescriptorstoragemages	storage texel buffer	
maxPerStageDescriptorUniformBuffers	uniform buffer	
maxrerstageDescriptorOnnormBuriers	uniform buffer dynamic	
may Dar Staga Dagarintar Staraga Pufforg	storage buffer	
maxPerStageDescriptorStorageBuffers	storage buffer dynamic	
maxPerStageDescriptorInputAttachments	input attachment ¹	

Table 14.4: Shader Resource Limits

1

Input attachments can only be used in the fragment shader stage

14.5.4 Offset and Stride Assignment

All variables with a storage class of **PushConstant** or **Uniform** must be explicitly laid out using the **Offset**, **ArrayStride**, and **MatrixStride** decorations. There are two different layouts requirements depending on the specific resources.

Standard Uniform Buffer Layout

Member variables of an **OpTypeStruct** with storage class of **Uniform** and a decoration of **Block** (uniform buffers) must be laid out according to the following rules.

- The **Offset** Decoration must be a multiple of its base alignment, computed recursively as follows:
 - a scalar of size N has a base alignment of N
 - a two-component vector, with components of size N, has a base alignment of 2N
 - a three- or four-component vector, with components of size N, has a base alignment of 4N
 - an array has a base alignment equal to the base alignment of its element type, rounded up to a multiple of 16
 - a structure has a base alignment equal to the largest base alignment of any of its members, rounded up to a multiple of 16
 - a row-major matrix of C columns has a base alignment equal to the base alignment of vector of C matrix components
 - a column-major matrix has a base alignment equal to the base alignment of the matrix column type
- Any ArrayStride or MatrixStride decoration must be an integer multiple of the base alignment of the array or matrix from above.
- The **Offset** Decoration of a member must not place it between the end of a structure or an array and the next multiple of the base alignment of that structure or array.
- The numeric order of Offset Decorations need not follow member declaration order.



Note

The std140 layout in GLSL satisfies these rules.

Standard Storage Buffer Layout

Member variables of an **OpTypeStruct** with a storage class of **PushConstant** (push constants), or a storage class of **Uniform** with a decoration of **BufferBlock** (storage buffers) must be laid out as above, except for array and structure base alignment which do not need to be rounded up to a multiple of 16.



Note

The std430 layout in GLSL satisfies these rules.

14.6 Built-In Variables

Built-in variables are accessed in shaders by declaring a variable decorated with a **BuiltIn** decoration. The meaning of each **BuiltIn** decoration is as follows. In the remainder of this section, the name of a built-in is used interchangeably with a term equivalent to a variable decorated with that particular built-in. Built-ins that represent integer values can be declared as either signed or unsigned 32-bit integers.

BaryCoordNoPerspAMD

The **BaryCoordNoPerspAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using linear interpolation at the pixel's center. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

BaryCoordNoPerspCentroidAMD

The **BaryCoordNoPerspCentroidAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using linear interpolation at the centroid. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

BaryCoordNoPerspSampleAMD

The **BaryCoordNoPerspCentroidAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using linear interpolation at each covered sample. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

BaryCoordPullModelAMD

The **BaryCoordPullModelAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain (1/W, 1/I, 1/J) evaluated at the pixel center and can be used to calculate gradients and then interpolate I, J, and W at any desired sample location.

BaryCoordSmoothAMD

The **BaryCoordSmoothAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using perspective interpolation at the pixel's center. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

BaryCoordSmoothCentroidAMD

The **BaryCoordSmoothCentroidAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using perspective interpolation at the centroid. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

BaryCoordSmoothSampleAMD

The **BaryCoordSmoothCentroidAMD** decoration can be used to decorate a fragment shader input variable. This variable will contain the (I,J) pair of the barycentric coordinates corresponding to the fragment evaluated using perspective interpolation at each covered sample. The K coordinate of the barycentric coordinates can be derived given the identity I + J + K = 1.0.

ClipDistance

Decorating a variable with the **ClipDistance** built-in decoration will make that variable contain the mechanism for controlling user clipping. **ClipDistance** is an array such that the ith element of the array specifies the clip distance for plane i. A clip distance of 0 means the vertex is on the plane, a positive distance means the vertex is inside the clip half-space, and a negative distance means the point is outside the clip half-space.

The **ClipDistance** decoration must be used only within vertex, fragment, tessellation control, tessellation evaluation, and geometry shaders.

In vertex shaders, any variable decorated with **ClipDistance** must be declared using the output storage class.

In fragment shaders, any variable decorated with ClipDistance must be declared using the input storage class.

In tessellation control, tessellation evaluation, or geometry shaders, any variable decorated with **ClipDistance** must not be in a storage class other than input or output.

Any variable decorated with **ClipDistance** must be declared as an array of 32-bit floating-point values.



Note

The array variable decorated with **ClipDistance** is explicitly sized by the shader.



Note

In the last vertex processing stage, these values will be linearly interpolated across the primitive and the portion of the primitive with interpolated distances less than 0 will be considered outside the clip volume. If ClipDistance is then used by a fragment shader, ClipDistance contains these linearly interpolated values.

CullDistance

Decorating a variable with the **CullDistance** built-in decoration will make that variable contain the mechanism for controlling user culling. If any member of this array is assigned a negative value for all vertices belonging to a primitive, then the primitive is discarded before rasterization.

The **CullDistance** decoration must be used only within vertex, fragment, tessellation control, tessellation evaluation, and geometry shaders.

In vertex shaders, any variable decorated with **CullDistance** must be declared using the output storage class.

In fragment shaders, any variable decorated with **CullDistance** must be declared using the input storage class.

In tessellation control, tessellation evaluation, or geometry shaders, any variable decorated with **CullDistance** must not be declared in a storage class other than input or output.

Any variable decorated with **CullDistance** must be declared as an array of 32-bit floating-point values.



Note

In fragment shaders, the values of the CullDistance array are linearly interpolated across each primitive.



Note

If **CullDistance** decorates an input variable, that variable will contain the corresponding value from the **CullDistance** decorated output variable from the previous shader stage.

${\tt FragCoord}$

Decorating a variable with the **FragCoord** built-in decoration will make that variable contain the framebuffer coordinate $(x, y, z, \frac{1}{w})$ of the fragment being processed. The (x, y) coordinate (0, 0) is the upper left corner of the upper left pixel in the framebuffer.

When sample shading is enabled, the *x* and *y* components of **FragCoord** reflect the location of the sample corresponding to the shader invocation.

When sample shading is not enabled, the x and y components of **FragCoord** reflect the location of the center of the pixel, (0.5, 0.5).

The z component of **FragCoord** is the interpolated depth value of the primitive.

The w component is the interpolated $\frac{1}{w}$.

The **FragCoord** decoration must be used only within fragment shaders.

The variable decorated with **FragCoord** must be declared using the input storage class.

The **Centroid** interpolation decoration is ignored on **FragCoord**.

The variable decorated with **FragCoord** must be declared as a four-component vector of 32-bit floating-point values.

FragDepth

Decorating a variable with the **FragDepth** built-in decoration will make that variable contain the new depth value for all samples covered by the fragment. This value will be used for depth testing and, if the depth test passes, any subsequent write to the depth/stencil attachment.

To write to **FragDepth**, a shader must declare the **DepthReplacing** execution mode. If a shader declares the **DepthReplacing** execution mode and there is an execution path through the shader that does not set **FragDepth**, then the fragment's depth value is undefined for executions of the shader that take that path.

The **FragDepth** decoration must be used only within fragment shaders.

The variable decorated with **FragDepth** must be declared using the output storage class.

The variable decorated with FragDepth must be declared as a scalar 32-bit floating-point value.

FrontFacing

Decorating a variable with the **FrontFacing** built-in decoration will make that variable contain whether a primitive is front or back facing. This variable is non-zero if the current fragment is considered to be part of a front-facing primitive and is zero if the fragment is considered to be part of a back-facing primitive.

The **FrontFacing** decoration must be used only within fragment shaders.

The variable decorated with **FrontFacing** must be declared using the input storage class.

The variable decorated with **FrontFacing** must be declared as a boolean.

GlobalInvocationId

Decorating a variable with the **GlobalInvocationId** built-in decoration will make that variable contain the location of the current invocation within the global workgroup. Each component is equal to the index of the local workgroup multiplied by the size of the local workgroup plus **LocalInvocationId**.

The **GlobalInvocationId** decoration must be used only within compute shaders.

The variable decorated with **GlobalInvocationId** must be declared using the input storage class.

The variable decorated with **GlobalInvocationId** must be declared as a three-component vector of 32-bit integers.

HelperInvocation

Decorating a variable with the **HelperInvocation** built-in decoration will make that variable contain whether the current invocation is a helper invocation. This variable is non-zero if the current fragment being shaded is a helper invocation and zero otherwise. A helper invocation is an invocation of the shader that is produced to satisfy internal requirements such as the generation of derivatives.

The **HelperInvocation** decoration must be used only within fragment shaders.

The variable decorated with **HelperInvocation** must be declared using the input storage class.

The variable decorated with **HelperInvocation** must be declared as a boolean.



Note

It is very likely that a helper invocation will have a value of **SampleMask** fragment shader input value that is zero.

InvocationId

Decorating a variable with the **InvocationId** built-in decoration will make that variable contain the index of the current shader invocation in a geometry shader, or the index of the output patch vertex in a tessellation control shader.

In a geometry shader, the index of the current shader invocation ranges from zero to the number of instances declared in the shader minus one. If the instance count of the geometry shader is one or is not specified, then **InvocationId** will be zero.

The **InvocationId** decoration must be used only within tessellation control and geometry shaders.

The variable decorated with **InvocationId** must be declared using the input storage class.

The variable decorated with **InvocationId** must be declared as a scalar 32-bit integer.

InstanceIndex

Decorating a variable with the **InstanceIndex** built-in decoration will make that variable contain the index of the instance that is being processed by the current vertex shader invocation. **InstanceIndex** begins at the <code>firstInstance</code> parameter to <code>vkCmdDraw</code> or <code>vkCmdDrawIndexed</code> or at the <code>firstInstance</code> member of a structure consumed by <code>vkCmdDrawIndirect</code> or <code>vkCmdDrawIndexedIndirect</code>.

The **InstanceIndex** decoration must be used only within vertex shaders.

The variable decorated with **InstanceIndex** must be declared using the input storage class.

The variable decorated with **InstanceIndex** must be declared as a scalar 32-bit integer.

Layer

Decorating a variable with the **Layer** built-in decoration will make that variable contain the select layer of a multi-layer framebuffer attachment.

In a geometry shader, any variable decorated with **Layer** can be written with the framebuffer layer index to which the primitive produced by the geometry shader will be directed. If a geometry shader entry-point's interface does not include a variable decorated with **Layer**, then the first layer is used. If a geometry shader entry-point's interface includes a variable decorated with **Layer**, it must write the same value to **Layer** for all output vertices of a given primitive.

In a fragment shader, a variable decorated with **Layer** contains the layer index of the primitive that the fragment invocation belongs to.

The **Layer** decoration must be used only within geometry and fragment shaders.

In a geometry shader, any variable decorated with **Layer** must be declared using the output storage class.

In a fragment shader, any variable decorated with **Layer** must be declared using the input storage class.

Any variable decorated with **Layer** must be declared as a scalar 32-bit integer.

LocalInvocationId

Decorating a variable with the **LocalInvocationId** built-in decoration will make that variable contain the location of the current compute shader invocation within the local workgroup. Each component ranges from zero through to the size of the workgroup in that dimension minus one.

The **LocalInvocationId** decoration must be used only within compute shaders.

The variable decorated with **LocalInvocationId** must be declared using the input storage class.

The variable decorated with **LocalInvocationId** must be declared as a three-component vector of 32-bit integers.



Note

If the size of the workgroup in a particular dimension is one, then the **LocalInvocationId** in that dimension will be zero. If the workgroup is effectively two-dimensional, then **LocalInvocationId**.z will be zero. If the workgroup is effectively one-dimensional, then both **LocalInvocationId**.y and **LocalInvocationId**.z will be zero.

NumWorkgroups

Decorating a variable with the **NumWorkgroups** built-in decoration will make that variable contain the number of local workgroups that are part of the dispatch that the invocation belongs to. Each component is equal to the values of the parameters passed into vkCmdDispatch or read from the VkDispatchIndirectCommand structure read through a call to vkCmdDispatchIndirect.

The **NumWorkgroups** decoration must be used only within compute shaders.

The variable decorated with **NumWorkgroups** must be declared using the input storage class.

The variable decorated with **NumWorkgroups** must be declared as a three-component vector of 32-bit integers.

PatchVertices

Decorating a variable with the **PatchVertices** built-in decoration will make that variable contain the number of vertices in the input patch being processed by the shader. A single tessellation control or tessellation evaluation shader can read patches of differing sizes, so the value of the **PatchVertices** variable may differ between patches.

The **PatchVertices** decoration must be used only within tessellation control and tessellation evaluation shaders.

The variable decorated with PatchVertices must be declared using the input storage class.

The variable decorated with **PatchVertices** must be declared as scalar 32-bit integer.

PointCoord

Decorating a variable with the **PointCoord** built-in decoration will make that variable contain the coordinate of the current fragment within the point being rasterized, normalized to the size of the point with origin in the upper left corner of the point, as described in Basic Point Rasterization. If the primitive the fragment shader invocation belongs to is not a point, then the variable decorated with **PointCoord** contains an undefined value.

The **PointCoord** decoration must be used only within fragment shaders.

The variable decorated with **PointCoord** must be declared using the input storage class.

The variable decorated with **PointCoord** must be declared as two-component vector of 32-bit floating-point values.



Note

Depending on how the point is rasterized, **PointCoord** may never reach (0,0) or (1,1).

PointSize

Decorating a variable with the **PointSize** built-in decoration will make that variable contain the size of point primitives. The value written to the variable decorated with **PointSize** by the last vertex processing stage in the pipeline is used as the framebuffer-space size of points produced by rasterization.

The **PointSize** decoration must be used only within vertex, tessellation control, tessellation evaluation, and geometry shaders.

In a vertex shader, any variable decorated with **PointSize** must be declared using the output storage class.

In a tessellation control, tessellation evaluation, or geometry shader, any variable decorated with **PointSize** must be declared using either the input or output storage class.

Any variable decorated with **PointSize** must be declared as a scalar 32-bit floating-point value.



Note

When **PointSize** decorates a variable in the input storage class, it contains the data written to the output variable decorated with **PointSize** from the previous shader stage.

Position

Decorating a variable with the **Position** built-in decoration will make that variable contain the position of the current vertex. In the last vertex processing stage, the value of the variable decorated with **Position** is used in subsequent primitive assembly, clipping, and rasterization operations.

The **Position** decoration must be used only within vertex, tessellation control, tessellation evaluation, and geometry shaders.

In a vertex shader, any variable decorated with **Position** must be declared using the output storage class.

In a tessellation control, tessellation evaluation, or geometry shader, any variable decorated with **Position** must not be declared in a storage class other than input or output.

Any variable decorated with **Position** must be declared as a four-component vector of 32-bit floating-point values.



Note

When **Position** decorates a variable in the input storage class, it contains the data written to the output variable decorated with **Position** from the previous shader stage.

PrimitiveId

Decorating a variable with the **PrimitiveId** built-in decoration will make that variable contain the index of the current primitive.

In tessellation control and tessellation evaluation shaders, it will contain the index of the patch within the current set of rendering primitives that correspond to the shader invocation.

In a geometry shader, it will contain the number of primitives presented as input to the shader since the current set of rendering primitives was started.

In a fragment shader, it will contain the primitive index written by the geometry shader if a geometry shader is present, or with the value that would have been presented as input to the geometry shader had it been present.

If a geometry shader is present and the fragment shader reads from an input variable decorated with **PrimitiveId**, then the geometry shader must write to an output variable decorated with **PrimitiveId** in all execution paths.

The **PrimitiveId** decoration must be used only within fragment, tessellation control, tessellation evaluation, and geometry shaders.

In a fragment, tessellation control or tessellation evaluation shader, any variable decorated with **PrimitiveId** must be declared using the output storage class.

In a geometry shader, any variable decorated with **PrimitiveId** must be declared using either the input or output storage class.

Any variable decorated with **PrimitiveId** must be declared as scalar 32-bit integer.



Note

When the **PrimitiveId** decoration is applied to an output variable in the geometry shader, the resulting value is seen through the **PrimitiveId** decorated input variable in the fragment shader.

SampleId

Decorating a variable with the **SampleId** built-in decoration will make that variable contain the zero-based index of the sample the invocation corresponds to. **SampleId** ranges from zero to the number of samples in the framebuffer minus one. If a fragment shader entry-point's interface includes an input variable decorated with **SampleId**, per-sample shading is enabled for draws that use that fragment shader.

The **SampleId** decoration must be used only within fragment shaders.

The variable decorated with **SampleId** must be declared using the input storage class.

The variable decorated with **SampleId** must be declared as a scalar 32-bit integer.

SampleMask

Decorating a variable with the **SampleMask** built-in decoration will make any variable contain the sample coverage mask for the current fragment shader invocation.

A variable in the input storage class decorated with **SampleMask** will contain a bitmask of the set of samples covered by the primitive generating the fragment during rasterization. It has a sample bit set if and only if the sample is considered covered for this fragment shader invocation. **SampleMask**[] is an array of integers. Bits are mapped to samples in a manner where bit B of mask M (SampleMask [M]) corresponds to sample $32 \times M + B$.

When state specifies multiple fragment shader invocations for a given fragment, the sample mask for any single fragment shader invocation specifies the subset of the covered samples for the fragment that correspond to the invocation. In this case, the bit corresponding to each covered sample will be set in exactly one fragment shader invocation.

A variable in the output storage class decorated with **SampleMask** is an array of integers forming a bit array in a manner similar an input variable decorated with **SampleMask**, but where each bit represents coverage as computed by the shader. Modifying the sample mask by writing zero to a bit of **SampleMask** causes the sample to be considered uncovered. However, setting sample mask bits to one will never enable samples not covered by the original primitive. If the fragment shader is being evaluated at any frequency other than per-fragment, bits of the sample mask not corresponding to the current fragment shader invocation are ignored. This array must be sized in the fragment shader either implicitly or explicitly, to be no larger than the implementation-dependent maximum sample-mask (as an array of 32-bit elements), determined by the maximum number of samples. If a fragment shader entry-point's interface includes an output variable decorated with **SampleMask**, the sample mask will be undefined for any array elements of any fragment shader invocations that fail to assign a value. If a fragment shader entry-point's interface does not include an output variable decorated with **SampleMask**, the sample mask has no effect on the processing of a fragment.

The **SampleMask** decoration must be used only within fragment shaders.

Any variable decorated with SampleMask must be declared using either the input or output storage class.

Any variable decorated with **SampleMask** must be declared as an array of 32-bit integers.

SamplePosition

Decorating a variable with the **SamplePosition** built-in decoration will make that variable contain the sub-pixel position of the sample being shaded. The top left of the pixel is considered to be at coordinate (0,0) and the bottom right of the pixel is considered to be at coordinate (1,1). If a fragment shader entry-point's interface includes an input variable decorated with **SamplePosition**, per-sample shading is enabled for draws that use that fragment shader.

The **SamplePosition** decoration must be used only within fragment shaders.

The variable decorated with **SamplePosition** must be declared using the input storage class.

The variable decorated with **SamplePosition** must be declared as a two-component vector of 32-bit floating-point values.

TessCoord

Decorating a variable with the **TessCoord** built-in decoration will make that variable contain the three-dimensional (u, v, w) barycentric coordinate of the tessellated vertex within the patch. u, v, and w are in the range [0,1] and vary linearly across the primitive being subdivided. For the tessellation modes of **Quads** or **IsoLines**, the third component is always zero.

The **TessCoord** decoration must be used only within tessellation evaluation shaders.

The variable decorated with **TessCoord** must be declared using the input storage class.

The variable decorated with **TessCoord** must be declared as three-component vector of 32-bit floating-point values.

TessLevelOuter

Decorating a variable with the **TessLevelOuter** built-in decoration will make that variable contain the outer tessellation levels for the current patch.

In tessellation control shaders, the variable decorated with **TessLevelOuter** can be written to which controls the tessellation factors for the resulting patch. These values are used by the tessellator to control primitive tessellation and can be read by tessellation evaluation shaders.

In tessellation evaluation shaders, the variable decorated with **TessLevelOuter** can read the values written by the tessellation control shader.

The **TessLevelOuter** decoration must be used only within tessellation control and tessellation evaluation shaders.

In a tessellation control shader, any variable decorated with **TessLevelOuter** must be declared using the output storage class.

In a tessellation evaluation shader, any variable decorated with **TessLevelOuter** must be declared using the input storage class.

Any variable decorated with **TessLevelOuter** must be declared as an array of size four, containing 32-bit floating-point values.

TessLevelInner

Decorating a variable with the **TessLevelInner** built-in decoration will make that variable contain the inner tessellation levels for the current patch.

In tessellation control shaders, the variable decorated with **TessLevelInner** can be written to, which controls the tessellation factors for the resulting patch. These values are used by the tessellator to control primitive tessellation and can be read by tessellation evaluation shaders.

In tessellation evaluation shaders, the variable decorated with **TessLevelInner** can read the values written by the tessellation control shader.

The **TessLevelInner** decoration must be used only within tessellation control and tessellation evaluation shaders.

In a tessellation control shader, any variable decorated with **TessLevelInner** must be declared using the output storage class.

In a tessellation evaluation shader, any variable decorated with **TessLevelInner** must be declared using the input storage class.

Any variable decorated with **TessLevelInner** must be declared as an array of size two, containing 32-bit floating-point values.

VertexIndex

Decorating a variable with the **VertexIndex** built-in decoration will make that variable contain the index of the vertex that is being processed by the current vertex shader invocation. For non-indexed draws, this variable begins at the <code>firstVertex</code> parameter to <code>vkCmdDraw</code> or the <code>firstVertex</code> member of a structure consumed by <code>vkCmdDrawIndirect</code> and increments by one for each vertex in the draw. For indexed draws, its value is the content of the index buffer for the vertex plus the <code>vertexOffset</code> parameter to <code>vkCmdDrawIndexed</code> or the <code>vertexOffset</code> member of the structure consumed by <code>vkCmdDrawIndexedIndirect</code>.

The **VertexIndex** decoration must be used only within vertex shaders.

The variable decorated with **VertexIndex** must be declared using the input storage class.

The variable decorated with **VertexIndex** must be declared as a scalar 32-bit integer.



Note

VertexIndex starts at the same starting value for each instance.

ViewportIndex

Decorating a variable with the **ViewportIndex** built-in decoration will make that variable contain the index of the viewport.

In a geometry shader, the variable decorated with **ViewportIndex** can be written to with the viewport index to which the primitive produced by the geometry shader will be directed. The selected viewport index is used to select the viewport transform and scissor rectangle. If a geometry shader entry-point's interface does not include a variable decorated with **ViewportIndex**, then the first viewport is used. If a geometry shader entry-point's interface includes a variable decorated with **ViewportIndex**, it must write the same value to **ViewportIndex** for all output vertices of a given primitive.

In a fragment shader, the variable decorated with **ViewportIndex** contains the viewport index of the primitive that the fragment invocation belongs to.

The **ViewportIndex** decoration must be used only within geometry and fragment shaders.

In a geometry shader, any variable decorated with **ViewportIndex** must be declared using the output storage class

In a fragment shader, any variable decorated with **ViewportIndex** must be declared using the input storage class.

Any variable decorated with **ViewportIndex** must be declared as a scalar 32-bit integer.

WorkgroupId

Decorating a variable with the **WorkgroupId** built-in decoration will make that variable contain the global workgroup that the current invocation is a member of. Each component ranges from zero to the values of the parameters passed into vkCmdDispatch or read from the VkDispatchIndirectCommand structure read through a call to vkCmdDispatchIndirect.

The **WorkgroupId** decoration must be used only within compute shaders.

The variable decorated with **WorkgroupId** must be declared using the input storage class.

The variable decorated with **WorkgroupId** must be declared as a three-component vector of 32-bit integers.

WorkgroupSize

Decorating a variable with the **WorkgroupSize** built-in decoration will make that variable contain the dimensions of a local workgroup. If an object is decorated with the **WorkgroupSize** decoration, this must take precedence over any execution mode set for **LocalSize**.

		t be a specialization cor t be declared as a three-	component vector of 32-bit	intege
	-			-

Chapter 15

Image Operations

15.1 Image Operations Overview

Image Operations are steps performed by SPIR-V image instructions, where those instructions which take an **OpTypeImage** (representing a VkImageView) or **OpTypeSampledImage** (representing a (VkImageView, VkSampler) pair) and texel coordinates as operands, and return a value based on one or more neighboring texture elements (*texels*) in the image.



Note

Texel is a term which is a combination of the words texture and element. Early interactive computer graphics supported texture operations on textures, a small subset of the image operations on images described here. The discrete samples remain essentially equivalent, however, so we retain the historical term texel to refer to them.

SPIR-V Image Instructions include the following functionality:

- OpImageSample* and OpImageSparseSample* read one or more neighboring texels of the image, and filter the texel values based on the state of the sampler.
 - Instructions with ImplicitLod in the name determine the level of detail used in the sampling operation based on the coordinates used in neighboring fragments.
 - Instructions with ExplicitLod in the name determine the level of detail used in the sampling operation based on additional coordinates.
 - Instructions with **Proj** in the name apply homogeneous projection to the coordinates.
- OpImageFetch and OpImageSparseFetch return a single texel of the image. No sampler is used.
- OpImage*Gather and OpImageSparse*Gather read neighboring texels and return a single component of each.
- OpImageRead (and OpImageSparseRead) and OpImageWrite read and write, respectively, a texel in the image. No sampler is used.
- Instructions with **Dref** in the name apply depth comparison on the texel values.
- Instructions with **Sparse** in the name additionally return a sparse residency code.

15.1.1 Texel Coordinate Systems

Images are addressed by texel coordinates. There are three texel coordinate systems:

- normalized texel coordinates [0.0, 1.0]
- unnormalized texel coordinates [0.0, width/height/depth]
- integer texel coordinates [0, width/height/depth)

SPIR-V OpImageFetch, OpImageSparseFetch, OpImageRead, OpImageSparseRead, and OpImageWrite instructions use integer texel coordinates. Other image instructions can use either normalized or unnormalized texel coordinates (selected by the unnormalizedCoordinates state of the sampler used in the instruction), but there are limitations on what operations, image state, and sampler state is supported. Normalized coordinates are logically converted to unnormalized as part of image operations, and certain steps are only performed on normalized coordinates. The array layer coordinate is always treated as unnormalized even when other coordinates are normalized.

Normalized texel coordinates are referred to as (s,t,r,q,a), with the coordinates having the following meanings:

- s: Coordinate in the first dimension of an image.
- t: Coordinate in the second dimension of an image.
- r: Coordinate in the third dimension of an image.
 - (s,t,r) are interpreted as a direction vector for Cube images.
- q: Fourth coordinate, for homogeneous (projective) coordinates.
- a: Coordinate for array layer.

The coordinates are extracted from the SPIR-V operand based on the dimensionality of the image variable and type of instruction. For **Proj** instructions, the components are in order (s, [t,] [r,] q) with t and r being conditionally present based on the **Dim** of the image. For non-**Proj** instructions, the coordinates are (s [,t] [,r] [,a]), with t and r being conditionally present based on the **Dim** of the image and a being conditionally present based on the **Arrayed** property of the image. Projective image instructions are not supported on **Arrayed** images.

Unnormalized texel coordinates are referred to as (u, v, w, a), with the coordinates having the following meanings:

- u: Coordinate in the first dimension of an image.
- v: Coordinate in the second dimension of an image.
- w: Coordinate in the third dimension of an image.
- a: Coordinate for array layer.

Only the u and v coordinates are directly extracted from the SPIR-V operand, because only 1D and 2D (non-**Arrayed**) dimensionalities support unnormalized coordinates. The components are in order (u [,v]), with v being conditionally present when the dimensionality is 2D. When normalized coordinates are converted to unnormalized coordinates, all four coordinates are used.

Integer texel coordinates are referred to as (i, j, k, l, n), and the first four in that order have the same meanings as unnormalized texel coordinates. They are extracted from the SPIR-V operand in order (i, [,j], [,k], [,l]), with j and k conditionally present based on the **Dim** of the image, and l conditionally present based on the **Arrayed** property of the image. n is the sample index and is taken from the **Sample** image operand.

For all coordinate types, unused coordinates are assigned a value of zero.



The Texel Coordinate Systems - For the example shown of an 8x4 texel two dimensional image.

- Normalized texel coordinates:
 - The s coordinate goes from 0.0 to 1.0, left to right.
 - The t coordinate goes from 0.0 to 1.0, top to bottom.
- Unnormalized texel coordinates:
 - The u coordinate goes from -1.0 to 9.0, left to right. The u coordinate within the range 0.0 to 8.0 is within the image, otherwise it is within the border.
 - The v coordinate goes from -1.0 to 5.0, top to bottom. The v coordinate within the range 0.0 to 4.0 is within the image, otherwise it is within the border.
- Integer texel coordinates:
 - The i coordinate goes from -1 to 8, left to right. The i coordinate within the range 0 to 7 addresses texels within the image, otherwise it addresses a border texel.
 - The j coordinate goes from -1 to 5, top to bottom. The j coordinate within the range 0 to 3 addresses texels within the image, otherwise it addresses a border texel.
- Also shown for linear filtering:
 - Given the unnormalized coordinates (u,v), the four texels selected are i0j0, i1j0, i0j1 and i1j1.
 - The weights α and β .
 - Given the offset Δ_i and Δ_j , the four texels selected by the offset are i0j0', i1j0', i0j1' and i1j1'.



The Texel Coordinate Systems - For the example shown of an 8x4 texel two dimensional image.

- Texel coordinates as above. Also shown for nearest filtering:
 - Given the unnormalized coordinates (u,v), the texel selected is ij.
 - Given the offset Δ_i and Δ_j , the texel selected by the offset is ij'.

15.2 Conversion Formulas

15.2.1 RGB to Shared Exponent Conversion

An RGB color (red, green, blue) is transformed to a shared exponent color $(red_{shared}, green_{shared}, blue_{shared}, exp_{shared})$ as follows:

First, the components (red, green, blue) are clamped to (red_{clamped}, green_{clamped}, blue_{clamped}) as:

$$red_{clamped} = \max(0, min(sharedexp_{max}, red))$$

 $green_{clamped} = \max(0, min(sharedexp_{max}, green))$
 $blue_{clamped} = \max(0, min(sharedexp_{max}, blue))$

Where:

$$N=9$$
 number of mantissa bits per component $B=15$ exponent bias $E_{max}=31$ maximum possible biased exponent value $sharedexp_{max}=\frac{(2^N-1)}{2^N}\times 2^{(E_{max}-B)}$



Note

NaN, if supported, is handled as in IEEE 754-2008 minNum() and maxNum(). That is the result is a NaN is mapped to zero.

The largest clamped component, *max_{clamped}* is determined:

$$max_{clamped} = max(red_{clamped}, green_{clamped}, blue_{clamped})$$

A preliminary shared exponent exp' is computed:

$$exp' = \begin{cases} \left\lfloor \log_2(max_{clamped}) \right\rfloor + (B+1) & \text{for } max_{clamped} > 2^{-(B+1)} \\ 0 & \text{for } max_{clamped} \leq 2^{-(B+1)} \end{cases}$$

The shared exponent exp_{shared} is computed:

$$max_{shared} = \left\lfloor \frac{max_{clamped}}{2^{(exp'-B-N)}} + \frac{1}{2} \right\rfloor$$

$$exp_{shared} = \begin{cases} exp' & \text{for } 0 \le max_{shared} < 2^N \\ exp' + 1 & \text{for } max_{shared} = 2^N \end{cases}$$

Finally, three integer values in the range 0 to 2^N are computed:

$$red_{shared} = \left \lfloor rac{red_{clamped}}{2^{(exp_{shared}-B-N)}} + rac{1}{2}
ight
floor \ green_{shared} = \left \lfloor rac{green_{clamped}}{2^{(exp_{shared}-B-N)}} + rac{1}{2}
ight
floor \ blue_{shared} = \left \lfloor rac{blue_{clamped}}{2^{(exp_{shared}-B-N)}} + rac{1}{2}
ight
floor$$

15.2.2 Shared Exponent to RGB

A shared exponent color $(red_{shared}, green_{shared}, blue_{shared}, exp_{shared})$ is transformed to an RGB color (red, green, blue) as follows:

$$red = red_{shared} \times 2^{(exp_{shared} - B - N)}$$
 $green = green_{shared} \times 2^{(exp_{shared} - B - N)}$ $blue = blue_{shared} \times 2^{(exp_{shared} - B - N)}$

Where:

$$N = 9$$
 number of mantissa bits per component $B = 15$ exponent bias

15.3 Texel Input Operations

Texel input instructions are SPIR-V image instructions that read from an image. *Texel input operations* are a set of steps that are performed on state, coordinates, and texel values while processing a texel input instruction, and which are common to some or all texel input instructions. They include the following steps, which are performed in the listed order:

- Validation operations
 - Instruction/Sampler/Image validation
 - Coordinate validation
 - Sparse validation
- Format conversion
- · Texel replacement
- Depth comparison
- · Conversion to RGBA
- Component swizzle

For texel input instructions involving multiple texels (for sampling or gathering), these steps are applied for each texel that is used in the instruction. Depending on the type of image instruction, other steps are conditionally performed between these steps or involving multiple coordinate or texel values.

15.3.1 Texel Input Validation Operations

Texel input validation operations inspect instruction/image/sampler state or coordinates, and in certain circumstances cause the texel value to be replaced or become undefined. There are a series of validations that the texel undergoes.

15.3.1.1 Instruction/Sampler/Image Validation

There are a number of cases where a SPIR-V instruction can mismatch with the sampler, the image, or both. There are a number of cases where the sampler can mismatch with the image. In such cases the value of the texel returned is undefined.

These cases include:

- The sampler borderColor is an integer type and the image format is not one of the VkFormat integer types or a stencil component of a depth/stencil format.
- The sampler borderColor is a float type and the image format is not one of the VkFormat float types or a depth component of a depth/stencil format.
- The sampler borderColor is one of the opaque black colors (VK_BORDER_COLOR_FLOAT_OPAQUE_BLACK or VK_BORDER_COLOR_INT_OPAQUE_BLACK) and the image VkComponentSwizzle for any of the VkComponentMapping components is not VK_COMPONENT_SWIZZLE_IDENTITY.
- If the instruction is OpImageRead or OpImageSparseRead and the shaderStorageImageReadWithoutFormat feature is not enabled, or the instruction is OpImageWrite and the shaderStorageImageWriteWithoutFormat feature is not enabled, then the SPIR-V Image Format must be compatible with the image view's format.

- The sampler unnormalizedCoordinates is VK_TRUE and any of the limitations of unnormalized coordinates are violated.
- The SPIR-V instruction is one of the OpImage*Dref* instructions and the sampler compareEnable is VK_FALSE
- The SPIR-V instruction is not one of the OpImage*Dref* instructions and the sampler compareEnable is VK_ TRUE
- The SPIR-V instruction is one of the **OpImage*Dref*** instructions and the image *format* is not one of the depth/stencil formats with a depth component, or the image aspect is not VK_IMAGE_ASPECT_DEPTH_BIT.
- The SPIR-V instruction's image variable's properties are not compatible with the image view:
 - Rules for viewType:

```
* VK_IMAGE_VIEW_TYPE_1D must have Dim = 1D, Arrayed = 0, MS = 0.
```

- * VK IMAGE VIEW TYPE 2D must have Dim = 2D, Arrayed = 0.
- * VK_IMAGE_VIEW_TYPE_3D must have Dim = 3D, Arrayed = 0, MS = 0.
- * VK_IMAGE_VIEW_TYPE_CUBE must have Dim = Cube, Arrayed = 0, MS = 0.
- * VK_IMAGE_VIEW_TYPE_1D_ARRAY must have Dim = 1D, Arrayed = 1, MS = 0.
- * VK_IMAGE_VIEW_TYPE_2D_ARRAY must have Dim = 2D, Arrayed = 1.
- * VK_IMAGE_VIEW_TYPE_CUBE_ARRAY must have Dim = Cube, Arrayed = 1, MS = 0.
- If the image was created with VkImageCreateInfo::samples equal to VK_SAMPLE_COUNT_1_BIT, the instruction must have MS = 0.
- If the image was created with VkImageCreateInfo::samples not equal to VK_SAMPLE_COUNT_1_BIT, the
 instruction must have MS = 1.

15.3.1.2 Integer Texel Coordinate Validation

Integer texel coordinates are validated against the size of the image level, and the number of layers and number of samples in the image. For SPIR-V instructions that use integer texel coordinates, this is performed directly on the integer coordinates. For instructions that use normalized or unnormalized texel coordinates, this is performed on the coordinates that result after conversion to integer texel coordinates.

If the integer texel coordinates satisfy any of the conditions

i < 0	$i \geq w_s$
j < 0	$j \geq h_s$
k < 0	$k \geq d_s$
l < 0	$l \ge layers$
n < 0	n > samples

where:

```
w_s= width of the image levelh_s= height of the image leveld_s= depth of the image levellayers= number of layers in the imagesamples= number of samples per texel in the image
```

then the texel fails integer texel coordinate validation.

There are four cases to consider:

1. Valid Texel Coordinates

• If the texel coordinates pass validation (that is, the coordinates lie within the image),

then the texel value comes from the value in image memory.

2. Border Texel

- If the texel coordinates fail validation, and
- If the read is the result of an image sample instruction or image gather instruction, and
- If the image is not a cube image,

then the texel is a border texel and texel replacement is performed.

3. Invalid Texel

- If the texel coordinates fail validation, and
- If the read is the result of an image fetch instruction, image read instruction, or atomic instruction,

then the texel is an invalid texel and texel replacement is performed.

4. Cube Map Edge or Corner

Otherwise the texel coordinates lie on the borders along the edges and corners of a cube map image, and Cube map edge handling is performed.

15.3.1.3 Cube Map Edge Handling

If the texel coordinates lie on the borders along the edges and corners of a cube map image, the following steps are performed. Note that this only occurs when using VK_FILTER_LINEAR filtering within a mip level, since VK_FILTER NEAREST is treated as using VK SAMPLER ADDRESS MODE CLAMP TO EDGE.

- Cube Map Edge Texel
 - If the texel lies along the border in either only i or only j

then the texel lies along an edge, so the coordinates (i, j) and the array layer l are transformed to select the adjacent texel from the appropriate neighboring face.

- Cube Map Corner Texel
 - If the texel lies along the border in both i and j

then the texel lies at a corner and there is no unique neighboring face from which to read that texel. The texel should be replaced by the average of the three values of the adjacent texels in each incident face. However, implementations may replace the cube map corner texel by other methods, subject to the constraint that if the three available samples have the same value, the replacement texel also has that value.

15.3.1.4 Sparse Validation

If the texel reads from an unbound region of a sparse image, the texel is a *sparse unbound texel*, and processing continues with texel replacement.

15.3.2 Format Conversion

Texels undergo a format conversion from the VkFormat of the image view to a vector of either floating point or signed or unsigned integer components, with the number of components based on the number of components present in the format.

- Color formats have one, two, three, or four components, according to the format.
- Depth/stencil formats are one component. The depth or stencil component is selected by the <code>aspectMask</code> of the image view.

Each component is converted based on its type and size (as defined in the Format Definition section for each VkFormat), using the appropriate equations in 16-Bit Floating-Point Numbers, Unsigned 11-Bit Floating-Point Numbers, Unsigned 10-Bit Floating-Point Numbers, Fixed-Point Data Conversion, and Shared Exponent to RGB. Signed integer components smaller than 32 bits are sign-extended.

If the image format is sRGB, the color components are first converted as if they are UNORM, and then sRGB to linear conversion is applied to the R, G, and B components as described in the "KHR_DF_TRANSFER_SRGB" section of the Khronos Data Format Specification. The A component, if present, is unchanged.

If the image view format is block-compressed, then the texel value is first decoded, then converted based on the type and number of components defined by the compressed format.

15.3.3 Texel Replacement

A texel is replaced if it is one (and only one) of:

- · a border texel,
- · an invalid texel, or
- a sparse unbound texel.

Border texels are replaced with a value based on the image format and the borderColor of the sampler. The border color is:

Sampler borderColor	Corresponding Border Color
VK_BORDER_COLOR_FLOAT_TRANSPARENT_BLACK	B = (0.0, 0.0, 0.0, 0.0)
VK_BORDER_COLOR_FLOAT_OPAQUE_BLACK	B = (0.0, 0.0, 0.0, 1.0)
VK_BORDER_COLOR_FLOAT_OPAQUE_WHITE	B = (1.0, 1.0, 1.0, 1.0)
VK_BORDER_COLOR_INT_TRANSPARENT_BLACK	B = (0,0,0,0)
VK_BORDER_COLOR_INT_OPAQUE_BLACK	B = (0,0,0,1)
VK_BORDER_COLOR_INT_OPAQUE_WHITE	B = (1, 1, 1, 1)

Table 15.1: Border Color B



Note

The names VK_BORDER_COLOR_*_TRANSPARENT_BLACK, VK_BORDER_COLOR_*_OPAQUE_BL ACK, and VK_BORDER_COLOR_*_OPAQUE_WHITE are meant to describe which components are zeros and ones in the vocabulary of compositing, and are not meant to imply that the numerical value of VK_BORDER_COLOR_INT_OPAQUE_WHITE is a saturating value for integers.

This is substituted for the texel value by replacing the number of components in the image format

Texel Aspect or FormatComponent AssignmentDepth aspect $D = (B_r)$ Stencil aspect $S = (B_r)$ One component color format $C_r = (B_r)$ Two component color format $C_{rg} = (B_r, B_g)$ Three component color format $C_{rgb} = (B_r, B_g, B_b)$ Four component color format $C_{rgba} = (B_r, B_g, B_b, B_a)$

Table 15.2: Border Texel Components After Replacement

If the read operation is from a buffer resource, and the robustBufferAccess feature is enabled, an invalid texel is replaced as described here.

If the robustBufferAccess feature is not enabled, the value of an invalid texel is undefined.

If the VkPhysicalDeviceSparseProperties property residencyNonResidentStrict is true, a sparse unbound texel is replaced with 0 or 0.0 values for integer and floating-point components of the image format, respectively.

If residencyNonResidentStrict is false, the read must be safe, but the value of the sparse unbound texel is undefined.

15.3.4 Depth Compare Operation

If the image view has a depth/stencil format, the depth component is selected by the <code>aspectMask</code>, and the operation is a <code>Dref</code> instruction, a depth comparison is performed. The value of the result *D* is 1.0 if the result of the compare operation is <code>true</code>, and 0.0 otherwise. The compare operation is selected by the <code>compareOp</code> member of the sampler.

$$D = 1.0$$

$$\begin{cases} D_{ref} \leq D & \text{for LEQUAL} \\ D_{ref} \geq D & \text{for GEQUAL} \\ D_{ref} < D & \text{for LESS} \\ D_{ref} > D & \text{for GREATER} \\ D_{ref} = D & \text{for EQUAL} \\ D_{ref} \neq D & \text{for NOTEQUAL} \\ true & \text{for ALWAYS} \\ false & \text{for NEVER} \\ \end{cases}$$

$$D = 0.0$$

where, in the depth comparison:

$$D_{ref} = shaderOp.D_{ref}$$
 (from optional SPIR-V operand)
 D texel depth value

15.3.5 Conversion to RGBA

The texel is expanded from one, two, or three to four components based on the image base color:

 Texel Aspect or Format
 RGBA Color

 Depth aspect
 $C_{rgba} = (D, 0, 0, one)$

 Stencil aspect
 $C_{rgba} = (S, 0, 0, one)$

 One component color format
 $C_{rgba} = (C_r, 0, 0, one)$

 Two component color format
 $C_{rgba} = (C_{rg}, 0, one)$

 Three component color format
 $C_{rgba} = (C_{rgb}, one)$

 Four component color format
 $C_{rgba} = C_{rgba}$

Table 15.3: Texel Color After Conversion To RGBA

where one = 1.0f for floating-point formats and depth aspects, and one = 1 for integer formats and stencil aspects.

15.3.6 Component Swizzle

All texel input instructions apply a *swizzle* based on the VkComponentSwizzle enums in the *components* member of the VkImageViewCreateInfo structure for the image being read. The swizzle can rearrange the components of the texel, or substitute zero and one for any components. It is defined as follows for the R component, and operates similarly for the other components.

$$C'_{rgba}[R] = \begin{cases} C_{rgba}[R] & \text{for RED swizzle} \\ C_{rgba}[G] & \text{for GREEN swizzle} \\ C_{rgba}[B] & \text{for BLUE swizzle} \\ C_{rgba}[A] & \text{for ALPHA swizzle} \\ 0 & \text{for ZERO swizzle} \\ one & \text{for ONE swizzle} \\ C_{rgba}[R] & \text{for IDENTITY swizzle} \end{cases}$$

where:

 $C_{rgba}[R]$ is the RED component $C_{rgba}[G]$ is the GREEN component $C_{rgba}[B]$ is the BLUE component $C_{rgba}[A]$ is the ALPHA component one = 1.0f one = 1

for floating point components for integer components

For each component this is applied to, the VK_COMPONENT_SWIZZLE_IDENTITY swizzle selects the corresponding component from C_{rgba} .

If the border color is one of the VK_BORDER_COLOR_*_OPAQUE_BLACK enums and the VkComponentSwizzle is not VK_COMPONENT_SWIZZLE_IDENTITY for all components (or the equivalent identity mapping), the value of the texel after swizzle is undefined.

15.3.7 Sparse Residency

OpImageSparse* instructions return a structure which includes a *residency code* indicating whether any texels accessed by the instruction are sparse unbound texels. This code can be interpreted by the **OpImageSparseTexelsResident** instruction which converts the residency code to a boolean value.

15.4 Texel Output Operations

Texel output instructions are SPIR-V image instructions that write to an image. Texel output operations are a set of steps that are performed on state, coordinates, and texel values while processing a texel output instruction, and which are common to some or all texel output instructions. They include the following steps, which are performed in the listed order:

- Validation operations
 - Format validation
 - Coordinate validation
 - Sparse validation
- Texel output format conversion

15.4.1 Texel Output Validation Operations

Texel output validation operations inspect instruction/image state or coordinates, and in certain circumstances cause the write to have no effect. There are a series of validations that the texel undergoes.

15.4.1.1 Texel Format Validation

If the image format of the **OpTypeImage** is not compatible with the VkImageView's format, the effect of the write on the image view's memory is undefined, but the write must not access memory outside of the image view.

15.4.2 Integer Texel Coordinate Validation

The integer texel coordinates are validated according to the same rules as for texel input coordinate validation. If the texel fails integer texel coordinate validation, then the write has no effect.

15.4.3 Sparse Texel Operation

If the texel attempts to write to an unbound region of a sparse image, the texel is a sparse unbound texel. In such a case, if the VkPhysicalDeviceSparseProperties property residencyNonResidentStrict is VK_TRUE, the sparse unbound texel write has no effect. If residencyNonResidentStrict is VK_FALSE, the effect of the write is undefined but must be safe. In addition, the write may have a side effect that is visible to other image instructions, but must not be written to any device memory allocation.

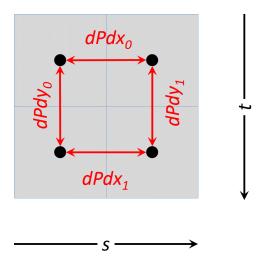
15.4.4 Texel Output Format Conversion

Texels undergo a format conversion from the floating point, signed, or unsigned integer type of the texel data to the VkFormat of the image view. Any unused components are ignored.

Each component is converted based on its type and size (as defined in the Format Definition section for each VkFormat), using the appropriate equations in 16-Bit Floating-Point Numbers and Fixed-Point Data Conversion.

15.5 Derivative Operations

SPIR-V derivative instructions include OpDPdx, OpDPdy, OpDPdxFine, OpDPdyFine, OpDPdxCoarse, and OpDPdyCoarse. Derivative instructions are only available in a fragment shader.



Derivatives are computed as if there is a 2x2 neighborhood of fragments for each fragment shader invocation. These neighboring fragments are used to compute derivatives with the assumption that the values of P in the neighborhood are piecewise linear. It is further assumed that the values of P in the neighborhood are locally continuous, therefore derivatives in non-uniform control flow are undefined.

$$\begin{aligned} dPdx_{i_1,j_0} &= dPdx_{i_0,j_0} \\ dPdx_{i_1,j_1} &= dPdx_{i_0,j_1} \end{aligned} &= P_{i_1,j_0} - P_{i_0,j_0} \\ dPdy_{i_0,j_1} &= dPdy_{i_0,j_0} \\ dPdy_{i_1,j_1} &= dPdy_{i_1,j_0} \end{aligned} &= P_{i_0,j_1} - P_{i_0,j_0} \\ dPdy_{i_1,j_1} &= dPdy_{i_1,j_0} \end{aligned}$$

The **Fine** derivative instructions must return the values above, for a group of fragments in a 2x2 neighborhood. Coarse derivatives may return only two values. In this case, the values should be:

$$dPdx = \begin{cases} dPdx_{i_0,j_0} & \text{preferred} \\ dPdx_{i_0,j_1} \end{cases}$$

$$dPdy = \begin{cases} dPdy_{i_0,j_0} & \text{preferred} \\ dPdy_{i_1,j_0} \end{cases}$$

OpDPdx and OpDPdy must return the same result as either OpDPdxFine or OpDPdxCoarse and either OpDPdyFine or OpDPdyCoarse, respectively. Implementations must make the same choice of either coarse or fine for both OpDPdx and OpDPdy, and implementations should make the choice that is more efficient to compute.

15.6 Normalized Texel Coordinate Operations

If the image sampler instruction provides normalized texel coordinates, some of the following operations are performed.

15.6.1 Projection Operation

For **Proj** image operations, the normalized texel coordinates (s,t,r,q,a) and (if present) the D_{ref} coordinate are transformed as follows:

$$s=rac{s}{q},$$
 for 1D, 2D, or 3D image $t=rac{t}{q},$ for 2D or 3D image $r=rac{r}{q},$ for 3D image $D_{ref}=rac{D_{ref}}{q},$ if provided

15.6.2 Derivative Image Operations

Derivatives are used for level-of-detail selection. These derivatives are either implicit (in an **ImplicitLod** image instruction in a fragment shader) or explicit (provided explicitly by shader to the image instruction in any shader).

For implicit derivatives image instructions, the derivatives of texel coordinates are calculated in the same manner as derivative operations above. That is:

$$\partial s/\partial x = dPdx(s),$$
 $\partial s/\partial y = dPdy(s),$ for 1D, 2D, Cube, or 3D image $\partial t/\partial x = dPdx(t),$ $\partial t/\partial y = dPdy(t),$ for 2D, Cube, or 3D image $\partial u/\partial x = dPdx(u),$ $\partial u/\partial y = dPdy(u),$ for Cube or 3D image

Partial derivatives not defined above for certain image dimensionalities are set to zero.

For explicit level-of-detail image instructions, if the optional SPIR-V operand *Grad* is provided, then the operand values are used for the derivatives. The number of components present in each derivative for a given image dimensionality matches the number of partial derivatives computed above.

If the optional SPIR-V operand Lod is provided, then derivatives are set to zero, the cube map derivative transformation is skipped, and the scale factor operation is skipped. Instead, the floating point scalar coordinate is directly assigned to λ_{base} as described in Level-of-Detail Operation.

15.6.3 Cube Map Face Selection and Transformations

For cube map image instructions, the (s,t,r) coordinates are treated as a direction vector (r_x,r_y,r_z) . The direction vector is used to select a cube map face. The direction vector is transformed to a per-face texel coordinate system (s_{face},t_{face}) . The direction vector is also used to transform the derivatives to per-face derivatives.

15.6.4 Cube Map Face Selection

The direction vector selects one of the cube map's faces based on the largest magnitude coordinate direction (the major axis direction). Since two or more coordinates can have identical magnitude, the implementation must have rules to disambiguate this situation.

The rules should have as the first rule that r_z wins over r_y and r_x , and the second rule that r_y wins over r_x . An implementation may choose other rules, but the rules must be deterministic and depend only on (r_x, r_y, r_z) .

The layer number (corresponding to a cube map face), the coordinate selections for s_c , t_c , r_c , and the selection of derivatives, are determined by the major axis direction as specified in the following two tables.

Major	Layer	Cube Map	S_C	t_c	r_c
Axis	Number	Face			
Direction					
$+r_x$	0	PositiveX	$-r_z$	$-r_y$	r_{χ}
$-r_{\chi}$	1	NegativeX	$+r_z$	$-r_y$	r_{χ}
$+r_y$	2	PositiveY	$+r_{x}$	$+r_z$	r_y
$-r_y$	3	NegativeY	$+r_{x}$	$-r_z$	r_y
$+r_z$	4	PositiveZ	$+r_{x}$	$-r_y$	r_z
$-r_z$	5	NegativeZ	$-r_{\chi}$	$-r_{v}$	r_z

Table 15.4: Cube map face and coordinate selection

Table 15.5: Cube map derivative selection

Major	$\partial s_c/\partial x$	$\partial s_c/\partial y$	$\partial t_c/\partial x$	$\partial t_c/\partial y$	$\partial r_c/\partial x$	$\partial r_c/\partial y$
Axis Di-						
rection						
$+r_{\chi}$	$-\partial r_z/\partial x$	$-\partial r_z/\partial y$	$-\partial r_y/\partial x$	$-\partial r_y/\partial y$	$+\partial r_x/\partial x$	$+\partial r_x/\partial y$
$-r_{x}$	$+\partial r_z/\partial x$	$+\partial r_z/\partial y$	$-\partial r_y/\partial x$	$-\partial r_y/\partial y$	$-\partial r_x/\partial x$	$-\partial r_x/\partial y$
$+r_y$	$+\partial r_x/\partial x$	$+\partial r_x/\partial y$	$+\partial r_z/\partial x$	$+\partial r_z/\partial y$	$+\partial r_y/\partial x$	$+\partial r_y/\partial y$
$-r_{y}$	$+\partial r_x/\partial x$	$+\partial r_x/\partial y$	$-\partial r_z/\partial x$	$-\partial r_z/\partial y$	$-\partial r_y/\partial x$	$-\partial r_{y}/\partial y$
$+r_z$	$+\partial r_x/\partial x$	$+\partial r_x/\partial y$	$-\partial r_y/\partial x$	$-\partial r_y/\partial y$	$+\partial r_z/\partial x$	$+\partial r_z/\partial y$
$-r_z$	$-\partial r_x/\partial x$	$-\partial r_x/\partial y$	$-\partial r_{\rm y}/\partial x$	$-\partial r_{\rm y}/\partial y$	$-\partial r_z/\partial x$	$-\partial r_z/\partial y$

15.6.5 Cube Map Coordinate Transformation

$$s_{face} = \frac{1}{2} \times \frac{s_c}{|r_c|} + \frac{1}{2}$$
$$t_{face} = \frac{1}{2} \times \frac{t_c}{|r_c|} + \frac{1}{2}$$

15.6.6 Cube Map Derivative Transformation

$$\frac{\partial s_{face}}{\partial x} = \frac{\partial}{\partial x} \left(\frac{1}{2} \times \frac{s_c}{|r_c|} + \frac{1}{2} \right)$$

$$\frac{\partial s_{face}}{\partial x} = \frac{1}{2} \times \frac{\partial}{\partial x} \left(\frac{s_c}{|r_c|} \right)$$

$$\frac{\partial s_{face}}{\partial x} = \frac{1}{2} \times \left(\frac{|r_c| \times \partial s_c / \partial x - s_c \times \partial r_c / \partial x}{(r_c)^2} \right)$$

$$\frac{\partial s_{face}}{\partial y} = \frac{1}{2} \times \left(\frac{|r_c| \times \partial s_c / \partial y - s_c \times \partial r_c / \partial y}{(r_c)^2} \right)$$

$$\frac{\partial t_{face}}{\partial x} = \frac{1}{2} \times \left(\frac{|r_c| \times \partial t_c / \partial x - t_c \times \partial r_c / \partial x}{(r_c)^2} \right)$$

$$\frac{\partial t_{face}}{\partial y} = \frac{1}{2} \times \left(\frac{|r_c| \times \partial t_c / \partial y - t_c \times \partial r_c / \partial y}{(r_c)^2} \right)$$

15.6.7 Scale Factor Operation, Level-of-Detail Operation and Image Level(s) Selection

Level-of-detail selection can be either explicit (provided explicitly by the image instruction) or implicit (determined from a scale factor calculated from the derivatives).

15.6.7.1 Scale Factor Operation

The magnitude of the derivatives are calculated by:

$$m_{ux} = |\partial s/\partial x| \times w_{base}$$
 $m_{vx} = |\partial t/\partial x| \times h_{base}$
 $m_{wx} = |\partial r/\partial x| \times d_{base}$
 $m_{uy} = |\partial s/\partial y| \times w_{base}$
 $m_{vy} = |\partial t/\partial y| \times h_{base}$
 $m_{wy} = |\partial r/\partial y| \times d_{base}$

where:

$$\partial t/\partial x = \partial t/\partial y = 0$$
 (for 1D image) $\partial r/\partial x = \partial r/\partial y = 0$ (for 1D, 2D or Cube image) $w_{base} = image.w$ $h_{base} = image.h$ $d_{base} = image.d$ (from image descriptor)

The *scale factors* (ρ_x, ρ_y) should be calculated by:

$$\rho_x = \sqrt{m_{ux}^2 + m_{vx}^2 + m_{wx}^2}$$

$$\rho_y = \sqrt{m_{uy}^2 + m_{vy}^2 + m_{wy}^2}$$

The ideal functions ρ_x and ρ_y may be approximated with functions f_x and f_y , subject to the following constraints:

 f_x is continuous and monotonically increasing in each of m_{ux} , m_{vx} , and m_{wx} f_y is continuous and monotonically increasing in each of m_{uy} , m_{vy} , and m_{wy}

$$\max(|m_{ux}|, |m_{vx}|, |m_{wx}|) \le f_x \le |m_{ux}| + |m_{vx}| + |m_{wx}|$$

$$\max(|m_{uy}|, |m_{vy}|, |m_{wy}|) \le f_y \le |m_{uy}| + |m_{vy}| + |m_{wy}|$$

The minimum and maximum scale factors (ρ_{min}, ρ_{max}) are determined by:

$$\rho_{max} = \max(\rho_x, \rho_y)$$

$$\rho_{min} = \min(\rho_x, \rho_y)$$

The sampling rate is determined by:

$$N = \min\left(\left\lceil\frac{\rho_{max}}{\rho_{min}}\right\rceil, max_{Aniso}\right)$$

where:

$$sampler.max_{Aniso} = maxAnisotropy$$
 (from sampler descriptor)
 $limits.max_{Aniso} = maxSamplerAnisotropy$ (from physical device limits)
 $max_{Aniso} = min(sampler.max_{Aniso}, limits.max_{Aniso})$

If $\rho_{max} = \rho_{min} = 0$, then all the partial derivatives are zero, the fragment's footprint in texel space is a point, and N should be treated as 1. If $\rho_{max} \neq 0$ and $\rho_{min} = 0$ then all partial derivatives along one axis are zero, the fragment's footprint in texel space is a line segment, and N should be treated as max_{Aniso} . However, anytime the footprint is small in texel space the implementation may use a smaller value of N, even when ρ_{min} is zero or close to zero.

An implementation may round N up to the nearest supported sampling rate.

If N = 1, sampling is isotropic. If N > 1, sampling is anisotropic.

15.6.7.2 Level-of-Detail Operation

The *level-of-detail* parameter λ is computed as follows:

where:

$$sampler.bias = mipLodBias \qquad (from sampler descriptor)$$

$$shaderOp.bias = \begin{cases} Bias & (from optional SPIR-V operand) \\ 0 & otherwise \end{cases}$$

$$sampler.lod_{min} = minLod \qquad (from optional SPIR-V operand)$$

$$shaderOp.lod_{min} = \begin{cases} MinLod & (from optional SPIR-V operand) \\ 0 & otherwise \end{cases}$$

$$lod_{min} = \max(sampler.lod_{min}, shaderOp.lod_{min})$$

(from sampler descriptor)

and maxSamplerLodBias is the value of the VkPhysicalDeviceLimits feature maxSamplerLodBias.

15.6.7.3 Image Level(s) Selection

 $lod_{max} = maxLod$

The image level(s) d, d_{hi}, and d_{lo} which texels are read from are selected based on the level-of-detail parameter, as follows. If the sampler's mipmapMode is VK_SAMPLER_MIPMAP_MODE_NEAREST, then level d is used:

$$d = \begin{cases} level_{base}, & \lambda \leq \frac{1}{2} \\ nearest(\lambda), & \lambda > \frac{1}{2}, level_{base} + \lambda \leq q + \frac{1}{2} \\ q, & \lambda > \frac{1}{2}, level_{base} + \lambda > q + \frac{1}{2} \end{cases}$$

where:

$$nearest(\lambda) = \begin{cases} \left[level_{base} + \lambda + \frac{1}{2}\right] - 1, & \text{preferred} \\ \left[level_{base} + \lambda + \frac{1}{2}\right], & \text{alternative} \end{cases}$$

and

$$q = levelCount - 1$$

levelCount is taken from the subresourceRange of the image view.

If the sampler's mipmapMode is VK_SAMPLER_MIPMAP_MODE_LINEAR, two neighboring levels are selected:

$$d_{hi} = egin{cases} q, & level_{base} + \lambda \geq q \ \lfloor level_{base} + \lambda
floor, & ext{otherwise} \ d_{lo} = egin{cases} q, & level_{base} + \lambda \geq q \ d_{hi} + 1, & ext{otherwise} \end{cases}$$

 δ is the fractional value used for linear filtering between levels.

$$\delta = \operatorname{frac}(\lambda)$$

15.6.8 (s,t,r,q,a) to (u,v,w,a) Transformation

The normalized texel coordinates are scaled by the image level dimensions and the array layer is selected. This transformation is performed once for each level (d or d_{hi} and d_{lo}) used in filtering.

$$u(x,y) = s(x,y) \times width_{level}$$

$$v(x,y) = \begin{cases} 0 & \text{for 1D images} \\ t(x,y) \times height_{level} & \text{otherwise} \end{cases}$$

$$w(x,y) = \begin{cases} 0 & \text{for 2D or Cube images} \\ r(x,y) \times depth_{level} & \text{otherwise} \end{cases}$$

$$a(x,y) = \begin{cases} a(x,y) & \text{for array images} \\ 0 & \text{otherwise} \end{cases}$$

Operations then proceed to Unnormalized Texel Coordinate Operations.

15.7 Unnormalized Texel Coordinate Operations

15.7.1 (u,v,w,a) to (i,j,k,l,n) Transformation And Array Layer Selection

The unnormalized texel coordinates are transformed to integer texel coordinates relative to the selected mipmap level.

The layer index 1 is computed as:

$$l = \text{clamp}(RNE(a), 0, layerCount - 1) + baseArrayLayer$$

where layerCount is the number of layers in the image subresource range of the image view, baseArrayLayer is the first layer from the subresource range, and where:

$$\text{RNE}(a) = \begin{cases} \text{roundTiesToEven}(a) & \text{preferred, from IEEE Std 754-2008 Floating-Point Arithmetic} \\ \left\lfloor a + \frac{1}{2} \right\rfloor & \text{alternative} \end{cases}$$

The sample index n is assigned the value zero.

Nearest filtering (VK_FILTER_NEAREST) computes the integer texel coordinates that the unnormalized coordinates lie within:

$$i = \lfloor u \rfloor$$
$$j = \lfloor v \rfloor$$
$$k = \lfloor w \rfloor$$

Linear filtering (VK_FILTER_LINEAR) computes a set of neighboring coordinates which bound the unnormalized coordinates. The integer texel coordinates are combinations of i_0 or i_1 , j_0 or j_1 , k_0 or k_1 , as well as weights α , β , and γ .

$$i_{0} = \left\lfloor u - \frac{1}{2} \right\rfloor$$

$$j_{0} = \left\lfloor v - \frac{1}{2} \right\rfloor$$

$$j_{1} = j_{0} + 1$$

$$k_{0} = \left\lfloor w - \frac{1}{2} \right\rfloor$$

$$k_{1} = k_{0} + 1$$

$$\alpha = \operatorname{frac}\left(u - \frac{1}{2}\right)$$

$$\beta = \operatorname{frac}\left(v - \frac{1}{2}\right)$$

$$\gamma = \operatorname{frac}\left(w - \frac{1}{2}\right)$$

Cubic filtering (VK_FILTER_CUBIC_IMG) computes a set of neighboring coordinates which bound the unnormalized coordinates. The integer texel coordinates are combinations of i_0, i_1, i_2 or i_3, j_0, j_1, j_2 or j_3 , as well as weights α and β .

$$i_0 = \left\lfloor u - \frac{3}{2} \right\rfloor$$

$$i_1 = i_0 + 1$$

$$i_2 = i_1 + 1$$

$$i_3 = i_2 + 1$$

$$j_0 = \left\lfloor u - \frac{3}{2} \right\rfloor$$

$$j_1 = j_0 + 1$$

$$j_2 = j_1 + 1$$

$$j_3 = j_2 + 1$$

$$\alpha = \operatorname{frac}\left(u - \frac{1}{2}\right)$$

$$\beta = \operatorname{frac}\left(v - \frac{1}{2}\right)$$

If the image instruction includes a *ConstOffset* operand, the constant offsets $(\Delta_i, \Delta_j, \Delta_k)$ are added to (i, j, k) components of the integer texel coordinates.

15.8 Image Sample Operations

15.8.1 Wrapping Operation

Cube images ignore the wrap modes specified in the sampler. Instead, if VK_FILTER_NEAREST is used within a mip level then VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE is used, and if VK_FILTER_LINEAR is used within a mip level then sampling at the edges is performed as described earlier in the Cube map edge handling section.

The first integer texel coordinate i is transformed based on the addressModeU parameter of the sampler.

$$i = \begin{cases} i \operatorname{mod} size & \text{for repeat} \\ (size-1) - \operatorname{mirror}((i \operatorname{mod}(2 \times size)) - size) & \text{for mirrored repeat} \\ \operatorname{clamp}(i, 0, size-1) & \text{for clamp to edge} \\ \operatorname{clamp}(i, -1, size) & \text{for clamp to border} \\ \operatorname{clamp}(\operatorname{mirror}(i), 0, size-1) & \text{for mirror clamp to edge} \end{cases}$$

where:

$$mirror(n) = \begin{cases} n & \text{for } n \ge 0\\ -(1+n) & \text{otherwise} \end{cases}$$

j (for 2D and Cube image) and k (for 3D image) are similarly transformed based on the addressModeV and addressModeW parameters of the sampler, respectively.

15.8.2 Texel Gathering

SPIR-V instructions with **Gather** in the name return a vector derived from a 2x2 rectangular region of texels in the base level of the image view. The rules for the VK_FILTER_LINEAR minification filter are applied to identify the four selected texels. Each texel is then converted to an RGBA value according to conversion to RGBA and then swizzled. A four-component vector is then assembled by taking the component indicated by the **Component** value in the instruction from the swizzled color value of the four texels:

$$au[R] = au_{i0j1}[level_{base}][comp] \ au[G] = au_{i1j1}[level_{base}][comp] \ au[B] = au_{i1j0}[level_{base}][comp] \ au[A] = au_{i0j0}[level_{base}][comp]$$

where:

$$\tau[level_{base}][comp] = \begin{cases} \tau[level_{base}][R], & \text{for } comp = 0 \\ \tau[level_{base}][G], & \text{for } comp = 1 \\ \tau[level_{base}][B], & \text{for } comp = 2 \\ \tau[level_{base}][A], & \text{for } comp = 3 \end{cases}$$

comp from SPIR-V operand Component

15.8.3 Texel Filtering

If λ is less than or equal to zero, the texture is said to be *magnified*, and the filter mode within a mip level is selected by the *magFilter* in the sampler. If λ is greater than zero, the texture is said to be *minified*, and the filter mode within a mip level is selected by the *minFilter* in the sampler.

Within a mip level, VK_FILTER_NEAREST filtering selects a single value using the (i, j, k) texel coordinates, with all texels taken from layer 1.

$$\tau[level] = \begin{cases} \tau_{ijk}[level], & \text{for 3D image} \\ \tau_{ij}[level], & \text{for 2D or Cube image} \\ \tau_{i}[level], & \text{for 1D image} \end{cases}$$

Within a mip level, VK_FILTER_LINEAR filtering computes a weighted average of 8 (for 3D), 4 (for 2D or Cube), or 2

(for 1D) texel values, using the weights computed earlier:

$$\tau_{3D}[level] = (1-\alpha)(1-\beta)(1-\gamma)\tau_{i0j0k0}[level] \\ + (\alpha)(1-\beta)(1-\gamma)\tau_{i1j0k0}[level] \\ + (1-\alpha)(\beta)(1-\gamma)\tau_{i0j1k0}[level] \\ + (\alpha)(\beta)(1-\gamma)\tau_{i0j1k0}[level] \\ + (\alpha)(\beta)(1-\gamma)\tau_{i1j1k0}[level] \\ + (1-\alpha)(1-\beta)(\gamma)\tau_{i0j0k1}[level] \\ + (\alpha)(1-\beta)(\gamma)\tau_{i0j1k1}[level] \\ + (1-\alpha)(\beta)(\gamma)\tau_{i0j1k1}[level] \\ + (\alpha)(\beta)(\gamma)\tau_{i1j1k1}[level]$$

$$\tau_{2D}[level] = (1-\alpha)(1-\beta)\tau_{i0j0}[level] \\ + (\alpha)(1-\beta)\tau_{i1j0}[level] \\ + (\alpha)(\beta)\tau_{i1j1}[level]$$

$$\tau_{1D}[level] = (1-\alpha)\tau_{i0}[level] \\ + (\alpha)\tau_{i1}[level]$$

$$\tau_{1D}[level], \text{ for 3D image} \\ \tau_{2D}[level], \text{ for 2D or Cube image} \\ \tau_{1D}[level], \text{ for 1D image}$$

Within a mip level, VK_FILTER_CUBIC_IMG filtering computes a weighted average of 16 (for 2D), or 4 (for 1D) texel values, using the weights computed during texel selection.

Catmull-Rom Spine interpolation of four points is defined by the equation:

$$cinterp(\tau_0, \tau_1, \tau_2, \tau_3, \omega) = \frac{1}{2} \begin{bmatrix} 1 & \omega & \omega^2 & \omega^3 \end{bmatrix} \cdot \begin{bmatrix} 0 & 2 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ 2 & -5 & 4 & 1 \\ -1 & 3 & -3 & 1 \end{bmatrix} \cdot \begin{bmatrix} \tau_0 \\ \tau_1 \\ \tau_2 \\ \tau_3 \end{bmatrix}$$

Using the values calculated in texel selection, this equation is applied to the four points in 1D images. For 2D images, the this equation is evaluated first for each row, and the result is then fed back into the equation and interpolated again:

$$\begin{split} \tau_{1D}[level] &= cinterp(\tau_{i0}[level], \tau_{i1}[level], \tau_{i2}[level], \tau_{i3}[level], \alpha) \\ \tau_{j0}[level] &= cinterp(\tau_{i0j0}[level], \tau_{i1j0}[level], \tau_{i2j0}[level], \tau_{i3j0}[level], \alpha) \\ \tau_{j1}[level] &= cinterp(\tau_{i0j1}[level], \tau_{i1j1}[level], \tau_{i2j1}[level], \tau_{i3j1}[level], \alpha) \\ \tau_{j2}[level] &= cinterp(\tau_{i0j2}[level], \tau_{i1j2}[level], \tau_{i2j2}[level], \tau_{i3j2}[level], \alpha) \\ \tau_{j3}[level] &= cinterp(\tau_{i0j3}[level], \tau_{i1j3}[level], \tau_{i2j3}[level], \tau_{i3j3}[level], \alpha) \\ \tau_{2D}[level] &= cinterp(\tau_{j0}[level], \tau_{j1}[level], \tau_{j2}[level], \tau_{j3}[level], \beta) \\ \tau_{[D}[level], \quad \text{for 2D image} \\ \tau_{1D}[level], \quad \text{for 1D image} \end{split}$$

Finally, mipmap filtering either selects a value from one mip level or computes a weighted average between neighboring mip levels:

$$au = egin{cases} au[d], & ext{for mip mode BASE or NEAREST} \ (1-\delta) au[d_{hi}] + \delta au[d_{lo}], & ext{for mip mode LINEAR} \end{cases}$$

15.8.4 Texel Anisotropic Filtering

Anisotropic filtering is enabled by the anisotropyEnable in the sampler. When enabled, the image filtering scheme accounts for a degree of anisotropy.

The particular scheme for anisotropic texture filtering is implementation dependent. Implementations should consider the magFilter, minFilter and mipmapMode of the sampler to control the specifics of the anisotropic filtering scheme used. In addition, implementations should consider minLod and maxLod of the sampler.

The following describes one particular approach to implementing anisotropic filtering for the 2D Image case, implementations may choose other methods:

Given a magFilter, minFilter of VK_FILTER_LINEAR and a mipmapMode of VK_SAMPLER_MIPMAP_MODE_NEAREST:

Instead of a single isotropic sample, N isotropic samples are be sampled within the image footprint of the image level d to approximate an anisotropic filter. The sum $\tau_{2Daniso}$ is defined using the single isotropic $\tau_{2D}(u,v)$ at level d.

$$\tau_{2Daniso} = \frac{1}{N} \sum_{i=1}^{N} \tau_{2D} \left(u \left(x - \frac{1}{2} + \frac{i}{N+1}, y \right), \left(v \left(x - \frac{1}{2} + \frac{i}{N+1} \right), y \right) \right), \quad \text{when } \rho_x > \rho_y$$

$$\tau_{2Daniso} = \frac{1}{N} \sum_{i=1}^{N} \tau_{2D} \left(u \left(x, y - \frac{1}{2} + \frac{i}{N+1} \right), \left(v \left(x, y - \frac{1}{2} + \frac{i}{N+1} \right) \right) \right), \quad \text{when } \rho_{y} \ge \rho_{y}$$

15.9 Image Operation Steps

Each step described in this chapter is performed by a subset of the image instructions:

- Texel Input Validation Operations, Format Conversion, Texel Replacement, Conversion to RGBA, and Component Swizzle: Performed by all instructions except OpImageWrite.
- Depth Comparison: Performed by OpImage*Dref instructions.
- All Texel output operations: Performed by **OpImageWrite**.
- Projection: Performed by all OpImage*Proj instructions.
- Derivative Image Operations, Cube Map Operations, Scale Factor Operation, Level-of-Detail Operation and Image Level(s) Selection, and Texel Anisotropic Filtering: Performed by all OpImageSample* and OpImageSparseSample* instructions.
- (s,t,r,q,a) to (u,v,w,a) Transformation, Wrapping, and (u,v,w,a) to (i,j,k,l,n) Transformation And Array Layer Selection: Performed by all OpImageSample, OpImageSparseSample, and OpImage*Gather instructions.
- Texel Gathering: Performed by **OpImage*Gather** instructions.
- Texel Filtering: Performed by all OpImageSample* and OpImageSparseSample* instructions.
- Sparse Residency: Performed by all OpImageSparse* instructions.

Chapter 16

Queries

Queries provide a mechanism to return information about the processing of a sequence of Vulkan commands. Query operations are asynchronous, and as such, their results are not returned immediately. Instead, their results, and their availability status, are stored in a Query Pool. The state of these queries can be read back on the host, or copied to a buffer object on the device.

The supported query types are Occlusion Queries, Pipeline Statistics Queries, and Timestamp Queries.

16.1 Query Pools

Queries are managed using *query pool* objects. Each query pool is a collection of a specific number of queries of a particular type.

Query pools are represented by VkQueryPool handles:

```
VK_DEFINE_NON_DISPATCHABLE_HANDLE(VkQueryPool)
```

To create a query pool, call:

- device is the logical device that creates the query pool.
- pCreateInfo is a pointer to an instance of the VkQueryPoolCreateInfo structure containing the number and type of queries to be managed by the pool.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.
- pQueryPool is a pointer to a VkQueryPool handle in which the resulting query pool object is returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkQueryPoolCreateInfo structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pQueryPool must be a pointer to a VkQueryPool handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkQueryPoolCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- *flags* is reserved for future use.
- queryType is the type of queries managed by the pool, and must be one of the values

```
typedef enum VkQueryType {
    VK_QUERY_TYPE_OCCLUSION = 0,
    VK_QUERY_TYPE_PIPELINE_STATISTICS = 1,
    VK_QUERY_TYPE_TIMESTAMP = 2,
} VkQueryType;
```

• queryCount is the number of queries managed by the pool.

• pipelineStatistics is a bitmask indicating which counters will be returned in queries on the new pool, as described below in Section 16.4. pipelineStatistics is ignored if queryType is not VK_QUERY_TYPE_PIPELINE_STATISTICS.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_QUERY_POOL_CREATE_INFO
- pNext must be NULL
- flags must be 0
- queryType must be a valid VkQueryType value
- If the pipeline statistics queries feature is not enabled, queryType must not be VK_QUERY_TYPE_PIPELINE_ STATISTICS
- If queryType is VK_QUERY_TYPE_PIPELINE_STATISTICS, pipelineStatistics must be a valid combination of VkQueryPipelineStatisticFlagBits values

To destroy a query pool, call:

- device is the logical device that destroys the query pool.
- *queryPool* is the query pool to destroy.
- pAllocator controls host memory allocation as described in the Memory Allocation chapter.

- device must be a valid VkDevice handle
- If queryPool is not VK_NULL_HANDLE, queryPool must be a valid VkQueryPool handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If queryPool is a valid handle, it must have been created, allocated, or retrieved from device
- All submitted commands that refer to queryPool must have completed execution
- If VkAllocationCallbacks were provided when *queryPool* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when queryPool was created, pAllocator must be NULL

• Host access to queryPool must be externally synchronized

16.2 Query Operation

The operation of queries is controlled by the commands vkCmdBeginQuery, vkCmdEndQuery, vkCmdResetQueryPool, vkCmdCopyQueryPoolResults, and vkCmdWriteTimestamp.

In order for a VkCommandBuffer to record query management commands, the queue family for which its VkCommandPool was created must support the appropriate type of operations (graphics, compute) suitable for the query type of a given query pool.

Each query in a query pool has a status that is either *unavailable* or *available*, and also has state to store the numerical results of a query operation of the type requested when the query pool was created. Resetting a query via vkCmdResetQueryPool sets the status to unavailable and makes the numerical results undefined. Performing a query operation with vkCmdBeginQuery and vkCmdEndQuery changes the status to available when the query finishes, and updates the numerical results. Both the availability status and numerical results are retrieved by calling either vkGetQueryPoolResults or vkCmdCopyQueryPoolResults.

All query commands execute in order and are guaranteed to see the effects of each other's memory accesses, with one significant exception: **vkCmdCopyQueryPoolResults** may execute before the results of **vkCmdEndQuery** are available. However, if VK_QUERY_RESULT_WAIT_BIT is used, then **vkCmdCopyQueryPoolResults** must reflect the result of any previously executed queries. Other sequences of commands, such as **vkCmdResetQueryPool** followed by **vkCmdBeginQuery**, must make the effects of the first command visible to the second command.

After query pool creation, each query is in an undefined state and must be reset prior to use. Queries must also be reset between uses. Using a query that has not been reset will result in undefined behavior.

To reset a range of queries in a query pool, call:

- commandBuffer is the command buffer into which this command will be recorded.
- queryPool is the handle of the query pool managing the queries being reset.
- firstQuery is the initial query index to reset.
- queryCount is the number of queries to reset.

When executed on a queue, this command sets the status of query indices firstQuery, firstQuery + queryCount - 1 to unavailable.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- queryPool must be a valid VkQueryPool handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance
- Both of commandBuffer, and queryPool must have been created, allocated, or retrieved from the same VkDevice
- firstQuery must be less than the number of queries in queryPool
- The sum of firstQuery and queryCount must be less than or equal to the number of queries in queryPool

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		COMPUTE

Once queries are reset and ready for use, query commands can be issued to a command buffer. Occlusion queries and pipeline statistics queries count events - drawn samples and pipeline stage invocations, respectively - resulting from commands that are recorded between a vkCmdBeginQuery command and a vkCmdEndQuery command within a specified command buffer, effectively scoping a set of drawing and/or compute commands. Timestamp queries write timestamps to a query pool.

A query must begin and end in the same command buffer, although if it is a primary command buffer, and the inherited queries feature is enabled, it can execute secondary command buffers during the query operation. For a secondary command buffer to be executed while a query is active, it must set the <code>occlusionQueryEnable</code>, <code>queryFlags</code>, and/or <code>pipelineStatistics</code> members of <code>VkCommandBufferInheritanceInfo</code> to conservative values, as described in the Command Buffer Recording section. A query must either begin and end inside the same subpass of a render pass instance, or must both begin and end outside of a render pass instance (i.e. contain entire render pass instances).

To begin a query, call:

- commandBuffer is the command buffer into which this command will be recorded.
- queryPool is the query pool that will manage the results of the query.
- query is the query index within the query pool that will contain the results.
- flags is a bitmask indicating constraints on the types of queries that can be performed. Bits which can be set include:

```
typedef enum VkQueryControlFlagBits {
    VK_QUERY_CONTROL_PRECISE_BIT = 0x00000001,
} VkQueryControlFlagBits;
```

If the <code>queryType</code> of the pool is <code>VK_QUERY_TYPE_OCCLUSION</code> and <code>flags</code> contains <code>VK_QUERY_CONTROL_PRECISE_BIT</code>, an implementation must return a result that matches the actual number of samples passed. This is described in more detail in Occlusion Oueries.

After beginning a query, that query is considered *active* within the command buffer it was called in until that same query is ended. Queries active in a primary command buffer when secondary command buffers are executed are considered active for those secondary command buffers.

- commandBuffer must be a valid VkCommandBuffer handle
- queryPool must be a valid VkQueryPool handle
- flags must be a valid combination of VkQueryControlFlagBits values
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- Both of commandBuffer, and queryPool must have been created, allocated, or retrieved from the same VkDevice
- The query identified by queryPool and query must currently not be active
- The query identified by queryPool and query must be unavailable
- If the precise occlusion queries feature is not enabled, or the <code>queryType</code> used to create <code>queryPool</code> was not <code>VK_QUERY_Type_OCCLUSION</code>, <code>flags</code> must not contain <code>VK_QUERY_CONTROL_PRECISE_BIT</code>
- queryPool must have been created with a queryType that differs from that of any other queries that have been made active, and are currently still active within commandBuffer
- query must be less than the number of queries in queryPool

- If the *queryType* used to create *queryPool* was VK_QUERY_TYPE_OCCLUSION, the VkCommandPool that *commandBuffer* was allocated from must support graphics operations
- If the <code>queryType</code> used to create <code>queryPool</code> was <code>VK_QUERY_TYPE_PIPELINE_STATISTICS</code> and any of the <code>pipelineStatistics</code> indicate graphics operations, the <code>VkCommandPool</code> that <code>commandBuffer</code> was allocated from must support graphics operations
- If the <code>queryType</code> used to create <code>queryPool</code> was <code>VK_QUERY_TYPE_PIPELINE_STATISTICS</code> and any of the <code>pipelineStatistics</code> indicate compute operations, the <code>VkCommandPool</code> that <code>commandBuffer</code> was allocated from must support compute operations

Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

To end a query after the set of desired draw or dispatch commands is executed, call:

- commandBuffer is the command buffer into which this command will be recorded.
- queryPool is the query pool that is managing the results of the query.
- query is the query index within the query pool where the result is stored.

As queries operate asynchronously, ending a query does not immediately set the query's status to available. A query is considered *finished* when the final results of the query are ready to be retrieved by vkGetQueryPoolResults and vkCmdCopyQueryPoolResults, and this is when the query's status is set to available.

Once a query is ended the query must finish in finite time, unless the state of the query is changed using other commands, e.g. by issuing a reset of the query.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- queryPool must be a valid VkQueryPool handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- Both of commandBuffer, and queryPool must have been created, allocated, or retrieved from the same VkDevice
- The query identified by queryPool and query must currently be active
- query must be less than the number of queries in queryPool

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

An application can retrieve results either by requesting they be written into application-provided memory, or by requesting they be copied into a VkBuffer. In either case, the layout in memory is defined as follows:

- The first query's result is written starting at the first byte requested by the command, and each subsequent query's result begins stride bytes later.
- Each query's result is a tightly packed array of unsigned integers, either 32- or 64-bits as requested by the command, storing the numerical results and, if requested, the availability status.
- If VK_QUERY_RESULT_WITH_AVAILABILITY_BIT is used, the final element of each query's result is an integer indicating whether the query's result is available, with any non-zero value indicating that it is available.

- Occlusion queries write one integer value the number of samples passed. Pipeline statistics queries write one integer value for each bit that is enabled in the <code>pipelineStatistics</code> when the pool is created, and the statistics values are written in bit order starting from the least significant bit. Timestamps write one integer value.
- If more than one query is retrieved and stride is not at least as large as the size of the array of integers corresponding to a single query, the values written to memory are undefined.

To retrieve status and results for a set of queries, call:

```
VkResult vkGetQueryPoolResults(
    VkDevice
                                                   device,
    VkQueryPool
                                                   queryPool,
    uint32_t
                                                   firstQuery,
    uint32_t
                                                   queryCount,
                                                   dataSize,
    size_t
    void*
                                                   pData,
    VkDeviceSize
                                                   stride,
    VkQueryResultFlags
                                                   flags);
```

- device is the logical device that owns the query pool.
- queryPool is the query pool managing the queries containing the desired results.
- firstQuery is the initial query index.
- queryCount is the number of queries. firstQuery and queryCount together define a range of queries.
- dataSize is the size in bytes of the buffer pointed to by pData.
- pData is a pointer to a user-allocated buffer where the results will be written
- stride is the stride in bytes between results for individual queries within pData.
- flags is a bitmask of VkQueryResultFlagBits specifying how and when results are returned. Bits which can be set include:

```
typedef enum VkQueryResultFlagBits {
    VK_QUERY_RESULT_64_BIT = 0x00000001,
    VK_QUERY_RESULT_WAIT_BIT = 0x00000002,
    VK_QUERY_RESULT_WITH_AVAILABILITY_BIT = 0x00000004,
    VK_QUERY_RESULT_PARTIAL_BIT = 0x00000008,
} VkQueryResultFlagBits;
```

- VK_QUERY_RESULT_64_BIT indicates the results will be written as an array of 64-bit unsigned integer values. If this bit is not set, the results will be written as an array of 32-bit unsigned integer values.
- VK_QUERY_RESULT_WAIT_BIT indicates that Vulkan will wait for each query's status to become available before retrieving its results.
- VK_QUERY_RESULT_WITH_AVAILABILITY_BIT indicates that the availability status accompanies the results.
- VK_QUERY_RESULT_PARTIAL_BIT indicates that returning partial results is acceptable.

If no bits are set in *flags*, and all requested queries are in the available state, results are written as an array of 32-bit unsigned integer values. The behavior when not all queries are available, is described below.

If $VK_QUERY_RESULT_64_BIT$ is not set and the result overflows a 32-bit value, the value may either wrap or saturate. Similarly, if $VK_QUERY_RESULT_64_BIT$ is set and the result overflows a 64-bit value, the value may either wrap or saturate.

If VK_QUERY_RESULT_WAIT_BIT is set, Vulkan will wait for each query to be in the available state before retrieving the numerical results for that query. In this case, **vkGetQueryPoolResults** is guaranteed to succeed and return VK_SUCCESS if the queries become available in a finite time (i.e. if they have been issued and not reset). If queries will never finish (e.g. due to being reset but not issued), then **vkGetQueryPoolResults** may not return in finite time.

If $VK_QUERY_RESULT_WAIT_BIT$ and $VK_QUERY_RESULT_PARTIAL_BIT$ are both not set then no result values are written to pData for queries that are in the unavailable state at the time of the call, and

vkGetQueryPoolResults returns VK_NOT_READY. However, availability state is still written to pData for those queries if VK_QUERY_RESULT_WITH_AVAILABILITY_BIT is set.

Note

Applications must take care to ensure that use of the VK_QUERY_RESULT_WAIT_BIT bit has the desired effect.



For example, if a query has been used previously and a command buffer records the commands **vkCmdRes etQueryPool**, **vkCmdBeginQuery**, and **vkCmdEndQuery** for that query, then the query will remain in the available state until the **vkCmdResetQueryPool** command executes on a queue. Applications can use fences or events to ensure that a query has already been reset before checking for its results or availability status. Otherwise, a stale value could be returned from a previous use of the query.

The above also applies when VK_QUERY_RESULT_WAIT_BIT is used in combination with VK_QUERY_RESULT_WITH_AVAILABILITY_BIT. In this case, the returned availability status may reflect the result of a previous use of the query unless the **vkCmdResetQueryPool** command has been executed since the last use of the query.



Note

Applications can double-buffer query pool usage, with a pool per frame, and reset queries at the end of the frame in which they are read.

If VK_QUERY_RESULT_PARTIAL_BIT is set, VK_QUERY_RESULT_WAIT_BIT is not set, and the query's status is unavailable, an intermediate result value between zero and the final result value is written to pData for that query.

VK_QUERY_RESULT_PARTIAL_BIT must not be used if the pool's queryType is VK_QUERY_TYPE_TIMESTAMP.

If VK_QUERY_RESULT_WITH_AVAILABILITY_BIT is set, the final integer value written for each query is non-zero if the query's status was available or zero if the status was unavailable. When VK_QUERY_RESULT_WITH_ AVAILABILITY_BIT is used, implementations must guarantee that if they return a non-zero availability value then the numerical results must be valid, assuming the results are not reset by a subsequent command.



Note

Satisfying this guarantee may require careful ordering by the application, e.g. to read the availability status before reading the results.

- device must be a valid VkDevice handle
- queryPool must be a valid VkQueryPool handle
- pData must be a pointer to an array of dataSize bytes
- flags must be a valid combination of VkQueryResultFlagBits values
- dataSize must be greater than 0
- queryPool must have been created, allocated, or retrieved from device
- firstQuery must be less than the number of queries in queryPool
- If VK_QUERY_RESULT_64_BIT is not set in flags then pData and stride must be multiples of 4
- If VK_QUERY_RESULT_64_BIT is set in flags then pData and stride must be multiples of 8
- The sum of firstQuery and queryCount must be less than or equal to the number of queries in queryPool
- dataSize must be large enough to contain the result of each query, as described here
- If the <code>queryType</code> used to create <code>queryPool</code> was <code>VK_QUERY_TYPE_TIMESTAMP</code>, <code>flags</code> must not contain <code>VK_QUERY_RESULT_PARTIAL_BIT</code>

Return Codes

Success

- VK SUCCESS
- VK NOT READY

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

To copy query statuses and numerical results directly to buffer memory, call:

```
void vkCmdCopyQueryPoolResults(
   VkCommandBuffer
                                                  commandBuffer,
   VkQueryPool
                                                  queryPool,
   uint32_t
                                                  firstQuery,
   uint32_t
                                                  queryCount,
    VkBuffer
                                                  dstBuffer,
   VkDeviceSize
                                                  dstOffset,
                                                  stride,
   VkDeviceSize
   VkQueryResultFlags
                                                  flags);
```

- commandBuffer is the command buffer into which this command will be recorded.
- queryPool is the query pool managing the queries containing the desired results.
- firstQuery is the initial query index.
- queryCount is the number of queries. firstQuery and queryCount together define a range of queries.
- dstBuffer is a VkBuffer object that will receive the results of the copy command.
- dstOffset is an offset into dstBuffer.
- stride is the stride in bytes between results for individual queries within dstBuffer. The required size of the backing memory for dstBuffer is determined as described above for vkGetQueryPoolResults.
- flags is a bitmask of VkQueryResultFlagBits specifying how and when results are returned.

vkCmdCopyQueryPoolResults is guaranteed to see the effect of previous uses of **vkCmdResetQueryPool** in the same queue, without any additional synchronization. Thus, the results will always reflect the most recent use of the query.

flags has the same possible values described above for the flags parameter of vkGetQueryPoolResults, but the different style of execution causes some subtle behavioral differences. Because **vkCmdCopyQueryPoolResults** executes in order with respect to other query commands, there is less ambiguity about which use of a query is being requested.

If no bits are set in *flags*, results for all requested queries in the available state are written as 32-bit unsigned integer values, and nothing is written for queries in the unavailable state.

If VK_QUERY_RESULT_64_BIT is set, the results are written as an array of 64-bit unsigned integer values as described for vkGetQueryPoolResults.

If VK_QUERY_RESULT_WAIT_BIT is set, the implementation will wait for each query's status to be in the available state before retrieving the numerical results for that query. This is guaranteed to reflect the most recent use of the query on the same queue, assuming that the query is not being simultaneously used by other queues. If the query does not become available in a finite amount of time (e.g. due to not issuing a query since the last reset), a VK_ERROR_DEVICE_LOST error may occur.

Similarly, if VK_QUERY_RESULT_WITH_AVAILABILITY_BIT is set and VK_QUERY_RESULT_WAIT_BIT is not set, the availability is guaranteed to reflect the most recent use of the query on the same queue, assuming that the query is not being simultaneously used by other queues. As with **vkGetQueryPoolResults**, implementations must guarantee that if they return a non-zero availability value, then the numerical results are valid.

If VK_QUERY_RESULT_PARTIAL_BIT is set, VK_QUERY_RESULT_WAIT_BIT is not set, and the query's status is unavailable, an intermediate result value between zero and the final result value is written for that query.

VK_QUERY_RESULT_PARTIAL_BIT must not be used if the pool's queryType is VK_QUERY_TYPE_TIMESTAMP.

vkCmdCopyQueryPoolResults is considered to be a transfer operation, and its writes to buffer memory must be synchronized using VK_PIPELINE_STAGE_TRANSFER_BIT and VK_ACCESS_TRANSFER_WRITE_BIT before using the results.

Valid Usage

• commandBuffer must be a valid VkCommandBuffer handle

- queryPool must be a valid VkQueryPool handle
- dstBuffer must be a valid VkBuffer handle
- flags must be a valid combination of VkQueryResultFlagBits values
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance
- Each of commandBuffer, dstBuffer, and queryPool must have been created, allocated, or retrieved from the same VkDevice
- dstOffset must be less than the size of dstBuffer
- firstQuery must be less than the number of queries in queryPool
- The sum of firstQuery and queryCount must be less than or equal to the number of queries in queryPool
- If VK_QUERY_RESULT_64_BIT is not set in flags then dstOffset and stride must be multiples of 4
- If VK_QUERY_RESULT_64_BIT is set in flags then dstOffset and stride must be multiples of 8
- dstBuffer must have enough storage, from dstOffset, to contain the result of each query, as described here
- dstBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_DST_BIT usage flag
- If the queryType used to create queryPool was VK_QUERY_TYPE_TIMESTAMP, flags must not contain VK_QUERY_RESULT_PARTIAL_BIT

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		COMPUTE

Rendering operations such as clears, MSAA resolves, attachment load/store operations, and blits may count towards the results of queries. This behavior is implementation-dependent and may vary depending on the path used within an

implementation. For example, some implementations have several types of clears, some of which may include vertices and some not.

16.3 Occlusion Queries

Occlusion queries track the number of samples that pass the per-fragment tests for a set of drawing commands. As such, occlusion queries are only available on queue families supporting graphics operations. The application can then use these results to inform future rendering decisions. An occlusion query is begun and ended by calling **vkCmdBeginQuery** and **vkCmdEndQuery**, respectively. When an occlusion query begins, the count of passing samples always starts at zero. For each drawing command, the count is incremented as described in Sample Counting. If flags does not contain VK_QUERY_CONTROL_PRECISE_BIT an implementation may generate any non-zero result value for the query if the count of passing samples is non-zero.

stc.

Note

Not setting VK_QUERY_CONTROL_PRECISE_BIT mode may be more efficient on some implementations, and should be used where it is sufficient to know a boolean result on whether any samples passed the perfragment tests. In this case, some implementations may only return zero or one, indifferent to the actual number of samples passing the per-fragment tests.

When an occlusion query finishes, the result for that query is marked as available. The application can then either copy the result to a buffer (via **vkCmdCopyQueryPoolResults**) or request it be put into host memory (via **vkGetQueryPoolResults**).



Note

If occluding geometry is not drawn first, samples can pass the depth test, but still not be visible in a final image.

16.4 Pipeline Statistics Queries

Pipeline statistics queries allow the application to sample a specified set of VkPipeline counters. These counters are accumulated by Vulkan for a set of either draw or dispatch commands while a pipeline statistics query is active. As such, pipeline statistics queries are available on queue families supporting either graphics or compute operations. Further, the availability of pipeline statistics queries is indicated by the <code>pipelineStatisticsQuery</code> member of the VkPhysicalDeviceFeatures object (see vkGetPhysicalDeviceFeatures and vkCreateDevice for detecting and requesting this query type on a VkDevice).

A pipeline statistics query is begun and ended by calling **vkCmdBeginQuery** and **vkCmdEndQuery**, respectively. When a pipeline statistics query begins, all statistics counters are set to zero. While the query is active, the pipeline type determines which set of statistics are available, but these must be configured on the query pool when it is created. If a statistic counter is issued on a command buffer that does not support the corresponding operation, that counter is undefined after the query has finished. At least one statistic counter relevant to the operations supported on the recording command buffer must be enabled.

The pipeline statistic counters are individually enabled for query pools with VkQueryPoolCreateInfo::pipelineStatistics, and for secondary command buffers with VkCommandBufferInheritanceInfo::pipelineStatistics.

Bits which can be set in pipelineStatistics include:

These bits have the following meanings:

- If VK_QUERY_PIPELINE_STATISTIC_INPUT_ASSEMBLY_VERTICES_BIT is set, queries managed by the pool will count the number of vertices processed by the input assembly stage. Vertices corresponding to incomplete primitives may contribute to the count.
- If VK_QUERY_PIPELINE_STATISTIC_INPUT_ASSEMBLY_PRIMITIVES_BIT is set, queries managed by the pool will count the number of primitives processed by the input assembly stage. If primitive restart is enabled, restarting the primitive topology has no effect on the count. Incomplete primitives may be counted.
- If VK_QUERY_PIPELINE_STATISTIC_VERTEX_SHADER_INVOCATIONS_BIT is set, queries managed by the pool will count the number of vertex shader invocations. This counter's value is incremented each time a vertex shader is invoked.
- If VK_QUERY_PIPELINE_STATISTIC_GEOMETRY_SHADER_INVOCATIONS_BIT is set, queries managed by the pool will count the number of geometry shader invocations. This counter's value is incremented each time a geometry shader is invoked. In the case of instanced geometry shaders, the geometry shader invocations count is incremented for each separate instanced invocation.
- If VK_QUERY_PIPELINE_STATISTIC_GEOMETRY_SHADER_PRIMITIVES_BIT is set, queries managed by the pool will count the number of primitives generated by geometry shader invocations. The counter's value is incremented each time the geometry shader emits a primitive. Restarting primitive topology using the SPIR-V instructions OpEndPrimitive or OpEndStreamPrimitive has no effect on the geometry shader output primitives count.
- If VK_QUERY_PIPELINE_STATISTIC_CLIPPING_INVOCATIONS_BIT is set, queries managed by the pool will count the number of primitives processed by the Primitive Clipping stage of the pipeline. The counter's value is incremented each time a primitive reaches the primitive clipping stage.
- If VK_QUERY_PIPELINE_STATISTIC_CLIPPING_PRIMITIVES_BIT is set, queries managed by the pool will count the number of primitives output by the Primitive Clipping stage of the pipeline. The counter's value is incremented each time a primitive passes the primitive clipping stage. The actual number of primitives output by the primitive clipping stage for a particular input primitive is implementation-dependent but must satisfy the following conditions:
 - If at least one vertex of the input primitive lies inside the clipping volume, the counter is incremented by one or more.
 - Otherwise, the counter is incremented by zero or more.

- If VK_QUERY_PIPELINE_STATISTIC_FRAGMENT_SHADER_INVOCATIONS_BIT is set, queries managed by the pool will count the number of fragment shader invocations. The counter's value is incremented each time the fragment shader is invoked.
- If VK_QUERY_PIPELINE_STATISTIC_TESSELLATION_CONTROL_SHADER_PATCHES_BIT is set, queries managed by the pool will count the number of patches processed by the tessellation control shader. The counter's value is incremented once for each patch for which a tessellation control shader is invoked.
- If VK_QUERY_PIPELINE_STATISTIC_TESSELLATION_EVALUATION_SHADER_INVOCATIONS_BIT is set, queries managed by the pool will count the number of invocations of the tessellation evaluation shader. The counter's value is incremented each time the tessellation evaluation shader is invoked.
- If VK_QUERY_PIPELINE_STATISTIC_COMPUTE_SHADER_INVOCATIONS_BIT is set, queries managed by the pool will count the number of compute shader invocations. The counter's value is incremented every time the compute shader is invoked. Implementations may skip the execution of certain compute shader invocations or execute additional compute shader invocations for implementation-dependent reasons as long as the results of rendering otherwise remain unchanged.

These values are intended to measure relative statistics on one implementation. Various device architectures will count these values differently. Any or all counters may be affected by the issues described in Query Operation.



Note

For example, tile-based rendering devices may need to replay the scene multiple times, affecting some of the counts.

If a pipeline has rasterizerDiscardEnable enabled, implementations may discard primitives after the final vertex processing stage. As a result, if rasterizerDiscardEnable is enabled, the clipping input and output primitives counters may not be incremented.

When a pipeline statistics query finishes, the result for that query is marked as available. The application can copy the result to a buffer (via **vkCmdCopyQueryPoolResults**), or request it be put into host memory (via **vkGetQueryPoolResults**).

16.5 Timestamp Queries

Timestamps provide applications with a mechanism for timing the execution of commands. A timestamp is an integer value generated by the VkPhysicalDevice. Unlike other queries, timestamps do not operate over a range, and so do not use vkCmdBeginQuery or vkCmdEndQuery. The mechanism is built around a set of commands that allow the application to tell the VkPhysicalDevice to write timestamp values to a query pool and then either read timestamp values on the host (using vkGetQueryPoolResults) or copy timestamp values to a VkBuffer (using vkCmdCopyQueryPoolResults). The application can then compute differences between timestamps to determine execution time.

The number of valid bits in a timestamp value is determined by the

VkQueueFamilyProperties::timestampValidBits property of the queue on which the timestamp is written. Timestamps are supported on any queue which reports a non-zero value for timestampValidBits via vkGetPhysicalDeviceQueueFamilyProperties. If the timestampComputeAndGraphics limit is VK_TRUE, timestamps are supported by every queue family that supports either graphics or compute operations (see VkQueueFamilyProperties).

The number of nanoseconds it takes for a timestamp value to be incremented by 1 can be obtained from VkPhysicalDeviceLimits::timestampPeriod after a call to vkGetPhysicalDeviceProperties.

To request a timestamp, call:

- commandBuffer is the command buffer into which the command will be recorded.
- pipelineStage is one of the VkPipelineStageFlagBits, specifying a stage of the pipeline.
- *queryPool* is the query pool that will manage the timestamp.
- query is the query within the query pool that will contain the timestamp.

vkCmdWriteTimestamp latches the value of the timer when all previous commands have completed executing as far as the specified pipeline stage, and writes the timestamp value to memory. When the timestamp value is written, the availability status of the query is set to available.



Note

If an implementation is unable to detect completion and latch the timer at any specific stage of the pipeline, it may instead do so at any logically later stage.

vkCmdCopyQueryPoolResults can then be called to copy the timestamp value from the query pool into buffer memory, with ordering and synchronization behavior equivalent to how other queries operate. Timestamp values can also be retrieved from the query pool using vkGetQueryPoolResults. As with other queries, the query must be reset using vkCmdResetQueryPool before requesting the timestamp value be written to it.

While **vkCmdWriteTimestamp** can be called inside or outside of a render pass instance, vkCmdCopyQueryPoolResults must only be called outside of a render pass instance.

- commandBuffer must be a valid VkCommandBuffer handle
- pipelineStage must be a valid VkPipelineStageFlagBits value
- queryPool must be a valid VkQueryPool handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- Both of commandBuffer, and queryPool must have been created, allocated, or retrieved from the same VkDevice
- The query identified by queryPool and query must be unavailable
- The command pool's queue family must support a non-zero timestampValidBits

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

Chapter 17

Clear Commands

17.1 Clearing Images Outside A Render Pass Instance

Color and depth/stencil images can be cleared outside a render pass instance using vkCmdClearColorImage or vkCmdClearDepthStencilImage, respectively. These commands are only allowed outside of a render pass instance.

To clear one or more subranges of a color image, call:

- commandBuffer is the command buffer into which the command will be recorded.
- *image* is the image to be cleared.
- *imageLayout* specifies the current layout of the image subresource ranges to be cleared, and must be VK_IMAGE_LAYOUT_GENERAL or VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL.
- pColor is a pointer to a VkClearColorValue structure that contains the values the image subresource ranges will be cleared to (see Section 17.3 below).
- rangeCount is the number of image subresource range structures in pRanges.
- pRanges points to an array of VkImageSubresourceRange structures that describe a range of mipmap levels, array layers, and aspects to be cleared, as described in Image Views. The aspectMask of all image subresource ranges must only include VK_IMAGE_ASPECT_COLOR_BIT.

Each specified range in pRanges is cleared to the value specified by pColor.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- image must be a valid VkImage handle
- imageLayout must be a valid VkImageLayout value
- pColor must be a pointer to a valid VkClearColorValue union
- pRanges must be a pointer to an array of rangeCount valid VkImageSubresourceRange structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance
- rangeCount must be greater than 0
- Both of commandBuffer, and image must have been created, allocated, or retrieved from the same VkDevice
- image must have been created with VK_IMAGE_USAGE_TRANSFER_DST_BIT usage flag
- imageLayout must specify the layout of the image subresource ranges of image specified in pRanges at the time this command is executed on a VkDevice
- imageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- The image range of any given element of *pRanges* must be an image subresource range that is contained within *image*
- image must not have a compressed or depth/stencil format

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		COMPUTE

To clear one or more subranges of a depth/stencil image, call:

- commandBuffer is the command buffer into which the command will be recorded.
- image is the image to be cleared.
- *imageLayout* specifies the current layout of the image subresource ranges to be cleared, and must be VK_IMAGE_LAYOUT_GENERAL or VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL.
- pDepthStencil is a pointer to a VkClearDepthStencilValue structure that contains the values the depth and stencil image subresource ranges will be cleared to (see Section 17.3 below).
- rangeCount is the number of image subresource range structures in pRanges.
- pRanges points to an array of VkImageSubresourceRange structures that describe a range of mipmap levels, array layers, and aspects to be cleared, as described in Image Views. The <code>aspectMask</code> of each image subresource range in <code>pRanges</code> can include VK_IMAGE_ASPECT_DEPTH_BIT if the image format has a depth component, and VK_IMAGE_ASPECT_STENCIL_BIT if the image format has a stencil component. <code>pDepthStencil</code> is a pointer to a VkClearDepthStencilValue structure that contains the values the image subresource ranges will be cleared to (see Section 17.3 below).

- commandBuffer must be a valid VkCommandBuffer handle
- image must be a valid VkImage handle
- imageLayout must be a valid VkImageLayout value
- pDepthStencil must be a pointer to a valid VkClearDepthStencilValue structure
- pRanges must be a pointer to an array of rangeCount valid VkImageSubresourceRange structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called outside of a render pass instance
- rangeCount must be greater than 0

- Both of commandBuffer, and image must have been created, allocated, or retrieved from the same VkDevice
- image must have been created with VK_IMAGE_USAGE_TRANSFER_DST_BIT usage flag
- imageLayout must specify the layout of the image subresource ranges of image specified in pRanges at the time this command is executed on a VkDevice
- imageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- The image range of any given element of pRanges must be an image subresource range that is contained within image
- image must have a depth/stencil format

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		

Clears outside render pass instances are treated as transfer operations for the purposes of memory barriers.

17.2 Clearing Images Inside A Render Pass Instance

To clear one or more regions of color and depth/stencil attachments inside a render pass instance, call:

• commandBuffer is the command buffer into which the command will be recorded.

- attachmentCount is the number of entries in the pAttachments array.
- pAttachments is a pointer to an array of VkClearAttachment structures defining the attachments to clear and the clear values to use.
- rectCount is the number of entries in the pRects array.
- pRects points to an array of VkClearRect structures defining regions within each selected attachment to clear.

vkCmdClearAttachments can clear multiple regions of each attachment used in the current subpass of a render pass instance. This command must be called only inside a render pass instance, and implicitly selects the images to clear based on the current framebuffer attachments and the command parameters.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pAttachments must be a pointer to an array of attachment Count valid VkClearAttachment structures
- pRects must be a pointer to an array of rectCount VkClearRect structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- attachmentCount must be greater than 0
- rectCount must be greater than 0
- If the aspectMask member of any given element of pAttachments contains VK_IMAGE_ASPECT_COLOR_BIT, the colorAttachment member of those elements must refer to a valid color attachment in the current subpass
- The rectangular region specified by a given element of pRects must be contained within the render area of the current render pass instance
- The layers specified by a given element of pRects must be contained within every attachment that pAttachments refers to

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

The VkClearRect structure is defined as:

```
typedef struct VkClearRect {
    VkRect2D    rect;
    uint32_t    baseArrayLayer;
    uint32_t    layerCount;
} VkClearRect;
```

- rect is the two-dimensional region to be cleared.
- baseArrayLayer is the first layer to be cleared.
- layerCount is the number of layers to clear.

The layers [baseArrayLayer, baseArrayLayer + layerCount) counting from the base layer of the attachment image view are cleared.

The VkClearAttachment structure is defined as:

- aspectMask is a mask selecting the color, depth and/or stencil aspects of the attachment to be cleared. aspectMask can include VK_IMAGE_ASPECT_COLOR_BIT for color attachments, VK_IMAGE_ASPECT_DEPTH_BIT for depth/stencil attachments with a depth component, and VK_IMAGE_ASPECT_STENCIL_BIT for depth/stencil attachments with a stencil component. If the subpass's depth/stencil attachment is VK_ATTACHMENT_UNUSED, then the clear has no effect.
- colorAttachment is only meaningful if VK_IMAGE_ASPECT_COLOR_BIT is set in aspectMask, in which case it is an index to the pColorAttachments array in the VkSubpassDescription structure of the current subpass which selects the color attachment to clear. If colorAttachment is VK_ATTACHMENT_UNUSED or is greater than or equal to VkSubpassDescription::colorAttachmentCount, then the clear has no effect.
- clearValue is the color or depth/stencil value to clear the attachment to, as described in Clear Values below.

No memory barriers are needed between **vkCmdClearAttachments** and preceding or subsequent draw or attachment clear commands in the same subpass.

The **vkCmdClearAttachments** command is not affected by the bound pipeline state.

Attachments can also be cleared at the beginning of a render pass instance by setting <code>loadOp</code> (or <code>stencilLoadOp</code>) of <code>VkAttachmentDescription</code> to <code>VK_ATTACHMENT_LOAD_OP_CLEAR</code>, as described for <code>vkCreateRenderPass</code>.

Valid Usage

- aspectMask must be a valid combination of VkImageAspectFlagBits values
- aspectMask must not be 0
- clearValue must be a valid VkClearValue union
- If aspectMask includes VK_IMAGE_ASPECT_COLOR_BIT, it must not include VK_IMAGE_ASPECT_DEPTH_BIT or VK_IMAGE_ASPECT_STENCIL_BIT
- aspectMask must not include VK_IMAGE_ASPECT_METADATA_BIT

17.3 Clear Values

The VkClearColorValue structure is defined as:

```
typedef union VkClearColorValue {
   float     float32[4];
   int32_t     int32[4];
   uint32_t     uint32[4];
} VkClearColorValue;
```

- float 32 are the color clear values when the format of the image or attachment is floating point, unorm, snorm, uscaled, packed float, or sRGB. Floating point values are automatically converted to the format of the image, with the clear value being treated as linear if the image is sRGB.
- int 32 are the color clear values when the format of the image or attachment is signed integer. Signed integer values are converted to the format of the image by casting to the smaller type (with negative 32-bit values mapping to negative values in the smaller type). If the integer clear value is not representable in the target type (e.g. would overflow in conversion to that type), the clear value is undefined.
- uint 32 are the color clear values when the format of the image or attachment is unsigned integer. Unsigned integer values are converted to the format of the image by casting to the integer type with fewer bits.

The four array elements of the clear color map to R, G, B, and A components of image formats, in order.

If the image has more than one sample, the same value is written to all samples for any pixels being cleared.

The VkClearDepthStencilValue structure is defined as:

```
typedef struct VkClearDepthStencilValue {
   float depth;
   uint32_t stencil;
} VkClearDepthStencilValue;
```

• depth is the clear value for the depth aspect of the depth/stencil attachment. It is a floating-point value which is automatically converted to the attachment's format.

• stencil is the clear value for the stencil aspect of the depth/stencil attachment. It is a 32-bit integer value which is converted to the attachment's format by taking the appropriate number of LSBs.

Valid Usage

• depth must be between 0.0 and 1.0, inclusive

The VkClearValue union is defined as:

- color specifies the color image clear values to use when clearing a color image or attachment.
- depthStencil specifies the depth and stencil clear values to use when clearing a depth/stencil image or attachment.

This union is used where part of the API requires either color or depth/stencil clear values, depending on the attachment, and defines the initial clear values in the VkRenderPassBeginInfo structure.

Valid Usage

• depthStencil must be a valid VkClearDepthStencilValue structure

17.4 Filling Buffers

To clear buffer data, call:

- commandBuffer is the command buffer into which the command will be recorded.
- *dstBuffer* is the buffer to be filled.

- dstOffset is the byte offset into the buffer at which to start filling, and must be a multiple of 4.
- size is the number of bytes to fill, and must be either a multiple of 4, or VK_WHOLE_SIZE to fill the range from offset to the end of the buffer. If VK_WHOLE_SIZE is used and the remaining size of the buffer is not a multiple of 4, then the nearest smaller multiple is used.
- data is the 4-byte word written repeatedly to the buffer to fill size bytes of data. The data word is written to memory according to the host endianness.

vkCmdFillBuffer is treated as "transfer" operation for the purposes of synchronization barriers. The VK_BUFFER_ USAGE_TRANSFER_DST_BIT must be specified in *usage* of VkBufferCreateInfo in order for the buffer to be compatible with **vkCmdFillBuffer**.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- dstBuffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- This command must only be called outside of a render pass instance
- Both of commandBuffer, and dstBuffer must have been created, allocated, or retrieved from the same VkDevice
- dstOffset must be less than the size of dstBuffer
- dstOffset must be a multiple of 4
- If size is not equal to VK_WHOLE_SIZE, size must be greater than 0
- If size is not equal to VK_WHOLE_SIZE, size must be less than or equal to the size of dstBuffer minus dstOffset
- If size is not equal to VK WHOLE SIZE, size must be a multiple of 4
- dstBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_DST_BIT usage flag

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		COMPUTE

17.5 Updating Buffers

To update buffer data inline in a command buffer, call:

- commandBuffer is the command buffer into which the command will be recorded.
- *dstBuffer* is a handle to the buffer to be updated.
- dstOffset is the byte offset into the buffer to start updating, and must be a multiple of 4.
- dataSize is the number of bytes to update, and must be a multiple of 4.
- pData is a pointer to the source data for the buffer update, and must be at least dataSize bytes in size.

dataSize must be less than or equal to 65536 bytes. For larger updates, applications can use buffer to buffer copies.

The source data is copied from the user pointer to the command buffer when the command is called.

vkCmdUpdateBuffer is only allowed outside of a render pass. This command is treated as "transfer" operation, for the purposes of synchronization barriers. The VK_BUFFER_USAGE_TRANSFER_DST_BIT must be specified in usage of VkBufferCreateInfo in order for the buffer to be compatible with **vkCmdUpdateBuffer**.

- commandBuffer must be a valid VkCommandBuffer handle
- dstBuffer must be a valid VkBuffer handle
- pData must be a pointer to an array of dataSize bytes
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support transfer, graphics, or compute operations

- This command must only be called outside of a render pass instance
- dataSize must be greater than 0
- Both of commandBuffer, and dstBuffer must have been created, allocated, or retrieved from the same VkDevice
- dstOffset must be less than the size of dstBuffer
- dataSize must be less than or equal to the size of dstBuffer minus dstOffset
- dstBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_DST_BIT usage flag
- dstOffset must be a multiple of 4
- dataSize must be less than or equal to 65536
- dataSize must be a multiple of 4

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	TRANSFER
Secondary		GRAPHICS
		COMPUTE



Note

The pData parameter was of type uint32_t* instead of void* prior to revision 1.0.19 of the Specification and VK_HEADER_VERSION 19 of vulkan.h. This was a historical anomaly, as the source data may be of other types.

Chapter 18

Copy Commands

An application can copy buffer and image data using several methods depending on the type of data transfer. Data can be copied between buffer objects with **vkCmdCopyBuffer** and a portion of an image can be copied to another image with **vkCmdCopyImage**. Image data can also be copied to and from buffer memory using **vkCmdCopyImageToBuffer** and **vkCmdCopyBufferToImage**. Image data can be blitted (with or without scaling and filtering) with **vkCmdBlitImage**. Multisampled images can be resolved to a non-multisampled image with **vkCmdResolveImage**.

18.1 Common Operation

Some rules for valid operation are common to all copy commands:

- Copy commands must be recorded outside of a render pass instance.
- For non-sparse resources, the union of the source regions in a given buffer or image must not overlap the union of the destination regions in the same buffer or image.
- For sparse resources, the set of bytes used by all the source regions must not intersect the set of bytes used by all the destination regions.
- Copy regions must be non-empty.
- Regions must not extend outside the bounds of the buffer or image level, except that regions of compressed images can extend as far as the dimension of the image level rounded up to a complete compressed texel block.
- Source image subresources must be in either the VK_IMAGE_LAYOUT_GENERAL or VK_IMAGE_LAYOUT_
 TRANSFER_SRC_OPTIMAL layout. Destination image subresources must be in either the VK_IMAGE_LAYOUT_
 GENERAL or VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL layout. As a consequence, if an image subresource is used as both source and destination of a copy, it must be in the VK_IMAGE_LAYOUT_GENERAL layout.
- Source images must have been created with the VK_IMAGE_USAGE_TRANSFER_SRC_BIT usage bit enabled and destination images must have been created with the VK_IMAGE_USAGE_TRANSFER_DST_BIT usage bit enabled.
- Source buffers must have been created with the VK_BUFFER_USAGE_TRANSFER_SRC_BIT usage bit enabled and destination buffers must have been created with the VK_BUFFER_USAGE_TRANSFER_DST_BIT usage bit enabled.

All copy commands are treated as "transfer" operations for the purposes of synchronization barriers.

18.2 Copying Data Between Buffers

To copy data between buffer objects, call:

- commandBuffer is the command buffer into which the command will be recorded.
- srcBuffer is the source buffer.
- dstBuffer is the destination buffer.
- regionCount is the number of regions to copy.
- pRegions is a pointer to an array of VkBufferCopy structures specifying the regions to copy.

Each region in pRegions is copied from the source buffer to the same region of the destination buffer. srcBuffer and dstBuffer can be the same buffer or alias the same memory, but the result is undefined if the copy regions overlap in memory.

- commandBuffer must be a valid VkCommandBuffer handle
- srcBuffer must be a valid VkBuffer handle
- dstBuffer must be a valid VkBuffer handle
- pRegions must be a pointer to an array of regionCount VkBufferCopy structures
- commandBuffer must be in the recording state
- The VkCommandPool that *commandBuffer* was allocated from must support transfer, graphics, or compute operations
- This command must only be called outside of a render pass instance
- regionCount must be greater than 0
- Each of commandBuffer, dstBuffer, and srcBuffer must have been created, allocated, or retrieved from the same VkDevice
- The size member of a given element of pRegions must be greater than 0
- The srcOffset member of a given element of pRegions must be less than the size of srcBuffer
- The dstOffset member of a given element of pRegions must be less than the size of dstBuffer

- The size member of a given element of pRegions must be less than or equal to the size of srcBuffer minus srcOffset
- The size member of a given element of pRegions must be less than or equal to the size of dstBuffer minus dstOffset.
- The union of the source regions, and the union of the destination regions, specified by the elements of pRegions, must not overlap in memory
- $\hbox{\color{red} \bullet $\it srcBuffer must have been created with $\tt VK_BUFFER_USAGE_TRANSFER_SRC_BIT$ usage flag}$
- dstBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_DST_BIT usage flag

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	TRANSFER
Secondary		GRAPHICS
		COMPUTE

The VkBufferCopy structure is defined as:

```
typedef struct VkBufferCopy {
   VkDeviceSize    srcOffset;
   VkDeviceSize    dstOffset;
   VkDeviceSize    size;
} VkBufferCopy;
```

- *srcOffset* is the starting offset in bytes from the start of *srcBuffer*.
- dstOffset is the starting offset in bytes from the start of dstBuffer.
- size is the number of bytes to copy.

18.3 Copying Data Between Images

vkCmdCopyImage performs image copies in a similar manner to a host memcpy. It does not perform general-purpose conversions such as scaling, resizing, blending, color-space conversion, or format conversions. Rather, it simply copies raw image data. **vkCmdCopyImage** can copy between images with different formats, provided the formats are compatible as defined below.

To copy data between image objects, call:

- commandBuffer is the command buffer into which the command will be recorded.
- srcImage is the source image.
- srcImageLayout is the current layout of the source image subresource.
- dstImage is the destination image.
- dstImageLayout is the current layout of the destination image subresource.
- regionCount is the number of regions to copy.
- pRegions is a pointer to an array of VkImageCopy structures specifying the regions to copy.

Each region in pRegions is copied from the source image to the same region of the destination image. srcImage and dstImage can be the same image or alias the same memory.

Copies are done layer by layer starting with baseArrayLayer member of srcSubresource for the source and dstSubresource for the destination. layerCount layers are copied to the destination image.

The formats of <code>srcImage</code> and <code>dstImage</code> must be compatible. Formats are considered compatible if their texel size in bytes is the same between both formats. For example, <code>VK_FORMAT_R8G8B8A8_UNORM</code> is compatible with <code>VK_FORMAT_R32_UINT</code> because both texels are 4 bytes in size. Depth/stencil formats must match exactly.

vkCmdCopyImage allows copying between size-compatible compressed and uncompressed internal formats. Formats are size-compatible if the texel size of the uncompressed format is equal to the compressed texel block size in bytes of the compressed format. Such a copy does not perform on-the-fly compression or decompression. When copying from an uncompressed format to a compressed format, each texel of uncompressed data of the source image is copied as a raw value to the corresponding compressed texel block of the destination image. When copying from a compressed format to an uncompressed format, each compressed texel block of the source image is copied as a raw value to the corresponding texel of uncompressed data in the destination image. Thus, for example, it is legal to copy between a 128-bit uncompressed format and a compressed format which has a 128-bit sized compressed texel block representing 4x4 texels (using 8 bits per texel), or between a 64-bit uncompressed format and a compressed format which has a 64-bit sized compressed texel block representing 4x4 texels (using 4 bits per texel).

When copying between compressed and uncompressed formats the extent members represent the texel dimensions of the source image and not the destination. When copying from a compressed image to an uncompressed image the image texel dimensions written to the uncompressed image will be source extent divided by the compressed texel block

dimensions. When copying from an uncompressed image to a compressed image the image texel dimensions written to the compressed image will be the source extent multiplied by the compressed texel block dimensions. In both cases the number of bytes read and the number of bytes written will be identical.

Copying to or from block-compressed images is typically done in multiples of the compressed texel block. For this reason the <code>extent</code> must be a multiple of the compressed texel block dimension. There is one exception to this rule which is required to handle compressed images created with dimensions that are not a multiple of the compressed texel block dimensions. If the <code>srcImage</code> is compressed and if <code>extent.width</code> is not a multiple of the compressed texel block width then (<code>extent.width</code>

<code>srcOffset.x</code>) must equal the image subresource width, if <code>extent.height</code> is not a multiple of the compressed texel block height then (<code>extent.height+srcOffset.y</code>) must equal the image subresource height and if <code>extent.depth</code> is not a multiple of the compressed texel block depth then (<code>extent.depth+srcOffset.z</code>) must equal the image subresource depth. Similarly, if the <code>dstImage</code> is compressed and if <code>extent.width</code> is not a multiple of the compressed texel block width then (<code>extent.width+dstOffset.x</code>) must equal the image subresource width, if <code>extent.height</code> is not a multiple of the compressed texel block height then (<code>extent.height+dstOffset.y</code>) must equal the image subresource height and if <code>extent.depth</code> is not a multiple of the compressed texel block depth then (<code>extent.depth dstOffset.z</code>) must equal the image subresource depth. This allows the last compressed texel block of the image in each non-multiple dimension to be included as a source or destination of the copy.

vkCmdCopyImage can be used to copy image data between multisample images, but both images must have the same number of samples.

- commandBuffer must be a valid VkCommandBuffer handle
- srcImage must be a valid VkImage handle
- srcImageLayout must be a valid VkImageLayout value
- dstImage must be a valid VkImage handle
- dstImageLayout must be a valid VkImageLayout value
- pRegions must be a pointer to an array of regionCount valid VkImageCopy structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support transfer, graphics, or compute operations
- This command must only be called outside of a render pass instance
- regionCount must be greater than 0
- Each of commandBuffer, dstImage, and srcImage must have been created, allocated, or retrieved from the same VkDevice
- The source region specified by a given element of pRegions must be a region that is contained within srcImage
- The destination region specified by a given element of pRegions must be a region that is contained within dst Image

- The union of all source regions, and the union of all destination regions, specified by the elements of pRegions, must not overlap in memory
- srcImage must have been created with VK_IMAGE_USAGE_TRANSFER_SRC_BIT usage flag
- srcImageLayout must specify the layout of the image subresources of srcImage specified in pRegions at the
 time this command is executed on a VkDevice
- srcImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- dstImage must have been created with VK_IMAGE_USAGE_TRANSFER_DST_BIT usage flag
- dstImageLayout must specify the layout of the image subresources of dstImage specified in pRegions at the time this command is executed on a VkDevice
- dstImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- The VkFormat of each of srcImage and dstImage must be compatible, as defined below
- The sample count of srcImage and dstImage must match

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	TRANSFER
Secondary		GRAPHICS
		COMPUTE

The VkImageCopy structure is defined as:

- srcSubresource and dstSubresource are VkImageSubresourceLayers structures specifying the image subresources of the images used for the source and destination image data, respectively.
- srcOffset and dstOffset select the initial x, y, and z offsets in texels of the sub-regions of the source and destination image data.
- extent is the size in texels of the source image to copy in width, height and depth. 1D images use only x and width. 2D images use x, y, width and height. 3D images use x, y, z, width, height and depth.

- srcSubresource must be a valid VkImageSubresourceLayers structure
- dstSubresource must be a valid VkImageSubresourceLayers structure
- The aspectMask member of srcSubresource and dstSubresource must match
- The layerCount member of srcSubresource and dstSubresource must match
- If either of the calling command's <code>srcImage</code> or <code>dstImage</code> parameters are of <code>VkImageTypeVK_IMAGE_TYPE_3D</code>, the <code>baseArrayLayer</code> and <code>layerCount</code> members of both <code>srcSubresource</code> and <code>dstSubresource</code> must be 0 and 1, respectively
- The aspectMask member of srcSubresource must specify aspects present in the calling command's srcImage
- The aspectMask member of dstSubresource must specify aspects present in the calling command's dstImage
- srcOffset.x and (extent.width + srcOffset.x) must both be greater than or equal to 0 and less than or equal to the source image subresource width
- srcOffset.y and (extent.height + srcOffset.y) must both be greater than or equal to 0 and less than or equal to the source image subresource height
- srcOffset.z and (extent.depth+srcOffset.z) must both be greater than or equal to 0 and less than or equal to the source image subresource depth
- dstOffset.x and (extent.width + dstOffset.x) must both be greater than or equal to 0 and less than or equal to the destination image subresource width
- dstOffset.y and (extent.height + dstOffset.y) must both be greater than or equal to 0 and less than or equal to the destination image subresource height
- dstOffset.z and (extent.depth+dstOffset.z) must both be greater than or equal to 0 and less than or equal to the destination image subresource depth
- If the calling command's *srcImage* is a compressed format image:
 - all members of srcoffset must be a multiple of the corresponding dimensions of the compressed texel block
 - extent.width must be a multiple of the compressed texel block width or (extent.width + srcOffset.x)
 must equal the source image subresource width

- extent.height must be a multiple of the compressed texel block height or (extent.height + srcOffset.
 y) must equal the source image subresource height
- extent.depth must be a multiple of the compressed texel block depth or (extent.depth + srcOffset.z) must equal the source image subresource depth
- If the calling command's dst Image is a compressed format image:
 - all members of dstOffset must be a multiple of the corresponding dimensions of the compressed texel block
 - extent.width must be a multiple of the compressed texel block width or (extent.width + dstOffset.x)
 must equal the destination image subresource width
 - extent.height must be a multiple of the compressed texel block height or (extent.height + dstOffset.
 y) must equal the destination image subresource height
 - extent.depth must be a multiple of the compressed texel block depth or (extent.depth + dstOffset.z)
 must equal the destination image subresource depth
- srcOffset, dstOffset, and extent must respect the image transfer granularity requirements of the queue family that it will be submitted against, as described in Physical Device Enumeration

The VkImageSubresourceLayers structure is defined as:

- aspectMask is a combination of VkImageAspectFlagBits, selecting the color, depth and/or stencil aspects to be copied.
- mipLevel is the mipmap level to copy from.
- baseArrayLayer and layerCount are the starting layer and number of layers to copy.

- $\bullet \ \textit{aspectMask must be a valid combination of VkImageAspectFlagBits \ \textit{values}}$
- aspectMask must not be 0
- If aspectMask contains VK_IMAGE_ASPECT_COLOR_BIT, it must not contain either of VK_IMAGE_ASPECT_DEPTH_BIT or VK_IMAGE_ASPECT_STENCIL_BIT
- aspectMask must not contain VK_IMAGE_ASPECT_METADATA_BIT
- mipLevel must be less than the mipLevels specified in VkImageCreateInfo when the image was created

• (baseArrayLayer+layerCount) must be less than or equal to the arrayLayers specified in VkImageCreateInfo when the image was created

18.4 Copying Data Between Buffers and Images

To copy data from a buffer object to an image object, call:

- commandBuffer is the command buffer into which the command will be recorded.
- srcBuffer is the source buffer.
- dstImage is the destination image.
- dstImageLayout is the layout of the destination image subresources for the copy.
- regionCount is the number of regions to copy.
- pRegions is a pointer to an array of VkBufferImageCopy structures specifying the regions to copy.

Each region in *pRegions* is copied from the specified region of the source buffer to the specified region of the destination image.

- commandBuffer must be a valid VkCommandBuffer handle
- srcBuffer must be a valid VkBuffer handle
- dstImage must be a valid VkImage handle
- dstImageLayout must be a valid VkImageLayout value
- pRegions must be a pointer to an array of regionCount valid VkBufferImageCopy structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support transfer, graphics, or compute operations
- This command must only be called outside of a render pass instance

- regionCount must be greater than 0
- Each of commandBuffer, dstImage, and srcBuffer must have been created, allocated, or retrieved from the same VkDevice
- The buffer region specified by a given element of pRegions must be a region that is contained within srcBuffer
- The image region specified by a given element of pRegions must be a region that is contained within dstImage
- The union of all source regions, and the union of all destination regions, specified by the elements of pRegions, must not overlap in memory
- srcBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_SRC_BIT usage flag
- dstImage must have been created with VK_IMAGE_USAGE_TRANSFER_DST_BIT usage flag
- dstImage must have a sample count equal to VK_SAMPLE_COUNT_1_BIT
- dstImageLayout must specify the layout of the image subresources of dstImage specified in pRegions at the time this command is executed on a VkDevice
- dstImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	TRANSFER
Secondary		GRAPHICS
		COMPUTE

To copy data from an image object to a buffer object, call:

void vkCmdCopyImageToBuffer(
 VkCommandBuffer
 VkImage
 VkImageLayout
 VkBuffer

commandBuffer, srcImage, srcImageLayout, dstBuffer,

- commandBuffer is the command buffer into which the command will be recorded.
- srcImage is the source image.
- srcImageLayout is the layout of the source image subresources for the copy.
- dstBuffer is the destination buffer.
- regionCount is the number of regions to copy.
- pRegions is a pointer to an array of VkBufferImageCopy structures specifying the regions to copy.

Each region in pRegions is copied from the specified region of the source image to the specified region of the destination buffer.

- commandBuffer must be a valid VkCommandBuffer handle
- srcImage must be a valid VkImage handle
- srcImageLayout must be a valid VkImageLayout value
- dstBuffer must be a valid VkBuffer handle
- pRegions must be a pointer to an array of regionCount valid VkBufferImageCopy structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support transfer, graphics, or compute operations
- This command must only be called outside of a render pass instance
- regionCount must be greater than 0
- Each of commandBuffer, dstBuffer, and srcImage must have been created, allocated, or retrieved from the same VkDevice
- The image region specified by a given element of pRegions must be a region that is contained within srcImage
- The buffer region specified by a given element of pRegions must be a region that is contained within dstBuffer
- The union of all source regions, and the union of all destination regions, specified by the elements of pRegions, must not overlap in memory
- $\hbox{\color{red} \bullet $\it srcImage must have been created with $\tt VK_IMAGE_USAGE_TRANSFER_SRC_BIT$ usage flag}$
- srcImage must have a sample count equal to VK_SAMPLE_COUNT_1_BIT

- srcImageLayout must specify the layout of the image subresources of srcImage specified in pRegions at the time this command is executed on a VkDevice
- srcImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- dstBuffer must have been created with VK_BUFFER_USAGE_TRANSFER_DST_BIT usage flag

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	TRANSFER
Secondary		GRAPHICS
		COMPUTE

For both vkCmdCopyBufferToImage and vkCmdCopyImageToBuffer, each element of pRegions is a structure defined as:

- bufferOffset is the offset in bytes from the start of the buffer object where the image data is copied from or to.
- bufferRowLength and bufferImageHeight specify the data in buffer memory as a subregion of a larger two- or three-dimensional image, and control the addressing calculations of data in buffer memory. If either of these values is zero, that aspect of the buffer memory is considered to be tightly packed according to the imageExtent.
- imageSubresource is a VkImageSubresourceLayers used to specify the specific image subresources of the image used for the source or destination image data.

- imageOffset selects the initial x, y, z offsets in texels of the sub-region of the source or destination image data.
- imageExtent is the size in texels of the image to copy in width, height and depth. 1D images use only x and width. 2D images use x, y, width and height. 3D images use x, y, z, width, height and depth.

When copying to or from a depth or stencil aspect, the data in buffer memory uses a layout that is a (mostly) tightly packed representation of the depth or stencil data. Specifically:

- data copied to or from the stencil aspect of any depth/stencil format is tightly packed with one VK_FORMAT_S8_ UINT value per texel.
- data copied to or from the depth aspect of a VK_FORMAT_D16_UNORM or VK_FORMAT_D16_UNORM_S8_UINT format is tightly packed with one VK_FORMAT_D16_UNORM value per texel.
- data copied to or from the depth aspect of a VK_FORMAT_D32_SFLOAT or VK_FORMAT_D32_SFLOAT_S8_UINT format is tightly packed with one VK_FORMAT_D32_SFLOAT value per texel.
- data copied to or from the depth aspect of a VK_FORMAT_X8_D24_UNORM_PACK32 or VK_FORMAT_D24_ UNORM_S8_UINT format is packed with one 32-bit word per texel with the D24 value in the LSBs of the word, and undefined values in the eight MSBs.



Note

To copy both the depth and stencil aspects of a depth/stencil format, two entries in pRegions can be used, where one specifies the depth aspect in imageSubresource, and the other specifies the stencil aspect.

Because depth or stencil aspect buffer to image copies may require format conversions on some implementations, they are not supported on queues that do not support graphics. When copying to a depth aspect, the data in buffer memory must be in the the range [0,1] or undefined results occur.

Copies are done layer by layer starting with image layer baseArrayLayer member of imageSubresource. layerCount layers are copied from the source image or to the destination image.

- imageSubresource must be a valid VkImageSubresourceLayers structure
- bufferOffset must be a multiple of the calling command's VkImage parameter's texel size
- bufferOffset must be a multiple of 4
- bufferRowLength must be 0, or greater than or equal to the width member of imageExtent
- bufferImageHeight must be 0, or greater than or equal to the height member of imageExtent
- imageOffset.x and (imageExtent.width + imageOffset.x) must both be greater than or equal to 0 and less than or equal to the image subresource width
- imageOffset.y and (imageExtent.height + imageOffset.y) must both be greater than or equal to 0 and less than or equal to the image subresource height

- imageOffset.z and (imageExtent.depth + imageOffset.z) must both be greater than or equal to 0 and less than or equal to the image subresource depth
- If the calling command's Vk Image parameter is a compressed format image:
 - bufferRowLength must be a multiple of the compressed texel block width
 - bufferImageHeight must be a multiple of the compressed texel block height
 - all members of imageOffset must be a multiple of the corresponding dimensions of the compressed texel block
 - bufferOffset must be a multiple of the compressed texel block size in bytes
 - imageExtent.width must be a multiple of the compressed texel block width or (imageExtent.width + imageOffset.x) must equal the image subresource width
 - imageExtent.height must be a multiple of the compressed texel block height or (imageExtent.height + imageOffset.y) must equal the image subresource height
 - imageExtent.depth must be a multiple of the compressed texel block depth or (imageExtent.depth + imageOffset.z) must equal the image subresource depth
- bufferOffset, bufferRowLength, bufferImageHeight and all members of imageOffset and imageExtent must respect the image transfer granularity requirements of the queue family that it will be submitted against, as described in Physical Device Enumeration
- The aspectMask member of imageSubresource must specify aspects present in the calling command's VkImage parameter
- The aspectMask member of imageSubresource must only have a single bit set
- If the calling command's VkImage parameter is of VkImageType VK_IMAGE_TYPE_3D, the baseArrayLayer and layerCount members of imageSubresource must be 0 and 1, respectively
- When copying to the depth aspect of an image subresource, the data in the source buffer must be in the range [0, 1]

Pseudocode for image/buffer addressing is:

```
rowLength = region->bufferRowLength;
if (rowLength == 0)
    rowLength = region->imageExtent.width;

imageHeight = region->bufferImageHeight;
if (imageHeight == 0)
    imageHeight = region->imageExtent.height;

texelSize = <texel size taken from the src/dstImage>;

address of (x,y,z) = region->bufferOffset + (((z * imageHeight) + y) * rowLength + x)  
    * texelSize;

where x,y,z range from (0,0,0) to region->imageExtent.{width,height,depth}.
```

Note that *imageOffset* does not affect addressing calculations for buffer memory. Instead, *bufferOffset* can be used to select the starting address in buffer memory.

For block-compression formats, all parameters are still specified in texels rather than compressed texel blocks, but the addressing math operates on whole compressed texel blocks. Pseudocode for compressed copy addressing is:

```
rowLength = region->bufferRowLength;
if (rowLength == 0)
    rowLength = region->imageExtent.width;

imageHeight = region->bufferImageHeight;
if (imageHeight == 0)
    imageHeight = region->imageExtent.height;

compressedTexelBlockSizeInBytes = <compressed texel block size taken from the src/ \( \to \)
    dstImage>;
rowLength /= compressedTexelBlockWidth;
imageHeight /= compressedTexelBlockHeight;

address of (x,y,z) = region->bufferOffset + (((z * imageHeight) + y) * rowLength + x) \( \to \)
    * compressedTexelBlockSizeInBytes;

where x,y,z range from (0,0,0) to region->imageExtent.{width/compressedTexelBlockWidth \( \to \)
    ,height/compressedTexelBlockHeight,depth/compressedTexelBlockDepth}.
```

Copying to or from block-compressed images is typically done in multiples of the compressed texel block. For this reason the <code>imageExtent</code> must be a multiple of the compressed texel block dimension. There is one exception to this rule which is required to handle compressed images created with dimensions that are not a multiple of the compressed texel block dimensions. If <code>imageExtent.width</code> is not a multiple of the compressed texel block width then (<code>imageExtent.width + imageOffset.x</code>) must equal the image subresource width, if <code>imageExtent.height</code> is not a multiple of the compressed texel block height then (<code>imageExtent.height</code> is not a multiple of the compressed texel block height then (<code>imageExtent.height</code>).

imageOffset.y) must equal the image subresource height and if imageExtent.depth is not a multiple of the compressed texel block depth then (imageExtent.depth + imageOffset.z) must equal the image subresource depth. This allows the last compressed texel block of the image in each non-multiple dimension to be included as a source or destination of the copy.

18.5 Image Copies with Scaling

To copy regions of a source image into a destination image, potentially performing format conversion, arbitrary scaling, and filtering, call:

```
void vkCmdBlitImage(
    VkCommandBuffer
                                                  commandBuffer,
    VkImage
                                                  srcImage,
    VkImageLayout
                                                  srcImageLayout,
    VkImage
                                                  dstImage,
    VkImageLayout
                                                  dstImageLayout,
    uint32_t
                                                  regionCount,
    const VkImageBlit*
                                                  pRegions,
    VkFilter
                                                  filter);
```

- commandBuffer is the command buffer into which the command will be recorded.
- srcImage is the source image.
- srcImageLayout is the layout of the source image subresources for the blit.

- dstImage is the destination image.
- dstImageLayout is the layout of the destination image subresources for the blit.
- regionCount is the number of regions to blit.
- pRegions is a pointer to an array of VkImageBlit structures specifying the regions to blit.
- filter is a VkFilter specifying the filter to apply if the blits require scaling.

vkCmdBlitImage must not be used for multisampled source or destination images. Use vkCmdResolveImage for this purpose.

As the sizes of the source and destination extents can differ in any dimension, texels in the source extent are scaled and filtered to the destination extent. Scaling occurs via the following operations:

• For each destination texel, the integer coordinate of that texel is converted to an unnormalized texture coordinate, using the effective inverse of the equations described in unnormalized to integer conversion:

$$u_{base} = i + \frac{1}{2}$$

$$v_{base} = j + \frac{1}{2}$$

$$w_{base} = k + \frac{1}{2}$$

• These base coordinates are then offset by the first destination offset:

$$u_{offset} = u_{base} - x_{dst_0}$$

 $v_{offset} = v_{base} - y_{dst_0}$
 $w_{offset} = w_{base} - z_{dst_0}$
 $a_{offset} = a - baseArrayCount_{dst}$

• The scale is determined from the source and destination regions, and applied to the offset coordinates:

$$scale_{u} = \frac{x_{src_{1}} - x_{src_{0}}}{x_{dst_{1}} - x_{dst_{0}}}$$

$$scale_{v} = \frac{y_{src_{1}} - y_{src_{0}}}{y_{dst_{1}} - y_{dst_{0}}}$$

$$scale_{w} = \frac{z_{src_{1}} - z_{src_{0}}}{z_{dst_{1}} - z_{dst_{0}}}$$

$$u_{scaled} = u_{offset} * scale_{u}$$

$$v_{scaled} = v_{offset} * scale_{v}$$

$$w_{scaled} = w_{offset} * scale_{w}$$

• Finally the source offset is added to the scaled coordinates, to determine the final unnormalized coordinates used to sample from srcImage:

```
u = u_{scaled} + x_{src_0}

v = v_{scaled} + y_{src_0}

w = w_{scaled} + z_{src_0}

q = mipLevel

a = a_{offset} + baseArrayCount_{src}
```

These coordinates are used to sample from the source image, as described in Image Operations chapter, with the filter mode equal to that of <code>filter</code>, a mipmap mode of <code>VK_SAMPLER_MIPMAP_MODE_NEAREST</code> and an address mode of <code>VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE</code>. Implementations must clamp at the edge of the source image, and may additionally clamp to the edge of the source region.



Note

Due to allowable rounding errors in the generation of the source texture coordinates, it is not always possible to guarantee exactly which source texels will be sampled for a given blit. As rounding errors are implementation dependent, the exact results of a blitting operation are also implementation dependent.

Blits are done layer by layer starting with the baseArrayLayer member of srcSubresource for the source and dstSubresource for the destination. layerCount layers are blitted to the destination image.

3D textures are blitted slice by slice. Slices in the source region bounded by srcoffsets[0].z and srcoffsets[1].z are copied to slices in the destination region bounded by dstoffsets[0].z and dstoffsets[1].z. For each destination slice, a source z coordinate is linearly interpolated between srcoffsets[0].z and srcoffsets[1].z. If the filter parameter is VK_FILTER_LINEAR then the value sampled from the source image is taken by doing linear filtering using the interpolated z coordinate. If filter parameter is VK_FILTER_NEAREST then value sampled from the source image is taken from the single nearest slice (with undefined rounding mode).

The following filtering and conversion rules apply:

- Integer formats can only be converted to other integer formats with the same signedness.
- No format conversion is supported between depth/stencil images the formats must match.
- Format conversions on unorm, snorm, unscaled and packed float formats of the copied aspect of the image are performed by first converting the pixels to float values.
- In case of sRGB source format, nonlinear RGB values are converted to linear representation prior to filtering.
- After filtering, the float values are first clamped and then cast to the destination image format. In case of sRGB destination format, linear RGB values are converted to nonlinear representation before writing the pixel to the image.

Signed and unsigned integers are converted by first clamping to the representable range of the destination format, then casting the value.

- commandBuffer must be a valid VkCommandBuffer handle
- srcImage must be a valid VkImage handle

- srcImageLayout must be a valid VkImageLayout value
- dstImage must be a valid VkImage handle
- dstImageLayout must be a valid VkImageLayout value
- pRegions must be a pointer to an array of regionCount valid VkImageBlit structures
- filter must be a valid VkFilter value
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called outside of a render pass instance
- regionCount must be greater than 0
- Each of commandBuffer, dstImage, and srcImage must have been created, allocated, or retrieved from the same VkDevice
- The source region specified by a given element of pRegions must be a region that is contained within srcImage
- The destination region specified by a given element of pRegions must be a region that is contained within dstImage
- The union of all destination regions, specified by the elements of pRegions, must not overlap in memory with any texel that may be sampled during the blit operation
- srcImage must use a format that supports VK_FORMAT_FEATURE_BLIT_SRC_BIT, which is indicated by VkFormatProperties::linearTilingFeatures (for linear tiled images) or VkFormatProperties::optimalTilingFeatures (for optimally tiled images) as returned by vkGetPhysicalDeviceFormatProperties
- srcImage must have been created with VK_IMAGE_USAGE_TRANSFER_SRC_BIT usage flag
- srcImageLayout must specify the layout of the image subresources of srcImage specified in pRegions at the time this command is executed on a VkDevice
- srcImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL or VK_IMAGE_LAYOUT GENERAL
- dstImage must use a format that supports VK_FORMAT_FEATURE_BLIT_DST_BIT, which is indicated by VkFormatProperties::linearTilingFeatures (for linear tiled images) or VkFormatProperties::optimalTilingFeatures (for optimally tiled images) as returned by vkGetPhysicalDeviceFormatProperties
- dstImage must have been created with VK_IMAGE_USAGE_TRANSFER_DST_BIT usage flag
- dstImageLayout must specify the layout of the image subresources of dstImage specified in pRegions at the time this command is executed on a VkDevice
- dstImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- The sample count of srcImage and dstImage must both be equal to VK_SAMPLE_COUNT_1_BIT
- If either of srcImage or dstImage was created with a signed integer VkFormat, the other must also have been created with a signed integer VkFormat

- If either of <code>srcImage</code> or <code>dstImage</code> was created with an unsigned integer <code>VkFormat</code>, the other must also have been created with an unsigned integer <code>VkFormat</code>
- If either of srcImage or dstImage was created with a depth/stencil format, the other must have exactly the same format
- If srcImage was created with a depth/stencil format, filter must be VK_FILTER_NEAREST
- srcImage must have been created with a samples value of VK SAMPLE COUNT 1 BIT
- dstImage must have been created with a samples value of VK_SAMPLE_COUNT_1_BIT
- If filter is VK_FILTER_LINEAR, srcImage must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- If filter is VK_FILTER_CUBIC_IMG, srcImage must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures(for an optimally tiled image) returned by
- If filter is VK_FILTER_CUBIC_IMG, srcImage must have a VkImageType of VK_IMAGE_TYPE_3D

Host access to commandBuffer must be externally synchronized

vkGetPhysicalDeviceFormatProperties

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		

The VkImageBlit structure is defined as:

```
VkImageSubresourceLayers dstSubresource;
  VkOffset3D dstOffsets[2];
} VkImageBlit;
```

- srcSubresource is the subresource to blit from.
- *srcOffsets* is an array of two VkOffset3D structures specifying the bounds of the source region within *srcSubresource*.
- dstSubresource is the subresource to blit into.
- dstOffsets is an array of two VkOffset3D structures specifying the bounds of the destination region within dstSubresource.

For each element of the pRegions array, a blit operation is performed the specified source and destination regions.

- srcSubresource must be a valid VkImageSubresourceLayers structure
- $\bullet \ \textit{dstSubresource} \ \textbf{must} \ \textbf{be} \ \textbf{a} \ \textbf{valid} \ \textbf{VkImageSubresourceLayers} \ \textbf{structure}$
- The aspectMask member of srcSubresource and dstSubresource must match
- The layerCount member of srcSubresource and dstSubresource must match
- If either of the calling command's <code>srcImage</code> or <code>dstImage</code> parameters are of <code>VkImageTypeVK_IMAGE_TYPE_3D</code>, the <code>baseArrayLayer</code> and <code>layerCount</code> members of both <code>srcSubresource</code> and <code>dstSubresource</code> must be 0 and 1, respectively
- The aspectMask member of srcSubresource must specify aspects present in the calling command's srcImage
- The aspectMask member of dstSubresource must specify aspects present in the calling command's dstImage
- The <code>layerCount</code> member of <code>dstSubresource</code> must be equal to the <code>layerCount</code> member of <code>srcSubresource</code>
- srcOffset[0].x and srcOffset[1].x must both be greater than or equal to 0 and less than or equal to the source image subresource width
- srcOffset[0].y and srcOffset[1].y must both be greater than or equal to 0 and less than or equal to the source image subresource height
- srcOffset[0].z and srcOffset[1].z must both be greater than or equal to 0 and less than or equal to the source image subresource depth
- dstOffset[0].x and dstOffset[1].x must both be greater than or equal to 0 and less than or equal to the destination image subresource width
- dstOffset[0].y and dstOffset[1].y must both be greater than or equal to 0 and less than or equal to the destination image subresource height
- dstOffset[0].z and dstOffset[1].z must both be greater than or equal to 0 and less than or equal to the destination image subresource depth

18.6 Resolving Multisample Images

To resolve a multisample image to a non-multisample image, call:

- commandBuffer is the command buffer into which the command will be recorded.
- srcImage is the source image.
- srcImageLayout is the layout of the source image subresources for the resolve.
- dstImage is the destination image.
- dstImageLayout is the layout of the destination image subresources for the resolve.
- regionCount is the number of regions to resolve.
- pRegions is a pointer to an array of VkImageResolve structures specifying the regions to resolve.

During the resolve the samples corresponding to each pixel location in the source are converted to a single sample before being written to the destination. If the source formats are floating-point or normalized types, the sample values for each pixel are resolved in an implementation-dependent manner. If the source formats are integer types, a single sample's value is selected for each pixel.

srcOffset and dstOffset select the initial x, y, and z offsets in texels of the sub-regions of the source and destination image data. extent is the size in texels of the source image to resolve in width, height and depth. 1D images use only x and width. 2D images use x, y, width and height. 3D images use x, y, z, width, height and depth.

Resolves are done layer by layer starting with baseArrayLayer member of srcSubresource for the source and dstSubresource for the destination. layerCount layers are resolved to the destination image.

- commandBuffer must be a valid VkCommandBuffer handle
- srcImage must be a valid VkImage handle
- srcImageLayout must be a valid VkImageLayout value
- dstImage must be a valid VkImage handle
- dstImageLayout must be a valid VkImageLayout value
- pRegions must be a pointer to an array of regionCount valid VkImageResolve structures

- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called outside of a render pass instance
- regionCount must be greater than 0
- Each of commandBuffer, dstImage, and srcImage must have been created, allocated, or retrieved from the same VkDevice
- The source region specified by a given element of pRegions must be a region that is contained within srcImage
- The destination region specified by a given element of pRegions must be a region that is contained within dst Image
- The union of all source regions, and the union of all destination regions, specified by the elements of pRegions, must not overlap in memory
- srcImage must have a sample count equal to any valid sample count value other than VK_SAMPLE_COUNT_1_ BIT
- dstImage must have a sample count equal to VK_SAMPLE_COUNT_1_BIT
- srcImageLayout must specify the layout of the image subresources of srcImage specified in pRegions at the time this command is executed on a VkDevice
- srcImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- dstImageLayout must specify the layout of the image subresources of dstImage specified in pRegions at the time this command is executed on a VkDevice
- dstImageLayout must be either of VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL or VK_IMAGE_LAYOUT_GENERAL
- If dstImage was created with tiling equal to VK_IMAGE_TILING_LINEAR, dstImage must have been created with a format that supports being a color attachment, as specified by the VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT flag in VkFormatProperties::linearTilingFeatures returned by vkGetPhysicalDeviceFormatProperties
- If dstImage was created with tiling equal to VK_IMAGE_TILING_OPTIMAL, dstImage must have been created with a format that supports being a color attachment, as specified by the VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	GRAPHICS
Secondary		

The VkImageResolve structure is defined as:

- srcSubresource and dstSubresource are VkImageSubresourceLayers structures specifying the image subresources of the images used for the source and destination image data, respectively. Resolve of depth/stencil images is not supported.
- srcOffset and dstOffset select the initial x, y, and z offsets in texels of the sub-regions of the source and destination image data.
- extent is the size in texels of the source image to resolve in width, height and depth. 1D images use only x and width. 2D images use x, y, width and height. 3D images use x, y, z, width, height and depth.

- srcSubresource must be a valid VkImageSubresourceLayers structure
- ullet dstSubresource must be a valid VkImageSubresourceLayers structure
- The aspectMask member of srcSubresource and dstSubresource must only contain VK_IMAGE_ ASPECT_COLOR_BIT
- The layerCount member of srcSubresource and dstSubresource must match
- If either of the calling command's <code>srcImage</code> or <code>dstImage</code> parameters are of <code>VkImageTypeVK_IMAGE_Type_3D</code>, the <code>baseArrayLayer</code> and <code>layerCount</code> members of both <code>srcSubresource</code> and <code>dstSubresource</code> must be 0 and 1, respectively

Chapter 19

Drawing Commands

Drawing commands (commands with **Draw** in the name) provoke work in a graphics pipeline. Drawing commands are recorded into a command buffer and when executed by a queue, will produce work which executes according to the currently bound graphics pipeline. A graphics pipeline must be bound to a command buffer before any drawing commands are recorded in that command buffer.

Each draw is made up of zero or more vertices and zero or more instances, which are processed by the device and result in the assembly of primitives. Primitives are assembled according to the pInputAssemblyState member of the VkGraphicsPipelineCreateInfo structure, which is of type VkPipelineInputAssemblyStateCreateInfo:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- topology is a VkPrimitiveTopology defining the primitive topology, as described below.
- primitiveRestartEnable controls whether a special vertex index value is treated as restarting the assembly of primitives. This enable only applies to indexed draws (vkCmdDrawIndexed and vkCmdDrawIndexedIndirect), and the special index value is either 0xFFFFFFFF when the indexType parameter of vkCmdBindIndexBuffer is equal to VK_INDEX_TYPE_UINT32, or 0xFFFF when indexType is equal to VK_INDEX_TYPE_UINT16. Primitive restart is not allowed for "list" topologies.

Restarting the assembly of primitives discards the most recent index values if those elements formed an incomplete primitive, and restarts the primitive assembly using the subsequent indices, but only assembling the immediately following element through the end of the originally specified elements. The primitive restart index value comparison is performed before adding the <code>vertexOffset</code> value to the index value.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_PIPELINE_INPUT_ASSEMBLY_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- topology must be a valid VkPrimitiveTopology value
- If topology is VK_PRIMITIVE_TOPOLOGY_POINT_LIST, VK_PRIMITIVE_TOPOLOGY_LINE_LIST, VK_PRIMITIVE_TOPOLOGY_LINE_LIST, VK_PRIMITIVE_TOPOLOGY_LINE_LIST_WITH_ ADJACENCY, VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY or VK_PRIMITIVE_ TOPOLOGY_PATCH_LIST, primitiveRestartEnable must be VK_FALSE
- If the geometry shaders feature is not enabled, <code>topology</code> must not be any of <code>VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY</code>, <code>VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY</code>, <code>VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY</code> or <code>VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY</code>

 TRIANGLE_STRIP_WITH_ADJACENCY
- If the tessellation shaders feature is not enabled, topology must not be VK_PRIMITIVE_TOPOLOGY_ PATCH_LIST

19.1 Primitive Topologies

Primitive topology determines how consecutive vertices are organized into primitives, and determines the type of primitive that is used at the beginning of the graphics pipeline. The effective topology for later stages of the pipeline is altered by tessellation or geometry shading (if either is in use) and depends on the execution modes of those shaders. Supported topologies are defined by VkPrimitiveTopology and include:

```
typedef enum VkPrimitiveTopology {
    VK_PRIMITIVE_TOPOLOGY_POINT_LIST = 0,
    VK_PRIMITIVE_TOPOLOGY_LINE_LIST = 1,
    VK_PRIMITIVE_TOPOLOGY_LINE_STRIP = 2,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST = 3,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP = 4,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN = 5,
    VK_PRIMITIVE_TOPOLOGY_LINE_LIST_WITH_ADJACENCY = 6,
    VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY = 7,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY = 8,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY = 9,
    VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY = 9,
    VK_PRIMITIVE_TOPOLOGY_PATCH_LIST = 10,
} VkPrimitiveTopology;
```

Each primitive topology, and its construction from a list of vertices, is summarized below.

19.1.1 Points

A series of individual points are specified with topology VK_PRIMITIVE_TOPOLOGY_POINT_LIST. Each vertex defines a separate point.

19.1.2 Separate Lines

Individual line segments, each defined by a pair of vertices, are specified with topology VK_PRIMITIVE_ TOPOLOGY_LINE_LIST. The first two vertices define the first segment, with subsequent pairs of vertices each defining one more segment. If the number of vertices is odd, then the last vertex is ignored.

19.1.3 Line Strips

A series of one or more connected line segments are specified with topology VK_PRIMITIVE_TOPOLOGY_LINE_STRIP. In this case, the first vertex specifies the first segment's start point while the second vertex specifies the first segment's endpoint and the second segment's start point. In general, the ith vertex (for i > 0) specifies the beginning of the ith segment and the end of the i - 1st. The last vertex specifies the end of the last segment. If only one vertex is specified, then no primitive is generated.

19.1.4 Triangle Strips

A triangle strip is a series of triangles connected along shared edges, and is specified with <code>topology VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP</code>. In this case, the first three vertices define the first triangle, and their order is significant. Each subsequent vertex defines a new triangle using that point along with the last two vertices from the previous triangle, as shown in figure Figure 19.1. If fewer than three vertices are specified, no primitive is produced. The order of vertices in successive triangles changes as shown in the figure, so that all triangle faces have the same orientation.

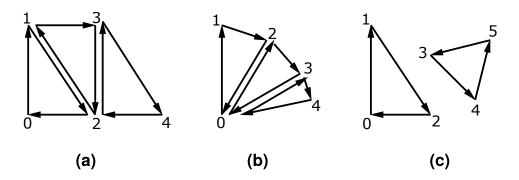


Figure 19.1: Triangle strips, fans, and lists

19.1.5 Triangle Fans

A triangle fan is specified with topology VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN. It is similar to a triangle strip, but changes the vertex replaced from the previous triangle as shown in figure Figure 19.1, so that all triangles in the fan share a common vertex.

19.1.6 Separate Triangles

Separate triangles are specified with topology VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST, as shown in figure Figure 19.1. In this case, vertices 3i, 3i+1, and 3i+2 vertices (in that order) determine a triangle for each $i=0,1,\ldots,n-1$, where there are 3n+k vertices drawn. k is either 0, 1, or 2; if k is not zero, the final k vertices are ignored.

19.1.7 Lines With Adjacency

Lines with adjacency are specified with topology VK_PRIMITIVE_TOPOLOGY_LINE_LIST_WITH_ ADJACENCY, and are independent line segments where each endpoint has a corresponding *adjacent* vertex that is accessible in a geometry shader. If a geometry shader is not active, the adjacent vertices are ignored.

A line segment is drawn from the 4i + 1st vertex to the 4i + 2nd vertex for each i = 0, 1, ..., n - 1, where there are 4n + k vertices. k is either 0, 1, 2, or 3; if k is not zero, the final k vertices are ignored. For line segment i, the 4ith and 4i + 3rd vertices are considered adjacent to the 4i + 1st and 4i + 2nd vertices, respectively, as shown in figure Figure 19.2.



Figure 19.2: Lines with adjacency

19.1.8 Line Strips With Adjacency

Line strips with adjacency are specified with topology VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ ADJACENCY and are similar to line strips, except that each line segment has a pair of adjacent vertices that are accessible in a geometry shader. If a geometry shader is not active, the adjacent vertices are ignored.

A line segment is drawn from the i + 1st vertex to the i + 2nd vertex for each i = 0, 1, ..., n - 1, where there are n + 3 vertices. If there are fewer than four vertices, all vertices are ignored. For line segment i, the ith and i + 3rd vertex are considered adjacent to the i + 1st and i + 2nd vertices, respectively, as shown in figure Figure 19.2.

19.1.9 Triangle List With Adjacency

Triangles with adjacency are specified with topology VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ ADJACENCY, and are similar to separate triangles except that each triangle edge has an adjacent vertex that is accessible in a geometry shader. If a geometry shader is not active, the adjacent vertices are ignored.

The 6ith, 6i + 2nd, and 6i + 4th vertices (in that order) determine a triangle for each i = 0, 1, ..., n - 1, where there are 6n + k vertices. k is either 0, 1, 2, 3, 4, or 5; if k is non-zero, the final k vertices are ignored. For triangle i, the 6i + 1st, 6i + 3rd, and 6i + 5th vertices are considered adjacent to edges from the 6ith to the 6i + 2nd, from the 6i + 4th to the 6ith vertices, respectively, as shown in figure Figure 19.3.

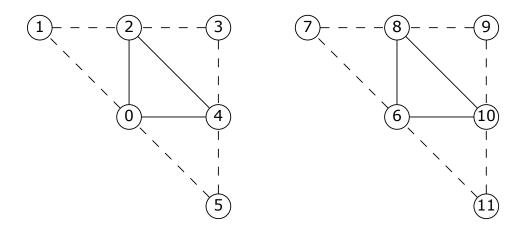


Figure 19.3: Triangles with adjacency

19.1.10 Triangle Strips With Adjacency

Triangle strips with adjacency are specified with topology VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY, and are similar to triangle strips except that each triangle edge has an adjacent vertex that is accessible in a geometry shader. If a geometry shader is not active, the adjacent vertices are ignored.

In triangle strips with adjacency, n triangles are drawn where there are 2(n+2)+k vertices. k is either 0 or 1; if k is 1, the final vertex is ignored. If there are fewer than 6 vertices, the entire primitive is ignored. Table 19.1 describes the vertices and order used to draw each triangle, and which vertices are considered adjacent to each edge of the triangle, as shown in figure Figure 19.4.

	Primitive V	ertices		Adjacent Ve	rtices	
Primitive	1st	2nd	3rd	1/2	2/3	3/1
only $(i = 0, n = 1)$	0	2	4	1	5	3
first $(i = 0)$	0	2	4	1	6	3
middle (i odd)	2i+2	2 <i>i</i>	2i+4	2i-2	2i+3	2i+6
middle (i even)	2i	2i+2	2i+4	2i-2	2i+6	2i + 3
last (i = n - 1, i odd)	2i+2	2 <i>i</i>	2i + 4	2i-2	2i + 3	2i + 5
last $(i = n - 1, i \text{ even})$	2i	2i+2	2i + 4	2i-2	2i+5	2i + 3

Table 19.1: Triangles generated by triangle strips with adjacency.



Figure 19.4: Triangle strips with adjacency

19.1.11 Separate Patches

Separate patches are specified with topology VK_PRIMITIVE_TOPOLOGY_PATCH_LIST. A patch is an ordered collection of vertices used for primitive tessellation. The vertices comprising a patch have no implied geometric ordering, and are used by tessellation shaders and the fixed-function tessellator to generate new point, line, or triangle primitives.

Each patch in the series has a fixed number of vertices, specified by the <code>patchControlPoints</code> member of the <code>VkPipelineTessellationStateCreateInfo</code> structure passed to <code>vkCreateGraphicsPipelines</code>. Once assembled and vertex shaded, these patches are provided as input to the tessellation control shader stage.

If the number of vertices in a patch is given by v, the vith through vi + v - 1st vertices (in that order) determine a patch for each $i = 0, 1, \dots, n-1$, where there are vn + k vertices. k is in the range [0, v-1]; if k is not zero, the final k vertices are ignored.

19.1.12 General Considerations For Polygon Primitives

Depending on the polygon mode, a *polygon primitive* generated from a drawing command with *topology* VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN, VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP, VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY, or VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY is rendered in one of several ways, such as outlining its border or filling its interior. The order of vertices in such a primitive is significant during polygon rasterization and fragment shading.

19.2 Programmable Primitive Shading

Once primitives are assembled, they proceed to the vertex shading stage of the pipeline. If the draw includes multiple instances, then the set of primitives is sent to the vertex shading stage multiple times, once for each instance.

It is undefined whether vertex shading occurs on vertices that are discarded as part of incomplete primitives, but if it does occur then it operates as if they were vertices in complete primitives and such invocations can have side effects.

Vertex shading receives two per-vertex inputs from the primitive assembly stage - the **vertexIndex** and the **instanceIndex**. How these values are generated is defined below, with each command.

Drawing commands fall roughly into two categories:

- Non-indexed drawing commands present a sequential vertexIndex to the vertex shader. The sequential index is
 generated automatically by the device (see Fixed-Function Vertex Processing for details on both specifying the vertex
 attributes indexed by vertexIndex, as well as binding vertex buffers containing those attributes to a command
 buffer). These commands are:
 - vkCmdDraw
 - vkCmdDrawIndirect
 - vkCmdDrawIndirectCountAMD.
- Indexed drawing commands read index values from an *index buffer* and use this to compute the **vertexIndex** value for the vertex shader. These commands are:
 - vkCmdDrawIndexed
 - vkCmdDrawIndexedIndirect
 - vkCmdDrawIndexedIndirectCountAMD.

To bind an index buffer to a command buffer, call:

- commandBuffer is the command buffer into which the command is recorded.
- buffer is the buffer being bound.
- offset is the starting offset in bytes within buffer used in index buffer address calculations.
- indexType selects whether indices are treated as 16 bits or 32 bits. Possible values include:

```
typedef enum VkIndexType {
    VK_INDEX_TYPE_UINT16 = 0,
    VK_INDEX_TYPE_UINT32 = 1,
} VkIndexType;
```

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- buffer must be a valid VkBuffer handle
- indexType must be a valid VkIndexType value
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- Both of buffer, and commandBuffer must have been created, allocated, or retrieved from the same VkDevice
- offset must be less than the size of buffer
- The sum of offset and the address of the range of VkDeviceMemory object that is backing buffer, must be a multiple of the type indicated by indexType
- buffer must have been created with the VK_BUFFER_USAGE_INDEX_BUFFER_BIT flag

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

The parameters for each drawing command are specified directly in the command or read from buffer memory, depending on the command. Drawing commands that source their parameters from buffer memory are known as *indirect* drawing commands.

All drawing commands interact with the Robust Buffer Access feature.

Primitives assembled by draw commands are considered to have an API order, which defines the order their fragments affect the framebuffer. When a draw command includes multiple instances, the lower numbered instances are earlier in API order. For non-indexed draws, primitives with lower numbered **vertexIndex** values are earlier in API order. For indexed draws, primitives assembled from lower index buffer addresses are earlier in API order.

To record a non-indexed draw, call:

- commandBuffer is the command buffer into which the command is recorded.
- vertexCount is the number of vertices to draw.
- instanceCount is the number of instances to draw.
- firstVertex is the index of the first vertex to draw.
- firstInstance is the instance ID of the first instance to draw.

When the command is executed, primitives are assembled using the current primitive topology and <code>vertexCount</code> consecutive vertex indices with the first <code>vertexIndex</code> value equal to <code>firstVertex</code>. The primitives are drawn <code>instanceCount</code> times with <code>instanceIndex</code> starting with <code>firstInstance</code> and increasing sequentially for each instance. The assembled primitives execute the currently bound graphics pipeline.

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_ GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_ POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound
- For a given vertex buffer binding, any attribute data fetched must be entirely contained within the corresponding vertex buffer binding, as described in Section 20.2

- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

To record an indexed draw, call:

- commandBuffer is the command buffer into which the command is recorded.
- indexCount is the number of vertices to draw.
- instanceCount is the number of instances to draw.
- firstIndex is the base index within the index buffer.
- vertexOffset is the value added to the vertex index before indexing into the vertex buffer.
- firstInstance is the instance ID of the first instance to draw.

When the command is executed, primitives are assembled using the current primitive topology and <code>indexCount</code> vertices whose indices are retrieved from the index buffer. The index buffer is treated as an array of tightly packed unsigned integers of size defined by the <code>vkCmdBindIndexBuffer::indexType</code> parameter with which the buffer was bound.

The first vertex index is at an offset of firstIndex*indexSize+offset within the currently bound index buffer, where offset is the offset specified by vkCmdBindIndexBuffer and indexSize is the byte size of the type specified by indexType. Subsequent index values are retrieved from consecutive locations in the index buffer. Indices are first compared to the primitive restart value, then zero extended to 32 bits (if the indexType is VK_INDEX_TYPE_UINT16) and have vertexOffset added to them, before being supplied as the vertexIndex value.

The primitives are drawn instanceCount times with instanceIndex starting with firstInstance and increasing sequentially for each instance. The assembled primitives execute the currently bound graphics pipeline.

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_ GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound
- For a given vertex buffer binding, any attribute data fetched must be entirely contained within the corresponding vertex buffer binding, as described in Section 20.2
- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- (indexSize * (firstIndex + indexCount) + offset) must be less than or equal to the size of the currently bound index buffer, with indexSize being based on the type specified by indexType, where the index buffer, indexType, and offset are specified via vkCmdBindIndexBuffer
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage

- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE ARRAY

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

To record a non-indexed indirect draw, call:

- commandBuffer is the command buffer into which the command is recorded.
- buffer is the buffer containing draw parameters.
- offset is the byte offset into buffer where parameters begin.
- drawCount is the number of draws to execute, and can be zero.
- stride is the byte stride between successive sets of draw parameters.

vkCmdDrawIndirect behaves similarly to vkCmdDraw except that the parameters are read by the device from a buffer during execution. *drawCount* draws are executed by the command, with parameters taken from *buffer* starting at *offset* and increasing by *stride* bytes for each successive draw. The parameters of each draw are encoded in an array of VkDrawIndirectCommand structures. If *drawCount* is less than or equal to one, *stride* is ignored.

- commandBuffer must be a valid VkCommandBuffer handle
- buffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- Both of buffer, and commandBuffer must have been created, allocated, or retrieved from the same VkDevice
- offset must be a multiple of 4
- If drawCount is greater than 1, stride must be a multiple of 4 and must be greater than or equal to sizeof(VkDrawIndirectCommand)
- If the multi-draw indirect feature is not enabled, drawCount must be 0 or 1
- If the drawIndirectFirstInstance feature is not enabled, all the firstInstance members of the VkDrawIndirectCommand structures accessed by this command must be 0
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_ POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound

- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- If drawCount is equal to 1, (offset + sizeof(VkDrawIndirectCommand)) must be less than or equal to the size of buffer
- If drawCount is greater than 1, (stride x (drawCount 1) + offset + sizeof(VkDrawIndirectCommand)) must be less than or equal to the size of buffer
- drawCount must be less than or equal to VkPhysicalDeviceLimits::maxDrawIndirectCount
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

The VkDrawIndirectCommand structure is defined as:

```
typedef struct VkDrawIndirectCommand {
   uint32_t    vertexCount;
   uint32_t    instanceCount;
   uint32_t    firstVertex;
   uint32_t    firstInstance;
} VkDrawIndirectCommand;
```

- vertexCount is the number of vertices to draw.
- instanceCount is the number of instances to draw.
- firstVertex is the index of the first vertex to draw.
- firstInstance is the instance ID of the first instance to draw.

The members of VkDrawIndirectCommand have the same meaning as the similarly named parameters of vkCmdDraw.

Valid Usage

- For a given vertex buffer binding, any attribute data fetched must be entirely contained within the corresponding vertex buffer binding, as described in Section 20.2
- If the drawIndirectFirstInstance feature is not enabled, firstInstance must be 0

To record a non-indexed draw call with a draw call count sourced from a buffer, call:

- commandBuffer is the command buffer into which the command is recorded.
- buffer is the buffer containing draw parameters.
- offset is the byte offset into buffer where parameters begin.
- countBuffer is the buffer containing the draw count.
- countBufferOffset is the byte offset into countBuffer where the draw count begins.
- maxDrawCount specifies the maximum number of draws that will be executed. The actual number of executed draw calls is the minimum of the count specified in countBuffer and maxDrawCount.
- stride is the byte stride between successive sets of draw parameters.

vkCmdDrawIndirectCountAMD behaves similar to vkCmdDrawIndirect except that the draw count is read by the device from a buffer during execution. The command will read an unsigned 32-bit integer from *countBuffer* located at *countBufferOffset* and use this as the draw count.

- commandBuffer must be a valid VkCommandBuffer handle
- buffer must be a valid VkBuffer handle
- countBuffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- Each of buffer, commandBuffer, and countBuffer must have been created, allocated, or retrieved from the same VkDevice
- offset must be a multiple of 4
- countBufferOffset must be a multiple of 4
- stride must be a multiple of 4 and must be greater than or equal to sizeof(VkDrawIndirectCommand)
- If maxDrawCount is greater than or equal to 1, (stride x (maxDrawCount 1) + offset + sizeof(VkDrawIndirectCommand)) must be less than or equal to the size of buffer

- If the drawIndirectFirstInstance feature is not enabled, all the <code>firstInstance</code> members of the <code>VkDrawIndirectCommand</code> structures accessed by this command must be **0**
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_ GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_ POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound
- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- If the count stored in <code>countBuffer</code> is equal to 1, (<code>offset + sizeof(VkDrawIndirectCommand))</code> must be less than or equal to the size of <code>buffer</code>
- If the count stored in <code>countBuffer</code> is greater than 1, (<code>stridex(drawCount-1)+offset+sizeof(VkDrawIndirectCommand))</code> must be less than or equal to the size of <code>buffer</code>
- The count stored in <code>countBuffer</code> must be less than or equal to <code>VkPhysicalDeviceLimits::maxDrawIndirectCount</code>
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set

- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

To record an indexed indirect draw, call:

- commandBuffer is the command buffer into which the command is recorded.
- buffer is the buffer containing draw parameters.
- offset is the byte offset into buffer where parameters begin.
- drawCount is the number of draws to execute, and can be zero.
- stride is the byte stride between successive sets of draw parameters.

vkCmdDrawIndexedIndirect behaves similarly to vkCmdDrawIndexed except that the parameters are read by the device from a buffer during execution. *drawCount* draws are executed by the command, with parameters taken from buffer starting at offset and increasing by stride bytes for each successive draw. The parameters of each draw are encoded in an array of VkDrawIndexedIndirectCommand structures. If *drawCount* is less than or equal to one, stride is ignored.

- commandBuffer must be a valid VkCommandBuffer handle
- buffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- Both of buffer, and commandBuffer must have been created, allocated, or retrieved from the same VkDevice
- offset must be a multiple of 4
- If drawCount is greater than 1, stride must be a multiple of 4 and must be greater than or equal to sizeof(VkDrawIndexedIndirectCommand)
- If the multi-draw indirect feature is not enabled, drawCount must be 0 or 1
- If the drawIndirectFirstInstance feature is not enabled, all the <code>firstInstance</code> members of the <code>VkDrawIndexedIndirectCommand</code> structures accessed by this command must be <code>0</code>
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound
- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- If drawCount is equal to 1, (offset + sizeof(VkDrawIndexedIndirectCommand)) must be less than or equal to the size of buffer

- If drawCount is greater than 1, (stride x (drawCount 1) + offset + sizeof(VkDrawIndexedIndirectCommand)) must be less than or equal to the size of buffer
- drawCount must be less than or equal to VkPhysicalDeviceLimits::maxDrawIndirectCount
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY

Host Synchronization

Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

The VkDrawIndexedIndirectCommand structure is defined as:

```
typedef struct VkDrawIndexedIndirectCommand {
    uint32_t    indexCount;
    uint32_t    instanceCount;
    uint32_t    firstIndex;
    int32_t    vertexOffset;
    uint32_t    firstInstance;
} VkDrawIndexedIndirectCommand;
```

- indexCount is the number of vertices to draw.
- instanceCount is the number of instances to draw.
- firstIndex is the base index within the index buffer.
- vertexOffset is the value added to the vertex index before indexing into the vertex buffer.
- firstInstance is the instance ID of the first instance to draw.

The members of VkDrawIndexedIndirectCommand have the same meaning as the similarly named parameters of vkCmdDrawIndexed.

Valid Usage

- For a given vertex buffer binding, any attribute data fetched must be entirely contained within the corresponding vertex buffer binding, as described in Section 20.2
- (indexSize * (firstIndex + indexCount) + offset) must be less than or equal to the size of the currently bound index buffer, with indexSize being based on the type specified by indexType, where the index buffer, indexType, and offset are specified via vkCmdBindIndexBuffer
- If the drawIndirectFirstInstance feature is not enabled, firstInstance must be 0

To record an indexed draw call with a draw call count sourced from a buffer, call:

- commandBuffer is the command buffer into which the command is recorded.
- buffer is the buffer containing draw parameters.
- offset is the byte offset into buffer where parameters begin.
- countBuffer is the buffer containing the draw count.
- countBufferOffset is the byte offset into countBuffer where the draw count begins.
- maxDrawCount specifies the maximum number of draws that will be executed. The actual number of executed draw calls is the minimum of the count specified in countBuffer and maxDrawCount.
- stride is the byte stride between successive sets of draw parameters.

vkCmdDrawIndexedIndirectCountAMD behaves similar to vkCmdDrawIndirectCountAMD except that the draw count is read by the device from a buffer during execution. The command will read an unsigned 32-bit integer from countBuffer located at countBufferOffset and use this as the draw count.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- buffer must be a valid VkBuffer handle
- countBuffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- This command must only be called inside of a render pass instance
- Each of buffer, commandBuffer, and countBuffer must have been created, allocated, or retrieved from the same VkDevice
- offset must be a multiple of 4
- countBufferOffset must be a multiple of 4
- stride must be a multiple of 4 and must be greater than or equal to sizeof(VkDrawIndirectCommand)

- If maxDrawCount is greater than or equal to 1, (stride x (maxDrawCount 1) + offset + sizeof(VkDrawIndirectCommand)) must be less than or equal to the size of buffer
- If the drawIndirectFirstInstance feature is not enabled, all the firstInstance members of the VkDrawIndexedIndirectCommand structures accessed by this command must be 0
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_ GRAPHICS, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_ POINT_GRAPHICS, a push constant value must have been set for VK_PIPELINE_BIND_POINT_GRAPHICS, with a VkPipelineLayout that is compatible for push constants, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- All vertex input bindings accessed via vertex input variables declared in the vertex shader entry point's interface must have valid buffers bound
- A valid graphics pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ GRAPHICS
- If the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS requires any dynamic state, that state must have been set on the current command buffer
- If count stored in <code>countBuffer</code> is equal to 1, (<code>offset + sizeof(VkDrawIndexedIndirectCommand))</code> must be less than or equal to the size of <code>buffer</code>
- If count stored in <code>countBuffer</code> is greater than 1, (<code>stride x (drawCount 1) + offset + sizeof(VkDrawIndexedIndirectCommand))</code> must be less than or equal to the size of <code>buffer</code>
- drawCount must be less than or equal to VkPhysicalDeviceLimits::maxDrawIndirectCount
- Every input attachment used by the current subpass must be bound to the pipeline via a descriptor set
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_GRAPHICS uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set

- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_GRAPHICS accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Inside	GRAPHICS
Secondary		

Chapter 20

Fixed-Function Vertex Processing

Some implementations have specialized fixed-function hardware for fetching and format-converting vertex input data from buffers, rather than performing the fetch as part of the vertex shader. Vulkan includes a vertex attribute fetch stage in the graphics pipeline in order to take advantage of this.

20.1 Vertex Attributes

Vertex shaders can define input variables, which receive *vertex attribute* data transferred from one or more VkBuffer(s) by drawing commands. Vertex shader input variables are bound to buffers via an indirect binding where the vertex shader associates a *vertex input attribute* number with each variable, vertex input attributes are associated to *vertex input bindings* on a per-pipeline basis, and vertex input bindings are associated with specific buffers on a per-draw basis via the **vkCmdBindVertexBuffers** command. Vertex input attribute and vertex input binding descriptions also contain format information controlling how data is extracted from buffer memory and converted to the format expected by the vertex shader.

There are VkPhysicalDeviceLimits::maxVertexInputAttributes number of vertex input attributes and VkPhysicalDeviceLimits::maxVertexInputBindings number of vertex input bindings (each referred to by zero-based indices), where there are at least as many vertex input attributes as there are vertex input bindings. Applications can store multiple vertex input attributes interleaved in a single buffer, and use a single vertex input binding to access those attributes.

In GLSL, vertex shaders associate input variables with a vertex input attribute number using the **location** layout qualifier. The **component** layout qualifier associates components of a vertex shader input variable with components of a vertex input attribute.

GLSL example

```
// Assign location M to variableName
layout (location=M, component=2) in vec2 variableName;

// Assign locations [N,N+L) to the array elements of variableNameArray
layout (location=N) in vec4 variableNameArray[L];
```

In SPIR-V, vertex shaders associate input variables with a vertex input attribute number using the **Location** decoration. The **Component** decoration associates components of a vertex shader input variable with components of a vertex input attribute. The **Location** and **Component** decorations are specified via the **OpDecorate** instruction.

SPIR-V example

```
%1 = OpExtInstImport "GLSL.std.450"
      . . .
      OpName %9 "variableName"
      OpName %15 "variableNameArray"
      OpDecorate %18 Builtin VertexIndex
      OpDecorate %19 Builtin InstanceIndex
      OpDecorate %9 Location M
      OpDecorate %9 Component 2
      OpDecorate %15 Location N
 %2 = OpTypeVoid
 %3 = OpTypeFunction %2
 %6 = OpTypeFloat 32
 %7 = OpTypeVector %6 2
 %8 = OpTypePointer Input %7
 %9 = OpVariable %8 Input
%10 = OpTypeVector %6 4
%11 = OpTypeInt 32 0
%12 = OpConstant %11 L
%13 = OpTypeArray %10 %12
%14 = OpTypePointer Input %13
%15 = OpVariable %14 Input
```

20.1.1 Attribute Location and Component Assignment

Vertex shaders allow **Location** and **Component** decorations on input variable declarations. The **Location** decoration specifies which vertex input attribute is used to read and interpret the data that a variable will consume. The **Component** decoration allows the location to be more finely specified for scalars and vectors, down to the individual components within a location that are consumed. The components within a location are 0, 1, 2, and 3. A variable starting at component N will consume components N, N+1, N+2, ... up through its size. For single precision types, it is invalid if the sequence of components gets larger than 3.

When a vertex shader input variable declared using a scalar or vector 32-bit data type is assigned a location, its value(s) are taken from the components of the input attribute specified with the corresponding

VkVertexInputAttributeDescription::location. The components used depend on the type of variable and the **Component** decoration specified in the variable declaration, as identified in Table 20.1. Any 32-bit scalar or vector input will consume a single location. For 32-bit data types, missing components are filled in with default values as described below.

32-bit data type	Component decoration	Components consumed
scalar	0 or unspecified	(x, o, o, o)
scalar	1	(o, y, o, o)
scalar	2	(o, o, z, o)
scalar	3	(o, o, o, w)
two-component vector	0 or unspecified	(x, y, o, o)
two-component vector	1	(o, y, z, o)
two-component vector	2	(o, o, z, w)

Table 20.1: Input attribute components accessed by 32-bit input variables

 32-bit data type
 Component decoration
 Components consumed

 three-component vector
 0 or unspecified
 (x, y, z, o)

 three-component vector
 1
 (o, y, z, w)

 four-component vector
 0 or unspecified
 (x, y, z, w)

Table 20.1: (continued)

Components indicated by 'o' are available for use by other input variables which are sourced from the same attribute, and if used, are either filled with the corresponding component from the input format (if present), or the default value.

When a vertex shader input variable declared using a 32-bit floating point matrix type is assigned a location i, its values are taken from consecutive input attributes starting with the corresponding

VkVertexInputAttributeDescription:: *location*. Such matrices are treated as an array of column vectors with values taken from the input attributes identified in Table 20.2. The

VkVertexInputAttributeDescription::format must be specified with a VkFormat that corresponds to the appropriate type of column vector. The Component decoration must not be used with matrix types.

Data type	Column vector type	Locations consumed	Components consumed
mat2	two-component vector	i, i+1	(x, y, o, o), (x, y, o, o)
mat2x3	three-component vector	i, i+1	(x, y, z, o), (x, y, z, o)
mat2x4	four-component vector	i, i+1	(x, y, z, w), (x, y, z, w)
mat3x2	two-component vector	i, i+1, i+2	(x, y, o, o), (x, y, o, o), (x, y, o, o)
mat3	three-component vector	i, i+1, i+2	(x, y, z, o), (x, y, z, o), (x, y, z, o)
mat3x4	four-component vector	i, i+1, i+2	(x, y, z, w), (x, y, z, w), (x, y, z, w)
mat4x2	two-component vector	i, i+1, i+2, i+3	(x, y, o, o), (x, y, o, o), (x, y, o, o), (x, y, o, o)
mat4x3	three-component vector	i, i+1, i+2, i+3	(x, y, z, o), (x, y, z, o), (x, y, z, o), (x, y, z, o)
mat4	four-component vector	i, i+1, i+2, i+3	(x, y, z, w), (x, y, z, w), (x, y, z, w), (x, y, z, w)

Table 20.2: Input attributes accessed by 32-bit input matrix variables

Components indicated by 'o' are available for use by other input variables which are sourced from the same attribute, and if used, are either filled with the corresponding component from the input (if present), or the default value.

When a vertex shader input variable declared using a scalar or vector 64-bit data type is assigned a location i, its values are taken from consecutive input attributes starting with the corresponding

VkVertexInputAttributeDescription::location. The locations and components used depend on the type of variable and the **Component** decoration specified in the variable declaration, as identified in Table 20.3. For 64-bit data types, no default attribute values are provided. Input variables must not use more components than provided by the attribute. Input attributes which have one- or two-component 64-bit formats will consume a single location. Input attributes which have three- or four-component 64-bit formats will consume two consecutive locations. A 64-bit scalar data type will consume two components, and a 64-bit two-component vector data type will consume all four components available within a location. A three- or four-component 64-bit data type must not specify a component. A three-component 64-bit data type will consume all four components of the first location and components 0 and 1 of the second location. This leaves components 2 and 3 available for other component-qualified declarations. A four-component 64-bit data type will consume all four components of the first location and all four components of the second location. It is invalid for a scalar or two-component 64-bit data type to specify a component of 1 or 3.

Table 20.3: Input attribute locations and components accessed by 64-bit input variables

Input format	Locations consumed	64-bit data type	Location decoration	Component decoration	32-bit components consumed
R64	i	scalar	i	0 or unspecified	(x, y, -, -)
		scalar	i	0 or unspecified	(x, y, o, o)
R64G64	i	scalar	i	2	(o, o, z, w)
		two-component vector	i	0 or unspecified	(x, y, z, w)
		scalar	i	0 or unspecified	(x, y, o, o), (o, o, -, -)
R64G64B64	i, i+1	scalar	i	2	(0, 0, z, w), (0, 0, -, -)
		scalar	i+1	0 or unspecified	(0, 0, 0, 0), (x, y, -, -)
		two-component vector	i	0 or unspecified	(x, y, z, w), (o, o, -, -)
		three-component vector	i	unspecified	(x, y, z, w), (x, y, -, -)
		scalar	i	0 or unspecified	(x, y, o, o), (o, o, o, o)
DC1GC1DC11C1		scalar	i	2	(o, o, z, w), (o, o, o, o)
R64G64B64A64	i, i+1	scalar	i+1	0 or unspecified	(o, o, o, o), (x, y, o, o)
		scalar	i+1	2	(0, 0, 0, 0), (0, 0, z, w)
		two-component vector	i	0 or unspecified	(x, y, z, w), (0, 0, 0, 0)
		two-component vector	i+1	0 or unspecified	(o, o, o, o), (x, y, z, w)
		three-component vector	i	unspecified	(x, y, z, w), (x, y, o, o)
		four-component vector	i	unspecified	(x, y, z, w), (x, y, z, w)

Components indicated by 'o' are available for use by other input variables which are sourced from the same attribute. Components indicated by '-' are not available for input variables as there are no default values provided for 64-bit data types, and there is no data provided by the input format.

When a vertex shader input variable declared using a 64-bit floating-point matrix type is assigned a location i, its values are taken from consecutive input attribute locations. Such matrices are treated as an array of column vectors with values taken from the input attributes as shown in Table 20.3. Each column vector starts at the location immediately following the last location of the previous column vector. The number of attributes and components assigned to each matrix is determined by the matrix dimensions and ranges from two to eight locations.

When a vertex shader input variable declared using an array type is assigned a location, its values are taken from consecutive input attributes starting with the corresponding

VkVertexInputAttributeDescription::location. The number of attributes and components assigned to each element are determined according to the data type of the array elements and Component decoration (if any)

specified in the declaration of the array, as described above. Each element of the array, in order, is assigned to consecutive locations, but all at the same specified component within each location.

Only input variables declared with the data types and component decorations as specified above are supported. *Location aliasing* is causing two variables to have the same location number. *Component aliasing* is assigning the same (or overlapping) component number for two location aliases. Location aliasing is allowed only if it does not cause component aliasing. Further, when location aliasing, the aliases sharing the location must all have the same SPIR-V floating-point component type or all have the same width integer-type components.

20.2 Vertex Input Description

Applications specify vertex input attribute and vertex input binding descriptions as part of graphics pipeline creation. The VkGraphicsPipelineCreateInfo::pVertexInputState points to a structure of type VkPipelineVertexInputStateCreateInfo.

The VkPipelineVertexInputStateCreateInfo structure is defined as:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- vertexBindingDescriptionCount is the number of vertex binding descriptions provided in pVertexBindingDescriptions.
- pVertexBindingDescriptions is a pointer to an array of VkVertexInputBindingDescription structures.
- vertexAttributeDescriptionCount is the number of vertex attribute descriptions provided in pVertexAttributeDescriptions.
- pVertexAttributeDescriptions is a pointer to an array of VkVertexInputAttributeDescription structures.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO
- pNext must be NULL

- flags must be 0
- If vertexBindingDescriptionCount is not 0, pVertexBindingDescriptions must be a pointer to an array of vertexBindingDescriptionCount valid VkVertexInputBindingDescription structures
- If vertexAttributeDescriptionCount is not 0, pVertexAttributeDescriptions must be a pointer to an array of vertexAttributeDescriptionCount valid VkVertexInputAttributeDescription structures
- vertexBindingDescriptionCount must be less than or equal to VkPhysicalDeviceLimits::maxVertexInputBindings
- vertexAttributeDescriptionCount must be less than or equal to VkPhysicalDeviceLimits::maxVertexInputAttributes
- For every binding specified by any given element of pVertexAttributeDescriptions, a VkVertexInputBindingDescription must exist in pVertexBindingDescriptions with the same value of binding
- All elements of pVertexBindingDescriptions must describe distinct binding numbers
- All elements of pVertexAttributeDescriptions must describe distinct attribute locations

Each vertex input binding is specified by an instance of the VkVertexInputBindingDescription structure.

The VkVertexInputBindingDescription structure is defined as:

- binding is the binding number that this structure describes.
- stride is the distance in bytes between two consecutive elements within the buffer.
- inputRate specifies whether vertex attribute addressing is a function of the vertex index or of the instance index. Possible values include:

```
typedef enum VkVertexInputRate {
    VK_VERTEX_INPUT_RATE_VERTEX = 0,
    VK_VERTEX_INPUT_RATE_INSTANCE = 1,
} VkVertexInputRate;
```

- $\hbox{--} {\tt VK_VERTEX_INPUT_RATE_VERTEX} \ indicates \ that \ vertex \ attribute \ addressing \ is \ a \ function \ of \ the \ vertex \ index.$
- VK_VERTEX_INPUT_RATE_INSTANCE indicates that vertex attribute addressing is a function of the instance index.

Valid Usage

- inputRate must be a valid VkVertexInputRate value
- binding must be less than VkPhysicalDeviceLimits::maxVertexInputBindings
- stride must be less than or equal to VkPhysicalDeviceLimits::maxVertexInputBindingStride

Each vertex input attribute is specified by an instance of the VkVertexInputAttributeDescription structure.

The VkVertexInputAttributeDescription structure is defined as:

```
typedef struct VkVertexInputAttributeDescription {
   uint32_t location;
   uint32_t binding;
   VkFormat format;
   uint32_t offset;
} VkVertexInputAttributeDescription;
```

- *location* is the shader binding location number for this attribute.
- binding is the binding number which this attribute takes its data from.
- format is the size and type of the vertex attribute data.
- offset is a byte offset of this attribute relative to the start of an element in the vertex input binding.

Valid Usage

- format must be a valid VkFormat value
- location must be less than VkPhysicalDeviceLimits::maxVertexInputAttributes
- binding must be less than VkPhysicalDeviceLimits::maxVertexInputBindings
- offset must be less than or equal to VkPhysicalDeviceLimits::maxVertexInputAttributeOffset
- format must be allowed as a vertex buffer format, as specified by the VK_FORMAT_FEATURE_VERTEX_ BUFFER_BIT flag in VkFormatProperties::bufferFeatures returned by vkGetPhysicalDeviceFormatProperties

To bind vertex buffers to a command buffer for use in subsequent draw commands, call:

- commandBuffer is the command buffer into which the command is recorded.
- firstBinding is the index of the first vertex input binding whose state is updated by the command.
- bindingCount is the number of vertex input bindings whose state is updated by the command.
- pBuffers is a pointer to an array of buffer handles.
- poffsets is a pointer to an array of buffer offsets.

The values taken from elements i of pBuffers and pOffsets replace the current state for the vertex input binding firstBinding + i, for i in [0, bindingCount). The vertex input binding is updated to start at the offset indicated by pOffsets[i] from the start of the buffer pBuffers[i]. All vertex input attributes that use each of these bindings will use these updated addresses in their address calculations for subsequent draw commands.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pBuffers must be a pointer to an array of bindingCount valid VkBuffer handles
- pOffsets must be a pointer to an array of bindingCount VkDeviceSize values
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- bindingCount must be greater than 0
- Both of commandBuffer, and the elements of pBuffers must have been created, allocated, or retrieved from the same VkDevice
- firstBinding must be less than VkPhysicalDeviceLimits::maxVertexInputBindings
- The sum of firstBinding and bindingCount must be less than or equal to VkPhysicalDeviceLimits::maxVertexInputBindings
- ullet All elements of poffsets must be less than the size of the corresponding element in pBuffers
- All elements of pBuffers must have been created with the VK_BUFFER_USAGE_VERTEX_BUFFER_BIT flag

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

The address of each attribute for each **vertexIndex** and **instanceIndex** is calculated as follows:

- Let attribDesc be the member of VkPipelineVertexInputStateCreateInfo::pVertexAttributeDescriptions with VkVertexInputAttributeDescription::location equal to the vertex input attribute number.
- Let bindingDesc be the member of VkPipelineVertexInputStateCreateInfo::pVertexBindingDescriptions with VkVertexInputAttributeDescription::binding equal to attribDesc.binding.
- Let **vertexIndex** be the index of the vertex within the draw (a value between *firstVertex* and *firstVertex*+vertexCount for **vkCmdDraw**, or a value taken from the index buffer for **vkCmdDrawIndexed**), and let **instanceIndex** be the instance number of the draw (a value between *firstInstance* and *firstInstance+instanceCount*).

For each attribute, raw data is extracted starting at attribAddress and is converted from the VkVertexInputAttributeDescription's format to either to floating-point, unsigned integer, or signed integer based on the base type of the format; the base type of the format must match the base type of the input variable in the shader. If format is a packed format, attribAddress must be a multiple of the size in bytes of the whole attribute data type as described in Packed Formats. Otherwise, attribAddress must be a multiple of the size in bytes of the component type indicated by format (see Formats). If the format does not include G, B, or A components, then those are filled with (0,0,1) as needed (using either 1.0f or integer 1 based on the format) for attributes that are not 64-bit data types. The number of components in the vertex shader input variable need not exactly match the number of components in the format. If the vertex shader has fewer components, the extra components are discarded.

20.3 Example

To create a graphics pipeline that uses the following vertex description:

```
struct Vertex
{
    float    x, y, z, w;
    uint8_t u, v;
};
```

The application could use the following set of structures:

```
const VkVertexInputBindingDescription binding =
                                              // binding
   sizeof(Vertex),
                                              // stride
   VK_VERTEX_INPUT_RATE_VERTEX
                                              // inputRate
};
const VkVertexInputAttributeDescription attributes[] =
                                             // location
       binding.binding,
                                             // binding
       VK_FORMAT_R32G32B32A32_SFLOAT,
                                             // format
                                              // offset
                                             // location
                                             // binding
       binding.binding,
                                             // format
       VK_FORMAT_R8G8_UNORM,
       4 * sizeof(float)
                                             // offset
   }
};
const VkPipelineVertexInputStateCreateInfo viInfo =
   VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_CREATE_INFO, // sType
   NULL,
                                // pNext
                                // flags
   0,
   1,
                                // vertexBindingDescriptionCount
                                // pVertexBindingDescriptions
   &binding,
                                // vertexAttributeDescriptionCount
   2,
   &attributes[0]
                                // pVertexAttributeDescriptions
};
```

Chapter 21

Tessellation

Tessellation involves three pipeline stages. First, a tessellation control shader transforms control points of a patch and can produce per-patch data. Second, a fixed-function tessellator generates multiple primitives corresponding to a tessellation of the patch in (u,v) or (u,v,w) parameter space. Third, a tessellation evaluation shader transforms the vertices of the tessellated patch, for example to compute their positions and attributes as part of the tessellated surface. The tessellator is enabled when the pipeline contains both a tessellation control shader and a tessellation evaluation shader.

21.1 Tessellator

If a pipeline includes both tessellation shaders (control and evaluation), the tessellator consumes each input patch (after vertex shading) and produces a new set of independent primitives (points, lines, or triangles). These primitives are logically produced by subdividing a geometric primitive (rectangle or triangle) according to the per-patch outer and inner tessellation levels written by the tessellation control shader. These levels are specified using the built-in variables **TessLevelOuter** and **TessLevelInner**, respectively. This subdivision is performed in an implementation-dependent manner. If no tessellation shaders are present in the pipeline, the tessellator is disabled and incoming primitives are passed through without modification.

The type of subdivision performed by the tessellator is specified by an **OpExecutionMode** instruction in the tessellation evaluation or tessellation control shader using one of execution modes **Triangles**, **Quads**, and **IsoLines**. Other tessellation-related execution modes can also be specified in either the tessellation control or tessellation evaluation shaders, and if they are specified in both then the modes must be the same.

Tessellation execution modes include:

- **Triangles**, **Quads**, and **IsoLines**. These control the type of subdivision and topology of the output primitives. One mode must be set in at least one of the tessellation shader stages.
- **VertexOrderCw** and **VertexOrderCcw**. These control the orientation of triangles generated by the tessellator. One mode must be set in at least one of the tessellation shader stages.
- **PointMode**. Controls generation of points rather than triangles or lines. This functionality defaults to disabled, and is enabled if either shader stage includes the execution mode.
- SpacingEqual, SpacingFractionalEven, and SpacingFractionalOdd. Controls the spacing of segments on the edges of tessellated primitives. One mode must be set in at least one of the tessellation shader stages.
- OutputVertices. Controls the size of the output patch of the tessellation control shader. One value must be set in at least one of the tessellation shader stages.

For triangles, the tessellator subdivides a triangle primitive into smaller triangles. For quads, the tessellator subdivides a rectangle primitive into smaller triangles. For isolines, the tessellator subdivides a rectangle primitive into a collection of line segments arranged in strips stretching across the rectangle in the u dimension (i.e. the coordinates in **TessCoord** are of the form (0,x) through (1,x) for all tessellation evaluation shader invocations that share a line).

Each vertex produced by the tessellator has an associated (u,v,w) or (u,v) position in a normalized parameter space, with parameter values in the range [0,1], as illustrated in figure Figure 21.1.



Figure 21.1: Domain parameterization for tessellation primitive modes

For triangles, the vertex's position is a barycentric coordinate (u,v,w), where u+v+w=1.0, and indicates the relative influence of the three vertices of the triangle on the position of the vertex. For quads and isolines, the position is a (u,v) coordinate indicating the relative horizontal and vertical position of the vertex relative to the subdivided rectangle. The subdivision process is explained in more detail in subsequent sections.

21.2 Tessellator Patch Discard

A patch is discarded by the tessellator if any relevant outer tessellation level is less than or equal to zero.

Patches will also be discarded if any relevant outer tessellation level corresponds to a floating-point NaN (not a number) in implementations supporting NaN.

No new primitives are generated and the tessellation evaluation shader is not executed for patches that are discarded. For **Quads**, all four outer levels are relevant. For **Triangles** and **IsoLines**, only the first three or two outer levels, respectively, are relevant. Negative inner levels will not cause a patch to be discarded; they will be clamped as described below.

21.3 Tessellator Spacing

Each of the tessellation levels is used to determine the number and spacing of segments used to subdivide a corresponding edge. The method used to derive the number and spacing of segments is specified by an **OpExecutionMode** in the tessellation control or tessellation evaluation shader using one of the identifiers **SpacingEqual**, **SpacingFractionalEven**, or **SpacingFractionalOdd**.

If **SpacingEqual** is used, the floating-point tessellation level is first clamped to [1, maxLevel], where maxLevel is the implementation-dependent maximum tessellation level

(VkPhysicalDeviceLimits::maxTessellationGenerationLevel). The result is rounded up to the nearest integer n, and the corresponding edge is divided into n segments of equal length in (u,v) space.

If **SpacingFractionalEven** is used, the tessellation level is first clamped to [2, maxLevel] and then rounded up to the nearest even integer n. If **SpacingFractionalOdd** is used, the tessellation level is clamped to [1, maxLevel - 1] and then rounded up to the nearest odd integer n. If n is one, the edge will not be subdivided. Otherwise, the corresponding edge will be divided into n-2 segments of equal length, and two additional segments of equal length that are typically shorter than the other segments. The length of the two additional segments relative to the others will decrease monotonically with n-f, where f is the clamped floating-point tessellation level. When n-f is zero, the additional segments will have equal length to the other segments. As n-f approaches 2.0, the relative length of the additional segments approaches zero. The two additional segments must be placed symmetrically on opposite sides of the subdivided edge. The relative location of these two segments is implementation-dependent, but must be identical for any pair of subdivided edges with identical values of f.

When the tessellator produces triangles (in the **Triangles** or **Quads** modes), the orientation of all triangles is specified with an **OpExecutionMode** of **VertexOrderCw** or **VertexOrderCcw** in the tessellation control or tessellation evaluation shaders. If the order is **VertexOrderCw**, the vertices of all generated triangles will have clockwise ordering in (u,v) or (u,v,w) space. If the order is **VertexOrderCcw**, the vertices will have counter-clockwise ordering.

The vertices of a triangle have counter-clockwise ordering if

$$a = u_0v_1 - u_1v_0 + u_1v_2 - u_2v_1 + u_2v_0 - u_0v_2$$

is positive, and clockwise ordering if a is negative. u_i and v_i are the u and v coordinates in normalized parameter space of the ith vertex of the triangle.



Note

The value *a* is proportional (with a positive factor) to the signed area of the triangle.

In **Triangles** mode, even though the vertex coordinates have a w value, it does not participate directly in the computation of a, being an affine combination of u and v.

For all primitive modes, the tessellator is capable of generating points instead of lines or triangles. If the tessellation control or tessellation evaluation shader specifies the <code>OpExecutionMode PointMode</code>, the primitive generator will generate one point for each distinct vertex produced by tessellation. Otherwise, the tessellator will produce a collection of line segments or triangles according to the primitive mode. When tessellating triangles or quads in point mode with fractional odd spacing, the tessellator may produce <code>interior vertices</code> that are positioned on the edge of the patch if an

inner tessellation level is less than or equal to one. Such vertices are considered distinct from vertices produced by subdividing the outer edge of the patch, even if there are pairs of vertices with identical coordinates.

The points, lines, or triangles produced by the tessellator are passed to subsequent pipeline stages in an implementation-dependent order.

21.4 Triangle Tessellation

If the tessellation primitive mode is **Triangles**, an equilateral triangle is subdivided into a collection of triangles covering the area of the original triangle. First, the original triangle is subdivided into a collection of concentric equilateral triangles. The edges of each of these triangles are subdivided, and the area between each triangle pair is filled by triangles produced by joining the vertices on the subdivided edges. The number of concentric triangles and the number of subdivisions along each triangle except the outermost is derived from the first inner tessellation level. The edges of the outermost triangle are subdivided independently, using the first, second, and third outer tessellation levels to control the number of subdivisions of the u = 0 (left), v = 0 (bottom), and w = 0 (right) edges, respectively. The second inner tessellation level and the fourth outer tessellation level have no effect in this mode.

If the first inner tessellation level and all three outer tessellation levels are exactly one after clamping and rounding, only a single triangle with (u,v,w) coordinates of (0,0,1), (1,0,0), and (0,1,0) is generated. If the inner tessellation level is one and any of the outer tessellation levels is greater than one, the inner tessellation level is treated as though it were originally specified as $1+\varepsilon$ and will result in a two- or three-segment subdivision depending on the tessellation spacing. When used with fractional odd spacing, the three-segment subdivision may produce *inner vertices* positioned on the edge of the triangle.

If any tessellation level is greater than one, tessellation begins by producing a set of concentric inner triangles and subdividing their edges. First, the three outer edges are temporarily subdivided using the clamped and rounded first inner tessellation level and the specified tessellation spacing, generating n segments. For the outermost inner triangle, the inner triangle is degenerate — a single point at the center of the triangle — if n is two. Otherwise, for each corner of the outer triangle, an inner triangle corner is produced at the intersection of two lines extended perpendicular to the corner's two adjacent edges running through the vertex of the subdivided outer edge nearest that corner. If n is three, the edges of the inner triangle are not subdivided and is the final triangle in the set of concentric triangles. Otherwise, each edge of the inner triangle is divided into n-2 segments, with the n-1 vertices of this subdivision produced by intersecting the inner edge with lines perpendicular to the edge running through the n-1 innermost vertices of the subdivision of the outer edge. Once the outermost inner triangle is subdivided, the previous subdivision process repeats itself, using the generated triangle as an outer triangle. This subdivision process is illustrated in Inner Triangle Tessellation.

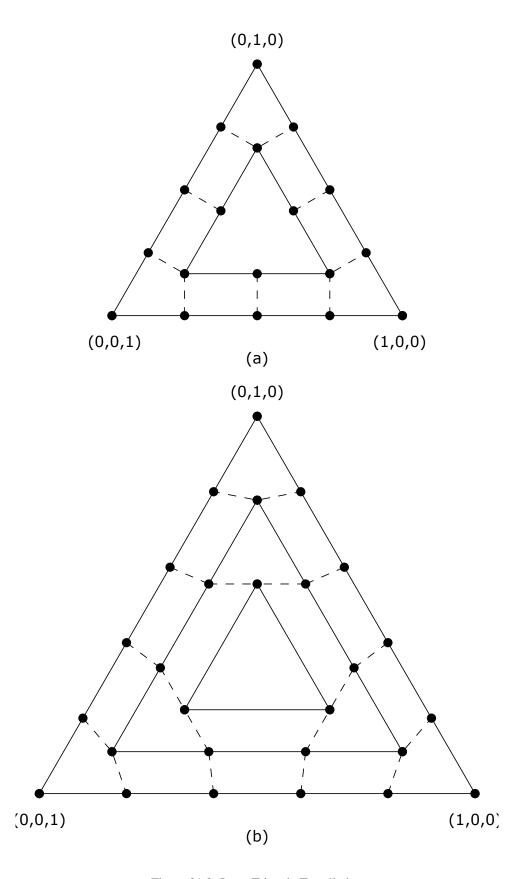


Figure 21.2: Inner Triangle Tessellation

Once all the concentric triangles are produced and their edges are subdivided, the area between each pair of adjacent inner triangles is filled completely with a set of non-overlapping triangles. In this subdivision, two of the three vertices of each triangle are taken from adjacent vertices on a subdivided edge of one triangle; the third is one of the vertices on the corresponding edge of the other triangle. If the innermost triangle is degenerate (i.e., a point), the triangle containing it is subdivided into six triangles by connecting each of the six vertices on that triangle with the center point. If the innermost triangle is not degenerate, that triangle is added to the set of generated triangles as-is.

After the area corresponding to any inner triangles is filled, the tessellator generates triangles to cover the area between the outermost triangle and the outermost inner triangle. To do this, the temporary subdivision of the outer triangle edge above is discarded. Instead, the u = 0, v = 0, and w = 0 edges are subdivided according to the first, second, and third outer tessellation levels, respectively, and the tessellation spacing. The original subdivision of the first inner triangle is retained. The area between the outer and first inner triangles is completely filled by non-overlapping triangles as described above. If the first (and only) inner triangle is degenerate, a set of triangles is produced by connecting each vertex on the outer triangle edges with the center point.

After all triangles are generated, each vertex in the subdivided triangle is assigned a barycentric (u,v,w) coordinate based on its location relative to the three vertices of the outer triangle.

The algorithm used to subdivide the triangular domain in (u,v,w) space into individual triangles is implementation-dependent. However, the set of triangles produced will completely cover the domain, and no portion of the domain will be covered by multiple triangles. The order in which the generated triangles passed to subsequent pipeline stages and the order of the vertices in those triangles are both implementation-dependent. However, when depicted in a manner similar to Inner Triangle Tessellation, the order of the vertices in the generated triangles will be either all clockwise or all counter-clockwise, according to the vertex order layout declaration.

21.5 Quad Tessellation

If the tessellation primitive mode is **Quads**, a rectangle is subdivided into a collection of triangles covering the area of the original rectangle. First, the original rectangle is subdivided into a regular mesh of rectangles, where the number of rectangles along the u=0 and u=1 (vertical) and v=0 and v=1 (horizontal) edges are derived from the first and second inner tessellation levels, respectively. All rectangles, except those adjacent to one of the outer rectangle edges, are decomposed into triangle pairs. The outermost rectangle edges are subdivided independently, using the first, second, third, and fourth outer tessellation levels to control the number of subdivisions of the u=0 (left), v=0 (bottom), u=1 (right), and v=1 (top) edges, respectively. The area between the inner rectangles of the mesh and the outer rectangle edges are filled by triangles produced by joining the vertices on the subdivided outer edges to the vertices on the edge of the inner rectangle mesh.

If both clamped inner tessellation levels and all four clamped outer tessellation levels are exactly one, only a single triangle pair covering the outer rectangle is generated. Otherwise, if either clamped inner tessellation level is one, that tessellation level is treated as though it were originally specified as $1+\varepsilon$ and will result in a two- or three-segment subdivision depending on the tessellation spacing. When used with fractional odd spacing, the three-segment subdivision may produce *inner vertices* positioned on the edge of the rectangle.

If any tessellation level is greater than one, tessellation begins by subdividing the u=0 and u=1 edges of the outer rectangle into m segments using the clamped and rounded first inner tessellation level and the tessellation spacing. The v=0 and v=1 edges are subdivided into n segments using the second inner tessellation level. Each vertex on the u=0 and v=0 edges are joined with the corresponding vertex on the u=1 and v=1 edges to produce a set of vertical and horizontal lines that divide the rectangle into a grid of smaller rectangles. The primitive generator emits a pair of non-overlapping triangles covering each such rectangle not adjacent to an edge of the outer rectangle. The boundary of the region covered by these triangles forms an inner rectangle, the edges of which are subdivided by the grid vertices that lie on the edge. If either m or n is two, the inner rectangle is degenerate, and one or both of the rectangle's edges consist of a single point. This subdivision is illustrated in Figure Inner Quad Tessellation.

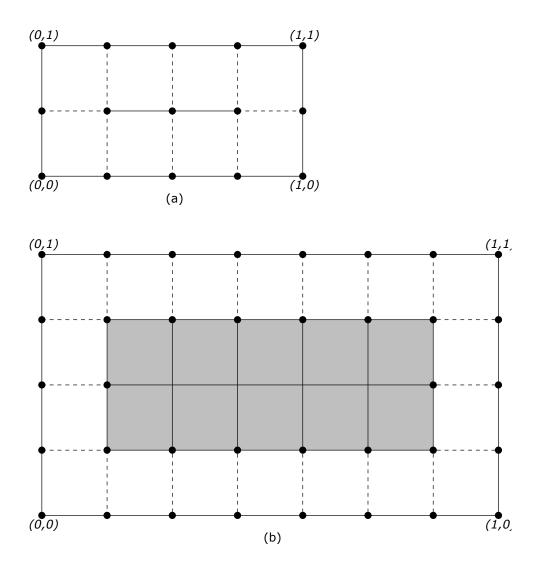


Figure 21.3: Inner Quad Tessellation

After the area corresponding to the inner rectangle is filled, the tessellator must produce triangles to cover the area between the inner and outer rectangles. To do this, the subdivision of the outer rectangle edge above is discarded. Instead, the u = 0, v = 0, u = 1, and v = 1 edges are subdivided according to the first, second, third, and fourth outer tessellation levels, respectively, and the tessellation spacing. The original subdivision of the inner rectangle is retained. The area between the outer and inner rectangles is completely filled by non-overlapping triangles. Two of the three vertices of each triangle are adjacent vertices on a subdivided edge of one rectangle; the third is one of the vertices on the corresponding edge of the other triangle. If either edge of the innermost rectangle is degenerate, the area near the corresponding outer edges is filled by connecting each vertex on the outer edge with the single vertex making up the *inner edge*.

The algorithm used to subdivide the rectangular domain in (u,v) space into individual triangles is implementation-dependent. However, the set of triangles produced will completely cover the domain, and no portion of the domain will be covered by multiple triangles. The order in which the generated triangles passed to subsequent pipeline stages and the order of the vertices in those triangles are both implementation-dependent. However, when depicted in a manner similar to Inner Quad Tessellation, the order of the vertices in the generated triangles will be either all clockwise or all counter-clockwise, according to the vertex order layout declaration.

21.6 Isoline Tessellation

If the tessellation primitive mode is **IsoLines**, a set of independent horizontal line segments is drawn. The segments are arranged into connected strips called *isolines*, where the vertices of each isoline have a constant v coordinate and u coordinates covering the full range [0,1]. The number of isolines generated is derived from the first outer tessellation level; the number of segments in each isoline is derived from the second outer tessellation level. Both inner tessellation levels and the third and fourth outer tessellation levels have no effect in this mode.

Each of the n isolines is then subdivided according to the second outer tessellation level and the tessellation spacing, resulting in m line segments. Each segment of each line is emitted by the tessellator.

The order in which the generated line segments are passed to subsequent pipeline stages and the order of the vertices in each generated line segment are both implementation-dependent.

21.7 Tessellation Pipeline State

The pTessellationState member of VkGraphicsPipelineCreateInfo points to a structure of type VkPipelineTessellationStateCreateInfo.

The VkPipelineTessellationStateCreateInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- patchControlPoints number of control points per patch.

Valid Usage

- stype must be VK STRUCTURE TYPE PIPELINE TESSELLATION STATE CREATE INFO
- pNext must be NULL

- flags must be 0
- patchControlPoints must be greater than zero and less than or equal to VkPhysicalDeviceLimits::maxTessellationPatchSize

Chapter 22

Geometry Shading

The geometry shader operates on a group of vertices and their associated data assembled from a single input primitive, and emits zero or more output primitives and the group of vertices and their associated data required for each output primitive. Geometry shading is enabled when a geometry shader is included in the pipeline.

22.1 Geometry Shader Input Primitives

Each geometry shader invocation has access to all vertices in the primitive (and their associated data), which are presented to the shader as an array of inputs. The input primitive type expected by the geometry shader is specified with an **OpExecutionMode** instruction in the geometry shader, and must be compatible with the primitive topology used by primitive assembly (if tessellation is not in use) or must match the type of primitive generated by the tessellation primitive generator (if tessellation is in use). Compatibility is defined below, with each input primitive type. The input primitive types accepted by a geometry shader are:

Points

Geometry shaders that operate on points use an **OpExecutionMode** instruction specifying the **InputPoints** input mode. Such a shader is valid only when the pipeline primitive topology is **VK_PRIMITIVE_TOPOLOGY_POINT_LIST** (if tessellation is not in use) or if tessellation is in use and the tessellation evaluation shader uses **PointMode**. There is only a single input vertex available for each geometry shader invocation. However, inputs to the geometry shader are still presented as an array, but this array has a length of one.

Lines

Geometry shaders that operate on line segments are generated by including an <code>OpExecutionMode</code> instruction with the <code>InputLines</code> mode. Such a shader is valid only for the <code>VK_PRIMITIVE_TOPOLOGY_LINE_LIST</code>, and <code>VK_PRIMITIVE_TOPOLOGY_LINE_STRIP</code> primitive topologies (if tessellation is not in use) or if tessellation is in use and the tessellation mode is <code>Isolines</code>. There are two input vertices available for each geometry shader invocation. The first vertex refers to the vertex at the beginning of the line segment and the second vertex refers to the vertex at the end of the line segment.

Lines with Adjacency

Geometry shaders that operate on line segments with adjacent vertices are generated by including an OpexecutionMode instruction with the InputLinesAdjacency mode. Such a shader is valid only for the VK_PRIMITIVE_TOPOLOGY_LINES_WITH_ADJACENCY and VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY primitive topologies and must not be used when tessellation is in use.

In this mode, there are four vertices available for each geometry shader invocation. The second vertex refers to attributes of the vertex at the beginning of the line segment and the third vertex refers to the vertex at the end of the

line segment. The first and fourth vertices refer to the vertices adjacent to the beginning and end of the line segment, respectively.

Triangles

Geometry shaders that operate on triangles are created by including an OpExecutionMode instruction with the Triangles mode. Such a shader is valid when the pipeline topology is VK_PRIMITIVE_TOPOLOGY_
TRIANGLE_LIST, VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP, or VK_PRIMITIVE_TOPOLOGY_
TRIANGLE_FAN (if tessellation is not in use) or when tessellation is in use and the tessellation mode is Triangles or Quads.

In this mode, there are three vertices available for each geometry shader invocation. The first, second, and third vertices refer to attributes of the first, second, and third vertex of the triangle, respectively.

Triangles with Adjacency

Geometry shaders that operate on triangles with adjacent vertices are created by including an OpExecutionMode instruction with the InputTrianglesAdjacency mode. Such a shader is valid when the pipeline topology is VK_PRIMITIVE_TOPOLOGY_TRIANGLES_WITH_ADJACENCY or VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY, and must not be used when tessellation is in use.

In this mode, there are six vertices available for each geometry shader invocation. The first, third and fifth vertices refer to attributes of the first, second and third vertex of the triangle, respectively. The second, fourth and sixth vertices refer to attributes of the vertices adjacent to the edges from the first to the second vertex, from the second to the third vertex, and from the third to the first vertex, respectively.

22.2 Geometry Shader Output Primitives

A geometry shader generates primitives in one of three output modes: points, line strips, or triangle strips. The primitive mode is specified in the shader using an **OpExecutionMode** instruction with the **OutputPoints**, **OutputLineStrip** or **OutputTriangleStrip** modes, respectively. Each geometry shader must include exactly one output primitive mode.

The vertices output by the geometry shader are assembled into points, lines, or triangles based on the output primitive type and the resulting primitives are then further processed as described in Chapter 24. If the number of vertices emitted by the geometry shader is not sufficient to produce a single primitive, vertices corresponding to incomplete primitives are not processed by subsequent pipeline stages. The number of vertices output by the geometry shader is limited to a maximum count specified in the shader.

The maximum output vertex count is specified in the shader using an **OpExecutionMode** instruction with the mode set to **OutputVertices** and the maximum number of vertices that will be produced by the geometry shader specified as a literal. Each geometry shader must specify a maximum output vertex count.

22.3 Multiple Invocations of Geometry Shaders

Geometry shaders can be invoked more than one time for each input primitive. This is known as *geometry shader instancing* and is requested by including an **OpExecutionMode** instruction with **mode** specified as **Invocations** and the number of invocations specified as an integer literal.

In this mode, the geometry shader will execute *n* times for each input primitive, where *n* is the number of invocations specified in the **OpExecutionMode** instruction. The instance number is available to each invocation as a built-in input using **InvocationId**.

22.4 Geometry Shader Primitive Ordering

Limited guarantees are provided for the relative ordering of primitives produced by a geometry shader.

- For instanced geometry shaders, the output primitives generated from each input primitive are passed to subsequent pipeline stages using the invocation number to order the primitives, from least to greatest.
- All output primitives generated from a given input primitive are passed to subsequent pipeline stages before any output primitives generated from subsequent input primitives.

Chapter 23

Fixed-Function Vertex Post-Processing

After programmable vertex processing, the following fixed-function operations are applied to vertices of the resulting primitives:

- Flatshading (see Flatshading).
- Primitive clipping, including client-defined half-spaces (see Primitive Clipping).
- Shader output attribute clipping (see Clipping Shader Outputs).
- Perspective division on clip coordinates (see Coordinate Transformations).
- Viewport mapping, including depth range scaling (see Controlling the Viewport).
- Front face determination for polygon primitives (see Basic Polygon Rasterization).

Next, rasterization is performed on primitives as described in chapter Rasterization.

23.1 Flat Shading

Flat shading a vertex output attribute means to assign all vertices of the primitive the same value for that output.

The output values assigned are those of the *provoking vertex* of the primitive. The provoking vertex depends on the primitive topology, and is generally the "first" vertex of the primitive. For primitives not processed by tessellation or geometry shaders, the provoking vertex is selected from the input vertices according to the following table.

Table 23.1: Provoking vertex selection

Primitive type of primitive <i>i</i>	Provoking vertex number
VK_PRIMITIVE_TOPOLOGY_POINT_LIST	i
VK_PRIMITIVE_TOPOLOGY_LINE_LIST	2 <i>i</i>
VK_PRIMITIVE_TOPOLOGY_LINE_STRIP	i
VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST	3 <i>i</i>
VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP	i
VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN	i+1
VK_PRIMITIVE_TOPOLOGY_LINE_LIST_WITH_ADJACENCY	4i+1
VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY	i+1
VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY	6 <i>i</i>
VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY	2 <i>i</i>

Flat shading is applied to those vertex attributes that match fragment input attributes which are decorated as **Flat**.

If a geometry shader is active, the output primitive topology is either points, line strips, or triangle strips, and the selection of the provoking vertex behaves according to the corresponding row of the table. If a tessellation evaluation shader is active and a geometry shader is not active, the provoking vertex is undefined but must be one of the vertices of the primitive.

23.2 Primitive Clipping

Primitives are culled against the *cull volume* and then clipped to the *clip volume*. In clip coordinates, the *view volume* is defined by:

$$-w_c \le x_c \le w_c$$

$$-w_c \le y_c \le w_c$$

$$0 \le z_c \le w_c$$

This view volume can be further restricted by as many as VkPhysicalDeviceLimits::maxClipDistances client-defined half-spaces.

The cull volume is the intersection of up to VkPhysicalDeviceLimits::maxCullDistances client-defined half-spaces (if no client-defined cull half-spaces are enabled, culling against the cull volume is skipped).

A shader must write a single cull distance for each enabled cull half-space to elements of the **CullDistance** array. If the cull distance for any enabled cull half-space is negative for all of the vertices of the primitive under consideration, the primitive is discarded. Otherwise the primitive is clipped against the clip volume as defined below.

The clip volume is the intersection of up to VkPhysicalDeviceLimits::maxClipDistances client-defined half-spaces with the view volume (if no client-defined clip half-spaces are enabled, the clip volume is the view volume).

A shader must write a single clip distance for each enabled clip half-space to elements of the **ClipDistance** array. Clip half-space *i* is then given by the set of points satisfying the inequality

$$c_i(P) \geq 0$$

where $c_i(P)$ is the clip distance i at point P. For point primitives, $c_i(P)$ is simply the clip distance for the vertex in question. For line and triangle primitives, per-vertex clip distances are interpolated using a weighted mean, with weights derived according to the algorithms described in sections Basic Line Segment Rasterization and Basic Polygon Rasterization, using the perspective interpolation equations.

The number of client-defined clip and cull half-spaces that are enabled is determined by the explicit size of the built-in arrays **ClipDistance** and **CullDistance**, respectively, declared as an output in the interface of the entry point of the final shader stage before clipping.

Depth clamping is enabled or disabled via the depthClampEnable enable of the

VkPipelineRasterizationStateCreateInfo structure. If depth clamping is enabled, the plane equation

$$0 \le z_c \le w_c$$

(see the clip volume definition above) is ignored by view volume clipping (effectively, there is no near or far plane clipping).

If the primitive under consideration is a point, then clipping passes it unchanged if it lies within the clip volume; otherwise, it is discarded.

If the primitive is a line segment, then clipping does nothing to it if it lies entirely within the clip volume, and discards it if it lies entirely outside the volume.

If part of the line segment lies in the volume and part lies outside, then the line segment is clipped and new vertex coordinates are computed for one or both vertices. A clipped line segment endpoint lies on both the original line segment and the boundary of the clip volume.

This clipping produces a value, $0 \le t \le 1$, for each clipped vertex. If the coordinates of a clipped vertex are **P** and the original vertices' coordinates are **P**₁ and **P**₂, then *t* is given by

$$\mathbf{P} = t\mathbf{P}_1 + (1-t)\mathbf{P}_2.$$

t is used to clip vertex output attributes as described in Clipping Shader Outputs.

If the primitive is a polygon, it passes unchanged if every one of its edges lie entirely inside the clip volume, and it is discarded if every one of its edges lie entirely outside the clip volume. If the edges of the polygon intersect the boundary of the clip volume, the intersecting edges are reconnected by new edges that lie along the boundary of the clip volume - in some cases requiring the introduction of new vertices into a polygon.

If a polygon intersects an edge of the clip volume's boundary, the clipped polygon must include a point on this boundary edge.

Primitives rendered with user-defined half-spaces must satisfy a complementarity criterion. Suppose a series of primitives is drawn where each vertex i has a single specified clip distance d_i (or a number of similarly specified clip distances, if multiple half-spaces are enabled). Next, suppose that the same series of primitives are drawn again with each such clip distance replaced by $-d_i$ (and the graphics pipeline is otherwise the same). In this case, primitives must not be missing any pixels, and pixels must not be drawn twice in regions where those primitives are cut by the clip planes.

23.3 Clipping Shader Outputs

Next, vertex output attributes are clipped. The output values associated with a vertex that lies within the clip volume are unaffected by clipping. If a primitive is clipped, however, the output values assigned to vertices produced by clipping are clipped.

Let the output values assigned to the two vertices P_1 and P_2 of an unclipped edge be c_1 and c_2 . The value of t (see Primitive Clipping) for a clipped point P is used to obtain the output value associated with P as

$$\mathbf{c} = t\mathbf{c}_1 + (1-t)\mathbf{c}_2.$$

(Multiplying an output value by a scalar means multiplying each of x, y, z, and w by the scalar.)

Since this computation is performed in clip space before division by w_c , clipped output values are perspective-correct.

Polygon clipping creates a clipped vertex along an edge of the clip volume's boundary. This situation is handled by noting that polygon clipping proceeds by clipping against one half-space at a time. Output value clipping is done in the

same way, so that clipped points always occur at the intersection of polygon edges (possibly already clipped) with the clip volume's boundary.

For vertex output attributes whose matching fragment input attributes are decorated with **NoPerspective**, the value of t used to obtain the output value associated with **P** will be adjusted to produce results that vary linearly in framebuffer space.

Output attributes of integer or unsigned integer type must always be flat shaded. Flat shaded attributes are constant over the primitive being rasterized (see Basic Line Segment Rasterization and Basic Polygon Rasterization), and no interpolation is performed. The output value \mathbf{c} is taken from either \mathbf{c}_1 or \mathbf{c}_2 , since flat shading has already occurred and the two values are identical.

23.4 Coordinate Transformations

Clip coordinates for a vertex result from shader execution, which yields a vertex coordinate Position.

Perspective division on clip coordinates yields *normalized device coordinates*, followed by a *viewport* transformation (see Controlling the Viewport) to convert these coordinates into *framebuffer coordinates*.

If a vertex in clip coordinates has a position given by

$$\left(\begin{array}{c} x_c \\ y_c \\ z_c \\ w_c \end{array}\right)$$

then the vertex's normalized device coordinates are

$$\begin{pmatrix} x_d \\ y_d \\ z_d \end{pmatrix} = \begin{pmatrix} \frac{x_c}{w_c} \\ \frac{y_c}{w_c} \\ \frac{z_c}{w_c} \end{pmatrix}$$

23.5 Controlling the Viewport

The viewport transformation is determined by the selected viewport's width and height in pixels, p_x and p_y , respectively, and its center (o_x, o_y) (also in pixels), as well as its depth range min and max determining a depth range scale value p_z

and a depth range bias value o_z (defined below). The vertex's framebuffer coordinates, $\begin{pmatrix} x_f \\ y_f \\ z_f \end{pmatrix}$, are given by

$$\begin{pmatrix} x_f \\ y_f \\ z_f \end{pmatrix} = \begin{pmatrix} \frac{p_x}{2} x_d + o_x \\ \frac{p_y}{2} y_d + o_y \\ p_z \times z_d + o_z \end{pmatrix}.$$

Multiple viewports are available, numbered zero up to VkPhysicalDeviceLimits::maxViewports minus one. The number of viewports used by a pipeline is controlled by the viewportCount member of the VkPipelineViewportStateCreateInfo structure used in pipeline creation.

The VkPipelineViewportStateCreateInfo structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- *flags* is reserved for future use.
- viewportCount is the number of viewports used by the pipeline.
- pViewports is a pointer to an array of VkViewport structures, defining the viewport transforms. If the viewport state is dynamic, this member is ignored.
- scissorCount is the number of scissors and must match the number of viewports.
- pScissors is a pointer to an array of VkRect2D structures which define the rectangular bounds of the scissor for the corresponding viewport. If the scissor state is dynamic, this member is ignored.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_PIPELINE_VIEWPORT_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- viewportCount must be greater than 0
- scissorCount must be greater than 0
- If the multiple viewports feature is not enabled, viewportCount must be 1
- If the multiple viewports feature is not enabled, scissorCount must be 1
- viewportCount must be between 1 and VkPhysicalDeviceLimits::maxViewports, inclusive
- scissorCount must be between 1 and VkPhysicalDeviceLimits::maxViewports, inclusive
- scissorCount and viewportCount must be identical

If a geometry shader is active and has an output variable decorated with **ViewportIndex**, the viewport transformation uses the viewport corresponding to the value assigned to **ViewportIndex** taken from an implementation-dependent vertex of each primitive. If **ViewportIndex** is outside the range zero to <code>viewportCount</code> minus one for a primitive, or if the geometry shader did not assign a value to **ViewportIndex** for all vertices of a primitive due to flow control, the results of the viewport transformation of the vertices of such primitives are undefined. If no geometry shader is active, or if the geometry shader does not have an output decorated with **ViewportIndex**, the viewport numbered zero is used by the viewport transformation.

A single vertex can be used in more than one individual primitive, in primitives such as VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP. In this case, the viewport transformation is applied separately for each primitive.

If the bound pipeline state object was not created with the VK_DYNAMIC_STATE_VIEWPORT dynamic state enabled, viewport transformation parameters are specified using the pViewports member of

VkPipelineViewportStateCreateInfo in the pipeline state object. If the pipeline state object was created with the VK_DYNAMIC_STATE_VIEWPORT dynamic state enabled, the viewport transformation parameters are dynamically set and changed with the command:

- commandBuffer is the command buffer into which the command will be recorded.
- firstViewport is the index of the first viewport whose parameters are updated by the command.
- viewportCount is the number of viewports whose parameters are updated by the command.
- pViewports is a pointer to an array of VkViewport structures specifying viewport parameters.

The viewport parameters taken from element i of pViewports replace the current state for the viewport index firstViewport + i, for i in [0, viewportCount).

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pViewports must be a pointer to an array of viewportCount valid VkViewport structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- viewportCount must be greater than 0
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_VIEWPORT dynamic state enabled
- firstViewport must be less than VkPhysicalDeviceLimits::maxViewports
- The sum of firstViewport and viewportCount must be between 1 and VkPhysicalDeviceLimits::maxViewports, inclusive

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

Both VkPipelineViewportStateCreateInfo and vkCmdSetViewport use VkViewport to set the viewport transformation parameters.

The VkViewport structure is defined as:

```
typedef struct VkViewport {
    float     x;
    float     y;
    float     width;
    float     height;
    float     minDepth;
    float     maxDepth;
} VkViewport;
```

- x and y are the viewport's upper left corner (x, y).
- width and height are the viewport's width and height, respectively.
- minDepth and maxDepth are the depth range for the viewport. It is valid for minDepth to be greater than or equal to maxDepth.

The framebuffer depth coordinate z_f may be represented using either a fixed-point or floating-point representation. However, a floating-point representation must be used if the depth/stencil attachment has a floating-point depth component. If an m-bit fixed-point representation is used, we assume that it represents each value $\frac{k}{2^m-1}$, where $k \in \{0, 1, \dots, 2^m - 1\}$, as k (e.g. 1.0 is represented in binary as a string of all ones).

The viewport parameters shown in the above equations are found from these values as

$$o_x = x + \frac{width}{2}$$

$$o_y = y + \frac{height}{2}$$

$$o_z = minDepth$$

$$p_x = width$$

$$p_y = height$$

$$p_z = maxDepth - minDepth.$$

The width and height of the implementation-dependent maximum viewport dimensions must be greater than or equal to the width and height of the largest image which can be created and attached to a framebuffer.

The floating-point viewport bounds are represented with an implementation-dependent precision.

Valid Usage

- width must be greater than 0.0 and less than or equal to VkPhysicalDeviceLimits::maxViewportDimensions[0]
- height must be greater than 0.0 and less than or equal to VkPhysicalDeviceLimits::maxViewportDimensions[1]
- $\bullet \ \ x \ and \ y \ must \ each \ be \ between \ \textit{viewportBoundsRange}[0] \ and \ \textit{viewportBoundsRange}[1], inclusive \\$
- x + width must be less than or equal to viewportBoundsRange[1]
- y + height must be less than or equal to viewportBoundsRange[1]
- minDepth must be between 0.0 and 1.0, inclusive
- maxDepth must be between 0.0 and 1.0, inclusive
- If the VK_AMD_negative_viewport_height extension is enabled, height can also be negative

Chapter 24

Rasterization

Rasterization is the process by which a primitive is converted to a two-dimensional image. Each point of this image contains associated data such as depth, color, or other attributes.

Rasterizing a primitive begins by determining which squares of an integer grid in framebuffer coordinates are occupied by the primitive, and assigning one or more depth values to each such square. This process is described below for points, lines, and polygons.

A grid square, including its (x, y) framebuffer coordinates, z (depth), and associated data added by fragment shaders, is called a fragment. A fragment is located by its upper left corner, which lies on integer grid coordinates.

Rasterization operations also refer to a fragment's sample locations, which are offset by subpixel fractional values from its upper left corner. The rasterization rules for points, lines, and triangles involve testing whether each sample location is inside the primitive. Fragments need not actually be square, and rasterization rules are not affected by the aspect ratio of fragments. Display of non-square grids, however, will cause rasterized points and line segments to appear fatter in one direction than the other.

We assume that fragments are square, since it simplifies antialiasing and texturing. After rasterization, fragments are processed by the early per-fragment tests, if enabled.

Several factors affect rasterization, including the members of VkPipelineRasterizationStateCreateInfo and VkPipelineMultisampleStateCreateInfo.

The VkPipelineRasterizationStateCreateInfo structure is defined as:

```
typedef struct VkPipelineRasterizationStateCreateInfo {
   VkStructureType
                                               sType;
   const void*
                                               pNext;
                                               flags;
   VkPipelineRasterizationStateCreateFlags
   VkBool32
                                               depthClampEnable;
   VkBool32
                                               rasterizerDiscardEnable;
   VkPolygonMode
                                               polygonMode;
   VkCullModeFlags
                                               cullMode;
   VkFrontFace
                                               frontFace;
   VkBool32
                                               depthBiasEnable;
   float
                                               depthBiasConstantFactor;
   float.
                                               depthBiasClamp;
   float
                                               depthBiasSlopeFactor;
                                               lineWidth;
} VkPipelineRasterizationStateCreateInfo;
```

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- depthClampEnable controls whether to clamp the fragment's depth values instead of clipping primitives to the z planes of the frustum, as described in Primitive Clipping.
- rasterizerDiscardEnable controls whether primitives are discarded immediately before the rasterization stage.
- polygonMode is the triangle rendering mode. See VkPolygonMode.
- cullMode is the triangle facing direction used for primitive culling. See VkCullModeFlagBits.
- frontFace is the front-facing triangle orientation to be used for culling. See VkFrontFace.
- depthBiasEnable controls whether to bias fragment depth values.
- depthBiasConstantFactor is a scalar factor controlling the constant depth value added to each fragment.
- depthBiasClamp is the maximum (or minimum) depth bias of a fragment.
- depthBiasSlopeFactor is a scalar factor applied to a fragment's slope in depth bias calculations.
- lineWidth is the width of rasterized line segments.

The application can also chain a VkPipelineRasterizationStateRasterizationOrderAMD structure to the VkPipelineRasterizationStateCreateInfo structure through its pNext member. This structure enables selecting the rasterization order to use when rendering with the corresponding graphics pipeline as described in Rasterization Order.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_CREATE_INFO
- pNext must be NULL, or a pointer to a valid instance of VkPipelineRasterizationStateRasterizationOrderAMD
- flags must be 0
- $\bullet \ \textit{polygonMode must be a valid} \ \texttt{VkPolygonMode value}$
- cullMode must be a valid combination of VkCullModeFlagBits values
- frontFace must be a valid VkFrontFace value
- If the depth clamping feature is not enabled, depthClampEnable must be VK_FALSE
- If the non-solid fill modes feature is not enabled, polygonMode must be VK POLYGON MODE FILL

The VkPipelineMultisampleStateCreateInfo structure is defined as:

```
typedef struct VkPipelineMultisampleStateCreateInfo {
   VkStructureType
   const void*
                                             pNext;
   VkPipelineMultisampleStateCreateFlags flags;
   VkSampleCountFlagBits
                                             rasterizationSamples;
   VkBool32
                                             sampleShadingEnable;
   float
                                             minSampleShading;
   const VkSampleMask*
                                             pSampleMask;
   VkBool32
                                             alphaToCoverageEnable;
                                             alphaToOneEnable;
   VkBool32
} VkPipelineMultisampleStateCreateInfo;
```

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- rasterizationSamples is a VkSampleCountFlagBits specifying the number of samples per pixel used in rasterization.
- sampleShadingEnable specifies that fragment shading executes per-sample if VK_TRUE, or per-fragment if VK_FALSE, as described in Sample Shading.
- minSampleShading is the minimum fraction of sample shading, as described in Sample Shading.
- pSampleMask is a bitmask of static coverage information that is ANDed with the coverage information generated during rasterization, as described in Sample Mask.
- alphaToCoverageEnable controls whether a temporary coverage value is generated based on the alpha component of the fragment's first color output as specified in the Multisample Coverage section.
- alphaToOneEnable controls whether the alpha component of the fragment's first color output is replaced with one as described in Multisample Coverage.

- sType must be VK_STRUCTURE_TYPE_PIPELINE_MULTISAMPLE_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- rasterizationSamples must be a valid VkSampleCountFlagBits value
- If pSampleMask is not NULL, pSampleMask must be a pointer to an array of $\lceil \frac{rasterizationSamples}{32} \rceil$ VkSampleMask values
- If the sample rate shading feature is not enabled, sampleShadingEnable must be VK_FALSE
- If the alpha to one feature is not enabled, alphaToOneEnable must be VK_FALSE
- minSampleShading must be in the range [0,1]

Rasterization only produces fragments corresponding to pixels in the framebuffer. Fragments which would be produced by application of any of the primitive rasterization rules described below but which lie outside the framebuffer are not produced, nor are they processed by any later stage of the pipeline, including any of the early per-fragment tests described in Early Per-Fragment Tests.

Surviving fragments are processed by fragment shaders. Fragment shaders determine associated data for fragments, and can also modify or replace their assigned depth values.

If the subpass for which this pipeline is being created uses color and/or depth/stencil attachments, then rasterizationSamples must be the same as the sample count for those subpass attachments. Otherwise, rasterizationSamples must follow the rules for a zero-attachment subpass.

24.1 Discarding Primitives Before Rasterization

Primitives are discarded before rasterization if the rasterizerDiscardEnable member of VkPipelineRasterizationStateCreateInfo is enabled. When enabled, primitives are discarded after they are processed by the last active shader stage in the pipeline before rasterization.

24.2 Rasterization Order

Within a subpass of a render pass instance, for a given (x,y,layer,sample) sample location, the following stages are guaranteed to execute in *rasterization order* for each separate primitive that includes that sample location:

- · depth bounds test
- stencil test, stencil op and stencil write
- · depth test and depth write
- · occlusion queries
- blending, logic op and color write

The application can select a graphics pipeline to use one of the following primitive rasterization ordering rules:

```
typedef enum VkRasterizationOrderAMD {
    VK_RASTERIZATION_ORDER_STRICT_AMD = 0,
    VK_RASTERIZATION_ORDER_RELAXED_AMD = 1,
} VkRasterizationOrderAMD;
```

- VK_RASTERIZATION_ORDER_STRICT_AMD indicates that primitive rasterization must follow API order.
- VK RASTERIZATION ORDER RELAXED AMD indicates that primitive rasterization may not follow API order.

The rasterization order to use for a graphics pipeline is specified by chaining a

VkPipelineRasterizationStateRasterizationOrderAMD structure through the pNext member to VkPipelineRasterizationStateCreateInfo.

The VkPipelineRasterizationStateRasterizationOrderAMD structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- rasterizationOrder is the primitive rasterization order to use.

- sType must be VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_RASTERIZATION_ORDER_AMD
- pNext must be NULL
- rasterizationOrder must be a valid VkRasterizationOrderAMD value

If the VK_AMD_rasterization_order device extension is not enabled or the application does not request a particular rasterization order through specifying a

VkPipelineRasterizationStateRasterizationOrderAMD structure then the rasterization order used by the graphics pipeline defaults to VK_RASTERIZATION_ORDER_STRICT_AMD.

24.3 Multisampling

Multisampling is a mechanism to antialias all Vulkan primitives: points, lines, and polygons. The technique is to sample all primitives multiple times at each pixel. Each sample in each framebuffer attachment has storage for a color, depth, and/or stencil value, such that per-fragment operations apply to each sample independently. The color sample values can be later *resolved* to a single color (see Resolving Multisample Images and the Render Pass chapter for more details on how to resolve multisample images to non-multisample images).

Vulkan defines rasterization rules for single-sample modes in a way that is equivalent to a multisample mode with a single sample in the center of each pixel.

Each fragment includes a coverage value with <code>rasterizationSamples</code> bits (see Sample Mask). Each fragment includes <code>rasterizationSamples</code> depth values and sets of associated data. An implementation may choose to assign the same associated data to more than one sample. The location for evaluating such associated data may be anywhere within the pixel including the pixel center or any of the sample locations. When <code>rasterizationSamples</code> is VK_ SAMPLE_COUNT_1_BIT, the pixel center must be used. The different associated data values need not all be evaluated at the same location. Each pixel fragment thus consists of integer x and y grid coordinates, <code>rasterizationSamples</code> depth values and sets of associated data, and a coverage value with <code>rasterizationSamples</code> bits.

It is understood that each pixel has rasterizationSamples locations associated with it. These locations are exact positions, rather than regions or areas, and each is referred to as a sample point. The sample points associated with a

pixel must be located inside or on the boundary of the unit square that is considered to bound the pixel. Furthermore, the relative locations of sample points may be identical for each pixel in the framebuffer, or they may differ. If the current pipeline includes a fragment shader with one or more variables in its interface decorated with **Sample** and **Input**, the data associated with those variables will be assigned independently for each sample. The values for each sample must be evaluated at the location of the sample. The data associated with any other variables not decorated with **Sample** and **Input** need not be evaluated independently for each sample.

If the <code>standardSampleLocations</code> member of <code>VkPhysicalDeviceFeatures</code> is <code>VK_TRUE</code>, then the sample counts <code>VK_SAMPLE_COUNT_1_BIT</code>, <code>VK_SAMPLE_COUNT_2_BIT</code>, <code>VK_SAMPLE_COUNT_4_BIT</code>, <code>VK_SAMPLE_COUNT_4_BIT</code>, <code>VK_SAMPLE_COUNT_8_BIT</code>, and <code>VK_SAMPLE_COUNT_16_BIT</code> have sample locations as listed in the following table, with the <code>ith</code> entry in the table corresponding to bit <code>i</code> in the sample masks. <code>VK_SAMPLE_COUNT_32_BIT</code> and <code>VK_SAMPLE_COUNT_64_BIT</code> do not have standard sample locations. Locations are defined relative to an origin in the upper left corner of the pixel.

VK_SAMPLE_	VK_SAMPLE_	VK_SAMPLE_	VK_SAMPLE_	VK_SAMPLE_
COUNT_1_BIT	COUNT_2_BIT	COUNT_4_BIT	COUNT_8_BIT	COUNT_16_BIT
(0.5, 0.5)	(0.25, 0.25)	(0.375, 0.125)	(0.5625, 0.3125)	(0.5625, 0.5625)
	(0.75, 0.75)	(0.875, 0.375)	(0.4375, 0.6875)	(0.4375, 0.3125)
		(0.125, 0.625)	(0.8125, 0.5625)	(0.3125, 0.625)
		(0.625, 0.875)	(0.3125, 0.1875)	(0.75, 0.4375)
			(0.1875, 0.8125)	(0.1875, 0.375)
			(0.0625, 0.4375)	(0.625, 0.8125)
			(0.6875, 0.9375)	(0.8125, 0.6875)
			(0.9375, 0.0625)	(0.6875, 0.1875)
				(0.375, 0.875)
				(0.5, 0.0625)
				(0.25, 0.125)
				(0.125, 0.75)
				(0.0, 0.5)
				(0.9375, 0.25)
				(0.875, 0.9375)
				(0.0625, 0.0)

Table 24.1: Standard sample locations

24.4 Sample Shading

Sample shading can be used to specify a minimum number of unique samples to process for each fragment. Sample shading is controlled by the <code>sampleShadingEnable</code> member of <code>VkPipelineMultisampleStateCreateInfo</code>. If <code>sampleShadingEnable</code> is <code>VK_FALSE</code>, sample shading is considered disabled and has no effect. Otherwise, an implementation must provide a minimum of max([minSampleShading × rasterizationSamples], 1) unique associated data for each fragment, where <code>minSampleShading</code> is the minimum fraction of sample shading and <code>rasterizationSamples</code> is the number of samples requested in <code>VkPipelineMultisampleStateCreateInfo</code>. These are associated with the samples in an implementation-dependent manner. When the sample shading fraction is 1.0, a separate set of associated data are evaluated for each sample, and each set of values is evaluated at the sample location.

24.5 Points

A point is drawn by generating a set of fragments in the shape of a square centered around the vertex of the point. Each vertex has an associated point size that controls the width/height of that square. The point size is taken from the (potentially clipped) shader built-in **PointSize** written by:

- the geometry shader, if active;
- the tessellation evaluation shader, if active and no geometry shader is active;
- the tessellation control shader, if active and no geometry or tessellation evaluation shader is active; or
- the vertex shader, otherwise

and clamped to the implementation-dependent point size range [pointSizeRange[0], pointSizeRange[1]]. If the value written to PointSize is less than or equal to zero, or if no value was written to PointSize, results are undefined.

Not all point sizes need be supported, but the size 1.0 must be supported. The range of supported sizes and the size of evenly-spaced gradations within that range are implementation-dependent. The range and gradations are obtained from the pointSizeRange and pointSizeGranularity members of VkPhysicalDeviceLimits. If, for instance, the size range is from 0.1 to 2.0 and the gradation size is 0.1, then the size 0.1, 0.2, ..., 1.9, 2.0 are supported. Additional point sizes may also be supported. There is no requirement that these sizes be equally spaced. If an unsupported size is requested, the nearest supported size is used instead.

24.5.1 Basic Point Rasterization

Point rasterization produces a fragment for each framebuffer pixel with one or more sample points that intersect a region centered at the point's (x_f, y_f) . This region is a square with side equal to the current point size. Coverage bits that correspond to sample points that intersect the region are 1, other coverage bits are 0.

All fragments produced in rasterizing a point are assigned the same associated data, which are those of the vertex corresponding to the point. However, the fragment shader built-in **PointCoord** contains point sprite texture coordinates. The *s* and *t* point sprite texture coordinates vary from zero to one across the point horizontally left-to-right and top-to-bottom, respectively. The following formulas are used to evaluate *s* and *t*:

$$s = \frac{1}{2} + \frac{\left(x_p - x_f\right)}{size}$$

$$t = \frac{1}{2} + \frac{\left(y_p - y_f\right)}{\text{size}}.$$

where size is the point's size, (x_p, y_p) is the location at which the point sprite coordinates are evaluated - this may be the framebuffer coordinates of the pixel center (i.e. at the half-integer) or the location of a sample, and (x_f, y_f) is the exact, unrounded framebuffer coordinate of the vertex for the point. When rasterizationSamples is $VK_SAMPLE_COUNT_1_BIT$, the pixel center must be used.

24.6 Line Segments

A line is drawn by generating a set of fragments overlapping a rectangle centered on the line segment. Each line segment has an associated width that controls the width of that rectangle.

The line width is set by the <code>lineWidth</code> property of <code>VkPipelineRasterizationStateCreateInfo</code> in the currently active pipeline if the pipeline was not created with <code>VK_DYNAMIC_STATE_LINE_WIDTH</code> enabled. Otherwise, the line width is set by calling <code>vkCmdSetLineWidth</code>:

- commandBuffer is the command buffer into which the command will be recorded.
- lineWidth is the width of rasterized line segments.

Valid	Usage
-------	-------

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_LINE_WIDTH dynamic state enabled
- If the wide lines feature is not enabled, lineWidth must be 1.0

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

Not all line widths need be supported for line segment rasterization, but width 1.0 antialiased segments must be provided. The range and gradations are obtained from the <code>lineWidthRange</code> and <code>lineWidthGranularity</code> members of <code>VkPhysicalDeviceLimits</code>. If, for instance, the size range is from 0.1 to 2.0 and the gradation size is 0.1, then the size 0.1, 0.2, ..., 1.9, 2.0 are supported. Additional line widths may also be supported. There is no requirement that these widths be equally spaced. If an unsupported width is requested, the nearest supported width is used instead.

24.6.1 Basic Line Segment Rasterization

Rasterized line segments produce fragments which intersect a rectangle centered on the line segment. Two of the edges are parallel to the specified line segment; each is at a distance of one-half the current width from that segment in directions perpendicular to the direction of the line. The other two edges pass through the line endpoints and are perpendicular to the direction of the specified line segment. Coverage bits that correspond to sample points that intersect the rectangle are 1, other coverage bits are 0.

Next we specify how the data associated with each rasterized fragment are obtained. Let $\mathbf{p}_r = (x_d, y_d)$ be the framebuffer coordinates at which associated data are evaluated. This may be the pixel center of a fragment or the location of a sample within the fragment. When rasterizationSamples is VK_SAMPLE_COUNT_1_BIT, the pixel center must be used. Let $\mathbf{p}_a = (x_a, y_a)$ and $\mathbf{p}_b = (x_b, y_b)$ be initial and final endpoints of the line segment, respectively. Set

$$t = \frac{(\mathbf{p}_r - \mathbf{p}_a) \cdot (\mathbf{p}_b - \mathbf{p}_a)}{\|\mathbf{p}_b - \mathbf{p}_a\|^2}$$

(Note that t = 0 at \mathbf{p}_a and t = 1 at \mathbf{p}_b . Also note that this calculation projects the vector from \mathbf{p}_a to \mathbf{p}_r onto the line, and thus computes the normalized distance of the fragment along the line.)

The value of an associated datum f for the fragment, whether it be a shader output or the clip w coordinate, is found as

$$f = \frac{(1-t)f_a/w_a + tf_b/w_b}{(1-t)/w_a + t/w_b}$$

EQUATION 24.1: line_perspective_interpolation

where f_a and f_b are the data associated with the starting and ending endpoints of the segment, respectively; w_a and w_b are the clip w coordinates of the starting and ending endpoints of the segments, respectively. However, depth values for lines must be interpolated by

$$z = (1 - t)z_a + tz_b$$

EQUATION 24.2: line_noperspective_interpolation

where z_a and z_b are the depth values of the starting and ending endpoints of the segment, respectively.

The **NoPerspective** and **Flat** interpolation decorations can be used with fragment shader inputs to declare how they are interpolated. When neither decoration is applied, interpolation is performed as described in Equation line_perspective_interpolation. When the **NoPerspective** decoration is used, interpolation is performed in the same fashion as for depth values, as described in Equation line_noperspective_interpolation. When the **Flat** decoration is used, no interpolation is performed, and outputs are taken from the corresponding input value of the provoking vertex corresponding to that primitive.

The above description documents the preferred method of line rasterization, and must be used when the implementation advertises the <code>strictLines</code> limit in <code>VkPhysicalDeviceLimits</code> as <code>VK_TRUE</code>.

When strictLines is VK_FALSE, the edges of the lines are generated as a parallelogram surrounding the original line. The major axis is chosen by noting the axis in which there is the greatest distance between the line start and end points. If the difference is equal in both directions then the X axis is chosen as the major axis. Edges 2 and 3 are aligned to the minor axis and are centered on the endpoints of the line as in Figure 24.1, and each is lineWidth long. Edges 0 and 1 are parallel to the line and connect the endpoints of edges 2 and 3. Coverage bits that correspond to sample points that intersect the parallelogram are 1, other coverage bits are 0.

Samples that fall exactly on the edge of the parallelogram follow the polygon rasterization rules.

Interpolation occurs as if the parallelogram was decomposed into two triangles where each pair of vertices at each end of the line has identical attributes.

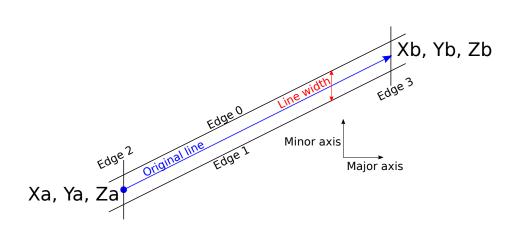


Figure 24.1: Non strict lines

24.7 Polygons

A polygon results from the decomposition of a triangle strip, triangle fan or a series of independent triangles. Like points and line segments, polygon rasterization is controlled by several variables in the VkPipelineRasterizationStateCreateInfo structure.

24.7.1 Basic Polygon Rasterization

The first step of polygon rasterization is to determine whether the triangle is *back-facing* or *front-facing*. This determination is made based on the sign of the (clipped or unclipped) polygon's area computed in framebuffer coordinates. One way to compute this area is:

$$a = -\frac{1}{2} \sum_{i=0}^{n-1} x_f^i y_f^{i \oplus 1} - x_f^{i \oplus 1} y_f^i$$

where x_f^i and y_f^i are the x and y framebuffer coordinates of the ith vertex of the n-vertex polygon (vertices are numbered starting at zero for the purposes of this computation) and $i \oplus 1$ is $(i+1) \mod n$.

The interpretation of the sign of a is determined by the

VkPipelineRasterizationStateCreateInfo::frontFace property of the currently active pipeline, which takes the following values:

```
typedef enum VkFrontFace {
    VK_FRONT_FACE_COUNTER_CLOCKWISE = 0,
    VK_FRONT_FACE_CLOCKWISE = 1,
} VkFrontFace;
```

If frontFace is set to VK_FRONT_FACE_COUNTER_CLOCKWISE, a triangle with positive area is considered front-facing. If it is set to VK_FRONT_FACE_CLOCKWISE, a triangle with negative area is considered front-facing. Any triangle which is not front-facing is back-facing, including zero-area triangles.

Once the orientation of triangles is determined, they are culled according to the setting of the VkPipelineRasterizationStateCreateInfo::cullMode property of the currently active pipeline, which takes the following values:

```
typedef enum VkCullModeFlagBits {
    VK_CULL_MODE_NONE = 0,
    VK_CULL_MODE_FRONT_BIT = 0x00000001,
    VK_CULL_MODE_BACK_BIT = 0x00000002,
    VK_CULL_MODE_FRONT_AND_BACK = 0x00000003,
} VkCullModeFlagBits;
```

If the <code>cullMode</code> is set to <code>VK_CULL_MODE_NONE</code> no triangles are discarded, if it is set to <code>VK_CULL_MODE_FRONT_BIT</code> front-facing triangles are discarded, if it is set to <code>VK_CULL_MODE_BACK_BIT</code> then back-facing triangles are discarded and if it is set to <code>VK_CULL_MODE_FRONT_AND_BACK</code> then all triangles are discarded. Following culling, fragments are produced for any triangles which have not been discarded.

The rule for determining which fragments are produced by polygon rasterization is called *point sampling*. The two-dimensional projection obtained by taking the x and y framebuffer coordinates of the polygon's vertices is formed. Fragments are produced for any pixels for which any sample points lie inside of this polygon. Coverage bits that correspond to sample points that satisfy the point sampling criteria are 1, other coverage bits are 0. Special treatment is given to a sample whose sample location lies on a polygon edge. In such a case, if two polygons lie on either side of a common edge (with identical endpoints) on which a sample point lies, then exactly one of the polygons must result in a covered sample for that fragment during rasterization. As for the data associated with each fragment produced by rasterizing a polygon, we begin by specifying how these values are produced for fragments in a triangle. Define *barycentric coordinates* for a triangle. Barycentric coordinates are a set of three numbers, a, b, and c, each in the range [0,1], with a+b+c=1. These coordinates uniquely specify any point p within the triangle or on the triangle's boundary as

$$p = ap_a + bp_b + cp_c$$

where p_a , p_b , and p_c are the vertices of the triangle. a, b, and c are determined by:

$$a = \frac{A(pp_bp_c)}{A(p_ap_bp_c)}, \quad b = \frac{A(pp_ap_c)}{A(p_ap_bp_c)}, \quad c = \frac{A(pp_ap_b)}{A(p_ap_bp_c)},$$

where A(lmn) denotes the area in framebuffer coordinates of the triangle with vertices l, m, and n.

Denote an associated datum at p_a , p_b , or p_c as f_a , f_b , or f_c , respectively. Then the value f of a datum at a fragment produced by rasterizing a triangle is given by:

$$f = \frac{af_a/w_a + bf_b/w_b + cf_c/w_c}{a/w_a + b/w_b + c/w_c}$$

EQUATION 24.3: triangle_perspective_interpolation

where w_a , w_b , and w_c are the clip w coordinates of p_a , p_b , and p_c , respectively. a, b, and c are the barycentric coordinates of the location at which the data are produced - this must be a pixel center or the location of a sample. When rasterizationSamples is $VK_SAMPLE_COUNT_1_BIT$, the pixel center must be used. Depth values for triangles must be interpolated by

$$z = az_a + bz_b + cz_c$$

EQUATION 24.4: triangle_noperspective_interpolation

where z_a , z_b , and z_c are the depth values of p_a , p_b , and p_c , respectively.

The **NoPerspective** and **Flat** interpolation decorations can be used with fragment shader inputs to declare how they are interpolated. When neither decoration is applied, interpolation is performed as described in Equation triangle_perspective_interpolation. When the **NoPerspective** decoration is used, interpolation is performed in the same fashion as for depth values, as described in Equation triangle_noperspective_interpolation. When the **Flat** decoration is used, no interpolation is performed, and outputs are taken from the corresponding input value of the provoking vertex corresponding to that primitive.

When the VK_AMD_shader_explicit_vertex_parameter device extension is enabled the **CustomInterpAMD** interpolation decoration can also be used with fragment shader inputs which indicate that the decorated inputs can only be accessed by the extended instruction **InterpolateAtVertexAMD** and allows accessing the value of the inputs for individual vertices of the primitive.

For a polygon with more than three edges, such as are produced by clipping a triangle, a convex combination of the values of the datum at the polygon's vertices must be used to obtain the value assigned to each fragment produced by the rasterization algorithm. That is, it must be the case that at every fragment

$$f = \sum_{i=1}^{n} a_i f_i$$

where *n* is the number of vertices in the polygon and f_i is the value of f at vertex i. For each i, $0 \le a_i \le 1$ and $\sum_{i=1}^{n} a_i = 1$. The values of a_i may differ from fragment to fragment, but at vertex i, $a_i = 1$ and $a_j = 0$ for $j \ne i$.

Note



One algorithm that achieves the required behavior is to triangulate a polygon (without adding any vertices) and then treat each triangle individually as already discussed. A scan-line rasterizer that linearly interpolates data along each edge and then linearly interpolates data across each horizontal span from edge to edge also satisfies the restrictions (in this case, the numerator and denominator of equation Equation triangle_perspective_interpolation are iterated independently and a division performed for each fragment).

24.7.2 Polygon Mode

The method of rasterization for polygons is determined by the

 $\label{thm:polygonMode} \begin{tabular}{l} Vk \verb"Pipeline" Rasterization State Create Info::polygon Mode property of the currently active pipeline, which takes the following values: \end{tabular}$

```
typedef enum VkPolygonMode {
   VK_POLYGON_MODE_FILL = 0,
   VK_POLYGON_MODE_LINE = 1,
   VK_POLYGON_MODE_POINT = 2,
} VkPolygonMode;
```

The polygonMode selects which method of rasterization is used for polygons. If polygonMode is VK_POLYGON_MODE_POINT, then the vertices of polygons are treated, for rasterization purposes, as if they had been drawn as points. VK_POLYGON_MODE_LINE causes polygon edges to be drawn as line segments. VK_POLYGON_MODE_FILL causes polygons to render using the polygon rasterization rules in this section.

Note that these modes affect only the final rasterization of polygons: in particular, a polygon's vertices are shaded and the polygon is clipped and possibly culled before these modes are applied.

24.7.3 Depth Bias

The depth values of all fragments generated by the rasterization of a polygon can be offset by a single value that is computed for that polygon. This behavior is controlled by the <code>depthBiasEnable</code>, <code>depthBiasConstantFactor</code>, <code>depthBiasClamp</code>, and <code>depthBiasSlopeFactor</code> members of

VkPipelineRasterizationStateCreateInfo, or by the corresponding parameters to the **vkCmdSetDepthBias** command if depth bias state is dynamic.

- commandBuffer is the command buffer into which the command will be recorded.
- depthBiasConstantFactor is a scalar factor controlling the constant depth value added to each fragment.
- depthBiasClamp is the maximum (or minimum) depth bias of a fragment.
- depthBiasSlopeFactor is a scalar factor applied to a fragment's slope in depth bias calculations.

If depthBiasEnable is VK_FALSE, no depth bias is applied and the fragment's depth values are unchanged.

depthBiasSlopeFactor scales the maximum depth slope of the polygon, and depthBiasConstantFactor scales an implementation-dependent constant that relates to the usable resolution of the depth buffer. The resulting values are summed to produce the depth bias value which is then clamped to a minimum or maximum value specified by depthBiasClamp. depthBiasSlopeFactor, depthBiasConstantFactor, and depthBiasClamp can each be positive, negative, or zero.

The maximum depth slope m of a triangle is

$$m = \sqrt{\left(\frac{\partial z_f}{\partial x_f}\right)^2 + \left(\frac{\partial z_f}{\partial y_f}\right)^2} \tag{24.1}$$

where (x_f, y_f, z_f) is a point on the triangle. m may be approximated as

$$m = \max(\left|\frac{\partial z_f}{\partial x_f}\right|, \left|\frac{\partial z_f}{\partial y_f}\right|). \tag{24.2}$$

The minimum resolvable difference r is an implementation-dependent parameter that depends on the depth buffer representation. It is the smallest difference in framebuffer coordinate z values that is guaranteed to remain distinct throughout polygon rasterization and in the depth buffer. All pairs of fragments generated by the rasterization of two polygons with otherwise identical vertices, but z_f values that differ by \$r\$, will have distinct depth values.

For fixed-point depth buffer representations, r is constant throughout the range of the entire depth buffer. For floating-point depth buffers, there is no single minimum resolvable difference. In this case, the minimum resolvable difference for a given polygon is dependent on the maximum exponent, e, in the range of z values spanned by the primitive. If n is the number of bits in the floating-point mantissa, the minimum resolvable difference, r, for the given primitive is defined as

$$r = 2^{e-n} (24.3)$$

If no depth buffer is present, *r* is undefined.

The bias value o for a polygon is

```
o = \begin{cases} m \times depthBiasSlopeFactor + r \times depthBiasConstantFactor & depthBiasClamp = 0 \text{ or } NaN \\ \min(m \times depthBiasSlopeFactor + r \times depthBiasConstantFactor, depthBiasClamp) & depthBiasClamp > 0 \\ \max(m \times depthBiasSlopeFactor + r \times depthBiasConstantFactor, depthBiasClamp) & depthBiasClamp < 0 \\ (24.4) \end{cases}
```

m is computed as described above. If the depth buffer uses a fixed-point representation, m is a function of depth values in the range [0, 1], and o is applied to depth values in the same range.

For fixed-point depth buffers, fragment depth values are always limited to the range [0,1] by clamping after depth bias addition is performed. Fragment depth values are clamped even when the depth buffer uses a floating-point representation.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_DEPTH_BIAS dynamic state enabled
- If the depth bias clamping feature is not enabled, depthBiasClamp must be 0.0

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

Chapter 25

Fragment Operations

25.1 Early Per-Fragment Tests

Once fragments are produced by rasterization, a number of per-fragment operations are performed prior to fragment shader execution. If a fragment is discarded during any of these operations, it will not be processed by any subsequent stage, including fragment shader execution.

Two fragment operations are performed in the following order:

- the scissor test (see Scissor Test)
- multisample fragment operations (see Sample Mask)

If early per-fragment operations are enabled by the fragment shader, these tests are also performed in the following order:

- the depth bounds tests (see Depth Bounds Tests)
- the stencil test (see Stencil Test)
- the depth test (see Depth Test)
- sample counting (see Sample Counting)

25.2 Scissor Test

The scissor test determines if a fragment's framebuffer coordinates (x_f, y_f) lie within the scissor rectangle corresponding to the viewport index (see Controlling the Viewport) used by the primitive that generated the fragment. If the pipeline state object is created without VK_DYNAMIC_STATE_SCISSOR enabled then the scissor rectangles are set by the VkPipelineViewportStateCreateInfo state of the pipeline state object. Otherwise, to dynamically set the scissor rectangles call:

- commandBuffer is the command buffer into which the command will be recorded.
- firstScissor is the index of the first scissor whose state is updated by the command.
- scissorCount is the number of scissors whose rectangles are updated by the command.
- pScissors is a pointer to an array of VkRect2D structures defining scissor rectangles.

The scissor rectangles taken from element i of pScissors replace the current state for the scissor index firstScissor + i, for i in [0, scissorCount).

Each scissor rectangle is described by a VkRect2D structure, with the offset.x and offset.y values determining the upper left corner of the scissor rectangle, and the extent.width and extent.height values determining the size in pixels.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pScissors must be a pointer to an array of scissorCount VkRect2D structures
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- scissorCount must be greater than 0
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_SCISSOR dynamic state enabled
- firstScissor must be less than VkPhysicalDeviceLimits::maxViewports
- The sum of firstScissor and scissorCount must be between 1 and VkPhysicalDeviceLimits::maxViewports, inclusive
- The x and y members of offset must be greater than or equal to 0
- Evaluation of (offset.x + extent.width) must not cause a signed integer addition overflow
- Evaluation of (offset.y + extent.height) must not cause a signed integer addition overflow

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

If offset. $x \le x_f < offset.x + extent.width$ and offset. $y \le y_f < offset.y + extent.height$ for the selected scissor rectangle, then the scissor test passes. Otherwise, the test fails and the fragment is discarded. For points, lines, and polygons, the scissor rectangle for a primitive is selected in the same manner as the viewport (see Controlling the Viewport). The scissor rectangles only apply to drawing commands, not to other commands like clears or copies.

It is legal for *offset.x* + *extent.width* or *offset.y* + *extent.height* to exceed the dimensions of the framebuffer - the scissor test still applies as defined above. Rasterization does not produce fragments outside of the framebuffer, so such fragments never have the scissor test performed on them.

The scissor test is always performed. Applications can effectively disable the scissor test by specifying a scissor rectangle that encompasses the entire framebuffer.

25.3 Sample Mask

This step modifies fragment coverage values based on the values in the pSampleMask array member of VkPipelineMultisampleStateCreateInfo, as described previously in section Section 9.2.

pSampleMask contains an array of static coverage information that is **ANDed** with the coverage information generated during rasterization. Bits that are zero disable coverage for the corresponding sample. Bit B of mask word M corresponds to sample $32 \times M + B$. The array is sized to a length of [rasterizationSamples/32] words. If pSampleMask is NULL, it is treated as if the mask has all bits enabled, i.e. no coverage is removed from fragments.

The elements of the sample mask array are of type VkSampleMask, each representing 32 bits of coverage information:

typedef uint32_t VkSampleMask;

25.4 Early Fragment Test Mode

The depth bounds test, stencil test, depth test, and occlusion query sample counting are performed before fragment shading if and only if early fragment tests are enabled by the fragment shader (see Early Fragment Tests). When early per-fragment operations are enabled, these operations are performed prior to fragment shader execution, and the stencil buffer, depth buffer, and occlusion query sample counts will be updated accordingly; these operations will not be performed again after fragment shader execution.

If a pipeline's fragment shader has early fragment tests disabled, these operations are performed only after fragment program execution, in the order described below. If a pipeline does not contain a fragment shader, these operations are performed only once.

If early fragment tests are enabled, any depth value computed by the fragment shader has no effect. Additionally, the depth test (including depth writes), stencil test (including stencil writes) and sample counting operations are performed even for fragments or samples that would be discarded after fragment shader execution due to per-fragment operations such as alpha-to-coverage tests, or due to the fragment being discarded by the shader itself.

25.5 Late Per-Fragment Tests

After programmable fragment processing, per-fragment operations are performed before blending and color output to the framebuffer.

A fragment is produced by rasterization with framebuffer coordinates of (x_f, y_f) and depth z, as described in Rasterization. The fragment is then modified by programmable fragment processing, which adds associated data as described in Shaders. The fragment is then further modified, and possibly discarded by the late per-fragment operations described in this chapter. Finally, if the fragment was not discarded, it is used to update the framebuffer at the fragment's framebuffer coordinates for any samples that remain covered.

The depth bounds test, stencil test, and depth test are performed for each pixel sample, rather than just once for each fragment. Stencil and depth operations are performed for a pixel sample only if that sample's fragment coverage bit is a value of 1 when the fragment executes the corresponding stage of the graphics pipeline. If the corresponding coverage bit is 0, no operations are performed for that sample. Failure of the depth bounds, stencil, or depth test results in termination of the processing of that sample by means of disabling coverage for that sample, rather than discarding of the fragment. If, at any point, a fragment's coverage becomes zero for all samples, then the fragment is discarded. All operations are performed on the depth and stencil values stored in the depth/stencil attachment of the framebuffer. The contents of the color attachments are not modified at this point.

The depth bounds test, stencil test, depth test, and occlusion query operations described in Depth Bounds Test, Stencil Test, Depth Test, Sample Counting are instead performed prior to fragment processing, as described in Early Fragment Test Mode, if requested by the fragment shader.

25.6 Multisample Coverage

If a fragment shader is active and its entry point's interface includes a built-in output variable decorated with <code>SampleMask</code>, the fragment coverage is <code>ANDed</code> with the bits of the sample mask to generate a new fragment coverage value. If such a fragment shader did not assign a value to <code>SampleMask</code> due to flow of control, the value <code>ANDed</code> with the fragment coverage is undefined. If no fragment shader is active, or if the active fragment shader does not include <code>SampleMask</code> in its interface, the fragment coverage is not modified.

Next, the fragment alpha and coverage values are modified based on the alphaToCoverageEnable and alphaToOneEnable members of the VkPipelineMultisampleStateCreateInfo structure.

All alpha values in this section refer only to the alpha component of the fragment shader output that has a **Location** and **Index** decoration of zero (see the Fragment Output Interface section). If that shader output has an integer or unsigned integer type, then these operations are skipped.

If alphaToCoverageEnable is enabled, a temporary coverage value is generated where each bit is determined by the fragment's alpha value. The temporary coverage value is then ANDed with the fragment coverage value to generate a new fragment coverage value.

No specific algorithm is specified for converting the alpha value to a temporary coverage mask. It is intended that the number of 1's in this value be proportional to the alpha value (clamped to [0,1]), with all 1's corresponding to a value of 1.0 and all 0's corresponding to 0.0. The algorithm may be different at different pixel locations.



Note

Using different algorithms at different pixel location may help to avoid artifacts caused by regular coverage sample locations.

Next, if alphaToOneEnable is enabled, each alpha value is replaced by the maximum representable alpha value for fixed-point color buffers, or by 1.0 for floating-point buffers. Otherwise, the alpha values are not changed.

25.7 Depth and Stencil Operations

Pipeline state controlling the depth bounds tests, stencil test, and depth test is specified through the members of the VkPipelineDepthStencilStateCreateInfo structure.

The VkPipelineDepthStencilStateCreateInfo structure is defined as:

```
typedef struct VkPipelineDepthStencilStateCreateInfo {
   VkStructureType
                                               sType;
   const void*
                                               pNext;
   VkPipelineDepthStencilStateCreateFlags
                                              flags;
   VkBool32
                                               depthTestEnable;
   VkBool32
                                               depthWriteEnable;
   VkCompareOp
                                               depthCompareOp;
   VkBool32
                                               depthBoundsTestEnable;
   VkBool32
                                               stencilTestEnable;
   VkStencilOpState
                                               front;
   VkStencilOpState
                                               back:
   float
                                               minDepthBounds;
   float
                                               maxDepthBounds;
} VkPipelineDepthStencilStateCreateInfo;
```

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- depthTestEnable controls whether depth testing is enabled.
- depthWriteEnable controls whether depth writes are enabled.
- depthCompareOp is the comparison operator used in the depth test.
- depthBoundsTestEnable controls whether depth bounds testing is enabled.
- stencilTestEnable controls whether stencil testing is enabled.
- front and back control the parameters of the stencil test.
- minDepthBounds and maxDepthBounds define the range of values used in the depth bounds test.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_PIPELINE_DEPTH_STENCIL_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- depthCompareOp must be a valid VkCompareOp value

- front must be a valid VkStencilOpState structure
- back must be a valid VkStencilOpState structure
- If the depth bounds testing feature is not enabled, depthBoundsTestEnable must be VK_FALSE

25.8 Depth Bounds Test

The depth bounds test conditionally disables coverage of a sample based on the outcome of a comparison between the value z_a in the depth attachment at location (x_f, y_f) (for the appropriate sample) and a range of values. The test is enabled or disabled by the depthBoundsTestEnable member of VkPipelineDepthStencilStateCreateInfo: If the pipeline state object is created without the $VK_DYNAMIC_STATE_DEPTH_BOUNDS$ dynamic state enabled then the range of values used in the depth bounds test are defined by the minDepthBounds and maxDepthBounds members of the VkPipelineDepthStencilStateCreateInfo structure. Otherwise, to dynamically set the depth bounds range values call:

- commandBuffer is the command buffer into which the command will be recorded.
- minDepthBounds is the lower bound of the range of depth values used in the depth bounds test.
- maxDepthBounds is the upper bound of the range.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_DEPTH_ BOUNDS dynamic state enabled
- minDepthBounds must be between 0.0 and 1.0, inclusive
- maxDepthBounds must be between 0.0 and 1.0, inclusive

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

If $minDepthBounds \le z_a \le maxDepthBounds$, then the depth bounds test passes. Otherwise, the test fails and the sample's coverage bit is cleared in the fragment. If there is no depth framebuffer attachment or if the depth bounds test is disabled, it is as if the depth bounds test always passes.

25.9 Stencil Test

The stencil test conditionally disables coverage of a sample based on the outcome of a comparison between the stencil value in the depth/stencil attachment at location (x_f, y_f) (for the appropriate sample) and a reference value. The stencil test also updates the value in the stencil attachment, depending on the test state, the stencil value and the stencil write masks. The test is enabled or disabled by the stencilTestEnable member of VkPipelineDepthStencilStateCreateInfo.

When disabled, the stencil test and associated modifications are not made, and the sample's coverage is not modified.

The stencil test is controlled with the front and back members of

VkPipelineDepthStencilStateCreateInfo which are of type VkStencilOpState.

The VkStencilOpState structure is defined as:

```
typedef struct VkStencilOpState {
   VkStencilOp    failOp;
   VkStencilOp    passOp;
   VkStencilOp    depthFailOp;
   VkCompareOp    compareOp;
   uint32_t    compareMask;
   uint32_t    writeMask;
   uint32_t    reference;
} VkStencilOpState;
```

- failop is the action performed on samples that fail the stencil test.
- passOp is the action performed on samples that pass both the depth and stencil tests.
- depthFailOp is the action performed on samples that pass the stencil test and fail the depth test.

- compareOp is the comparison operator used in the stencil test.
- compareMask selects the bits of the unsigned integer stencil values participating in the stencil test.
- writeMask selects the bits of the unsigned integer stencil values updated by the stencil test in the stencil framebuffer attachment.
- reference is an integer reference value that is used in the unsigned stencil comparison.

- failOp must be a valid VkStencilOp value
- passOp must be a valid VkStencilOp value
- depthFailOp must be a valid VkStencilOp value
- compareOp must be a valid VkCompareOp value

There are two sets of stencil-related state, the front stencil state set and the back stencil state set. Stencil tests and writes use the front set of stencil state when processing fragments rasterized from non-polygon primitives (points and lines) and front-facing polygon primitives while the back set of stencil state is used when processing fragments rasterized from back-facing polygon primitives. For the purposes of stencil testing, a primitive is still considered a polygon even if the polygon is to be rasterized as points or lines due to the current VkPolygonMode. Whether a polygon is front- or back-facing is determined in the same manner used for face culling (see Basic Polygon Rasterization).

The operation of the stencil test is also affected by the <code>compareMask</code>, <code>writeMask</code>, and <code>reference</code> members of <code>VkStencilOpState</code> set in the pipeline state object if the pipeline state object is created without the <code>VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK</code>, <code>VK_DYNAMIC_STATE_STENCIL_WRITE_MASK</code>, and <code>VK_DYNAMIC_STATE_STENCIL_REFERENCE</code> dynamic states enabled, respectively.

If the pipeline state object is created with the VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK dynamic state enabled, then to dynamically set the stencil compare mask call:

- commandBuffer is the command buffer into which the command will be recorded.
- faceMask is a bitmask specifying the set of stencil state for which to update the compare mask. Bits which can be set include:

```
typedef enum VkStencilFaceFlagBits {
   VK_STENCIL_FACE_FRONT_BIT = 0x00000001,
   VK_STENCIL_FACE_BACK_BIT = 0x00000002,
   VK_STENCIL_FRONT_AND_BACK = 0x00000003,
} VkStencilFaceFlagBits;
```

- VK_STENCIL_FACE_FRONT_BIT indicates that only the front set of stencil state is updated.
- VK_STENCIL_FACE_BACK_BIT indicates that only the back set of stencil state is updated.
- VK_STENCIL_FRONT_AND_BACK is the combination of VK_STENCIL_FACE_FRONT_BIT and VK_STENCIL_FACE_BACK_BIT and indicates that both sets of stencil state are updated.
- compareMask is the new value to use as the stencil compare mask.

- commandBuffer must be a valid VkCommandBuffer handle
- faceMask must be a valid combination of VkStencilFaceFlagBits values
- faceMask must not be 0
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK dynamic state enabled

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

If the pipeline state object is created with the VK_DYNAMIC_STATE_STENCIL_WRITE_MASK dynamic state enabled, then to dynamically set the stencil write mask call:

- commandBuffer is the command buffer into which the command will be recorded.
- faceMask is a bitmask of VkStencilFaceFlagBits specifying the set of stencil state for which to update the write mask, as described above for vkCmdSetStencilCompareMask.
- writeMask is the new value to use as the stencil write mask.

- commandBuffer must be a valid VkCommandBuffer handle
- faceMask must be a valid combination of VkStencilFaceFlagBits values
- faceMask must not be 0
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_STENCIL_ WRITE_MASK dynamic state enabled

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

If the pipeline state object is created with the VK_DYNAMIC_STATE_STENCIL_REFERENCE dynamic state enabled, then to dynamically set the stencil reference value call:

- commandBuffer is the command buffer into which the command will be recorded.
- faceMask is a bitmask of VkStencilFaceFlagBits specifying the set of stencil state for which to update the reference value, as described above for vkCmdSetStencilCompareMask.
- reference is the new value to use as the stencil reference value.

- commandBuffer must be a valid VkCommandBuffer handle
- faceMask must be a valid combination of VkStencilFaceFlagBits values
- faceMask must not be 0
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_STENCIL_ REFERENCE dynamic state enabled

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

reference is an integer reference value that is used in the unsigned stencil comparison. Stencil comparison clamps the reference value to $[0,2^s-1]$, where s is the number of bits in the stencil framebuffer attachment. The s least significant bits of compareMask are bitwise **ANDed** with both the reference and the stored stencil value, and the resulting masked values are those that participate in the comparison controlled by compareOp. Let R be the masked reference value and S be the masked stored stencil value.

compareOp is a symbolic constant that determines the stencil comparison function:

```
typedef enum VkCompareOp {
    VK_COMPARE_OP_NEVER = 0,
    VK_COMPARE_OP_LESS = 1,
    VK_COMPARE_OP_EQUAL = 2,
    VK_COMPARE_OP_LESS_OR_EQUAL = 3,
    VK_COMPARE_OP_GREATER = 4,
    VK_COMPARE_OP_NOT_EQUAL = 5,
    VK_COMPARE_OP_REATER_OR_EQUAL = 6,
    VK_COMPARE_OP_ALWAYS = 7,
} VkCompareOp;
```

- VK COMPARE OP NEVER: the test never passes.
- VK_COMPARE_OP_LESS: the test passes when R < S.
- VK_COMPARE_OP_EQUAL: the test passes when R = S.
- VK_COMPARE_OP_LESS_OR_EQUAL: the test passes when $R \leq S$.
- VK_COMPARE_OP_GREATER: the test passes when R > S.
- VK_COMPARE_OP_NOT_EQUAL: the test passes when $R \neq S$.
- VK_COMPARE_OP_GREATER_OR_EQUAL: the test passes when $R \geq S$.
- VK_COMPARE_OP_ALWAYS: the test always passes.

As described earlier, the failop, passop, and depthFailop members of VkStencilOpState indicate what happens to the stored stencil value if this or certain subsequent tests fail or pass. Each enum is of type VkStencilOp, which is defined as:

```
typedef enum VkStencilOp {
    VK_STENCIL_OP_KEEP = 0,
    VK_STENCIL_OP_ZERO = 1,
    VK_STENCIL_OP_REPLACE = 2,
    VK_STENCIL_OP_INCREMENT_AND_CLAMP = 3,
    VK_STENCIL_OP_DECREMENT_AND_CLAMP = 4,
    VK_STENCIL_OP_INVERT = 5,
    VK_STENCIL_OP_INCREMENT_AND_WRAP = 6,
    VK_STENCIL_OP_DECREMENT_AND_WRAP = 7,
} VkStencilOp;
```

The possible values are:

- VK_STENCIL_OP_KEEP keeps the current value.
- VK_STENCIL_OP_ZERO sets the value to 0.
- VK_STENCIL_OP_REPLACE sets the value to reference.
- VK_STENCIL_OP_INCREMENT_AND_CLAMP increments the current value and clamps to the maximum representable unsigned value.
- VK_STENCIL_OP_DECREMENT_AND_CLAMP decrements the current value and clamps to 0.
- VK_STENCIL_OP_INVERT bitwise-inverts the current value.

- VK_STENCIL_OP_INCREMENT_AND_WRAP increments the current value and wraps to 0 when the maximum value
 would have been exceeded.
- VK_STENCIL_OP_DECREMENT_AND_WRAP decrements the current value and wraps to the maximum possible value when the value would go below 0.

For purposes of increment and decrement, the stencil bits are considered as an unsigned integer.

If the stencil test fails, the sample's coverage bit is cleared in the fragment. If there is no stencil framebuffer attachment, stencil modification cannot occur, and it is as if the stencil tests always pass.

If the stencil test passes, the writeMask member of the VkStencilOpState structures controls how the updated stencil value is written to the stencil framebuffer attachment.

The least significant s bits of writeMask, where s is the number of bits in the stencil framebuffer attachment, specify an integer mask. Where a 1 appears in this mask, the corresponding bit in the stencil value in the depth/stencil attachment is written; where a 0 appears, the bit is not written. The writeMask value uses either the front-facing or back-facing state based on the facing-ness of the fragment. Fragments generated by front-facing primitives use the front mask and fragments generated by back-facing primitives use the back mask.

25.10 Depth Test

The depth test conditionally disables coverage of a sample based on the outcome of a comparison between the fragment's depth value at the sample location and the sample's depth value in the depth/stencil attachment at location (x_f, y_f) . The comparison is enabled or disabled with the depthTestEnable member of the

VkPipelineDepthStencilStateCreateInfo structure. When disabled, the depth comparison and subsequent possible updates to the value of the depth component of the depth/stencil attachment are bypassed and the fragment is passed to the next operation. The stencil value, however, can be modified as indicated above as if the depth test passed. If enabled, the comparison takes place and the depth/stencil attachment value can subsequently be modified.

The comparison is specified with the depthCompareOp member of

VkPipelineDepthStencilStateCreateInfo. Let z_f be the incoming fragment's depth value for a sample, and let z_a be the depth/stencil attachment value in memory for that sample. The depth test passes under the following conditions:

- VK_COMPARE_OP_NEVER: the test never passes.
- VK_COMPARE_OP_LESS: the test passes when $z_f < z_a$.
- VK_COMPARE_OP_EQUAL: the test passes when $z_f = z_a$.
- VK_COMPARE_OP_LESS_OR_EQUAL: the test passes when $z_f \leq z_a$.
- VK_COMPARE_OP_GREATER: the test passes when $z_f > z_a$.
- VK_COMPARE_OP_NOT_EQUAL: the test passes when $z_f \neq z_a$.
- VK_COMPARE_OP_GREATER_OR_EQUAL: the test passes when $z_f \ge z_a$.
- VK COMPARE OP ALWAYS: the test always passes.

If depth clamping (see Primitive Clipping) is enabled, before the incoming fragment's z_f is compared to z_a , z_f is clamped to $[\min(n, f), \max(n, f)]$, where n and f are the $\min Depth$ and $\max Depth$ depth range values of the viewport used by this fragment, respectively.

If the depth test fails, the sample's coverage bit is cleared in the fragment. The stencil value at the sample's location is updated according to the function currently in effect for depth test failure.

If the depth test passes, the sample's (possibly clamped) z_f value is conditionally written to the depth framebuffer attachment based on the depthWriteEnable member of VkPipelineDepthStencilStateCreateInfo. If depthWriteEnable is VK_TRUE the value is written, and if it is VK_FALSE the value is not written. The stencil value at the sample's location is updated according to the function currently in effect for depth test success.

If there is no depth framebuffer attachment, it is as if the depth test always passes.

25.11 Sample Counting

Occlusion queries use query pool entries to track the number of samples that pass all the per-fragment tests. The mechanism of collecting an occlusion query value is described in Occlusion Queries.

The occlusion query sample counter increments by one for each sample with a coverage value of 1 in each fragment that survives all the per-fragment tests, including scissor, sample mask, alpha to coverage, stencil, and depth tests.

Chapter 26

The Framebuffer

26.1 Blending

Blending combines the incoming "source" fragment's R, G, B, and A values with the "destination" R, G, B, and A values of each sample stored in the framebuffer at the fragment's (x_f, y_f) location. Blending is performed for each pixel sample, rather than just once for each fragment.

Source and destination values are combined according to the blend operation, quadruplets of source and destination weighting factors determined by the blend factors, and a blend constant, to obtain a new set of R, G, B, and A values, as described below.

Blending is computed and applied separately to each color attachment used by the subpass, with separate controls for each attachment.

Prior to performing the blend operation, signed and unsigned normalized fixed-point color components undergo an implied conversion to floating-point as specified by Conversion from Normalized Fixed-Point to Floating-Point.

Blending computations are treated as if carried out in floating-point, and will be performed with a precision and dynamic range no lower than that used to represent destination components.

Blending applies only to fixed-point and floating-point color attachments. If the color attachment has an integer format, blending is not applied.

The pipeline blend state is included in the VkPipelineColorBlendStateCreateInfo structure during graphics pipeline creation:

The VkPipelineColorBlendStateCreateInfo structure is defined as:

```
typedef struct VkPipelineColorBlendStateCreateInfo {
   VkStructureType
                                                  sType;
   const void*
                                                  pNext;
   VkPipelineColorBlendStateCreateFlags
                                                  flags;
                                                  logicOpEnable;
   VkBool32
   VkLogicOp
                                                  logicOp;
                                                  attachmentCount;
   const VkPipelineColorBlendAttachmentState*
                                                 pAttachments;
                                                  blendConstants[4];
VkPipelineColorBlendStateCreateInfo;
```

• sType is the type of this structure.

- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- logicOpEnable controls whether to apply Logical Operations.
- logicOp selects which logical operation to apply.
- attachmentCount is the number of VkPipelineColorBlendAttachmentState elements in pAttachments. This value must equal the colorAttachmentCount for the subpass in which this pipeline is used.
- pAttachments: is a pointer to array of per target attachment states.
- blendConstants is an array of four values used as the R, G, B, and A components of the blend constant that are used in blending, depending on the blend factor.

Each element of the pAttachments array is a VkPipelineColorBlendAttachmentState structure specifying per-target blending state for each individual color attachment. If the independent blending feature is not enabled on the device, all VkPipelineColorBlendAttachmentState elements in the pAttachments array must be identical.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_PIPELINE_COLOR_BLEND_STATE_CREATE_INFO
- pNext must be NULL
- flags must be 0
- If attachmentCount is not 0, pAttachments must be a pointer to an array of attachmentCount valid VkPipelineColorBlendAttachmentState structures
- If the independent blending feature is not enabled, all elements of pAttachments must be identical
- If the logic operations feature is not enabled, <code>logicOpEnable</code> must be <code>VK_FALSE</code>
- If logicOpEnable is VK_TRUE, logicOp must be a valid VkLogicOp value

The VkPipelineColorBlendAttachmentState structure is defined as:

```
typedef struct VkPipelineColorBlendAttachmentState {
   VkBool32
                            blendEnable;
   VkBlendFactor
                            srcColorBlendFactor;
                        srcColorBlendFactor;
dstColorBlendFactor;
   VkBlendFactor
                            colorBlendOp;
   VkBlend0p
                       srcAlphaBlendFactor;
dstAlphaBlendFactor;
   VkBlendFactor
   VkBlendFactor
   VkBlend0p
                             alphaBlendOp;
    VkColorComponentFlags colorWriteMask;
} VkPipelineColorBlendAttachmentState;
```

• blendEnable controls whether blending is enabled for the corresponding color attachment. If blending is not enabled, the source fragment's color for that attachment is passed through unmodified.

- srcColorBlendFactor selects which blend factor is used to determine the source factors S_r, S_g, S_b .
- dstColorBlendFactor selects which blend factor is used to determine the destination factors D_r, D_g, D_b .
- colorBlendOp selects which blend operation is used to calculate the RGB values to write to the color attachment.
- srcAlphaBlendFactor selects which blend factor is used to determine the source factor S_a.
- dstAlphaBlendFactor selects which blend factor is used to determine the destination factor D_a.
- alphaBlendOp selects which blend operation is use to calculate the alpha values to write to the color attachment.
- colorWriteMask is a bitmask selecting which of the R, G, B, and/or A components are enabled for writing, as
 described later in this chapter.

- srcColorBlendFactor must be a valid VkBlendFactor value
- dstColorBlendFactor must be a valid VkBlendFactor value
- colorBlendOp must be a valid VkBlendOp value
- srcAlphaBlendFactor must be a valid VkBlendFactor value
- dstAlphaBlendFactor must be a valid VkBlendFactor value
- alphaBlendOp must be a valid VkBlendOp value
- colorWriteMask must be a valid combination of VkColorComponentFlagBits values
- If the dual source blending feature is not enabled, <code>srcColorBlendFactor</code> must not be <code>VK_BLEND_FACTOR_SRC1_COLOR</code>, <code>VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR</code>, <code>VK_BLEND_FACTOR_SRC1_ALPHA</code>, or <code>VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA</code>
- If the dual source blending feature is not enabled, <code>dstColorBlendFactor</code> must not be VK_BLEND_FACTOR_ SRC1_COLOR, VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR, VK_BLEND_FACTOR_SRC1_ALPHA, or VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA
- If the dual source blending feature is not enabled, <code>srcAlphaBlendFactor</code> must not be <code>VK_BLEND_FACTOR_SRC1_COLOR</code>, <code>VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR</code>, <code>VK_BLEND_FACTOR_SRC1_ALPHA</code>, or <code>VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA</code>
- If the dual source blending feature is not enabled, <code>dstAlphaBlendFactor</code> must not be VK_BLEND_FACTOR_ SRC1_COLOR, VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR, VK_BLEND_FACTOR_SRC1_ALPHA, or VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA

26.1.1 Blend Factors

The source and destination color and alpha blending factors are selected from the enum:

```
typedef enum VkBlendFactor {
   VK\_BLEND\_FACTOR\_ZERO = 0,
   VK_BLEND_FACTOR_ONE = 1,
   VK\_BLEND\_FACTOR\_SRC\_COLOR = 2,
   VK_BLEND_FACTOR_ONE_MINUS_SRC_COLOR = 3,
   VK\_BLEND\_FACTOR\_DST\_COLOR = 4,
   VK_BLEND_FACTOR_ONE_MINUS_DST_COLOR = 5,
   VK_BLEND_FACTOR_SRC_ALPHA = 6,
   VK_BLEND_FACTOR_ONE_MINUS_SRC_ALPHA = 7,
   VK_BLEND_FACTOR_DST_ALPHA = 8,
   VK_BLEND_FACTOR_ONE_MINUS_DST_ALPHA = 9,
   VK_BLEND_FACTOR_CONSTANT_COLOR = 10,
   VK_BLEND_FACTOR_ONE_MINUS_CONSTANT_COLOR = 11,
   VK_BLEND_FACTOR_CONSTANT_ALPHA = 12,
   VK_BLEND_FACTOR_ONE_MINUS_CONSTANT_ALPHA = 13,
   VK_BLEND_FACTOR_SRC_ALPHA_SATURATE = 14,
   VK_BLEND_FACTOR_SRC1_COLOR = 15,
   VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR = 16,
   VK_BLEND_FACTOR_SRC1_ALPHA = 17,
   VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA = 18,
} VkBlendFactor;
```

The semantics of each enum value is described in the table below:

Table 26.1: Blend Factors

VkBlendFactor	RGB Blend Factors	Alpha Blend
	(S_r, S_g, S_b) or (D_r, D_g, D_b)	Factor (S _a or
		D_a)
VK_BLEND_FACTOR_ZERO	(0,0,0)	0
VK_BLEND_FACTOR_ONE	(1,1,1)	1
VK_BLEND_FACTOR_SRC_COLOR	(R_{s0}, G_{s0}, B_{s0})	A_{s0}
VK_BLEND_FACTOR_ONE_MINUS_SRC_COLOR	$(1-R_{s0},1-G_{s0},1-B_{s0})$	$1 - A_{s0}$
VK_BLEND_FACTOR_DST_COLOR	(R_d, G_d, B_d)	A_d
VK_BLEND_FACTOR_ONE_MINUS_DST_COLOR	$(1-R_d, 1-G_d, 1-B_d)$	$1-A_d$
VK_BLEND_FACTOR_SRC_ALPHA	(A_{s0},A_{s0},A_{s0})	A_{s0}
VK_BLEND_FACTOR_ONE_MINUS_SRC_ALPHA	$(1-A_{s0},1-A_{s0},1-A_{s0})$	$1 - A_{s0}$
VK_BLEND_FACTOR_DST_ALPHA	(A_d, A_d, A_d)	A_d
VK_BLEND_FACTOR_ONE_MINUS_DST_ALPHA	$(1-A_d, 1-A_d, 1-A_d)$	$1-A_d$
VK_BLEND_FACTOR_CONSTANT_COLOR	(R_c,G_c,B_c)	A_c
VK_BLEND_FACTOR_ONE_MINUS_CONSTANT_COLOR	$(1-R_c, 1-G_c, 1-B_c)$	$1-A_c$
VK_BLEND_FACTOR_CONSTANT_ALPHA	(A_c, A_c, A_c)	A_c
VK_BLEND_FACTOR_ONE_MINUS_CONSTANT_ALPHA	$(1-A_c, 1-A_c, 1-A_c)$	$1-A_c$
VK_BLEND_FACTOR_SRC_ALPHA_SATURATE	$(f,f,f); f = \min(A_{s0}, 1 - A_d)$	1
VK_BLEND_FACTOR_SRC1_COLOR	(R_{s1},G_{s1},B_{s1})	A_{s1}
VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR	$(1-R_{s1},1-G_{s1},1-B_{s1})$	$1-A_{s1}$
VK_BLEND_FACTOR_SRC1_ALPHA	(A_{s1},A_{s1},A_{s1})	A_{s1}
VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA	$(1-A_{s1},1-A_{s1},1-A_{s1})$	$1-A_{s1}$

In this table, the following conventions are used:

- R_{s0} , G_{s0} , B_{s0} and A_{s0} represent the first source color R, G, B, and A components, respectively, for the fragment output location corresponding to the color attachment being blended.
- R_{s1} , G_{s1} , B_{s1} and A_{s1} represent the second source color R, G, B, and A components, respectively, used in dual source blending modes, for the fragment output location corresponding to the color attachment being blended.
- R_d , G_d , B_d and A_d represent the R, G, B, and A components of the destination color. That is, the color currently in the corresponding color attachment for this fragment/sample.
- R_c , G_c , B_c and A_c represent the blend constant R, G, B, and A components, respectively.

If the pipeline state object is created without the VK_DYNAMIC_STATE_BLEND_CONSTANTS dynamic state enabled then the "blend constant" (R_c, G_c, B_c, A_c) is specified via the blendConstants member of VkPipelineColorBlendStateCreateInfo.

Otherwise, to dynamically set and change the blend constant, call:

- commandBuffer is the command buffer into which the command will be recorded.
- blendConstants is an array of four values specifying the R, G, B, and A components of the blend constant color used in blending, depending on the blend factor.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics operations
- The currently bound graphics pipeline must have been created with the VK_DYNAMIC_STATE_BLEND_ CONSTANTS dynamic state enabled

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		

26.1.2 Dual-Source Blending

Blend factors that use the secondary color input $(R_{s1}, G_{s1}, B_{s1}, A_{s1})$ (VK_BLEND_FACTOR_SRC1_COLOR, VK_BLEND_FACTOR_SRC1_ALPHA, and VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR, VK_BLEND_FACTOR_SRC1_ALPHA, and VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA) may consume hardware resources that could otherwise be used for rendering to multiple color attachments. Therefore, the number of color attachments that can be used in a framebuffer may be lower when using dual-source blending.

Dual-source blending is only supported if the dualSrcBlend feature is enabled.

The maximum number of color attachments that can be used in a subpass when using dual-source blending functions is implementation-dependent and is reported as the <code>maxFragmentDualSrcAttachments</code> member of <code>VkPhysicalDeviceLimits</code>.

When using a fragment shader with dual-source blending functions, the color outputs are bound to the first and second inputs of the blender using the **Index** decoration, as described in Fragment Output Interface. If the second color input to the blender is not written in the shader, or if no output is bound to the second input of a blender, the result of the blending operation is not defined.

26.1.3 Blend Operations

Once the source and destination blend factors have been selected, they along with the source and destination components are passed to the blending operation. The blending operations are selected from the following enum, with RGB and alpha components potentially using different blend operations:

```
typedef enum VkBlendOp {
    VK_BLEND_OP_ADD = 0,
    VK_BLEND_OP_SUBTRACT = 1,
    VK_BLEND_OP_REVERSE_SUBTRACT = 2,
    VK_BLEND_OP_MIN = 3,
    VK_BLEND_OP_MAX = 4,
} VkBlendOp;
```

The semantics of each enum value is described in the table below:

VkBlendOp	RGB Components	Alpha Component
VK_BLEND_OP_ADD	$R = R_{s0} \times S_r + R_d \times D_r$	$A = A_{s0} \times S_a + A_d \times D_a$
	$G = G_{s0} \times S_g + G_d \times D_g$	
	$B = B_{s0} \times S_b + B_d \times D_b$	
VK_BLEND_OP_SUBTRACT	$R = R_{s0} \times S_r - R_d \times D_r$	$A = A_{s0} \times S_a - A_d \times D_a$
	$G = G_{s0} \times S_g - G_d \times D_g$	
	$B = B_{s0} \times S_b - B_d \times D_b$	
VK_BLEND_OP_REVERSE_SUBTRACT	$R = R_d \times D_r - R_{s0} \times S_r$	$A = A_d \times D_a - A_{s0} \times S_a$
	$G = G_d \times D_g - G_{s0} \times S_g$	
	$B = B_d \times D_b - B_{s0} \times S_b$	
VK_BLEND_OP_MIN	$R = \min(R_{s0}, R_d)$	$A = \min(A_{s0}, A_d)$
	$G = \min(G_{s0}, G_d)$	
	$B = \min(B_{s0}, B_d)$	
VK_BLEND_OP_MAX	$R = \max(R_{s0}, R_d)$	$A = \max(A_{s0}, A_d)$
	$G = \max(G_{s0}, G_d)$	
	$B = \max(B_{s0}, B_d)$	

Table 26.2: Blend Operations

In this table, the following conventions are used:

- R_{s0} , G_{s0} , B_{s0} and A_{s0} represent the first source color R, G, B, and A components, respectively.
- R_d , G_d , B_d and A_d represent the R, G, B, and A components of the destination color. That is, the color currently in the corresponding color attachment for this fragment/sample.
- S_r, S_g, S_b and S_a represent the source blend factor R, G, B, and A components, respectively.
- D_r, D_g, D_b and D_a represent the destination blend factor R, G, B, and A components, respectively.

The blending operation produces a new set of values R, G, B and A, which are written to the framebuffer attachment. If blending is not enabled for this attachment, then R, G, B and A are assigned R_{s0} , G_{s0} , B_{s0} and A_{s0} , respectively.

If the color attachment is fixed-point, the components of the source and destination values and blend factors are each clamped to [0,1] or [-1,1] respectively for an unsigned normalized or signed normalized color attachment prior to evaluating the blend operations. If the color attachment is floating-point, no clamping occurs.

The colorWriteMask member of VkPipelineColorBlendAttachmentState determines whether the final color values R, G, B and A are written to the framebuffer attachment. colorWriteMask is any combination of the following bits:

```
typedef enum VkColorComponentFlagBits {
    VK_COLOR_COMPONENT_R_BIT = 0x00000001,
    VK_COLOR_COMPONENT_G_BIT = 0x00000002,
    VK_COLOR_COMPONENT_B_BIT = 0x00000004,
    VK_COLOR_COMPONENT_A_BIT = 0x00000008,
} VkColorComponentFlagBits;
```

If VK_COLOR_COMPONENT_R_BIT is set, then the *R* value is written to color attachment for the appropriate sample, otherwise the value in memory is unmodified. The VK_COLOR_COMPONENT_B_BIT, VK_COLOR_COMPONENT_B_

BIT, and VK_COLOR_COMPONENT_A_BIT bits similarly control writing of the G, B, and A values. The colorWriteMask is applied regardless of whether blending is enabled.

If the numeric format of a framebuffer attachment uses sRGB encoding, the R, G, and B destination color values (after conversion from fixed-point to floating-point) are considered to be encoded for the sRGB color space and hence are linearized prior to their use in blending. Each R, G, and B component is converted from nonlinear to linear as described in the "KHR_DF_TRANSFER_SRGB" section of the Khronos Data Format Specification. If the format is not sRGB, no linearization is performed.

If the numeric format of a framebuffer attachment uses sRGB encoding, then the final R, G and B values are converted into the nonlinear sRGB representation before being written to the framebuffer attachment as described in the "KHR_DF_TRANSFER_SRGB" section of the Khronos Data Format Specification.

If the framebuffer color attachment numeric format is not sRGB encoded then the resulting c_s values for R, G and B are unmodified. The value of A is never sRGB encoded. That is, the alpha component is always stored in memory as linear.

If the framebuffer color attachment is VK_ATTACHMENT_UNUSED, no writes are performed through that attachment. Framebuffer color attachments greater than or equal to VkSubpassDescription::colorAttachmentCount perform no writes.

26.2 Logical Operations

The application can enable a *logical operation* between the fragment's color values and the existing value in the framebuffer attachment. This logical operation is applied prior to updating the framebuffer attachment. Logical operations are applied only for signed and unsigned integer and normalized integer framebuffers. Logical operations are not applied to floating-point or sRGB format color attachments.

Logical operations are controlled by the <code>logicOpEnable</code> and <code>logicOp</code> members of <code>VkPipelineColorBlendStateCreateInfo</code>. If <code>logicOpEnable</code> is <code>VK_TRUE</code>, then a logical operation selected by <code>logicOp</code> is applied between each color attachment and the fragment's corresponding output value, and blending of all attachments is treated as if it were disabled. Any attachments using color formats for which logical operations are not supported simply pass through the color values unmodified. The logical operation is applied independently for each of the red, green, blue, and alpha components. The <code>logicOp</code> is selected from the following operations:

```
typedef enum VkLogicOp {
   VK\_LOGIC\_OP\_CLEAR = 0,
   VK\_LOGIC\_OP\_AND = 1,
    VK LOGIC OP AND REVERSE = 2,
    VK\_LOGIC\_OP\_COPY = 3,
    VK_LOGIC_OP_AND_INVERTED = 4,
    VK_LOGIC_OP_NO_OP = 5,
    VK\_LOGIC\_OP\_XOR = 6,
    VK\_LOGIC\_OP\_OR = 7,
    VK\_LOGIC\_OP\_NOR = 8,
    VK LOGIC OP EQUIVALENT = 9,
    VK_LOGIC_OP_INVERT = 10,
    VK_LOGIC_OP_OR_REVERSE = 11,
    VK_LOGIC_OP_COPY_INVERTED = 12,
    VK_LOGIC_OP_OR_INVERTED = 13,
    VK_LOGIC_OP_NAND = 14,
    VK LOGIC OP SET = 15,
 VkLogicOp;
```

The logical operations supported by Vulkan are summarized in the following table in which

- ¬ is bitwise invert,
- \wedge is bitwise and,
- \vee is bitwise or,
- \oplus is bitwise exclusive or,
- s is the fragment's R_{s0} , G_{s0} , B_{s0} or A_{s0} component value for the fragment output corresponding to the color attachment being updated, and
- d is the color attachment's R, G, B or A component value:

Table 26.3: Logical Operations

Mode	Operation
VK_LOGIC_OP_CLEAR	0
VK_LOGIC_OP_AND	$s \wedge d$
VK_LOGIC_OP_AND_REVERSE	$s \wedge \neg d$
VK_LOGIC_OP_COPY	S
VK_LOGIC_OP_AND_INVERTED	$\neg s \wedge d$
VK_LOGIC_OP_NO_OP	d
VK_LOGIC_OP_XOR	$s \oplus d$
VK_LOGIC_OP_OR	$s \lor d$
VK_LOGIC_OP_NOR	$\neg(s \lor d)$
VK_LOGIC_OP_EQUIVALENT	$\neg (s \oplus d)$
VK_LOGIC_OP_INVERT	$\neg d$
VK_LOGIC_OP_OR_REVERSE	$s \lor \neg d$
VK_LOGIC_OP_COPY_INVERTED	$\neg s$
VK_LOGIC_OP_OR_INVERTED	$\neg s \lor d$
VK_LOGIC_OP_NAND	$\neg(s \land d)$
VK_LOGIC_OP_SET	all 1s

The result of the logical operation is then written to the color attachment as controlled by the component write mask, described in Blend Operations.

Chapter 27

Dispatching Commands

Dispatching commands (commands with **Dispatch** in the name) provoke work in a compute pipeline. Dispatching commands are recorded into a command buffer and when executed by a queue, will produce work which executes according to the currently bound compute pipeline. A compute pipeline must be bound to a command buffer before any dispatch commands are recorded in that command buffer.

To record a dispatch, call:

- commandBuffer is the command buffer into which the command will be recorded.
- x is the number of local workgroups to dispatch in the X dimension.
- y is the number of local workgroups to dispatch in the Y dimension.
- z is the number of local workgroups to dispatch in the Z dimension.

When the command is executed, a global workgroup consisting of $x \times y \times z$ local workgroups is assembled.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support compute operations
- This command must only be called outside of a render pass instance
- x must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[0]

- y must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[1]
- z must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[2]
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_COMPUTE, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_COMPUTE, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- A valid compute pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_COMPUTE, a push constant value must have been set for VK_PIPELINE_BIND_POINT_COMPUTE, with a VkPipelineLayout that is compatible for push constants with the one used to create the current VkPipeline, as described in Section 13.2.2.1
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_COMPUTE accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_COMPUTE accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties

• Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	COMPUTE
Secondary		

To record an indirect command dispatch, call:

- commandBuffer is the command buffer into which the command will be recorded.
- buffer is the buffer containing dispatch parameters.
- offset is the byte offset into buffer where parameters begin.

vkCmdDispatchIndirect behaves similarly to vkCmdDispatch except that the parameters are read by the device from a buffer during execution. The parameters of the dispatch are encoded in a VkDispatchIndirectCommand structure taken from buffer starting at offset.

Valid Usage

• commandBuffer must be a valid VkCommandBuffer handle

- buffer must be a valid VkBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support compute operations
- This command must only be called outside of a render pass instance
- Both of buffer, and commandBuffer must have been created, allocated, or retrieved from the same VkDevice
- For each set *n* that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_COMPUTE, a descriptor set must have been bound to *n* at VK_PIPELINE_BIND_POINT_COMPUTE, with a VkPipelineLayout that is compatible for set *n*, with the VkPipelineLayout used to create the current VkPipeline, as described in Section 13.2.2.1
- Descriptors in each bound descriptor set, specified via **vkCmdBindDescriptorSets**, must be valid if they are statically used by the currently bound VkPipeline object, specified via **vkCmdBindPipeline**
- A valid compute pipeline must be bound to the current command buffer with VK_PIPELINE_BIND_POINT_ COMPUTE
- buffer must have been created with the VK_BUFFER_USAGE_INDIRECT_BUFFER_BIT bit set
- offset must be a multiple of 4
- The sum of offset and the size of VkDispatchIndirectCommand must be less than or equal to the size of buffer
- For each push constant that is statically used by the VkPipeline currently bound to VK_PIPELINE_BIND_POINT_COMPUTE, a push constant value must have been set for VK_PIPELINE_BIND_POINT_COMPUTE, with a VkPipelineLayout that is compatible for push constants with the one used to create the current VkPipeline, as described in Section 13.2.2.1
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used to sample from any VkImage with a VkImageView of the type VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, VK_IMAGE_VIEW_TYPE_1D_ARRAY, VK_IMAGE_VIEW_TYPE_2D_ARRAY or VK_IMAGE_VIEW_TYPE_CUBE_ARRAY, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions with ImplicitLod, Dref or Proj in their name, in any shader stage
- If any VkSampler object that is accessed from a shader by the VkPipeline currently bound to VK_ PIPELINE_BIND_POINT_COMPUTE uses unnormalized coordinates, it must not be used with any of the SPIR-V OpImageSample* or OpImageSparseSample* instructions that includes a LOD bias or any offset values, in any shader stage
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_COMPUTE accesses a uniform buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set
- If the robust buffer access feature is not enabled, and any shader stage in the VkPipeline object currently bound to VK_PIPELINE_BIND_POINT_COMPUTE accesses a storage buffer, it must not access values outside of the range of that buffer specified in the currently bound descriptor set

- Any VkImageView being sampled with VK_FILTER_LINEAR as a result of this command must be of a format which supports linear filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must be of a format which supports cubic filtering, as specified by the VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG flag in VkFormatProperties::linearTilingFeatures (for a linear image) or VkFormatProperties::optimalTilingFeatures (for an optimally tiled image) returned by vkGetPhysicalDeviceFormatProperties
- Any VkImageView being sampled with VK_FILTER_CUBIC_IMG as a result of this command must not have a VkImageViewType of VK_IMAGE_VIEW_TYPE_3D, VK_IMAGE_VIEW_TYPE_CUBE, or VK_IMAGE_VIEW_TYPE_CUBE ARRAY

Host Synchronization

• Host access to commandBuffer must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Outside	COMPUTE
Secondary		

The VkDispatchIndirectCommand structure is defined as:

```
typedef struct VkDispatchIndirectCommand {
   uint32_t     x;
   uint32_t    y;
   uint32_t    z;
} VkDispatchIndirectCommand;
```

- x is the number of local workgroups to dispatch in the X dimension.
- y is the number of local workgroups to dispatch in the Y dimension.
- z is the number of local workgroups to dispatch in the Z dimension.

The members of $\mbox{VkDispatchIndirectCommand}$ structure have the same meaning as the similarly named parameters of $\mbox{vkCmdDispatch}$.

Valid Usage

- x must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[0]
- y must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[1]
- z must be less than or equal to VkPhysicalDeviceLimits::maxComputeWorkGroupCount[2]

Chapter 28

Sparse Resources

As documented in Resource Memory Association, VkBuffer and VkImage resources in Vulkan must be bound completely and contiguously to a single VkDeviceMemory object. This binding must be done before the resource is used, and the binding is immutable for the lifetime of the resource.

Sparse resources relax these restrictions and provide these additional features:

- Sparse resources can be bound non-contiguously to one or more VkDeviceMemory allocations.
- Sparse resources can be re-bound to different memory allocations over the lifetime of the resource.
- Sparse resources can have descriptors generated and used orthogonally with memory binding commands.

28.1 Sparse Resource Features

Sparse resources have several features that must be enabled explicitly at resource creation time. The features are enabled by including bits in the *flags* parameter of VkImageCreateInfo or VkBufferCreateInfo. Each feature also has one or more corresponding feature enables specified in VkPhysicalDeviceFeatures.

- Sparse binding is the base feature, and provides the following capabilities:
 - Resources can be bound at some defined (sparse block) granularity.
 - The entire resource must be bound to memory before use regardless of regions actually accessed.
 - No specific mapping of image region to memory offset is defined, i.e. the location that each texel corresponds to in memory is implementation-dependent.
 - Sparse buffers have a well-defined mapping of buffer range to memory range, where an offset into a range of the
 buffer that is bound to a single contiguous range of memory corresponds to an identical offset within that range of
 memory.
 - Requested via the VK_IMAGE_CREATE_SPARSE_BINDING_BIT and VK_BUFFER_CREATE_SPARSE_BINDING_BIT bits.
 - A sparse image created using VK_IMAGE_CREATE_SPARSE_BINDING_BIT (but not VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT) supports all formats that non-sparse usage supports, and supports both VK_IMAGE_TILING_OPTIMAL and VK_IMAGE_TILING_LINEAR tiling.
- Sparse Residency builds on (and requires) the sparseBinding feature. It includes the following capabilities:

- Resources do not have to be completely bound to memory before use on the device.
- Images have a prescribed sparse image block layout, allowing specific rectangular regions of the image to be bound to specific offsets in memory allocations.
- Consistency of access to unbound regions of the resource is defined by the absence or presence of
 VkPhysicalDeviceSparseProperties::residencyNonResidentStrict. If this property is present,
 accesses to unbound regions of the resource are well defined and behave as if the data bound is populated with all
 zeros; writes are discarded. When this property is absent, accesses are considered safe, but reads will return
 undefined values.
- Requested via the VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT and VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT bits.
- is advertised on a finer grain via the following features:
 - * sparseResidencyBuffer: Support for creating VkBuffer objects with the VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT.
 - * sparseResidencyImage2D: Support for creating 2D single-sampled VkImage objects with VK_IMAGE_ CREATE SPARSE RESIDENCY BIT.
 - * sparseResidencyImage3D: Support for creating 3D VkImage objects with VK_IMAGE_CREATE_SPARSE_ RESIDENCY_BIT.
 - * sparseResidency2Samples: Support for creating 2D VkImage objects with 2 samples and VK_IMAGE_ CREATE_SPARSE_RESIDENCY_BIT.
 - * sparseResidency4Samples: Support for creating 2D VkImage objects with 4 samples and VK_IMAGE_ CREATE SPARSE RESIDENCY BIT.
 - * sparseResidency8Samples: Support for creating 2D VkImage objects with 8 samples and VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.
 - * sparseResidency16Samples: Support for creating 2D VkImage objects with 16 samples and VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.

Implementations supporting <code>sparseResidencyImage2D</code> are only required to support sparse 2D, single-sampled images. Support is not required for sparse 3D and MSAA images and is enabled via <code>sparseResidencyImage3D</code>, <code>sparseResidency2Samples</code>, <code>sparseResidency4Samples</code>, <code>sparseResidency8Samples</code>, and <code>sparseResidency16Samples</code>.

- A sparse image created using VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT supports all non-compressed color formats with power-of-two texel size that non-sparse usage supports. Additional formats may also be supported and can be queried via vkGetPhysicalDeviceSparseImageFormatProperties. VK_IMAGE_TILING_LINEAR tiling is not supported.
- Sparse aliasing provides the following capability that can be enabled per resource:

Allows physical memory ranges to be shared between multiple locations in the same sparse resource or between multiple sparse resources, with each binding of a memory location observing a consistent interpretation of the memory contents.

See Sparse Memory Aliasing for more information.

28.2 Sparse Buffers and Fully-Resident Images

Both VkBuffer and VkImage objects created with the VK_IMAGE_CREATE_SPARSE_BINDING_BIT or VK_BUFFER_CREATE_SPARSE_BINDING_BIT bits can be thought of as a linear region of address space. In the VkImage case if VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT is not used, this linear region is entirely opaque, meaning that there is no application-visible mapping between pixel location and memory offset.

Unless VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT or VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT are also used, the entire resource must be bound to one or more VkDeviceMemory objects before use.

28.2.1 Sparse Buffer and Fully-Resident Image Block Size

The sparse block size in bytes for sparse buffers and fully-resident images is reported as VkMemoryRequirements::alignment.alignment represents both the memory alignment requirement and the binding granularity (in bytes) for sparse resources.

28.3 Sparse Partially-Resident Buffers

VkBuffer objects created with the VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT bit allow the buffer to be made only partially resident. Partially resident VkBuffer objects are allocated and bound identically to VkBuffer objects using only the VK_BUFFER_CREATE_SPARSE_BINDING_BIT feature. The only difference is the ability for some regions of the buffer to be unbound during device use.

28.4 Sparse Partially-Resident Images

VkImage objects created with the VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT bit allow specific rectangular regions of the image called sparse image blocks to be bound to specific ranges of memory. This allows the application to manage residency at either image subresource or sparse image block granularity. Each image subresource (outside of the mip tail) starts on a sparse block boundary and has dimensions that are integer multiples of the corresponding dimensions of the sparse image block.

Note



Applications can use these types of images to control level-of-detail based on total memory consumption. If memory pressure becomes an issue the application can unbind and disable specific mipmap levels of images without having to recreate resources or modify pixel data of unaffected levels.

The application can also use this functionality to access subregions of the image in a "megatexture" fashion. The application can create a large image and only populate the region of the image that is currently being used in the scene.

28.4.1 Accessing Unbound Regions

The following member of VkPhysicalDeviceSparseProperties affects how data in unbound regions of sparse resources are handled by the implementation:

• residencyNonResidentStrict

If this property is not present, reads of unbound regions of the image will return undefined values. Both reads and writes are still considered *safe* and will not affect other resources or populated regions of the image.

If this property is present, all reads of unbound regions of the image will behave as if the region was bound to memory populated with all zeros; writes will be discarded.

Formatted accesses to unbound memory may still alter some component values in the natural way for those accesses, e.g. substituting a value of one for alpha in formats that do not have an alpha component.

Example: Reading the alpha component of an unbacked VK_FORMAT_R8_UNORM image will return a value of 1.0f.

See Physical Device Enumeration for instructions for retrieving physical device properties.

Implementor's Note

For hardware that cannot natively handle access to unbound regions of a resource, the implementation may allocate and bind memory to the unbound regions. Reads and writes to unbound regions will access the implementation-managed memory instead of causing a hardware fault.

Given that reads of unbound regions are undefined in this scenario, implementations may use the same physical memory for unbound regions of multiple resources within the same process.

28.4.2 Mip Tail Regions

Sparse images created using VK_IMAGE_CREATE_SPARSE_BINDING_BIT (without also using VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT) have no specific mapping of image region or image subresource to memory offset defined, so the entire image can be thought of as a linear opaque address region. However, images created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT do have a prescribed sparse image block layout, and hence each image subresource must start on a sparse block boundary. Within each array layer, the set of mip levels that have a smaller size than the sparse block size in bytes are grouped together into a *mip tail region*.

If the VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT flag is present in the flags member of VkSparseImageFormatProperties, for the image's format, then any mip level which has dimensions that are not integer multiples of the corresponding dimensions of the sparse image block, and all subsequent mip levels, are also included in the mip tail region.

The following member of VkPhysicalDeviceSparseProperties may affect how the implementation places mip levels in the mip tail region:

• residencyAlignedMipSize

Each mip tail region is bound to memory as an opaque region (i.e. must be bound using a VkSparseImageOpaqueMemoryBindInfo structure) and may be of a size greater than or equal to the sparse block size in bytes. This size is guaranteed to be an integer multiple of the sparse block size in bytes.

An implementation may choose to allow each array-layer's mip tail region to be bound to memory independently or require that all array-layer's mip tail regions be treated as one. This is dictated by VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT in VkSparseImageMemoryRequirements::flags.

The following diagrams depict how VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT and VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT alter memory usage and requirements.

Mip Level 0 Mip Level 1 Mip Level 2 Mip Tail Mip Tail Mip Tail Sparse Memory Block

Arrayed Sparse Image

Figure 28.1: Sparse Image

In the absence of VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT and VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT, each array layer contains a mip tail region containing pixel data for all mip levels smaller than the sparse image block in any dimension.

Mip levels that are as large or larger than a sparse image block in all dimensions can be bound individually. Right-edges and bottom-edges of each level are allowed to have partially used sparse blocks. Any bound partially-used-sparse-blocks must still have their full sparse block size in bytes allocated in memory.

Arrayed Sparse Image

VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT



Figure 28.2: Sparse Image with Single Mip Tail

When $VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT$ is present all array layers will share a single mip tail region.

Arrayed Sparse Image

VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT



Figure 28.3: Sparse Image with Aligned Mip Size



Note

The mip tail regions are presented here in 2D arrays simply for figure size reasons. Each mip tail is logically a single array of sparse blocks with an implementation-dependent mapping of pixels to sparse blocks.

When VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT is present the first mip level that would contain partially used sparse blocks begins the mip tail region. This level and all subsequent levels are placed in the mip tail. Only the first *N* mip levels whose dimensions are an exact multiple of the sparse image block dimensions can be bound and unbound on a sparse block basis.

Arrayed Sparse Image

VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT

VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT



Figure 28.4: Sparse Image with Aligned Mip Size and Single Mip Tail



Note

The mip tail region is presented here in a 2D array simply for figure size reasons. It is logically a single array of sparse blocks with an implementation-dependent mapping of pixels to sparse blocks.

When both VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT and VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT are present the constraints from each of these flags are in effect.

28.4.3 Standard Sparse Image Block Shapes

Standard sparse image block shapes define a standard set of dimensions for sparse image blocks that depend on the format of the image. Layout of pixels within a sparse image block is implementation dependent. All currently defined standard sparse image block shapes are 64 KB in size.

For block-compressed formats (e.g. VK_FORMAT_BC5_UNORM_BLOCK), the pixel size is the size of the compressed texel block (128-bit for BC5) thus the dimensions of the standard sparse image block shapes apply in terms of compressed texel blocks.



Note

For block-compressed formats, the dimensions of a sparse image block in terms of texels can be calculated by multiplying the sparse image block dimensions by the compressed texel block dimensions.

Table 28.1: Standard Sparse Image Block Shapes (Single Sample)

PIXEL SIZE (bits)	Block Shape (2D)	Block Shape (3D)
8-Bit	$256 \times 256 \times 1$	$64 \times 32 \times 32$
16-Bit	$256 \times 128 \times 1$	$32 \times 32 \times 32$
32-Bit	$128 \times 128 \times 1$	$32 \times 32 \times 16$
64-Bit	$128 \times 64 \times 1$	$32 \times 16 \times 16$
128-Bit	$64 \times 64 \times 1$	$16 \times 16 \times 16$

Table 28.2: Standard Sparse Image Block Shapes (MSAA)

PIXEL SIZE (bits)	Block Shape (2X)	Block Shape (4X)	Block Shape (8X)	Block Shape (16X)
8-Bit	$128 \times 256 \times 1$	$128 \times 128 \times 1$	64 × 128 × 1	$64 \times 64 \times 1$
16-Bit	$128 \times 128 \times 1$	$128 \times 64 \times 1$	$64 \times 64 \times 1$	$64 \times 32 \times 1$
32-Bit	$64 \times 128 \times 1$	$64 \times 64 \times 1$	$32 \times 64 \times 1$	$32 \times 32 \times 1$
64-Bit	$64 \times 64 \times 1$	$64 \times 32 \times 1$	$32 \times 32 \times 1$	$32 \times 16 \times 1$
128-Bit	$32 \times 64 \times 1$	$32 \times 32 \times 1$	$16 \times 32 \times 1$	$16 \times 16 \times 1$

Implementations that support the standard sparse image block shape for all applicable formats may advertise the following VkPhysicalDeviceSparseProperties:

- residencyStandard2DBlockShape
- residencyStandard2DMultisampleBlockShape
- residencyStandard3DBlockShape

Reporting each of these features does *not* imply that all possible image types are supported as sparse. Instead, this indicates that no supported sparse image of the corresponding type will use custom sparse image block dimensions for any formats that have a corresponding standard sparse image block shape.

28.4.4 Custom Sparse Image Block Shapes

An implementation that does not support a standard image block shape for a particular sparse partially-resident image may choose to support a custom sparse image block shape for it instead. The dimensions of such a custom sparse image block shape are reported in VkSparseImageFormatProperties::imageGranularity. As with standard sparse image block shapes, the size in bytes of the custom sparse image block shape will be reported in VkMemoryRequirements::alignment.

Custom sparse image block dimensions are reported through

 $\label{lem:condition} \mbox{\bf vkGetPhysicalDeviceSparseImageFormatProperties} \ \mbox{and} \\ \mbox{\bf vkGetImageSparseMemoryRequirements}.$

An implementation must not support both the standard sparse image block shape and a custom sparse image block shape for the same image. The standard sparse image block shape must be used if it is supported.

28.4.5 Multiple Aspects

Partially resident images are allowed to report separate sparse properties for different aspects of the image. One example is for depth/stencil images where the implementation separates the depth and stencil data into separate planes. Another reason for multiple aspects is to allow the application to manage memory allocation for implementation-private *metadata* associated with the image. See the figure below:

Multiple Aspect Sparse Image

Mip Level 0 Mip Level 1 Mip Level 2 Mip Level 3 Mip Tail Mip Tail Stencil Aspect Metadata Aspect Mip Tail Mip Tail Mip Tail Image Pixel Data Sparse Memory Block

Figure 28.5: Multiple Aspect Sparse Image



Note

The mip tail regions are presented here in 2D arrays simply for figure size reasons. Each mip tail is logically a single array of sparse blocks with an implementation-dependent mapping of pixels to sparse blocks.

In the figure above the depth, stencil, and metadata aspects all have unique sparse properties. The per-pixel stencil data is $^{1}/_{4}$ the size of the depth data, hence the stencil sparse blocks include 4x the number of pixels. The sparse block size in bytes for all of the aspects is identical and defined by VkMemoryRequirements::alignment.

28.4.5.1 Metadata

The metadata aspect of an image has the following constraints:

- All metadata is reported in the mip tail region of the metadata aspect.
- All metadata must be bound prior to device use of the sparse image.

28.5 Sparse Memory Aliasing

By default sparse resources have the same aliasing rules as non-sparse resources. See Memory Aliasing for more information.

VkDevice objects that have the sparseResidencyAliased feature enabled are able to use the VK_BUFFER_CREATE_SPARSE_ALIASED_BIT and VK_IMAGE_CREATE_SPARSE_ALIASED_BIT flags for resource creation. These flags allow resources to access physical memory bound into multiple locations within one or more sparse resources in a *data consistent* fashion. This means that reading physical memory from multiple aliased locations will return the same value.

Care must be taken when performing a write operation to aliased physical memory. Memory dependencies must be used to separate writes to one alias from reads or writes to another alias. Writes to aliased memory that are not properly guarded against accesses to different aliases will have undefined results for all accesses to the aliased memory.

Applications that wish to make use of data consistent sparse memory aliasing must abide by the following guidelines:

- All sparse resources that are bound to aliased physical memory must be created with the VK_BUFFER_CREATE_ SPARSE_ALIASED_BIT / VK_IMAGE_CREATE_SPARSE_ALIASED_BIT flag.
- All resources that access aliased physical memory must interpret the memory in the same way. This implies the following:
 - Buffers and images cannot alias the same physical memory in a data consistent fashion. The physical memory ranges must be used exclusively by buffers or used exclusively by images for data consistency to be guaranteed.
 - Memory in sparse image mip tail regions cannot access aliased memory in a data consistent fashion.
 - Sparse images that alias the same physical memory must have compatible formats and be using the same sparse image block shape in order to access aliased memory in a data consistent fashion.

Failure to follow any of the above guidelines will require the application to abide by the normal, non-sparse resource aliasing rules. In this case memory cannot be accessed in a data consistent fashion.



Note

Enabling sparse resource memory aliasing can be a way to lower physical memory use, but it may reduce performance on some implementations. An application developer can test on their target HW and balance the memory / performance trade-offs measured.

28.6 Sparse Resource Implementation Guidelines

This section is Informative. It is included to aid in implementors' understanding of sparse resources.

Device Virtual Address The basic <code>sparseBinding</code> feature allows the resource to reserve its own device virtual address range at resource creation time rather than relying on a bind operation to set this. Without any other creation flags, no other constraints are relaxed compared to normal resources. All pages must be bound to physical memory before the device accesses the resource.

The sparse residency features allow sparse resources to be used even when not all pages are bound to memory. Hardware that supports access to unbound pages without causing a fault may support residencyNonResidentStrict.

Not faulting on access to unbound pages is not enough to support <code>sparseResidencyNonResidentStrict</code>. An implementation must also guarantee that reads after writes to unbound regions of the resource always return data for the read as if the memory contains zeros. Depending on the cache implementation of the hardware this may not always be possible.

Hardware that does not fault, but does not guarantee correct read values will not require dummy pages, but also must not support <code>sparseResidencyNonResidentStrict</code>.

Hardware that cannot access unbound pages without causing a fault will require the implementation to bind the entire device virtual address range to physical memory. Any pages that the application does not bind to memory may be bound to one (or more) "dummy" physical page(s) allocated by the implementation. Given the following properties:

- A process must not access memory from another process
- · Reads return undefined values

It is sufficient for each host process to allocate these dummy pages and use them for all resources in that process. Implementations may allocate more often (per instance, per device, or per resource).

Binding Memory The byte size reported in VkMemoryRequirements::size must be greater than or equal to the amount of physical memory required to fully populate the resource. Some hardware requires "holes" in the device virtual address range that are never accessed. These holes may be included in the size reported for the resource.

Including or not including the device virtual address holes in the resource size will alter how the implementation provides support for VkSparseImageOpaqueMemoryBindInfo. This operation must be supported for all sparse images, even ones created with VK IMAGE CREATE SPARSE RESIDENCY BIT.

- If the holes are included in the size, this bind function becomes very easy. In most cases the resourceOffset is simply a device virtual address offset and the implementation does not require any sophisticated logic to determine what device virtual address to bind. The cost is that the application can allocate more physical memory for the resource than it needs.
- If the holes are not included in the size, the application can allocate less physical memory than otherwise for the resource. However, in this case the implementation must account for the holes when mapping resourceOffset to the actual device virtual address intended to be mapped.



Note

If the application always uses <code>VkSparseImageMemoryBindInfo</code> to bind memory for the non-tail mip levels, any holes that are present in the resource size may never be bound.

Since VkSparseImageMemoryBindInfo uses pixel locations to determine which device virtual addresses to bind, it is impossible to bind device virtual address holes with this operation.

Binding Metadata Memory All metadata for sparse images have their own sparse properties and are embedded in the mip tail region for said properties. See the Multiaspect section for details.

Given that metadata is in a mip tail region, and the mip tail region must be reported as contiguous (either globally or per-array-layer), some implementations will have to resort to complicated offset \rightarrow device virtual address mapping for handling VkSparseImageOpaqueMemoryBindInfo.

To make this easier on the implementation, the VK_SPARSE_MEMORY_BIND_METADATA_BIT explicitly denotes when metadata is bound with VkSparseImageOpaqueMemoryBindInfo. When this flag is not present, the resourceOffset may be treated as a strict device virtual address offset.

When VK_SPARSE_MEMORY_BIND_METADATA_BIT is present, the resourceOffset must have been derived explicitly from the <code>imageMipTailOffset</code> in the sparse resource properties returned for the metadata aspect. By manipulating the value returned for <code>imageMipTailOffset</code>, the <code>resourceOffset</code> does not have to correlate directly to a device virtual address offset, and may instead be whatever values makes it easiest for the implementation to derive the correct device virtual address.

28.7 Sparse Resource API

The APIs related to sparse resources are grouped into the following categories:

- Physical Device Features
- Physical Device Sparse Properties
- Sparse Image Format Properties
- Sparse Resource Creation
- Sparse Resource Memory Requirements
- Binding Resource Memory

28.7.1 Physical Device Features

Some sparse-resource related features are reported and enabled in VkPhysicalDeviceFeatures. These features must be supported and enabled on the VkDevice object before applications can use them. See Physical Device Features for information on how to get and set enabled device features, and for more detailed explanations of these features.

28.7.1.1 Sparse Physical Device Features

- sparseBinding: Support for creating VkBuffer and VkImage objects with the VK_BUFFER_CREATE_ SPARSE BINDING BIT and VK IMAGE CREATE SPARSE BINDING BIT flags, respectively.
- sparseResidencyBuffer: Support for creating VkBuffer objects with the VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT flag.
- sparseResidencyImage2D: Support for creating 2D single-sampled VkImage objects with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.
- sparseResidencyImage3D: Support for creating 3D VkImage objects with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.
- sparseResidency2Samples: Support for creating 2D VkImage objects with 2 samples and VK_IMAGE_CREATE SPARSE RESIDENCY BIT.
- sparseResidency4Samples: Support for creating 2D VkImage objects with 4 samples and VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.

- sparseResidency8Samples: Support for creating 2D VkImage objects with 8 samples and VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.
- sparseResidency16Samples: Support for creating 2D VkImage objects with 16 samples and VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT.
- sparseResidencyAliased: Support for creating VkBuffer and VkImage objects with the VK_BUFFER_ CREATE_SPARSE_ALIASED_BIT and VK_IMAGE_CREATE_SPARSE_ALIASED_BIT flags, respectively.

28.7.2 Physical Device Sparse Properties

Some features of the implementation are not possible to disable, and are reported to allow applications to alter their sparse resource usage accordingly. These read-only capabilities are reported in the

VkPhysicalDeviceProperties::sparseProperties member, which is a structure of type VkPhysicalDeviceSparseProperties.

The VkPhysicalDeviceSparseProperties structure is defined as:

```
typedef struct VkPhysicalDeviceSparseProperties {
   VkBool32    residencyStandard2DBlockShape;
   VkBool32    residencyStandard2DMultisampleBlockShape;
   VkBool32    residencyStandard3DBlockShape;
   VkBool32    residencyAlignedMipSize;
   VkBool32    residencyNonResidentStrict;
} VkPhysicalDeviceSparseProperties;
```

- residencyStandard2DBlockShape is VK_TRUE if the physical device will access all single-sample 2D sparse resources using the standard sparse image block shapes (based on image format), as described in the Standard Sparse Image Block Shapes (Single Sample) table. If this property is not supported the value returned in the imageGranularity member of the VkSparseImageFormatProperties structure for single-sample 2D images is not required to match the standard sparse image block dimensions listed in the table.
- residencyStandard2DMultisampleBlockShape is VK_TRUE if the physical device will access all multisample 2D sparse resources using the standard sparse image block shapes (based on image format), as described in the Standard Sparse Image Block Shapes (MSAA) table. If this property is not supported, the value returned in the imageGranularity member of the VkSparseImageFormatProperties structure for multisample 2D images is not required to match the standard sparse image block dimensions listed in the table.
- residencyStandard3DBlockShape is VK_TRUE if the physical device will access all 3D sparse resources using the standard sparse image block shapes (based on image format), as described in the Standard Sparse Image Block Shapes (Single Sample) table. If this property is not supported, the value returned in the imageGranularity member of the VkSparseImageFormatProperties structure for 3D images is not required to match the standard sparse image block dimensions listed in the table.
- residencyAlignedMipSize is VK_TRUE if images with mip level dimensions that are not integer multiples of the corresponding dimensions of the sparse image block may be placed in the mip tail. If this property is not reported, only mip levels with dimensions smaller than the imageGranularity member of the VkSparseImageFormatProperties structure will be placed in the mip tail. If this property is reported the implementation is allowed to return VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT in the flags member of VkSparseImageFormatProperties, indicating that mip level dimensions that are not integer multiples of the corresponding dimensions of the sparse image block will be placed in the mip tail.
- residencyNonResidentStrict specifies whether the physical device can consistently access non-resident regions of a resource. If this property is VK_TRUE, access to non-resident regions of resources will be guaranteed to return values as if the resource were populated with 0; writes to non-resident regions will be discarded.

28.7.3 Sparse Image Format Properties

Given that certain aspects of sparse image support, including the sparse image block dimensions, may be implementation-dependent, vkGetPhysicalDeviceSparseImageFormatProperties can be used to query for sparse image format properties prior to resource creation. This command is used to check whether a given set of sparse image parameters is supported and what the sparse image block shape will be.

28.7.3.1 Sparse Image Format Properties API

The VkSparseImageFormatProperties structure is defined as:

- aspectMask is a bitmask of VkImageAspectFlagBits specifying which aspects of the image the properties apply to.
- imageGranularity is the width, height, and depth of the sparse image block in texels or compressed texel blocks.
- flags is a bitmask specifying additional information about the sparse resource. Bits which can be set include:

```
typedef enum VkSparseImageFormatFlagBits {
   VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT = 0x00000001,
   VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT = 0x00000002,
   VK_SPARSE_IMAGE_FORMAT_NONSTANDARD_BLOCK_SIZE_BIT = 0x00000004,
} VkSparseImageFormatFlagBits;
```

- If VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT is set, the image uses a single mip tail region for all array layers.
- If VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT is set, the first mip level whose dimensions are not integer multiples of the corresponding dimensions of the sparse image block begins the mip tail region.
- If VK_SPARSE_IMAGE_FORMAT_NONSTANDARD_BLOCK_SIZE_BIT is set, the image uses non-standard sparse image block dimensions, and the <code>imageGranularity</code> values do not match the standard sparse image block dimensions for the given pixel format.

vkGetPhysicalDeviceSparseImageFormatProperties returns an array of

VkSparseImageFormatProperties. Each element will describe properties for one set of image aspects that are bound simultaneously in the image. This is usually one element for each aspect in the image, but for interleaved depth/stencil images there is only one element describing the combined aspects.

```
void vkGetPhysicalDeviceSparseImageFormatProperties(
                                                 physicalDevice,
   VkPhysicalDevice
   VkFormat
                                                 format,
   VkImageType
                                                 type,
   VkSampleCountFlagBits
                                                  samples,
   VkImageUsageFlags
                                                 usage,
   VkImageTiling
                                                 tiling,
   uint32_t*
                                                 pPropertyCount,
   VkSparseImageFormatProperties*
                                                 pProperties);
```

- physicalDevice is the physical device from which to query the sparse image capabilities.
- format is the image format.
- type is the dimensionality of image.
- samples is the number of samples per pixel as defined in VkSampleCountFlagBits.
- usage is a bitmask describing the intended usage of the image.
- tiling is the tiling arrangement of the data elements in memory.
- pPropertyCount is a pointer to an integer related to the number of sparse format properties available or queried, as described below.
- pProperties is either NULL or a pointer to an array of VkSparseImageFormatProperties structures.

If pProperties is NULL, then the number of sparse format properties available is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If pPropertyCount is less than the number of sparse format properties available, at most pPropertyCount structures will be written.

If VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT is not supported for the given arguments, pPropertyCount will be set to zero upon return, and no data will be written to pProperties.

Multiple aspects are returned for depth/stencil images that are implemented as separate planes by the implementation. The depth and stencil data planes each have unique VkSparseImageFormatProperties data.

Depth/stencil images with depth and stencil data interleaved into a single plane will return a single VkSparseImageFormatProperties structure with the aspectMask set to VK_IMAGE_ASPECT_DEPTH_BIT | VK_IMAGE_ASPECT_STENCIL_BIT.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- format must be a valid VkFormat value
- type must be a valid VkImageType value
- samples must be a valid VkSampleCountFlagBits value
- usage must be a valid combination of VkImageUsageFlagBits values
- usage must not be 0
- tiling must be a valid VkImageTiling value
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkSparseImageFormatProperties structures
- samples must be a bit value that is set in VkImageFormatProperties::sampleCounts returned by vkGetPhysicalDeviceImageFormatProperties with format, type, tiling, and usage equal to those in this command and flags equal to the value that is set in sname::VkImageCreateInfo::pname::flags when the image is created

28.7.4 Sparse Resource Creation

Sparse resources require that one or more sparse feature flags be specified (as part of the VkPhysicalDeviceFeatures structure described previously in the Physical Device Features section) at CreateDevice time. When the appropriate device features are enabled, the VK_BUFFER_CREATE_SPARSE_* and VK_IMAGE_CREATE_SPARSE_* flags can be used. See vkCreateBuffer and vkCreateImage for details of the resource creation APIs.

Note



Specifying VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT or VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT requires specifying VK_BUFFER_CREATE_SPARSE_BINDING_BIT or VK_IMAGE_CREATE_SPARSE_BINDING_BIT, respectively, as well. This means that resources must be created with the appropriate *_SPARSE_BINDING_BIT to be used with the sparse binding command (vkQueueBindSparse).

28.7.5 Sparse Resource Memory Requirements

Sparse resources have specific memory requirements related to binding sparse memory. These memory requirements are reported differently for VkBuffer objects and VkImage objects.

28.7.5.1 Buffer and Fully-Resident Images

Buffers (both fully and partially resident) and fully-resident images can be bound to memory using only the data from VkMemoryRequirements. For all sparse resources the VkMemoryRequirements::alignment member denotes both the bindable sparse block size in bytes and required alignment of VkDeviceMemory.

28.7.5.2 Partially Resident Images

Partially resident images have a different method for binding memory. As with buffers and fully resident images, the VkMemoryRequirements::alignment field denotes the bindable sparse block size in bytes for the image.

Requesting sparse memory requirements for VkImage objects using **vkGetImageSparseMemoryRequirements** will return an array of one or more VkSparseImageMemoryRequirements structures. Each structure describes the sparse memory requirements for a group of aspects of the image.

The sparse image must have been created using the VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT flag to retrieve valid sparse image memory requirements.

28.7.5.3 Sparse Image Memory Requirements

The VkSparseImageMemoryRequirements structure is defined as:

```
typedef struct VkSparseImageMemoryRequirements {
    VkSparseImageFormatProperties formatProperties;
    uint32_t imageMipTailFirstLod;
    VkDeviceSize imageMipTailSize;
    VkDeviceSize imageMipTailOffset;
    VkDeviceSize imageMipTailStride;
} VkSparseImageMemoryRequirements;
```

- formatProperties.aspectMask is the set of aspects of the image that this sparse memory requirement applies to. This will usually have a single aspect specified. However, depth/stencil images may have depth and stencil data interleaved in the same sparse block, in which case both VK_IMAGE_ASPECT_DEPTH_BIT and VK_IMAGE_ASPECT_STENCIL_BIT would be present.
- formatProperties.imageGranularity describes the dimensions of a single bindable sparse image block in pixel units. For aspect VK_IMAGE_ASPECT_METADATA_BIT, all dimensions will be zero pixels. All metadata is located in the mip tail region.
- formatProperties.flags is a bitmask of VkSparseImageFormatFlagBits:
 - If VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT is set the image uses a single mip tail region for all array layers.
 - If VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT is set the dimensions of mip levels must be integer multiples of the corresponding dimensions of the sparse image block for levels not located in the mip tail.
 - If VK_SPARSE_IMAGE_FORMAT_NONSTANDARD_BLOCK_SIZE_BIT is set the image uses non-standard sparse image block dimensions. The <code>formatProperties.imageGranularity</code> values do not match the standard sparse image block dimension corresponding to the image's pixel format.
- imageMipTailFirstLod is the first mip level at which image subresources are included in the mip tail region.
- imageMipTailSize is the memory size (in bytes) of the mip tail region. If formatProperties.flags contains VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT, this is the size of the whole mip tail, otherwise this is the size of the mip tail of a single array layer. This value is guaranteed to be a multiple of the sparse block size in bytes.
- imageMipTailOffset is the opaque memory offset used with VkSparseImageOpaqueMemoryBindInfo to bind the mip tail region(s).
- imageMipTailStride is the offset stride between each array-layer's mip tail, if formatProperties.flags does not contain VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT (otherwise the value is undefined).

To query sparse memory requirements for an image, call:

- device is the logical device that owns the image.
- *image* is the Vk Image object to get the memory requirements for.
- pSparseMemoryRequirementCount is a pointer to an integer related to the number of sparse memory requirements available or queried, as described below.
- pSparseMemoryRequirements is either NULL or a pointer to an array of VkSparseImageMemoryRequirements structures.

If pSparseMemoryRequirements is NULL, then the number of sparse memory requirements available is returned in pSparseMemoryRequirementCount. Otherwise, pSparseMemoryRequirementCount must point to a variable set by the user to the number of elements in the pSparseMemoryRequirements array, and on return the variable is overwritten with the number of structures actually written to pSparseMemoryRequirements. If

pSparseMemoryRequirementCount is less than the number of sparse memory requirements available, at most pSparseMemoryRequirementCount structures will be written.

If the image was not created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT then pSparseMemoryRequirementCount will be set to zero and pSparseMemoryRequirements will not be written to.



Note

It is legal for an implementation to report a larger value in VkMemoryRequirements::size than would be obtained by adding together memory sizes for all VkSparseImageMemoryRequirements returned by vkGetImageSparseMemoryRequirements. This may occur when the hardware requires unused padding in the address range describing the resource.

Valid Usage

- device must be a valid VkDevice handle
- image must be a valid VkImage handle
- pSparseMemoryRequirementCount must be a pointer to a uint32_t value
- If the value referenced by pSparseMemoryRequirementCount is not 0, and pSparseMemoryRequirements is not NULL, pSparseMemoryRequirements must be a pointer to an array of pSparseMemoryRequirementCount VkSparseImageMemoryRequirements structures
- image must have been created, allocated, or retrieved from device

28.7.6 Binding Resource Memory

Non-sparse resources are backed by a single physical allocation prior to device use (via **vkBindImageMemory** or **vkBindBufferMemory**), and their backing must not be changed. On the other hand, sparse resources can be bound to memory non-contiguously and these bindings can be altered during the lifetime of the resource.

Note



It is important to note that freeing a VkDeviceMemory object with **vkFreeMemory** will not cause resources (or resource regions) bound to the memory object to become unbound. Access to resources that are bound to memory objects that have been freed will result in undefined behavior, potentially including application termination.

Implementations must ensure that no access to physical memory owned by the system or another process will occur in this scenario. In other words, accessing resources bound to freed memory may result in application termination, but must not result in system termination or in reading non-process-accessible memory.

Sparse memory bindings execute on a queue that includes the VK_QUEUE_SPARSE_BINDING_BIT bit. Applications must use synchronization primitives to guarantee that other queues do not access ranges of memory concurrently with a binding change. Accessing memory in a range while it is being rebound results in undefined behavior. It is valid to access other ranges of the same resource while a bind operation is executing.



Note

Implementations must provide a guarantee that simultaneously binding sparse blocks while another queue accesses those same sparse blocks via a sparse resource must not access memory owned by another process or otherwise corrupt the system.

While some implementations may include VK_QUEUE_SPARSE_BINDING_BIT support in queue families that also include graphics and compute support, other implementations may only expose a VK_QUEUE_SPARSE_BINDING_BIT-only queue family. In either case, applications must use synchronization primitives to explicitly request any ordering dependencies between sparse memory binding operations and other graphics/compute/transfer operations, as sparse binding operations are not automatically ordered against command buffer execution, even within a single queue.

When binding memory explicitly for the VK_IMAGE_ASPECT_METADATA_BIT the application must use the VK_SPARSE_MEMORY_BIND_METADATA_BIT in the VkSparseMemoryBind::flags field when binding memory. Binding memory for metadata is done the same way as binding memory for the mip tail, with the addition of the VK_SPARSE_MEMORY_BIND_METADATA_BIT flag.

Binding the mip tail for any aspect must only be performed using VkSparseImageOpaqueMemoryBindInfo. If formatProperties.flags contains VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT, then it can be bound with a single VkSparseMemoryBind structure, with resourceOffset = imageMipTailOffset and size = imageMipTailSize.

If formatProperties.flags does not contain VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT then the offset for the mip tail in each array layer is given as:

```
arrayMipTailOffset = imageMipTailOffset + arrayLayer * imageMipTailStride;
```

and the mip tail can be bound with layerCount VkSparseMemoryBind structures, each using size = imageMipTailSize and resourceOffset = arrayMipTailOffset as defined above.

Sparse memory binding is handled by the following APIs and related data structures.

28.7.6.1 Sparse Memory Binding Functions

The VkSparseMemoryBind structure is defined as:

- resourceOffset is the offset into the resource.
- size is the size of the memory region to be bound.
- memory is the VkDeviceMemory object that the range of the resource is bound to. If memory is VK_NULL_HANDLE, the range is unbound.
- memoryOffset is the offset into the VkDeviceMemory object to bind the resource range to. If memory is VK_NULL HANDLE, this value is ignored.
- flags is a bitmask specifying usage of the binding operation. Bits which can be set include:

```
typedef enum VkSparseMemoryBindFlagBits {
    VK_SPARSE_MEMORY_BIND_METADATA_BIT = 0x00000001,
} VkSparseMemoryBindFlagBits;
```

- VK_SPARSE_MEMORY_BIND_METADATA_BIT indicates that the memory being bound is only for the metadata aspect.

The binding range [resourceOffset, resourceOffset + size) has different constraints based on flags. If flags contains VK_SPARSE_MEMORY_BIND_METADATA_BIT, the binding range must be within the mip tail region of the metadata aspect. This metadata region is defined by:

```
metadataRegion = [imageMipTailOffset + imageMipTailStride \times n, \\ imageMipTailOffset + imageMipTailStride \times n + imageMipTailSize)
```

Where imageMipTailOffset, imageMipTailSize, and imageMipTailStride values are from the VkSparseImageMemoryRequirements that correspond to the metadata aspect of the image. The term n is a valid array layer index for the image.

imageMipTailStride is considered to be zero for aspects where
VkSparseImageMemoryRequirements::formatProperties.flags contains VK_SPARSE_IMAGE_
FORMAT SINGLE MIPTAIL BIT.

If flags does not contain VK_SPARSE_MEMORY_BIND_METADATA_BIT, the binding range must be within the range [0, VkMemoryRequirements :: size).

Valid Usage

- If memory is not VK NULL HANDLE, memory must be a valid VkDeviceMemory handle
- flags must be a valid combination of VkSparseMemoryBindFlagBits values
- If memory is not VK_NULL_HANDLE, memory and memoryOffset must match the memory requirements of the resource, as described in section Section 11.6
- If memory is not VK_NULL_HANDLE, memory must not have been created with a memory type that reports VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT bit set
- size must be greater than 0
- resourceOffset must be less than the size of the resource
- size must be less than or equal to the size of the resource minus resourceOffset
- memoryOffset must be less than the size of memory
- size must be less than or equal to the size of memory minus memoryOffset

Memory is bound to VkBuffer objects created with the VK_BUFFER_CREATE_SPARSE_BINDING_BIT flag using the following structure:

- buffer is the VkBuffer object to be bound.
- bindCount is the number of VkSparseMemoryBind structures in the pBinds array.
- pBinds is a pointer to array of VkSparseMemoryBind structures.

Valid Usage

- buffer must be a valid VkBuffer handle
- pBinds must be a pointer to an array of bindCount valid VkSparseMemoryBind structures
- bindCount must be greater than 0

Memory is bound to opaque regions of VkImage objects created with the VK_IMAGE_CREATE_SPARSE_BINDING_BIT flag using the following structure:

- image is the VkImage object to be bound.
- \bullet bindCount is the number of VkSparseMemoryBind structures in the pBinds array.
- pBinds is a pointer to array of VkSparseMemoryBind structures.

Valid Usage

- image must be a valid VkImage handle
- pBinds must be a pointer to an array of bindCount valid VkSparseMemoryBind structures
- bindCount must be greater than 0
- For any given element of pBinds, if the flags member of that element contains VK_SPARSE_MEMORY_BIND_ METADATA_BIT, the binding range defined must be within the mip tail region of the metadata aspect of image

Note

This operation is normally used to bind memory to fully-resident sparse images or for mip tail regions of partially resident images. However, it can also be used to bind memory for the entire binding range of partially resident images.



In case flags does not contain VK_SPARSE_MEMORY_BIND_METADATA_BIT, the resourceOffset is in the range [0, VkMemoryRequirements:: size). This range includes data from all aspects of the image, including metadata. For most implementations this will probably mean that the resourceOffset is a simple device address offset within the resource. It is possible for an application to bind a range of memory that includes both resource data and metadata. However, the application would not know what part of the image the memory is used for, or if any range is being used for metadata.

When <code>flags</code> contains <code>VK_SPARSE_MEMORY_BIND_METADATA_BIT</code>, the binding range specified must be within the mip tail region of the metadata aspect. In this case the <code>resourceOffset</code> is not required to be a simple device address offset within the resource. However, it <code>is</code> defined to be within <code>[imageMipTailOffset, imageMipTailOffset]</code> for the metadata aspect. See <code>VkSparseMemoryBind</code> for the full constraints on binding region with this flag present.

Memory can be bound to sparse image blocks of VkImage objects created with the VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT flag using the following structure:

- image is the VkImage object to be bound
- bindCount is the number of VkSparseImageMemoryBind structures in pBinds array
- pBinds is a pointer to array of VkSparseImageMemoryBind structures

Valid Usage

- image must be a valid VkImage handle
- pBinds must be a pointer to an array of bindCount valid VkSparseImageMemoryBind structures
- bindCount must be greater than 0

The VkSparseImageMemoryBind structure is defined as:

- subresource is the aspectMask and region of interest in the image.
- offset are the coordinates of the first texel within the image subresource to bind.
- extent is the size in texels of the region within the image subresource to bind. The extent must be a multiple of the sparse image block dimensions, except when binding sparse image blocks along the edge of an image subresource it can instead be such that any coordinate of offset + extent equals the corresponding dimensions of the image subresource.
- memory is the VkDeviceMemory object that the sparse image blocks of the image are bound to. If memory is VK_NULL_HANDLE, the sparse image blocks are unbound.
- memoryOffset is an offset into VkDeviceMemory object. If memory is VK NULL HANDLE, this value is ignored.
- flags are sparse memory binding flags.

- subresource must be a valid VkImageSubresource structure
- If memory is not VK_NULL_HANDLE, memory must be a valid VkDeviceMemory handle
- flags must be a valid combination of VkSparseMemoryBindFlagBits values
- If the sparse aliased residency feature is not enabled, and if any other resources are bound to ranges of memory, the range of memory being bound must not overlap with those bound ranges
- memory and memoryOffset must match the memory requirements of the calling command's image, as described in section Section 11.6
- subresource must be a valid image subresource for image (see Section 11.5)
- offset.x must be a multiple of the sparse image block width (VkSparseImageFormatProperties::imageGranularity.width) of the image
- extent.width must either be a multiple of the sparse image block width of the image, or else extent.width + offset.x must equal the width of the image subresource
- offset.y must be a multiple of the sparse image block height (VkSparseImageFormatProperties::imageGranularity.height) of the image
- extent.height must either be a multiple of the sparse image block height of the image, or else extent. height + offset.y must equal the height of the image subresource
- offset.z must be a multiple of the sparse image block depth
 (VkSparseImageFormatProperties::imageGranularity.depth) of the image
- extent.depth must either be a multiple of the sparse image block depth of the image, or else extent.depth + offset.z must equal the depth of the image subresource

To submit sparse binding operations to a queue, call:

- queue is the queue that the sparse binding operations will be submitted to.
- bindInfoCount is the number of elements in the pBindInfo array.
- pBindInfo is an array of VkBindSparseInfo structures, each specifying a sparse binding submission batch.
- fence is an optional handle to a fence to be signaled. If fence is not VK_NULL_HANDLE, it defines a fence signal operation.

vkQueueBindSparse is a queue submission command, with each batch defined by an element of *pBindInfo* as an instance of the VkBindSparseInfo structure.

Within a batch, a given range of a resource must not be bound more than once. Across batches, if a range is to be bound to one allocation and offset and then to another allocation and offset, then the application must guarantee (usually using semaphores) that the binding operations are executed in the correct order, as well as to order binding operations against the execution of command buffer submissions.

As no operation to vkQueueBindSparse causes any pipeline stage to access memory, synchronization primitives used in this command effectively only define execution dependencies.

Additional information about fence and semaphore operation is described in the synchronization chapter.

- queue must be a valid VkQueue handle
- If bindInfoCount is not 0, pBindInfo must be a pointer to an array of bindInfoCount valid VkBindSparseInfo structures
- If fence is not VK_NULL_HANDLE, fence must be a valid VkFence handle
- The queue must support sparse binding operations
- Both of fence, and queue that are valid handles must have been created, allocated, or retrieved from the same VkDevice
- fence must be unsignaled
- fence must not be associated with any other queue command that has not yet completed execution on that queue

Host Synchronization

- Host access to queue must be externally synchronized
- Host access to pBindInfo[].pWaitSemaphores[] must be externally synchronized
- Host access to pBindInfo[].pSignalSemaphores[] must be externally synchronized
- Host access to pBindInfo[].pBufferBinds[].buffer must be externally synchronized
- Host access to pBindInfo[].pImageOpaqueBinds[].image must be externally synchronized
- Host access to pBindInfo[].pImageBinds[].image must be externally synchronized
- Host access to fence must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
-	-	SPARSE_BINDING

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST

The $\mbox{VkBindSparseInfo}$ structure is defined as:

- sType is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- wait SemaphoreCount is the number of semaphores upon which to wait before executing the sparse binding operations for the batch.
- pWaitSemaphores is a pointer to an array of semaphores upon which to wait on before the sparse binding operations for this batch begin execution. If semaphores to wait on are provided, they define a semaphore wait operation.
- bufferBindCount is the number of sparse buffer bindings to perform in the batch.
- pBufferBinds is a pointer to an array of VkSparseBufferMemoryBindInfo structures.
- imageOpaqueBindCount is the number of opaque sparse image bindings to perform.
- pImageOpaqueBinds is a pointer to an array of VkSparseImageOpaqueMemoryBindInfo structures, indicating opaque sparse image bindings to perform.
- imageBindCount is the number of sparse image bindings to perform.
- pImageBinds is a pointer to an array of VkSparseImageMemoryBindInfo structures, indicating sparse image bindings to perform.
- signalSemaphoreCount is the number of semaphores to be signaled once the sparse binding operations specified by the structure have completed execution.
- pSignalSemaphores is a pointer to an array of semaphores which will be signaled when the sparse binding operations for this batch have completed execution. If semaphores to be signaled are provided, they define a semaphore signal operation.

- stype must be VK_STRUCTURE_TYPE_BIND_SPARSE_INFO
- pNext must be NULL
- If waitSemaphoreCount is not 0, pWaitSemaphores must be a pointer to an array of waitSemaphoreCount valid VkSemaphore handles
- If bufferBindCount is not 0, pBufferBinds must be a pointer to an array of bufferBindCount valid VkSparseBufferMemoryBindInfo structures

- If imageOpaqueBindCount is not 0, pImageOpaqueBinds must be a pointer to an array of imageOpaqueBindCount valid VkSparseImageOpaqueMemoryBindInfo structures
- If imageBindCount is not 0, pImageBinds must be a pointer to an array of imageBindCount valid VkSparseImageMemoryBindInfo structures
- If signalSemaphoreCount is not 0, pSignalSemaphores must be a pointer to an array of signalSemaphoreCount valid VkSemaphore handles
- Both of the elements of pSignalSemaphores, and the elements of pWaitSemaphores that are valid handles must have been created, allocated, or retrieved from the same VkDevice

28.8 Examples

The following examples illustrate basic creation of sparse images and binding them to physical memory.

28.8.1 Basic Sparse Resources

This basic example creates a normal VkImage object but uses fine-grained memory allocation to back the resource with multiple memory ranges.

```
VkDevice
                      device;
Vk0ueue
                     queue;
VkImage
                     sparseImage;
VkMemoryRequirements memoryRequirements = {};
// ...
// Allocate image object
const VkImageCreateInfo sparseImageInfo =
   VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO,
                                          // sType
   NULT.
                                          // pNext
   VK_IMAGE_CREATE_SPARSE_BINDING_BIT | ..., // flags
};
vkCreateImage(device, &sparseImageInfo, &sparseImage);
// Get memory requirements
vkGetImageMemoryRequirements(
   device,
   sparseImage,
   &memoryRequirements);
// Bind memory in fine-grained fashion, find available memory ranges
// from potentially multiple VkDeviceMemory pools.
// (Illustration purposes only, can be optimized for perf)
while (memoryRequirements.size && bindCount < MAX_CHUNKS)</pre>
```

```
VkSparseMemoryBind* pBind = &binds[bindCount];
   pBind->resourceOffset = offset;
   AllocateOrGetMemoryRange(
       device,
       &memoryRequirements,
       &pBind->memory,
        &pBind->memoryOffset,
        &pBind->size);
   // memory ranges must be sized as multiples of the alignment
    assert(IsMultiple(pBind->size, memoryRequirements.alignment));
    assert(IsMultiple(pBind->memoryOffset, memoryRequirements.alignment));
   memoryRequirements.size -= pBind->size;
   offset
                           += pBind->size;
   bindCount++;
// Ensure all image has backing
if (memoryRequirements.size)
{
   // Error condition - too many chunks
const VkSparseImageOpaqueMemoryBindInfo opaqueBindInfo =
   sparseImage,
                                                // image
   bindCount,
                                                // bindCount
   binds
                                                // pBinds
};
const VkBindSparseInfo bindSparseInfo =
   VK_STRUCTURE_TYPE_BIND_SPARSE_INFO,
                                               // sType
   NULL,
                                                // pNext
   . . .
                                                // imageOpaqueBindCount
   1,
   &opaqueBindInfo,
                                                // pImageOpaqueBinds
};
// vkQueueBindSparse is application synchronized per queue object.
AcquireQueueOwnership(queue);
// Actually bind memory
vkQueueBindSparse(queue, 1, &bindSparseInfo, VK_NULL_HANDLE);
ReleaseQueueOwnership(queue);
```

28.8.2 Advanced Sparse Resources

This more advanced example creates an arrayed color attachment / texture image and binds only LOD zero and the required metadata to physical memory.

```
VkDevice
                                   device;
VkQueue
                                   queue;
VkImage
                                   sparseImage;
                                   memoryRequirements = {};
VkMemoryRequirements
uint32_t
                                   sparseRequirementsCount = 0;
VkSparseMemoryBind
                                   binds[MY_IMAGE_ARRAY_SIZE] = {};
VkSparseImageMemoryBind
                                   imageBinds[MY_IMAGE_ARRAY_SIZE] = {};
uint32_t
                                   bindCount = 0;
// Allocate image object (both renderable and sampleable)
const VkImageCreateInfo sparseImageInfo =
   VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO,
                                               // sType
                                               // pNext
   VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT | ..., // flags
   VK_FORMAT_R8G8B8A8_UNORM,
                                              // format
   MY_IMAGE_ARRAY_SIZE,
                                              // arrayLayers
   VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT |
   VK_IMAGE_USAGE_SAMPLED_BIT,
                                               // usage
};
vkCreateImage(device, &sparseImageInfo, &sparseImage);
// Get memory requirements
vkGetImageMemoryRequirements(
   device,
   sparseImage,
   &memoryRequirements);
// Get sparse image aspect properties
vkGetImageSparseMemoryRequirements(
   device,
   sparseImage,
   &sparseRequirementsCount,
   NULL);
pSparseReqs = (VkSparseImageMemoryRequirements*)
   malloc(sparseRequirementsCount * sizeof(VkSparseImageMemoryRequirements));
vkGetImageSparseMemoryRequirements(
   device,
   sparseImage,
   &sparseRequirementsCount,
   pSparseReqs);
// Bind LOD level 0 and any required metadata to memory
for (uint32_t i = 0; i < sparseRequirementsCount; ++i)</pre>
    if (pSparseReqs[i].formatProperties.aspectMask &
       VK_IMAGE_ASPECT_METADATA_BIT)
    {
       // Metadata must not be combined with other aspects
```

```
assert(pSparseReqs[i].formatProperties.aspectMask ==
           VK_IMAGE_ASPECT_METADATA_BIT);
    if (pSparseReqs[i].formatProperties.flags &
        VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT)
    {
        VkSparseMemoryBind* pBind = &binds[bindCount];
        pBind->memorySize = pSparseReqs[i].imageMipTailSize;
        bindCount++;
        // ... Allocate memory range
        pBind->resourceOffset = pSparseReqs[i].imageMipTailOffset;
        pBind->memoryOffset = /* allocated memoryOffset */;
        pBind->memory = /* allocated memory */;
        pBind->flags = VK_SPARSE_MEMORY_BIND_METADATA_BIT;
    }
    else
        // Need a mip tail region per array layer.
        for (uint32_t a = 0; a < sparseImageInfo.arrayLayers; ++a)</pre>
        {
            VkSparseMemoryBind* pBind = &binds[bindCount];
            pBind->memorySize = pSparseReqs[i].imageMipTailSize;
            bindCount++;
            // ... Allocate memory range
            pBind->resourceOffset = pSparseReqs[i].imageMipTailOffset +
                                    (a * pSparseReqs[i].imageMipTailStride);
            pBind->memoryOffset = /* allocated memoryOffset */;
            pBind->memory = /* allocated memory */
            pBind->flags = VK_SPARSE_MEMORY_BIND_METADATA_BIT;
}
else
    // resource data
   VkExtent3D lod0BlockSize =
        AlignedDivide(
            sparseImageInfo.extent.width,
            pSparseReqs[i].formatProperties.imageGranularity.width);
        AlignedDivide(
            sparseImageInfo.extent.height,
            pSparseRegs[i].formatProperties.imageGranularity.height);
        AlignedDivide(
            sparseImageInfo.extent.depth,
            pSparseReqs[i].formatProperties.imageGranularity.depth);
    size_t totalBlocks =
        lod0BlockSize.width *
        lod0BlockSize.height *
        lod0BlockSize.depth;
```

```
VkDeviceSize lod0MemSize = totalBlocks * memoryRequirements.alignment;
        // Allocate memory for each array layer
        for (uint32_t a = 0; a < sparseImageInfo.arrayLayers; ++a)</pre>
            // ... Allocate memory range
            VkSparseImageMemoryBind* pBind = &imageBinds[a];
            pBind->subresource.aspectMask = pSparseReqs[i].formatProperties.aspectMask ←
            pBind->subresource.mipLevel = 0;
            pBind->subresource.arrayLayer = a;
            pBind->offset = (VkOffset3D) {0, 0, 0};
            pBind->extent = sparseImageInfo.extent;
            pBind->memoryOffset = /* allocated memoryOffset */;
            pBind->memory = /* allocated memory */;
            pBind->flags = 0;
    free (pSparseReqs);
const VkSparseImageOpaqueMemoryBindInfo opaqueBindInfo =
    sparseImage,
                                                 // image
    bindCount,
                                                 // bindCount
   binds
                                                 // pBinds
};
const VkSparseImageMemoryBindInfo imageBindInfo =
    sparseImage,
                                                // image
    sparseImageInfo.arrayLayers,
                                                // bindCount
    imageBinds
                                                 // pBinds
};
const VkBindSparseInfo bindSparseInfo =
   VK_STRUCTURE_TYPE_BIND_SPARSE_INFO,
                                               // sType
   NULL,
                                                // pNext
    . . .
                                                // imageOpaqueBindCount
    1,
    &opaqueBindInfo,
                                                // pImageOpaqueBinds
                                                 // imageBindCount
    1,
    &imageBindInfo,
                                                 // pImageBinds
    . . .
};
// vkQueueBindSparse is application synchronized per queue object.
AcquireQueueOwnership(queue);
// Actually bind memory
vkQueueBindSparse(queue, 1, &bindSparseInfo, VK_NULL_HANDLE);
```

eleaseQueueOwnership(queue);	

Chapter 29

Window System Integration (WSI)

This chapter discusses the window system integration (WSI) between the Vulkan API and the various forms of displaying the results of rendering to a user. Since the Vulkan API can be used without displaying results, WSI is provided through the use of optional Vulkan extensions. This chapter provides an overview of WSI. See the appendix for additional details of each WSI extension, including which extensions must be enabled in order to use each of the functions described in this chapter.

29.1 WSI Platform

A platform is an abstraction for a window system, OS, etc. Some examples include MS Windows, Android, and Wayland. The Vulkan API may be integrated in a unique manner for each platform.

The Vulkan API does not define any type of platform object. Platform-specific WSI extensions are defined, which contain platform-specific functions for using WSI. Use of these extensions is guarded by preprocessor symbols as defined in the Window System-Specific Header Control appendix.

In order for an application to be compiled to use WSI with a given platform, it must #define the appropriate preprocessor symbol prior to including the "vulkan.h" header file. Each platform-specific extension is an instance extension. The application must enable instance extensions with **vkCreateInstance** before using them.

29.2 WSI Surface

A VkSurfaceKHR object abstracts a native platform surface or window object for use with Vulkan. The VK_KHR_ surface extension declares the VkSurfaceKHR object, and provides a function for destroying VkSurfaceKHR objects. Separate platform-specific extensions each provide a function for creating a VkSurfaceKHR object for the respective platform. From the application's perspective this is an opaque handle, just like the handles of other Vulkan objects.



Note

On certain platforms, the Vulkan loader and ICDs may have conventions that treat the handle as a pointer to a struct that contains the platform-specific information about the surface. This will be described in the documentation for the loader-ICD interface, and in the "vk_icd.h" header file of the LoaderAndTools source-code repository. This does not affect the loader-layer interface; layers may wrap VkSurfaceKHR objects.

29.2.1 Android Platform

To create a VkSurfaceKHR object for an Android native window, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkAndroidSurfaceCreateInfoKHR structure containing parameters affecting the creation of the surface object.
- pallocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

During the lifetime of a surface created using a particular ANativeWindow handle any attempts to create another surface for the same ANativeWindow and any attempts to connect to the same ANativeWindow through other platform mechanisms will fail.



NOTE

In particular, only one VkSurfaceKHR can exist at a time for a given window. Similarly, a native window cannot be used by both a VkSurfaceKHR and EGLSurface simultaneously.

If successful, **vkCreateAndroidSurfaceKHR** increments the ANativeWindow's reference count, and **vkDestroySurfaceKHR** will decrement it.

On Android, when a swapchain's <code>imageExtent</code> does not match the surface's <code>currentExtent</code>, the swapchain images will be scaled to the surface's dimensions during presentation. <code>minImageExtent</code> is (1,1), and <code>maxImageExtent</code> is the maximum image size supported by the consumer. For the system compositor, <code>currentExtent</code> is the window size (i.e. the consumer's preferred size).

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkAndroidSurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_NATIVE_WINDOW_IN_USE_KHR

The VkAndroidSurfaceCreateInfoKHR structure is defined as:

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_ANDROID_SURFACE_CREATE_INFO_ KHR
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- window is a pointer to the ANativeWindow to associate the surface with.

Valid Usage

- stype must be VK_STRUCTURE_TYPE_ANDROID_SURFACE_CREATE_INFO_KHR
- pNext must be NULL
- flags must be 0
- window must be a pointer to a ANativeWindow value
- window must not be in a connected state

29.2.2 Mir Platform

To create a VkSurfaceKHR object for a Mir window, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkMirSurfaceCreateInfoKHR structure containing parameters affecting the creation of the surface object.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkMirSurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkMirSurfaceCreateInfoKHR structure is defined as:

```
MirConnection* connection;
MirSurface* mirSurface;
} VkMirSurfaceCreateInfoKHR;
```

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_MIR_SURFACE_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- connection and surface are pointers to the MirConnection and MirSurface for the window to associate the surface with.

- • sType must be $VK_STRUCTURE_TYPE_MIR_SURFACE_CREATE_INFO_KHR$
- pNext must be NULL
- flags must be 0
- connection must be a pointer to a MirConnection value
- mirSurface must be a pointer to a MirSurface value

On Mir, when a swapchain's <code>imageExtent</code> does not match the surface's <code>currentExtent</code>, the swapchain images will be scaled to the surface's dimensions during presentation. <code>minImageExtent</code> is (1,1), and <code>maxImageExtent</code> is the maximum supported surface size.

29.2.3 Wayland Platform

To create a VkSurfaceKHR object for a Wayland surface, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkWaylandSurfaceCreateInfoKHR structure containing parameters affecting the creation of the surface object.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkWaylandSurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkWaylandSurfaceCreateInfoKHR structure is defined as:

- *sType* is the type of this structure and must be VK_STRUCTURE_TYPE_WAYLAND_SURFACE_CREATE_INFO_ KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- display and surface are pointers to the Wayland wl_display and wl_surface to associate the surface with.

- stype must be VK STRUCTURE TYPE WAYLAND SURFACE CREATE INFO KHR
- pNext must be NULL
- flags must be 0
- display must be a pointer to a wl_display value
- surface must be a pointer to a wl_surface value

On Wayland, currentExtent is undefined (0,0). Whatever the application sets a swapchain's imageExtent to will be the size of the window, after the first image is presented. minImageExtent is (1,1), and maxImageExtent is the maximum supported surface size.

Some Vulkan functions may send protocol over the specified **wl_display** connection when using a swapchain or presentable images created from a VkSurfaceKHR referring to a **wl_surface**. Applications must therefore ensure that both the **wl_display** and the **wl_surface** remain valid for the lifetime of any VkSwapchainKHR objects created from a particular **wl_display** and **wl_surface**. Also, calling vkQueuePresentKHR will result in Vulkan sending **wl_surface**.commit requests to the underlying **wl_surface** of each VkSwapchainKHR objects referenced by pPresentInfo. Therefore, if the application wishes to synchronize any window changes with a particular frame, such requests must be sent to the Wayland display server prior to calling vkQueuePresentKHR.

29.2.4 Win32 Platform

To create a VkSurfaceKHR object for a Win32 window, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkWin32SurfaceCreateInfoKHR structure containing parameters affecting the creation of the surface object.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

Valid Usage

• instance must be a valid VkInstance handle

- pCreateInfo must be a pointer to a valid VkWin32SurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkWin32SurfaceCreateInfoKHR structure is defined as:

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_WIN32_SURFACE_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- hinstance and hwnd are the Win32 HINSTANCE and HWND for the window to associate the surface with.

Valid Usage

- stype must be $VK_STRUCTURE_TYPE_WIN32_SURFACE_CREATE_INFO_KHR$
- pNext must be NULL
- flags must be 0

With Win32, minImageExtent, maxImageExtent, and currentExtent are the window size. Therefore, a swapchain's imageExtent must match the window's size.

29.2.5 XCB Platform

To create a VkSurfaceKHR object for an X11 window, using the XCB client-side library, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkXcbSurfaceCreateInfoKHR structure containing parameters affecting the creation of the surface object.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

Valid Usage

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkXcbSurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkXcbSurfaceCreateInfoKHR structure is defined as:

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_XCB_SURFACE_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use.
- connection is a pointer to an xcb_connection_t to the X server.
- window is the xcb_window_t for the X11 window to associate the surface with.

- sType must be VK_STRUCTURE_TYPE_XCB_SURFACE_CREATE_INFO_KHR
- pNext must be NULL
- flags must be 0
- connection must be a pointer to a xcb_connection_t value

With Xcb, minImageExtent, maxImageExtent, and currentExtent are the window size. Therefore, a swapchain's imageExtent must match the window's size.

Some Vulkan functions may send protocol over the specified xcb connection when using a swapchain or presentable images created from a VkSurface referring to an xcb window. Applications must therefore ensure the xcb connection is available to Vulkan for the duration of any functions that manipulate such swapchains or their presentable images, and any functions that build or queue command buffers that operate on such presentable images. Specifically, applications using Vulkan with xcb-based swapchains must

 Avoid holding a server grab on an xcb connection while waiting for Vulkan operations to complete using a swapchain derived from a different xcb connection referring to the same X server instance. Failing to do so may result in deadlock.

29.2.6 Xlib Platform

To create a VkSurfaceKHR object for an X11 window, using the Xlib client-side library, call:

- instance is the instance to associate the surface with.
- pCreateInfo is a pointer to an instance of the VkXlibSurfaceCreateInfoKHR structure containing the parameters affecting the creation of the surface object.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface object is returned.

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkXlibSurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkXlibSurfaceCreateInfoKHR structure is defined as:

- *sType* is the type of this structure and must be VK_STRUCTURE_TYPE_XLIB_SURFACE_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.

- flags is reserved for future use.
- dpy is a pointer to an Xlib Display connection to the X server.
- window is an Xlib Window to associate the surface with.

- sType must be VK_STRUCTURE_TYPE_XLIB_SURFACE_CREATE_INFO_KHR
- pNext must be NULL
- flags must be 0
- dpy must be a pointer to a Display value

With Xlib, minImageExtent, maxImageExtent, and currentExtent are the window size. Therefore, a swapchain's imageExtent must match the window's size.

Some Vulkan functions may send protocol over the specified Xlib Display connection when using a swapchain or presentable images created from a VkSurface referring to an Xlib window. Applications must therefore ensure the display connection is available to Vulkan for the duration of any functions that manipulate such swapchains or their presentable images, and any functions that build or queue command buffers that operate on such presentable images. Specifically, applications using Vulkan with Xlib-based swapchains must

- Call XInitThreads() before calling any other Xlib functions if they intend to use Vulkan in multiple threads, or use Vulkan and Xlib in separate threads.
- Avoid holding a server grab on a display connection while waiting for Vulkan operations to complete using a swapchain derived from a different display connection referring to the same X server instance. Failing to do so may result in deadlock.

29.2.7 Platform-Independent Information

Once created, VkSurfaceKHR objects can be used in this and other extensions, in particular the $VK_KHR_swapchain$ extension.

Several WSI functions return VK_ERROR_SURFACE_LOST_KHR if the surface becomes no longer available. After such an error, the surface (and any child swapchain, if one exists) should be destroyed, as there is no way to restore them to a not-lost state. Applications may attempt to create a new VkSurfaceKHR using the same native platform window object, but whether such re-creation will succeed is platform-dependent and may depend on the reason the surface became unavailable. A lost surface does not otherwise cause devices to be lost.

To destroy a VkSurfaceKHR object, call:

- instance is the instance used to create the surface.
- surface is the surface to destroy.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).

Destroying a VkSurfaceKHR merely severs the connection between Vulkan and the native surface, and doesn't imply destroying the native surface, closing a window, or similar behavior.

Valid Usage

- instance must be a valid VkInstance handle
- If surface is not VK_NULL_HANDLE, surface must be a valid VkSurfaceKHR handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- If surface is a valid handle, it must have been created, allocated, or retrieved from instance
- All VkSwapchainKHR objects created for surface must have been destroyed prior to destroying surface
- If VkAllocationCallbacks were provided when *surface* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when surface was created, pAllocator must be NULL

Host Synchronization

• Host access to surface must be externally synchronized

29.3 Presenting Directly to Display Devices

In some environments applications can also present Vulkan rendering directly to display devices without using an intermediate windowing system. This can be useful for embedded applications, or implementing the rendering/presentation backend of a windowing system using Vulkan. The VK_EXT_KHR_display extension provides the functionality necessary to enumerate display devices and create VkSurface objects that target displays.

29.3.1 Display Enumeration

Various functions are provided for enumerating the available display devices present on a Vulkan physical device. To query information about the available displays, call:

- physicalDevice is a physical device.
- pPropertyCount is a pointer to an integer related to the number of display devices available or queried, as described below.
- pProperties is either NULL or a pointer to an array of VkDisplayPropertiesKHR structures.

If pProperties is NULL, then the number of display devices available for physicalDevice is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If the value of pPropertyCount is less than the number of display devices for physicalDevice, at most pPropertyCount structures will be written. If pPropertyCount is smaller than the number of display devices available for physicalDevice, VK_INCOMPLETE will be returned instead of VK_SUCCESS to indicate that not all the available values were returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkDisplayPropertiesKHR structures

Return Codes

Success

- VK SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDisplayPropertiesKHR structure is defined as:

```
typedef struct VkDisplayPropertiesKHR {
   VkDisplayKHR
                                  display;
   const char*
                                  displayName;
   VkExtent2D
                                  physicalDimensions;
   VkExtent2D
                                  physicalResolution;
   VkSurfaceTransformFlagsKHR
                               supportedTransforms;
   VkBool32
                                  planeReorderPossible;
   VkBool32
                                  persistentContent;
} VkDisplayPropertiesKHR;
```

- display is a handle that is used to refer to the display described here. This handle will be valid for the lifetime of the Vulkan instance.
- displayName is a pointer to a NULL-terminated string containing the name of the display. Generally, this will be the name provided by the display's EDID. It can be NULL if no suitable name is available.
- physicalDimensions describes the physical width and height of the visible portion of the display, in millimeters.
- physicalResolution describes the physical, native, or preferred resolution of the display.



Note

For devices which have no natural value to return here, implementations should return the maximum resolution supported.

- *supportedTransforms* tells which transforms are supported by this display. This will contain one or more of the bits from VkSurfaceTransformFlagsKHR.
- planeReorderPossible tells whether the planes on this display can have their z order changed. If this is VK_TRUE, the application can re-arrange the planes on this display in any order relative to each other.
- persistentContent tells whether the display supports self-refresh/internal buffering. If this is true, the application can submit persistent present operations on swapchains created against this display.



Note

Persistent presents may have higher latency, and may use less power when the screen content is updated infrequently, or when only a portion of the screen needs to be updated in most frames.

- display must be a valid VkDisplayKHR handle
- displayName must be a null-terminated string
- supportedTransforms must be a valid combination of VkSurfaceTransformFlagBitsKHR values

29.3.1.1 Display Planes

Images are presented to individual planes on a display. Devices must support at least one plane on each display. Planes can be stacked and blended to composite multiple images on one display. Devices may support only a fixed stacking order and fixed mapping between planes and displays, or they may allow arbitrary application specified stacking orders and mappings between planes and displays. To query the properties of device display planes, call:

- physicalDevice is a physical device.
- pPropertyCount is a pointer to an integer related to the number of display planes available or queried, as described below
- pProperties is either NULL or a pointer to an array of VkDisplayPlanePropertiesKHR structures.

If pProperties is NULL, then the number of display planes available for physicalDevice is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If the value of pPropertyCount is less than the number of display planes for physicalDevice, at most pPropertyCount structures will be written.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkDisplayPlanePropertiesKHR structures

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDisplayPlanePropertiesKHR structure is defined as:

- *currentDisplay* is the handle of the display the plane is currently associated with. If the plane is not currently attached to any displays, this will be VK NULL HANDLE.
- currentStackIndex is the current z-order of the plane. This will be between 0 and the value returned by vkGetPhysicalDeviceDisplayPlanePropertiesKHR() in pPropertyCount.

Valid Usage

• currentDisplay must be a valid VkDisplayKHR handle

To determine which displays a plane is usable with, call

- physicalDevice is a physical device.
- planeIndex is the plane which the application wishes to use, and must be in the range [0,physicaldeviceplanecount 1].
- pDisplayCount is a pointer to an integer related to the number of display planes available or queried, as described below.
- pDisplays is either NULL or a pointer to an array of VkDisplayKHR structures.

If pDisplays is NULL, then the number of displays usable with the specified planeIndex for physicalDevice is returned in pDisplayCount. Otherwise, pDisplayCount must point to a variable set by the user to the number of elements in the pDisplays array, and on return the variable is overwritten with the number of structures actually written to pDisplays. If the value of pDisplayCount is less than the number of display planes for physicalDevice, at most pDisplayCount structures will be written. If pDisplayCount is smaller than the number of displays usable with the specified planeIndex for physicalDevice, VK_INCOMPLETE will be returned instead of VK_SUCCESS to indicate that not all the available values were returned.

- physicalDevice must be a valid VkPhysicalDevice handle
- pDisplayCount must be a pointer to a uint32_t value
- If the value referenced by pDisplayCount is not 0, and pDisplays is not NULL, pDisplays must be a pointer to an array of pDisplayCount VkDisplayKHR handles
- planeIndex must be less than the number of display planes supported by the device as determined by calling vkGetPhysicalDeviceDisplayPlanePropertiesKHR

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

Additional properties of displays are queried using specialized query functions.

29.3.1.2 Display Modes

Each display has one or more supported modes associated with it by default. These built-in modes are queried by calling:

- physicalDevice is the physical device associated with display.
- display is the display to query.
- pPropertyCount is a pointer to an integer related to the number of display modes available or queried, as described below.
- pProperties is either NULL or a pointer to an array of VkDisplayModePropertiesKHR structures.

If pProperties is NULL, then the number of display modes available on the specified <code>display</code> for <code>physicalDevice</code> is returned in <code>pPropertyCount</code>. Otherwise, <code>pPropertyCount</code> must point to a variable set by the user to the number of elements in the <code>pProperties</code> array, and on return the variable is overwritten with the number of structures actually written to <code>pProperties</code>. If the value of <code>pPropertyCount</code> is less than the number of display modes for <code>physicalDevice</code>, at most <code>pPropertyCount</code> structures will be written. If <code>pPropertyCount</code> is smaller than the number of display modes available on the specified <code>display</code> for <code>physicalDevice</code>, <code>VK_INCOMPLETE</code> will be returned instead of <code>VK_SUCCESS</code> to indicate that not all the available values were returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- display must be a valid VkDisplayKHR handle
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkDisplayModePropertiesKHR structures

Return Codes

Success

- VK SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDisplayModePropertiesKHR structure is defined as:

- displayMode is a handle to the display mode described in this structure. This handle will be valid for the lifetime of the Vulkan instance.
- parameters is a VkDisplayModeParametersKHR structure describing the display parameters associated with displayMode.

• displayMode must be a valid VkDisplayModeKHR handle

The $\mbox{VkDisplayModeParameters}\mbox{KHR}$ structure is defined as:

```
typedef struct VkDisplayModeParametersKHR {
    VkExtent2D    visibleRegion;
    uint32_t    refreshRate;
} VkDisplayModeParametersKHR;
```

- visibleRegion is the 2D extents of the visible region.
- refreshRate is a uint32_t that is the number of times the display is refreshed each second multiplied by 1000.



Note

For example, a 60Hz display mode would report a refreshRate of 60,000.

Additional modes may also be created by calling:

- physicalDevice is the physical device associated with display.
- display is the display to create an additional mode for.
- pCreateInfo is a VkDisplayModeCreateInfoKHR structure describing the new mode to create.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pMode returns the handle of the mode created.

- physicalDevice must be a valid VkPhysicalDevice handle
- display must be a valid VkDisplayKHR handle
- pCreateInfo must be a pointer to a valid VkDisplayModeCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pMode must be a pointer to a VkDisplayModeKHR handle

Host Synchronization

• Host access to display must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_INITIALIZATION_FAILED

The VkDisplayModeCreateInfoKHR structure is defined as:

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_DISPLAY_MODE_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use, and must be zero.

• parameters is a VkDisplayModeParametersKHR structure describing the display parameters to use in creating the new mode. If the parameters are not compatible with the specified display, the implementation must return VK_ERROR_INITIALIZATION_FAILED.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DISPLAY_MODE_CREATE_INFO_KHR
- pNext must be NULL
- flags must be 0
- The width and height members of the visibleRegion member of parameters must be greater than 0
- The refreshRate member of parameters must be greater than 0

Applications that wish to present directly to a display must select which layer, or "plane" of the display they wish to target, and a mode to use with the display. Each display supports at least one plane. The capabilities of a given mode and plane combination are determined by calling:

- physicalDevice is the physical device associated with display
- mode is the display mode the application intends to program when using the specified plane. Note this parameter also implicitly specifies a display.
- planeIndex is the plane which the application intends to use with the display, and is less than the number of display planes supported by the device.
- pCapabilities is a pointer to a VkDisplayPlaneCapabilitiesKHR structure in which the capabilities are returned.

- physicalDevice must be a valid VkPhysicalDevice handle
- mode must be a valid VkDisplayModeKHR handle
- pCapabilities must be a pointer to a VkDisplayPlaneCapabilitiesKHR structure

Host Synchronization

• Host access to mode must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The $\mbox{VkDisplayPlaneCapabilities}\mbox{KHR}$ structure is defined as:

```
typedef struct VkDisplayPlaneCapabilitiesKHR {
   VkDisplayPlaneAlphaFlagsKHR supportedAlpha;
   VkOffset2D
                                  minSrcPosition;
   VkOffset2D
                                  maxSrcPosition;
   VkExtent2D
                                  minSrcExtent;
   VkExtent2D
                                  maxSrcExtent;
   VkOffset2D
                                  minDstPosition;
   VkOffset2D
                                  maxDstPosition;
   VkExt.ent.2D
                                  minDstExtent;
   VkExtent2D
                                   maxDstExtent;
} VkDisplayPlaneCapabilitiesKHR;
```

- supportedAlpha is a bitmask of VkDisplayPlaneAlphaFlagBitsKHR describing the supported alpha blending modes.
- minSrcPosition is the minimum source rectangle offset supported by this plane using the specified mode.
- maxSrcPosition is the maximum source rectangle offset supported by this plane using the specified mode. The x and y components of maxSrcPosition must each be greater than or equal to the x and y components of minSrcPosition, respectively.
- minSrcExtent is the minimum source rectangle size supported by this plane using the specified mode.
- maxSrcExtent is the maximum source rectangle size supported by this plane using the specified mode.
- minDstPosition, maxDstPosition, minDstExtent, maxDstExtent all have similar semantics to their corresponding "Src" equivalents, but apply to the output region within the mode rather than the input region within the source image. Unlike the "Src" offsets, minDstPosition and maxDstPosition may contain negative values.

The minimum and maximum position and extent fields describe the hardware limits, if any, as they apply to the specified display mode and plane. Vendors may support displaying a subset of a swapchain's presentable images on the specified display plane. This is expressed by returning <code>minSrcPosition</code>, <code>maxSrcPosition</code>, <code>minSrcExtent</code>, and <code>maxSrcExtent</code> values that indicate a range of possible positions and sizes may be used to specify the region within the presentable images that source pixels will be read from when creating a swapchain on the specified display mode and plane.

Vendors may also support mapping the presentable images' content to a subset or superset of the visible region in the specified display mode. This is expressed by returning <code>minDstPosition</code>, <code>maxDstPosition</code>, <code>minDstExtent</code> and <code>maxDstExtent</code> values that indicate a range of possible positions and sizes may be used to describe the region within the display mode that the source pixels will be mapped to.

Other vendors may support only a 1-1 mapping between pixels in the presentable images and the display mode. This may be indicated by returning (0,0) for minSrcPosition, maxSrcPosition, minDstPosition, and maxDstPosition, and (display mode width, display mode height) for minSrcExtent, maxSrcExtent, minDstExtent, and maxDstExtent.

These values indicate the limits of the hardware's individual fields. Not all combinations of values within the offset and extent ranges returned in VkDisplayPlaneCapabilitiesKHR are guaranteed to be supported. Vendors may still fail presentation requests that specify unsupported combinations.

Valid Usage

• supportedAlpha must be a valid combination of VkDisplayPlaneAlphaFlaqBitsKHR values

29.3.2 Display Surfaces

A complete display configuration includes a mode, one or more display planes and any parameters describing their behavior, and parameters describing some aspects of the images associated with those planes. Display surfaces describe the configuration of a single plane within a complete display configuration. To create a VkSurfaceKHR structure for a display surface, call:

- instance is the instance corresponding to the physical device the targeted display is on.
- pCreateInfo:pCreateInfo is a pointer to an instance of the VkDisplaySurfaceCreateInfoKHR structure specifying which mode, plane, and other parameters to use, as described below.
- pAllocator is the allocator used for host memory allocated for the surface object when there is no more specific allocator available (see Memory Allocation).
- pSurface points to a VkSurfaceKHR handle in which the created surface is returned.

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkDisplaySurfaceCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSurface must be a pointer to a VkSurfaceKHR handle

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDisplaySurfaceCreateInfoKHR structure is defined as:

```
typedef struct VkDisplaySurfaceCreateInfoKHR {
   VkStructureType
                                    sType;
                                   pNext;
   const void*
   VkDisplaySurfaceCreateFlagsKHR flags;
   VkDisplayModeKHR
                                   displayMode;
   uint32_t
                                   planeIndex;
   uint32_t
                                   planeStackIndex;
   VkSurfaceTransformFlagBitsKHR
                                   transform;
                                   globalAlpha;
   VkDisplayPlaneAlphaFlagBitsKHR alphaMode;
   VkExtent2D
                                    imageExtent;
} VkDisplaySurfaceCreateInfoKHR;
```

- *sType* is the type of this structure and must be VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_ KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use, and must be zero.
- displayMode is the mode to use when displaying this surface.

- planeIndex is the plane on which this surface appears.
- planeStackIndex is the z-order of the plane.
- transform is the transform to apply to the images as part of the scanout operation.
- globalAlpha is the global alpha value. This value is ignored if alphaMode is not VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR.
- alphaMode is the type of alpha blending to use.
- imageSize The size of the presentable images to use with the surface.



Note

Creating a display surface must not modify the state of the displays, planes, or other resources it names. For example, it must not apply the specified mode to be set on the associated display. Application of display configuration occurs as a side effect of presenting to a display surface.

- sType must be VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR
- pNext must be NULL
- flags must be 0
- displayMode must be a valid VkDisplayModeKHR handle
- transform must be a valid VkSurfaceTransformFlagBitsKHR value
- alphaMode must be a valid VkDisplayPlaneAlphaFlagBitsKHR value
- planeIndex must be less than the number of display planes supported by the device as determined by calling vkGetPhysicalDeviceDisplayPlanePropertiesKHR
- If the planeReorderPossible member of the VkDisplayPropertiesKHR structure returned by vkGetPhysicalDeviceDisplayPropertiesKHR for the display corresponding to displayMode is VK_TRUE then planeStackIndex must be less than the number of display planes supported by the device as determined by calling vkGetPhysicalDeviceDisplayPlanePropertiesKHR; otherwise planeStackIndex must equal the currentStackIndex member of VkDisplayPlanePropertiesKHR returned by vkGetPhysicalDeviceDisplayPlanePropertiesKHR for the display plane corresponding to displayMode
- If alphaMode is VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR then globalAlpha must be between 0 and 1, inclusive
- alphaMode must be 0 or one of the bits present in the supportedAlpha member of VkDisplayPlaneCapabilitiesKHR returned by **vkGetDisplayPlaneCapabilitiesKHR** for the display plane corresponding to displayMode
- The width and height members of imageExtent must be less than the maxImageDimensions2D member of VkPhysicalDeviceLimits

Types of alpha blending supported by or used on a display are defined by the bitmask VkDisplayPlaneAlphaFlagBitsKHR, which contains the following values:

```
typedef enum VkDisplayPlaneAlphaFlagBitsKHR {
    VK_DISPLAY_PLANE_ALPHA_OPAQUE_BIT_KHR = 0x00000001,
    VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR = 0x00000002,
    VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR = 0x00000004,
    VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_PREMULTIPLIED_BIT_KHR = 0x00000008,
} VkDisplayPlaneAlphaFlagBitsKHR;
```

These values are described as follows:

- VK_DISPLAY_PLANE_ALPHA_OPAQUE_BIT_KHR: The source image will be treated as opaque.
- VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR: A global alpha value must be specified that will be applied to all pixels in the source image.
- VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR: The alpha value will be determined by the alpha channel of the source image's pixels. If the source format contains no alpha values, no blending will be applied. The source alpha values are not premultiplied into the source image's other color channels.
- VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_PREMULTIPLIED_BIT_KHR: This is equivalent to VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR except the source alpha values are assumed to be premultiplied into the source image's other color channels.

29.4 Querying for WSI Support

Not all physical devices will include WSI support. Within a physical device, not all queue families will support presentation. WSI support and compatibility can be determined in a platform-neutral manner (which determines support for presentation to a particular surface object) and additionally may be determined in platform-specific manners (which determine support for presentation on the specified physical device but don't guarantee support for presentation to a particular surface object).

To determine whether a queue family of a physical device supports presentation to a given surface, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family.
- surface is the surface.
- pSupported is a pointer to a VkBool32, which is set to VK TRUE to indicate support, and VK FALSE otherwise.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- surface must be a valid VkSurfaceKHR handle
- pSupported must be a pointer to a VkBool32 value
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_SURFACE_LOST_KHR

29.4.1 Android Platform

On Android, all physical devices and queue families must be capable of presentation with any native window. As a result there is no Android-specific query for these capabilities.

29.4.2 Mir Platform

To determine whether a queue family of a physical device supports presentation to the Mir compositor, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family index.
- connection is a pointer to the MirConnection, and identifies the desired Mir compositor.

This platform-specific function can be called prior to creating a surface.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- connection must be a pointer to a MirConnection value
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

29.4.3 Wayland Platform

To determine whether a queue family of a physical device supports presentation to a Wayland compositor, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family index.
- display is a pointer to the wl_display associated with a Wayland compositor.

This platform-specific function can be called prior to creating a surface.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- display must be a pointer to a wl_display value
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

29.4.4 Win32 Platform

To determine whether a queue family of a physical device supports presentation to the Microsoft Windows desktop, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family index.

This platform-specific function can be called prior to creating a surface.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

29.4.5 XCB Platform

To determine whether a queue family of a physical device supports presentation to an X11 server, using the XCB client-side library, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family index.
- connection is a pointer to an xcb_connection_t to the X server. visual_id is an X11 visual (xcb_visualid_t).

This platform-specific function can be called prior to creating a surface.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- connection must be a pointer to a xcb_connection_t value
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

29.4.6 Xlib Platform

To determine whether a queue family of a physical device supports presentation to an X11 server, using the Xlib client-side library, call:

- physicalDevice is the physical device.
- queueFamilyIndex is the queue family index.
- dpy is a pointer to an Xlib Display connection to the server.
- visualId is an X11 visual (VisualID).

This platform-specific function can be called prior to creating a surface.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- dpy must be a pointer to a Display value
- queueFamilyIndex must be less than pQueueFamilyPropertyCount returned by
 vkGetPhysicalDeviceQueueFamilyProperties for the given physicalDevice

29.5 Surface Queries

To query the basic capabilities of a surface, needed in order to create a swapchain, call:

- physicalDevice is the physical device that will be associated with the swapchain to be created, as described for vkCreateSwapchainKHR.
- surface is the surface that will be associated with the swapchain.
- pSurfaceCapabilities is a pointer to an instance of the VkSurfaceCapabilitiesKHR structure in which the capabilities are returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- surface must be a valid VkSurfaceKHR handle
- pSurfaceCapabilities must be a pointer to a VkSurfaceCapabilitiesKHR structure

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_SURFACE_LOST_KHR

The VkSurfaceCapabilitiesKHR structure is defined as:

```
typedef struct VkSurfaceCapabilitiesKHR {
   uint32_t
                                    minImageCount;
   uint32_t
                                    maxImageCount;
   VkExtent2D
                                   currentExtent;
   VkExtent2D
                                   minImageExtent;
   VkExtent2D
                                   maxImageExtent;
   uint32_t
                                   maxImageArrayLayers;
   VkSurfaceTransformFlagsKHR
                                   supportedTransforms;
   VkSurfaceTransformFlagBitsKHR
                                   currentTransform;
   VkCompositeAlphaFlagsKHR
                                   supportedCompositeAlpha;
   VkImageUsageFlags
                                    supportedUsageFlags;
} VkSurfaceCapabilitiesKHR;
```

- minImageCount is the minimum number of images the specified device supports for a swapchain created for the surface.
- maxImageCount is the maximum number of images the specified device supports for a swapchain created for the surface. A value of 0 means that there is no limit on the number of images, though there may be limits related to the total amount of memory used by swapchain images.
- currentExtent is the current width and height of the surface, or the special value (0xFFFFFFFFF,0xFFFFFFFFF) indicating that the surface size will be determined by the extent of a swapchain targeting the surface.

- minImageExtent contains the smallest valid swapchain extent for the surface on the specified device.
- maxImageExtent contains the largest valid swapchain extent for the surface on the specified device.
- maxImageArrayLayers is the maximum number of layers swapchain images can have for a swapchain created for this device and surface.
- supportedTransforms is a bitmask of VkSurfaceTransformFlagBitsKHR, describing the presentation transforms supported for the surface on the specified device.
- currentTransform is a bitmask of VkSurfaceTransformFlagBitsKHR, describing the surface's current transform relative to the presentation engine's natural orientation.
- supportedCompositeAlpha is a bitmask of VkCompositeAlphaFlagBitsKHR, representing the alpha compositing modes supported by the presentation engine for the surface on the specified device. Opaque composition can be achieved in any alpha compositing mode by either using a swapchain image format that has no alpha component, or by ensuring that all pixels in the swapchain images have an alpha value of 1.0.
- supportedUsageFlags is a bitmask of VkImageUsageFlagBits representing the ways the application can use the presentable images of a swapchain created for the surface on the specified device. VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT must be included in the set but implementations may support additional usages.



Note

Formulas such as min(N, maxImageCount) are not correct, since maxImageCount may be zero.

Valid Usage

- supportedTransforms must be a valid combination of VkSurfaceTransformFlagBitsKHR values
- currentTransform must be a valid VkSurfaceTransformFlagBitsKHR value
- supportedCompositeAlpha must be a valid combination of VkCompositeAlphaFlagBitsKHR values
- supportedUsageFlags must be a valid combination of VkImageUsageFlagBits values

The supportedTransforms and currentTransform members are of type VkSurfaceTransformFlaqBitsKHR, which contains the following values:

```
typedef enum VkSurfaceTransformFlagBitsKHR {
    VK_SURFACE_TRANSFORM_IDENTITY_BIT_KHR = 0x00000001,
    VK_SURFACE_TRANSFORM_ROTATE_90_BIT_KHR = 0x00000002,
    VK_SURFACE_TRANSFORM_ROTATE_180_BIT_KHR = 0x00000004,
    VK_SURFACE_TRANSFORM_ROTATE_270_BIT_KHR = 0x00000008,
    VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_BIT_KHR = 0x00000010,
    VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_90_BIT_KHR = 0x00000020,
    VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_180_BIT_KHR = 0x000000040,
    VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_270_BIT_KHR = 0x000000080,
    VK_SURFACE_TRANSFORM_INHERIT_BIT_KHR = 0x00000100,
} VkSurfaceTransformFlagBitsKHR;
```

These values are described as follows:

- VK_SURFACE_TRANSFORM_IDENTITY_BIT_KHR: The image content is presented without being transformed.
- VK_SURFACE_TRANSFORM_ROTATE_90_BIT_KHR: The image content is rotated 90 degrees clockwise.
- VK_SURFACE_TRANSFORM_ROTATE_180_BIT_KHR: The image content is rotated 180 degrees clockwise.
- VK_SURFACE_TRANSFORM_ROTATE_270_BIT_KHR: The image content is rotated 270 degrees clockwise.
- VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_BIT_KHR: The image content is mirrored horizontally.
- VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_90_BIT_KHR: The image content is mirrored horizontally, then rotated 90 degrees clockwise.
- VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_180_BIT_KHR: The image content is mirrored horizontally, then rotated 180 degrees clockwise.
- VK_SURFACE_TRANSFORM_HORIZONTAL_MIRROR_ROTATE_270_BIT_KHR: The image content is mirrored horizontally, then rotated 270 degrees clockwise.
- VK_SURFACE_TRANSFORM_INHERIT_BIT_KHR: The presentation transform is not specified, and is instead determined by platform-specific considerations and mechanisms outside Vulkan.

The <code>supportedCompositeAlpha</code> member is of type <code>VkCompositeAlphaFlagBitsKHR</code>, which contains the following values:

```
typedef enum VkCompositeAlphaFlagBitsKHR {
    VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR = 0x00000001,
    VK_COMPOSITE_ALPHA_PRE_MULTIPLIED_BIT_KHR = 0x00000002,
    VK_COMPOSITE_ALPHA_POST_MULTIPLIED_BIT_KHR = 0x000000004,
    VK_COMPOSITE_ALPHA_INHERIT_BIT_KHR = 0x000000008,
} VkCompositeAlphaFlagBitsKHR;
```

These values are described as follows:

- VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR: The alpha channel, if it exists, of the images is ignored in the compositing process. Instead, the image is treated as if it has a constant alpha of 1.0.
- VK_COMPOSITE_ALPHA_PRE_MULTIPLIED_BIT_KHR: The alpha channel, if it exists, of the images is respected in the compositing process. The non-alpha channels of the image are expected to already be multiplied by the alpha channel by the application.
- VK_COMPOSITE_ALPHA_POST_MULTIPLIED_BIT_KHR: The alpha channel, if it exists, of the images is respected in the compositing process. The non-alpha channels of the image are not expected to already be multiplied by the alpha channel by the application; instead, the compositor will multiply the non-alpha channels of the image by the alpha channel during compositing.
- VK_COMPOSITE_ALPHA_INHERIT_BIT_KHR: The way in which the presentation engine treats the alpha channel in the images is unknown to the Vulkan API. Instead, the application is responsible for setting the composite alpha blending mode using native window system commands. If the application does not set the blending mode using native window system commands, then a platform-specific default will be used.

To query the supported swapchain format-color space pairs for a surface, call:

- physicalDevice is the physical device that will be associated with the swapchain to be created, as described for vkCreateSwapchainKHR.
- surface is the surface that will be associated with the swapchain.
- pSurfaceFormatCount is a pointer to an integer related to the number of format pairs available or queried, as described below.
- pSurfaceFormats is either NULL or a pointer to an array of VkSurfaceFormatKHR structures.

If pSurfaceFormats is NULL, then the number of format pairs supported for the given surface is returned in pSurfaceFormatCount. Otherwise, pSurfaceFormatCount must point to a variable set by the user to the number of elements in the pSurfaceFormats array, and on return the variable is overwritten with the number of structures actually written to pSurfaceFormats. If the value of pSurfaceFormatCount is less than the number of format pairs supported, at most pSurfaceFormatCount structures will be written. If pSurfaceFormatCount is smaller than the number of format pairs supported for the given surface, VK_INCOMPLETE will be returned instead of VK_SUCCESS to indicate that not all the available values were returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- surface must be a valid VkSurfaceKHR handle
- pSurfaceFormatCount must be a pointer to a uint32_t value
- If the value referenced by <code>pSurfaceFormatCount</code> is not 0, and <code>pSurfaceFormats</code> is not <code>NULL</code>, <code>pSurfaceFormats</code> must be a pointer to an array of <code>pSurfaceFormatCount</code> <code>VkSurfaceFormatKHR</code> structures

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_SURFACE_LOST_KHR

The VkSurfaceFormatKHR structure is defined as:

- format is a VkFormat that is compatible with the specified surface.
- colorSpace is a presentation VkColorSpaceKHR that is compatible with the surface.

Valid Usage

- format must be a valid VkFormat value
- colorSpace must be a valid VkColorSpaceKHR value

While the format of a presentable image refers to the encoding of each pixel, the colorSpace determines how the presentation engine interprets the pixel values. A color space in this document refers to a specific combination of color model (i.e. RGB, YUV, HSL etc.), the actual color space (defined by the chromaticities of its primaries and a white point in CIE Lab), and a transfer function that is applied before storing or transmitting color data in the given color space.

The VkColorSpaceKHR is defined as follows:

```
typedef enum VkColorSpaceKHR {
    VK_COLOR_SPACE_SRGB_NONLINEAR_KHR = 0,
} VkColorSpaceKHR;
```

• VK_COLOR_SPACE_SRGB_NONLINEAR_KHR: The presentation engine supports the sRGB color space.



Note

If pSurfaceFormats includes just one entry, whose value for format is VK_FORMAT_UNDEFINED, surface has no preferred format. In this case, the application can use any valid VkFormat value.



Note

In the initial release of the VK_KHR_surface and VK_KHR_swapchain extensions, the token VK_COLORSPA CE_SRGB_NONLINEAR_KHR was used. Starting in the May 13, 2016 updates to the extension branches, matching release 1.0.13 of the core API specification, VK_COLOR_SPACE_SRGB_NONLINEAR_KHR is used instead for consistency with Vulkan naming rules. The older enum is still available for backwards compatibility.

To query the supported presentation modes for a surface, call:

- physicalDevice is the physical device that will be associated with the swapchain to be created, as described for vkCreateSwapchainKHR.
- surface is the surface that will be associated with the swapchain.
- pPresentModeCount is a pointer to an integer related to the number of presentation modes available or queried, as described below.
- pPresentModes is either NULL or a pointer to an array of VkPresentModeKHR structures.

If <code>pPresentModes</code> is <code>NULL</code>, then the number of presentation modes supported for the given <code>surface</code> is returned in <code>pPresentModeCount</code>. Otherwise, <code>pPresentModeCount</code> must point to a variable set by the user to the number of elements in the <code>pPresentModes</code> array, and on return the variable is overwritten with the number of values actually written to <code>pPresentModes</code>. If the value of <code>pPresentModeCount</code> is less than the number of presentation modes supported, at most <code>pPresentModeCount</code> values will be written. If <code>pPresentModeCount</code> is smaller than the number of presentation modes supported for the given <code>surface</code>, <code>VK_INCOMPLETE</code> will be returned instead of <code>VK_SUCCESS</code> to indicate that not all the available values were returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- surface must be a valid VkSurfaceKHR handle
- pPresentModeCount must be a pointer to a uint32_t value
- If the value referenced by pPresentModeCount is not 0, and pPresentModes is not NULL, pPresentModes must be a pointer to an array of pPresentModeCount VkPresentModeKHR values

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

```
• VK_ERROR_OUT_OF_HOST_MEMORY
```

- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_SURFACE_LOST_KHR

The definition of VkPresentModeKHR is:

```
typedef enum VkPresentModeKHR {
    VK_PRESENT_MODE_IMMEDIATE_KHR = 0,
    VK_PRESENT_MODE_MAILBOX_KHR = 1,
    VK_PRESENT_MODE_FIFO_KHR = 2,
    VK_PRESENT_MODE_FIFO_RELAXED_KHR = 3,
} VkPresentModeKHR;
```

- VK_PRESENT_MODE_IMMEDIATE_KHR: The presentation engine does not wait for a vertical blanking period to
 update the current image, meaning this mode may result in visible tearing. No internal queuing of presentation
 requests is needed, as the requests are applied immediately.
- VK_PRESENT_MODE_MAILBOX_KHR: The presentation engine waits for the next vertical blanking period to update
 the current image. Tearing cannot be observed. An internal single-entry queue is used to hold pending presentation
 requests. If the queue is full when a new presentation request is received, the new request replaces the existing entry,
 and any images associated with the prior entry become available for re-use by the application. One request is removed
 from the queue and processed during each vertical blanking period in which the queue is non-empty.
- VK_PRESENT_MODE_FIFO_KHR: The presentation engine waits for the next vertical blanking period to update the current image. Tearing cannot be observed. An internal queue is used to hold pending presentation requests. New requests are appended to the end of the queue, and one request is removed from the beginning of the queue and processed during each vertical blanking period in which the queue is non-empty. This is the only value of presentMode that is required to be supported.
- VK_PRESENT_MODE_FIFO_RELAXED_KHR: The presentation engine generally waits for the next vertical blanking period to update the current image. If a vertical blanking period has already passed since the last update of the current image then the presentation engine does not wait for another vertical blanking period for the update, meaning this mode may result in visible tearing in this case. This mode is useful for reducing visual stutter with an application that will mostly present a new image before the next vertical blanking period, but may occasionally be late, and present a new image just after the next vertical blanking period. An internal queue is used to hold pending presentation requests. New requests are appended to the end of the queue, and one request is removed from the beginning of the queue and processed during or after each vertical blanking period in which the queue is non-empty.

not8

Note

For reference, the mode indicated by VK_PRESENT_MODE_FIFO_KHR is equivalent to the behavior of {wgl|g|X|eg|}SwapBuffers with a swap interval of 1, while the mode indicated by VK_PRESENT_MODE_FIFO_RELAXED_KHR is equivalent to the behavior of {wgl|g|X}SwapBuffers with a swap interval of -1 (from the {WGL|GLX}_EXT_swap_control_tear extensions).

29.6 WSI Swapchain

A VkSwapchainKHR object (a.k.a. swapchain) provides the ability to present rendering results to a surface. A swapchain is an abstraction for an array of presentable images that are associated with a surface. The swapchain images

are represented by VkImage objects created by the platform. One image (which can be an array image for multiview/stereoscopic-3D surfaces) is displayed at a time, but multiple images can be queued for presentation. An application renders to the image, and then queues the image for presentation to the surface. A native window cannot be associated with more than one swapchain at a time. Further, swapchains cannot be created for native windows that have a non-Vulkan graphics API surface associated with them.

The presentation engine is an abstraction for the platform's compositor or hardware/software display engine.



Note

The presentation engine may be synchronous or asynchronous with respect to the application and/or logical device

Some implementations may use the device's graphics queue or dedicated presentation hardware to perform presentation.

The presentable images of a swapchain are owned by the presentation engine. An application can acquire use of a presentable image from the presentation engine. Use of a presentable image must occur only after the image is returned by **vkAcquireNextImageKHR**, and before it is presented by **vkQueuePresentKHR**. This includes transitioning the image layout and rendering commands.

An application can acquire use of a presentable image with **vkAcquireNextImageKHR**. After acquiring a presentable image and before modifying it, the application must use a synchronization primitive to ensure that the presentation engine has finished reading from the image. The application can then transition the image's layout, queue rendering commands to it, etc. Finally, the application presents the image with **vkQueuePresentKHR**, which releases the acquisition of the image.

The presentation engine controls the order in which presentable images are acquired for use by the application.



Note

This allows the platform to handle situations which require out-of-order return of images after presentation. At the same time, it allows the application to generate command buffers referencing all of the images in the swapchain at initialization time, rather than in its main loop.

How this all works is described below.

To create a swapchain, call:

- device is the device to create the swapchain for.
- pCreateInfo is a pointer to an instance of the VkSwapchainCreateInfoKHR structure specifying the parameters of the created swapchain.
- pAllocator is the allocator used for host memory allocated for the swapchain object when there is no more specific allocator available (see Memory Allocation).
- pSwapchain is a pointer to a VkSwapchainKHR handle in which the created swapchain object will be returned.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfo must be a pointer to a valid VkSwapchainCreateInfoKHR structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSwapchain must be a pointer to a VkSwapchainKHR handle

Host Synchronization

- Host access to pCreateInfo.surface must be externally synchronized
- Host access to pCreateInfo.oldSwapchain must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST
- VK_ERROR_SURFACE_LOST_KHR
- VK_ERROR_NATIVE_WINDOW_IN_USE_KHR

The VkSwapchainCreateInfoKHR structure is defined as:

```
minImageCount;
   uint32_t
   VkFormat.
                                     imageFormat;
   VkColorSpaceKHR
                                     imageColorSpace;
   VkExtent2D
                                     imageExtent;
   uint32_t
                                     imageArrayLayers;
   VkImageUsageFlags
                                     imageUsage;
   VkSharingMode
                                     imageSharingMode;
   uint32_t
                                     queueFamilyIndexCount;
   const uint32_t*
                                     pQueueFamilyIndices;
   VkSurfaceTransformFlagBitsKHR
                                    preTransform;
   VkCompositeAlphaFlagBitsKHR
                                    compositeAlpha;
   VkPresentModeKHR
                                     presentMode;
   VkBool32
                                     clipped;
   VkSwapchainKHR
                                     oldSwapchain;
} VkSwapchainCreateInfoKHR;
```

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- flags is reserved for future use, and must be zero.
- surface is the surface that the swapchain will present images to.
- minImageCount is the minimum number of presentable images that the application needs. The platform will either create the swapchain with at least that many images, or will fail to create the swapchain.
- imageFormat is a VkFormat that is valid for swapchains on the specified surface.
- imageColorSpace is a VkColorSpaceKHR that is valid for swapchains on the specified surface.
- *imageExtent* is the size (in pixels) of the swapchain. Behavior is platform-dependent when the image extent does not match the surface's *currentExtent* as returned by **vkGetPhysicalDeviceSurfaceCapabilitiesKHR**.
- imageArrayLayers is the number of views in a multiview/stereo surface. For non-stereoscopic-3D applications, this value is 1.
- *imageUsage* is a bitmask of VkImageUsageFlagBits, indicating how the application will use the swapchain's presentable images.
- imageSharingMode is the sharing mode used for the images of the swapchain.
- queueFamilyIndexCount is the number of queue families having access to the images of the swapchain in case imageSharingMode is VK_SHARING_MODE_CONCURRENT.
- pQueueFamilyIndices is an array of queue family indices having access to the images of the swapchain in case imageSharingMode is VK_SHARING_MODE_CONCURRENT.
- preTransform is a bitmask of VkSurfaceTransformFlagBitsKHR, describing the transform, relative to the presentation engine's natural orientation, applied to the image content prior to presentation. If it does not match the currentTransform value returned by vkGetPhysicalDeviceSurfaceCapabilitiesKHR, the presentation engine will transform the image content as part of the presentation operation.
- compositeAlpha is a bitmask of VkCompositeAlphaFlagBitsKHR, indicating the alpha compositing mode to use when this surface is composited together with other surfaces on certain window systems.
- presentMode is the presentation mode the swapchain will use. A swapchain's present mode determines how incoming present requests will be processed and queued internally.

- clipped indicates whether the Vulkan implementation is allowed to discard rendering operations that affect regions of
 the surface which aren't visible.
 - If set to VK_TRUE, the presentable images associated with the swapchain may not own all of their pixels. Pixels in the presentable images that correspond to regions of the target surface obscured by another window on the desktop or subject to some other clipping mechanism will have undefined content when read back. Pixel shaders may not execute for these pixels, and thus any side affects they would have had will not occur.
 - If set to VK_FALSE, presentable images associated with the swapchain will own all the pixels they contain. Setting
 this value to VK_TRUE does not guarantee any clipping will occur, but allows more optimal presentation methods to
 be used on some platforms.



Note

Applications should set this value to VK_TRUE if they do not expect to read back the content of presentable images before presenting them or after reacquiring them and if their pixel shaders do not have any side effects that require them to run for all pixels in the presentable image.

• oldSwapchain, if not VK_NULL_HANDLE, specifies the swapchain that will be replaced by the new swapchain being created. The new swapchain will be a descendant of oldSwapchain. Further, any descendants of the new swapchain will also be descendants of oldSwapchain. Upon calling vkCreateSwapchainKHR with a oldSwapchain that is not VK_NULL_HANDLE, any images not acquired by the application may be freed by the implementation, which may occur even if creation of the new swapchain fails. The application must destroy the old swapchain to free all memory associated with the old swapchain. The application must wait for the completion of any outstanding rendering to images it currently has acquired at the time the swapchain is destroyed. The application can continue to present any images it acquired and has not yet presented using the old swapchain, as long as it has not entered a state that causes it to return VK_ERROR_OUT_OF_DATE_KHR. However, the application cannot acquire any more images from the old swapchain regardless of whether or not creation of the new swapchain succeeds.

Valid Usage

- sType must be $VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR$
- pNext must be NULL
- flags must be 0
- surface must be a valid VkSurfaceKHR handle
- imageFormat must be a valid VkFormat value
- imageColorSpace must be a valid VkColorSpaceKHR value
- imageUsage must be a valid combination of VkImageUsageFlagBits values
- imageUsage must not be 0
- imageSharingMode must be a valid VkSharingMode value
- preTransform must be a valid VkSurfaceTransformFlagBitsKHR value

- compositeAlpha must be a valid VkCompositeAlphaFlagBitsKHR value
- presentMode must be a valid VkPresentModeKHR value
- If oldSwapchain is not VK_NULL_HANDLE, oldSwapchain must be a valid VkSwapchainKHR handle
- If oldSwapchain is a valid handle, it must have been created, allocated, or retrieved from surface
- surface must be a surface that is supported by the device as determined using vkGetPhysicalDeviceSurfaceSupportKHR
- The native window referred to by *surface* must not already be associated with a swapchain other than *oldSwapchain*, or with a non-Vulkan graphics API surface
- minImageCount must be greater than or equal to the value returned in the minImageCount member of the VkSurfaceCapabilitiesKHR structure returned by
 - $\textbf{vkGetPhysicalDeviceSurfaceCapabilitiesKHR}\ for\ the\ surface$
- minImageCount must be less than or equal to the value returned in the maxImageCount member of the VkSurfaceCapabilitiesKHR structure returned by
 - **vkGetPhysicalDeviceSurfaceCapabilitiesKHR** for the surface if the returned maxImageCount is not zero
- imageFormat and imageColorspace must match the format and colorSpace members, respectively, of one of the VkSurfaceFormatKHR structures returned by vkGetPhysicalDeviceSurfaceFormatsKHR for the surface
- imageExtent must be between minImageExtent and maxImageExtent, inclusive, where minImageExtent and maxImageExtent are members of the VkSurfaceCapabilitiesKHR structure returned by vkGetPhysicalDeviceSurfaceCapabilitiesKHR for the surface
- imageArrayLayers must be greater than 0 and less than or equal to the maxImageArrayLayers member of the VkSurfaceCapabilitiesKHR structure returned by
- ${\tt vkGetPhysicalDeviceSurfaceCapabilitiesKHR}\ for\ the\ surface$
- imageUsage must be a subset of the supported usage flags present in the supportedUsageFlags member of the VkSurfaceCapabilitiesKHR structure returned by
 - vkGetPhysicalDeviceSurfaceCapabilitiesKHR for the surface
- If imageSharingMode is VK_SHARING_MODE_CONCURRENT, pQueueFamilyIndices must be a pointer to an array of queueFamilyIndexCount uint32_t values
- If imageSharingMode is VK_SHARING_MODE_CONCURRENT, queueFamilyIndexCount must be greater than 1
- preTransform must be one of the bits present in the supportedTransforms member of the VkSurfaceCapabilitiesKHR structure returned by
 - vkGetPhysicalDeviceSurfaceCapabilitiesKHR for the surface
- compositeAlpha must be one of the bits present in the supportedCompositeAlpha member of the VkSurfaceCapabilitiesKHR structure returned by
 - vkGetPhysicalDeviceSurfaceCapabilitiesKHR for the surface
- presentMode must be one of the VkPresentModeKHR values returned by vkGetPhysicalDeviceSurfacePresentModesKHR for the surface

As mentioned above, if **vkCreateSwapchainKHR** succeeds, it will return a handle to a swapchain that contains an array of at least *minImageCount* presentable images.

While acquired by the application, swapchain images can be used in any way that equivalent non-swapchain images can be used. A swapchain image is equivalent to a non-swapchain image created with the following VkImageCreateInfo parameters:

VkImageCreateInfo Field	Value
flags	0
imageType	VK_IMAGE_TYPE_2D
format	pCreateInfo->imageFormat
extent	{pCreateInfo->imageExtent.width,
	<pre>pCreateInfo->imageExtent.height, 1}</pre>
mipLevels	1
arrayLayers	pCreateInfo->imageArrayLayers
samples	VK_SAMPLE_COUNT_1_BIT
tiling	VK_IMAGE_TILING_OPTIMAL
usage	pCreateInfo->imageUsage
sharingMode	pCreateInfo->imageSharingMode
queueFamilyIndexCount	pCreateInfo->queueFamilyIndexCount
<i>pQueueFamilyIndices</i>	pCreateInfo->pQueueFamilyIndices
initialLayout	VK_IMAGE_LAYOUT_UNDEFINED

The VkSurfaceKHR associated with a swapchain must not be destroyed until after the swapchain is destroyed.

Like core functions, several WSI functions, including vkCreateSwapchainKHR return VK_ERROR_DEVICE_ LOST if the logical device was lost. See Lost Device. As with most core objects, VkSwapchainKHR is a child of the device and is affected by the lost state; it must be destroyed before destroying the VkDevice. However, VkSurfaceKHR is not a child of any VkDevice and is not otherwise affected by the lost device. After successfully recreating a VkDevice, the same VkSurfaceKHR can be used to create a new VkSwapchainKHR, provided the previous one was destroyed.



Note

As mentioned in Lost Device, after a lost device event, the <code>VkPhysicalDevice</code> may also be lost. If other <code>VkPhysicalDevice</code> are available, they can be used together with the same <code>VkSurfaceKHR</code> to create the new <code>VkSwapchainKHR</code>, however the application must query the surface capabilities again, because they may differ on a per-physical device basis.

To destroy a swapchain object call:

- device is the VkDevice associated with swapchain.
- swapchain is the swapchain to destroy.
- pAllocator is the allocator used for host memory allocated for the swapchain object when there is no more specific allocator available (see Memory Allocation).

swapchain and all associated VkImage handles are destroyed, and must not be acquired or used any more by the application. The memory of each VkImage will only be freed after that image is no longer used by the platform. For example, if one image of the swapchain is being displayed in a window, the memory for that image may not be freed until the window is destroyed, or another swapchain is created for the window. Destroying the swapchain does not invalidate the parent VkSurfaceKHR, and a new swapchain can be created with it.

If a swapchain associated with a display surface is destroyed and there are no valid descendants of that swapchain, the implementation must either revert any display resources modified by presenting images with the swapchain to their state prior to the first present performed with the swapchain and its ancestors, or leave such resources in their current state.

Valid Usage

- device must be a valid VkDevice handle
- If swapchain is not VK_NULL_HANDLE, swapchain must be a valid VkSwapchainKHR handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- All uses of presentable images acquired from swapchain must have completed execution
- If VkAllocationCallbacks were provided when *swapchain* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when swapchain was created, pAllocator must be NULL

Host Synchronization

• Host access to swapchain must be externally synchronized

When the VK_KHR_display_swapchain extension is enabled, multiple swapchains that share presentable images are created by calling:

- device is the device to create the swapchains for.
- swapchainCount is the number of swapchains to create.
- pCreateInfos is a pointer to an array of VkSwapchainCreateInfoKHR structures specifying the parameters of the created swapchains.

- pAllocator is the allocator used for host memory allocated for the swapchain objects when there is no more specific allocator available (see Memory Allocation).
- pSwapchains is a pointer to an array of VkSwapchainKHR handles in which the created swapchain objects will be returned.

vkCreateSharedSwapchains is similar to vkCreateSwapchainKHR, except that it takes an array of VkSwapchainCreateInfoKHR structures, and returns an array of swapchain objects.

The swapchain creation parameters that affect the properties and number of presentable images must match between all the swapchains. If the displays used by any of the swapchains do not use the same presentable image layout or are incompatible in a way that prevents sharing images, swapchain creation will fail with the result code VK_ERROR_INCOMPATIBLE_DISPLAY_KHR. If any error occurs, no swapchains will be created. Images presented to multiple swapchains must be re-acquired from all of them before transitioning away from VK_IMAGE_LAYOUT_PRESENT_SRC_KHR. After destroying one or more of the swapchains, the remaining swapchains and the presentable images can continue to be used.

Valid Usage

- device must be a valid VkDevice handle
- pCreateInfos must be a pointer to an array of swapchainCount valid VkSwapchainCreateInfoKHR structures
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pSwapchains must be a pointer to an array of swapchainCount VkSwapchainKHR handles
- swapchainCount must be greater than 0

Host Synchronization

- Host access to pCreateInfos[].surface must be externally synchronized
- Host access to pCreateInfos[].oldSwapchain must be externally synchronized

Return	Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK ERROR OUT OF DEVICE MEMORY
- VK ERROR INCOMPATIBLE DISPLAY KHR
- VK_ERROR_DEVICE_LOST
- VK_ERROR_SURFACE_LOST_KHR

To obtain the array of presentable images associated with a swapchain, call:

- device is the device associated with swapchain.
- swapchain is the swapchain to query.
- pSwapchainImageCount is a pointer to an integer related to the number of swapchain images available or queried, as described below.
- pSwapchainImages is either NULL or a pointer to an array of VkImage handles.

If <code>pSwapchainImages</code> is <code>NULL</code>, then the number of presentable images for <code>swapchain</code> is returned in <code>pSwapchainImageCount</code>. Otherwise, <code>pSwapchainImageCount</code> must point to a variable set by the user to the number of elements in the <code>pSwapchainImages</code> array, and on return the variable is overwritten with the number of structures actually written to <code>pSwapchainImages</code>. If the value of <code>pSwapchainImageCount</code> is less than the number of presentable images for <code>swapchain</code>, at most <code>pSwapchainImageCount</code> structures will be written. If <code>pSwapchainImageCount</code> is smaller than the number of presentable images for <code>swapchain</code>, <code>VK_INCOMPLETE</code> will be returned instead of <code>VK_SUCCESS</code> to indicate that not all the available values were returned.

Valid Usage

- device must be a valid VkDevice handle
- swapchain must be a valid VkSwapchainKHR handle
- pSwapchainImageCount must be a pointer to a uint32_t value
- If the value referenced by pSwapchainImageCount is not 0, and pSwapchainImages is not NULL, pSwapchainImages must be a pointer to an array of pSwapchainImageCount VkImage handles

Return Codes

Success

- VK_SUCCESS
- VK INCOMPLETE

Failure

- VK ERROR OUT OF HOST MEMORY
- VK ERROR OUT OF DEVICE MEMORY



Note

By knowing all presentable images used in the swapchain, the application can create command buffers that reference these images prior to entering its main rendering loop.

The implementation will have already allocated and bound the memory backing the VkImages returned by **vkGetSwapchainImagesKHR**. The memory for each image will not alias with the memory for other images or with any VkDeviceMemory object. As such, performing any operation affecting the binding of memory to a presentable image results in undefined behavior. All presentable images are initially in the VK_IMAGE_LAYOUT_UNDEFINED layout, thus before using presentable images, the application must transition them to a valid layout for the intended use.

Further, the lifetime of presentable images is controlled by the implementation so destroying a presentable image with vkDestroyImage results in undefined behavior. See vkDestroySwapchainKHR for further details on the lifetime of presentable images.

To acquire an available presentable image to use, and retrieve the index of that image, call:

- device is the device associated with swapchain.
- swapchain is the swapchain from which an image is being acquired.
- timeout indicates how long the function waits, in nanoseconds, if no image is available.
- semaphore is VK_NULL_HANDLE or a semaphore to signal.
- fence is VK NULL HANDLE or a fence to signal.
- pImageIndex is a pointer to a uint32_t that is set to the index of the next image to use (i.e. an index into the array of images returned by vkGetSwapchainImagesKHR).

Valid Usage

- device must be a valid VkDevice handle
- swapchain must be a valid VkSwapchainKHR handle
- If semaphore is not VK_NULL_HANDLE, semaphore must be a valid VkSemaphore handle
- If fence is not VK_NULL_HANDLE, fence must be a valid VkFence handle
- pImageIndex must be a pointer to a uint32_t value
- If semaphore is a valid handle, it must have been created, allocated, or retrieved from device
- If fence is a valid handle, it must have been created, allocated, or retrieved from device
- If semaphore is not VK_NULL_HANDLE it must be unsignaled
- If fence is not VK_NULL_HANDLE it must be unsignaled and must not be associated with any other queue command that has not yet completed execution on that queue

Host Synchronization

- Host access to swapchain must be externally synchronized
- Host access to semaphore must be externally synchronized
- Host access to fence must be externally synchronized

Return Codes

Success

- VK_SUCCESS
- VK_TIMEOUT
- VK_NOT_READY
- VK SUBOPTIMAL KHR

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST
- VK_ERROR_OUT_OF_DATE_KHR
- VK_ERROR_SURFACE_LOST_KHR

When successful, **vkAcquireNextImageKHR** acquires a presentable image that the application can use, and sets pImageIndex to the index of that image. The presentation engine may not have finished reading from the image at the time it is acquired, so the application must use semaphore and/or fence to ensure that the image layout and contents are not modified until the presentation engine reads have completed.

As mentioned above, the presentation engine controls the order in which presentable images are made available to the application. This allows the platform to handle special situations. The order in which images are acquired is implementation-dependent. Images may be acquired in a seemingly random order that is not a simple round-robin.

If a swapchain has enough presentable images, applications can acquire multiple images without an intervening **vkQueuePresentKHR**. Applications can present images in a different order than the order in which they were acquired.

If timeout is 0, **vkAcquireNextImageKHR** will not block, but will either succeed or return VK_NOT_READY. If timeout is UINT64_MAX, the function will not return until an image is acquired from the presentation engine. Other values for timeout will cause the function to return when an image becomes available, or when the specified number of nanoseconds have passed (in which case it will return VK_TIMEOUT). An error can also cause **vkAcquireNextImageKHR** to return early.

Note



As mentioned above, the presentation engine may be asynchronous with respect to the application and/or logical device. **vkAcquireNextImageKHR** may return as soon as it can identify which image will be acquired, and can guarantee that <code>semaphore</code> and <code>fence</code> will be signaled by the presentation engine; and may not successfully return sooner. The application uses <code>timeout</code> to specify how long **vkAcquireNextImageKHR** waits for an image to become acquired.

Applications cannot rely on **vkAcquireNextImageKHR** blocking in order to meter their rendering speed. Various factors can interrupt **vkAcquireNextImageKHR** from blocking.



Note

For example, if an error occurs, **vkAcquireNextImageKHR** may return even though no image is available. As another example, some presentation engines are able to enqueue an unbounded number of presentation and acquire next image operations such that **vkAcquireNextImageKHR** never needs to wait for completion of outstanding present operations before returning.

The availability of presentable images is influenced by factors such as the implementation of the presentation engine, the VkPresentModeKHR being used, the number of images in the swapchain, the number of images that the application has acquired at any given time, and the performance of the application. The value of

VkSurfaceCapabilitiesKHR::minImageCount indicates how many images must be in the swapchain in order for **vkAcquireNextImageKHR** to acquire an image if the application currently has no acquired images.

Let *n* be the total number of images in the swapchain, *m* be the value of

VkSurfaceCapabilitiesKHR::minImageCount, and a be the number of presentable images that the application has currently acquired (i.e. images acquired with **vkAcquireNextImageKHR**, but not yet presented with **vkQueuePresentKHR**). **vkAcquireNextImageKHR** can always succeed if $a \le n - m$ at the time **vkAcquireNextImageKHR** is called. **vkAcquireNextImageKHR** should not be called if a > n - m with a timeout of **UINT64_MAX**; in such a case, **vkAcquireNextImageKHR** may block indefinitely.

Note



For example, if the <code>minImageCount</code> member of <code>VkSurfaceCapabilitiesKHR</code> is 2, and the application creates a swapchain with 2 presentable images, the application can acquire one image, and must present it before trying to acquire another image.

If we modify this example so that the application wishes to acquire up to 3 presentable images simultaneously, it must request a minimum image count of 4 when creating the swapchain.

If semaphore is not VK_NULL_HANDLE, the semaphore must be unsignaled and not have any uncompleted signal or wait operations pending. It will become signaled when the application can use the image. Queue operations that access the image contents must wait until the semaphore signals; typically applications should include the semaphore in the <code>pWaitSemaphores</code> list for the queue submission that transitions the image away from the VK_IMAGE_LAYOUT_PRESENT_SRC_KHR layout. Use of the semaphore allows rendering operations to be recorded and submitted before the presentation engine has completed its use of the image.

If fence is not equal to **VK_NULL_HANDLE**, the fence must be unsignaled and not have any uncompleted signal operations pending. It will become signaled when the application can use the image. Applications can use this to meter their frame generation work to match the presentation rate.

semaphore and fence must not both be equal to **VK_NULL_HANDLE**. An application must wait until either the semaphore or fence is signaled before using the presentable image.

semaphore and fence may already be signaled when **vkAcquireNextImageKHR** returns, if the image is being acquired for the first time, or if the presentable image is immediately ready for use.

A successful call to **vkAcquireNextImageKHR** counts as a signal operation on *semaphore* for the purposes of queue forward-progress requirements. The semaphore is guaranteed to signal, so a wait operation can be queued for the semaphore without risk of deadlock.

The vkCmdWaitEvents or vkCmdPipelineBarrier used to transition the image away from VK_IMAGE_ LAYOUT_PRESENT_SRC_KHR layout must have dstStageMask and dstAccessMask parameters set based on the next use of the image. The srcAccessMask must include VK_ACCESS_MEMORY_READ_BIT to ensure that all prior reads by the presentation engine are complete before the image layout transition occurs. The application must use implicit ordering guarantees and execution dependencies to prevent the image transition from occurring before the semaphore passed to vkAcquireNextImageKHR has signaled.

Note

When the swapchain image will be written by some stage S, the recommended idiom for ensuring the semaphore signals before the transition occurs is:



- The batch that contains the transition includes the image-acquire semaphore in the list of semaphores to wait for, with a wait stage mask that includes *S*.
- The pipeline barrier that performs the transition includes S in both the srcStageMask and dstStageMask.

This causes the pipeline barrier to wait at S until the semaphore signals before performing the transition and memory barrier, while allowing earlier pipeline stages of subsequent commands to proceed.

After a successful return, the image indicated by pImageIndex will still be in the VK_IMAGE_LAYOUT_PRESENT_ SRC_KHR layout if it was previously presented, or in the VK_IMAGE_LAYOUT_UNDEFINED layout if this is the first time it has been acquired.

The possible return values for **vkAcquireNextImageKHR**() depend on the *timeout* provided:

- VK SUCCESS is returned if an image became available.
- VK_ERROR_SURFACE_LOST_KHR if the surface becomes no longer available.
- VK NOT READY is returned if timeout is zero and no image was available.
- VK_TIMEOUT is returned if timeout is greater than zero and less than UINT64_MAX, and no image became available within the time allowed.
- VK_SUBOPTIMAL_KHR is returned if an image became available, and the swapchain no longer matches the surface properties exactly, but can still be used to present to the surface successfully.



Note

This may happen, for example, if the platform surface has been resized but the platform is able to scale the presented images to the new size to produce valid surface updates. It is up to applications to decide whether they prefer to continue using the current swapchain indefinitely or temporarily in this state, or to re-create the swapchain to better match the platform surface properties.

• VK_ERROR_OUT_OF_DATE_KHR is returned if the surface has changed in such a way that it is no longer compatible with the swapchain, and further presentation requests using the swapchain will fail. Applications must query the new surface properties and recreate their swapchain if they wish to continue presenting to the surface.

If the native surface and presented image sizes no longer match, presentation may fail. If presentation does succeed, parts of the native surface may be undefined, parts of the presented image may have been clipped before presentation, and/or the image may have been scaled (uniformly or not uniformly) before presentation. It is the application's responsibility to detect surface size changes and react appropriately. If presentation fails because of a mismatch in the surface and presented image sizes, a VK_ERROR_OUT_OF_DATE_KHR error will be returned.

Before an application can present an image, the image's layout must be transitioned to the VK_IMAGE_LAYOUT_PRESENT_SRC_KHR layout. The **vkCmdWaitEvents** or **vkCmdPipelineBarrier** that perform the transition must have <code>srcStageMask</code> and <code>srcAccessMask</code> parameters set based on the preceding use of the image. The <code>dstAccessMask</code> must include VK_ACCESS_MEMORY_READ_BIT indicating all prior accesses indicated in <code>srcAccessMask</code> from stages in <code>srcStageMask</code> are to be made available to reads by the presentation engine. Any value of <code>dstStageMask</code> is valid, but should be set to VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT to avoid delaying subsequent commands that don't access the image.

After queueing all rendering commands and transitioning the image to the correct layout, to queue an image for presentation, call:

- queue is a queue that is capable of presentation to the target surface's platform on the same device as the image's swapchain.
- pPresentInfo is a pointer to an instance of the VkPresentInfoKHR structure specifying the parameters of the presentation.

Valid Usage

- queue must be a valid VkQueue handle
- pPresentInfo must be a pointer to a valid VkPresentInfoKHR structure
- Any given element of pSwapchains member of pPresentInfo must be a swapchain that is created for a surface for which presentation is supported from queue as determined using a call to

vkGetPhysicalDeviceSurfaceSupportKHR

• If more than one member of *pSwapchains* was created from a display surface, all display surfaces referenced that refer to the same display must use the same display mode

Host Synchronization

- Host access to queue must be externally synchronized
- Host access to pPresentInfo.pWaitSemaphores[] must be externally synchronized
- Host access to pPresentInfo.pSwapchains[] must be externally synchronized

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
-	-	Any

Return Codes

Success

- VK SUCCESS
- VK_SUBOPTIMAL_KHR

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_DEVICE_LOST
- VK_ERROR_OUT_OF_DATE_KHR
- VK_ERROR_SURFACE_LOST_KHR

The VkPresentInfoKHR structure is defined as:

- sType is the type of this structure and must be VK STRUCTURE TYPE PRESENT INFO KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- waitSemaphoreCount is the number of semaphores to wait for before issuing the present request. The number may be zero.
- pWaitSemaphores, if not VK_NULL_HANDLE, is an array of VkSemaphore objects with waitSemaphoreCount entries, and specifies the semaphores to wait for before issuing the present request.
- swapchainCount is the number of swapchains being presented to by this command.
- pSwapchains is an array of VkSwapchainKHR objects with swapchainCount entries. A given swapchain must not appear in this list more than once.
- pImageIndices is an array of indices into the array of each swapchain's presentable images, with swapchainCount entries. Each entry in this array identifies the image to present on the corresponding entry in the pSwapchains array.
- pResults is an array of VkResult typed elements with swapchainCount entries. Applications that don't need per-swapchain results can use NULL for pResults. If non-NULL, each entry in pResults will be set to the VkResult for presenting the swapchain corresponding to the same index in pSwapchains.

Valid Usage

• sType must be VK_STRUCTURE_TYPE_PRESENT_INFO_KHR

- pNext must be NULL
- If waitSemaphoreCount is not 0, and pWaitSemaphores is not NULL, pWaitSemaphores must be a pointer to an array of waitSemaphoreCount valid VkSemaphore handles
- pSwapchains must be a pointer to an array of swapchainCount valid VkSwapchainKHR handles
- pImageIndices must be a pointer to an array of swapchainCount uint32_t values
- If pResults is not NULL, pResults must be a pointer to an array of swapchainCount VkResult values
- swapchainCount must be greater than 0
- Any given element of pImageIndices must be the index of a presentable image acquired from the swapchain specified by the corresponding element of the pSwapchains array
- Any given element of VkSemaphore in pWaitSemaphores must refer to a prior signal of that VkSemaphore that will not be consumed by any other wait on that semaphore

When the VK_KHR_display_swapchain extension is enabled additional fields can be specified when presenting an image to a swapchain by setting VkPresentInfoKHR::pNext to point to an instance of the VkDisplayPresentInfoKHR structure.

The VkDisplayPresentInfoKHR structure is defined as:

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_DISPLAY_PRESENT_INFO_KHR.
- pNext is NULL or a pointer to an extension-specific structure.
- srcRect is a rectangular region of pixels to present. It must be a subset of the image being presented. If VkDisplayPresentInfoKHR is not specified, this region will be assumed to be the entire presentable image.
- dstRect is a rectangular region within the visible region of the swapchain's display mode. If VkDisplayPresentInfoKHR is not specified, this region will be assumed to be the entire visible region of the visible region of the swapchain's mode. If the specified rectangle is a subset of the display mode's visible region, content from display planes below the swapchain's plane will be visible outside the rectangle. If there are no planes below the swapchain's, the area outside the specified rectangle will be black. If portions of the specified rectangle are outside of the display's visible region, pixels mapping only to those portions of the rectangle will be discarded.
- persistent: If this is VK_TRUE, the display engine will enable buffered mode on displays that support it. This allows the display engine to stop sending content to the display until a new image is presented. The display will instead maintain a copy of the last presented image. This allows less power to be used, but may increase presentation latency. If VkDisplayPresentInfoKHR is not specified, persistent mode will not be used.

If the extent of the srcRect and dstRect are not equal, the presented pixels will be scaled accordingly.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DISPLAY_PRESENT_INFO_KHR
- pNext must be NULL
- srcRect must specify a rectangular region that is a subset of the image being presented
- dstRect must specify a rectangular region that is a subset of the visibleRegion parameter of the display mode the swapchain being presented uses
- If the persistentContent member of the VkDisplayPropertiesKHR structure returned by **vkGetPhysicalDeviceDisplayPropertiesKHR** for the display the present operation targets then persistent must be VK_FALSE

vkQueuePresentKHR, releases the acquisition of the images referenced by *imageIndices*. A presented images must not be used again before it has been reacquired using **vkAcquireNextImageKHR**.

The processing of the presentation happens in issue order with other queue operations, but semaphores have to be used to ensure that prior rendering and other commands in the specified queue complete before the presentation begins. The presentation command itself does not delay processing of subsequent commands on the queue, however, presentation requests sent to a particular queue are always performed in order. Exact presentation timing is controlled by the semantics of the presentation engine and native platform in use.

If an image is presented to a swapchain created from a display surface, the mode of the associated display will be updated, if necessary, to match the mode specified when creating the display surface. The mode switch and presentation of the specified image will be performed as one atomic operation.

The result codes VK_ERROR_OUT_OF_DATE_KHR and VK_SUBOPTIMAL_KHR have the same meaning when returned by **vkQueuePresentKHR** as they do when returned by **vkAcquireNextImageKHR**(). If multiple swapchains are presented, the result code is determined applying the following rules in order:

- If the device is lost, VK_ERROR_DEVICE_LOST is returned.
- If any of the target surfaces are no longer available the error VK_ERROR_SURFACE_LOST_KHR is returned.
- If any of the presents would have a result of VK_ERROR_OUT_OF_DATE_KHR if issued separately then VK_ERROR_OUT_OF_DATE_KHR is returned.
- If any of the presents would have a result of VK_SUBOPTIMAL_KHR if issued separately then VK_SUBOPTIMAL_KHR is returned.
- Otherwise VK_SUCCESS is returned.

Presentation is a read-only operation that will not affect the content of the presentable images. Upon reacquiring the image and transitioning it away from the VK_IMAGE_LAYOUT_PRESENT_SRC_KHR layout, the contents will be the same as they were prior to transitioning the image to the present source layout and presenting it. However, if a mechanism other than Vulkan is used to modify the platform window associated with the swapchain, the content of all presentable images in the swapchain becomes undefined.

Chapter 30

Extended Functionality

Additional functionality may be provided by layers or extensions. A layer cannot add or modify Vulkan commands, while an extension may do so.

The set of layers to enable is specified when creating an instance, and those layers are able to intercept any Vulkan command dispatched to that instance or any of its child objects.

Extensions can operate at either the instance or device scope. Enabled instance extensions are able to affect the operation of the instance and any of its child objects, while device extensions may only be available on a subset of physical devices, must be individually enabled per-device, and only affect the operation of the devices where they are enabled.

Examples of these might be:

- Whole API validation is an example of a layer.
- Debug capabilities might make a good instance extension.
- A layer that provides hardware-specific performance telemetry and analysis could be a layer that is only active for devices created from compatible physical devices.
- Functions to allow an application to use additional hardware features beyond the core would be a good candidate for a device extension.

30.1 Layers

When a layer is enabled, it inserts itself into the call chain for Vulkan commands the layer is interested in. A common use of layers is to validate application behavior during development. For example, the implementation will not check that Vulkan enums used by the application fall within allowed ranges. Instead, a validation layer would do those checks and flag issues. This avoids a performance penalty during production use of the application because those layers would not be enabled in production.

Vulkan layers may wrap object handles (i.e. return a different handle value to the application than that generated by the implementation). This is generally discouraged, as it increases the probability of incompatibilities with new extensions. The validation layers wrap handles in order to track the proper use and destruction of each object. See the "Vulkan Loader Specification and Architecture Overview" document for additional information.

To query the available layers, call:

- pPropertyCount is a pointer to an integer related to the number of layer properties available or queried, as described below.
- pProperties is either NULL or a pointer to an array of VkLayerProperties structures.

If pProperties is NULL, then the number of layer properties available is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If pPropertyCount is less than the number of layer properties available, at most pPropertyCount structures will be written. If pPropertyCount is smaller than the number of layers available, VK_INCOMPLETE will be returned instead of VK_SUCCESS, to indicate that not all the available layer properties were returned.

The list of available layers may change at any time due to actions outside of the Vulkan implementation, so two calls to **vkEnumerateInstanceLayerProperties** with the same parameters may return different results, or retrieve different pPropertyCount values or pProperties contents. Once an instance has been created, the layers enabled for that instance will continue to be enabled and valid for the lifetime of that instance, even if some of them become unavailable for future instances.

Valid Usage

- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkLayerProperties structures

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK ERROR OUT OF HOST MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkLayerProperties structure is defined as:

- layerName is a null-terminated UTF-8 string specifying the name of the layer. Use this name in the ppEnabledLayerNames array passed in the VkInstanceCreateInfo structure to enable this layer for an instance.
- specVersion is the Vulkan version the layer was written to, encoded as described in the API Version Numbers and Semantics section.
- implementationVersion is the version of this layer. It is an integer, increasing with backward compatible changes.
- description is a null-terminated UTF-8 string providing additional details that can be used by the application to identify the layer.

To enable a layer, the name of the layer should be added to the ppEnabledLayerNames member of VkInstanceCreateInfo when creating a VkInstance.

Loader implementations may provide mechanisms outside the Vulkan API for enabling specific layers. Layers enabled through such a mechanism are *implicitly enabled*, while layers enabled by including the layer name in the <code>ppEnabledLayerNames</code> member of <code>VkInstanceCreateInfo</code> are *explicitly enabled*. Except where otherwise specified, implicitly enabled and explicitly enabled layers differ only in the way they are enabled. Explicitly enabling a layer that is implicitly enabled has no additional effect.

30.1.1 Device Layer Deprecation

Previous versions of this specification distinguished between instance and device layers. Instance layers were only able to intercept commands that operate on VkInstance and VkPhysicalDevice, except they were not able to intercept vkCreateDevice. Device layers were enabled for individual devices when they were created, and could only intercept commands operating on that device or its child objects.

Device-only layers are now deprecated, and this specification no longer distinguishes between instance and device layers. Layers are enabled during instance creation, and are able to intercept all commands operating on that instance or any of its child objects. At the time of deprecation there were no known device-only layers and no compelling reason to create one.

In order to maintain compatibility with implementations released prior to device-layer deprecation, applications should still enumerate and enable device layers. The behavior of **vkEnumerateDeviceLayerProperties** and valid usage of the <code>ppEnabledLayerNames</code> member of VkDeviceCreateInfo maximizes compatibility with applications written to work with the previous requirements.

To enumerate device layers, call:

pPropertyCount is a pointer to an integer related to the number of layer properties available or queried.

• pProperties is either NULL or a pointer to an array of VkLayerProperties structures.

If pProperties is NULL, then the number of layer properties available is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If pPropertyCount is less than the number of layer properties available, at most pPropertyCount structures will be written. If pPropertyCount is smaller than the number of layers available, VK_INCOMPLETE will be returned instead of VK_SUCCESS, to indicate that not all the available layer properties were returned.

The list of layers enumerated by **vkEnumerateDeviceLayerProperties** must be exactly the sequence of layers enabled for the instance. The members of VkLayerProperties for each enumerated layer must be the same as the properties when the layer was enumerated by **vkEnumerateInstanceLayerProperties**.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkLayerProperties structures

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The ppEnabledLayerNames and enabledLayerCount members of VkDeviceCreateInfo are deprecated and their values must be ignored by implementations. However, for compatibility, only an empty list of layers or a list that exactly matches the sequence enabled at instance creation time are valid, and validation layers should issue diagnostics for other cases.

Regardless of the enabled layer list provided in $\protect{VkDeviceCreateInfo}$, the sequence of layers active for a device will be exactly the sequence of layers enabled when the parent instance was created.

30.2 Extensions

Extensions may define new Vulkan commands, structures, and enumerants. For compilation purposes, the interfaces defined by registered extensions, including new structures and enumerants as well as function pointer types for new commands, are defined in the Khronos-supplied vulkan.h together with the core API. However, commands defined by extensions may not be available for static linking - in which case function pointers to these commands should be queried at runtime as described in Section 3.1. Extensions may be provided by layers as well as by a Vulkan implementation.

Because extensions may extend or change the behavior of the Vulkan API, extension authors should add support for their extensions to the Khronos validation layers. This is especially important for new commands whose parameters have been wrapped by the validation layers. See the "Vulkan Loader Specification and Architecture Overview" document for additional information.

To query the available instance extensions, call:

- pLayerName is either NULL or a pointer to a null-terminated UTF-8 string naming the layer to retrieve extensions from
- pPropertyCount is a pointer to an integer related to the number of extension properties available or queried, as described below.
- pProperties is either NULL or a pointer to an array of VkExtensionProperties structures.

When pLayerName parameter is NULL, only extensions provided by the Vulkan implementation or by implicitly enabled layers are returned. When pLayerName is the name of a layer, the instance extensions provided by that layer are returned.

If pProperties is NULL, then the number of extensions properties available is returned in pPropertyCount. Otherwise, pPropertyCount must point to a variable set by the user to the number of elements in the pProperties array, and on return the variable is overwritten with the number of structures actually written to pProperties. If pPropertyCount is less than the number of extension properties available, at most pPropertyCount structures will be written. If pPropertyCount is smaller than the number of extensions available, VK_INCOMPLETE will be returned instead of VK_SUCCESS, to indicate that not all the available properties were returned.

Because the list of available layers may change externally between calls to

vkEnumerateInstanceExtensionProperties, two calls may retrieve different results if a pLayerName is available in one call but not in another. The extensions supported by a layer may also change between two calls, e.g. if the layer implementation is replaced by a different version between those calls.

Valid Usage

- If pLayerName is not NULL, pLayerName must be a null-terminated string
- pPropertyCount must be a pointer to a uint32_t value

- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkExtensionProperties structures
- If pLayerName is not NULL, it must be the name of a layer returned by vkEnumerateInstanceLayerProperties

Return Codes

Success

- VK SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_LAYER_NOT_PRESENT

To enable an instance extension, the name of the extension should be added to the <code>ppEnabledExtensionNames</code> member of <code>VkInstanceCreateInfo</code> when creating a <code>VkInstance</code>.

Enabling an extension does not change behavior of functionality exposed by the core Vulkan API or any other extension, other than making valid the use of the commands, enums and structures defined by that extension.

To query the extensions available to a given physical device, call:

- physicalDevice is the physical device that will be queried.
- pLayerName is either NULL or a pointer to a null-terminated UTF-8 string naming the layer to retrieve extensions from
- pPropertyCount is a pointer to an integer related to the number of extension properties available or queried, and is treated in the same fashion as the vkEnumerateInstanceExtensionProperties::pPropertyCount parameter.
- pProperties is either NULL or a pointer to an array of VkExtensionProperties structures.

When playerName parameter is NULL, only extensions provided by the Vulkan implementation or by implicitly enabled layers are returned. When playerName is the name of a layer, the device extensions provided by that layer are returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- If pLayerName is not NULL, pLayerName must be a null-terminated string
- pPropertyCount must be a pointer to a uint32_t value
- If the value referenced by pPropertyCount is not 0, and pProperties is not NULL, pProperties must be a pointer to an array of pPropertyCount VkExtensionProperties structures
- If pLayerName is not NULL, it must be the name of a layer returned by vkEnumerateDeviceLayerProperties

Return Codes

Success

- VK_SUCCESS
- VK_INCOMPLETE

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_LAYER_NOT_PRESENT

The VkExtensionProperties structure is defined as:

- extensionName is a null-terminated string specifying the name of the extension.
- specVersion is the version of this extension. It is an integer, incremented with backward compatible changes.

30.2.1 Instance Extensions and Device Extensions

Because an instance extension can affect the operation of an instance and any of its child objects, the decision to expose functionality as an instance extension or as a device extension is not always clear. This section provides some guidelines and rules for when to expose new functionality as an instance extension, device extension, or both.

The decision is influenced by whether extension functionality affects instance-level objects (e.g. instances and physical devices) and commands, or device-level objects (e.g. logical devices, queues, and command buffers) and commands, or both.

In some cases, the decision is clear:

- Functionality that is restricted to the instance-level must be implemented as an instance extension.
- Functionality that is restricted to the device-level must be implemented as a device extension.

In other cases, the decision is not so clear:

- Global functionality that affects the entire Vulkan API, including instance and device-level objects and commands, should be an instance extension.
- Device-level functionality that contains physical-device queries, can be implemented as an instance extension. If some part of an instance extension's functionality might not be available on all physical devices, the extension should provide a query to determine which physical devices provide the functionality.
- For a set of global functionality that provides new instance-level and device-level commands, and can be enabled for a subset of devices, it is recommended that the functionality be partitioned across two extensions—one for the instance-level functionality, and one for the device-specific functionality. In this latter case, it is generally recommended that the two extensions have unique names.

Examples of instance extensions include:

- Logging of debug messages by any enabled layers for all Vulkan commands.
- Functionality creating new objects which are direct children of an instance.
- Functionality creating new objects which are direct children of a physical device and intended to work with any logical device created from the physical device.
- Functionality adding new instance-level Vulkan commands that do not affect any device-level commands.

Note



Instance extensions generally require support in the Vulkan loader. This is especially true for commands that are dispatched from instances and physical devices. Additional information about supporting instance-level commands may be found in the extensions "Vulkan Loader Specification and Architecture Overview" document, Please see the "Architectural overview of layers and loader" section for information about how both instance-level and device-level commands are supported and dispatched.

Chapter 31

Features, Limits, and Formats

Vulkan is designed to support a wide range of hardware and as such there are a number of features, limits, and formats which are not supported on all hardware. Features describe functionality that is not required and which must be explicitly enabled. Limits describe implementation-dependent minimums, maximums, and other device characteristics that an application may need to be aware of. Supported buffer and image formats may vary across implementations. A minimum set of format features are guaranteed, but others must be explicitly queried before use to ensure they are supported by the implementation.



Note on extensibility

The features and limits are reported via basic structures (that is VkPhysicalDeviceFeatures and VkPhysicalDeviceFeatures). It is expected that when new features or limits are added in a future Vulkan version, new structure(s) and entry point(s) will be added as necessary to query these. New functionality added by extensions is not expected to modify the core feature and limit structures.

31.1 Features

The Specification defines a set of fine-grained features that are not required, but may be supported by a Vulkan implementation. Support for features is reported and enabled on a per-feature basis. Features are properties of the physical device.

To query supported features, call:

- physicalDevice is the physical device from which to query the supported features.
- pFeatures is a pointer to a VkPhysicalDeviceFeatures structure in which the physical device features are returned. For each feature, a value of VK_TRUE indicates that the feature is supported on this physical device, and VK_FALSE indicates that the feature is not supported.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- pFeatures must be a pointer to a VkPhysicalDeviceFeatures structure

Fine-grained features used by a logical device must be enabled at VkDevice creation time. If a feature is enabled that the physical device does not support, VkDevice creation will fail. If an application uses a feature without enabling it at VkDevice creation time, the device behavior is undefined. The validation layer will warn if features are used without being enabled.

The fine-grained features are enabled by passing a pointer to the VkPhysicalDeviceFeatures structure via the pEnabledFeatures member of the VkDeviceCreateInfo structure that is passed into the **vkCreateDevice** call. If a member of pEnabledFeatures is set to VK_TRUE or VK_FALSE, then the device will be created with the indicated feature enabled or disabled, respectively.

If an application wishes to enable all features supported by a device, it can simply pass in the VkPhysicalDeviceFeatures structure that was previously returned by **vkGetPhysicalDeviceFeatures**. To disable an individual feature, the application can set the desired member to VK_FALSE in the same structure. Setting <code>pEnabledFeatures</code> to NULL is equivalent to setting all members of the structure to VK_FALSE.



Note

Some features, such as <code>robustBufferAccess</code>, may incur a run-time performance cost. Application writers should carefully consider the implications of enabling all supported features.

The VkPhysicalDeviceFeatures structure is defined as:

```
typedef struct VkPhysicalDeviceFeatures {
   VkBool32 robustBufferAccess;
   VkBool32 fullDrawIndexUint32;
   VkBool32 imageCubeArray;
   VkBool32 independentBlend;
   VkBool32 geometryShader;
   VkBool32 tessellationShader;
   VkBool32 sampleRateShading;
   VkBool32 dualSrcBlend;
   VkBool32 logicOp;
   VkBool32 multiDrawIndirect;
   VkBool32
             drawIndirectFirstInstance;
   VkBool32 depthClamp;
   VkBool32 depthBiasClamp;
   VkBool32 fillModeNonSolid;
   VkBool32 depthBounds;
   VkBool32 wideLines;
   VkBool32 largePoints;
   VkBool32 alphaToOne;
   VkBool32 multiViewport;
   VkBool32 samplerAnisotropy;
   VkBool32 textureCompressionETC2;
   VkBool32
              textureCompressionASTC_LDR;
   VkBool32 textureCompressionBC;
```

```
VkBool32 occlusionQueryPrecise;
   VkBool32 pipelineStatisticsQuery;
   VkBool32 vertexPipelineStoresAndAtomics;
   VkBool32 fragmentStoresAndAtomics;
   VkBool32 shaderTessellationAndGeometryPointSize;
   VkBool32 shaderImageGatherExtended;
   VkBool32
              shaderStorageImageExtendedFormats;
   VkBool32
              shaderStorageImageMultisample;
            shaderStorageImageReadWithoutFormat;
   VkBool32
   VkBool32 shaderStorageImageWriteWithoutFormat;
   VkBool32 shaderUniformBufferArrayDynamicIndexing;
   VkBool32 shaderSampledImageArrayDynamicIndexing;
   VkBool32 shaderStorageBufferArrayDynamicIndexing;
   VkBool32 shaderStorageImageArrayDynamicIndexing;
   VkBool32 shaderClipDistance;
   VkBool32 shaderCullDistance;
   VkBool32 shaderFloat64;
   VkBool32 shaderInt64;
   VkBool32 shaderInt16;
   VkBool32 shaderResourceResidency;
   VkBool32
              shaderResourceMinLod;
            sparseBinding;
   VkBool32
   VkBool32 sparseResidencyBuffer;
   VkBool32 sparseResidencyImage2D;
   VkBool32 sparseResidencyImage3D;
   VkBool32 sparseResidency2Samples;
   VkBool32 sparseResidency4Samples;
   VkBool32 sparseResidency8Samples;
   VkBool32 sparseResidency16Samples;
   VkBool32 sparseResidencyAliased;
   VkBoo132
              variableMultisampleRate;
   VkBool32
              inheritedQueries;
} VkPhysicalDeviceFeatures;
```

The members of the VkPhysicalDeviceFeatures structure describe the following features:

- robustBufferAccess indicates that accesses to buffers are bounds-checked against the range of the buffer descriptor (as determined by VkDescriptorBufferInfo::range, VkBufferViewCreateInfo::range, or the size of the buffer). Out of bounds accesses must not cause application termination, and the effects of shader loads, stores, and atomics must conform to an implementation-dependent behavior as described below.
 - A buffer access is considered to be out of bounds if any of the following are true:
 - * The pointer was formed by **OpImageTexelPointer** and the coordinate is less than zero or greater than or equal to the number of whole elements in the bound range.
 - * The pointer was not formed by **OpImageTexelPointer** and the object pointed to is not wholly contained within the bound range.



Note

If a SPIR-V **OpLoad** instruction loads a structure and the tail end of the structure is out of bounds, then all members of the structure are considered out of bounds even if the members at the end are not statically used.

* If any buffer access in a given SPIR-V block is determined to be out of bounds, then any other access of the same type (load, store, or atomic) in the same SPIR-V block that accesses an address less than 16 bytes away from the out of bounds address may also be considered out of bounds.

- Out-of-bounds buffer loads will return any of the following values:
 - * Values from anywhere within the memory range(s) bound to the buffer (possibly including bytes of memory past the end of the buffer, up to the end of the bound range).
 - * Zero values, or (0,0,0,x) vectors for vector reads where x is a valid value represented in the type of the vector components and may be any of:
 - · 0, 1, or the maximum representable positive integer value, for signed or unsigned integer components
 - · 0.0 or 1.0, for floating-point components
- Out-of-bounds writes may modify values within the memory range(s) bound to the buffer, but must not modify any other memory.
- Out-of-bounds atomics may modify values within the memory range(s) bound to the buffer, but must not modify any other memory, and return an undefined value.
- Vertex input attributes are considered out of bounds if the address of the attribute plus the size of the attribute is
 greater than the size of the bound buffer. Further, if any vertex input attribute using a specific vertex input binding is
 out of bounds, then all vertex input attributes using that vertex input binding for that vertex shader invocation are
 considered out of bounds.
 - * If a vertex input attribute is out of bounds, it will be assigned one of the following values:
 - · Values from anywhere within the memory range(s) bound to the buffer, converted according to the format of the attribute.
 - · Zero values, format converted according to the format of the attribute.
 - · Zero values, or (0,0,0,x) vectors, as described above.
- If *robustBufferAccess* is not enabled, out of bounds accesses may corrupt any memory within the process and cause undefined behavior up to and including application termination.
- fullDrawIndexUint32 indicates the full 32-bit range of indices is supported for indexed draw calls when using a VkIndexType of VK_INDEX_TYPE_UINT32. maxDrawIndexedIndexValue is the maximum index value that may be used (aside from the primitive restart index, which is always 2³²-1 when the VkIndexType is VK_INDEX_TYPE_UINT32). If this feature is supported, maxDrawIndexedIndexValue must be 2³²-1; otherwise it must be no smaller than 2²⁴-1. See maxDrawIndexedIndexValue.
- imageCubeArray indicates whether image views with a VkImageViewType of VK_IMAGE_VIEW_TYPE_ CUBE_ARRAY can be created, and that the corresponding **SampledCubeArray** and **ImageCubeArray** SPIR-V capabilities can be used in shader code.
- independentBlend indicates whether the VkPipelineColorBlendAttachmentState settings are controlled independently per-attachment. If this feature is not enabled, the VkPipelineColorBlendAttachmentState settings for all color attachments must be identical. Otherwise, a different VkPipelineColorBlendAttachmentState can be provided for each bound color attachment.
- geometryShader indicates whether geometry shaders are supported. If this feature is not enabled, the VK_SHADER_ STAGE_GEOMETRY_BIT and VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT enum values must not be used. This also indicates whether shader modules can declare the **Geometry** capability.
- tessellationShader indicates whether tessellation control and evaluation shaders are supported. If this feature is not enabled, the VK_SHADER_STAGE_TESSELLATION_CONTROL_BIT, VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT, VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT, VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT, and VK_STRUCTURE_TYPE_PIPELINE_TESSELLATION_STATE_CREATE_INFO enum values must not be used. This also indicates whether shader modules can declare the Tessellation capability.

- sampleRateShading indicates whether per-sample shading and multisample interpolation are supported. If this feature is not enabled, the sampleShadingEnable member of the VkPipelineMultisampleStateCreateInfo structure must be set to VK_FALSE and the minSampleShading member is ignored. This also indicates whether shader modules can declare the SampleRateShading capability.
- dualSrcBlend indicates whether blend operations which take two sources are supported. If this feature is not enabled, the VK_BLEND_FACTOR_SRC1_COLOR, VK_BLEND_FACTOR_ONE_MINUS_SRC1_COLOR, VK_BLEND_FACTOR_SRC1_ALPHA, and VK_BLEND_FACTOR_ONE_MINUS_SRC1_ALPHA enum values must not be used as source or destination blending factors. See Section 26.1.2.
- logicOp indicates whether logic operations are supported. If this feature is not enabled, the logicOpEnable member of the VkPipelineColorBlendStateCreateInfo structure must be set to VK_FALSE, and the logicOp member is ignored.
- multiDrawIndirect indicates whether multiple draw indirect is supported. If this feature is not enabled, the drawCount parameter to the vkCmdDrawIndirect and vkCmdDrawIndexedIndirect commands must be 0 or 1. The maxDrawIndirectCount member of the VkPhysicalDeviceLimits structure must also be 1 if this feature is not supported. See maxDrawIndirectCount.
- drawIndirectFirstInstance indicates whether indirect draw calls support the firstInstance parameter. If this feature is not enabled, the firstInstance member of all VkDrawIndirectCommand and VkDrawIndexedIndirectCommand structures that are provided to the vkCmdDrawIndirect and vkCmdDrawIndexedIndirect commands must be 0.
- depthClamp indicates whether depth clamping is supported. If this feature is not enabled, the depthClampEnable member of the VkPipelineRasterizationStateCreateInfo structure must be set to VK_FALSE. Otherwise, setting depthClampEnable to VK_TRUE will enable depth clamping.
- depthBiasClamp indicates whether depth bias clamping is supported. If this feature is not enabled, the depthBiasClamp member of the VkPipelineRasterizationStateCreateInfo structure must be set to 0.0 unless the VK_DYNAMIC_STATE_DEPTH_BIAS dynamic state is enabled, and the depthBiasClamp parameter to vkCmdSetDepthBias must be set to 0.0.
- fillModeNonSolid indicates whether point and wireframe fill modes are supported. If this feature is not enabled, the VK_POLYGON_MODE_POINT and VK_POLYGON_MODE_LINE enum values must not be used.
- depthBounds indicates whether depth bounds tests are supported. If this feature is not enabled, the depthBoundsTestEnable member of the VkPipelineDepthStencilStateCreateInfo structure must be set to VK_FALSE. When depthBoundsTestEnable is set to VK_FALSE, the minDepthBounds and maxDepthBounds members of the VkPipelineDepthStencilStateCreateInfo structure are ignored.
- wideLines indicates whether lines with width other than 1.0 are supported. If this feature is not enabled, the lineWidth member of the VkPipelineRasterizationStateCreateInfo structure must be set to 1.0 unless the VK_DYNAMIC_STATE_LINE_WIDTH dynamic state is enabled, and the lineWidth parameter to vkCmdSetLineWidth must be set to 1.0. When this feature is supported, the range and granularity of supported line widths are indicated by the lineWidthRange and lineWidthGranularity members of the VkPhysicalDeviceLimits structure, respectively.
- largePoints indicates whether points with size greater than 1.0 are supported. If this feature is not enabled, only a point size of 1.0 written by a shader is supported. The range and granularity of supported point sizes are indicated by the pointSizeRange and pointSizeGranularity members of the VkPhysicalDeviceLimits structure, respectively.
- alphaToOne indicates whether the implementation is able to replace the alpha value of the color fragment output from the fragment shader with the maximum representable alpha value for fixed-point colors or 1.0 for floating-point colors.

If this feature is not enabled, then the <code>alphaToOneEnable</code> member of the VkPipelineMultisampleStateCreateInfo structure must be set to VK_FALSE. Otherwise setting <code>alphaToOneEnable</code> to VK_TRUE will enable alpha-to-one behavior.

- multiViewport indicates whether more than one viewport is supported. If this feature is not enabled, the viewportCount and scissorCount members of the VkPipelineViewportStateCreateInfo structure must be set to 1. Similarly, the viewportCount parameter to the vkCmdSetViewport command and the scissorCount parameter to the vkCmdSetScissor command must be 1, and the firstViewport parameter to the vkCmdSetViewport command and the firstScissor parameter to the vkCmdSetScissor command must be 0.
- samplerAnisotropy indicates whether anisotropic filtering is supported. If this feature is not enabled, the maxAnisotropy member of the VkSamplerCreateInfo structure must be 1.0.
- textureCompressionETC2 indicates whether the ETC2 and EAC compressed texture formats are supported. If this feature is not enabled, the following formats must not be used to create images:
 - VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK
 - VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK
 - VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK
 - VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK
 - VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK
 - VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK
 - VK FORMAT EAC R11 UNORM BLOCK
 - VK FORMAT EAC R11 SNORM BLOCK
 - VK_FORMAT_EAC_R11G11_UNORM_BLOCK
 - VK_FORMAT_EAC_R11G11_SNORM_BLOCK

vkGetPhysicalDeviceFormatProperties is used to check for the supported properties of individual formats.

- textureCompressionASTC_LDR indicates whether the ASTC LDR compressed texture formats are supported. If this feature is not enabled, the following formats must not be used to create images:
 - VK_FORMAT_ASTC_4x4_UNORM_BLOCK
 - VK_FORMAT_ASTC_4x4_SRGB_BLOCK
 - VK_FORMAT_ASTC_5x4_UNORM_BLOCK
 - VK_FORMAT_ASTC_5x4_SRGB_BLOCK
 - VK_FORMAT_ASTC_5x5_UNORM_BLOCK
 - VK_FORMAT_ASTC_5x5_SRGB_BLOCK
 - VK_FORMAT_ASTC_6x5_UNORM_BLOCK
 - VK_FORMAT_ASTC_6x5_SRGB_BLOCK
 - VK_FORMAT_ASTC_6x6_UNORM_BLOCK
 - VK_FORMAT_ASTC_6x6_SRGB_BLOCK
 - VK_FORMAT_ASTC_8x5_UNORM_BLOCK
 - VK_FORMAT_ASTC_8x5_SRGB_BLOCK
 - VK_FORMAT_ASTC_8x6_UNORM_BLOCK

- VK_FORMAT_ASTC_8x6_SRGB_BLOCK
- VK_FORMAT_ASTC_8x8_UNORM_BLOCK
- VK_FORMAT_ASTC_8x8_SRGB_BLOCK
- VK_FORMAT_ASTC_10x5_UNORM_BLOCK
- VK_FORMAT_ASTC_10x5_SRGB_BLOCK
- VK_FORMAT_ASTC_10x6_UNORM_BLOCK
- VK_FORMAT_ASTC_10x6_SRGB_BLOCK
- VK_FORMAT_ASTC_10x8_UNORM_BLOCK
- VK_FORMAT_ASTC_10x8_SRGB_BLOCK
- VK_FORMAT_ASTC_10x10_UNORM_BLOCK
- VK_FORMAT_ASTC_10x10_SRGB_BLOCK
- VK FORMAT ASTC 12x10 UNORM BLOCK
- VK_FORMAT_ASTC_12x10_SRGB_BLOCK
- VK_FORMAT_ASTC_12x12_UNORM_BLOCK
- VK_FORMAT_ASTC_12x12_SRGB_BLOCK

vkGetPhysicalDeviceFormatProperties is used to check for the supported properties of individual formats.

- textureCompressionBC indicates whether the BC compressed texture formats are supported. If this feature is not enabled, the following formats must not be used to create images:
 - VK_FORMAT_BC1_RGB_UNORM_BLOCK
 - VK_FORMAT_BC1_RGB_SRGB_BLOCK
 - VK_FORMAT_BC1_RGBA_UNORM_BLOCK
 - VK_FORMAT_BC1_RGBA_SRGB_BLOCK
 - VK_FORMAT_BC2_UNORM_BLOCK
 - VK_FORMAT_BC2_SRGB_BLOCK
 - VK FORMAT BC3 UNORM BLOCK
 - VK_FORMAT_BC3_SRGB_BLOCK
 - VK_FORMAT_BC4_UNORM_BLOCK
 - VK_FORMAT_BC4_SNORM_BLOCK
 - VK_FORMAT_BC5_UNORM_BLOCK
 - VK_FORMAT_BC5_SNORM_BLOCK
 - VK_FORMAT_BC6H_UFLOAT_BLOCK
 - VK_FORMAT_BC6H_SFLOAT_BLOCK
 - VK_FORMAT_BC7_UNORM_BLOCK
 - VK_FORMAT_BC7_SRGB_BLOCK

vkGetPhysicalDeviceFormatProperties is used to check for the supported properties of individual formats.

- occlusionQueryPrecise indicates whether occlusion queries returning actual sample counts are supported. Occlusion queries are created in a VkQueryPool by specifying the queryType of VK_QUERY_TYPE_
 OCCLUSION in the VkQueryPoolCreateInfo structure which is passed to vkCreateQueryPool. If this feature is enabled, queries of this type can enable VK_QUERY_CONTROL_PRECISE_BIT in the flags parameter to vkCmdBeginQuery. If this feature is not supported, the implementation supports only boolean occlusion queries. When any samples are passed, boolean queries will return a non-zero result value, otherwise a result value of zero is returned. When this feature is enabled and VK_QUERY_CONTROL_PRECISE_BIT is set, occlusion queries will report the actual number of samples passed.
- pipelineStatisticsQuery indicates whether the pipeline statistics queries are supported. If this feature is not enabled, queries of type VK_QUERY_TYPE_PIPELINE_STATISTICS cannot be created, and none of the VkQueryPipelineStatisticFlagBits bits can be set in the pipelineStatistics member of the VkQueryPoolCreateInfo structure.
- vertexPipelineStoresAndAtomics indicates whether storage buffers and images support stores and atomic operations in the vertex, tessellation, and geometry shader stages. If this feature is not enabled, all storage image, storage texel buffers, and storage buffer variables used by these stages in shader modules must be decorated with the NonWriteable decoration (or the readonly memory qualifier in GLSL).
- fragmentStoresAndAtomics indicates whether storage buffers and images support stores and atomic operations in the fragment shader stage. If this feature is not enabled, all storage image, storage texel buffers, and storage buffer variables used by the fragment stage in shader modules must be decorated with the NonWriteable decoration (or the readonly memory qualifier in GLSL).
- shaderTessellationAndGeometryPointSize indicates whether the **PointSize** built-in decoration is available in the tessellation control, tessellation evaluation, and geometry shader stages. If this feature is not enabled, members decorated with the **PointSize** built-in decoration must not be read from or written to and all points written from a tessellation or geometry shader will have a size of 1.0. This also indicates whether shader modules can declare the **TessellationPointSize** capability for tessellation control and evaluation shaders, or if the shader modules can declare the **GeometryPointSize** capability for geometry shaders. An implementation supporting this feature must also support one or both of the tessellationShader or geometryShader features.
- shaderImageGatherExtended indicates whether the extended set of image gather instructions are available in shader code. If this feature is not enabled, the OpImage*Gather instructions do not support the Offset and ConstOffsets operands. This also indicates whether shader modules can declare the ImageGatherExtended capability.
- shaderStorageImageExtendedFormats indicates whether the extended storage image formats are available in shader code. If this feature is not enabled, the formats requiring the **StorageImageExtendedFormats** capability are not supported for storage images. This also indicates whether shader modules can declare the **StorageImageExtendedFormats** capability.
- shaderStorageImageMultisample indicates whether multisampled storage images are supported. If this feature is not enabled, images that are created with a usage that includes VK_IMAGE_USAGE_STORAGE_BIT must be created with samples equal to VK_SAMPLE_COUNT_1_BIT. This also indicates whether shader modules can declare the StorageImageMultisample capability.
- shaderStorageImageReadWithoutFormat indicates whether storage images require a format qualifier to be specified when reading from storage images. If this feature is not enabled, the OpImageRead instruction must not have an OpTypeImage of Unknown. This also indicates whether shader modules can declare the StorageImageReadWithoutFormat capability.
- shaderStorageImageWriteWithoutFormat indicates whether storage images require a format qualifier to be specified when writing to storage images. If this feature is not enabled, the OpImageWrite instruction must not have an OpTypeImage of Unknown. This also indicates whether shader modules can declare the StorageImageWriteWithoutFormat capability.

- shaderUniformBufferArrayDynamicIndexing indicates whether arrays of uniform buffers can be indexed by dynamically uniform integer expressions in shader code. If this feature is not enabled, resources with a descriptor type of VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC must be indexed only by constant integral expressions when aggregated into arrays in shader code. This also indicates whether shader modules can declare the **UniformBufferArrayDynamicIndexing** capability.
- shaderSampledImageArrayDynamicIndexing indicates whether arrays of samplers or sampled images can be indexed by dynamically uniform integer expressions in shader code. If this feature is not enabled, resources with a descriptor type of VK_DESCRIPTOR_TYPE_SAMPLER, VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_
 SAMPLER, or VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE must be indexed only by constant integral expressions when aggregated into arrays in shader code. This also indicates whether shader modules can declare the SampledImageArrayDynamicIndexing capability.
- shaderStorageBufferArrayDynamicIndexing indicates whether arrays of storage buffers can be indexed by dynamically uniform integer expressions in shader code. If this feature is not enabled, resources with a descriptor type of VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC must be indexed only by constant integral expressions when aggregated into arrays in shader code. This also indicates whether shader modules can declare the **StorageBufferArrayDynamicIndexing** capability.
- shaderStorageImageArrayDynamicIndexing indicates whether arrays of storage images can be indexed by dynamically uniform integer expressions in shader code. If this feature is not enabled, resources with a descriptor type of VK_DESCRIPTOR_TYPE_STORAGE_IMAGE must be indexed only by constant integral expressions when aggregated into arrays in shader code. This also indicates whether shader modules can declare the StorageImageArrayDynamicIndexing capability.
- shaderClipDistance indicates whether clip distances are supported in shader code. If this feature is not enabled, any members decorated with the ClipDistance built-in decoration must not be read from or written to in shader modules. This also indicates whether shader modules can declare the ClipDistance capability.
- shaderCullDistance indicates whether cull distances are supported in shader code. If this feature is not enabled, any members decorated with the **CullDistance** built-in decoration must not be read from or written to in shader modules. This also indicates whether shader modules can declare the **CullDistance** capability.
- shaderFloat64 indicates whether 64-bit floats (doubles) are supported in shader code. If this feature is not enabled, 64-bit floating-point types must not be used in shader code. This also indicates whether shader modules can declare the **Float64** capability.
- shaderInt 64 indicates whether 64-bit integers (signed and unsigned) are supported in shader code. If this feature is not enabled, 64-bit integer types must not be used in shader code. This also indicates whether shader modules can declare the **Int64** capability.
- shaderInt16 indicates whether 16-bit integers (signed and unsigned) are supported in shader code. If this feature is not enabled, 16-bit integer types must not be used in shader code. This also indicates whether shader modules can declare the **Int16** capability.
- shaderResourceResidency indicates whether image operations that return resource residency information are supported in shader code. If this feature is not enabled, the OpImageSparse* instructions must not be used in shader code. This also indicates whether shader modules can declare the SparseResidency capability. The feature requires at least one of the sparseResidency* features to be supported.
- shaderResourceMinLod indicates whether image operations that specify the minimum resource level-of-detail (LOD) are supported in shader code. If this feature is not enabled, the **MinLod** image operand must not be used in shader code. This also indicates whether shader modules can declare the **MinLod** capability.
- sparseBinding indicates whether resource memory can be managed at opaque sparse block level instead of at the object level. If this feature is not enabled, resource memory must be bound only on a per-object basis using the

vkBindBufferMemory and vkBindImageMemory commands. In this case, buffers and images must not be created with VK_BUFFER_CREATE_SPARSE_BINDING_BIT and VK_IMAGE_CREATE_SPARSE_BINDING_BIT set in the flags member of the VkBufferCreateInfo and VkImageCreateInfo structures, respectively. Otherwise resource memory can be managed as described in Sparse Resource Features.

- sparseResidencyBuffer indicates whether the device can access partially resident buffers. If this feature is not enabled, buffers must not be created with VK_BUFFER_CREATE_SPARSE_RESIDENCY_BIT set in the flags member of the VkBufferCreateInfo structure.
- sparseResidencyImage2D indicates whether the device can access partially resident 2D images with 1 sample per pixel. If this feature is not enabled, images with an imageType of VK_IMAGE_TYPE_2D and samples set to VK_ SAMPLE_COUNT_1_BIT must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the flags member of the VkImageCreateInfo structure.
- sparseResidencyImage3D indicates whether the device can access partially resident 3D images. If this feature is not enabled, images with an imageType of VK_IMAGE_TYPE_3D must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the flags member of the VkImageCreateInfo structure.
- sparseResidency2Samples indicates whether the physical device can access partially resident 2D images with 2 samples per pixel. If this feature is not enabled, images with an <code>imageType</code> of VK_IMAGE_TYPE_2D and <code>samples</code> set to VK_SAMPLE_COUNT_2_BIT must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the <code>flags</code> member of the VkImageCreateInfo structure.
- sparseResidency4Samples indicates whether the physical device can access partially resident 2D images with 4 samples per pixel. If this feature is not enabled, images with an <code>imageType</code> of VK_IMAGE_TYPE_2D and <code>samples</code> set to VK_SAMPLE_COUNT_4_BIT must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the <code>flags</code> member of the VkImageCreateInfo structure.
- sparseResidency8Samples indicates whether the physical device can access partially resident 2D images with 8 samples per pixel. If this feature is not enabled, images with an <code>imageType</code> of VK_IMAGE_TYPE_2D and <code>samples</code> set to VK_SAMPLE_COUNT_8_BIT must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the <code>flags</code> member of the VkImageCreateInfo structure.
- sparseResidency16Samples indicates whether the physical device can access partially resident 2D images with 16 samples per pixel. If this feature is not enabled, images with an <code>imageType</code> of VK_IMAGE_TYPE_2D and <code>samples</code> set to VK_SAMPLE_COUNT_16_BIT must not be created with VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT set in the <code>flags</code> member of the VkImageCreateInfo structure.
- sparseResidencyAliased indicates whether the physical device can correctly access data aliased into multiple locations. If this feature is not enabled, the VK_BUFFER_CREATE_SPARSE_ALIASED_BIT and VK_IMAGE_CREATE_SPARSE_ALIASED_BIT enum values must not be used in flags members of the VkBufferCreateInfo and VkImageCreateInfo structures, respectively.
- variableMultisampleRate indicates whether all pipelines that will be bound to a command buffer during a subpass with no attachments must have the same value for VkPipelineMultisampleStateCreateInfo::rasterizationSamples. If set to VK_TRUE, the implementation supports variable multisample rates in a subpass with no attachments. If set to VK_FALSE, then all pipelines bound in such a subpass must have the same multisample rate. This has no effect in situations where a subpass uses any attachments.
- inheritedQueries indicates whether a secondary command buffer may be executed while a query is active.

Valid Usage

• If any member of this structure is VK_FALSE, as returned by vkGetPhysicalDeviceFeatures, then it must be VK_FALSE when passed as part of the VkDeviceCreateInfo struct when creating a device

31.1.1 Feature Requirements

All Vulkan graphics implementations must support the following features:

• robustBufferAccess.

All other features are not required by the Specification.

31.2 Limits

There are a variety of implementation-dependent limits.

The VkPhysicalDeviceLimits are properties of the physical device. These are available in the <code>limits</code> member of the VkPhysicalDeviceProperties structure which is returned from <code>vkGetPhysicalDeviceProperties</code>.

The VkPhysicalDeviceLimits structure is defined as:

```
typedef struct VkPhysicalDeviceLimits {
   uint32 t
                        maxImageDimension1D;
   uint32_t
                        maxImageDimension2D;
   uint32_t
                       maxImageDimension3D;
   uint32_t
                       maxImageDimensionCube;
   uint32_t
                       maxImageArrayLayers;
   uint32_t
                       maxTexelBufferElements;
   uint32_t
                       maxUniformBufferRange;
   uint32_t
                       maxStorageBufferRange;
                       maxPushConstantsSize;
   uint32_t
   uint32_t
                        maxMemoryAllocationCount;
   uint32_t
                        maxSamplerAllocationCount;
   VkDeviceSize
                       bufferImageGranularity;
   VkDeviceSize
                        sparseAddressSpaceSize;
   uint32 t
                       maxBoundDescriptorSets;
   uint32_t
                        maxPerStageDescriptorSamplers;
                       maxPerStageDescriptorUniformBuffers;
   uint32_t
   uint32_t
                       maxPerStageDescriptorStorageBuffers;
   uint32_t
                       maxPerStageDescriptorSampledImages;
   uint32_t
                       maxPerStageDescriptorStorageImages;
   uint32_t
                       maxPerStageDescriptorInputAttachments;
   uint32_t
                       maxPerStageResources;
   uint32_t
                        maxDescriptorSetSamplers;
   uint32_t
                        maxDescriptorSetUniformBuffers;
   uint32_t
                         maxDescriptorSetUniformBuffersDynamic;
   uint32_t
                         maxDescriptorSetStorageBuffers;
   uint32_t
                         maxDescriptorSetStorageBuffersDynamic;
   uint32_t
                         maxDescriptorSetSampledImages;
```

```
uint32_t
                      maxDescriptorSetStorageImages;
uint32 t
                      maxDescriptorSetInputAttachments;
uint32_t
                      maxVertexInputAttributes;
                      maxVertexInputBindings;
11int32 t
uint32 t
                      maxVertexInputAttributeOffset;
uint32 t
                      maxVertexInputBindingStride;
uint32_t
                      maxVertexOutputComponents;
uint32 t
                      maxTessellationGenerationLevel;
uint32_t
                      maxTessellationPatchSize;
uint32 t
                      maxTessellationControlPerVertexInputComponents;
uint32_t
                      maxTessellationControlPerVertexOutputComponents;
uint32 t
                      maxTessellationControlPerPatchOutputComponents;
                      maxTessellationControlTotalOutputComponents;
uint32_t
uint32_t
                      maxTessellationEvaluationInputComponents;
uint32 t
                      maxTessellationEvaluationOutputComponents;
uint32_t
                      maxGeometryShaderInvocations;
uint32 t
                      maxGeometryInputComponents;
uint32 t
                      maxGeometryOutputComponents;
uint32_t
                      maxGeometryOutputVertices;
uint32_t
                      maxGeometryTotalOutputComponents;
uint32_t
                      maxFragmentInputComponents;
uint32_t
                      maxFragmentOutputAttachments;
uint32 t
                      maxFragmentDualSrcAttachments;
                      maxFragmentCombinedOutputResources;
uint32 t
uint32 t
                      maxComputeSharedMemorySize;
                     maxComputeWorkGroupCount[3];
uint32_t
uint32_t
                     maxComputeWorkGroupInvocations;
                    maxComputeWorkGroupSize[3];
uint32 t
                     subPixelPrecisionBits;
uint32 t
uint32 t
                     subTexelPrecisionBits;
uint32 t
                     mipmapPrecisionBits;
                     maxDrawIndexedIndexValue;
uint32_t
uint32_t
                      maxDrawIndirectCount;
float
                      maxSamplerLodBias;
float
                      maxSamplerAnisotropy;
uint32 t
                     maxViewports;
uint32_t
                     maxViewportDimensions[2];
float.
                     viewportBoundsRange[2];
uint32_t
                     viewportSubPixelBits;
size_t
                     minMemoryMapAlignment;
VkDeviceSize
                     minTexelBufferOffsetAlignment;
                    minUniformBufferOffsetAlignment;
VkDeviceSize
VkDeviceSize
                     minStorageBufferOffsetAlignment;
int32 t
                     minTexelOffset;
uint32_t
                     maxTexelOffset;
                      minTexelGatherOffset;
int32_t
uint32_t
                      maxTexelGatherOffset;
float
                      minInterpolationOffset;
float
                      maxInterpolationOffset;
uint32_t
                      subPixelInterpolationOffsetBits;
                      maxFramebufferWidth;
uint32 t
                      maxFramebufferHeight;
uint32_t
uint32_t
                      maxFramebufferLayers;
VkSampleCountFlags framebufferColorSampleCounts;
                      framebufferDepthSampleCounts;
VkSampleCountFlags
VkSampleCountFlags
                      framebufferStencilSampleCounts;
VkSampleCountFlags framebufferNoAttachmentsSampleCounts;
```

```
uint32_t
                        maxColorAttachments;
   VkSampleCountFlags sampledImageColorSampleCounts;
   VkSampleCountFlags sampledImageIntegerSampleCounts;
   VkSampleCountFlags sampledImageDepthSampleCounts;
   VkSampleCountFlags sampledImageStencilSampleCounts;
   VkSampleCountFlags storageImageSampleCounts;
                       maxSampleMaskWords;
   uint32_t
   VkBool32
                        timestampComputeAndGraphics;
   float
                        timestampPeriod;
                      maxClipDistances;
   uint32 t
                      maxCullDistances;
   uint32_t
                      maxCombinedClipAndCullDistances;
   uint32_t
   uint32_t
                      discreteQueuePriorities;
                      pointSizeRange[2];
   float
   float
                      lineWidthRange[2];
   float
                      pointSizeGranularity;
   float
                      lineWidthGranularity;
                    strictLines;
   VkBool32
   VkBool32
                       standardSampleLocations;
                      optimalBufferCopyOffsetAlignment;
   VkDeviceSize
   VkDeviceSize
                       optimalBufferCopyRowPitchAlignment;
   VkDeviceSize
                       nonCoherentAtomSize;
} VkPhysicalDeviceLimits;
```

- maxImageDimension1D is the maximum dimension (width) of an image created with an imageType of VK_IMAGE TYPE 1D.
- maxImageDimension2D is the maximum dimension (width or height) of an image created with an imageType of VK_IMAGE_TYPE_2D and without VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT set in flags.
- maxImageDimension3D is the maximum dimension (width, height, or depth) of an image created with an imageType of VK_IMAGE_TYPE_3D.
- maxImageDimensionCube is the maximum dimension (width or height) of an image created with an imageType of VK_IMAGE_TYPE_2D and with VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT set in flags.
- maxImageArrayLayers is the maximum number of layers (arrayLayers) for an image.
- maxTexelBufferElements is the maximum number of addressable texels for a buffer view created on a buffer which was created with the VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT or VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT set in the usage member of the VkBufferCreateInfo structure.
- maxUniformBufferRange is the maximum value that can be specified in the range member of any VkDescriptorBufferInfo structures passed to a call to vkUpdateDescriptorSets for descriptors of type VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC.
- maxStorageBufferRange is the maximum value that can be specified in the range member of any VkDescriptorBufferInfo structures passed to a call to vkUpdateDescriptorSets for descriptors of type VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC.
- maxPushConstantsSize is the maximum size, in bytes, of the pool of push constant memory. For each of the push constant ranges indicated by the pPushConstantRanges member of the VkPipelineLayoutCreateInfo structure, offset + size must be less than or equal to this limit.
- maxMemoryAllocationCount is the maximum number of device memory allocations, as created by vkAllocateMemory, which can simultaneously exist.

- maxSamplerAllocationCount is the maximum number of sampler objects, as created by vkCreateSampler, which can simultaneously exist on a device.
- bufferImageGranularity is the granularity, in bytes, at which buffer or linear image resources, and optimal image resources can be bound to adjacent offsets in the same VkDeviceMemory object without aliasing. See Buffer-Image Granularity for more details.
- sparseAddressSpaceSize is the total amount of address space available, in bytes, for sparse memory resources. This is an upper bound on the sum of the size of all sparse resources, regardless of whether any memory is bound to them.
- maxBoundDescriptorSets is the maximum number of descriptor sets that can be simultaneously used by a pipeline. All DescriptorSet decorations in shader modules must have a value less than maxBoundDescriptorSets. See Section 13.2.
- maxPerStageDescriptorSamplers is the maximum number of samplers that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_SAMPLER or VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER count against this limit. A descriptor is accessible to a shader stage when the stageFlags member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. See Section 13.1.2 and Section 13.1.4.
- maxPerStageDescriptorUniformBuffers is the maximum number of uniform buffers that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC count against this limit. A descriptor is accessible to a shader stage when the stageFlags member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. See Section 13.1.7 and Section 13.1.9.
- maxPerStageDescriptorStorageBuffers is the maximum number of storage buffers that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_STORAGE_BUFFER or VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC count against this limit. A descriptor is accessible to a pipeline shader stage when the <code>stageFlags</code> member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. See Section 13.1.8 and Section 13.1.10.
- maxPerStageDescriptorSampledImages is the maximum number of sampled images that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, or VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER count against this limit. A descriptor is accessible to a pipeline shader stage when the stageFlags member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. See Section 13.1.4, Section 13.1.3, and Section 13.1.5.
- maxPerStageDescriptorStorageImages is the maximum number of storage images that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, or VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER count against this limit. A descriptor is accessible to a pipeline shader stage when the stageFlags member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. See Section 13.1.1, and Section 13.1.6.
- maxPerStageDescriptorInputAttachments is the maximum number of input attachments that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT count against this limit. A descriptor is accessible to a pipeline shader stage when the stageFlags member of the VkDescriptorSetLayoutBinding structure has the bit for that shader stage set. These are only supported for the fragment stage. See Section 13.1.11.
- maxPerStageResources is the maximum number of resources that can be accessible to a single shader stage in a pipeline layout. Descriptors with a type of VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER, VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER, VK_

DESCRIPTOR_TYPE_UNIFORM_BUFFER, VK_DESCRIPTOR_TYPE_STORAGE_BUFFER, VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC, VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC, or VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT count against this limit. For the fragment shader stage the framebuffer color attachments also count against this limit.

- maxDescriptorSetSamplers is the maximum number of samplers that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_ DESCRIPTOR_TYPE_SAMPLER or VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER count against this limit. See Section 13.1.2 and Section 13.1.4.
- maxDescriptorSetUniformBuffers is the maximum number of uniform buffers that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC count against this limit. See Section 13.1.7 and Section 13.1.9.
- maxDescriptorSetUniformBuffersDynamic is the maximum number of dynamic uniform buffers that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers.

 Descriptors with a type of VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC count against this limit. See Section 13.1.9.
- maxDescriptorSetStorageBuffers is the maximum number of storage buffers that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC count against this limit. See Section 13.1.8 and Section 13.1.10.
- maxDescriptorSetStorageBuffersDynamic is the maximum number of dynamic storage buffers that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers.

 Descriptors with a type of VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC count against this limit. See Section 13.1.10.
- maxDescriptorSetSampledImages is the maximum number of sampled images that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER, VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE, or VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER count against this limit. See Section 13.1.4, Section 13.1.3, and Section 13.1.5.
- maxDescriptorSetStorageImages is the maximum number of storage images that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_DESCRIPTOR_TYPE_STORAGE_IMAGE, or VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER count against this limit. See Section 13.1.1, and Section 13.1.6.
- maxDescriptorSetInputAttachments is the maximum number of input attachments that can be included in descriptor bindings in a pipeline layout across all pipeline shader stages and descriptor set numbers. Descriptors with a type of VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT count against this limit. See Section 13.1.11.
- maxVertexInputAttributes is the maximum number of vertex input attributes that can be specified for a graphics pipeline. These are described in the array of VkVertexInputAttributeDescription structures that are provided at graphics pipeline creation time via the pVertexAttributeDescriptions member of the VkPipelineVertexInputStateCreateInfo structure. See Section 20.1 and Section 20.2.
- maxVertexInputBindings is the maximum number of vertex buffers that can be specified for providing vertex attributes to a graphics pipeline. These are described in the array of VkVertexInputBindingDescription structures that are provided at graphics pipeline creation time via the pVertexBindingDescriptions member of the VkPipelineVertexInputStateCreateInfo structure. The binding member of VkVertexInputBindingDescription must be less than this limit. See Section 20.2.

- maxVertexInputAttributeOffset is the maximum vertex input attribute offset that can be added to the vertex input binding stride. The offset member of the VkVertexInputAttributeDescription structure must be less than or equal to this limit. See Section 20.2.
- maxVertexInputBindingStride is the maximum vertex input binding stride that can be specified in a vertex input binding. The stride member of the VkVertexInputBindingDescription structure must be less than or equal to this limit. See Section 20.2.
- maxVertexOutputComponents is the maximum number of components of output variables which can be output by a vertex shader. See Section 8.5.
- maxTessellationGenerationLevel is the maximum tessellation generation level supported by the fixed-function tessellation primitive generator. See Chapter 21.
- maxTessellationPatchSize is the maximum patch size, in vertices, of patches that can be processed by the tessellation control shader and tessellation primitive generator. The patchControlPoints member of the VkPipelineTessellationStateCreateInfo structure specified at pipeline creation time and the value provided in the OutputVertices execution mode of shader modules must be less than or equal to this limit. See Chapter 21.
- maxTessellationControlPerVertexInputComponents is the maximum number of components of input variables which can be provided as per-vertex inputs to the tessellation control shader stage.
- maxTessellationControlPerVertexOutputComponents is the maximum number of components of per-vertex output variables which can be output from the tessellation control shader stage.
- maxTessellationControlPerPatchOutputComponents is the maximum number of components of per-patch output variables which can be output from the tessellation control shader stage.
- maxTessellationControlTotalOutputComponents is the maximum total number of components of per-vertex and per-patch output variables which can be output from the tessellation control shader stage.
- maxTessellationEvaluationInputComponents is the maximum number of components of input variables which can be provided as per-vertex inputs to the tessellation evaluation shader stage.
- maxTessellationEvaluationOutputComponents is the maximum number of components of per-vertex output variables which can be output from the tessellation evaluation shader stage.
- maxGeometryShaderInvocations is the maximum invocation count supported for instanced geometry shaders. The value provided in the **Invocations** execution mode of shader modules must be less than or equal to this limit. See Chapter 22.
- maxGeometryInputComponents is the maximum number of components of input variables which can be provided as inputs to the geometry shader stage.
- maxGeometryOutputComponents is the maximum number of components of output variables which can be output from the geometry shader stage.
- maxGeometryOutputVertices is the maximum number of vertices which can be emitted by any geometry shader.
- maxGeometryTotalOutputComponents is the maximum total number of components of output, across all emitted vertices, which can be output from the geometry shader stage.
- maxFragmentInputComponents is the maximum number of components of input variables which can be provided as inputs to the fragment shader stage.
- maxFragmentOutputAttachments is the maximum number of output attachments which can be written to by the fragment shader stage.

- maxFragmentDualSrcAttachments is the maximum number of output attachments which can be written to by the fragment shader stage when blending is enabled and one of the dual source blend modes is in use. See Section 26.1.2 and dualSrcBlend.
- maxFragmentCombinedOutputResources is the total number of storage buffers, storage images, and output buffers which can be used in the fragment shader stage.
- maxComputeSharedMemorySize is the maximum total storage size, in bytes, of all variables declared with the **WorkgroupLocal** storage class in shader modules (or with the **shared** storage qualifier in GLSL) in the compute shader stage.
- maxComputeWorkGroupCount[3] is the maximum number of local workgroups that can be dispatched by a single dispatch command. These three values represent the maximum number of local workgroups for the X, Y, and Z dimensions, respectively. The x, y, and z parameters to the vkCmdDispatch command, or members of the VkDispatchIndirectCommand structure must be less than or equal to the corresponding limit. See Chapter 27.
- maxComputeWorkGroupInvocations is the maximum total number of compute shader invocations in a single local workgroup. The product of the X, Y, and Z sizes as specified by the **LocalSize** execution mode in shader modules and by the object decorated by the **WorkgroupSize** decoration must be less than or equal to this limit.
- maxComputeWorkGroupSize[3] is the maximum size of a local compute workgroup, per dimension. These three values represent the maximum local workgroup size in the X, Y, and Z dimensions, respectively. The x, y, and z sizes specified by the **LocalSize** execution mode and by the object decorated by the **WorkgroupSize** decoration in shader modules must be less than or equal to the corresponding limit.
- subPixelPrecisionBits is the number of bits of subpixel precision in framebuffer coordinates x_f and y_f . See Chapter 24.
- subTexelPrecisionBits is the number of bits of precision in the division along an axis of an image used for minification and magnification filters. 2^{subTexelPrecisionBits} is the actual number of divisions along each axis of the image represented. The filtering hardware will snap to these locations when computing the filtered results.
- mipmapPrecisionBits is the number of bits of division that the LOD calculation for mipmap fetching get snapped to when determining the contribution from each mip level to the mip filtered results. 2^{mipmapPrecisionBits} is the actual number of divisions.



Note

For example, if this value is 2 bits then when linearly filtering between two levels, each level could: contribute: 0%, 33%, 66%, or 100% (this is just an example and the amount of contribution should be covered by different equations in the spec).

- maxDrawIndexedIndexValue is the maximum index value that can be used for indexed draw calls when using 32-bit indices. This excludes the primitive restart index value of 0xFFFFFFF. See fullDrawIndexUint32.
- maxDrawIndirectCount is the maximum draw count that is supported for indirect draw calls. See multiDrawIndirect.
- maxSamplerLodBias is the maximum absolute sampler level of detail bias. The sum of the mipLodBias member of the VkSamplerCreateInfo structure and the Bias operand of image sampling operations in shader modules (or 0 if no Bias operand is provided to an image sampling operation) are clamped to the range [-maxSamplerLodBias, +maxSamplerLodBias]. See [samplers-mipLodBias].
- maxSamplerAnisotropy is the maximum degree of sampler anisotropy. The maximum degree of anisotropic filtering used for an image sampling operation is the minimum of the maxAnisotropy member of the VkSamplerCreateInfo structure and this limit. See [samplers-maxAnisotropy].

- maxViewports is the maximum number of active viewports. The viewportCount member of the VkPipelineViewportStateCreateInfo structure that is provided at pipeline creation must be less than or equal to this limit.
- maxViewportDimensions[2] are the maximum viewport dimensions in the X (width) and Y (height) dimensions, respectively. The maximum viewport dimensions must be greater than or equal to the largest image which can be created and used as a framebuffer attachment. See Controlling the Viewport.
- viewportBoundsRange[2] is the [minimum, maximum] range that the corners of a viewport must be contained in. This range must be at least

 $[-2 \times \textit{maxViewportDimensions}, 2 \times \textit{maxViewportDimensions} - 1].$

See Controlling the Viewport.

Note



The intent of the viewportBoundsRange limit is to allow a maximum sized viewport to be arbitrarily shifted relative to the output target as long as at least some portion intersects. This would give a bounds limit of $[-maxViewportDimensions+1,2\times maxViewportDimensions-1]$, which would allow all possible non-empty-set intersections of the output target and the viewport. Since these numbers are typically powers of two, picking the signed number range using the smallest possible number of bits ends up with the specified range.

- *viewportSubPixelBits* is the number of bits of subpixel precision for viewport bounds. The subpixel precision that floating-point viewport bounds are interpreted at is given by this limit.
- minMemoryMapAlignment is the minimum required alignment, in bytes, of host visible memory allocations within the host address space. When mapping a memory allocation with vkMapMemory, subtracting offset bytes from the returned pointer will always produce an integer multiple of this limit. See Section 10.2.1.
- minTexelBufferOffsetAlignment is the minimum required alignment, in bytes, for the offset member of the VkBufferViewCreateInfo structure for texel buffers. When a buffer view is created for a buffer which was created with VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT or VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT set in the usage member of the VkBufferCreateInfo structure, the offset must be an integer multiple of this limit.
- minUniformBufferOffsetAlignment is the minimum required alignment, in bytes, for the offset member of the VkDescriptorBufferInfo structure for uniform buffers. When a descriptor of type VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER or VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC is updated, the offset must be an integer multiple of this limit. Similarly, dynamic offsets for uniform buffers must be multiples of this limit.
- minStorageBufferOffsetAlignment is the minimum required alignment, in bytes, for the offset member of the VkDescriptorBufferInfo structure for storage buffers. When a descriptor of type VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC is updated, the offset must be an integer multiple of this limit. Similarly, dynamic offsets for storage buffers must be multiples of this limit.
- minTexelOffset is the minimum offset value for the ConstOffset image operand of any of the OpImageSample* or OpImageFetch* image instructions.
- maxTexelOffset is the maximum offset value for the ConstOffset image operand of any of the OpImageSample* or OpImageFetch* image instructions.
- minTexelGatherOffset is the minimum offset value for the Offset or ConstOffsets image operands of any of the OpImage*Gather image instructions.
- maxTexelGatherOffset is the maximum offset value for the Offset or ConstOffsets image operands of any of the OpImage*Gather image instructions.

- minInterpolationOffset is the minimum negative offset value for the **offset** operand of the **InterpolateAtOffset** extended instruction.
- maxInterpolationOffset is the maximum positive offset value for the offset operand of the InterpolateAtOffset extended instruction.
- subPixelInterpolationOffsetBits is the number of subpixel fractional bits that the **x** and **y** offsets to the **InterpolateAtOffset** extended instruction may be rounded to as fixed-point values.
- maxFramebufferWidth is the maximum width for a framebuffer. The width member of the VkFramebufferCreateInfo structure must be less than or equal to this limit.
- maxFramebufferHeight is the maximum height for a framebuffer. The height member of the VkFramebufferCreateInfo structure must be less than or equal to this limit.
- maxFramebufferLayers is the maximum layer count for a layered framebuffer. The layers member of the VkFramebufferCreateInfo structure must be less than or equal to this limit.
- framebufferColorSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the color sample counts that are supported for all framebuffer color attachments.
- framebufferDepthSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the supported depth sample counts for all framebuffer depth/stencil attachments, when the format includes a depth component.
- framebufferStencilSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the supported stencil sample counts for all framebuffer depth/stencil attachments, when the format includes a stencil component.
- framebufferNoAttachmentsSampleCounts is a bitmask of VkSampleCountFlagBits bits indicating the supported sample counts for a framebuffer with no attachments.
- maxColorAttachments is the maximum number of color attachments that can be used by a subpass in a render pass. The colorAttachmentCount member of the VkSubpassDescription structure must be less than or equal to this limit.
- sampledImageColorSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the sample counts supported for all 2D images created with VK_IMAGE_TILING_OPTIMAL, usage containing VK_IMAGE_USAGE_SAMPLED_BIT, and a non-integer color format.
- sampledImageIntegerSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the sample counts supported for all 2D images created with VK_IMAGE_TILING_OPTIMAL, usage containing VK_IMAGE_USAGE_SAMPLED_BIT, and an integer color format.
- sampledImageDepthSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the sample counts supported for all 2D images created with VK_IMAGE_TILING_OPTIMAL, usage containing VK_IMAGE_USAGE_SAMPLED_BIT, and a depth format.
- sampledImageStencilSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the sample supported for all 2D images created with VK_IMAGE_TILING_OPTIMAL, usage containing VK_IMAGE_USAGE_SAMPLED_BIT, and a stencil format.
- storageImageSampleCounts is a bitmask¹ of VkSampleCountFlagBits bits indicating the sample counts supported for all 2D images created with VK_IMAGE_TILING_OPTIMAL, and usage containing VK_IMAGE_USAGE STORAGE BIT.
- maxSampleMaskWords is the maximum number of array elements of a variable decorated with the SampleMask built-in decoration.

- timestampComputeAndGraphics indicates support for timestamps on all graphics and compute queues. If this limit is set to VK_TRUE, all queues that advertise the VK_QUEUE_GRAPHICS_BIT or VK_QUEUE_COMPUTE_BIT in the VkQueueFamilyProperties::queueFlags support VkQueueFamilyProperties::timestampValidBits of at least 36. See Timestamp Queries.
- timestampPeriod is the number of nanoseconds required for a timestamp query to be incremented by 1. See Timestamp Queries.
- maxClipDistances is the maximum number of clip distances that can be used in a single shader stage. The size of any array declared with the **ClipDistance** built-in decoration in a shader module must be less than or equal to this limit.
- maxCullDistances is the maximum number of cull distances that can be used in a single shader stage. The size of any array declared with the CullDistance built-in decoration in a shader module must be less than or equal to this limit.
- maxCombinedClipAndCullDistances is the maximum combined number of clip and cull distances that can be used in a single shader stage. The sum of the sizes of any pair of arrays declared with the ClipDistance and CullDistance built-in decoration used by a single shader stage in a shader module must be less than or equal to this limit.
- discreteQueuePriorities is the number of discrete priorities that can be assigned to a queue based on the value of each member of VkDeviceQueueCreateInfo::pQueuePriorities. This must be at least 2, and levels must be spread evenly over the range, with at least one level at 1.0, and another at 0.0. See Section 4.3.4.
- pointSizeRange[2] is the range [minimum, maximum] of supported sizes for points. Values written to variables decorated with the **PointSize** built-in decoration are clamped to this range.
- lineWidthRange[2] is the range [minimum, maximum] of supported widths for lines. Values specified by the lineWidth member of the VkPipelineRasterizationStateCreateInfo or the lineWidth parameter to **vkCmdSetLineWidth** are clamped to this range.
- pointSizeGranularity is the granularity of supported point sizes. Not all point sizes in the range defined by pointSizeRange are supported. This limit specifies the granularity (or increment) between successive supported point sizes.
- lineWidthGranularity is the granularity of supported line widths. Not all line widths in the range defined by lineWidthRange are supported. This limit specifies the granularity (or increment) between successive supported line widths.
- strictLines indicates whether lines are rasterized according to the preferred method of rasterization. If set to VK_FALSE, lines may be rasterized under a relaxed set of rules. If set to VK_TRUE, lines are rasterized as per the strict definition. See Basic Line Segment Rasterization.
- standardSampleLocations indicates whether rasterization uses the standard sample locations as documented in Multisampling. If set to VK_TRUE, the implementation uses the documented sample locations. If set to VK_FALSE, the implementation may use different sample locations.
- optimalBufferCopyOffsetAlignment is the optimal buffer offset alignment in bytes for **vkCmdCopyBufferToImage** and **vkCmdCopyImageToBuffer**. The per texel alignment requirements are still enforced, this is just an additional alignment recommendation for optimal performance and power.
- optimalBufferCopyRowPitchAlignment is the optimal buffer row pitch alignment in bytes for **vkCmdCopyBufferToImage** and **vkCmdCopyImageToBuffer**. Row pitch is the number of bytes between texels with the same X coordinate in adjacent rows (Y coordinates differ by one). The per texel alignment requirements are still enforced, this is just an additional alignment recommendation for optimal performance and power.

• nonCoherentAtomSize is the size and alignment in bytes that bounds concurrent access to host-mapped device memory.

1

For all bitmasks of type VkSampleCountFlags above, possible values include:

```
typedef enum VkSampleCountFlagBits {
    VK_SAMPLE_COUNT_1_BIT = 0x00000001,
    VK_SAMPLE_COUNT_2_BIT = 0x00000002,
    VK_SAMPLE_COUNT_4_BIT = 0x00000004,
    VK_SAMPLE_COUNT_8_BIT = 0x00000008,
    VK_SAMPLE_COUNT_16_BIT = 0x00000010,
    VK_SAMPLE_COUNT_32_BIT = 0x00000020,
    VK_SAMPLE_COUNT_64_BIT = 0x000000040,
} VkSampleCountFlagBits;
```

The sample count limits defined above represent the minimum supported sample counts for each image type. Individual images may support additional sample counts, which are queried using vkGetPhysicalDeviceImageFormatProperties as described in Supported Sample Counts.

31.2.1 Limit Requirements

The following table specifies the required minimum/maximum for all Vulkan graphics implementations. Where a limit corresponds to a fine-grained device feature which is optional, the feature name is listed with two required limits, one when the feature is supported and one when it is not supported. If an implementation supports a feature, the limits reported are the same whether or not the feature is enabled.

Type	Limit	Feature
uint32_t	maxImageDimension1D	-
uint32_t	maxImageDimension2D	-
uint32_t	maxImageDimension3D	-
uint32_t	maxImageDimensionCube	-
uint32_t	maxImageArrayLayers	-
uint32_t	maxTexelBufferElements	-
uint32_t	maxUniformBufferRange	-
uint32_t	maxStorageBufferRange	-
uint32_t	maxPushConstantsSize	-
uint32_t	maxMemoryAllocationCount	-
uint32_t	maxSamplerAllocationCount	-
VkDeviceSize	bufferImageGranularity	-
VkDeviceSize	sparseAddressSpaceSize	sparseBinding
uint32_t	maxBoundDescriptorSets	-
uint32_t	maxPerStageDescriptorSamplers	-
uint32_t	maxPerStageDescriptorUniformBuffers	-
uint32_t	maxPerStageDescriptorStorageBuffers	-
uint32_t	maxPerStageDescriptorSampledImages	-
uint32_t	maxPerStageDescriptorStorageImages	-
uint32_t	maxPerStageDescriptorInputAttachments	-
uint32_t	maxPerStageResources	-

Table 31.1: Required Limit Types

Table 31.1: (continued)

uimi32_t maxDescriptorSetUniformBuffers - uimi32_t maxDescriptorSetUniformBuffersDynamic - uimi32_t maxDescriptorSetStorageBuffers - uimi32_t maxDescriptorSetStorageBuffersDynamic - uimi32_t maxDescriptorSetStorageBuffersDynamic - uimi32_t maxDescriptorSetStorageImages - uimi32_t maxDescriptorSetInputAttrabutes - uimi32_t maxVertexInputBindings - uimi32_t maxVertexInputBindings - uimi32_t maxVertexInputBindings - uimi32_t maxVertexInputBindingstride - uimi32_t maxTessellationGenerationLevel tessellationShader uimi32_t maxTessellationControlPerVertexInputComponents tessellationShader uimi32_t maxTessellationControlPerVertexOutputComponents tessellationShader uimi32_t maxTessellationControlPerPertexOutputComponents tessellationShader uimi32_t maxTessellationEvaluationInputComponents tessellationShader uimi32_t maxTessellationEvaluationOutputComponents tessell	Туре	Limit	Feature
uint32_t maxDescriptorSetUniformBuffersDynamic - uint32_t maxDescriptorSetStorageBuffersDynamic - uint32_t maxDescriptorSetStorageBuffersDynamic - uint32_t maxDescriptorSetStorageImages - uint32_t maxDescriptorSetInputAttributes - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexInputBindingStride - uint32_t maxVertexInputBindingStride - uint32_t maxTessellationCenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatehOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationDutputComponents tessellationShader uint32_t maxGeometryAbaderInvocations geometryShader uint32_t maxGeometryInputC	uint32_t	maxDescriptorSetSamplers	-
uint32_t maxDescriptorSetStorageBuffers - uint32_t maxDescriptorSetStorageBuffersDynamic - uint32_t maxDescriptorSetStorageImages - uint32_t maxDescriptorSetInputAttrachments - uint32_t maxDescriptorSetInputAttributeS - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationDutputComponents tessellationShader uint32_t maxTessellationEvaluationDutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents	uint32_t	maxDescriptorSetUniformBuffers	-
uint32_t maxDescriptorSetStorageBuffersDynamic - uint32_t maxDescriptorSetStorageBuages - uint32_t maxDescriptorSetStorageBuages - uint32_t maxVertexInputAttributes - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutput Vertices	uint32_t	maxDescriptorSetUniformBuffersDynamic	-
uint32_t maxDescriptorSetSampledImages - uint32_t maxDescriptorSetStorageImages - uint32_t maxDescriptorSetInputAttachments - uint32_t maxVertexInputAttributes - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices	uint32_t	maxDescriptorSetStorageBuffers	-
uint32_t maxDescriptorSetInputAttachments - uint32_t maxDescriptorSetInputAttachments - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingS - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t	uint32_t	maxDescriptorSetStorageBuffersDynamic	-
uint32_t maxDescriptorSetInputAttachments - uint32_t maxDescriptorSetInputAttachments - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingS - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t	uint32_t	maxDescriptorSetSampledImages	-
uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxFragmentDoutputAttachments - uint32_t	uint32_t		-
uint32_t maxVertexInputBindings - uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentOutputAtachments -	uint32_t	maxDescriptorSetInputAttachments	-
uint32_t maxVertexInputBindingStride - uint32_t maxVertexOutputComponents - uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryDutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxFragmentDuptComponents geometryShader uint32_t maxFragmentOutputAttachments -	uint32_t	maxVertexInputAttributes	-
uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPexteXOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxComputeWorkGroupCount -	uint32_t	maxVertexInputBindings	-
uint32_t maxVertexOutputComponents - uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPexteXOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxComputeWorkGroupCount -	uint32_t	maxVertexInputAttributeOffset	-
uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentOutputComponents geometryShader uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments dualSrcBlend uint32_t maxComputeWorkGroupCount -	uint32_t		-
uint32_t maxTessellationGenerationLevel tessellationShader uint32_t maxTessellationControlPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentOutputComponents geometryShader uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments dualSrcBlend uint32_t maxComputeWorkGroupCount -	uint32_t		-
uint32_t maxTessellationOntrolPerVertexInputComponents tessellationShader uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryShader partyShader geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxFragmentInputComponents geometryShader uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t	uint32 t		tessellationShader
uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentDualSreAttachments - uint32_t maxFragmentDualSreAttachments - uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - 3 × uint32_t subFexelPrecision		maxTessellationPatchSize	tessellationShader
uint32_t maxTessellationControlPerVertexOutputComponents tessellationShader uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxFragmentInputComponents geometryShader uint32_t maxFragmentOutputComponents - uint32_t maxFragmentDualSreAttachments dualSrcBlend uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupCount - uint32_t subFexelPrecisionBits			
uint32_t maxTessellationControlPerPatchOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryTotalOutputComponents - uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments dualSrcBlend uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupCount - uint32_t subPixelPrecisionBits - uint32_t subPixelPrecisionBits - u			tessellationShader
uint32_t maxTessellationControlTotalOutputComponents tessellationShader uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryTotalOutputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments dualSrcBlend uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t subPixelPrecisionBits - uint32_t subPixelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDa			
uint32_t maxTessellationEvaluationInputComponents tessellationShader uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryTotalOutputComponents geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments dualSrcBlend uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerAnisotr			tessellationShader
uint32_t maxTessellationEvaluationOutputComponents tessellationShader uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryTotalOutputComponents geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subPixelPrecisionBits - uint32_t mindal maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndexedIndexValue multiDrawIndexUint32 uint32_t maxDrawIndexedIndexValue multiDrawIndex Uint32 uint32_t maxSamplerAniso		maxTessellationEvaluationInputComponents	tessellationShader
uint32_t maxGeometryShaderInvocations geometryShader uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryTotalOutputComponents geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments dualSrcBlend uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subPixelPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewport -			tessellationShader
uint32_t maxGeometryInputComponents geometryShader uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments - uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t		geometryShader
uint32_t maxGeometryOutputComponents geometryShader uint32_t maxGeometryOutputVertices geometryShader uint32_t maxGeometryTotalOutputComponents geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments dualSrcBlend uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subFixelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t		geometryShader
uint32_tmaxGeometryOutputVerticesgeometryShaderuint32_tmaxGeometryTotalOutputComponentsgeometryShaderuint32_tmaxFragmentInputComponents-uint32_tmaxFragmentOutputAttachments-uint32_tmaxFragmentDualSrcAttachmentsdualSrcBlenduint32_tmaxFragmentCombinedOutputResources-uint32_tmaxComputeSharedMemorySize-3 × uint32_tmaxComputeWorkGroupCount-uint32_tmaxComputeWorkGroupInvocations-3 × uint32_tsubPixelPrecisionBits-uint32_tsubPixelPrecisionBits-uint32_tsubTexelPrecisionBits-uint32_tmipmapPrecisionBits-uint32_tmipmapPrecisionBits-uint32_tmaxDrawIndexedIndex ValuefullDrawIndex Uint32uint32_tmaxDrawIndirectCountmultiDrawIndirectfloatmaxSampler LodBias-floatmaxSampler Anisotropysampler Anisotropyuint32_tmaxViewportsmultiViewport2 × uint32_tmaxViewportDimensions-	uint32_t		geometryShader
uint32_t maxGeometryTotalOutputComponents geometryShader uint32_t maxFragmentInputComponents - uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentDualSrcAttachments dualSrcBlend uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndex Value fullDrawIndex Uint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t		
uint32_t maxFragmentOutputAttachments - uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxGeometryTotalOutputComponents	geometryShader
uint32_tmaxFragmentDualSrcAttachmentsdualSrcBlenduint32_tmaxFragmentCombinedOutputResources-uint32_tmaxComputeSharedMemorySize-3 × uint32_tmaxComputeWorkGroupCount-uint32_tmaxComputeWorkGroupInvocations-3 × uint32_tmaxComputeWorkGroupSize-uint32_tsubPixelPrecisionBits-uint32_tsubTexelPrecisionBits-uint32_tmipmapPrecisionBits-uint32_tmaxDrawIndexedIndexValuefullDrawIndexUint32uint32_tmaxDrawIndirectCountmultiDrawIndirectfloatmaxSamplerLodBias-floatmaxSamplerAnisotropysamplerAnisotropyuint32_tmaxViewportsmultiViewport2 × uint32_tmaxViewportDimensions-	uint32_t	maxFragmentInputComponents	-
uint32_t maxFragmentCombinedOutputResources - uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxFragmentOutputAttachments	-
uint32_t maxComputeSharedMemorySize - 3 × uint32_t maxComputeWorkGroupInvocations - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxFragmentDualSrcAttachments	dualSrcBlend
3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t		-
3 × uint32_t maxComputeWorkGroupCount - uint32_t maxComputeWorkGroupSize - 3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxComputeSharedMemorySize	-
3 × uint32_t maxComputeWorkGroupSize - uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	3 × uint32_t		-
uint32_t subPixelPrecisionBits - uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxComputeWorkGroupInvocations	-
uint32_t subTexelPrecisionBits - uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	3 × uint32_t	maxComputeWorkGroupSize	-
uint32_t mipmapPrecisionBits - uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	subPixelPrecisionBits	-
uint32_t maxDrawIndexedIndexValue fullDrawIndexUint32 uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	subTexelPrecisionBits	-
uint32_t maxDrawIndirectCount multiDrawIndirect float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	mipmapPrecisionBits	-
float maxSamplerLodBias - float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxDrawIndexedIndexValue	fullDrawIndexUint32
float maxSamplerAnisotropy samplerAnisotropy uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	uint32_t	maxDrawIndirectCount	multiDrawIndirect
uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -	float	maxSamplerLodBias	-
uint32_t maxViewports multiViewport 2 × uint32_t maxViewportDimensions -			samplerAnisotropy
2 × uint32_t maxViewportDimensions -			
	2 × uint32_t		-
2 × float viewportBoundsRange -	$2 \times \text{float}$	viewportBoundsRange	-
uint32_t viewportSubPixelBits -			-
size_t minMemoryMapAlignment -			-
VkDeviceSize minTexelBufferOffsetAlignment -			-
VkDeviceSize minUniformBufferOffsetAlignment -			-
VkDeviceSize minStorageBufferOffsetAlignment -			-

Table 31.1: (continued)

Type	Limit	Feature
int32_t	minTexelOffset	-
uint32_t	maxTexelOffset	-
int32_t	minTexelGatherOffset	shaderImageGatherExtended
uint32_t	maxTexelGatherOffset	shaderImageGatherExtended
float	minInterpolationOffset	sampleRateShading
float	maxInterpolationOffset	sampleRateShading
uint32_t	subPixelInterpolationOffsetBits	sampleRateShading
uint32_t	maxFramebufferWidth	-
uint32_t	maxFramebufferHeight	-
uint32_t	maxFramebufferLayers	-
VkSampleCountFlags	framebufferColorSampleCounts	-
VkSampleCountFlags	framebufferDepthSampleCounts	-
VkSampleCountFlags	framebufferStencilSampleCounts	-
VkSampleCountFlags	framebufferNoAttachmentsSampleCounts	-
uint32_t	maxColorAttachments	-
VkSampleCountFlags	sampledImageColorSampleCounts	-
VkSampleCountFlags	sampledImageIntegerSampleCounts	-
VkSampleCountFlags	sampledImageDepthSampleCounts	-
VkSampleCountFlags	sampledImageStencilSampleCounts	-
VkSampleCountFlags	storageImageSampleCounts	shaderStorageImageMultisample
uint32_t	maxSampleMaskWords	-
VkBool32	timestampComputeAndGraphics	-
float	timestampPeriod	-
uint32_t	maxClipDistances	shaderClipDistance
uint32_t	maxCullDistances	shaderCullDistance
uint32_t	maxCombinedClipAndCullDistances	shaderCullDistance
uint32_t	discreteQueuePriorities	-
$2 \times \text{float}$	pointSizeRange	largePoints
$2 \times \text{float}$	lineWidthRange	wideLines
float	pointSizeGranularity	largePoints
float	lineWidthGranularity	wideLines
VkBool32	strictLines	-
VkBool32	standardSampleLocations	-
VkDeviceSize	optimalBufferCopyOffsetAlignment	-
VkDeviceSize	optimalBufferCopyRowPitchAlignment	-
VkDeviceSize	nonCoherentAtomSize	-

Table 31.2: Required Limits

Limit	Unsupported Limit	Supported Limit	Limit Type ¹
maxImageDimension1D	-	4096	min
maxImageDimension2D	-	4096	min
maxImageDimension3D	-	256	min
maxImageDimensionCube	-	4096	min
maxImageArrayLayers	-	256	min

Table 31.2: (continued)

Limit	Unsupported Limit	Supported Limit	Limit Type ¹
maxTexelBufferElements	-	65536	min
maxUniformBufferRange	-	16384	min
maxStorageBufferRange	_	2^{27}	min
maxPushConstantsSize	-	128	min
maxMemoryAllocationCount	-	4096	min
maxSamplerAllocationCount	-	4000	min
bufferImageGranularity	-	131072	max
sparseAddressSpaceSize	0	2 ³¹	min
maxBoundDescriptorSets	-	4	min
maxPerStageDescriptorSamplers	-	16	min
maxPerStageDescriptorUniformBuffers	-	12	min
maxPerStageDescriptorStorageBuffers	-	4	min
maxPerStageDescriptorSampledImages	_	16	min
maxPerStageDescriptorStorageImages	-	4	min
maxPerStageDescriptorInputAttachments	_	4	min
maxPerStageResources	_	128 2	min
maxDescriptorSetSamplers	_	96	min,
•			6×PerStage
maxDescriptorSetUniformBuffers	-	72	min,
			6×PerStage
maxDescriptorSetUniformBuffersDynamic	-	8	min
maxDescriptorSetStorageBuffers	-	24	min, 6×PerStage
maxDescriptorSetStorageBuffersDynamic	-	4	min
maxDescriptorSetSampledImages	-	96	min,
			6×PerStage
maxDescriptorSetStorageImages	-	24	min, 6×PerStage
maxDescriptorSetInputAttachments	_	4	min
maxVertexInputAttributes	-	16	min
max Vertex Input Bindings	-	16	min
maxVertexInputAttributeOffset	-	2047	min
maxVertexInputBindingStride	-	2048	min
maxVertexOutputComponents	-	64	min
maxTessellationGenerationLevel	0	64	min
maxTessellationPatchSize	0	32	
maxTessellationControlPerVertexInputComponents	0	64	min
maxTessellationControlPerVertexOutputComponents	0	64	min
maxTessellationControlPerPatchOutputComponents	0	120	min
maxTessellationControlTotalOutputComponents	0	2048	min
maxTessellationEvaluationInputComponents	0	64	min
maxTessellationEvaluationOutputComponents	0	64	min min
maxGeometryShaderInvocations	0	32	
	0	64	min
maxGeometryInputComponents	0	64	min
maxGeometryOutputComponents		I .	min
maxGeometryOutputVertices	0	256	min
maxGeometryTotalOutputComponents	0	1024	min
maxFragmentInputComponents	-	64	min

Table 31.2: (continued)

Limit	Unsupported Limit	Supported Limit	Limit Type ¹
maxFragmentOutputAttachments	-	4	min
maxFragmentDualSrcAttachments	0	1	min
maxFragmentCombinedOutputResources	-	4	min
maxComputeSharedMemorySize	-	16384	min
maxComputeWorkGroupCount	-	(65535,65535,65535)	min
maxComputeWorkGroupInvocations	-	128	min
maxComputeWorkGroupSize	-	(128,128,64)	min
subPixelPrecisionBits	-	4	min
subTexelPrecisionBits	-	4	min
mipmapPrecisionBits	-	4	min
maxDrawIndexedIndexValue	2 ²⁴ -1	2 ³² -1	min
maxDrawIndirectCount	1	2 ¹⁶ -1	min
maxSamplerLodBias	-	2	min
maxSamplerAnisotropy	1	16	min
maxViewports	1	16	min
maxViewportDimensions	-	(4096,4096) ³	min
viewportBoundsRange	_	(-8192,8191) ⁴	(max,min)
viewportSubPixelBits		0	min
minMemoryMapAlignment		64	min
minTexelBufferOffsetAlignment		256	max
minUniformBufferOffsetAlignment		256	max
minStorageBufferOffsetAlignment		256	max
minTexelOffset		-8	max
maxTexelOffset		7	min
minTexelGatherOffset	0	-8	max
maxTexelGatherOffset	0	7	min
minInterpolationOffset	0.0	-0.5 5	max
maxInterpolationOffset	0.0	0.5 - (1 ULP) ⁵	min
subPixelInterpolationOffsetBits	0.0	4.5	min
maxFramebufferWidth	-	4096	min
maxFramebufferHeight		4096	min
maxFramebufferLayers	-	256	
framebufferColorSampleCounts	-		min
Transcourrer Colors ample Counts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
S 1 SS . D 1 S 1 S		COUNT_4_BIT)	•
framebufferDepthSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
C 1 CC. C 1C 1 C		COUNT_4_BIT)	•
framebufferStencilSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
		COUNT_4_BIT)	•
framebufferNoAttachmentsSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
		COUNT_4_BIT)	

Table 31.2: (continued)

Limit	Unsupported	Supported Limit	Limit Type ¹
	Limit		
maxColorAttachments	-	4	min
sampledImageColorSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
		COUNT_4_BIT)	
sampledImageIntegerSampleCounts	-	VK_SAMPLE_	min
		COUNT_1_BIT	
sampledImageDepthSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
		COUNT_4_BIT)	
sampledImageStencilSampleCounts	-	(VK_SAMPLE_	min
		COUNT_1_BIT	
		VK_SAMPLE_	
		COUNT_4_BIT)	
storageImageSampleCounts	VK_	(VK_SAMPLE_	min
	SAMPLE_	COUNT_1_BIT	
	COUNT_1_	VK_SAMPLE_	
	BIT	COUNT_4_BIT)	
maxSampleMaskWords	-	1	min
timestampComputeAndGraphics	-	-	implementation
			dependent
timestampPeriod	-	-	duration
maxClipDistances	0	8	min
maxCullDistances	0	8	min
maxCombinedClipAndCullDistances	0	8	min
discreteQueuePriorities	-	2	min
pointSizeRange	(1.0,1.0)	(1.0,64.0 - ULP) ⁶	(max,min)
lineWidthRange	(1.0,1.0)	$(1.0,8.0 - ULP)^7$	(max,min)
pointSizeGranularity	0.0	1.0 6	max, fixed point
1			increment
lineWidthGranularity	0.0	1.0 7	max, fixed point
			increment
strictLines	-	-	implementation
			dependent
standardSampleLocations	-	_	implementation
			dependent
optimalBufferCopyOffsetAlignment	-	_	recommendation
optimalBufferCopyRowPitchAlignment		_	recommendation
nonCoherentAtomSize	-	256	max
nonconcient tomoize		230	IIIux

The **Limit Type** column indicates the limit is either the minimum limit all implementations must support or the maximum limit all implementations must support. For bitmasks a minimum limit is the least bits all implementations must set, but they may have additional bits set beyond this minimum.

2

The maxPerStageResources must be at least the smallest of the following:

- the sum of the maxPerStageDescriptorUniformBuffers, maxPerStageDescriptorStorageBuffers, maxPerStageDescriptorSampledImages, maxPerStageDescriptorStorageImages, maxPerStageDescriptorInputAttachments, maxColorAttachments limits, or
- 128.

It may not be possible to reach this limit in every stage.

3

See maxViewportDimensions for the required relationship to other limits.

4

See viewportBoundsRange for the required relationship to other limits.

5

The values minInterpolationOffset and maxInterpolationOffset describe the closed interval of supported interpolation offsets: [minInterpolationOffset, maxInterpolationOffset]. The ULP is determined by subPixelInterpolationOffsetBits. If subPixelInterpolationOffsetBits is 4, this provides increments of $(1/2^4) = 0.0625$, and thus the range of supported interpolation offsets would be [-0.5, 0.4375].

6

The point size ULP is determined by pointSizeGranularity. If the pointSizeGranularity is 0.125, the range of supported point sizes must be at least [1.0, 63.875].

7

The line width ULP is determined by <code>lineWidthGranularity</code>. If the <code>lineWidthGranularity</code> is 0.0625, the range of supported line widths must be at least [1.0, 7.9375].

31.3 Formats

The features for the set of formats (VkFormat) supported by the implementation are queried individually using the vkGetPhysicalDeviceFormatProperties command.

31.3.1 Format Definition

The available formats are defined by the VkFormat enumeration:

```
typedef enum VkFormat {
   VK_FORMAT_UNDEFINED = 0,
   VK_FORMAT_R4G4_UNORM_PACK8 = 1,
   VK_FORMAT_R4G4B4A4_UNORM_PACK16 = 2,
   VK_FORMAT_B4G4R4A4_UNORM_PACK16 = 3,
   VK_FORMAT_R5G6B5_UNORM_PACK16 = 4,
   VK_FORMAT_B5G6R5_UNORM_PACK16 = 5,
   VK_FORMAT_R5G5B5A1_UNORM_PACK16 = 6,
   VK_FORMAT_B5G5R5A1_UNORM_PACK16 = 7,
   VK_FORMAT_B1G5B5_UNORM_PACK16 = 8,
   VK_FORMAT_A1R5G5B5_UNORM_PACK16 = 8,
   VK_FORMAT_R8_UNORM = 9,
   VK_FORMAT_R8_UNORM = 10,
   VK_FORMAT_R8_USCALED = 11,
```

```
VK_FORMAT_R8_SSCALED = 12,
VK_FORMAT_R8_UINT = 13,
VK_FORMAT_R8_SINT = 14,
VK_FORMAT_R8_SRGB = 15,
VK_FORMAT_R8G8_UNORM = 16,
VK_FORMAT_R8G8_SNORM = 17,
VK_FORMAT_R8G8_USCALED = 18,
VK_FORMAT_R8G8_SSCALED = 19,
VK_FORMAT_R8G8_UINT = 20,
VK_FORMAT_R8G8_SINT = 21,
VK_FORMAT_R8G8_SRGB = 22,
VK_FORMAT_R8G8B8_UNORM = 23,
VK_FORMAT_R8G8B8_SNORM = 24,
VK_FORMAT_R8G8B8_USCALED = 25,
VK_FORMAT_R8G8B8_SSCALED = 26,
VK_FORMAT_R8G8B8_UINT = 27,
VK_FORMAT_R8G8B8_SINT = 28,
VK_FORMAT_R8G8B8_SRGB = 29,
VK\_FORMAT\_B8G8R8\_UNORM = 30,
VK_FORMAT_B8G8R8_SNORM = 31,
VK_FORMAT_B8G8R8_USCALED = 32,
VK_FORMAT_B8G8R8_SSCALED = 33,
VK FORMAT B8G8R8 UINT = 34,
VK_FORMAT_B8G8R8_SINT = 35,
VK_FORMAT_B8G8R8_SRGB = 36,
VK_FORMAT_R8G8B8A8_UNORM = 37,
VK_FORMAT_R8G8B8A8_SNORM = 38,
VK_FORMAT_R8G8B8A8_USCALED = 39,
VK_FORMAT_R8G8B8A8_SSCALED = 40,
VK_FORMAT_R8G8B8A8_UINT = 41,
VK_FORMAT_R8G8B8A8_SINT = 42,
VK_FORMAT_R8G8B8A8_SRGB = 43,
VK_FORMAT_B8G8R8A8_UNORM = 44,
VK_FORMAT_B8G8R8A8_SNORM = 45,
VK_FORMAT_B8G8R8A8_USCALED = 46,
VK FORMAT B8G8R8A8 SSCALED = 47,
VK_FORMAT_B8G8R8A8_UINT = 48,
VK_FORMAT_B8G8R8A8_SINT = 49,
VK_FORMAT_B8G8R8A8_SRGB = 50,
VK_FORMAT_A8B8G8R8_UNORM_PACK32 = 51,
VK_FORMAT_A8B8G8R8_SNORM_PACK32 = 52,
VK_FORMAT_A8B8G8R8_USCALED_PACK32 = 53,
VK_FORMAT_A8B8G8R8_SSCALED_PACK32 = 54,
VK_FORMAT_A8B8G8R8_UINT_PACK32 = 55,
VK_FORMAT_A8B8G8R8_SINT_PACK32 = 56,
VK_FORMAT_A8B8G8R8_SRGB_PACK32 = 57
VK_FORMAT_A2R10G10B10_UNORM_PACK32 = 58,
VK_FORMAT_A2R10G10B10_SNORM_PACK32 = 59,
VK_FORMAT_A2R10G10B10_USCALED_PACK32 = 60,
VK_FORMAT_A2R10G10B10_SSCALED_PACK32 = 61,
VK_FORMAT_A2R10G10B10_UINT_PACK32 = 62,
VK_FORMAT_A2R10G10B10_SINT_PACK32 = 63,
VK_FORMAT_A2B10G10R10_UNORM_PACK32 = 64,
VK_FORMAT_A2B10G10R10_SNORM_PACK32 = 65,
VK_FORMAT_A2B10G10R10_USCALED_PACK32 = 66,
VK_FORMAT_A2B10G10R10_SSCALED_PACK32 = 67,
VK_FORMAT_A2B10G10R10_UINT_PACK32 = 68,
```

```
VK_FORMAT_A2B10G10R10_SINT_PACK32 = 69,
VK_FORMAT_R16_UNORM = 70,
VK_FORMAT_R16_SNORM = 71,
VK_FORMAT_R16_USCALED = 72,
VK_FORMAT_R16_SSCALED = 73,
VK_FORMAT_R16_UINT = 74,
VK_FORMAT_R16_SINT = 75,
VK_FORMAT_R16_SFLOAT = 76,
VK_FORMAT_R16G16_UNORM = 77
VK_FORMAT_R16G16_SNORM = 78,
VK_FORMAT_R16G16_USCALED = 79,
VK_FORMAT_R16G16_SSCALED = 80,
VK_FORMAT_R16G16_UINT = 81,
VK_FORMAT_R16G16_SINT = 82,
VK_FORMAT_R16G16_SFLOAT = 83,
VK_FORMAT_R16G16B16_UNORM = 84,
VK_FORMAT_R16G16B16_SNORM = 85,
VK_FORMAT_R16G16B16_USCALED = 86,
VK_FORMAT_R16G16B16_SSCALED = 87,
VK_FORMAT_R16G16B16_UINT = 88,
VK_FORMAT_R16G16B16_SINT = 89
VK_FORMAT_R16G16B16_SFLOAT = 90,
VK FORMAT R16G16B16A16 UNORM = 91,
VK_FORMAT_R16G16B16A16_SNORM = 92,
VK_FORMAT_R16G16B16A16_USCALED = 93
VK_FORMAT_R16G16B16A16_SSCALED = 94
VK_FORMAT_R16G16B16A16_UINT = 95,
VK_FORMAT_R16G16B16A16_SINT = 96
VK_FORMAT_R16G16B16A16_SFLOAT = 97,
VK FORMAT R32 UINT = 98,
VK_FORMAT_R32_SINT = 99,
VK_FORMAT_R32_SFLOAT = 100,
VK_FORMAT_R32G32_UINT = 101,
VK_FORMAT_R32G32_SINT = 102
VK_FORMAT_R32G32_SFLOAT = 103,
VK FORMAT R32G32B32 UINT = 104,
VK_FORMAT_R32G32B32_SINT = 105,
VK_FORMAT_R32G32B32_SFLOAT = 106,
VK_FORMAT_R32G32B32A32_UINT = 107,
VK_FORMAT_R32G32B32A32_SINT = 108,
VK_FORMAT_R32G32B32A32_SFLOAT = 109
VK_FORMAT_R64_UINT = 110,
VK_FORMAT_R64_SINT = 111,
VK_FORMAT_R64_SFLOAT = 112,
VK_FORMAT_R64G64_UINT = 113,
VK_FORMAT_R64G64_SINT = 114,
VK_FORMAT_R64G64_SFLOAT = 115,
VK_FORMAT_R64G64B64_UINT = 116,
VK_FORMAT_R64G64B64_SINT = 117,
VK_FORMAT_R64G64B64_SFLOAT = 118,
VK_FORMAT_R64G64B64A64_UINT = 119,
VK_FORMAT_R64G64B64A64_SINT = 120,
VK_FORMAT_R64G64B64A64_SFLOAT = 121,
VK_FORMAT_B10G11R11_UFLOAT_PACK32 = 122,
VK_FORMAT_E5B9G9R9_UFLOAT_PACK32 = 123,
VK_FORMAT_D16_UNORM = 124,
VK_FORMAT_X8_D24_UNORM_PACK32 = 125,
```

```
VK_FORMAT_D32_SFLOAT = 126
VK_FORMAT_S8_UINT = 127,
VK_FORMAT_D16_UNORM_S8_UINT = 128,
VK_FORMAT_D24_UNORM_S8_UINT = 129,
VK_FORMAT_D32_SFLOAT_S8_UINT = 130,
VK_FORMAT_BC1_RGB_UNORM_BLOCK = 131,
VK_FORMAT_BC1_RGB_SRGB_BLOCK = 132,
VK_FORMAT_BC1_RGBA_UNORM_BLOCK = 133,
VK_FORMAT_BC1_RGBA_SRGB_BLOCK = 134,
VK FORMAT BC2 UNORM BLOCK = 135,
VK_FORMAT_BC2_SRGB_BLOCK = 136,
VK_FORMAT_BC3_UNORM_BLOCK = 137,
VK_FORMAT_BC3_SRGB_BLOCK = 138,
VK_FORMAT_BC4_UNORM_BLOCK = 139,
VK_FORMAT_BC4_SNORM_BLOCK = 140,
VK_FORMAT_BC5_UNORM_BLOCK = 141,
VK_FORMAT_BC5_SNORM_BLOCK = 142,
VK_FORMAT_BC6H_UFLOAT_BLOCK = 143,
VK_FORMAT_BC6H_SFLOAT_BLOCK = 144,
VK_FORMAT_BC7_UNORM_BLOCK = 145,
VK_FORMAT_BC7_SRGB_BLOCK = 146,
VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK = 147,
VK FORMAT ETC2 R8G8B8 SRGB BLOCK = 148,
VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK = 149,
VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK = 150,
VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK = 151,
VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK = 152,
VK_FORMAT_EAC_R11_UNORM_BLOCK = 153,
VK_FORMAT_EAC_R11_SNORM_BLOCK = 154,
VK_FORMAT_EAC_R11G11_UNORM_BLOCK = 155,
VK_FORMAT_EAC_R11G11_SNORM_BLOCK = 156,
VK_FORMAT_ASTC_4x4_UNORM_BLOCK = 157,
VK_FORMAT_ASTC_4x4_SRGB_BLOCK = 158,
VK_FORMAT_ASTC_5x4_UNORM_BLOCK = 159,
VK_FORMAT_ASTC_5x4_SRGB_BLOCK = 160,
VK_FORMAT_ASTC_5x5_UNORM_BLOCK = 161,
VK_FORMAT_ASTC_5x5_SRGB_BLOCK = 162,
VK_FORMAT_ASTC_6x5_UNORM_BLOCK = 163,
VK_FORMAT_ASTC_6x5_SRGB_BLOCK = 164,
VK_FORMAT_ASTC_6x6_UNORM_BLOCK = 165,
VK_FORMAT_ASTC_6x6_SRGB_BLOCK = 166,
VK_FORMAT_ASTC_8x5_UNORM_BLOCK = 167,
VK_FORMAT_ASTC_8x5_SRGB_BLOCK = 168,
VK_FORMAT_ASTC_8x6_UNORM_BLOCK = 169,
VK_FORMAT_ASTC_8x6_SRGB_BLOCK = 170,
VK_FORMAT_ASTC_8x8_UNORM_BLOCK = 171,
VK_FORMAT_ASTC_8x8_SRGB_BLOCK = 172,
VK_FORMAT_ASTC_10x5_UNORM_BLOCK = 173,
VK_FORMAT_ASTC_10x5_SRGB_BLOCK = 174,
VK_FORMAT_ASTC_10x6_UNORM_BLOCK = 175,
VK_FORMAT_ASTC_10x6_SRGB_BLOCK = 176,
VK_FORMAT_ASTC_10x8_UNORM_BLOCK = 177,
VK_FORMAT_ASTC_10x8_SRGB_BLOCK = 178,
VK_FORMAT_ASTC_10x10_UNORM_BLOCK = 179
VK_FORMAT_ASTC_10x10_SRGB_BLOCK = 180,
VK_FORMAT_ASTC_12x10_UNORM_BLOCK = 181,
VK_FORMAT_ASTC_12x10_SRGB_BLOCK = 182
```

```
VK_FORMAT_ASTC_12x12_UNORM_BLOCK = 183,
   VK_FORMAT_ASTC_12x12_SRGB_BLOCK = 184,
} VkFormat;
```

VK FORMAT UNDEFINED

The format is not specified.

VK_FORMAT_R4G4_UNORM_PACK8

A two-component, 8-bit packed unsigned normalized format that has a 4-bit R component in bits 4..7, and a 4-bit G component in bits 0..3.

VK FORMAT R4G4B4A4 UNORM PACK16

A four-component, 16-bit packed unsigned normalized format that has a 4-bit R component in bits 12..15, a 4-bit G component in bits 8..11, a 4-bit B component in bits 4..7, and a 4-bit A component in bits 0..3.

VK_FORMAT_B4G4R4A4_UNORM_PACK16

A four-component, 16-bit packed unsigned normalized format that has a 4-bit B component in bits 12..15, a 4-bit G component in bits 8..11, a 4-bit R component in bits 4..7, and a 4-bit A component in bits 0..3.

VK_FORMAT_R5G6B5_UNORM_PACK16

A three-component, 16-bit packed unsigned normalized format that has a 5-bit R component in bits 11..15, a 6-bit G component in bits 5..10, and a 5-bit B component in bits 0..4.

VK_FORMAT_B5G6R5_UNORM_PACK16

A three-component, 16-bit packed unsigned normalized format that has a 5-bit B component in bits 11..15, a 6-bit G component in bits 5..10, and a 5-bit R component in bits 0..4.

VK FORMAT R5G5B5A1 UNORM PACK16

A four-component, 16-bit packed unsigned normalized format that has a 5-bit R component in bits 11..15, a 5-bit G component in bits 6..10, a 5-bit B component in bits 1..5, and a 1-bit A component in bit 0.

VK FORMAT B5G5R5A1 UNORM PACK16

A four-component, 16-bit packed unsigned normalized format that has a 5-bit B component in bits 11..15, a 5-bit G component in bits 6..10, a 5-bit R component in bits 1..5, and a 1-bit A component in bit 0.

VK_FORMAT_A1R5G5B5_UNORM_PACK16

A four-component, 16-bit packed unsigned normalized format that has a 1-bit A component in bit 15, a 5-bit R component in bits 10..14, a 5-bit G component in bits 5..9, and a 5-bit B component in bits 0..4.

VK_FORMAT_R8_UNORM

A one-component, 8-bit unsigned normalized format that has a single 8-bit R component.

VK_FORMAT_R8_SNORM

A one-component, 8-bit signed normalized format that has a single 8-bit R component.

VK FORMAT R8 USCALED

A one-component, 8-bit unsigned scaled integer format that has a single 8-bit R component.

VK_FORMAT_R8_SSCALED

A one-component, 8-bit signed scaled integer format that has a single 8-bit R component.

VK FORMAT R8 UINT

A one-component, 8-bit unsigned integer format that has a single 8-bit R component.

VK_FORMAT_R8_SINT

A one-component, 8-bit signed integer format that has a single 8-bit R component.

VK FORMAT R8 SRGB

A one-component, 8-bit unsigned normalized format that has a single 8-bit R component stored with sRGB nonlinear encoding.

VK FORMAT R8G8 UNORM

A two-component, 16-bit unsigned normalized format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK FORMAT R8G8 SNORM

A two-component, 16-bit signed normalized format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK_FORMAT_R8G8_USCALED

A two-component, 16-bit unsigned scaled integer format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK FORMAT R8G8 SSCALED

A two-component, 16-bit signed scaled integer format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK FORMAT R8G8 UINT

A two-component, 16-bit unsigned integer format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK FORMAT R8G8 SINT

A two-component, 16-bit signed integer format that has an 8-bit R component in byte 0, and an 8-bit G component in byte 1.

VK FORMAT R8G8 SRGB

A two-component, 16-bit unsigned normalized format that has an 8-bit R component stored with sRGB nonlinear encoding in byte 0, and an 8-bit G component stored with sRGB nonlinear encoding in byte 1.

VK FORMAT R8G8B8 UNORM

A three-component, 24-bit unsigned normalized format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK FORMAT R8G8B8 SNORM

A three-component, 24-bit signed normalized format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK_FORMAT_R8G8B8_USCALED

A three-component, 24-bit unsigned scaled format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK_FORMAT_R8G8B8_SSCALED

A three-component, 24-bit signed scaled format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK_FORMAT_R8G8B8_UINT

A three-component, 24-bit unsigned integer format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK_FORMAT_R8G8B8_SINT

A three-component, 24-bit signed integer format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, and an 8-bit B component in byte 2.

VK FORMAT R8G8B8 SRGB

A three-component, 24-bit unsigned normalized format that has an 8-bit R component stored with sRGB nonlinear encoding in byte 0, an 8-bit G component stored with sRGB nonlinear encoding in byte 1, and an 8-bit B component stored with sRGB nonlinear encoding in byte 2.

VK FORMAT B8G8R8 UNORM

A three-component, 24-bit unsigned normalized format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK_FORMAT_B8G8R8_SNORM

A three-component, 24-bit signed normalized format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK FORMAT B8G8R8 USCALED

A three-component, 24-bit unsigned scaled format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK FORMAT B8G8R8 SSCALED

A three-component, 24-bit signed scaled format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK_FORMAT_B8G8R8_UINT

A three-component, 24-bit unsigned integer format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK FORMAT B8G8R8 SINT

A three-component, 24-bit signed integer format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, and an 8-bit R component in byte 2.

VK FORMAT B8G8R8 SRGB

A three-component, 24-bit unsigned normalized format that has an 8-bit B component stored with sRGB nonlinear encoding in byte 0, an 8-bit G component stored with sRGB nonlinear encoding in byte 1, and an 8-bit R component stored with sRGB nonlinear encoding in byte 2.

VK FORMAT R8G8B8A8 UNORM

A four-component, 32-bit unsigned normalized format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT R8G8B8A8 SNORM

A four-component, 32-bit signed normalized format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK_FORMAT_R8G8B8A8_USCALED

A four-component, 32-bit unsigned scaled format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK_FORMAT_R8G8B8A8_SSCALED

A four-component, 32-bit signed scaled format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK_FORMAT_R8G8B8A8_UINT

A four-component, 32-bit unsigned integer format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK_FORMAT_R8G8B8A8_SINT

A four-component, 32-bit signed integer format that has an 8-bit R component in byte 0, an 8-bit G component in byte 1, an 8-bit B component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT R8G8B8A8 SRGB

A four-component, 32-bit unsigned normalized format that has an 8-bit R component stored with sRGB nonlinear encoding in byte 0, an 8-bit G component stored with sRGB nonlinear encoding in byte 1, an 8-bit B component stored with sRGB nonlinear encoding in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 UNORM

A four-component, 32-bit unsigned normalized format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 SNORM

A four-component, 32-bit signed normalized format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 USCALED

A four-component, 32-bit unsigned scaled format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 SSCALED

A four-component, 32-bit signed scaled format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK_FORMAT_B8G8R8A8_UINT

A four-component, 32-bit unsigned integer format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 SINT

A four-component, 32-bit signed integer format that has an 8-bit B component in byte 0, an 8-bit G component in byte 1, an 8-bit R component in byte 2, and an 8-bit A component in byte 3.

VK FORMAT B8G8R8A8 SRGB

A four-component, 32-bit unsigned normalized format that has an 8-bit B component stored with sRGB nonlinear encoding in byte 0, an 8-bit G component stored with sRGB nonlinear encoding in byte 1, an 8-bit R component stored with sRGB nonlinear encoding in byte 2, and an 8-bit A component in byte 3.

VK FORMAT A8B8G8R8 UNORM PACK32

A four-component, 32-bit packed unsigned normalized format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK FORMAT A8B8G8R8 SNORM PACK32

A four-component, 32-bit packed signed normalized format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK_FORMAT_A8B8G8R8_USCALED_PACK32

A four-component, 32-bit packed unsigned scaled integer format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK_FORMAT_A8B8G8R8_SSCALED_PACK32

A four-component, 32-bit packed signed scaled integer format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK_FORMAT_A8B8G8R8_UINT_PACK32

A four-component, 32-bit packed unsigned integer format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK_FORMAT_A8B8G8R8_SINT_PACK32

A four-component, 32-bit packed signed integer format that has an 8-bit A component in bits 24..31, an 8-bit B component in bits 16..23, an 8-bit G component in bits 8..15, and an 8-bit R component in bits 0..7.

VK FORMAT A8B8G8R8 SRGB PACK32

A four-component, 32-bit packed unsigned normalized format that has an 8-bit A component in bits 24..31, an 8-bit B component stored with sRGB nonlinear encoding in bits 16..23, an 8-bit G component stored with sRGB nonlinear encoding in bits 8..15, and an 8-bit R component stored with sRGB nonlinear encoding in bits 0..7.

VK FORMAT A2R10G10B10 UNORM PACK32

A four-component, 32-bit packed unsigned normalized format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2R10G10B10 SNORM PACK32

A four-component, 32-bit packed signed normalized format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2R10G10B10 USCALED PACK32

A four-component, 32-bit packed unsigned scaled integer format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2R10G10B10 SSCALED PACK32

A four-component, 32-bit packed signed scaled integer format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2R10G10B10 UINT PACK32

A four-component, 32-bit packed unsigned integer format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2R10G10B10 SINT PACK32

A four-component, 32-bit packed signed integer format that has a 2-bit A component in bits 30..31, a 10-bit R component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit B component in bits 0..9.

VK FORMAT A2B10G10R10 UNORM PACK32

A four-component, 32-bit packed unsigned normalized format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK FORMAT A2B10G10R10 SNORM PACK32

A four-component, 32-bit packed signed normalized format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK FORMAT A2B10G10R10 USCALED PACK32

A four-component, 32-bit packed unsigned scaled integer format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK_FORMAT_A2B10G10R10_SSCALED_PACK32

A four-component, 32-bit packed signed scaled integer format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK_FORMAT_A2B10G10R10_UINT_PACK32

A four-component, 32-bit packed unsigned integer format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK_FORMAT_A2B10G10R10_SINT_PACK32

A four-component, 32-bit packed signed integer format that has a 2-bit A component in bits 30..31, a 10-bit B component in bits 20..29, a 10-bit G component in bits 10..19, and a 10-bit R component in bits 0..9.

VK_FORMAT_R16_UNORM

A one-component, 16-bit unsigned normalized format that has a single 16-bit R component.

VK_FORMAT_R16_SNORM

A one-component, 16-bit signed normalized format that has a single 16-bit R component.

VK FORMAT R16 USCALED

A one-component, 16-bit unsigned scaled integer format that has a single 16-bit R component.

VK FORMAT R16 SSCALED

A one-component, 16-bit signed scaled integer format that has a single 16-bit R component.

VK FORMAT R16 UINT

A one-component, 16-bit unsigned integer format that has a single 16-bit R component.

VK_FORMAT_R16_SINT

A one-component, 16-bit signed integer format that has a single 16-bit R component.

VK FORMAT R16 SFLOAT

A one-component, 16-bit signed floating-point format that has a single 16-bit R component.

VK_FORMAT_R16G16_UNORM

A two-component, 32-bit unsigned normalized format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16_SNORM

A two-component, 32-bit signed normalized format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK FORMAT R16G16 USCALED

A two-component, 32-bit unsigned scaled integer format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16_SSCALED

A two-component, 32-bit signed scaled integer format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16_UINT

A two-component, 32-bit unsigned integer format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16_SINT

A two-component, 32-bit signed integer format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16_SFLOAT

A two-component, 32-bit signed floating-point format that has a 16-bit R component in bytes 0..1, and a 16-bit G component in bytes 2..3.

VK_FORMAT_R16G16B16_UNORM

A three-component, 48-bit unsigned normalized format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK_FORMAT_R16G16B16_SNORM

A three-component, 48-bit signed normalized format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK_FORMAT_R16G16B16_USCALED

A three-component, 48-bit unsigned scaled integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK_FORMAT_R16G16B16_SSCALED

A three-component, 48-bit signed scaled integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK FORMAT R16G16B16 UINT

A three-component, 48-bit unsigned integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK FORMAT R16G16B16 SINT

A three-component, 48-bit signed integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK FORMAT R16G16B16 SFLOAT

A three-component, 48-bit signed floating-point format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, and a 16-bit B component in bytes 4..5.

VK_FORMAT_R16G16B16A16_UNORM

A four-component, 64-bit unsigned normalized format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 SNORM

A four-component, 64-bit signed normalized format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 USCALED

A four-component, 64-bit unsigned scaled integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 SSCALED

A four-component, 64-bit signed scaled integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 UINT

A four-component, 64-bit unsigned integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 SINT

A four-component, 64-bit signed integer format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK FORMAT R16G16B16A16 SFLOAT

A four-component, 64-bit signed floating-point format that has a 16-bit R component in bytes 0..1, a 16-bit G component in bytes 2..3, a 16-bit B component in bytes 4..5, and a 16-bit A component in bytes 6..7.

VK_FORMAT_R32_UINT

A one-component, 32-bit unsigned integer format that has a single 32-bit R component.

VK_FORMAT_R32_SINT

A one-component, 32-bit signed integer format that has a single 32-bit R component.

VK FORMAT R32 SFLOAT

A one-component, 32-bit signed floating-point format that has a single 32-bit R component.

VK_FORMAT_R32G32_UINT

A two-component, 64-bit unsigned integer format that has a 32-bit R component in bytes 0..3, and a 32-bit G component in bytes 4..7.

VK_FORMAT_R32G32_SINT

A two-component, 64-bit signed integer format that has a 32-bit R component in bytes 0..3, and a 32-bit G component in bytes 4..7.

VK FORMAT R32G32 SFLOAT

A two-component, 64-bit signed floating-point format that has a 32-bit R component in bytes 0..3, and a 32-bit G component in bytes 4..7.

VK FORMAT R32G32B32 UINT

A three-component, 96-bit unsigned integer format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, and a 32-bit B component in bytes 8..11.

VK FORMAT R32G32B32 SINT

A three-component, 96-bit signed integer format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, and a 32-bit B component in bytes 8..11.

VK_FORMAT_R32G32B32_SFLOAT

A three-component, 96-bit signed floating-point format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, and a 32-bit B component in bytes 8..11.

VK FORMAT R32G32B32A32 UINT

A four-component, 128-bit unsigned integer format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, a 32-bit B component in bytes 8..11, and a 32-bit A component in bytes 12..15.

VK FORMAT R32G32B32A32 SINT

A four-component, 128-bit signed integer format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, a 32-bit B component in bytes 8..11, and a 32-bit A component in bytes 12..15.

VK FORMAT R32G32B32A32 SFLOAT

A four-component, 128-bit signed floating-point format that has a 32-bit R component in bytes 0..3, a 32-bit G component in bytes 4..7, a 32-bit B component in bytes 8..11, and a 32-bit A component in bytes 12..15.

VK FORMAT R64 UINT

A one-component, 64-bit unsigned integer format that has a single 64-bit R component.

VK FORMAT R64 SINT

A one-component, 64-bit signed integer format that has a single 64-bit R component.

VK FORMAT R64 SFLOAT

A one-component, 64-bit signed floating-point format that has a single 64-bit R component.

VK_FORMAT_R64G64_UINT

A two-component, 128-bit unsigned integer format that has a 64-bit R component in bytes 0..7, and a 64-bit G component in bytes 8..15.

VK FORMAT R64G64 SINT

A two-component, 128-bit signed integer format that has a 64-bit R component in bytes 0..7, and a 64-bit G component in bytes 8..15.

VK_FORMAT_R64G64_SFLOAT

A two-component, 128-bit signed floating-point format that has a 64-bit R component in bytes 0..7, and a 64-bit G component in bytes 8..15.

VK_FORMAT_R64G64B64_UINT

A three-component, 192-bit unsigned integer format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, and a 64-bit B component in bytes 16..23.

VK_FORMAT_R64G64B64_SINT

A three-component, 192-bit signed integer format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, and a 64-bit B component in bytes 16..23.

VK FORMAT R64G64B64 SFLOAT

A three-component, 192-bit signed floating-point format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, and a 64-bit B component in bytes 16..23.

VK_FORMAT_R64G64B64A64_UINT

A four-component, 256-bit unsigned integer format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, a 64-bit B component in bytes 16..23, and a 64-bit A component in bytes 24..31.

VK FORMAT R64G64B64A64 SINT

A four-component, 256-bit signed integer format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, a 64-bit B component in bytes 16..23, and a 64-bit A component in bytes 24..31.

VK_FORMAT_R64G64B64A64_SFLOAT

A four-component, 256-bit signed floating-point format that has a 64-bit R component in bytes 0..7, a 64-bit G component in bytes 8..15, a 64-bit B component in bytes 16..23, and a 64-bit A component in bytes 24..31.

VK_FORMAT_B10G11R11_UFLOAT_PACK32

A three-component, 32-bit packed unsigned floating-point format that has a 10-bit B component in bits 22..31, an 11-bit G component in bits 11..21, an 11-bit R component in bits 0..10. See Section 2.7.4 and Section 2.7.3.

VK FORMAT E5B9G9R9 UFLOAT PACK32

A three-component, 32-bit packed unsigned floating-point format that has a 5-bit shared exponent in bits 27..31, a 9-bit B component mantissa in bits 18..26, a 9-bit G component mantissa in bits 9..17, and a 9-bit R component mantissa in bits 0..8.

VK FORMAT D16 UNORM

A one-component, 16-bit unsigned normalized format that has a single 16-bit depth component.

VK FORMAT X8 D24 UNORM PACK32

A two-component, 32-bit format that has 24 unsigned normalized bits in the depth component and, optionally, 8 bits that are unused.

VK_FORMAT_D32_SFLOAT

A one-component, 32-bit signed floating-point format that has 32-bits in the depth component.

VK FORMAT S8 UINT

A one-component, 8-bit unsigned integer format that has 8-bits in the stencil component.

VK_FORMAT_D16_UNORM_S8_UINT

A two-component, 24-bit format that has 16 unsigned normalized bits in the depth component and 8 unsigned integer bits in the stencil component.

VK FORMAT D24 UNORM S8 UINT

A two-component, 32-bit packed format that has 8 unsigned integer bits in the stencil component, and 24 unsigned normalized bits in the depth component.

VK_FORMAT_D32_SFLOAT_S8_UINT

A two-component format that has 32 signed float bits in the depth component and 8 unsigned integer bits in the stencil component. There are optionally 24-bits that are unused.

VK_FORMAT_BC1_RGB_UNORM_BLOCK

A three-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data. This format has no alpha and is considered opaque.

VK FORMAT BC1 RGB SRGB BLOCK

A three-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data with sRGB nonlinear encoding. This format has no alpha and is considered opaque.

VK FORMAT BC1 RGBA UNORM BLOCK

A four-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data, and provides 1 bit of alpha.

VK FORMAT BC1 RGBA SRGB BLOCK

A four-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data with sRGB nonlinear encoding, and provides 1 bit of alpha.

VK_FORMAT_BC2_UNORM_BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values.

VK FORMAT BC2 SRGB BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values with sRGB nonlinear encoding.

VK_FORMAT_BC3_UNORM_BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values.

VK FORMAT BC3 SRGB BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values with sRGB nonlinear encoding.

VK FORMAT BC4 UNORM BLOCK

A one-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized red texel data.

VK_FORMAT_BC4_SNORM_BLOCK

A one-component, block-compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of signed normalized red texel data.

VK_FORMAT_BC5_UNORM_BLOCK

A two-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RG texel data with the first 64 bits encoding red values followed by 64 bits encoding green values.

VK FORMAT BC5 SNORM BLOCK

A two-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of signed normalized RG texel data with the first 64 bits encoding red values followed by 64 bits encoding green values.

VK_FORMAT_BC6H_UFLOAT_BLOCK

A three-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned floating-point RGB texel data.

VK_FORMAT_BC6H_SFLOAT_BLOCK

A three-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of signed floating-point RGB texel data.

VK_FORMAT_BC7_UNORM_BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data.

VK FORMAT BC7 SRGB BLOCK

A four-component, block-compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ETC2 R8G8B8 UNORM BLOCK

A three-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data. This format has no alpha and is considered opaque.

VK FORMAT ETC2 R8G8B8 SRGB BLOCK

A three-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data with sRGB nonlinear encoding. This format has no alpha and is considered opaque.

VK FORMAT ETC2 R8G8B8A1 UNORM BLOCK

A four-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data, and provides 1 bit of alpha.

VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK

A four-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGB texel data with sRGB nonlinear encoding, and provides 1 bit of alpha.

VK FORMAT ETC2 R8G8B8A8 UNORM BLOCK

A four-component, ETC2 compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values.

VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK

A four-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with the first 64 bits encoding alpha values followed by 64 bits encoding RGB values with sRGB nonlinear encoding applied.

VK_FORMAT_EAC_R11_UNORM_BLOCK

A one-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized red texel data.

VK_FORMAT_EAC_R11_SNORM_BLOCK

A one-component, ETC2 compressed format where each 64-bit compressed texel block encodes a 4x4 rectangle of signed normalized red texel data.

VK_FORMAT_EAC_R11G11_UNORM_BLOCK

A two-component, ETC2 compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RG texel data with the first 64 bits encoding red values followed by 64 bits encoding green values.

VK_FORMAT_EAC_R11G11_SNORM_BLOCK

A two-component, ETC2 compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of signed normalized RG texel data with the first 64 bits encoding red values followed by 64 bits encoding green values.

VK_FORMAT_ASTC_4x4_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_4x4_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 4x4 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ASTC 5x4 UNORM BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 5x4 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_5x4_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 5x4 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ASTC 5x5 UNORM BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 5x5 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_5x5_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 5x5 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ASTC 6x5 UNORM BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 6x5 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_6x5_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 6x5 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_6x6_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 6x6 rectangle of unsigned normalized RGBA texel data.

VK FORMAT ASTC 6x6 SRGB BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 6x6 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_8x5_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x5 rectangle of unsigned normalized RGBA texel data.

VK FORMAT ASTC 8x5 SRGB BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x5 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_8x6_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x6 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_8x6_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x6 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_8x8_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x8 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_8x8_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes an 8x8 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ASTC 10x5 UNORM BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x5 rectangle of unsigned normalized RGBA texel data.

VK FORMAT ASTC 10x5 SRGB BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x5 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_10x6_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x6 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_10x6_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x6 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_10x8_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x8 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_10x8_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x8 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_10x10_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x10 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_10x10_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 10x10 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK_FORMAT_ASTC_12x10_UNORM_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 12x10 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_12x10_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 12x10 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

VK FORMAT ASTC 12x12 UNORM BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 12x12 rectangle of unsigned normalized RGBA texel data.

VK_FORMAT_ASTC_12x12_SRGB_BLOCK

A four-component, ASTC compressed format where each 128-bit compressed texel block encodes a 12x12 rectangle of unsigned normalized RGBA texel data with sRGB nonlinear encoding applied to the RGB components.

31.3.1.1 Packed Formats

For the purposes of address alignment when accessing buffer memory containing vertex attribute or texel data, the following formats are considered *packed* - whole texels or attributes are stored in a single data element, rather than individual components occupying a single data element:

• Packed into 8-bit data types:

- VK_FORMAT_R4G4_UNORM_PACK8
- Packed into 16-bit data types:
 - VK_FORMAT_R4G4B4A4_UNORM_PACK16
 - VK_FORMAT_B4G4R4A4_UNORM_PACK16
 - VK_FORMAT_R5G6B5_UNORM_PACK16
 - VK_FORMAT_B5G6R5_UNORM_PACK16
 - VK_FORMAT_R5G5B5A1_UNORM_PACK16
 - VK FORMAT B5G5R5A1 UNORM PACK16
 - VK_FORMAT_A1R5G5B5_UNORM_PACK16

• Packed into 32-bit data types:

- VK_FORMAT_A8B8G8R8_UNORM_PACK32
- VK_FORMAT_A8B8G8R8_SNORM_PACK32
- VK_FORMAT_A8B8G8R8_USCALED_PACK32
- VK_FORMAT_A8B8G8R8_SSCALED_PACK32
- VK FORMAT A8B8G8R8 UINT PACK32
- VK_FORMAT_A8B8G8R8_SINT_PACK32
- VK FORMAT A8B8G8R8 SRGB PACK32
- VK_FORMAT_A2R10G10B10_UNORM_PACK32
- VK_FORMAT_A2R10G10B10_SNORM_PACK32
- VK_FORMAT_A2R10G10B10_USCALED_PACK32
- VK_FORMAT_A2R10G10B10_SSCALED_PACK32
- VK_FORMAT_A2R10G10B10_UINT_PACK32
- VK FORMAT A2R10G10B10 SINT PACK32
- VK_FORMAT_A2B10G10R10_UNORM_PACK32
- VK_FORMAT_A2B10G10R10_SNORM_PACK32
- VK_FORMAT_A2B10G10R10_USCALED_PACK32
- VK_FORMAT_A2B10G10R10_SSCALED_PACK32
- VK FORMAT A2B10G10R10 UINT PACK32
- VK_FORMAT_A2B10G10R10_SINT_PACK32
- VK_FORMAT_B10G11R11_UFLOAT_PACK32
- VK_FORMAT_E5B9G9R9_UFLOAT_PACK32
- VK_FORMAT_X8_D24_UNORM_PACK32

31.3.1.2 Identification of Formats

A "format" is represented by a single enum value. The name of a format is usually built up by using the following pattern:

etext:VK_FORMAT_{component-format|compression-scheme}_{numeric-format}

The component-format specifies either the size of the R, G, B, and A components (if they are present) in the case of a color format, or the size of the depth (D) and stencil (S) components (if they are present) in the case of a depth/stencil format (see below). An X indicates a component that is unused, but may be present for padding.

Numeric format Description The components are unsigned normalized values in the range [0,1]UNORM SNORM The components are signed normalized values in the range [-1,1] USCALED The components are unsigned integer values that get converted to floating-point in the range [0,2ⁿ-1] The components are signed integer values that get converted to floating-point in the SSCALED range $[-2^{n-1}, 2^{n-1}-1]$ The components are unsigned integer values in the range $[0,2^{n}-1]$ UTNT The components are signed integer values in the range $[-2^{n-1}, 2^{n-1}-1]$ SINT UFLOAT The components are unsigned floating-point numbers (used by packed, shared exponent, and some compressed formats) SFLOAT The components are signed floating-point numbers The R, G, and B components are unsigned normalized values that represent values using SRGB sRGB nonlinear encoding, while the A component (if one exists) is a regular unsigned normalized value

Table 31.3: Interpretation of Numeric Format

The suffix _PACKnn indicates that the format is packed into an underlying type with nn bits.

The suffix _BLOCK indicates that the format is a block-compressed format, with the representation of multiple pixels encoded interdependently within a region.

Compression	Description
scheme	
BC	Block Compression. See Section B.1.
ETC2	Ericsson Texture Compression. See Section B.2.
EAC	ETC2 Alpha Compression. See Section B.2.
ASTC	Adaptive Scalable Texture Compression (LDR Profile). See Section B.3.

Table 31.4: Interpretation of Compression Scheme

31.3.1.3 Representation

Color formats must be represented in memory in exactly the form indicated by the format's name. This means that promoting one format to another with more bits per component and/or additional components must not occur for color formats. Depth/stencil formats have more relaxed requirements as discussed below.

The representation of non-packed formats is that the first component specified in the name of the format is in the lowest memory addresses and the last component specified is in the highest memory addresses. See Byte mappings for non-packed/compressed color formats. The in-memory ordering of bytes within a component is determined by the host endianness.

Table 31.5: Byte mappings for non-packed/compressed color formats

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	← Byte						
R						'			,							VK_FORMAT_R8_*						
R	G															VK_FORMAT_R8G8_						
R	G	В												VK_FORMAT_R8G8B8_								
В	G	R												VK_FORMAT_B8G8R8_								
R	G	В	Α													VK_FORMAT_R8G8B8A8_*						
В	G	R	Α													VK_FORMAT_B8G8R8A8_*						
F	3															VK_FORMAT_R16_*						
F	3	(J													VK_FORMAT_R16G16_						
F	3	(J	I	В											VK_FORMAT_R16G16B16_*						
F	3	(J	I	В	1	4									VK_FORMAT_R16G16B16A16_*						
	I	3														VK_FORMAT_R32_*						
	I	3				G										VK_FORMAT_R32G32_*						
	I	3			-	G			I	3						VK_FORMAT_R32G32B32_*						
	I	ξ			-	G			I	3			A VK_FORMAT_R32G32B32A32_									
			F	₹							VK_FORMAT_R64											
			F	-												VK_FORMAT_R64G64_*						
	VK_FORMAT_R64G64B64_* as VK_FORMAT_R64G64_* but with B in bytes 16-23																					
	VK_FORMAT_R64G64B64A64_* as VK_FORMAT_R64G64B64_* but with A in bytes 24-31																					

Packed formats store multiple components within one underlying type. The bit representation is that the first component specified in the name of the format is in the most-significant bits and the last component specified is in the least-significant bits of the underlying type. The in-memory ordering of bytes comprising the underlying type is determined by the host endianness.

Table 31.6: Bit mappings for packed 8-bit formats

$\textbf{Bit} \rightarrow$	7	6	5	4	3	2	1	0
VK_FORMAT_R4G4_UNORM_PACK8	R_3	R_2	R_1	R_0	G_3	G_2	G_1	G_0

Table 31.7: Bit mappings for packed 16-bit formats

$\mathbf{Bit} ightarrow$	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
VK_FORMAT_R4G4B4A4_	R ₃	R_2	R_1	R_0	G ₃	G_2	G_1	G_0	B_3	B_2	B_1	B_0	A_3	A_2	A_1	A_0
UNORM_PACK16																
VK_FORMAT_B4G4R4A4_	B ₃	B_2	B_1	B_0	G ₃	G_2	G_1	G_0	R_3	R_2	R_1	R_0	A_3	A_2	A_1	A_0
UNORM_PACK16																
VK_FORMAT_R5G6B5_UNORM_	R ₄	R_3	R ₂	R_1	R_0	G_5	G_4	G ₃	G_2	G_1	G_0	B ₄	B_3	B_2	B_1	B_0
PACK16																
VK_FORMAT_B5G6R5_UNORM_	B ₄	B ₃	B_2	B_1	B_0	G_5	G_4	G ₃	G_2	G_1	G_0	R ₄	R_3	R_2	R_1	R_0
PACK16																
VK_FORMAT_R5G5B5A1_	R ₄	R ₃	R ₂	R_1	R_0	G_4	G ₃	G_2	G_1	G_0	B_4	B ₃	B_2	B_1	B_0	A_0
UNORM_PACK16																

 $Bit \rightarrow$ 15 13 12 11 10 9 4 3 2 1 0 14 8 7 6 5 $\overline{\mathbf{G}}_3$ $\overline{R_1}$ \overline{A}_0 VK FORMAT B5G5R5A1 B_4 B_3 B_2 B_1 B_0 G_4 G_2 G_1 G_0 R_4 R_3 R_2 R_0 UNORM_PACK16 \overline{A}_0 \overline{R}_3 \overline{R}_2 \overline{G}_3 \overline{G}_0 VK FORMAT A1R5G5B5 R_4 R_1 R_0 G_4 G_2 G_1 B_4 B_3 B_2 B_1 B_0 UNORM_PACK16

Table 31.7: (continued)

Table 31.8: Bit mappings for packed 32-bit formats

31 30 29	28 27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
							VI	K_F	ORI	TAN	_A8	3B8	G8I	R8_	*_F	AC	K32											
A ₇ A ₆ A ₅	A ₄ A ₃	A_2	A_1	A_0	B ₇	B ₆	B ₅	B_4	B_3	B_2	B_1	B_0	G ₇	G_6	G_5	G_4	G ₃	$\overline{G_2}$	G_1	G_0	R_7	R ₆	R_5	R_4	R_3	R_2	R_1	$\overline{R_0}$
						7	VK_	FOI	RMA	T_2	12R	100	310	B10)*	_P#	CK:	32										
A ₁ A ₀ R ₉	R ₈ R ₇	R ₆	R_5	R_4	R ₃	R_2	R_1	R_0	G ₉	G_8	G_7	G_6	G_5	G_4	G_3	G_2	G_1	G_0	B ₉	B_8	B ₇	B ₆	B ₅	B_4	B_3	B_2	B_1	$\overline{\mathrm{B}_{\mathrm{0}}}$
						7	VK_	FOI	RMA	T_2	12B	100	310	R10)_*	_P#	CK.	32										
A ₁ A ₀ B ₉	B ₈ B ₇	B ₆	B ₅	B_4	B_3	B_2	B_1	B_0	G ₉	G_8	G_7	G_6	G_5	G_4	G_3	G_2	G_1	G_0	R ₉	R_8	R_7	R ₆	R_5	R_4	R_3	R_2	R_1	R_0
						VK	_F	ORM	AT_	_B1	0G1	L1R	11_	UF	LOP	\ T _:	PAC	K3	2									
B ₉ B ₈ B ₇	B ₆ B ₅	B_4	B_3	B_2	B_1	B_0	G_{10}	G_9	G_8	G_7	G_6	G_5	G_4	G_3	G_2	G_1	G_0	R ₁₀	R ₉	R_8	R_7	R_6	R_5	R_4	R_3	R_2	R_1	R_0
						V	K_F	ORI	TAN	_E5	5B9	G9I	ર9_	UFI	OA	T_F	ACI	K32										
E ₄ E ₃ E ₂	E_1 E_0	B_8	B ₇	B ₆	B ₅	B_4	B_3	B_2	B_1	B_0	G_8	G_7	G_6	G_5	G_4	G_3	G_2	G_1	G_0	R ₈	R_7	R ₆	R_5	R_4	R_3	R_2	R_1	R_0
							VK_	_FO	RM/	\T_	X8_	_D2	4 _t	JNO	RM_	PA	CK3	2										
X ₇ X ₆ X ₅	$X_4 X_3$	X_2	X_1	X_0	D ₂₃	$\overline{\mathbf{D}_{22}}$	$_{2}D_{2}$	D ₂₀	$_{0}\overline{\mathrm{D}_{19}}$	$_{9}\overline{\mathrm{D}_{18}}$	$_{3}D_{12}$	$_{7}D_{10}$	$_{6}D_{1}$	5D14	1D13	$\overline{D_{12}}$	$_{2}D_{11}$	D_{10}	D9	D_8	D_7	D_6	D_5	D_4	D_3	D_2	D_1	$\overline{\mathrm{D}_{\mathrm{0}}}$

31.3.1.4 Depth/Stencil Formats

Depth/stencil formats are considered opaque and need not be stored in the exact number of bits per texel or component ordering indicated by the format enum. However, implementations must not substitute a different depth or stencil precision than that described in the format (e.g. D16 must not be implemented as D24 or D32).

31.3.1.5 Format Compatibility Classes

Uncompressed color formats are *compatible* with each other if they occupy the same number of bits per data element. Compressed color formats are compatible with each other if the only difference between them is the numerical type of the uncompressed pixels (e.g. signed vs. unsigned, or SRGB vs. UNORM encoding). Each depth/stencil format is only compatible with itself. In the following table, all the formats in the same row are compatible.

Table 31.9: Compatible formats

Class	Formats
8-bit	VK_FORMAT_R4G4_UNORM_PACK8,
	VK_FORMAT_R8_UNORM,
	VK_FORMAT_R8_SNORM,
	VK_FORMAT_R8_USCALED,
	VK_FORMAT_R8_SSCALED,
	VK_FORMAT_R8_UINT,
	VK_FORMAT_R8_SINT,
	VK_FORMAT_R8_SRGB
16-bit	VK_FORMAT_R4G4B4A4_UNORM_PACK16,
	VK_FORMAT_B4G4R4A4_UNORM_PACK16,
	VK_FORMAT_R5G6B5_UNORM_PACK16,
	VK_FORMAT_B5G6R5_UNORM_PACK16,
	VK_FORMAT_R5G5B5A1_UNORM_PACK16,
	VK_FORMAT_B5G5R5A1_UNORM_PACK16,
	VK_FORMAT_A1R5G5B5_UNORM_PACK16,
	VK_FORMAT_R8G8_UNORM,
	VK_FORMAT_R8G8_SNORM,
	VK_FORMAT_R8G8_USCALED,
	VK_FORMAT_R8G8_SSCALED,
	VK_FORMAT_R8G8_UINT,
	VK_FORMAT_R8G8_SINT,
	VK_FORMAT_R8G8_SRGB,
	VK_FORMAT_R16_UNORM,
	VK_FORMAT_R16_SNORM,
	VK_FORMAT_R16_USCALED,
	VK_FORMAT_R16_SSCALED,
	VK_FORMAT_R16_UINT,
	VK_FORMAT_R16_SINT,
	VK_FORMAT_R16_SFLOAT
24-bit	VK_FORMAT_R8G8B8_UNORM,
	VK_FORMAT_R8G8B8_SNORM,
	VK_FORMAT_R8G8B8_USCALED,
	VK_FORMAT_R8G8B8_SSCALED,
	VK_FORMAT_R8G8B8_UINT,
	VK_FORMAT_R8G8B8_SINT,
	VK_FORMAT_R8G8B8_SRGB,
	VK_FORMAT_B8G8R8_UNORM,
	VK_FORMAT_B8G8R8_SNORM,
	VK_FORMAT_B8G8R8_USCALED,
	VK_FORMAT_B8G8R8_SSCALED,
	VK_FORMAT_B8G8R8_UINT,
	VK_FORMAT_B8G8R8_SINT,
	VK_FORMAT_B8G8R8_SRGB

Table 31.9: (continued)

Class	Formats
32-bit	VK_FORMAT_R8G8B8A8_UNORM,
	VK_FORMAT_R8G8B8A8_SNORM,
	VK_FORMAT_R8G8B8A8_USCALED,
	VK_FORMAT_R8G8B8A8_SSCALED,
	VK_FORMAT_R8G8B8A8_UINT,
	VK_FORMAT_R8G8B8A8_SINT,
	VK_FORMAT_R8G8B8A8_SRGB,
	VK_FORMAT_B8G8R8A8_UNORM,
	VK_FORMAT_B8G8R8A8_SNORM,
	VK_FORMAT_B8G8R8A8_USCALED,
	VK_FORMAT_B8G8R8A8_SSCALED,
	VK_FORMAT_B8G8R8A8_UINT,
	VK_FORMAT_B8G8R8A8_SINT,
	VK_FORMAT_B8G8R8A8_SRGB,
	VK_FORMAT_A8B8G8R8_UNORM_PACK32,
	VK_FORMAT_A8B8G8R8_SNORM_PACK32,
	VK_FORMAT_A8B8G8R8_USCALED_PACK32,
	VK_FORMAT_A8B8G8R8_SSCALED_PACK32,
	VK_FORMAT_A8B8G8R8_UINT_PACK32,
	VK_FORMAT_A8B8G8R8_SINT_PACK32,
	VK_FORMAT_A8B8G8R8_SRGB_PACK32,
	VK_FORMAT_A2R10G10B10_UNORM_PACK32,
	VK_FORMAT_A2R10G10B10_SNORM_PACK32,
	VK_FORMAT_A2R10G10B10_USCALED_PACK32,
	VK_FORMAT_A2R10G10B10_SSCALED_PACK32,
	VK_FORMAT_A2R10G10B10_UINT_PACK32,
	VK_FORMAT_A2R10G10B10_SINT_PACK32,
	VK_FORMAT_A2B10G10R10_UNORM_PACK32,
	VK_FORMAT_A2B10G10R10_SNORM_PACK32,
	VK_FORMAT_A2B10G10R10_USCALED_PACK32,
	VK_FORMAT_A2B10G10R10_SSCALED_PACK32,
	VK_FORMAT_A2B10G10R10_UINT_PACK32,
	VK_FORMAT_A2B10G10R10_SINT_PACK32,
	VK_FORMAT_R16G16_UNORM,
	VK_FORMAT_R16G16_SNORM,
	VK_FORMAT_R16G16_USCALED,
	VK_FORMAT_R16G16_SSCALED,
	VK_FORMAT_R16G16_UINT,
	VK_FORMAT_R16G16_SINT,
	VK_FORMAT_R16G16_SFLOAT,
	VK_FORMAT_R32_UINT,
	VK_FORMAT_R32_SINT,
	VK_FORMAT_R32_SFLOAT,
	VK_FORMAT_B10G11R11_UFLOAT_PACK32,
	VK_FORMAT_E5B9G9R9_UFLOAT_PACK32

Table 31.9: (continued)

Class	Formats
48-bit	VK_FORMAT_R16G16B16_UNORM,
	VK_FORMAT_R16G16B16_SNORM,
	VK_FORMAT_R16G16B16_USCALED,
	VK_FORMAT_R16G16B16_SSCALED,
	VK_FORMAT_R16G16B16_UINT,
	VK_FORMAT_R16G16B16_SINT,
	VK_FORMAT_R16G16B16_SFLOAT
64-bit	VK_FORMAT_R16G16B16A16_UNORM,
	VK_FORMAT_R16G16B16A16_SNORM,
	VK_FORMAT_R16G16B16A16_USCALED,
	VK_FORMAT_R16G16B16A16_SSCALED,
	VK_FORMAT_R16G16B16A16_UINT,
	VK_FORMAT_R16G16B16A16_SINT,
	VK_FORMAT_R16G16B16A16_SFLOAT,
	VK_FORMAT_R32G32_UINT,
	VK_FORMAT_R32G32_SINT,
	VK_FORMAT_R32G32_SFLOAT,
	VK_FORMAT_R64_UINT,
	VK_FORMAT_R64_SINT,
	VK_FORMAT_R64_SFLOAT
96-bit	VK_FORMAT_R32G32B32_UINT,
	VK_FORMAT_R32G32B32_SINT,
	VK_FORMAT_R32G32B32_SFLOAT
128-bit	VK_FORMAT_R32G32B32A32_UINT,
	VK_FORMAT_R32G32B32A32_SINT,
	VK_FORMAT_R32G32B32A32_SFLOAT,
	VK_FORMAT_R64G64_UINT,
	VK_FORMAT_R64G64_SINT,
	VK_FORMAT_R64G64_SFLOAT
192-bit	VK_FORMAT_R64G64B64_UINT,
	VK_FORMAT_R64G64B64_SINT,
	VK_FORMAT_R64G64B64_SFLOAT
256-bit	VK_FORMAT_R64G64B64A64_UINT,
	VK_FORMAT_R64G64B64A64_SINT,
P.G1 P.GP	VK_FORMAT_R64G64B64A64_SFLOAT
BC1_RGB	VK_FORMAT_BC1_RGB_UNORM_BLOCK,
DC1 DCD4	VK_FORMAT_BC1_RGB_SRGB_BLOCK
BC1_RGBA	VK_FORMAT_BC1_RGBA_UNORM_BLOCK,
DC2	VK_FORMAT_BC1_RGBA_SRGB_BLOCK
BC2	VK_FORMAT_BC2_UNORM_BLOCK,
DC2	VK_FORMAT_BC2_SRGB_BLOCK
BC3	VK_FORMAT_BC3_UNORM_BLOCK,
DC4	VK_FORMAT_BC3_SRGB_BLOCK
BC4	VK_FORMAT_BC4_UNORM_BLOCK,
DC5	VK_FORMAT_BC4_SNORM_BLOCK
BC5	VK_FORMAT_BC5_UNORM_BLOCK,
DCGII	VK_FORMAT_BC5_SNORM_BLOCK
ВС6Н	VK_FORMAT_BC6H_UFLOAT_BLOCK,
	VK_FORMAT_BC6H_SFLOAT_BLOCK

Table 31.9: (continued)

Class	Formats
BC7	VK_FORMAT_BC7_UNORM_BLOCK,
	VK_FORMAT_BC7_SRGB_BLOCK
ETC2_RGB	VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK,
	VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK
ETC2 RGBA	VK FORMAT ETC2 R8G8B8A1 UNORM BLOCK,
LTC2_KGB/Y	VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK
ETC2 EAC RGBA	VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK,
LTCZ_L/IC_RGB/Y	VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK
EAC_R	VK_FORMAT_EAC_R11_UNORM_BLOCK,
LAC_K	VK_FORMAT_EAC_R11_SNORM_BLOCK VK_FORMAT_EAC_R11_SNORM_BLOCK
EAC RG	VK_FORMAT_EAC_R11G11_UNORM_BLOCK,
LAC_KO	VK_FORMAT_EAC_R11G11_ONORM_BLOCK, VK_FORMAT_EAC_R11G11_SNORM_BLOCK
ASTC 4x4	
ASTC_4X4	VK_FORMAT_ASTC_4x4_UNORM_BLOCK,
ASTC 5x4	VK_FORMAT_ASTC_4x4_SRGB_BLOCK
ASTC_3X4	VK_FORMAT_ASTC_5x4_UNORM_BLOCK,
ASTC 5x5	VK_FORMAT_ASTC_5x4_SRGB_BLOCK
AS1C_5x5	VK_FORMAT_ASTC_5x5_UNORM_BLOCK,
ASTC 6x5	VK_FORMAT_ASTC_5x5_SRGB_BLOCK
ASIC_6x5	VK_FORMAT_ASTC_6x5_UNORM_BLOCK,
A CITICA (VK_FORMAT_ASTC_6x5_SRGB_BLOCK
ASTC_6x6	VK_FORMAT_ASTC_6x6_UNORM_BLOCK,
1 GT G 0 7	VK_FORMAT_ASTC_6x6_SRGB_BLOCK
ASTC_8x5	VK_FORMAT_ASTC_8x5_UNORM_BLOCK,
	VK_FORMAT_ASTC_8x5_SRGB_BLOCK
ASTC_8x6	VK_FORMAT_ASTC_8x6_UNORM_BLOCK,
	VK_FORMAT_ASTC_8x6_SRGB_BLOCK
ASTC_8x8	VK_FORMAT_ASTC_8x8_UNORM_BLOCK,
	VK_FORMAT_ASTC_8x8_SRGB_BLOCK
ASTC_10x5	VK_FORMAT_ASTC_10x5_UNORM_BLOCK,
	VK_FORMAT_ASTC_10x5_SRGB_BLOCK
ASTC_10x6	VK_FORMAT_ASTC_10x6_UNORM_BLOCK,
	VK_FORMAT_ASTC_10x6_SRGB_BLOCK
ASTC_10x8	VK_FORMAT_ASTC_10x8_UNORM_BLOCK,
	VK_FORMAT_ASTC_10x8_SRGB_BLOCK
ASTC_10x10	VK_FORMAT_ASTC_10x10_UNORM_BLOCK,
	VK_FORMAT_ASTC_10x10_SRGB_BLOCK
ASTC_12x10	VK_FORMAT_ASTC_12x10_UNORM_BLOCK,
	VK_FORMAT_ASTC_12x10_SRGB_BLOCK
ASTC_12x12	VK_FORMAT_ASTC_12x12_UNORM_BLOCK,
	VK_FORMAT_ASTC_12x12_SRGB_BLOCK
D16	VK_FORMAT_D16_UNORM
D24	VK_FORMAT_X8_D24_UNORM_PACK32
D32	VK_FORMAT_D32_SFLOAT
S8	VK_FORMAT_S8_UINT
D16S8	VK_FORMAT_D16_UNORM_S8_UINT
D24S8	VK_FORMAT_D24_UNORM_S8_UINT
D32S8	VK_FORMAT_D32_SFLOAT_S8_UINT

31.3.2 Format Properties

To query supported format features which are properties of the physical device, call:

- physicalDevice is the physical device from which to query the format properties.
- format is the format whose properties are queried.
- pFormatProperties is a pointer to a VkFormatProperties structure in which physical device properties for format are returned.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- format must be a valid VkFormat value
- pFormatProperties must be a pointer to a VkFormatProperties structure

The VkPhysicalDeviceLimits structure is defined as:

- linearTilingFeatures describes the features supported by VK_IMAGE_TILING_LINEAR.
- optimalTilingFeatures describes the features supported by VK_IMAGE_TILING_OPTIMAL.
- bufferFeatures describes the features supported by buffers.

Supported features are described as a set of VkFormatFeatureFlagBits:

```
typedef enum VkFormatFeatureFlagBits {
   VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT = 0x00000001,
   VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT = 0x00000002,
   VK_FORMAT_FEATURE_STORAGE_IMAGE_ATOMIC_BIT = 0x000000004,
   VK_FORMAT_FEATURE_UNIFORM_TEXEL_BUFFER_BIT = 0x000000008,
   VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_BIT = 0x00000010,
   VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_ATOMIC_BIT = 0x000000020,
   VK_FORMAT_FEATURE_VERTEX_BUFFER_BIT = 0x000000040,
   VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT = 0x000000080,
```

```
VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BLEND_BIT = 0x00000100,
VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT = 0x00000200,
VK_FORMAT_FEATURE_BLIT_SRC_BIT = 0x00000400,
VK_FORMAT_FEATURE_BLIT_DST_BIT = 0x00000800,
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT = 0x00001000,
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG = 0x00002000,
} VkFormatFeatureFlagBits;
```

The linearTilingFeatures and optimalTilingFeatures members of the VkFormatProperties structure describe what features are supported by VK_IMAGE_TILING_LINEAR and VK_IMAGE_TILING_OPTIMAL images, respectively.

The following bits may be set in <code>linearTilingFeatures</code> and <code>optimalTilingFeatures</code>, indicating they are supported by images or image views created with the queried <code>vkGetPhysicalDeviceFormatProperties::format:</code>

VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT

VkImageView can be sampled from. See sampled images section.

VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT

VkImageView can be used as storage image. See storage images section.

VK FORMAT FEATURE STORAGE IMAGE ATOMIC BIT

VkImageView can be used as storage image that supports atomic operations.

VK FORMAT FEATURE COLOR ATTACHMENT BIT

VkImageView can be used as a framebuffer color attachment and as an input attachment.

VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BLEND_BIT

VkImageView can be used as a framebuffer color attachment that supports blending and as an input attachment.

VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT

VkImageView can be used as a framebuffer depth/stencil attachment and as an input attachment.

VK FORMAT FEATURE BLIT SRC BIT

VkImage can be used as srcImage for the vkCmdBlitImage command.

VK_FORMAT_FEATURE_BLIT_DST_BIT

VkImage can be used as dstImage for the vkCmdBlitImage command.

VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT

If VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT is also set, VkImageView can be used with a sampler that has either of magFilter or minFilter set to VK_FILTER_LINEAR, or mipmapMode set to VK_SAMPLER_MIPMAP_MODE_LINEAR. If VK_FORMAT_FEATURE_BLIT_SRC_BIT is also set, VkImage can be used as the srcImage to vkCmdBlitImage with a filter of VK_FILTER_LINEAR. This bit must only be exposed for formats that also support the VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT or VK_FORMAT_FEATURE_BLIT_SRC_BIT.

If the format being queried is a depth/stencil format, this bit only indicates that the depth aspect (not the stencil aspect) of an image of this format supports linear filtering, and that linear filtering of the depth aspect is supported whether depth compare is enabled in the sampler or not. If this bit is not present, linear filtering with depth compare disabled is unsupported and linear filtering with depth compare enabled is supported, but may compute the filtered value in an implementation-dependent manner which differs from the normal rules of linear filtering. The resulting value must be in the range [0,1] and should be proportional to, or a weighted average of, the number of comparison passes or failures.

VK FORMAT FEATURE SAMPLED IMAGE FILTER CUBIC BIT IMG

VkImage can be used with a sampler that has either of magFilter or minFilter set to VK_FILTER_CUBIC_IMG, or be the source image for a blit with filter set to VK_FILTER_CUBIC_IMG. This bit must only be exposed for formats that also support the VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT. If the format being queried is a depth/stencil format, this only indicates that the depth aspect is cubic filterable.

The following features may appear in bufferFeatures, indicating they are supported by buffers or buffer views created with the queried vkGetPhysicalDeviceFormatProperties::format:

VK FORMAT FEATURE UNIFORM TEXEL BUFFER BIT

Format can be used to create a VkBufferView that can be bound to a VK_DESCRIPTOR_TYPE_UNIFORM_ TEXEL_BUFFER descriptor.

VK FORMAT FEATURE STORAGE TEXEL BUFFER BIT

Format can be used to create a VkBufferView that can be bound to a VK_DESCRIPTOR_TYPE_STORAGE_ TEXEL_BUFFER descriptor.

VK FORMAT FEATURE STORAGE TEXEL BUFFER ATOMIC BIT

Atomic operations are supported on VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER with this format.

VK FORMAT FEATURE VERTEX BUFFER BIT

Format can be used as a vertex attribute format (VkVertexInputAttributeDescription::format).



Note

If no format feature flags are supported, then the only possible use would be image transfers - which alone are not useful. As such, if no format feature flags are supported, the format itself is not supported, and images of that format cannot be created.

If format is a block-compression format, then buffers must not support any features for the format.

31.3.3 Required Format Support

Implementations must support at least the following set of features on the listed formats. For images, these features must be supported for every VkImageType (including arrayed and cube variants) unless otherwise noted. These features are supported on existing formats without needing to advertise an extension or needing to explicitly enable them. Support for additional functionality beyond the requirements listed here is queried using the vkGetPhysicalDeviceFormatProperties command.

The following tables show which feature bits must be supported for each format.

Table 31.10: Key for format feature tables

✓	This feature must be supported on the named format
†	This feature must be supported on at least some of the named
	formats, with more information in the table where the symbol
	appears

Table 31.11: Feature bits in optimalTilingFeatures

VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT
VK_FORMAT_FEATURE_BLIT_SRC_BIT
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT
VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT
VK_FORMAT_FEATURE_STORAGE_IMAGE_ATOMIC_BIT
VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT
VK_FORMAT_FEATURE_BLIT_DST_BIT
VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BLEND_BIT
VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT

Table 31.12: Feature bits in bufferFeatures

VK_FORMAT_FEATURE_VERTEX_BUFFER_BIT
VK_FORMAT_FEATURE_UNIFORM_TEXEL_BUFFER_BIT
VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_BIT
VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_ATOMIC_BIT

Table 31.13: Mandatory format support: sub-byte channels

VK_FORMAT_FEA	TURE	S	ror.	AGE	_TE	XEI	L_BI	JFF.	ER_	ATC)IMC	С_В	IT
VK_FORMAT												ΙΤ	
VK_FORMAT_FF	EATU	RE_	UNI	FOE	RM_	TEX	EL_	BUF	FEI	R_B	ΙT		
VK_FORM										ΙT			
VK_FORMAT_FEATURE_DEP									ΙT				
VK_FORMAT_FEATURE_COLC								ΙT					
VK_FORMAT_F	'EAT	URE	_BL	IT_	_DS	Г_В	ΙT						
VK_FORMAT_FEATURE_COL						ΙT						↓	*
VK_FORMAT_FEATURE_STORAGE_IM					ΙT				l .i.	↓	*		
VK_FORMAT_FEATURE_STORAG				ΙT			1	↓	*				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI			ΙT		l .i.	↓	*						
VK_FORMAT_FEATURE_BLIT_SE		ΙT		↓	*								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_	BIT	↓	*										
Format	<u> </u>												
VK_FORMAT_UNDEFINED													
VK_FORMAT_R4G4_UNORM_PACK8													
VK_FORMAT_R4G4B4A4_UNORM_PACK16													
VK_FORMAT_B4G4R4A4_UNORM_PACK16	/	1	1										
VK_FORMAT_R5G6B5_UNORM_PACK16	/	/	✓			/	√	√					
VK_FORMAT_B5G6R5_UNORM_PACK16													
VK_FORMAT_R5G5B5A1_UNORM_PACK16													
VK_FORMAT_B5G5R5A1_UNORM_PACK16													
VK_FORMAT_A1R5G5B5_UNORM_PACK16	/	/	/			/	✓	/					

Table 31.14: Mandatory format support: 1-3 byte-sized channels

VK_FORMAT_FEAT	rurf	: S:	ror	AGE	ТЕ	EXEI	L B	UFF	ER	ATO)MT(; B	ΙΤ
VK_FORMAT													_
VK_FORMAT_FE													
VK_FORMA								_					
VK_FORMAT_FEATURE_DEP1													
VK_FORMAT_FEATURE_COLO													
VK_FORMAT_F									1				١.
VK FORMAT FEATURE COLO				_							١,	↓	↓
VK_FORMAT_FEATURE_STORAGE_IMA	GE	ATC)MI	СВ	IT		1		١.		1	'	
VK_FORMAT_FEATURE_STORAGE						1			↓				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI	VEAI	R_B	ΙΤ		1.		\	'					
VK_FORMAT_FEATURE_BLIT_SR	СВ	ΙΤ		↓	↓	'							
VK_FORMAT_FEATURE_SAMPLED_IMAGE_B		Ţ	+	'									
Format	T	•											
VK FORMAT R8 UNORM	1	1	1			1	1	1		1	1		
VK FORMAT R8 SNORM	1		1							1			
VK_FORMAT_R8_USCALED		-											
VK FORMAT R8 SSCALED													
VK_FORMAT_R8_UINT	1	1				1	1			1	1		
VK FORMAT R8 SINT	1	1				1	1			1	1		
VK_FORMAT_R8_SRGB													
VK_FORMAT_R8G8_UNORM	1	1	1			1	1	1		1	1		
VK_FORMAT_R8G8_SNORM	1	1	1							1	1		
VK_FORMAT_R8G8_USCALED													
VK_FORMAT_R8G8_SSCALED													
VK_FORMAT_R8G8_UINT	1	1				1	1			1	1		
VK_FORMAT_R8G8_SINT	1	1				1	1			1	1		
VK_FORMAT_R8G8_SRGB													
VK_FORMAT_R8G8B8_UNORM													
VK_FORMAT_R8G8B8_SNORM													
VK_FORMAT_R8G8B8_USCALED													
VK_FORMAT_R8G8B8_SSCALED													
VK_FORMAT_R8G8B8_UINT													
VK_FORMAT_R8G8B8_SINT													
VK_FORMAT_R8G8B8_SRGB													
VK_FORMAT_B8G8R8_UNORM													
VK_FORMAT_B8G8R8_SNORM													
VK_FORMAT_B8G8R8_USCALED													
VK_FORMAT_B8G8R8_SSCALED													
VK_FORMAT_B8G8R8_UINT													
VK_FORMAT_B8G8R8_SINT													
VK_FORMAT_B8G8R8_SRGB													

Table 31.15: Mandatory format support: 4 byte-sized channels

VK_FORMAT_FEAT	TURE	S	ΓOR	AGE	_TE	EXE	L_B	UFF	ER_	ATC	MI	C_B	ΙΤ
VK_FORMAT_	_FEA	ATU:	RE_	STC)RA(GE_	TEX	EL_	BUE	FEI	R_B	ΙΤ	
VK_FORMAT_FE	ATU!	RE_	UNI	FOE	RM_	TEX	EL_	BUE	FEI	R_B	ΙT		
VK_FORMA	T_F	EAT	URE	_V	ERT	EX_	BUE	FE.	R_B	ΙT			
VK_FORMAT_FEATURE_DEPI	H_S	TEI	ICI	L_A	TTA	CHN	IEN'	Г_В	ΙT				
VK_FORMAT_FEATURE_COLOR	R_AI	TA	СНМ	ENT	_BI	LENI	D_B	ΙT					
VK_FORMAT_F	EAT	JRE	_BL	IT_	_DS'	T_B	ΙT						l , '
VK_FORMAT_FEATURE_COLO	R_A	TTA	СНМ	IEN:	Г_В	ΙT						↓	+
VK_FORMAT_FEATURE_STORAGE_IMA	GE_	ATO)IMC	С_В	ΙT				١,		+		
VK_FORMAT_FEATURE_STORAGE	E_IM	IAGI	E_B	ΙΤ			١,	↓	↓				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI	NEAI	R_B	ΙT		١,		\						
VK_FORMAT_FEATURE_BLIT_SR	С_В	ΙT		↓	↓								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_B	IT	\downarrow	+										
Format	\	Ċ											
VK_FORMAT_R8G8B8A8_UNORM	1	1	1	1		1	1	1		1	1	1	
VK_FORMAT_R8G8B8A8_SNORM	1	1	1	1						1	1	1	
VK_FORMAT_R8G8B8A8_USCALED													
VK_FORMAT_R8G8B8A8_SSCALED													
VK_FORMAT_R8G8B8A8_UINT	1	1		1		1	1			1	1	1	
VK_FORMAT_R8G8B8A8_SINT	1	1		√		1	1			1	√	1	
VK_FORMAT_R8G8B8A8_SRGB	1	1	1			1	1	1					
VK_FORMAT_B8G8R8A8_UNORM	1	1	1			1	1	1		1	1		
VK_FORMAT_B8G8R8A8_SNORM													
VK_FORMAT_B8G8R8A8_USCALED													
VK_FORMAT_B8G8R8A8_SSCALED													
VK_FORMAT_B8G8R8A8_UINT													
VK_FORMAT_B8G8R8A8_SINT													
VK_FORMAT_B8G8R8A8_SRGB	1	1	1			1	1	1					
VK_FORMAT_A8B8G8R8_UNORM_PACK32	1	1	1			1	1	1		1	1	1	
VK_FORMAT_A8B8G8R8_SNORM_PACK32	1	1	1							1	1	1	
VK_FORMAT_A8B8G8R8_USCALED_PACK32													
VK_FORMAT_A8B8G8R8_SSCALED_PACK32													
VK_FORMAT_A8B8G8R8_UINT_PACK32	1	1				1	1			1	✓	1	
VK_FORMAT_A8B8G8R8_SINT_PACK32	1	\				1	1			1	✓	1	
VK_FORMAT_A8B8G8R8_SRGB_PACK32	1	1	✓			1	1	1					

Table 31.16: Mandatory format support: 10-bit channels

VK_FORMAT_FEA	TURE	: S:	ror.	AGE	TF	XEI	L B	UFF	ER	ATO)MT(; B	ΙΤ
VK_FORMAT													
VK_FORMAT_F													
VK_FORM													
VK_FORMAT_FEATURE_DEP	TH_S	TEN	ICI:	L_A	TTA	CHN	IEN:	Г_В	ΙT				
VK_FORMAT_FEATURE_COLO	R_AT	ГТА	СНМ	ENT	_BI	ENI	D_B	ΙT					
VK_FORMAT_I	EAT	JRE	_BL	IT_	_DS:	Г_В	ΙT						١,
VK_FORMAT_FEATURE_COL	OR_A	TTA	CHN	IEN:	Г_В	ΙT						↓	+
VK_FORMAT_FEATURE_STORAGE_IM	AGE_	ATC)IMC	С_В	ΙT				↓	1	+		
VK_FORMAT_FEATURE_STORAG				ΙT				↓	*				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI			ΙΤ		↓	↓	+						
VK_FORMAT_FEATURE_BLIT_S		ΙT		↓	*								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_	BIT	↓	*										
Format	↓ ↓												
VK_FORMAT_A2R10G10B10_UNORM_PACK32													
VK_FORMAT_A2R10G10B10_SNORM_PACK32													
VK_FORMAT_A2R10G10B10_USCALED_PACK32													
VK_FORMAT_A2R10G10B10_SSCALED_PACK32													
VK_FORMAT_A2R10G10B10_UINT_PACK32													
VK_FORMAT_A2R10G10B10_SINT_PACK32													
VK_FORMAT_A2B10G10R10_UNORM_PACK32	/	✓	✓			/	/	/		/	/		
VK_FORMAT_A2B10G10R10_SNORM_PACK32													
VK_FORMAT_A2B10G10R10_USCALED_PACK32													
VK_FORMAT_A2B10G10R10_SSCALED_PACK32						L_							_
VK_FORMAT_A2B10G10R10_UINT_PACK32	/	1				/	/				1		<u> </u>
VK_FORMAT_A2B10G10R10_SINT_PACK32													

Table 31.17: Mandatory format support: 16-bit channels

VK_FORMAT_FEA	TURE	S	ror.	AGE	_TE	XEI	B	UFF	ER_	ATC	MIC	B	IT
VK_FORMAT	Γ_FE	ATU	RE_	STC	RAC	GE_	ГЕХ	EL_	BUE	FEI	R_B	ΙΤ	
VK_FORMAT_F	EATU:	RE_	UNI	FOF	RM_	ГЕХ	EL_	BUE	FEI	R_B	ΙT		
VK_FORM	AT_F	EAT	URE	C_VI	ERT	EX_	BUE	FE.	R_B	ΙT			
VK_FORMAT_FEATURE_DEP	TH_S	TEI	ICI:	L_A	TTA	СНМ	IEN:	Г_В	ΙT				
VK_FORMAT_FEATURE_COLO	DR_AT	TA	СНМ	ENT	_BI	ENI	D_B	ΙT					
VK_FORMAT_I	FEAT	JRE	_BL	IT_	DS:	Г_В	ΙΤ		1				١,
VK_FORMAT_FEATURE_COL	OR_A	TTA	CHN	IEN:	Г_В	ΙΤ						\downarrow	↓
VK_FORMAT_FEATURE_STORAGE_IM	AGE_	ATO)IMC	С_В	ΙΤ				١,		↓		
VK_FORMAT_FEATURE_STORAG	SE_IN	1AGI	E_B	ΙT					↓				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_L	INEAI	R_B	ΙT		١,		+	,					
VK_FORMAT_FEATURE_BLIT_S	RC_B	ΙT		↓	↓								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_	BIT	1	\	i i									
Format	\perp	'											
VK_FORMAT_R16_UNORM										1			
VK FORMAT R16 SNORM										1			
VK FORMAT R16 USCALED													
VK_FORMAT_R16_SSCALED													
VK_FORMAT_R16_UINT	1	1				1	1			1	1		
VK_FORMAT_R16_SINT	1	1				1	1			1	1		
VK_FORMAT_R16_SFLOAT	1	1	1			1	1	1		1	1		
VK_FORMAT_R16G16_UNORM										1			
VK_FORMAT_R16G16_SNORM										1			
VK_FORMAT_R16G16_USCALED													
VK_FORMAT_R16G16_SSCALED													
VK_FORMAT_R16G16_UINT	1	1				1	1			1	1		
VK_FORMAT_R16G16_SINT	1	1				1	1			1	1		
VK_FORMAT_R16G16_SFLOAT	1	1	1			1	√	1		1	1		
VK_FORMAT_R16G16B16_UNORM													
VK_FORMAT_R16G16B16_SNORM													
VK_FORMAT_R16G16B16_USCALED													
VK_FORMAT_R16G16B16_SSCALED													
VK_FORMAT_R16G16B16_UINT													
VK_FORMAT_R16G16B16_SINT													
VK_FORMAT_R16G16B16_SFLOAT													
VK_FORMAT_R16G16B16A16_UNORM										1			
VK_FORMAT_R16G16B16A16_SNORM										√			
VK_FORMAT_R16G16B16A16_USCALED													
VK_FORMAT_R16G16B16A16_SSCALED													
VK_FORMAT_R16G16B16A16_UINT	/			1		1	✓			✓	1	✓	
VK_FORMAT_R16G16B16A16_SINT	/	1		1		1	✓			1	✓	✓	
VK_FORMAT_R16G16B16A16_SFLOAT	√	✓	✓	✓		✓	✓	/		✓	✓	√	

Table 31.18: Mandatory format support: 32-bit channels

VK_FORMAT_FEA	ת ידוד ה		r O D	۸۲۶	тъ	יעני	. Di	יםיםו	r D	7 T/	NAT (¬ p	тт
VK_FORMA'													T T
VK_FORMAT_F												11	
VK_FORM													
VK_FORMAT_FEATURE_DEF													
VK_FORMAT_FEATURE_COL(
VK_FORMAT_													
VK_FORMAT_FEATURE_COL													↓
VK_FORMAT_FEATURE_STORAGE_IM										↓	\downarrow	*	
VK_FORMAT_FEATURE_STORAG								↓	↓	*			
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_L						↓	↓	*					
VK_FORMAT_FEATURE_BLIT_S				1	↓	*							
VK_FORMAT_FEATURE_SAMPLED_IMAGE_		П	\	~									
Format		· ·											
VK FORMAT R32 UINT	1	1		1	1	1	1			1	/	1	1
VK_FORMAT_R32_SINT	1	1		1	1	1	1			1	1	1	1
VK_FORMAT_R32_SFLOAT	1	1		1		1	1			1	/	1	
THE DODGE DOGGO HIND		-	-										
VK_FORMAT_R32G32_UINT	1	1		1		1	✓			1	1	1	
VK_FORMAT_R32G32_UINT VK_FORMAT_R32G32_SINT	1	✓ ✓		✓ ✓		✓ ✓	1			√	✓ ✓	✓ ✓	
		•					•			✓ ✓	✓ ✓	✓ ✓	
VK_FORMAT_R32G32_SINT	1	•		√			•			√ √ √	√ √ √	√ √ √	
VK_FORMAT_R32G32_SINT VK_FORMAT_R32G32_SFLOAT	1	•		√			•			\frac{1}{\sqrt{1}}	√ √ √	✓ ✓ ✓	
VK_FORMAT_R32G32_SINT VK_FORMAT_R32G32_SFLOAT VK_FORMAT_R32G32B32_UINT VK_FORMAT_R32G32B32_SINT VK_FORMAT_R32G32B32_SFLOAT	1	•		√			•			\frac{1}{\sqrt{1}}	✓ ✓ ✓	✓ ✓ ✓	
VK_FORMAT_R32G32_SINT VK_FORMAT_R32G32_SFLOAT VK_FORMAT_R32G32B32_UINT VK_FORMAT_R32G32B32_SINT VK_FORMAT_R32G32B32_SFLOAT VK_FORMAT_R32G32B32A32_UINT	1	•		√			•			\frac{1}{\sqrt{1}}	\frac{1}{1}	\frac{1}{\sqrt{1}}	
VK_FORMAT_R32G32_SINT VK_FORMAT_R32G32_SFLOAT VK_FORMAT_R32G32B32_UINT VK_FORMAT_R32G32B32_SINT VK_FORMAT_R32G32B32_SFLOAT	<i>J</i>	√ ✓		√ ✓		✓ ✓	✓ ✓			\frac{1}{\sqrt{1}}	\frac{1}{\sqrt{1}}	\frac{1}{\sqrt{1}}	

Table 31.19: Mandatory format support: 64-bit/uneven channels and depth/stencil

VK FORMAT FEAT	rure	S	ror.	AGE	TF	XEI	ı B	UFF	ER	ATC)MT(СВ	IT
VK_FORMAT													
VK_FORMAT_FE													1
VK_FORMA								_				<u> </u>	
 VK_FORMAT_FEATURE_DEP1							_						
VK_FORMAT_FEATURE_COLO										1			
VK_FORMAT_F													١,
VK_FORMAT_FEATURE_COLO												↓	↓
VK_FORMAT_FEATURE_STORAGE_IMA	GE_	ATO)IMC	C_B	IT				١,		+		
VK_FORMAT_FEATURE_STORAGE	E_IN	/IAGI	E_B	ΙΤ		1		↓	↓				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI	NEAI	R_B	ΙT		1 .		+	ľ					
VK_FORMAT_FEATURE_BLIT_SR	С_В	ΙT		\downarrow	↓								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_B	IT	\downarrow	+										
Format	\downarrow	1											
VK_FORMAT_R64_UINT													
VK_FORMAT_R64_SINT													
VK_FORMAT_R64_SFLOAT													
VK_FORMAT_R64G64_UINT													
VK_FORMAT_R64G64_SINT													
VK_FORMAT_R64G64_SFLOAT													
VK_FORMAT_R64G64B64_UINT													
VK_FORMAT_R64G64B64_SINT													
VK_FORMAT_R64G64B64_SFLOAT													
VK_FORMAT_R64G64B64A64_UINT													
VK_FORMAT_R64G64B64A64_SINT													
VK_FORMAT_R64G64B64A64_SFLOAT													
VK_FORMAT_B10G11R11_UFLOAT_PACK32	1	/	✓								✓		
VK_FORMAT_E5B9G9R9_UFLOAT_PACK32	1	1	√										
VK_FORMAT_D16_UNORM	/	/							/				
VK_FORMAT_X8_D24_UNORM_PACK32									†				
VK_FORMAT_D32_SFLOAT	✓	1							†				
VK_FORMAT_S8_UINT													
VK_FORMAT_D16_UNORM_S8_UINT													
VK_FORMAT_D24_UNORM_S8_UINT									†				
VK_FORMAT_D32_SFLOAT_S8_UINT					L_				†	L			$oxed{oxed}$
VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT													
VK_FORMAT_X8_D24_UNORM_PACK32 and VK_FORMAT_	D32	_SF	LOP	ΔТ, а	and 1	nust	be s	supp	orte	d foi	at l	east	

one of VK_FORMAT_D24_UNORM_S8_UINT and VK_FORMAT_D32_SFLOAT_S8_UINT.

Table 31.20: Mandatory format support: BC compressed formats with

VK_FORMAT_FEA													ΙT
VK_FORMAT				-					_			IT	
VK_FORMAT_FE	ATU	RE_	UNI	FOF	RM_	ΓΕΧ	EL_	BUE	FEI	R_B	ΙT		
VK_FORMA	AT_F	EAI	URE	E_V	ERT	EX_	BUE	FEI	R_B	ΙT			
VK_FORMAT_FEATURE_DEP:	ΓH_S	STEI	NCI:	L_A	TTA	CHM	IEN:	Г_В	ΙT				
VK_FORMAT_FEATURE_COLO	R_AT	ГТА	СНМ	ENT	_BI	ENI)_B	ΙT					
VK_FORMAT_F	EAT	URE	_BL	IT_	DS:	Г_В	ΙT						
VK_FORMAT_FEATURE_COLO	R_A	TTA	CHM	IEN:	Г_В	ΙT					\downarrow	↓	*
VK_FORMAT_FEATURE_STORAGE_IMA	AGE_	AT()IMC	С_В	ΙT					↓	+		
VK_FORMAT_FEATURE_STORAG	E_IN	/IAG	E_B	ΙT				↓	+				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LI	NEAI	R_B	ΙT		↓	↓	+						
VK_FORMAT_FEATURE_BLIT_SR	.C_B	ΙT	1	↓	+								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_E	ВІТ	1	1										
Format	1												
VK_FORMAT_BC1_RGB_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC1_RGB_SRGB_BLOCK	†	†	†										
VK_FORMAT_BC1_RGBA_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC1_RGBA_SRGB_BLOCK	†	†	†										
VK_FORMAT_BC2_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC2_SRGB_BLOCK	†	†	†										
VK_FORMAT_BC3_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC3_SRGB_BLOCK	†	†	†										
VK_FORMAT_BC4_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC4_SNORM_BLOCK	†	†	†										
VK_FORMAT_BC5_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC5_SNORM_BLOCK	†	†	†										
VK_FORMAT_BC6H_UFLOAT_BLOCK	†	†	†										
VK_FORMAT_BC6H_SFLOAT_BLOCK	†	†	†										
VK_FORMAT_BC7_UNORM_BLOCK	†	†	†										
VK_FORMAT_BC7_SRGB_BLOCK	†	†	†										
The VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT, VK_										IT a	nd V	K_	
FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR													
optimalTilingFeatures for all the formats in at least one of	of: th	is ta	ble, '	Tabl	e 31	.21,	or T	able	31.	22.			

Table 31.21: Mandatory format support: ETC2 and EAC compressed formats with $\mbox{VkImageType}$ $\mbox{VK_IMAGE_TYPE_2D}$

VK_FORMAT_FEA	TURI	S	ΓOR	AGE	_TE	CXEI	_BI	UFF	ER_	ATC	OMIC	С_В	ΙΤ
VK_FORMA	Γ_FE	ATU	RE_	STC	RAC	GE_	ΓEΧ	EL_	BUE	FEI	R_B	ΙΤ	
VK_FORMAT_F	EATU	RE_	UNI	FOF	RM_'	TEX:	EL_	BUF	'FEI	R_B	ΙT		
VK_FORM	AT_F	'EAI	URE	_V	ERT	EX_	BUF	FEE	R_B	ΙT			
VK_FORMAT_FEATURE_DEF									ΙT				
VK_FORMAT_FEATURE_COLO								ΙT					
VK_FORMAT_							ΙΤ						l ,
VK_FORMAT_FEATURE_COL						ΙT						↓	*
VK_FORMAT_FEATURE_STORAGE_IM					ΙT				↓	↓	*		
VK_FORMAT_FEATURE_STORAG				ΙT			.1.	↓	*				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_L			IT		l .i.	↓	*						
VK_FORMAT_FEATURE_BLIT_S		IT	1.	↓	*								
VK_FORMAT_FEATURE_SAMPLED_IMAGE_	BIT	↓	~										
Format	<u></u>												
VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK	†	†	†										
VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK	†	†	†										
VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK	†	†	†										
VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK	†	†	†										
VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK	†	†	†										<u></u>
VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK	†	†	†										
VK_FORMAT_EAC_R11_UNORM_BLOCK	†	†	†										
VK_FORMAT_EAC_R11_SNORM_BLOCK	†	†	†										
VK_FORMAT_EAC_R11G11_UNORM_BLOCK	†	†	†										\perp
VK_FORMAT_EAC_R11G11_SNORM_BLOCK	†	†	†										
The VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT, VK_										T a	nd V	K_	
FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEA													
optimalTilingFeatures for all the formats in at least one	of: th	is tal	ble, '	Fabl	e 31	.20,	or T	able	31.	22.			

Table 31.22: Mandatory format support: ASTC LDR compressed formats with $\mbox{VkImageType}$ $\mbox{VK_IMAGE_TYPE_2D}$

VK_FORMAT_FEA													IT
VK_FORMA	T_FE	UTA	RE_	STO)RA(GE_	TEX	EL_	BUE	FEI	R_B	ΙT	
VK_FORMAT_F	EATU	RE_	UNI	FOI	RM_	TEX	EL_	BUE	FEI	R_B	ΙT		
VK_FORM	IAT_F	EAT	URE	_V	ERT	EX_	BUE	FE	R_B	ΙT			
VK_FORMAT_FEATURE_DEE	TH_S	TEI	NCI	L_A	TTA	CHN	1ΕΝ΄	T_B	ΙT				
VK_FORMAT_FEATURE_COL(OR_AT	ГТА	СНМ	ENI	_BI	LENI	D_B	ΙT					
VK_FORMAT_	FEAT	URE	_BL	IT_	_DS'	T_B	ΙT						١,
VK_FORMAT_FEATURE_COL	OR_A	TTA	CHM	IEN'	Т_В	ΙT						↓	+
VK_FORMAT_FEATURE_STORAGE_IM	IAGE_	ATO	OMIC	С_В	ΙT				١,	↓	+		
VK_FORMAT_FEATURE_STORAG	GE_IN	1AG1	E_B	ΙT				↓	↓				
VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_L	INEA	R_B	ΙΤ		1.	↓	↓						
VK_FORMAT_FEATURE_BLIT_S	RC_B	ΙT		↓	↓	'							
VK_FORMAT_FEATURE_SAMPLED_IMAGE_	BIT	Ţ	↓	,									
Format		'											
VK_FORMAT_ASTC_4x4_UNORM_BLOCK	†	†	†										
VK_FORMAT_ASTC_4x4_SRGB_BLOCK	+	†	+										
VK_FORMAT_ASTC_5x4_UNORM_BLOCK	+	†	†										
VK_FORMAT_ASTC_5x4_SRGB_BLOCK	†	+	+										
VK_FORMAT_ASTC_5x5_UNORM_BLOCK	†	†	+										
VK_FORMAT_ASTC_5x5_SRGB_BLOCK	+	+	+										
VK_FORMAT_ASTC_6x5_UNORM_BLOCK	+	+	+										
VK_FORMAT_ASTC_6x5_SRGB_BLOCK	†	†	+										\vdash
VK_FORMAT_ASTC_6x6_UNORM_BLOCK	+	+	+										
VK_FORMAT_ASTC_6x6_SRGB_BLOCK	†	+	+										
VK_FORMAT_ASTC_8x5_UNORM_BLOCK	†	+	+										\vdash
VK_FORMAT_ASTC_8x5_SRGB_BLOCK	+	+	†										
VK_FORMAT_ASTC_8x6_UNORM_BLOCK	+	+	†										\vdash
VK_FORMAT_ASTC_8x6_SRGB_BLOCK	+	+	+										
VK_FORMAT_ASTC_8x8_UNORM_BLOCK	+	+	+										_
VK_FORMAT_ASTC_8x8_SRGB_BLOCK	†	+	+										_
VK_FORMAT_ASTC_10x5_UNORM_BLOCK	†	+	†										\vdash
VK_FORMAT_ASTC_10x5_SRGB_BLOCK	†	†	†										
VK_FORMAT_ASTC_10x6_UNORM_BLOCK	†	†	†										
VK_FORMAT_ASTC_10x6_SRGB_BLOCK	†	+	†										
VK_FORMAT_ASTC_10x8_UNORM_BLOCK	+	+	†										
VK_FORMAT_ASTC_10x8_SRGB_BLOCK	+	+	+										_
VK_FORMAT_ASTC_10x10_UNORM_BLOCK	+	+	†										-
VK_FORMAT_ASTC_10x10_SNGB_BLOCK	+	+	+										
VK_FORMAT_ASTC_12x10_UNORM_BLOCK	+	+	†										
VK_FORMAT_ASTC_12x10_SRGB_BLOCK	†	†	†										-
VK_FORMAT_ASTC_12x12_UNORM_BLOCK	†	+	†										
VK_FORMAT_ASTC_12x12_ONORM_BLOCK VK_FORMAT_ASTC_12x12_SRGB_BLOCK	†	+	†										-
The VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT, VK_	'	<u>' '</u> 72 T	'	<u> </u>	IRF IRF	RT.	L TT	SRC	 R1	<u> </u> ГТ эт	nd 🗤	K	
FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEA										L	ıu v	-,—	
optimalTilingFeatures for all the formats in at least one										21			
opermatititing eacutes for an the formats in at least one	o1. u1.	is ia	J10,	I aUI	. U J I	.20,	OI I	aoic	, , , 1	<u>~ 1 .</u>			

 $\label{thm:commutation} $$ VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG \ must \ be supported for the following formats: $$ VK_FORMAT_R4G4_UNORM_PACK8 $$ VK_FORMAT_R4G4B4A4_UNORM_PACK16 $$ VK_FORMAT_R4G4B4A4_UNORM_PACK16 $$ VK_FORMAT_R4G4B4A4_UNORM_PACK16 $$ VK_FORMAT_R4G4B4A4_UNORM_PACK16 $$ VK_FORMAT_R4G4B4A4_UNORM_PACK16 $$ VK_FORMAT_R4G4B4A4 $$ VK_FORMAT $$ VK_FORMAT_R4G4B4A4 $$ VK_FORMAT $$ VK_FOR$

If ETC2 compressed formats are supported, the following additional formats must support VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_CUBIC_BIT_IMG: *VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK *VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK *VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK *VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK *VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK *VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK *VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK

31.4 Additional Image Capabilities

In addition to the minimum capabilities described in the previous sections (Limits and Formats), implementations may support additional capabilities for certain types of images. For example, larger dimensions or additional sample counts for certain image types, or additional capabilities for *linear* tiling format images.

To guery additional capabilities specific to image types, call:

- physicalDevice is the physical device from which to query the image capabilities.
- format is the image format, corresponding to VkImageCreateInfo::format.
- type is the image type, corresponding to VkImageCreateInfo::imageType.
- tiling is the image tiling, corresponding to VkImageCreateInfo::tiling.
- usage is the intended usage of the image, corresponding to VkImageCreateInfo::usage.
- flags is a bitmask describing additional parameters of the image, corresponding to VkImageCreateInfo::flags.
- pImageFormatProperties points to an instance of the VkImageFormatProperties structure in which capabilities are returned.

The format, type, tiling, usage, and flags parameters correspond to parameters that would be consumed by vkCreateImage.

If format is not a supported image format, or if the combination of format, type, tiling, usage, and flags is not supported for images, then **vkGetPhysicalDeviceImageFormatProperties** returns VK_ERROR_FORMAT_NOT_SUPPORTED.

The limitations on an image format that are reported by **vkGetPhysicalDeviceImageFormatProperties** have the following property: if **usage1** and **usage2** of type VkImageUsageFlags are such that the bits set in **usage1** are a subset of the bits set in **usage2**, and **flags1** and **flags2** of type VkImageCreateFlags are such that the bits set in **flags1** are a subset of the bits set in **flags2**, then the limitations for **usage1** and **flags1** must be no more strict than the limitations for **usage2** and **flags2**, for all values of *format*, *type*, and *tiling*.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- format must be a valid VkFormat value
- type must be a valid VkImageType value
- tiling must be a valid VkImageTiling value
- usage must be a valid combination of VkImageUsageFlagBits values
- usage must not be 0
- flags must be a valid combination of VkImageCreateFlagBits values
- pImageFormatProperties must be a pointer to a VkImageFormatProperties structure

Return Codes

Success

• VK SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK ERROR OUT OF DEVICE MEMORY
- VK ERROR FORMAT NOT SUPPORTED

The VkImageFormatProperties structure is defined as:

```
VkSampleCountFlags sampleCounts;
VkDeviceSize maxResourceSize;
} VkImageFormatProperties;
```

- maxExtent are the maximum image dimensions. See the Allowed Extent Values section below for how these values are constrained by type.
- maxMipLevels is the maximum number of mipmap levels. maxMipLevels must either be equal to 1 (valid only if tiling is VK_IMAGE_TILING_LINEAR) or be equal to \[log_2(max(width, height, depth)) \] + 1 where width, height, and depth are taken from the corresponding members of maxExtent.
- maxArrayLayers is the maximum number of array layers. maxArrayLayers must either be equal to 1 or be greater than or equal to the maxImageArrayLayers member of VkPhysicalDeviceLimits. A value of 1 is valid only if tiling is VK IMAGE TILING LINEAR or if type is VK IMAGE TYPE 3D.
- sampleCounts is a bitmask of VkSampleCountFlagBits specifying all the supported sample counts for this image as described below.
- maxResourceSize is an upper bound on the total image size in bytes, inclusive of all image subresources. Implementations may have an address space limit on total size of a resource, which is advertised by this property. maxResourceSize must be at least 2³¹.



Note

There is no mechanism to query the size of an image before creating it, to compare that size against maxResou rceSize. If an application attempts to create an image that exceeds this limit, the creation will fail or the image will be invalid. While the advertised limit must be at least 2^{31} , it may not be possible to create an image that approaches that size, particularly for $VK_IMAGE_TYPE_1D$.

If the combination of parameters to **vkGetPhysicalDeviceImageFormatProperties** is not supported by the implementation for use in vkCreateImage, then all members of VkImageFormatProperties will be filled with zero.

To determine the image capabilities compatible with an external memory handle type, call:

```
VkResult vkGetPhysicalDeviceExternalImageFormatPropertiesNV(
    VkPhysicalDevice
                                                 physicalDevice,
    VkFormat
                                                 format,
    VkImageType
                                                 type,
    VkImageTiling
                                                 tiling,
    VkImageUsageFlags
                                                 usage,
    VkImageCreateFlags
                                                 flags,
    VkExternalMemoryHandleTypeFlagsNV
                                                 externalHandleType,
    VkExternalImageFormatPropertiesNV*
                                                 pExternalImageFormatProperties);
```

- physicalDevice is the physical device from which to query the image capabilities
- format is the image format, corresponding to VkImageCreateInfo::format.
- type is the image type, corresponding to VkImageCreateInfo::imageType.
- tiling is the image tiling, corresponding to VkImageCreateInfo::tiling.
- usage is the intended usage of the image, corresponding to VkImageCreateInfo::usage.

- flags is a bitmask describing additional parameters of the image, corresponding to VkImageCreateInfo::flags.
- externalHandleType is either one of the bits from VkExternalMemoryHandleTypeFlagBitsNV, or 0.
- pExternalImageFormatProperties points to an instance of the VkExternalImageFormatPropertiesNV structure in which capabilities are returned.

If <code>externalHandleType</code> is 0, <code>pExternalImageFormatProperties::imageFormatProperties</code> will return the same values as a call to <code>vkGetPhysicalDeviceImageFormatProperties</code>, and the other members of <code>pExternalImageFormatProperties</code> will all be 0. Otherwise, they are filled in as described for <code>VkExternalImageFormatPropertiesNV</code>.

Valid Usage

- physicalDevice must be a valid VkPhysicalDevice handle
- format must be a valid VkFormat value
- type must be a valid VkImageType value
- tiling must be a valid VkImageTiling value
- usage must be a valid combination of VkImageUsageFlagBits values
- usage must not be 0
- flags must be a valid combination of VkImageCreateFlagBits values
- flags must not be 0
- externalHandleType must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- externalHandleType must not be 0
- pExternalImageFormatProperties must be a pointer to a VkExternalImageFormatPropertiesNV structure

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY
- VK_ERROR_FORMAT_NOT_SUPPORTED

The VkExternalImageFormatPropertiesNV structure is defined as:

- imageFormatProperties will be filled in as when calling vkGetPhysicalDeviceImageFormatProperties, but the values returned may vary depending on the external handle type requested.
- externalMemoryFeatures is a bitmask of VkExternalMemoryFeatureFlagBitsNV indicating properties of the external memory handle type (vkGetPhysicalDeviceExternalImageFormatPropertiesNV::externalHandleType) being queried, or 0 if the external memory handle type is 0.
- exportFromImportedHandleTypes is a bitmask of VkExternalMemoryHandleTypeFlagBitsNV containing a bit set for every external handle type that may be used to create memory from which the handles of the type specified in vkGetPhysicalDeviceExternalImageFormatPropertiesNV::externalHandleType can be exported, or 0 if the external memory handle type is 0.
- compatibleHandleTypes is a bitmask of VkExternalMemoryHandleTypeFlagBitsNV containing a bit set for every external handle type that may be specified simultaneously with the handle type specified by vkGetPhysicalDeviceExternalImageFormatPropertiesNV::externalHandleType when calling vkAllocateMemory, or 0 if the external memory handle type is 0. compatibleHandleTypes will always contain vkGetPhysicalDeviceExternalImageFormatPropertiesNV::externalHandleType

Valid Usage

- imageFormatProperties must be a valid VkImageFormatProperties structure
- externalMemoryFeatures must be a valid combination of VkExternalMemoryFeatureFlagBitsNV values
- externalMemoryFeatures must not be 0
- exportFromImportedHandleTypes must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- exportFromImportedHandleTypes must not be 0
- compatibleHandleTypes must be a valid combination of VkExternalMemoryHandleTypeFlagBitsNV values
- compatibleHandleTypes must not be 0

The bits in VkExternalMemoryFeatureFlagBitsNV are defined as follows:

```
typedef enum VkExternalMemoryFeatureFlagBitsNV {
    VK_EXTERNAL_MEMORY_FEATURE_DEDICATED_ONLY_BIT_NV = 0x00000001,
    VK_EXTERNAL_MEMORY_FEATURE_EXPORTABLE_BIT_NV = 0x00000002,
    VK_EXTERNAL_MEMORY_FEATURE_IMPORTABLE_BIT_NV = 0x00000004,
} VkExternalMemoryFeatureFlagBitsNV;
```

- VK_EXTERNAL_MEMORY_FEATURE_DEDICATED_ONLY_BIT_NV: External memory of the specified type must be created as a dedicated allocation when used in the manner specified.
- VK_EXTERNAL_MEMORY_FEATURE_EXPORTABLE_BIT_NV: The implementation supports exporting handles of the specified type.
- VK_EXTERNAL_MEMORY_FEATURE_IMPORTABLE_BIT_NV: The implementation supports importing handles of the specified type.

31.4.1 Supported Sample Counts

vkGetPhysicalDeviceImageFormatProperties returns a bitmask of VkSampleCountFlagBits in sampleCounts specifying the supported sample counts for the image parameters.

sampleCounts will be set to VK_SAMPLE_COUNT_1_BIT if at least one of the following conditions is true:

- tiling is VK_IMAGE_TILING_LINEAR
- type is not VK_IMAGE_TYPE_2D
- flags contains VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT
- The VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT flag in VkFormatProperties::optimalTilingFeatures returned by vkGetPhysicalDeviceFormatProperties is not set

Otherwise, the bits set in <code>sampleCounts</code> will be the sample counts supported for the specified values of <code>usage</code> and <code>format</code>. For each bit set in <code>usage</code>, the supported sample counts relate to the limits in <code>VkPhysicalDeviceLimits</code> as follows:

- If usage includes VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT, a superset of VkPhysicalDeviceLimits::framebufferColorSampleCounts
- If usage includes VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT, and format includes a depth aspect, a superset of VkPhysicalDeviceLimits::framebufferDepthSampleCounts
- If usage includes VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT, and format includes a stencil aspect, a superset of VkPhysicalDeviceLimits::framebufferStencilSampleCounts
- If usage includes VK_IMAGE_USAGE_SAMPLED_BIT, and format includes a color aspect, a superset of VkPhysicalDeviceLimits::sampledImageColorSampleCounts
- If usage includes VK_IMAGE_USAGE_SAMPLED_BIT, and format includes a depth aspect, a superset of VkPhysicalDeviceLimits::sampledImageDepthSampleCounts
- If usage includes VK_IMAGE_USAGE_SAMPLED_BIT, and format is an integer format, a superset of VkPhysicalDeviceLimits::sampledImageIntegerSampleCounts
- If usage includes VK_IMAGE_USAGE_STORAGE_BIT, a superset of VkPhysicalDeviceLimits::storageImageSampleCounts

If multiple bits are set in usage, sampleCounts will be the intersection of the per-usage values described above.

31.4.2 Allowed Extent Values Based On Image Type

For VK_IMAGE_TYPE_1D:

- $\bullet \ maxExtent.width \leq VkPhysicalDeviceLimits.maxImageDimension1D$
- *maxExtent.height* = 1
- maxExtent.depth = 1

For VK_IMAGE_TYPE_2D:

- $\bullet \ maxExtent.width \leq VkPhysicalDeviceLimits.maxImageDimension2D$
- $maxExtent.height \le VkPhysicalDeviceLimits.maxImageDimension2D$
- maxExtent.depth = 1

For VK_IMAGE_TYPE_3D:

- $maxExtent.width \le VkPhysicalDeviceLimits.maxImageDimension3D$
- $\bullet \ maxExtent.height \leq VkPhysicalDeviceLimits.maxImageDimension3D$
- $maxExtent.depth \le VkPhysicalDeviceLimits.maxImageDimension3D$

Chapter 32

Debugging

Given the complexity of Vulkan there is a strong need for verbose debugging information to aid the application developer in tracking down errors in the application's use of Vulkan, particularly in combination with an external debugger or profiler.

This extension adds five new commands that allow a flexible way for debugging and validation layers to receive annotation and debug information. An application enables this extension by including "VK_EXT_debug_marker" in the list of <code>ppEnabledExtensionNames</code> at <code>vkCreateDevice</code>. The Vulkan implementation including layers indicating support will return VK_EXT_debug_marker in the extension list queried via <code>vkEnumerateDeviceExtensionProperties</code>.

Details of the commands are listed in two sections, "Object Annotation" and "Command Buffer Markers" respectively. Object Annotation allows the application to associate a name or binary data with a Vulkan object, while command buffer markers provide the developer with a way of associating logical elements of the scene with commands in the command buffer.

32.1 Object Annotation

The commands in this section allow application developers to associate user-defined information with Vulkan objects at will.

An object can be given a user-friendly name by calling:

- device is the device that created the object.
- pNameInfo is a pointer to an instance of the VkDebugMarkerObjectNameInfoEXT structure specifying the parameters of the name to set on the object.

Valid Usage

- device must be a valid VkDevice handle
- pNameInfo must be a pointer to a VkDebugMarkerObjectNameInfoEXT structure
- pNameInfo.object must be a Vulkan object

Host Synchronization

• Host access to pNameInfo.object must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDebugMarkerObjectNameInfoEXT structure is defined as:

- *sType* is the type of this structure and must be VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_NAME_INFO_EXT.
- pNext is NULL or a pointer to an extension-specific structure.
- objectType is the type of the object to be named. Object types are defined by VkDebugReportObjectTypeEXT.
- object is the object to be named.

• pObjectName is a null-terminated UTF-8 string specifying the name to apply to object.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_NAME_INFO_EXT
- pNext must be NULL
- objectType must be a valid VkDebugReportObjectTypeEXT value
- pObjectName must be a null-terminated string

Applications may change the name associated with an object simply by calling

vkDebugMarkerSetObjectNameEXT again with a new string. To remove a previously set name, *pName* should be set to an empty string.

In addition to setting a name for an object, debugging and validation layers may have uses for additional binary data on a per-object basis that has no other place in the Vulkan API. For example, a VkShaderModule could have additional debugging data attached to it to aid in offline shader tracing.

Arbitrary data can be attached to an object by calling:

- device is the device that created the object.
- pTagInfo is a pointer to an instance of the VkDebugMarkerObjectTagInfoEXT structure specifying the parameters of the tag to attach to the object.

Valid Usage

- device must be a valid VkDevice handle
- pTagInfo must be a pointer to a VkDebugMarkerObjectTagInfoEXT structure
- pTagInfo.object must be a Vulkan object
- pTagInfo.tagName must not be 0

Host Synchronization

• Host access to pTagInfo.object must be externally synchronized

Return Codes

Success

• VK_SUCCESS

Failure

- VK_ERROR_OUT_OF_HOST_MEMORY
- VK_ERROR_OUT_OF_DEVICE_MEMORY

The VkDebugMarkerObjectTagInfoEXT structure is defined as:

- *sType* is the type of this structure and must be VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_TAG_INFO_EXT.
- pNext is NULL or a pointer to an extension-specific structure.
- object Type is the type of the object to be named. Object types are defined by VkDebugReportObjectTypeEXT.
- object is the object to be tagged.
- tagName is a numerical identifier of the tag.
- tagSize is the number of bytes of data to attach to the object.
- pTaq is an array of taqSize bytes containing the data to be associated with the object.

Valid Usage

- stype must be VK STRUCTURE TYPE DEBUG MARKER OBJECT TAG INFO EXT
- pNext must be NULL
- objectType must be a valid VkDebugReportObjectTypeEXT value
- pTag must be a pointer to an array of tagSize bytes
- tagSize must be greater than 0

The tagName parameter gives a name or identifier to the type of data being tagged. This can be used by debugging layers to easily filter for only data that can be used by that implementation.

Both of the above structs contain a <code>objectType</code> member of type <code>VkDebugReportObjectTypeEXT</code> to determine what kind of object is passed in the <code>object</code> parameter. The enum is defined as:

```
typedef enum VkDebugReportObjectTypeEXT {
   VK_DEBUG_REPORT_OBJECT_TYPE_UNKNOWN_EXT = 0,
   VK_DEBUG_REPORT_OBJECT_TYPE_INSTANCE_EXT = 1,
   VK_DEBUG_REPORT_OBJECT_TYPE_PHYSICAL_DEVICE_EXT = 2,
   VK_DEBUG_REPORT_OBJECT_TYPE_DEVICE_EXT = 3,
   VK_DEBUG_REPORT_OBJECT_TYPE_QUEUE_EXT = 4,
   VK_DEBUG_REPORT_OBJECT_TYPE_SEMAPHORE_EXT = 5,
   VK_DEBUG_REPORT_OBJECT_TYPE_COMMAND_BUFFER_EXT = 6,
   VK DEBUG REPORT OBJECT TYPE FENCE EXT = 7,
   VK_DEBUG_REPORT_OBJECT_TYPE_DEVICE_MEMORY_EXT = 8,
   VK_DEBUG_REPORT_OBJECT_TYPE_BUFFER_EXT = 9,
   VK_DEBUG_REPORT_OBJECT_TYPE_IMAGE_EXT = 10,
   VK_DEBUG_REPORT_OBJECT_TYPE_EVENT_EXT = 11,
   VK_DEBUG_REPORT_OBJECT_TYPE_QUERY_POOL_EXT = 12,
   VK_DEBUG_REPORT_OBJECT_TYPE_BUFFER_VIEW_EXT = 13,
   VK_DEBUG_REPORT_OBJECT_TYPE_IMAGE_VIEW_EXT = 14,
   VK_DEBUG_REPORT_OBJECT_TYPE_SHADER_MODULE_EXT = 15,
   VK_DEBUG_REPORT_OBJECT_TYPE_PIPELINE_CACHE_EXT = 16,
   VK_DEBUG_REPORT_OBJECT_TYPE_PIPELINE_LAYOUT_EXT = 17,
   VK_DEBUG_REPORT_OBJECT_TYPE_RENDER_PASS_EXT = 18,
   VK_DEBUG_REPORT_OBJECT_TYPE_PIPELINE_EXT = 19,
   VK DEBUG REPORT OBJECT TYPE DESCRIPTOR SET LAYOUT EXT = 20,
   VK_DEBUG_REPORT_OBJECT_TYPE_SAMPLER_EXT = 21,
   VK_DEBUG_REPORT_OBJECT_TYPE_DESCRIPTOR_POOL_EXT = 22,
   VK_DEBUG_REPORT_OBJECT_TYPE_DESCRIPTOR_SET_EXT = 23,
   VK_DEBUG_REPORT_OBJECT_TYPE_FRAMEBUFFER_EXT = 24,
   VK_DEBUG_REPORT_OBJECT_TYPE_COMMAND_POOL_EXT = 25,
   VK_DEBUG_REPORT_OBJECT_TYPE_SURFACE_KHR_EXT = 26,
   VK_DEBUG_REPORT_OBJECT_TYPE_SWAPCHAIN_KHR_EXT = 27,
   VK_DEBUG_REPORT_OBJECT_TYPE_DEBUG_REPORT_EXT = 28,
} VkDebugReportObjectTypeEXT;
```

The possible values are:

- VK_DEBUG_REPORT_OBJECT_TYPE_UNKNOWN_EXT is an unknown object.
- VK_DEBUG_REPORT_OBJECT_TYPE_INSTANCE_EXT is a VkInstance.

- VK_DEBUG_REPORT_OBJECT_TYPE_PHYSICAL_DEVICE_EXT is a VkPhysicalDevice.
- VK_DEBUG_REPORT_OBJECT_TYPE_DEVICE_EXT is a VkDevice.
- VK_DEBUG_REPORT_OBJECT_TYPE_QUEUE_EXT is a VkQueue.
- VK_DEBUG_REPORT_OBJECT_TYPE_SEMAPHORE_EXT is a VkSemaphore.
- VK_DEBUG_REPORT_OBJECT_TYPE_COMMAND_BUFFER_EXT is a VkCommandBuffer.
- VK_DEBUG_REPORT_OBJECT_TYPE_FENCE_EXT is a VkFence.
- VK_DEBUG_REPORT_OBJECT_TYPE_DEVICE_MEMORY_EXT is a VkDeviceMemory.
- VK_DEBUG_REPORT_OBJECT_TYPE_BUFFER_EXT is a VkBuffer.
- VK_DEBUG_REPORT_OBJECT_TYPE_IMAGE_EXT is a VkImage.
- VK_DEBUG_REPORT_OBJECT_TYPE_EVENT_EXT is a VkEvent.
- VK_DEBUG_REPORT_OBJECT_TYPE_QUERY_POOL_EXT is a VkQueryPool.
- VK_DEBUG_REPORT_OBJECT_TYPE_BUFFER_VIEW_EXT is a VkBufferView.
- VK_DEBUG_REPORT_OBJECT_TYPE_IMAGE_VIEW_EXT is a VkImageView.
- VK_DEBUG_REPORT_OBJECT_TYPE_SHADER_MODULE_EXT is a VkShaderModule.
- VK DEBUG REPORT OBJECT TYPE PIPELINE CACHE EXT is a VkPipelineCache.
- VK_DEBUG_REPORT_OBJECT_TYPE_PIPELINE_LAYOUT_EXT is a VkPipelineLayout.
- VK_DEBUG_REPORT_OBJECT_TYPE_RENDER_PASS_EXT is a VkRenderPass.
- VK_DEBUG_REPORT_OBJECT_TYPE_PIPELINE_EXT is a VkPipeline.
- VK_DEBUG_REPORT_OBJECT_TYPE_DESCRIPTOR_SET_LAYOUT_EXT is a VkDescriptorSetLayout.
- VK DEBUG REPORT OBJECT TYPE SAMPLER EXT is a VkSampler.
- VK_DEBUG_REPORT_OBJECT_TYPE_DESCRIPTOR_POOL_EXT is a VkDescriptorPool.
- VK_DEBUG_REPORT_OBJECT_TYPE_DESCRIPTOR_SET_EXT is a VkDescriptorSet.
- VK_DEBUG_REPORT_OBJECT_TYPE_FRAMEBUFFER_EXT is a VkFramebuffer.
- VK_DEBUG_REPORT_OBJECT_TYPE_COMMAND_POOL_EXT is a VkCommandPool.
- VK_DEBUG_REPORT_OBJECT_TYPE_SURFACE_KHR_EXT is a VkSurfaceKHR.
- VK_DEBUG_REPORT_OBJECT_TYPE_SWAPCHAIN_KHR_EXT is a VkSwapchainKHR.
- VK_DEBUG_REPORT_OBJECT_TYPE_DEBUG_REPORT_EXT is a VkDebugReportCallbackEXT.

32.2 Command Buffer Markers

Typical Vulkan applications will submit many command buffers in each frame, with each command buffer containing a large number of individual commands. Being able to logically annotate regions of command buffers that belong together as well as hierarchically subdivide the frame is important to a developer's ability to navigate the commands viewed holistically.

The marker commands **vkCmdDebugMarkerBeginEXT** and **vkCmdDebugMarkerEndEXT** define regions of a series of commands that are grouped together, and they can be nested to create a hierarchy. The **vkCmdDebugMarkerInsertEXT** command allows insertion of a single label within a command buffer.

A marker region can be opened by calling:

- commandBuffer is the command buffer into which the command is recorded.
- pMarkerInfo is a pointer to an instance of the VkDebugMarkerMarkerInfoEXT structure specifying the parameters of the marker region to open.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pMarkerInfo must be a pointer to a VkDebuqMarkerMarkerInfoEXT structure
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

The VkDebugMarkerMarkerInfoEXT structure is defined as:

```
typedef struct VkDebugMarkerMarkerInfoEXT {
    VkStructureType    sType;
    const void*    pNext;
```

```
const char* pMarkerName;
float color[4];
} VkDebugMarkerMarkerInfoEXT;
```

- sType is the type of this structure and must be VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT.
- pNext is NULL or a pointer to an extension-specific structure.
- pMarkerName is a pointer to a null-terminated UTF-8 string that contains the name of the marker.
- color is an optional RGBA color value that can be associated with the marker. A particular implementation may choose to ignore this color value. The values contain RGBA values in order, in the range 0.0 to 1.0. If all elements in color are set to 0.0 then it is ignored.

Valid Usage

- sType must be VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT
- pNext must be NULL
- pMarkerName must be a null-terminated string

A marker region can be closed by calling:

• commandBuffer is the command buffer into which the command is recorded.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations
- There must be an outstanding vkCmdDebugMarkerBeginEXT command prior to the vkCmdDebugMarkerEndEXT on the queue that commandBuffer is submitted to
- If the matching vkCmdDebugMarkerBeginEXT command was in a secondary command buffer, the vkCmdDebugMarkerEndEXT must be in the same commandBuffer

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

An application may open a marker region in one command buffer and close it in another, or otherwise split marker regions across multiple command buffers or multiple queue submissions. When viewed from the linear series of submissions to a single queue, the calls to **vkCmdDebugMarkerBeginEXT** and **vkCmdDebugMarkerEndEXT** must be matched and balanced.

Any calls to **vkCmdDebugMarkerBeginEXT** within a secondary command buffer must have a matching **vkCmdDebugMarkerEndEXT** in that same command buffer, and marker regions begun outside of the secondary command buffer must not be ended inside it.

A single marker label can be inserted into a command buffer by calling:

- commandBuffer is the command buffer into which the command is recorded.
- pMarkerInfo is a pointer to an instance of the VkDebugMarkerMarkerInfoEXT structure specifying the parameters of the marker to insert.

Valid Usage

- commandBuffer must be a valid VkCommandBuffer handle
- pMarkerInfo must be a pointer to a VkDebugMarkerMarkerInfoEXT structure
- commandBuffer must be in the recording state
- The VkCommandPool that commandBuffer was allocated from must support graphics, or compute operations

Command Properties

Command Buffer Levels	Render Pass Scope	Supported Queue Types
Primary	Both	GRAPHICS
Secondary		COMPUTE

Chapter 33

Glossary

The terms defined in this section are used consistently throughout this Specification and may be used with or without capitalization.

Accessible (Descriptor Binding)

A descriptor binding is accessible to a shader stage if that stage is included in the <code>stageFlags</code> of the descriptor binding. Descriptors using that binding can only be used by stages in which they are accessible.

Adjacent Vertex

A vertex in an adjacency primitive topology that is not part of a given primitive, but is accessible in geometry shaders.

Aliased Range (Memory)

A range of a device memory allocation that is bound to multiple resources simultaneously.

API Order

A set of ordering rules that govern how primitives in draw commands affect the framebuffer.

Aspect (Image)

An image may contain multiple kinds, or aspects, of data for each pixel, where each aspect is used in a particular way by the pipeline and may be stored differently or separately from other aspects. For example, the color components of an image format make up the color aspect of the image, and may be used as a framebuffer color attachment. Some operations, like depth testing, operate only on specific aspects of an image. Others operations, like image/buffer copies, only operate on one aspect at a time.

Attachment (Render Pass)

A zero-based integer index name used in render pass creation to refer to a framebuffer attachment that is accessed by one or more subpasses. The index also refers to an attachment description which includes information about the properties of the image view that will later be attached.

Available

See Memory Dependency.

Back-Facing

See Facingness.

Ratch

A single structure submitted to a queue as part of a queue submission command, describing a set of queue operations to execute.

Backwards Compatibility

A given version of the API is backwards compatible with an earlier version if an application, relying only on valid behavior and functionality defined by the earlier specification, is able to correctly run against each version without any modification. This assumes no active attempt by that application to not run when it detects a different version.

Full Compatibility

A given version of the API is fully compatible with another version if an application, relying only on valid behavior and functionality defined by either of those specifications, is able to correctly run against each version without any modification. This assumes no active attempt by that application to not run when it detects a different version.

Binding (Memory)

An association established between a range of a resource object and a range of a memory object. These associations determine the memory locations affected by operations performed on elements of a resource object. Memory bindings are established using the <code>vkBindBufferMemory</code> command for non-sparse buffer objects, using the <code>vkBindImageMemory</code> command for non-sparse image objects, and using the <code>vkQueueBindSparse</code> command for sparse resources.

Blend Constant

Four floating point (RGBA) values used as an input to blending.

Blending

Arithmetic operations between a fragment color value and a value in a color attachment that produce a final color value to be written to the attachment.

Buffer

A resource that represents a linear array of data in device memory. Represented by a VkBuffer object.

Buffer View

An object that represents a range of a specific buffer, and state that controls how the contents are interpreted. Represented by a VkBufferView object.

Built-In Variable

A variable decorated in a shader, where the decoration makes the variable take values provided by the execution environment or values that are generated by fixed-function pipeline stages.

Built-In Interface Block

A block defined in a shader that contains only variables decorated with built-in decorations, and is used to match against other shader stages.

Clip Coordinates

The homogeneous coordinate space that vertex positions (**Position** decoration) are written in by vertex processing stages.

Clip Distance

A built-in output from vertex processing stages that defines a clip half-space against which the primitive is clipped.

Clip Volume

The intersection of the view volume with all clip half-spaces.

Color Attachment

A subpass attachment point, or image view, that is the target of fragment color outputs and blending.

Combined Image Sampler

A descriptor type that includes both a sampled image and a sampler.

Command Buffer

An object that records commands to be submitted to a queue. Represented by a VkCommandBuffer object.

Command Pool

An object that command buffer memory is allocated from, and that owns that memory. Command pools aid multithreaded performance by enabling different threads to use different allocators, without internal synchronization on each use. Represented by a VkCommandPool object.

Compatible Allocator

When allocators are compatible, allocations from each allocator can be freed by the other allocator.

Compatible Image Formats

When formats are compatible, images created with one of the formats can have image views created from it using any of the compatible formats.

Compatible Queues

Queues within a queue family. Compatible queues have identical properties.

Component (Format)

A distinct part of a format. Depth, stencil, and color channels (e.g. R, G, B, A), are all separate components.

Compressed Texel Block

An element of an image having a block-compressed format, comprising a rectangular block of texel values that are encoded as a single value in memory. Compressed texel blocks of a particular block-compressed format have a corresponding width, height, and depth that define the dimensions of these elements in units of texels, and a size in bytes of the encoding in memory.

Cull Distance

A built-in output from vertex processing stages that defines a cull half-space where the primitive is rejected if all vertices have a negative value for the same cull distance.

Cull Volume

The intersection of the view volume with all cull half-spaces.

Decoration (SPIR-V)

Auxiliary information such as built-in variables, stream numbers, invariance, interpolation type, relaxed precision, etc., added to variables or structure-type members through decorations.

Depth/Stencil Attachment

A subpass attachment point, or image view, that is the target of depth and/or stencil test operations and writes.

Depth/Stencil Format

A VkFormat that includes depth and/or stencil components.

Depth/Stencil Image (or ImageView)

A VkImage (or VkImageView) with a depth/stencil format.

Derivative Group

A set of fragment shader invocations that cooperate to compute derivatives, including implicit derivatives for sampled image operations.

Descriptor

Information about a resource or resource view written into a descriptor set that is used to access the resource or view from a shader.

Descriptor Binding

An entry in a descriptor set layout corresponding to zero or more descriptors of a single descriptor type in a set. Defined by a VkDescriptorSetLayoutBinding structure.

Descriptor Pool

An object that descriptor sets are allocated from, and that owns the storage of those descriptor sets. Descriptor pools aid multithreaded performance by enabling different threads to use different allocators, without internal synchronization on each use. Represented by a VkDescriptorPool object.

Descriptor Set

An object that resource descriptors are written into via the API, and that can be bound to a command buffer such that the descriptors contained within it can be accessed from shaders. Represented by a VkDescriptorSet object.

Descriptor Set Layout

An object that defines the set of resources (types and counts) and their relative arrangement (in the binding namespace) within a descriptor set. Used when allocating descriptor sets and when creating pipeline layouts. Represented by a VkDescriptorSetLayout object.

Device

The processor(s) and execution environment that perform tasks requested by the application via the Vulkan API.

Device Memory

Memory accessible to the device. Represented by a VkDeviceMemory object.

Device-Level Object

Logical device objects and their child objects For example, VkDevice, VkQueue, and VkCommandBuffer objects are device-level objects.

Device-Local Memory

Memory that is connected to the device, and may be more performant for device access than host-local memory.

Dispatchable Handle

A handle of a pointer handle type which may be used by layers as part of intercepting API commands. The first argument to each Vulkan command is a dispatchable handle type.

Dispatching Commands

Commands that provoke work using a compute pipeline. Includes vkCmdDispatch and vkCmdDispatchIndirect.

Drawing Commands

 $\label{local_commands} Commands \ that \ provoke \ work \ using \ a \ graphics \ pipeline. \ Includes \ vkCmdDraw, \ vkCmdDrawIndexed, \ vkCmdDrawIndexedIndirect.$

Duration (Command)

The duration of a Vulkan command refers to the interval between calling the command and its return to the caller.

Dynamic Storage Buffer

A storage buffer whose offset is specified each time the storage buffer is bound to a command buffer via a descriptor set.

Dynamic Uniform Buffer

A uniform buffer whose offset is specified each time the uniform buffer is bound to a command buffer via a descriptor set.

Explicitly-Enabled Layer

A layer enabled by the application by adding it to the enabled layer list in vkCreateInstance or vkCreateDevice.

Event

A synchronization primitive that is signaled when execution of previous commands complete through a specified set of pipeline stages. Events can be waited on by the device and polled by the host. Represented by a VkEvent object.

Executable State (Command Buffer)

A command buffer that has ended recording commands and can be executed. See also Initial State and Recording State

Execution Dependency

A dependency that guarantees that certain pipeline stages' work for a first set of commands has completed execution before certain pipeline stages' work for a second set of commands begins execution. This is accomplished via pipeline barriers, subpass dependencies, events, or implicit ordering operations.

Execution Dependency Chain

A sequence of execution dependencies that transitively act as an execution dependency.

External synchronization

A type of synchronization required of the application, where parameters defined to be externally synchronized must not be used simultaneously in multiple threads.

Facingness (Polygon)

A classification of a polygon as either front-facing or back-facing, depending on the orientation (winding order) of its vertices.

Fence

A synchronization primitive that is signaled when a set of batches or sparse binding operations complete execution on a queue. Fences can be waited on by the host. Represented by a VkFence object.

Flat Shading

A property of a vertex attribute that causes the value from a single vertex (the provoking vertex) to be used for all vertices in a primitive, and for interpolation of that attribute to return that single value unaltered.

Fragment Input Attachment Interface

A fragment shader entry point's variables with **UniformConstant** storage class and a decoration of **InputAttachmentIndex**, which receive values from input attachments.

Fragment Output Interface

A fragment shader entry point's variables with **Output** storage class, which output to color and/or depth/stencil attachments.

Framebuffer

A collection of image views and a set of dimensions that, in conjunction with a render pass, define the inputs and outputs used by drawing commands. Represented by a VkFramebuffer object.

Framebuffer Attachment

One of the image views used in a framebuffer.

Framebuffer Coordinates

A coordinate system in which adjacent pixels' coordinates differ by 1 in x and/or y, with (0,0) in the upper left corner and pixel centers at half-integers.

Front-Facing

See Facingness.

Global Workgroup

A collection of local workgroups dispatched by a single dispatch command.

Handle

An opaque integer or pointer value used to refer to a Vulkan object. Each object type has a unique handle type.

Happen-after

A command happens-after a dependency if they are separated by an execution dependency chain, with the command included in the destination of the last dependency of the chain. A memory barrier makes visible memory writes to commands that happen-after it.

Happen-before

A command happens-before a dependency if they are separated by an execution dependency chain, with the command included in the source of the first dependency of the chain. A memory barrier makes available memory writes of commands that happen-before it.

Helper Invocation

A fragment shader invocation that is created solely for the purposes of evaluating derivatives for use in non-helper fragment shader invocations, and which does not have side effects.

Host

The processor(s) and execution environment that the application runs on, and that the Vulkan API is exposed on.

Host Memory

Memory not accessible to the device, used to store implementation data structures.

Host-Accessible Subresource

A buffer, or a linear image subresource in either the VK_IMAGE_LAYOUT_PREINITIALIZED or VK_IMAGE_LAYOUT_GENERAL layout. Host-accessible subresources have a well-defined addressing scheme which can be used by the host.

Host-Local Memory

Memory that is not local to the device, and may be less performant for device access than device-local memory.

Host-Visible Memory

Device memory that can be mapped on the host and can be read and written by the host.

Image

A resource that represents a multi-dimensional formatted interpretation of device memory. Represented by a VkImage object.

Image Subresource

A specific mipmap level and layer of an image.

Image Subresource Range

A set of image subresources that are contiguous mipmap levels and layers.

Image View

An object that represents an image subresource range of a specific image, and state that controls how the contents are interpreted. Represented by a VkImageView object.

Immutable Sampler

A sampler descriptor provided at descriptor set layout creation time, and that is used for that binding in all descriptor sets allocated from the layout, and cannot be changed.

Implicitly-Enabled Layer

A layer enabled by a loader-defined mechanism outside the Vulkan API, rather than explicitly by the application during instance or device creation.

Index Buffer

A buffer bound via vkCmdBindIndexBuffer which is the source of index values used to fetch vertex attributes for a vkCmdDrawIndexed or vkCmdDrawIndexedIndirect command.

Indirect Commands

Drawing or dispatching commands that source some of their parameters from structures in buffer memory. Includes vkCmdDrawIndirect, vkCmdDrawIndexedIndirect, and vkCmdDispatchIndirect.

Initial State (Command Buffer)

A command buffer that has not begun recording commands. See also Recorded State and Executable State.

Input Attachment

A descriptor type that represents an image view, and supports unfiltered read-only access in a shader, only at the fragment's location in the view.

Instance

The top-level Vulkan object, which represents the application's connection to the implementation. Represented by a VkInstance object.

Instance-Level Object

High-level Vulkan objects, which are not logical devices, nor children of logical devices. For example, VkInstance and VkPhysicalDevice objects are instance-level objects.

Internal Synchronization

A type of synchronization required of the implementation, where parameters not defined to be externally synchronized may require internal mutexing to avoid multithreaded race conditions.

Invocation (Shader)

A single execution of an entry point in a SPIR-V module. For example, a single vertex's execution of a vertex shader or a single fragment's execution of a fragment shader.

Invocation Group

A set of shader invocations that are executed in parallel and that must execute the same control flow path in order for control flow to be considered dynamically uniform.

Local Workgroup

A collection of compute shader invocations invoked by a single dispatch command, which share shared memory and can synchronize with each other.

Logical Device

An object that represents the application's interface to the physical device. The logical device is the parent of most Vulkan objects. Represented by a VkDevice object.

Logical Operation

Bitwise operations between a fragment color value and a value in a color attachment, that produce a final color value to be written to the attachment.

Lost Device

A state that a logical device may be in as a result of hardware errors or other exceptional conditions.

Mappable

See Host-Visible Memory.

Memory Dependency

A sequence of operations that makes writes available, performs an execution dependency between the writes and subsequent accesses, and makes available writes visible to later accesses. In order for the effects of a write to be coherent with later accesses, it must be made available from the old access type and then made visible to the new access type.

Memory Heap

A region of memory from which device memory allocations can be made.

Memory Type

An index used to select a set of memory properties (e.g. mappable, cached) for a device memory allocation.

Mip Tail Region

The set of mipmap levels of a sparse residency texture that are too small to fill a sparse block, and that must all be bound to memory collectively and opaquely.

Non-Dispatchable Handle

A handle of an integer handle type. Handle values may not be unique, even for two objects of the same type.

Normalized

A value that is interpreted as being in the range [0,1] as a result of being implicitly divided by some other value.

Normalized Device Coordinates

A coordinate space after perspective division is applied to clip coordinates, and before the viewport transformation converts to framebuffer coordinates.

Overlapped Range (Aliased Range)

The aliased range of a device memory allocation that intersects a given image subresource of an image or range of a buffer.

Packed Format

A format whose components are stored as a single data element in memory, with their relative locations defined within that element.

Physical Device

An object that represents a single device in the system. Represented by a VkPhysicalDevice object.

Pipeline

An object that controls how graphics or compute work is executed on the device. A pipeline includes one or more shaders, as well as state controlling any non-programmable stages of the pipeline. Represented by a VkPipeline object.

Pipeline Barrier

An execution and/or memory dependency recorded as an explicit command in a command buffer, that forms a dependency between the previous and subsequent commands.

Pipeline Cache

An object that can be used to collect and retrieve information from pipelines as they are created, and can be populated with previously retrieved information in order to accelerate pipeline creation. Represented by a VkPipelineCache object.

Pipeline Layout

An object that defines the set of resources (via a collection of descriptor set layouts) and push constants used by pipelines that are created using the layout. Used when creating a pipeline and when binding descriptor sets and setting push constant values. Represented by a VkPipelineLayout object.

Point Sampling (Rasterization)

A rule that determines whether a fragment sample location is covered by a polygon primitive by testing whether the sample location is in the interior of the polygon in framebuffer-space, or on the boundary of the polygon according to the tie-breaking rules.

Preserve Attachment

One of a list of attachments in a subpass description that is not read or written by the subpass, but that is read or written on earlier and later subpasses and whose contents must be preserved through this subpass.

Primary Command Buffer

A command buffer that can execute secondary command buffers, and can be submitted directly to a queue.

Primitive Topology

State that controls how vertices are assembled into primitives, e.g. as lists of triangles, strips of lines, etc..

Provoking Vertex

The vertex in a primitive from which flat shaded attribute values are taken. This is generally the "first" vertex in the primitive, and depends on the primitive topology.

Push Constants

A small bank of values writable via the API and accessible in shaders. Push constants allow the application to set values used in shaders without creating buffers or modifying and binding descriptor sets for each update.

Push Constant Interface

The set of variables with **PushConstant** storage class that are statically used by a shader entry point, and which receive values from push constant commands.

Query Pool

An object that contains a number of query entries and their associated state and results. Represented by a VkQueryPool object.

Queue

An object that executes command buffers and sparse binding operations on a device. Represented by a VkQueue object.

Queue Family

A set of queues that have common properties and support the same functionality, as advertised in VkQueueFamilyProperties.

Queue Operation

A unit of work to be executed by a specific queue on a device, submitted via a queue submission command. Each queue submission command details the specific queue operations that occur as a result of calling that command. Queue operations typically include work that is specific to each command, and synchronization tasks.

Queue Submission

Zero or more batches and an optional fence to be signaled, passed to a command for execution on a queue. See the Devices and Queues chapter for more information.

Recording State (Command Buffer)

A command buffer that is ready to record commands. See also Initial State and Executable State.

Render Pass

An object that represents a set of framebuffer attachments and phases of rendering using those attachments. Represented by a VkRenderPass object.

Render Pass Instance

A use of a render pass in a command buffer.

Reset (Command Buffer)

Resetting a command buffer discards any previously recorded commands and puts a command buffer in the initial state.

Residency Code

An integer value returned by sparse image instructions, indicating whether any sparse unbound texels were accessed.

Resolve Attachment

A subpass attachment point, or image view, that is the target of a multisample resolve operation from the corresponding color attachment at the end of the subpass.

Sampled Image

A descriptor type that represents an image view, and supports filtered (sampled) and unfiltered read-only access in a shader.

Sampler

An object that contains state that controls how sampled image data is sampled (or filtered) when accessed in a shader. Also a descriptor type describing the object. Represented by a VkSampler object.

Secondary Command Buffer

A command buffer that can be executed by a primary command buffer, and must not be submitted directly to a queue.

Self-Dependency

A subpass dependency from a subpass to itself, i.e. with <code>srcSubpass</code> equal to <code>dstSubpass</code>. A self-dependency is not automatically performed during a render pass instance, rather a subset of it can be performed via <code>vkCmdPipelineBarrier</code> during the subpass.

Semaphore

A synchronization primitive that supports signal and wait operations, and can be used to synchronize operations within a queue or across queues. Represented by a VkSemaphore object.

Shader

Instructions selected (via an entry point) from a shader module, which are executed in a shader stage.

Shader Code

A stream of instructions used to describe the operation of a shader.

Shader Module

A collection of shader code, potentially including several functions and entry points, that is used to create shaders in pipelines. Represented by a VkShaderModule object.

Shader Stage

A stage of the graphics or compute pipeline that executes shader code.

Side Effect

A store to memory or atomic operation on memory from a shader invocation.

Sparse Block

An element of a sparse resource that can be independently bound to memory. Sparse blocks of a particular sparse resource have a corresponding size in bytes that they use in the bound memory.

Sparse Image Block

A sparse block in a sparse partially-resident image. In addition to the sparse block size in bytes, sparse image blocks have a corresponding width, height, and depth that define the dimensions of these elements in units of texels or compressed texel blocks, the latter being used in case of sparse images having a block-compressed format.

Sparse Unbound Texel

A texel read from a region of a sparse texture that does not have memory bound to it.

Static Use

An object in a shader is statically used by a shader entry point if any function in the entry point's call tree contains an instruction using the object. Static use is used to constrain the set of descriptors used by a shader entry point.

Storage Buffer

A descriptor type that represents a buffer, and supports reads, writes, and atomics in a shader.

Storage Image

A descriptor type that represents an image view, and supports unfiltered loads, stores, and atomics in a shader.

Storage Texel Buffer

A descriptor type that represents a buffer view, and supports unfiltered, formatted reads, writes, and atomics in a shader.

Subpass

A phase of rendering within a render pass, that reads and writes a subset of the attachments.

Subpass Dependency

An execution and/or memory dependency between two subpasses described as part of render pass creation, and automatically performed between subpasses in a render pass instance. A subpass dependency limits the overlap of execution of the pair of subpasses, and can provide guarantees of memory coherence between accesses in the subpasses.

Subpass Description

Lists of attachment indices for input attachments, color attachments, depth/stencil attachment, resolve attachments, and preserve attachments used by the subpass in a render pass.

Subset (Self-Dependency)

A subset of a self-dependency is a pipeline barrier performed during the subpass of the self-dependency, and whose stage masks and access masks each contain a subset of the bits set in the identically named mask in the self-dependency.

Texel Coordinate System

One of three coordinate systems (normalized, unnormalized, integer) that define how texel coordinates are interpreted in an image or a specific mipmap level of an image.

Uniform Texel Buffer

A descriptor type that represents a buffer view, and supports unfiltered, formatted, read-only access in a shader.

Uniform Buffer

A descriptor type that represents a buffer, and supports read-only access in a shader.

Unnormalized

A value that is interpreted according to its conventional interpretation, and is not normalized.

User-Defined Variable Interface

A shader entry point's variables with **Input** or **Output** storage class that are not built-in variables.

Vertex Input Attribute

A graphics pipeline resource that produces input values for the vertex shader by reading data from a vertex input binding and converting it to the attribute's format.

Vertex Input Binding

A graphics pipeline resource that is bound to a buffer and includes state that affects addressing calculations within that buffer.

Vertex Input Interface

A vertex shader entry point's variables with **Input** storage class, which receive values from vertex input attributes.

Vertex Processing Stages

A set of shader stages that comprises the vertex shader, tessellation control shader, tessellation evaluation shader, and geometry shader stages.

View Volume

A subspace in homogeneous coordinates, corresponding to post-projection x and y values between -1 and +1, and z values between 0 and +1.

Viewport Transformation

A transformation from normalized device coordinates to framebuffer coordinates, based on a viewport rectangle and depth range.

Visible

See Memory Dependency.

Chapter 34

Src

G

Common Abbreviations

Abbreviations and acronyms are sometimes used in the Specification and the API where they are considered clear and commonplace, and are defined here:

Source **Dst** Destination Min Minimum Max Maximum Rect Rectangle Info Information LOD Level of Detail ID Identifier **UUID** Universally Unique Identifier Op Operation R Red color component

Green color component

В

Blue color component

A

Alpha color component

Chapter 35

Prefixes

Prefixes are used in the API to denote specific semantic meaning of Vulkan names, or as a label to avoid name clashes, and are explained here:

VK/Vk/vk

Vulkan namespace

All types, commands, enumerants and defines in this specification are prefixed with these two characters.

PFN/pfn

Function Pointer

Denotes that a type is a function pointer, or that a variable is of a pointer type.

p

Pointer

Variable is a pointer.

vkCmd

Commands that record commands in command buffers

These API commands do not result in immediate processing on the device. Instead, they record the requested action in a command buffer for execution when the command buffer is submitted to a queue.

S

Structure

Used to denote the VK_STRUCTURE_TYPE* member of each structure in sType

Appendix A

Vulkan Environment for SPIR-V

Shaders for Vulkan are defined by the Khronos SPIR-V Specification as well as the Khronos SPIR-V Extended Instructions for GLSL Specification. This appendix defines additional SPIR-V requirements applying to Vulkan shaders.

A.1 Required Versions and Formats

A Vulkan 1.0 implementation must support the 1.0 version of SPIR-V and the 1.0 version of the SPIR-V Extended Instructions for GLSL.

A SPIR-V module passed into vkCreateShaderModule is interpreted as a series of 32-bit words in host endianness, with literal strings packed as described in section 2.2 of the SPIR-V Specification. The first few words of the SPIR-V module must be a magic number and a SPIR-V version number, as described in section 2.3 of the SPIR-V Specification.

A.2 Capabilities

Implementations must support the following capability operands declared by **OpCapability**:

- Matrix
- Shader
- InputAttachment
- Sampled1D
- Image1D
- SampledBuffer
- ImageBuffer
- ImageQuery
- DerivativeControl

Implementations may support features that are not required by the Specification, as described in the Features chapter. If such a feature is supported, then any capability operand(s) corresponding to that feature must also be supported.

Table A.1: SPIR-V Capabilities which are not required, and corresponding feature names

SPIR-V OpCapability	Vulkan feature name
Geometry	geometryShader
Tessellation	tessellationShader
Float64	shaderFloat64
Int64	shaderInt64
Int16	shaderInt16
TessellationPointSize	shaderTessellationAndGeometryPointSize
GeometryPointSize	shaderTessellationAndGeometryPointSize
ImageGatherExtended	shaderImageGatherExtended
StorageImageMultisample	shaderStorageImageMultisample
UniformBufferArrayDynamicIndexing	shaderUniformBufferArrayDynamicIndexing
SampledImageArrayDynamicIndexing	shaderSampledImageArrayDynamicIndexing
StorageBufferArrayDynamicIndexing	shaderStorageBufferArrayDynamicIndexing
StorageImageArrayDynamicIndexing	shaderStorageImageArrayDynamicIndexing
ClipDistance	shaderClipDistance
CullDistance	shaderCullDistance
ImageCubeArray	imageCubeArray
SampleRateShading	sampleRateShading
SparseResidency	shaderResourceResidency
MinLod	shaderResourceMinLod
SampledCubeArray	imageCubeArray
ImageMSArray	shaderStorageImageMultisample
StorageImageExtendedFormats	shaderStorageImageExtendedFormats
InterpolationFunction	sampleRateShading
StorageImageReadWithoutFormat	shaderStorageImageReadWithoutFormat
StorageImageWriteWithoutFormat	shaderStorageImageWriteWithoutFormat
MultiViewport	multiViewport

The application can pass a SPIR-V module to vkCreateShaderModule that uses the SPV_AMD_shader_explicit_vertex_parameter SPIR-V extension.

The application can pass a SPIR-V module to vkCreateShaderModule that uses the SPV_AMD_gcn_shader SPIR-V extension.

The application can pass a SPIR-V module to vkCreateShaderModule that uses the SPV_AMD_shader_trinary_minmax SPIR-V extension.

The application must not pass a SPIR-V module containing any of the following to vkCreateShaderModule:

- any OpCapability not listed above,
- an unsupported capability, or
- a capability which corresponds to a Vulkan feature which has not been enabled.

A.3 Validation Rules within a Module

A SPIR-V module passed to vkCreateShaderModule must conform to the following rules:

- Every entry point must have no return value and accept no arguments.
- Recursion: The static function-call graph for an entry point must not contain cycles.
- The **Logical** addressing model must be selected.
- Scope for execution must be limited to:
 - Workgroup
 - Subgroup
- Scope for memory must be limited to:
 - Device
 - Workgroup
 - Invocation
- The OriginLowerLeft execution mode must not be used; fragment entry points must declare OriginUpperLeft.
- The PixelCenterInteger execution mode must not be used. Pixels are always centered at half-integer coordinates.
- · Images
 - OpTypeImage must declare a scalar 32-bit float or 32-bit integer type for the "Sampled Type".
 (RelaxedPrecision can be applied to a sampling instruction and to the variable holding the result of a sampling instruction.)
 - **OpSampledImage** must only consume an "Image" operand whose type has its "Sampled" operand set to 1.
 - The "(u, v)" coordinates used for a **SubpassData** must be the <id> of a constant vector (0, 0), or if a layer coordinate is used, must be a vector that was formed with constant 0 for the "u" and "v" components.
 - The "Depth" operand of **OpTypeImage** is ignored.
- Decorations
 - The GLSLShared and GLSLPacked decorations must not be used.
 - The **Flat**, **NoPerspective**, **Sample**, and **Centroid** decorations must not be used on variables with storage class other than **Input** or on variables used in the interface of non-fragment shader entry points.
 - The Patch decoration must not be used on variables in the interface of a vertex, geometry, or fragment shader stage's entry point.
- OpTypeRuntimeArray must only be used for the last member of an OpTypeStruct in the Uniform storage class.
- Linkage: See Shader Interfaces for additional linking and validation rules.
- · Compute Shaders
 - For each compute shader entry point, either a LocalSize execution mode or an object decorated with the WorkgroupSize decoration must be specified.

A.4 Precision and Operation of SPIR-V Instructions

The following rules apply to both single and double-precision floating point instructions:

- Positive and negative infinities and positive and negative zeros are generated as dictated by IEEE 754, but subject to the precisions allowed in the following table.
- Dividing a non-zero by a zero results in the appropriately signed IEEE 754 infinity.
- Any denormalized value input into a shader or potentially generated by any instruction in a shader may be flushed to 0.
- The rounding mode cannot be set and is undefined.
- NaNs may not be generated. Instructions that operate on a NaN may not result in a NaN.
- Support for signaling NaNs is optional and exceptions are never raised.

The precision of double-precision instructions is at least that of single precision. For single precision (32 bit) instructions, precisions are required to be at least as follows, unless decorated with RelaxedPrecision:

Table A.2: Precision of core SPIR-V Instructions

Instruction	Precision
OpFAdd	Correctly rounded.
OpFSub	Correctly rounded.
OpFMul	Correctly rounded.
OpFOrdEqual, OpFUnordEqual	Correct result.
OpFOrdLessThan, OpFUnordLessThan	Correct result.
OpFOrdGreaterThan, OpFUnordGreaterThan	Correct result.
OpFOrdLessThanEqual,	Correct result.
OpFUnordLessThanEqual	
OpFOrdGreaterThanEqual,	Correct result.
OpFUnordGreaterThanEqual	
OpFDiv	2.5 ULP for b in the range [2 ⁻¹²⁶ , 2 ¹²⁶].
conversions between types	Correctly rounded.

Table A.3: Precision of GLSL.std.450 Instructions

Instruction	Precision
fma()	Inherited from OpFMul followed by OpFAdd.
exp(x), exp2(x)	$(3+2\times x)$ ULP.
log(), log2()	3 ULP outside the range [0.5, 2.0]. Absolute error <
	2^{-21} inside the range [0.5, 2.0].
pow(x, y)	Inherited from $exp2(y \times log2(x))$.
sqrt()	Inherited from 1.0 / inversesqrt().
<pre>inversesqrt()</pre>	2 ULP.

GLSL.std.450 extended instructions specifically defined in terms of the above instructions inherit the above errors. GLSL.std.450 extended instructions not listed above and not defined in terms of the above have undefined precision. These include, for example, the trigonometric functions and determinant.

For the **OpSRem** and **OpSMod** instructions, if either operand is negative the result is undefined.



Note

While the **OpSRem** and **OpSMod** instructions are supported by the Vulkan environment, they require non-negative values and thus do not enable additional functionality beyond what **OpUMod** provides.

Compatibility Between SPIR-V Image Formats And Vulkan Formats

Images which are read from or written to by shaders must have SPIR-V image formats compatible with the Vulkan image formats backing the image under the circumstances described for texture image validation. The compatibile formats are:

Table A.4: SPIR-V and Vulkan Image Format Compatibility

SPIR-V Image Format	Compatible Vulkan Format
Rgba32f	VK_FORMAT_R32G32B32A32_SFLOAT
Rgba16f	VK_FORMAT_R16G16B16A16_SFLOAT
R32f	VK_FORMAT_R32_SFLOAT
Rgba8	VK_FORMAT_R8G8B8A8_UNORM
Rgba8Snorm	VK_FORMAT_R8G8B8A8_SNORM
Rg32f	VK_FORMAT_R32G32_SFLOAT
Rg16f	VK_FORMAT_R16G16_SFLOAT
R11fG11fB10f	VK_FORMAT_B10G11R11_UFLOAT_PACK32
R16f	VK_FORMAT_R16_SFLOAT
Rgba16	VK_FORMAT_R16G16B16A16_UNORM
Rgb10A2	VK_FORMAT_A2B10G10R10_UNORM_PACK32
Rg16	VK_FORMAT_R16G16_UNORM
Rg8	VK_FORMAT_R8G8_UNORM
R16	VK_FORMAT_R16_UNORM
R8	VK_FORMAT_R8_UNORM
Rgba16Snorm	VK_FORMAT_R16G16B16A16_SNORM
Rg16Snorm	VK_FORMAT_R16G16_SNORM
Rg8Snorm	VK_FORMAT_R8G8_SNORM
R16Snorm	VK_FORMAT_R16_SNORM
R8Snorm	VK_FORMAT_R8_SNORM
Rgba32i	VK_FORMAT_R32G32B32A32_SINT
Rgba16i	VK_FORMAT_R16G16B16A16_SINT
Rgba8i	VK_FORMAT_R8G8B8A8_SINT
R32i	VK_FORMAT_R32_SINT
Rg32i	VK_FORMAT_R32G32_SINT
Rg16i	VK_FORMAT_R16G16_SINT
Rg8i	VK_FORMAT_R8G8_SINT
R16i	VK_FORMAT_R16_SINT
R8i	VK_FORMAT_R8_SINT
Rgba32ui	VK_FORMAT_R32G32B32A32_UINT
Rgba16ui	VK_FORMAT_R16G16B16A16_UINT
Rgba8ui	VK_FORMAT_R8G8B8A8_UINT

Table A.4: (continued)

SPIR-V Image Format	Compatible Vulkan Format
R32ui	VK_FORMAT_R32_UINT
Rgb10a2ui	VK_FORMAT_A2B10G10R10_UINT_PACK32
Rg32ui	VK_FORMAT_R32G32_UINT
Rg16ui	VK_FORMAT_R16G16_UINT
Rg8ui	VK_FORMAT_R8G8_UINT
R16ui	VK_FORMAT_R16_UINT
R8ui	VK_FORMAT_R8_UINT

Appendix B

Compressed Image Formats

The compressed texture formats used by Vulkan are described in the specifically identified sections of the Khronos Data Format Specification, version 1.1.

Unless otherwise described, the quantities encoded in these compressed formats are treated as normalized, unsigned values.

Those formats listed as sRGB-encoded have in-memory representations of R, G and B components which are nonlinearly-encoded as R', G', and B'; any alpha component is unchanged. As part of filtering, the nonlinear R', G', and B' values are converted to linear R, G, and B components; any alpha component is unchanged. The conversion between linear and nonlinear encoding is performed as described in the "KHR_DF_TRANSFER_SRGB" section of the Khronos Data Format Specification.

B.1 Block-Compressed Image Formats

Table B.1: Mapping of Vulkan BC formats to descriptions

VkFormat	Khronos Data Format Specification	
	description	
Formats described in the "S3TC Compressed		
VK_FORMAT_BC1_RGB_UNORM_BLOCK	BC1 with no alpha	
VK_FORMAT_BC1_RGB_SRGB_BLOCK	BC1 with no alpha, sRGB-encoded	
VK_FORMAT_BC1_RGBA_UNORM_BLOCK	BC1 with alpha	
VK_FORMAT_BC1_RGBA_SRGB_BLOCK	BC1 with alpha, sRGB-encoded	
VK_FORMAT_BC2_UNORM_BLOCK	BC2	
VK_FORMAT_BC2_SRGB_BLOCK	BC2, sRGB-encoded	
VK_FORMAT_BC3_UNORM_BLOCK	BC3	
VK_FORMAT_BC3_SRGB_BLOCK	BC3, sRGB-encoded	
Formats described in the "RGTC Compresse	d Texture Image Formats" chapter	
VK_FORMAT_BC4_UNORM_BLOCK	BC4 unsigned	
VK_FORMAT_BC4_SNORM_BLOCK	BC4 signed	
VK_FORMAT_BC5_UNORM_BLOCK	BC5 unsigned	
VK_FORMAT_BC5_SNORM_BLOCK	BC5 signed	
Formats described in the "BPTC Compressed Texture Image Formats" chapter		
VK_FORMAT_BC6H_UFLOAT_BLOCK	BC6H (unsigned version)	
VK_FORMAT_BC6H_SFLOAT_BLOCK	BC6H (signed version)	
VK_FORMAT_BC7_UNORM_BLOCK	BC7	
VK_FORMAT_BC7_SRGB_BLOCK	BC7, sRGB-encoded	

B.2 ETC Compressed Image Formats

The following formats are described in the "ETC2 Compressed Texture Image Formats" chapter of the Khronos Data Format Specification.

Table B.2: Mapping of Vulkan ETC formats to descriptions

VkFormat	Khronos Data Format Specification description	
VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK	RGB ETC2	
VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK	RGB ETC2 with sRGB encoding	
VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK	RGB ETC2 with punch-through alpha	
VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK	RGB ETC2 with punch-through alpha and sRGB	
VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK	RGBA ETC2	
VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK	RGBA ETC2 with sRGB encoding	
VK_FORMAT_EAC_R11_UNORM_BLOCK	Unsigned R11 EAC	
VK_FORMAT_EAC_R11_SNORM_BLOCK	Signed R11 EAC	
VK_FORMAT_EAC_R11G11_UNORM_BLOCK	Unsigned RG11 EAC	
VK_FORMAT_EAC_R11G11_SNORM_BLOCK	Signed RG11 EAC	

B.3 ASTC Compressed Image Formats

ASTC formats are described in the "ASTC Compressed Texture Image Formats" chapter of the Khronos Data Format Specification.

Table B.3: Mapping of Vulkan ASTC formats to descriptions

VkFormat	Compressed texel block dimen- sions	sRGB-encoded
VK_FORMAT_ASTC_4x4_UNORM_BLOCK	4×4	No
VK_FORMAT_ASTC_4x4_SRGB_BLOCK	4×4	Yes
VK_FORMAT_ASTC_5x4_UNORM_BLOCK	5 × 4	No
VK_FORMAT_ASTC_5x4_SRGB_BLOCK	5 × 4	Yes
VK_FORMAT_ASTC_5x5_UNORM_BLOCK	5 × 5	No
VK_FORMAT_ASTC_5x5_SRGB_BLOCK	5 × 5	Yes
VK_FORMAT_ASTC_6x5_UNORM_BLOCK	6×5	No
VK_FORMAT_ASTC_6x5_SRGB_BLOCK	6×5	Yes
VK_FORMAT_ASTC_6x6_UNORM_BLOCK	6×6	No
VK_FORMAT_ASTC_6x6_SRGB_BLOCK	6×6	Yes
VK_FORMAT_ASTC_8x5_UNORM_BLOCK	8 × 5	No
VK_FORMAT_ASTC_8x5_SRGB_BLOCK	8 × 5	Yes
VK_FORMAT_ASTC_8x6_UNORM_BLOCK	8×6	No
VK_FORMAT_ASTC_8x6_SRGB_BLOCK	8 × 6	Yes
VK_FORMAT_ASTC_8x8_UNORM_BLOCK	8 × 8	No
VK_FORMAT_ASTC_8x8_SRGB_BLOCK	8 × 8	Yes
VK_FORMAT_ASTC_10x5_UNORM_BLOCK	10×5	No
VK_FORMAT_ASTC_10x5_SRGB_BLOCK	10×5	Yes
VK_FORMAT_ASTC_10x6_UNORM_BLOCK	10×6	No
VK_FORMAT_ASTC_10x6_SRGB_BLOCK	10×6	Yes
VK_FORMAT_ASTC_10x8_UNORM_BLOCK	10×8	No
VK_FORMAT_ASTC_10x8_SRGB_BLOCK	10×8	Yes
VK_FORMAT_ASTC_10x10_UNORM_BLOCK	10×10	No
VK_FORMAT_ASTC_10x10_SRGB_BLOCK	10×10	Yes
VK_FORMAT_ASTC_12x10_UNORM_BLOCK	12×10	No
VK_FORMAT_ASTC_12x10_SRGB_BLOCK	12×10	Yes
VK_FORMAT_ASTC_12x12_UNORM_BLOCK	12×12	No
VK_FORMAT_ASTC_12x12_SRGB_BLOCK	12×12	Yes

Appendix C

Layers & Extensions

Extensions to the Vulkan API can be defined by authors, groups of authors, and the Khronos Vulkan Working Group. In order not to compromise the readability of the Vulkan Specification, the core Specification does not incorporate most extensions. The online registry of extensions is available at URL

http://www.khronos.org/registry/vulkan/

and allows generating versions of the Specification incorporating different extensions.

Most of the content previously in this appendix does not specify **use** of specific Vulkan extensions and layers, but rather specifies the processes by which extensions and layers are created. As of version 1.0.21 of the Vulkan Specification, this content has been migrated to the Vulkan Documentation and Extensions document. Authors creating extensions and layers must follow the mandatory procedures in that document.

The remainder of this appendix documents each registered and published extension at a high level. Extensions are grouped as Khronos, multivendor, and then by vendor alphabetically.

C.1 VK_KHR_sampler_mirror_clamp_to_edge

```
Name String
```

VK_KHR_sampler_mirror_clamp_to_edge

Extension Type

Device extension

Registered Extension Number

15

Status

Final

Last Modified Date

2016-02-16

Revision

1

Dependencies

• This extension is written against version 1.0. of the Vulkan API.

Contributors

• Tobias Hector, Imagination Technologies

Contacts

• Tobias Hector (tobias.hector@imgtec.com)

VK_KHR_sampler_mirror_clamp_to_edge extends the set of sampler address modes to include an additional mode (VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE) that effectively uses a texture map twice as large as the original image in which the additional half of the new image is a mirror image of the original image.

This new mode relaxes the need to generate images whose opposite edges match by using the original image to generate a matching "mirror image". This mode allows the texture to be mirrored only once in the negative s, t, and r directions.

C.1.1 New Enum Constants

- Extending VkSamplerAddressMode:
 - VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE

C.1.2 Example

Creating a sampler with the new address mode in each dimension

```
VkSamplerCreateInfo createInfo =
{
    VK_STRUCTURE_TYPE_SAMPLER_CREATE_INFO // sType
    // Other members set to application-desired values
};

createInfo.addressModeU = VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE;
createInfo.addressModeV = VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE;
createInfo.addressModeW = VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE;

VkSampler sampler;
VkResult result = vkCreateSampler(
    device,
    &createInfo,
    &sampler);
```

C.1.3 Version History

- Revision 1, 2016-02-16 (Tobias Hector)
 - Initial draft

C.2 Window System Integration (WSI) Extensions

C.2.1 Editors

Significant specification work was performed by the following editors:

- Ian Elliott, LunarG
- Jesse Hall, Google
- James Jones, NVIDIA
- · Daniel Rakos, AMD

Relevant information is now documented for each extension.

C.2.2 VK_KHR_surface

Name String

VK_KHR_surface

Extension Type

Instance extension

Registered Extension Number

Last Modified Date

2016-14-01

Revision

25

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0 of the Vulkan API.

Contributors

- Patrick Doane, Blizzard
- Ian Elliott, LunarG
- Jesse Hall, Google
- James Jones, NVIDIA
- David Mao, AMD
- Norbert Nopper, Freescale
- Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- Graham Sellers, AMD

- · Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG
- Jason Ekstrand, Intel

Contacts

- · James Jones, NVIDIA
- · Ian Elliott, LunarG

The VK_KHR_surface extension is an instance extension. It introduces VkSurfaceKHR objects, which abstract native platform surface or window objects for use with Vulkan. It also provides a way to determine whether a queue family in a physical device supports presenting to particular surface.

Separate extensions for each each platform provide the mechanisms for creating VkSurfaceKHR objects, but once created they may be used in this and other platform-independent extensions, in particular the $VK_KHR_swapchain$ extension.

C.2.2.1 New Object Types

• VkSurfaceKHR

C.2.2.2 New Enum Constants

- Extending VkResult:
 - VK_ERROR_SURFACE_LOST_KHR
 - VK_ERROR_NATIVE_WINDOW_IN_USE_KHR

C.2.2.3 New Enums

- VkSurfaceTransformFlagBitsKHR
- VkPresentModeKHR
- VkColorSpaceKHR
- VkCompositeAlphaFlagBitsKHR

C.2.2.4 New Structures

- VkSurfaceCapabilitiesKHR
- VkSurfaceFormatKHR

C.2.2.5 New Functions

- vkDestroySurfaceKHR
- vkGetPhysicalDeviceSurfaceSupportKHR
- vkGetPhysicalDeviceSurfaceCapabilitiesKHR
- vkGetPhysicalDeviceSurfaceFormatsKHR
- vkGetPhysicalDeviceSurfacePresentModesKHR

C.2.2.6 Examples



Note

The example code for the VK_KHR_surface and VK_KHR_swapchain extensions will be removed from future versions of this appendix. The WSI example code was ported to the cube demo, that is shipped with the official Khronos SDK, and has been kept up-to-date in that location. There is little reason to maintain this example code in the appendix as well.

Example 1

Show how to obtain function pointers for WSI commands. Most applications should not need to do this, because the functions for the WSI extensions that are relevant for a given platform are exported by the official loader for that platform (e.g. the official Khronos loader for Microsoft Windows and Linux, and the Android loader). Other examples, will directly call the relevant command, not through a function pointer.

```
extern VkInstance instance;
// Obtain function pointers for the VK_KHR_surface commands:
PFN_vkGetPhysicalDeviceSurfaceSupportKHR pfnGetPhysicalDeviceSurfaceSupportKHR =
              (PFN_vkGetPhysicalDeviceSurfaceSupportKHR)vkGetInstanceProcAddr(instance,
                                                                                                                                                                                                                                                vkGetPhysicalDeviceSurfaceSupp
PFN_vkGetPhysicalDeviceSurfaceCapabilitiesKHR
           pfnGetPhysicalDeviceSurfaceCapabilitiesKHR =
              (PFN\_vkGetPhysicalDeviceSurfaceCapabilitiesKHR) vkGetInstanceProcAddr (instance, and all other processing the process of the
                                                                                                                                                                                                                                                                vkGetPhysicalDeviceSurfac
PFN_vkGetPhysicalDeviceSurfaceFormatsKHR pfnGetPhysicalDeviceSurfaceFormatsKHR =
              (PFN_vkGetPhysicalDeviceSurfaceFormatsKHR)vkGetInstanceProcAddr(instance,
                                                                                                                                                                                                                                                vkGetPhysicalDeviceSurfaceForm
                                                                                                                                                                                                                                                ");
PFN_vkGetPhysicalDeviceSurfacePresentModesKHR \leftrightarrow
           pfnGetPhysicalDeviceSurfacePresentModesKHR =
              (PFN_vkGetPhysicalDeviceSurfacePresentModesKHR)vkGetInstanceProcAddr(instance,
                                                                                                                                                                                                                                                                 vkGetPhysicalDeviceSurfac
```

Example 2

Pick which queues on a physical device to use for graphics and present operations for a given surface, and create a device with those queues.

```
extern VkInstance instance;
extern VkPhysicalDevice physicalDevice;
extern VkSurfaceKHR surface;
extern void Error(const char *);

uint32_t queueFamilyCount;
vkGetPhysicalDeviceQueueFamilyProperties(physicalDevice,
```

```
&queueFamilyCount, NULL);
VkQueueFamilyProperties* pMainQueueInfo =
    (VkQueueFamilyProperties*)malloc(queueFamilyCount * sizeof( <math>\leftarrow
       VkQueueFamilyProperties));
VkBool32* pSupportsPresent =
    (VkBool32 *)malloc(queueFamilyCount * sizeof(VkBool32));
vkGetPhysicalDeviceQueueFamilyProperties(physicalDevice, &queueFamilyCount,
                                          pMainQueueInfo);
for (uint32_t i = 0; i < queueFamilyCount; ++i)</pre>
    vkGetPhysicalDeviceSurfaceSupportKHR(physicalDevice, i, surface,
                                          &pSupportsPresent[i]);
// Search for a graphics and a present queue in the array of queue
// families, try to find one that supports both
uint32_t graphicsQueueFamilyIndex = UINT32_MAX;
uint32_t presentQueueFamilyIndex = UINT32_MAX;
for (uint32_t i = 0; i < queueFamilyCount; ++i)</pre>
    if ((pMainQueueInfo[i].queueFlags & VK_QUEUE_GRAPHICS_BIT) != 0)
    {
        if (graphicsQueueFamilyIndex == UINT32_MAX)
            graphicsQueueFamilyIndex = i;
        if (pSupportsPresent[i] == VK_TRUE)
            graphicsQueueFamilyIndex = i;
            presentQueueFamilyIndex = i;
            break;
if (presentQueueFamilyIndex == UINT32_MAX)
    // If didn't find a queue that supports both graphics and present, then
    // find a separate present queue.
    for (size_t i = 0; i < queueFamilyCount; ++i)</pre>
        if (pSupportsPresent[i] == VK_TRUE)
            presentQueueFamilyIndex = i;
            break;
// Free the temporary queue info allocations
free (pMainQueueInfo);
free (pSupportsPresent);
// Generate error if could not find both a graphics and a present queue
if (graphicsQueueFamilyIndex == UINT32_MAX ||
   presentQueueFamilyIndex == UINT32_MAX)
   Error("Could not find a graphics and a present queue");
```

```
// Put together the list of requested queues
const float queuePriority = 1.0f;
const VkDeviceQueueCreateInfo requestedQueues[] =
        VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO, // sType
                                                    // flags
                                                    // queueFamilyIndex
        graphicsQueueFamilyIndex,
                                                    // queueCount
        1.
                                                    // pQueuePriorities
        &queuePriority
    },
        VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO, // sType
                                                    // pNext
                                                    // flags
        Ο,
                                                    // queueFamilyIndex
        presentQueueFamilyIndex,
                                                   // queueCount
                                                   // pQueuePriorities
        &queuePriority
};
uint32_t requestedQueueCount = 2;
if (graphicsQueueFamilyIndex == presentQueueFamilyIndex)
    // We need only a single queue if the graphics queue is also the
    // present queue
   requestedQueueCount = 1;
}
// Create a device and request access to the specified queues
const VkDeviceCreateInfo deviceCreateInfo =
   VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO,
                                                   // sType
   NULL,
                                                    // pNext
   Ο,
                                                    // flags
   requestedQueueCount,
                                                    // queueCreateInfoCount
   &requestedQueues,
                                                    // pQueueCreateInfos
   . . .
};
VkDevice device;
vkCreateDevice(physicalDevice, &deviceCreateInfo, &device);
// Acquire graphics and present queue handle from device
VkQueue graphicsQueue, presentQueue;
vkGetDeviceQueue(device, graphicsQueueFamilyIndex, 0, &graphicsQueue);
vkGetDeviceQueue(device, presentQueueFamilyIndex, 0, &presentQueue);
```

C.2.2.7 Issues

1) Should this extension include a method to query whether a physical device supports presenting to a specific window or native surface on a given platform?

RESOLVED: Yes. Without this, applications would need to create a device instance to determine whether a particular window can be presented to. Knowing that a device supports presentation to a platform in general is not sufficient, as a single machine might support multiple seats, or instances of the platform that each use different underlying physical devices. Additionally, some platforms, such as X windows, different drivers and devices might be used for different windows depending on which section of the desktop they exist on.

2) Should the vkGetSurfacePropertiesKHR(), vkGetSurfaceFormatsKHR(), and vkGetSurfacePresentModesKHR() functions from VK_KHR_swapchain be modified to operate on physical devices and moved to this extension to implement the resolution of issue 1?

RESOLVED: No, separate query functions are needed, as the purposes served are similar but incompatible. The vkGetSurface*KHR functions return information that could potentially depend on an initialized device. For example, the formats supported for presentation to the surface might vary depending on which device extensions are enabled. The query introduced to resolve issue 1 should be used only to query generic driver or platform properties. The physical device parameter is intended to serve only as an identifier rather than a stateful object.

3) Should Vulkan include support Xlib or XCB as the API for accessing the X Window System platform?

RESOLVED: Both. XCB is a more modern and efficient API, but Xlib usage is deeply ingrained in many applications and likely will remain in use for the foreseeable future. Not all drivers necessarily need to support both, but including both as options in the core specification will probably encourage support, which should in turn eases adoption of the Vulkan API in older codebases. Additionally, the performance improvements possible with XCB likely won't have a measurable impact on the performance of Vulkan presentation and other minimal window system interactions defined here.

4) Should the GBM platform be included in the list of platform enums?

RESOLVED: Deferred, and will be addressed with a platform-specific extension to be written in the future.

C.2.2.8 Version History

- Revision 1, 2015-05-20 (James Jones)
 - Initial draft, based on LunarG KHR spec, other KHR specs, patches attached to bugs.
- Revision 2, 2015-05-22 (Ian Elliott)
 - Created initial Description section.
 - Removed query for whether a platform requires the use of a queue for presentation, since it was decided that presentation will always be modeled as being part of the queue.

- Fixed typos and other minor mistakes.
- Revision 3, 2015-05-26 (Ian Elliott)
 - Improved the Description section.
- Revision 4, 2015-05-27 (James Jones)
 - Fixed compilation errors in example code.
- Revision 5, 2015-06-01 (James Jones)
 - Added issues 1 and 2 and made related spec updates.
- Revision 6, 2015-06-01 (James Jones)
 - Merged the platform type mappings table previously removed from VK_KHR_swapchain with the platform description table in this spec.
 - Added issues 3 and 4 documenting choices made when building the initial list of native platforms supported.
- Revision 7, 2015-06-11 (Ian Elliott)
 - Updated table 1 per input from the KHR TSG.
 - Updated issue 4 (GBM) per discussion with Daniel Stone. He will create a platform-specific extension sometime in the future.
- Revision 8, 2015-06-17 (James Jones)
 - Updated enum-extending values using new convention.
 - Fixed the value of VK_SURFACE_PLATFORM_INFO_TYPE_SUPPORTED_KHR.
- Revision 9, 2015-06-17 (James Jones)
 - Rebased on Vulkan API version 126.
- Revision 10, 2015-06-18 (James Jones)
 - Marked issues 2 and 3 resolved.
- Revision 11, 2015-06-23 (Ian Elliott)
 - Examples now show use of function pointers for extension functions.
 - Eliminated extraneous whitespace.
- Revision 12, 2015-07-07 (Daniel Rakos)
 - Added error section describing when each error is expected to be reported.
 - Replaced the term "queue node index" with "queue family index" in the spec as that is the agreed term to be used in the latest version of the core header and spec.
 - Replaced bool32_t with VkBool32.
- Revision 13, 2015-08-06 (Daniel Rakos)
 - Updated spec against latest core API header version.

- Revision 14, 2015-08-20 (Ian Elliott)
 - Renamed this extension and all of its enumerations, types, functions, etc. This makes it compliant with the proposed standard for Vulkan extensions.
 - Switched from "revision" to "version", including use of the VK_MAKE_VERSION macro in the header file.
 - Did miscellaneous cleanup, etc.
- Revision 15, 2015-08-20 (Ian Elliott—porting a 2015-07-29 change from James Jones)
 - Moved the surface transform enums here from VK_WSI_swapchain so they could be re-used by VK_WSI_display.
- Revision 16, 2015-09-01 (James Jones)
 - Restore single-field revision number.
- Revision 17, 2015-09-01 (James Jones)
 - Fix example code compilation errors.
- Revision 18, 2015-09-26 (Jesse Hall)
 - Replaced VkSurfaceDescriptionKHR with the VkSurfaceKHR object, which is created via layered extensions.
 Added VkDestroySurfaceKHR.
- Revision 19, 2015-09-28 (Jesse Hall)
 - Renamed from VK_EXT_KHR_swapchain to VK_EXT_KHR_surface.
- Revision 20, 2015-09-30 (Jeff Vigil)
 - Add error result VK_ERROR_SURFACE_LOST_KHR.
- Revision 21, 2015-10-15 (Daniel Rakos)
 - Updated the resolution of issue #2 and include the surface capability queries in this extension.
 - Renamed SurfaceProperties to SurfaceCapabilities as it better reflects that the values returned are the capabilities of the surface on a particular device.
 - Other minor cleanup and consistency changes.
- Revision 22, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_surface to VK_KHR_surface.
- Revision 23, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to vkDestroySurfaceKHR.
- Revision 24, 2015-11-10 (Jesse Hall)
 - Removed VkSurfaceTransformKHR. Use VkSurfaceTransformFlagBitsKHR instead.
 - Rename VkSurfaceCapabilitiesKHR member maxImageArraySize to maxImageArrayLayers.
- Revision 25, 2016-01-14 (James Jones)

- Moved VK_ERROR_NATIVE_WINDOW_IN_USE_KHR from the VK_KHR_android_surface to the VK_KHR_ surface extension.
- 2016-08-23 (Ian Elliott)
 - Update the example code, to not have so many characters per line, and to split out a new example to show how to obtain function pointers.
- 2016-08-25 (Ian Elliott)
 - A note was added at the beginning of the example code, stating that it will be removed from future versions of the appendix.

C.2.3 VK_KHR_swapchain

Name String

VK_KHR_swapchain

Extension Type

Device extension

Registered Extension Number

2

Last Modified Date

2016-05-04

Revision

68

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Ian Elliott, LunarG
- Jesse Hall, Google
- Mathias Heyer, NVIDIA
- James Jones, NVIDIA
- · David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- · Jeff Vigil, Qualcomm

- Chia-I Wu, LunarG
- · Jason Ekstrand, Intel
- Matthaeus G. Chajdas, AMD
- Ray Smith, ARM

Contacts

- · James Jones, NVIDIA
- · Ian Elliott, LunarG

The VK_KHR_swapchain extension is the device-level companion to the VK_KHR_surface extension. It introduces VkSwapchainKHR objects, which provide the ability to present rendering results to a surface.

C.2.3.1 New Object Types

• VkSwapchainKHR

C.2.3.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR
 - VK_STRUCTURE_TYPE_PRESENT_INFO_KHR
- Extending VkImageLayout:
 - VK_IMAGE_LAYOUT_PRESENT_SRC_KHR
- Extending VkResult:
 - VK_SUBOPTIMAL_KHR
 - VK_ERROR_OUT_OF_DATE_KHR

C.2.3.3 New Enums

None

C.2.3.4 New Structures

- VkSwapchainCreateInfoKHR
- VkPresentInfoKHR

C.2.3.5 New Functions

- vkCreateSwapchainKHR
- vkDestroySwapchainKHR
- vkGetSwapchainImagesKHR
- vkAcquireNextImageKHR
- vkQueuePresentKHR

C.2.3.6 Issues

1) Does this extension allow the application to specify the memory backing of the presentable images?

RESOLVED: No. Unlike standard images, the implementation will allocate the memory backing of the presentable image.

2) What operations are allowed on presentable images?

RESOLVED: This is determined by the imageUsageFlags specified when creating the presentable image's swapchain.

3) Does this extension support MSAA presentable images?

RESOLVED: No. Presentable images are always single-sampled. Multi-sampled rendering must use regular images. To present the rendering results the application must manually resolve the multi-sampled image to a single-sampled presentable image prior to presentation.

4) Does this extension support stereo/multi-view presentable images?

RESOLVED: Yes. The number of views associated with a presentable image is determined by the imageArraySize specified when creating a swapchain. All presentable images in a given swapchain use the same array size.

5) Are the layers of stereo presentable images half-sized?

RESOLVED: No. The image extents always match those requested by the application.

6) Do the "present" and "acquire next image" commands operate on a queue? If not, do they need to include explicit semaphore objects to interlock them with queue operations?

RESOLVED: The present command operates on a queue. The image ownership operation it represents happens in order with other operations on the queue, so no explicit semaphore object is required to synchronize its actions. Applications may want to acquire the next image in separate threads from those in which they manage their queue, or in multiple threads. To make such usage easier, the acquire next image command takes a semaphore to signal as a method of explicit synchronization. The application must later queue a wait for this semaphore before queuing execution of any commands using the image.

7) Does vkAcquireNextImageKHR() block if no images are available?

RESOLVED: The command takes a timeout parameter. Special values for the timeout are 0, which makes the call a non-blocking operation, and UINT64_MAX, which blocks indefinitely. Values in between will block for up to the specified time. The call will return when an image becomes available or an error occurs. It may, but is not required to, return before the specified timeout expires if the swapchain becomes out of date.

8) Can multiple presents be queued using one QueuePresent call?

RESOLVED: Yes. VkPresentInfoKHR contains a list of swapchains and corresponding image indices that will be presented. When supported, all presentations queued with a single vkQueuePresentKHR call will be applied atomically as one operation. The same swapchain must not appear in the list more than once. Later extensions may provide applications stronger guarantees of atomicity for such present operations, and/or allow them to query whether atomic presentation of a particular group of swapchains is possible.

9) How do the presentation and acquire next image functions notify the application the targeted surface has changed?

RESOLVED: Two new result codes are introduced for this purpose:

VK_SUBOPTIMAL_KHR - Presentation will still succeed, subject to the window resize behavior, but the swapchain is no longer configured optimally for the surface it targets. Applications should query updated surface information and recreate their swapchain at the next convenient opportunity.

VK_ERROR_OUT_OF_DATE_KHR - Failure. The swapchain is no longer compatible with the surface it targets. The application must query updated surface information and recreate the swapchain before presentation will succeed.

These can be returned by both vkAcquireNextImageKHR and vkQueuePresentKHR.

10) Does the vkAcquireNextImageKHR command return a semaphore to the application via an output parameter, or accept a semaphore to signal from the application as an object handle parameter?

RESOLVED: Accept a semaphore to signal as an object handle. This avoids the need to specify whether the application must destroy the semaphore or whether it is owned by the swapchain, and if the latter, what its lifetime is and whether it can be re-used for other operations once it is received from vkAcquireNextImageKHR.

11) What types of swapchain queuing behavior should be exposed? Options include swap interval specification, mailbox/most recent Vs. FIFO queue management, targeting specific vertical blank intervals or absolute times for a given present operation, and probably others. For some of these, whether they are specified at swapchain creation time or as per-present parameters needs to be decided as well.

RESOLVED: The base swapchain extension will expose 3 possible behaviors (of which, FIFO will always be supported):

-Immediate present: Does not wait for vertical blanking period to update the current image, likely resulting in visible tearing. No internal queue is used. Present requests are applied immediately.

- -Mailbox queue: Waits for the next vertical blanking period to update the current image. No tearing should be observed. An internal single-entry queue is used to hold pending presentation requests. If the queue is full when a new presentation request is received, the new request replaces the existing entry, and any images associated with the prior entry become available for re-use by the application.
- -FIFO queue: Waits for the next vertical blanking period to update the current image. No tearing should be observed. An internal queue containing (numSwapchainImages 1) entries is used to hold pending presentation requests. New requests are appended to the end of the queue, and one request is removed from the beginning of the queue and processed during each vertical blanking period in which the queue is non-empty

Not all surfaces will support all of these modes, so the modes supported will be returned using a surface info query. All surfaces must support the FIFO queue mode. Applications must choose one of these modes up front when creating a swapchain. Switching modes can be accomplished by recreating the swapchain.

12) Can VK_PRESENT_MODE_MAILBOX_KHR provide non-blocking guarantees for vkAcquireNextImageKHR()? If so, what is the proper criteria?

RESOLVED: Yes. The difficulty is not immediately obvious here. Naively, if at least 3 images are requested, mailbox mode should always have an image available for the application if the application does not own any images when the call to vkAcquireNextImageKHR() was made. However, some presentation engines may have more than one "current" image, and would still need to block in some cases. The right requirement appears to be that if the application allocates the surface's minimum number of images + 1 then it is guaranteed non-blocking behavior when it does not currently own any images.

13) Is there a way to create and initialize a new swapchain for a surface that has generated a VK_SUBOPTIMAL_KHR return code while still using the old swapchain?

RESOLVED: Not as part of this specification. This could be useful to allow the application to create an "optimal" replacement swapchain and rebuild all its command buffers using it in a background thread at a low priority while continuing to use the "suboptimal" swapchain in the main thread. It could probably use the same "atomic replace" semantics proposed for recreating direct—to—device swapchains without incurring a mode switch. However, after discussion, it was determined some platforms probably could not support concurrent swapchains for the same surface though, so this will be left out of the base KHR extensions. A future extension could add this for platforms where it is supported.

14) Should there be a special value for VkSurfacePropertiesKHR::maxImageCount to indicate there are no practical limits on the number of images in a swapchain?

RESOLVED: Yes. There where often be cases where there is no practical limit to the number of images in a swapchain other than the amount of available resources (I.e., memory) in the system. Trying to derive a hard limit from things like memory size is prone to failure. It is better in such cases to leave it to applications to figure such soft limits out via trial/failure iterations.

15) Should there be a special value for VkSurfacePropertiesKHR::currentExtent to indicate the size of the platform surface is undefined?

RESOLVED: Yes. On some platforms (Wayland, for example), the surface size is defined by the images presented to it rather than the other way around.

16) Should there be a special value for VkSurfacePropertiesKHR::maxImageExtent to indicate there is no practical limit on the surface size?

RESOLVED: No. It seems unlikely such a system would exist. O could be used to indicate the platform places no limits on the extents beyond those imposed by Vulkan for normal images, but this query could just as easily return those same limits, so a special "unlimited" value doesn't seem useful for this field.

17) How should surface rotation and mirroring be exposed to applications? How do they specify rotation and mirroring transforms applied prior to presentation?

RESOLVED: Applications can query both the supported and current transforms of a surface. Both are specified relative to the device's "natural" display rotation and direction. The supported transforms indicates which orientations the presentation engine accepts images in. For example, a presentation engine that does not support transforming surfaces as part of presentation, and which is presenting to a surface that is displayed with a 90-degree rotation, would return only one supported transform bit: VK_SURFACE_TRANSFORM_ROT90_BIT_KHR. Applications must transform their rendering by the transform they specify when creating the swapchain in preTransform field.

18) Can surfaces ever not support VK_MIRROR_NONE? Can they support vertical and horizontal mirroring simultaneously? Relatedly, should VK_MIRROR_NONE[_BIT] be zero, or bit one, and should applications be allowed to specify multiple pre and current mirror transform bits, or exactly one?

RESOLVED: Since some platforms may not support presenting with a transform other than the native window's current transform, and prerotation/mirroring are specified relative to the device's natural rotation and direction, rather than relative to the surface's current rotation and direction, it is necessary to express lack of support for no mirroring. To allow this, the MIRROR_NONE enum must occupy a bit in the flags. Since MIRROR_NONE must be a bit in the bitmask rather than a bitmask with no values set, allowing more than one bit to be set in the bitmask would make it possible to describe undefined transforms such as VK_MIRROR_NONE_BIT | VK_MIRROR_HORIZONTAL_BIT, or a transform that includes both "no

mirroring" and "horizontal mirroring simultaneously. Therefore, it is desirable to allow specifying all supported mirroring transforms using only one bit. The question then becomes, should there be a VK_MIRROR_HORIZONTAL_AND_VERTICAL_BIT to represent a simultaneous horizontal and vertical mirror transform? However, such a transform is equivalent to a 180 degree rotation, so presentation engines and applications that wish to support or use such a transform can express it through rotation instead. Therefore, 3 exclusive bits are sufficient to express all needed mirroring transforms.

19) Should support for sRGB be required?

RESOLVED: In the advent of UHD and HDR display devices, proper color space information is vital to the display pipeline represented by the swapchain. The app can discover the supported format/color-space pairs and select a pair most suited to its rendering needs. Currently only the sRGB color space is supported, future extensions may provide support for more color spaces. See issues 23) and 24).

20) Is there a mechanism to modify or replace an existing swapchain with one targeting the same surface?

RESOLVED: Yes. This is described above in the text.

21) Should there be a way to set prerotation and mirroring using native APIs when presenting using a Vulkan swapchain?

RESOLVED: Yes. The transforms that can be expressed in this extension are a subset of those possible on native platforms. If a platform exposes a method to specify the transform of presented images for a given surface using native methods and exposes more transforms or other properties for surfaces than Vulkan supports, it might be impossible, difficult, or inconvenient to set some of those properties using Vulkan KHR extensions and some using the native interfaces. To avoid overwriting properties set using native commands when presenting using a Vulkan swapchain, the application can set the pretransform to "inherit", in which case the current native properties will be used, or if none are available, a platform-specific default will be used. Platforms that do not specify a reasonable default or do not provide native mechanisms to specify such transforms should not include the inherit bits in the supportedTransform bitmask they return in VkSurfacePropertiesKHR.

22) Should the content of presentable images be clipped by objects obscuring their target surface?

RESOLVED: Applications can choose which behavior they prefer. Allowing the content to be clipped could enable more optimal presentation methods on some platforms, but some applications might rely on the content of presentable images to perform techniques such as partial updates or motion blurs.

23) What is the purpose of specifying a VkColorSpaceKHR along with VkFormat when creating a swapchain?

RESOLVED: While Vulkan itself is color space agnostic (e.g. even the meaning of R, G, B and A can be freely defined by the rendering application), the swapchain eventually will have to present the images on a display device with specific color reproduction characteristics. If any color space transformations are necessary before an image can be displayed, the color space of the presented image must be known to the swapchain. A swapchain will only support a restricted set of color format and -space pairs. This set can be discovered via vkGetSurfaceInfoKHR. As it can be expected that most display devices support the sRGB color space, at least one format/color-space pair has to be exposed, where the color space is VK_COLOR_SPACE_SRGB_NONLINEAR.

24) How are sRGB formats and the sRGB color space related?

RESOLVED: While Vulkan exposes a number of SRGB texture formats, using such formats does not guarantee working in a specific color space. It merely means that the hardware can directly support applying the non-linear transfer functions defined by the sRGB standard color space when reading from or writing to images of that these formats. Still, it is unlikely that a swapchain will expose a _SRGB format along with any color space other than VK_COLOR_SPACE_SRGB_NONLINEAR.

On the other hand, non-SRGB formats will be very likely exposed in pair with a SRGB color space. This means, the hardware will not apply any transfer function when reading from or writing to such images, yet they will still be presented on a device with sRGB display characteristics. In this case the application is responsible for applying the transfer function, for instance by using shader math.

25) How are the lifetime of surfaces and swapchains targeting them related?

RESOLVED: A surface must outlive any swapchains targeting it. A VkSurfaceKHR owns the binding of the native window to the Vulkan driver.

26) How can the client control the way the alpha channel of swap chain images is treated by the presentation engine during compositing?

RESOLVED: We should add new enum values to allow the client to negotiate with the presentation engine on how to treat image alpha values during the compositing process. Since not all platforms can practically control this through the Vulkan driver, a value of INHERIT is provided like for surface transforms.

27) Is vkCreateSwapchainKHR() the right function to return VK_ERROR_NATIVE_WINDOW_IN_USE_KHR, or should the various platform- specific VkSurface factory functions catch this error earlier?

RESOLVED: For most platforms, the VkSurface structure is a simple container holding the data that identifies a native window or

other object representing a surface on a particular platform. For the surface factory functions to return this error, they would likely need to register a reference on the native objects with the native display server some how, and ensure no other such references exist. Surfaces were not intended to be that heavyweight.

Swapchains are intended to be the objects that directly manipulate native windows and communicate with the native presentation mechanisms. Swapchains will already need to communicate with the native display server to negotiate allocation and/or presentation of presentable images for a native surface. Therefore, it makes more sense for swapchain creation to be the point at which native object exclusivity is enforced. Platforms may choose to enforce further restrictions on the number of VkSurface objects that may be created for the same native window if such a requirement makes sense on a particular platform, but a global requirement is only sensible at the swapchain level.

C.2.3.7 Examples



Note

The example code for the VK_KHR_surface and VK_KHR_swapchain extensions will be removed from future versions of this appendix. The WSI example code was ported to the cube demo, that is shipped with the official Khronos SDK, and has been kept up-to-date in that location. There is little reason to maintain this example code in the appendix as well.

Example 1

Show how to obtain function pointers for WSI commands. Most applications shouldn't need to do this, because the functions for the WSI extensions that are relevant for a given platform are exported by the official loader for that platform (e.g. the official Khronos loader for Microsoft Windows and Linux, and the Android loader). Other examples, will directly call the relevant command, not through a function pointer.

Create a swapchain for a surface on a particular instance of a native platform.

```
extern VkPhysicalDevice physicalDevice;
extern VkDevice device;
extern VkSurfaceKHR surface;
// Check the surface properties and formats
VkSurfaceCapabilitiesKHR surfCapabilities;
vkGetPhysicalDeviceSurfaceCapabilitiesKHR(physicalDevice, surface,
                                           &surfCapabilities);
uint32 t formatCount;
vkGetPhysicalDeviceSurfaceFormatsKHR(physicalDevice, surface,
                                     &formatCount, NULL);
VkSurfaceFormatKHR* pSurfFormats =
    (VkSurfaceFormatKHR*) malloc(formatCount * sizeof(VkSurfaceFormatKHR));
vkGetPhysicalDeviceSurfaceFormatsKHR(physicalDevice, surface,
                                     &formatCount, pSurfFormats);
uint32_t presentModeCount;
vkGetPhysicalDeviceSurfacePresentModesKHR(physicalDevice, surface,
                                          &presentModeCount, NULL);
VkPresentModeKHR* pPresentModes =
    (VkPresentModeKHR*) malloc(presentModeCount * sizeof(VkPresentModeKHR));
vkGetPhysicalDeviceSurfacePresentModesKHR(physicalDevice, surface,
                                           &presentModeCount, pPresentModes);
VkExtent2D swapchainExtent;
// width and height are either both 0xFFFFFFFF, or both not 0xFFFFFFFF.
if (surfCapabilities.currentExtent.width == 0xFFFFFFFF)
    // If the surface size is undefined, the size is set to the size
    // of the images requested, which must fit within the minimum and
    // maximum values.
    swapchainExtent.width = 320;
    swapchainExtent.height = 320;
    if (swapchainExtent.width < surfCapabilities.minImageExtent.width)</pre>
        swapchainExtent.width = surfCapabilities.minImageExtent.width;
    else if (swapchainExtent.width > surfCapabilities.maxImageExtent.width)
        swapchainExtent.width = surfCapabilities.maxImageExtent.width;
    if (swapchainExtent.height < surfCapabilities.minImageExtent.height)</pre>
        swapchainExtent.height = surfCapabilities.minImageExtent.height;
    else if (swapchainExtent.height > surfCapabilities.maxImageExtent.height)
        swapchainExtent.height = surfCapabilities.maxImageExtent.height;
else
```

```
// If the surface size is defined, the swapchain size must match
    swapchainExtent = surfCapabilities.currentExtent;
// Application desires to simultaneously acquire 2 images (which is one
// more than "surfCapabilities.minImageCount").
uint32_t desiredNumOfSwapchainImages = surfCapabilities.minImageCount + 1;
if ((surfCapabilities.maxImageCount > 0) &&
    (desiredNumOfSwapchainImages > surfCapabilities.maxImageCount))
{
    // Application must settle for fewer images than desired:
    desiredNumOfSwapchainImages = surfCapabilities.maxImageCount;
}
VkFormat swapchainFormat;
// If the format list includes just one entry of VK_FORMAT_UNDEFINED,
// the surface has no preferred format. Otherwise, at least one
// supported format will be returned (assuming that the
// vkGetPhysicalDeviceSurfaceSupportKHR function, in the
// VK_KHR_surface extension returned support for the surface).
if ((formatCount == 1) && (pSurfFormats[0].format == VK_FORMAT_UNDEFINED))
    swapchainFormat = VK_FORMAT_R8G8B8_UNORM;
else
{
    assert(formatCount >= 1);
    swapchainFormat = pSurfFormats[0].format;
VkColorSpaceKHR swapchainColorSpace = pSurfFormats[0].colorSpace;
// If mailbox mode is available, use it, as it is the lowest-latency non-
// tearing mode. If not, fall back to FIFO which is always available.
VkPresentModeKHR swapchainPresentMode = VK_PRESENT_MODE_FIFO_KHR;
for (size_t i = 0; i < presentModeCount; ++i)</pre>
    if (pPresentModes[i] == VK_PRESENT_MODE_MAILBOX_KHR)
    {
        swapchainPresentMode = VK_PRESENT_MODE_MAILBOX_KHR;
        break;
    }
}
const VkSwapchainCreateInfoKHR createInfo =
    VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR,
                                                    // sType
   NULL,
                                                    // pNext
                                                    // flags
    Ο,
                                                    // surface
    surface,
    desiredNumOfSwapchainImages,
                                                    // minImageCount
                                                    // imageFormat
    swapchainFormat,
    swapchainColorSpace,
                                                    // imageColorSpace
                                                    // imageExtent
    swapchainExtent,
                                                    // imageArrayLayers
   1,
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,
                                                   // imageUsage
    VK_SHARING_MODE_EXCLUSIVE,
                                                   // imageSharingMode
    Ο,
                                                    // queueFamilyIndexCount
                                                    // pQueueFamilyIndices
   NULL,
    surfCapabilities.currentTransform,
                                              // preTransform
```

Obtaining the persistent images of a swapchain

```
extern VkDevice device;
uint32_t swapchainImageCount;
vkGetSwapchainImagesKHR(device, swapchain, &swapchainImageCount, NULL);

VkImage* pSwapchainImages = (VkImage*)malloc(swapchainImageCount * sizeof(VkImage) \( \to \);
vkGetSwapchainImagesKHR(device, swapchain, &swapchainImageCount, pSwapchainImages) \( \to \);

vkGetSwapchainImagesKHR(device, swapchain, &swapchainImageCount, pSwapchainImages) \( \to \);
```

Example 4

Simple rendering and presenting using separate graphics and present queues.

```
extern VkDevice device;
// Construct command buffers rendering to the presentable images
VkCmdBuffer cmdBuffers[swapchainImageCount];
VkImageView views[swapchainImageCount];
extern VkCmdBufferBeginInfo beginInfo;
extern uint32_t graphicsQueueFamilyIndex, presentQueueFamilyIndex;
for (size_t i = 0; i < swapchainImageCount; ++i)</pre>
    const VkImageViewCreateInfo viewInfo =
        VK_STRUCTURE_TYPE_IMAGE_VIEW_CREATE_INFO,
                                                   // sType
                                                    // pNext
        NULL,
                                                    // flags
                                                    // image
        pSwapchainImages[i],
                                                    // viewType
        VK_IMAGE_VIEW_TYPE_2D,
                                                    // format
        swapchainFormat,
    };
    vkCreateImageView(device, &viewInfo, &views[i]);
    vkBeginCommandBuffer(cmdBuffers[i], &beginInfo);
    // Need to transition image from presentable state before being able
    // to render
    const VkImageMemoryBarrier acquireImageBarrier =
        VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER, // sType
```

```
// pNext
       NULL,
       VK_ACCESS_MEMORY_READ_BIT,
                                                   // srcAccessMask
       VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,
                                                  // dstAccessMask
       VK_IMAGE_LAYOUT_PRESENT_SRC_KHR,
                                                  // oldLayout
       VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL, // newLayout
                                                  // srcQueueFamilyIndex
       presentQueueFamilyIndex,
                                                   // dstQueueFamilyIndex
       graphicsQueueFamilyIndex,
       pSwapchainImages[i].image,
                                                   // image
    };
    vkCmdPipelineBarrier(
       cmdBuffers[i],
       VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
       VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
       1,
       &acquireImageBarrier)
   // ... Render to views[i] ...
    // Need to transition image into presentable state before being able
    // to present
    const VkImageMemoryBarrier presentImageBarrier =
    {
       VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER,
                                                  // sType
                                                   // pNext
       NULL,
       VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,
                                                  // srcAccessMask
       VK_ACCESS_MEMORY_READ_BIT,
                                                  // dstAccessMask
       VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL, // oldLayout
       VK_IMAGE_LAYOUT_PRESENT_SRC_KHR,
                                                  // newLayout
       graphicsQueueFamilyIndex,
                                                  // srcQueueFamilyIndex
                                                  // dstQueueFamilyIndex
       presentQueueFamilyIndex,
                                                  // image
       pSwapchainImages[i].image,
    };
    vkCmdPipelineBarrier(
       cmdBuffers[i],
       VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
       VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT,
       VK_FALSE,
       1,
       &presentImageBarrier);
   vkEndCommandBuffer(cmdBuffers[i]);
}
const VkSemaphoreCreateInfo semaphoreCreateInfo =
   VK_STRUCTURE_TYPE_SEMAPHORE_CREATE_INFO,
                                               // sType
   NULL,
                                               // pNext
    0
                                               // flags
};
```

```
VkSemaphore imageAcquiredSemaphore;
vkCreateSemaphore(device,
                  &semaphoreCreateInfo,
                  &imageAcquiredSemaphore);
VkSemaphore renderingCompleteSemaphore;
vkCreateSemaphore(device,
                  &semaphoreCreateInfo,
                  &renderingCompleteSemaphore);
VkResult result;
do
    uint32_t imageIndex = UINT32_MAX;
    // Get the next available swapchain image
    result = vkAcquireNextImageKHR(
        device,
        swapchain,
        UINT64_MAX,
        imageAcquiredSemaphore,
        VK_NULL_HANDLE,
        &imageIndex);
    // Swapchain cannot be used for presentation if failed to acquired
    // new image.
    if (result < 0)</pre>
        break;
    // Submit rendering work to the graphics queue
    const VkPipelineStageFlags waitDstStageMask =
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT;
    const VkSubmitInfo submitInfo =
        VK_STRUCTURE_TYPE_SUBMIT_INFO,
                                                 // sType
        NULL,
                                                 // pNext
                                                 // waitSemaphoreCount
        1,
        &imageAcquiredSemaphore,
                                                 // pWaitSemaphores
                                                 // pWaitDstStageMasks
        &waitDstStageMask,
                                                 // commandBufferCount
        &cmdBuffers[imageIndex],
                                                 // pCommandBuffers
                                                 // signalSemaphoreCount
        &renderingCompleteSemaphore
                                                 // pSignalSemaphores
    vkQueueSubmit(graphicsQueue, 1, &submitInfo, VK_NULL_HANDLE);
    // Submit present operation to present queue
    const VkPresentInfoKHR presentInfo =
    {
        VK_STRUCTURE_TYPE_PRESENT_INFO_KHR,
                                                 // sType
        NULL,
                                                 // pNext
        1,
                                                 // waitSemaphoreCount
        &renderingCompleteSemaphore,
                                                 // pWaitSemaphores
                                                 // swapchainCount
        &swapchain,
                                                 // pSwapchains
                                                 // pImageIndices
        &imageIndex,
        NULL
                                                 // pResults
```

```
};

result = vkQueuePresentKHR(presentQueue, &presentInfo);
} while (result >= 0);
```

Handle VK_ERROR_OUT_OF_DATE_KHR by recreating the swapchain and VK_SUBOPTIMAL_KHR by checking whether recreation is needed.

```
extern VkInstance instance;
extern VkPhysicalDevice physicalDevice;
extern VkDevice device;
extern VkQueue presentQueue;
extern VkSurfaceKHR surface;
// Contains code to build the application's reusable command buffers.
extern void CreateCommandBuffers(VkDevice device,
                                const VkImage* swapchainImages);
// Returns non-zero if the application considers the swapchain out
// of date even though it is still compatible with the platform
// surface it is targeting.
extern int MustRecreateSwapchain(VkDevice device,
                                 VkSwapchainKHR swapchain,
                                 VkSurfaceKHR surface);
extern void CreateSwapchain(VkDevice device,
                            VkSurfaceKHR surface,
                            VkSwapchainKHR oldSwapchain,
                            VkSwapchainKHR* pSwapchain);
    . . .
    // Most of this function is identical to what is in Example 2.
    // One differences is in the initialization of the following struct.
    const VkSwapchainCreateInfoKHR createInfo =
        VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR, // sType
        NULL,
                                                      // pNext
                                                      // flags
        Ο,
                                                      // surface
        surface,
        desiredNumOfSwapchainImages,
                                                      // minImageCount
                                                      // imageFormat
        swapchainFormat,
        swapchainColorSpace,
                                                      // imageColorSpace
                                                      // imageExtent
        swapchainExtent,
                                                      // imageArrayLayers
        VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,
                                                      // imageUsage
        VK_SHARING_MODE_EXCLUSIVE,
                                                      // imageSharingMode
                                                      // queueFamilyIndexCount
                                                      // pQueueFamilyIndices
        NULL,
        surfCapabilities.currentTransform,
                                                      // preTransform
                                                      // compositeAlpha
        VK_COMPOSITE_ALPHA_INHERIT_BIT_KHR,
        swapchainPresentMode,
                                                      // presentMode
        VK_TRUE,
                                                      // clipped
```

```
// If the caller specified an existing swapchain, replace it with
        // the new swapchain being created here. This will allow a
        // seamless transition between the old and new swapchains on
        // platforms that support it.
                                                       // oldSwapchain
        oldSwapchain
    } ;
    . . .
    // Another difference is in the need to destroy the old swapchain, if
   // it exists.
   11
    // Note: destroying the swapchain also cleans up all its associated
    // presentable images once the platform is done with them.
   if (oldSwapchain != VK_NULL_HANDLE)
        vkDestroySwapchainKHR(device, oldSwapchain, NULL);
void mainLoop(void)
   VkSwapchainKHR swapchain = VK_NULL_HANDLE;
    const VkSemaphoreCreateInfo semaphoreCreateInfo =
        VK_STRUCTURE_TYPE_SEMAPHORE_CREATE_INFO,
                                                     // sType
        NULL,
                                                     // pNext
        0
                                                     // flags
    };
   VkSemaphore imageAcquiredSemaphore;
    vkCreateSemaphore(device,
                      &semaphoreCreateInfo,
                      &imageAcquiredSemaphore);
   VkSemaphore renderingCompleteSemaphore;
    vkCreateSemaphore(device,
                      &semaphoreCreateInfo,
                      &renderingCompleteSemaphore);
    while (1)
        VkResult result;
        CreateSwapchain(device, surface, swapchain, &swapchain);
        uint32_t swapchainImageCount;
        vkGetSwapchainImagesKHR(device, swapchain, &swapchainImageCount, NULL);
        VkImage* pSwapchainImages =
            (VkImage*) malloc(swapchainImageCount * sizeof(VkImage));
        vkGetSwapchainImagesKHR (device, swapchain,
                                &swapchainImageCount, pSwapchainImages);
        CreateCommandBuffers(device, pSwapchainImages);
        free(pSwapchainImages);
```

```
while (1)
    uint32_t imageIndex;
    // Get the next available swapchain image
    result = vkAcquireNextImageKHR(
       device,
        swapchain,
        UINT64 MAX,
        imageAcquiredSemaphore,
        VK_NULL_HANDLE,
        &imageIndex);
    if (result == VK_ERROR_OUT_OF_DATE_KHR)
        // swapchain is out of date. Needs to be recreated for
        // defined results.
       break;
    else if (result == VK_SUBOPTIMAL_KHR)
        \ensuremath{//} Ignore this result here. If any expensive
        // preprocessing work has already been done for this
        // frame, it is likely in the application's interest to
        // continue processing this frame with the current
        // swapchain rather than recreate it and waste the
        // preprocessing work.
    }
    else if (result < 0)</pre>
        // Unhandled error. Abort.
        return;
    // Submit rendering work to the graphics queue
    const VkPipelineStageFlags waitDstStageMask =
       VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT;
    const VkSubmitInfo submitInfo =
        VK_STRUCTURE_TYPE_SUBMIT_INFO,
                                           // sType
       NULL,
                                            // pNext
                                           // waitSemaphoreCount
        &imageAcquiredSemaphore,
                                           // pWaitSemaphores
        &waitDstStageMask,
                                            // pWaitDstStageMasks
                                            // commandBufferCount
        1,
        &cmdBuffers[imageIndex],
                                            // pCommandBuffers
        1,
                                            // signalSemaphoreCount
        &renderingCompleteSemaphore
                                            // pSignalSemaphores
    } ;
    vkQueueSubmit(graphicsQueue, 1, &submitInfo, VK_NULL_HANDLE);
    // Submit present operation to present queue
    const VkPresentInfoKHR presentInfo =
    {
        VK_STRUCTURE_TYPE_PRESENT_INFO_KHR, // sType
        NULL,
                               // pNext
```

```
// waitSemaphoreCount
            &renderingCompleteSemaphore,
                                                 // pWaitSemaphores
            1,
                                                 // swapchainCount
                                                 // pSwapchains
            &swapchain,
                                                 // pImageIndices
            &imageIndex,
            NULL
                                                 // pResults
        };
        result = vkQueuePresentKHR (presentQueue, &presentInfo);
        if (result == VK_ERROR_OUT_OF_DATE_KHR)
            // The swapchain is out of date, and must be recreated for
            // defined results.
            break;
        else if (result == VK_SUBOPTIMAL_KHR)
            // Something has changed about the native surface since
            // the swapchain was created, but it is still compatible
            // with the swapchain. The app must choose whether it
            // wants to create a more up to date swapchain before it
            // begins processing the next frame.
            if (MustRecreateSwapchain(device, swapchain, surface))
                break;
        }
        else if (result < 0)</pre>
            // Unhandled error. Abort.
            return;
    }
}
```

Meter a CPU thread based on presentation rate. Note this will only limit the thread to the actual presentation rate when using VK_PRESENT_MODE_IFFO_KHR. When using VK_PRESENT_MODE_IMMEDIATE_KHR, presentation rate will be limited only by rendering rate, so this example will be equivalent to waiting on a command buffer fence. When using VK_PRESENT_MODE_MAILBOX_KHR, some presented images may be skipped if a newer image is available by the time the presentation target is ready to process a new frame, so this code would again be limited only by the rendering rate. Applications using mailbox mode should use some other mechanism to meter their rendering, and it is assumed applications using immediate mode have no desire to limit the rate of their rendering based on presentation rate.

```
static const int FRAME_LAG = 2
VkPresentInfoKHR presentInfo;
const VkSemaphoreCreateInfo semaphoreCreateInfo =
    VK_STRUCTURE_TYPE_SEMAPHORE_CREATE_INFO,
                                                 // sType
    NULL,
                                                 // pNext
                                                 // flags
};
VkSemaphore imageAcquiredSemaphore;
vkCreateSemaphore(device,
                  &semaphoreCreateInfo,
                  &imageAcquiredSemaphore);
VkSemaphore renderingCompleteSemaphore;
vkCreateSemaphore(device,
                  &semaphoreCreateInfo,
                  &renderingCompleteSemaphore);
VkFence fences[FRAME_LAG];
bool fencesInited[FRAME_LAG];
int frameIdx = 0;
int imageIdx = 0;
int waitFrame;
CreateFences (device, fences, FRAME_LAG);
for (int i = 0; i < FRAME_LAG; ++i)</pre>
    fencesInited[i] = false;
while (1) {
    if (fencesInited[frameIdx])
        // Ensure no more than FRAME_LAG presentations are outstanding
        vkWaitForFences(device, 1, &fences[frameIdx], VK_TRUE, UINT64_MAX);
        vkResetFences(device, 1, &fences[frameIdx]);
    }
    vkAcquireNextImageKHR(
        device,
        fifoSwapchain,
        UINT64_MAX,
        imageAcquiredSemaphore,
        fences[frameIdx],
        &imageIdx);
    fencesInited[frameIdx] = true;
    // Generate a command buffer for the current frame.
    GenFrameCmdBuffer(cmdBuffer);
    // Submit rendering work to the graphics queue
    const VkPipelineStageFlags waitDstStageMask =
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT;
    const VkSubmitInfo submitInfo =
```

```
VK_STRUCTURE_TYPE_SUBMIT_INFO,
                                            // sType
   NULL,
                                            // pNext
                                            // waitSemaphoreCount
   &imageAcquiredSemaphore,
                                            // pWaitSemaphores
   &waitDstStageMask,
                                            // pWaitDstStageMasks
                                            // commandBufferCount
                                            // pCommandBuffers
   &cmdBuffer,
                                            // signalSemaphoreCount
                                            // pSignalSemaphores
   &renderingCompleteSemaphore
};
vkQueueSubmit(queue, 1, &submitInfo, VK_NULL_HANDLE);
// Submit present operation to present queue
const VkPresentInfoKHR presentInfo =
{
   VK_STRUCTURE_TYPE_PRESENT_INFO_KHR,
                                           // sType
   NULL.
                                            // pNext
                                            // waitSemaphoreCount
   1,
   &renderingCompleteSemaphore,
                                            // pWaitSemaphores
                                            // swapchainCount
   &fifoSwapchain,
                                            // pSwapchains
                                            // pImageIndices
   &imageIdx,
   NULL
                                            // pResults
};
result = vkQueuePresentKHR(queue, &presentInfo);
frameIdx += 1;
frameIdx %= FRAME_LAG;
```

C.2.3.8 Version History

- Revision 1, 2015-05-20 (James Jones)
 - Initial draft, based on LunarG KHR spec, other KHR specs, patches attached to bugs.
- Revision 2, 2015-05-22 (Ian Elliott)
 - Made many agreed-upon changes from 2015-05-21 KHR TSG meeting. This includes using only a queue for presentation, and having an explicit function to acquire the next image.
 - Fixed typos and other minor mistakes.
- Revision 3, 2015-05-26 (Ian Elliott)
 - Improved the Description section.
 - Added or resolved issues that were found in improving the Description. For example, pSurfaceDescription is used consistently, instead of sometimes using pSurface.
- Revision 4, 2015-05-27 (James Jones)
 - Fixed some grammatical errors and typos
 - Filled in the description of imageUseFlags when creating a swapchain.

- Added a description of swapInterval.
- Replaced the paragraph describing the order of operations on a queue for image ownership and presentation.
- Revision 5, 2015-05-27 (James Jones)
 - Imported relevant issues from the (abandoned) vk_wsi_persistent_swapchain_images extension.
 - Added issues 6 and 7, regarding behavior of the acquire next image and present commands with respect to queues.
 - Updated spec language and examples to align with proposed resolutions to issues 6 and 7.
- Revision 6, 2015-05-27 (James Jones)
 - Added issue 8, regarding atomic presentation of multiple swapchains
 - Updated spec language and examples to align with proposed resolution to issue 8.
- Revision 7, 2015-05-27 (James Jones)
 - Fixed compilation errors in example code, and made related spec fixes.
- Revision 8, 2015-05-27 (James Jones)
 - Added issue 9, and the related VK SUBOPTIMAL KHR result code.
 - Renamed VK_OUT_OF_DATE_KHR to VK_ERROR_OUT_OF_DATE_KHR.
- Revision 9, 2015-05-27 (James Jones)
 - Added inline proposed resolutions (marked with [JRJ]) to some XXX questions/issues. These should be moved to the issues section in a subsequent update if the proposals are adopted.
- Revision 10, 2015-05-28 (James Jones)
 - Converted vkAcquireNextImageKHR back to a non-queue operation that uses a VkSemaphore object for explicit synchronization.
 - Added issue 10 to determine whether vkAcquireNextImageKHR generates or returns semaphores, or whether it
 operates on a semaphore provided by the application.
- Revision 11, 2015-05-28 (James Jones)
 - Marked issues 6, 7, and 8 resolved.
 - Renamed VkSurfaceCapabilityPropertiesKHR to VkSurfacePropertiesKHR to better convey the mutable nature of the info it contains.
- Revision 12, 2015-05-28 (James Jones)
 - Added issue 11 with a proposed resolution, and the related issue 12.
 - Updated various sections of the spec to match the proposed resolution to issue 11.
- Revision 13, 2015-06-01 (James Jones)
 - Moved some structures to VK_EXT_KHR_swap_chain to resolve the spec's issues 1 and 2.
- Revision 14, 2015-06-01 (James Jones)
 - Added code for example 4 demonstrating how an application might make use of the two different present and acquire next image KHR result codes.

- Added issue 13.
- Revision 15, 2015-06-01 (James Jones)
 - Added issues 14 16 and related spec language.
 - Fixed some spelling errors.
 - Added language describing the meaningful return values for vkAcquireNextImageKHR and vkQueuePresentKHR.
- Revision 16, 2015-06-02 (James Jones)
 - Added issues 17 and 18, as well as related spec language.
 - Removed some erroneous text added by mistake in the last update.
- Revision 17, 2015-06-15 (Ian Elliott)
 - Changed special value from "-1" to "0" so that the data types can be unsigned.
- Revision 18, 2015-06-15 (Ian Elliott)
 - Clarified the values of VkSurfacePropertiesKHR::minImageCount and the timeout parameter of the vkAcquireNextImageKHR function.
- Revision 19, 2015-06-17 (James Jones)
 - Misc. cleanup. Removed resolved inline issues and fixed typos.
 - Fixed clarification of VkSurfacePropertiesKHR::minImageCount made in version 18.
 - Added a brief "Image Ownership" definition to the list of terms used in the spec.
- Revision 20, 2015-06-17 (James Jones)
 - Updated enum-extending values using new convention.
- Revision 21, 2015-06-17 (James Jones)
 - Added language describing how to use VK_IMAGE_LAYOUT_PRESENT_SOURCE_KHR.
 - Cleaned up an XXX comment regarding the description of which queues vkQueuePresentKHR can be used on.
- Revision 22, 2015-06-17 (James Jones)
 - Rebased on Vulkan API version 126.
- Revision 23, 2015-06-18 (James Jones)
 - Updated language for issue 12 to read as a proposed resolution.
 - Marked issues 11, 12, 13, 16, and 17 resolved.
 - Temporarily added links to the relevant bugs under the remaining unresolved issues.
 - Added issues 19 and 20 as well as proposed resolutions.
- Revision 24, 2015-06-19 (Ian Elliott)
 - Changed special value for VkSurfacePropertiesKHR::currentExtent back to "-1" from "0". This value will never need to be unsigned, and "0" is actually a legal value.
- Revision 25, 2015-06-23 (Ian Elliott)

- Examples now show use of function pointers for extension functions.
- Eliminated extraneous whitespace.
- Revision 26, 2015-06-25 (Ian Elliott)
 - Resolved Issues 9 & 10 per KHR TSG meeting.
- Revision 27, 2015-06-25 (James Jones)
 - Added oldSwapchain member to VkSwapchainCreateInfoKHR.
- Revision 28, 2015-06-25 (James Jones)
 - Added the "inherit" bits to the rotatation and mirroring flags and the associated issue 21.
- Revision 29, 2015-06-25 (James Jones)
 - Added the "clipped" flag to VkSwapchainCreateInfoKHR, and the associated issue 22.
 - Specified that presenting an image does not modify it.
- Revision 30, 2015-06-25 (James Jones)
 - Added language to the spec that clarifies the behavior of vkCreateSwapchainKHR() when the oldSwapchain field of VkSwapchainCreateInfoKHR is not NULL.
- Revision 31, 2015-06-26 (Ian Elliott)
 - Example of new VkSwapchainCreateInfoKHR members, "oldSwapchain" and "clipped".
 - Example of using VkSurfacePropertiesKHR::{minlmax}ImageCount to set VkSwapchainCreateInfoKHR::minImageCount.
 - Rename vkGetSurfaceInfoKHR()'s 4th param to "pDataSize", for consistency with other functions.
 - Add macro with C-string name of extension (just to header file).
- Revision 32, 2015-06-26 (James Jones)
 - Minor adjustments to the language describing the behavior of "oldSwapchain"
 - Fixed the version date on my previous two updates.
- Revision 33, 2015-06-26 (Jesse Hall)
 - Add usage flags to VkSwapchainCreateInfoKHR
- Revision 34, 2015-06-26 (Ian Elliott)
 - Rename vkQueuePresentKHR()'s 2nd param to "pPresentInfo", for consistency with other functions.
- Revision 35, 2015-06-26 (Jason Ekstrand)
 - Merged the VkRotationFlagBitsKHR and VkMirrorFlagBitsKHR enums into a single VkSurfaceTransformFlagBitsKHR enum.
- Revision 36, 2015-06-26 (Jason Ekstrand)

- Added a VkSurfaceTransformKHR enum that isn't a bitmask. Each value in VkSurfaceTransformKHR corresponds
 directly to one of the bits in VkSurfaceTransformFlagBitsKHR so transforming from one to the other is easy.
 Having a separate enum means that currentTransform and preTransform are now unambiguous by definition.
- Revision 37, 2015-06-29 (Ian Elliott)
 - Corrected one of the signatures of vkAcquireNextImageKHR, which had the last two parameters switched from what it is elsewhere in the specification and header files.
- Revision 38, 2015-06-30 (Ian Elliott)
 - Corrected a typo in description of the vkGetSwapchainInfoKHR() function.
 - Corrected a typo in header file comment for VkPresentInfoKHR::sType.
- Revision 39, 2015-07-07 (Daniel Rakos)
 - Added error section describing when each error is expected to be reported.
 - Replaced bool32_t with VkBool32.
- Revision 40, 2015-07-10 (Ian Elliott)
 - Updated to work with version 138 of the "vulkan.h" header. This includes declaring the VkSwapchainKHR type using the new VK_DEFINE_NONDISP_HANDLE macro, and no longer extending VkObjectType (which was eliminated).
- Revision 41 2015-07-09 (Mathias Heyer)
 - Added color space language.
- Revision 42, 2015-07-10 (Daniel Rakos)
 - Updated query mechanism to reflect the convention changes done in the core spec.
 - Removed "queue" from the name of VK_STRUCTURE_TYPE_QUEUE_PRESENT_INFO_KHR to be consistent with the established naming convention.
 - Removed reference to the no longer existing VkObjectType enum.
- Revision 43, 2015-07-17 (Daniel Rakos)
 - Added support for concurrent sharing of swapchain images across queue families.
 - Updated sample code based on recent changes
- Revision 44, 2015-07-27 (Ian Elliott)
 - Noted that support for VK_PRESENT_MODE_FIFO_KHR is required. That is ICDs may optionally support IMMEDIATE and MAILBOX, but must support FIFO.
- Revision 45, 2015-08-07 (Ian Elliott)
 - Corrected a typo in spec file (type and variable name had wrong case for the imageColorSpace member of the VkSwapchainCreateInfoKHR struct).
 - Corrected a typo in header file (last parameter in PFN_vkGetSurfacePropertiesKHR was missing "KHR" at the end of type: VkSurfacePropertiesKHR).
- Revision 46, 2015-08-20 (Ian Elliott)

- Renamed this extension and all of its enumerations, types, functions, etc. This makes it compliant with the proposed standard for Vulkan extensions.
- Switched from "revision" to "version", including use of the VK_MAKE_VERSION macro in the header file.
- Made improvements to several descriptions.
- Changed the status of several issues from PROPOSED to RESOLVED, leaving no unresolved issues.
- Resolved several TODOs, did miscellaneous cleanup, etc.
- Revision 47, 2015-08-20 (Ian Elliott—porting a 2015-07-29 change from James Jones)
 - Moved the surface transform enums to VK_WSI_swapchain so they could be re-used by VK_WSI_display.
- Revision 48, 2015-09-01 (James Jones)
 - Various minor cleanups.
- Revision 49, 2015-09-01 (James Jones)
 - Restore single-field revision number.
- Revision 50, 2015-09-01 (James Jones)
 - Update Example #4 to include code that illustrates how to use the oldSwapchain field.
- Revision 51, 2015-09-01 (James Jones)
 - Fix example code compilation errors.
- Revision 52, 2015-09-08 (Matthaeus G. Chajdas)
 - Corrected a typo.
- Revision 53, 2015-09-10 (Alon Or-bach)
 - Removed underscore from SWAP_CHAIN left in VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_ KHR.
- Revision 54, 2015-09-11 (Jesse Hall)
 - Described the execution and memory coherence requirements for image transitions to and from VK_IMAGE_ LAYOUT_PRESENT_SOURCE_KHR.
- Revision 55, 2015-09-11 (Ray Smith)
 - Added errors for destroying and binding memory to presentable images
- Revision 56, 2015-09-18 (James Jones)
 - Added fence argument to vkAcquireNextImageKHR
 - Added example of how to meter a CPU thread based on presentation rate.
- Revision 57, 2015-09-26 (Jesse Hall)
 - Replace VkSurfaceDescriptionKHR with VkSurfaceKHR.
 - Added issue 25 with agreed resolution.

- Revision 58, 2015-09-28 (Jesse Hall)
 - Renamed from VK EXT KHR device swapchain to VK EXT KHR swapchain.
- Revision 59, 2015-09-29 (Ian Elliott)
 - Changed vkDestroySwapchainKHR() to return void.
- Revision 60, 2015-10-01 (Jeff Vigil)
 - Added error result VK_ERROR_SURFACE_LOST_KHR.
- Revision 61, 2015-10-05 (Jason Ekstrand)
 - Added the VkCompositeAlpha enum and corresponding structure fields.
- Revision 62, 2015-10-12 (Daniel Rakos)
 - Added VK_PRESENT_MODE_FIFO_RELAXED_KHR.
- Revision 63, 2015-10-15 (Daniel Rakos)
 - Moved surface capability queries to VK_EXT_KHR_surface.
- Revision 64, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_swapchain to VK_KHR_swapchain.
- Revision 65, 2015-10-28 (Ian Elliott)
 - Added optional pResult member to VkPresentInfoKHR, so that per-swapchain results can be obtained from vkQueuePresentKHR().
- Revision 66, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to create and destroy functions.
 - Updated resource transition language.
 - Updated sample code.
- Revision 67, 2015-11-10 (Jesse Hall)
 - Add reserved flags bitmask to VkSwapchainCreateInfoKHR.
 - Modify naming and member ordering to match API style conventions, and so the VkSwapchainCreateInfoKHR image property members mirror corresponding VkImageCreateInfo members but with an *image* prefix.
 - Make VkPresentInfoKHR::pResults non-const; it's an output array parameter.
 - Make pPresentInfo parameter to vkQueuePresentKHR const.
- Revision 68, 2016-04-05 (Ian Elliott)
 - Moved the "validity" include for vkAcquireNextImage to be in its proper place, after the prototype and list of parameters.
 - Clarified language about presentable images, including how they are acquired, when applications can and can't use
 them, etc. As part of this, removed language about "ownership" of presentable images, and replaced it with
 more-consistent language about presentable images being "acquired" by the application.

- 2016-08-23 (Ian Elliott)
 - Update the example code, to use the final API command names, to not have so many characters per line, and to split
 out a new example to show how to obtain function pointers. This code is more similar to the LunarG "cube" demo
 program.
- 2016-08-25 (Ian Elliott)
 - A note was added at the beginning of the example code, stating that it will be removed from future versions of the appendix.

C.2.4 VK_KHR_display

Name String

VK_KHR_display

Extension Type

Instance extension

Registered Extension Number

3

Status

Draft.

Last Modified Date

11/10/2015

Revision

21

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- James Jones, NVIDIA
- Norbert Nopper, Freescale
- · Jeff Vigil, Qualcomm
- · Daniel Rakos, AMD

Contacts

- James Jones (jajones at nvidia.com)
- Norbert Nopper (Norbert.Nopper at freescale.com)

This extension provides the API to enumerate displays and available modes on a given device.

C.2.4.1 New Object Types

- VkDisplayKHR
- VkDisplayModeKHR

C.2.4.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_DISPLAY_MODE_CREATE_INFO_KHR
 - VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR

C.2.4.3 New Enums

• VkDisplayPlaneAlphaFlagBitsKHR

C.2.4.4 New Structures

- VkDisplayPropertiesKHR
- VkDisplayModeParametersKHR
- VkDisplayModePropertiesKHR
- VkDisplayModeCreateInfoKHR
- VkDisplayPlanePropertiesKHR
- VkDisplayPlaneCapabilitiesKHR
- VkDisplaySurfaceCreateInfoKHR

C.2.4.5 New Functions

- vkGetPhysicalDeviceDisplayPropertiesKHR
- vkGetPhysicalDeviceDisplayPlanePropertiesKHR
- vkGetDisplayPlaneSupportedDisplaysKHR
- vkGetDisplayModePropertiesKHR
- vkCreateDisplayModeKHR
- vkGetDisplayPlaneCapabilitiesKHR
- vkCreateDisplayPlaneSurfaceKHR

C.2.4.6 Issues

1) Which properties of a mode should be fixed in the mode info Vs. settable in some other function when setting the mode? E.g., do we need to double the size of the mode pool to include both stereo and non-stereo modes? YUV and RGB scanout even if they both take RGB input images? BGR Vs. RGB input? etc.

PROPOSED RESOLUTION: Many modern displays support at most a handful of resolutions and timings natively. Other "modes" are expected to be supported using scaling hardware on the display engine or GPU. Other properties, such as rotation and mirroring shouldn't require duplicating hardware modes just to express all combinations. Further, these properties may be implemented on a per-display or per-overlay granularity.

To avoid the exponential growth of modes as mutable properties are added, as was the case with EGLConfig/WGL pixel formats/GLXFBConfig, this specification should separate out hardware properties and configurable state into separate objects. Modes and overlay planes will express capabilities of the hardware, while a separate structure will allow applications to configure scaling, rotation, mirroring, color keys, LUT values, alpha masks, etc. for a given swapchain independent of the mode in use. Constraints on these settings will be established by properties of the immutable objects.

Note the resolution of this issue may affect issue (5) as well.

2) What properties of a display itself are useful?

PROPOSED RESOLUTION: This issue is too broad. It was meant to prompt general discussion, but resolving this issue amounts to completing this specification. All interesting properties should be included. The issue will remain as a placeholder since removing it would make it hard to parse existing discussion notes that refer to issues by number.

3) How are multiple overlay planes within a display or mode enumerated?

PROPOSED RESOLUTION: They are referred to by an index. Each display will report the number of overlay planes it contains.

4) Should swapchains be created relative to a mode or a display?

PROPOSED RESOLUTION: When using this extension, swapchains are created relative to a mode and a plane. The mode implies the display object the swapchain will present to. If the specified mode is not the display's current mode, the new mode will be applied when the first image is presented to the swapchain, and the default operating system mode, if any, will be restored when the swapchain is destroyed.

5) Should users query generic ranges from displays and construct their own modes explicitly using those constraints rather than querying a fixed set of modes (Most monitors only have one real "mode" these days, even though many support relatively arbitrary scaling, either on the monitor side or in the GPU display engine, making "modes" something of a relic/compatibility construct).

PROPOSED RESOLUTION: Expose both. Display info structures will expose a set of predefined modes, as well as any attributes necessary to construct a customized mode.

6) Is it fine if we return the display and display mode handles in the structure used to query their properties?

PROPOSED RESOLUTION: Yes.

7) Is there a possibility that not all displays of a device work with all of the present queues of a device? If yes, how do we determine which displays work with which present queues?

PROPOSED RESOLUTION: No known hardware has such limitations, but determining such limitations is supported automatically using the existing VK_EXT_KHR_surface and VK_EXT_KHR_swapchain query mechanisms.

8) Should all presentation need to be done relative to an overlay plane, or can a display mode + display be used alone to target an output?

PROPOSED RESOLUTION: Require specifying a plane explicitly.

9) Should displays have an associated window system display, such as an HDC or Display*?

PROPOSED RESOLUTION: No. Displays are independent of any windowing system in use on the system. Further, neither HDC nor Display* refer to a physical display object.

10) Are displays queried from a physical GPU or from a device instance?

PROPOSED RESOLUTION: Developers prefer to query modes directly from the physical GPU so they can use display information as an input to their device selection algorithms prior to device creation. This avoids the need to create dummy device instances to enumerate displays.

This preference must be weighed against the extra initialization that must be done by driver vendors prior to device instance creation to support this usage.

11) Should displays and/or modes be dispatchable objects? If functions are to take displays, overlays, or modes as their first parameter, they must be dispatchable objects as defined in Khronos bug 13529. If they aren't added to the list of dispatchable objects, functions operating on them must take some higher-level object as their first parameter. There is no performance case against making them dispatchable objects, but they would be the first extension objects to be dispatchable.

PROPOSED RESOLUTION: Do not make displays or modes dispatchable. They will dispatch based on their associated physical device.

12) Should hardware cursor capabilities be exposed?

PROPOSED RESOLUTION: Defer. This could be a separate extension on top of the base WSI specs.

if they are one physical display device to an end user, but may internally be implemented as two side-by-side displays using the same display engine (and sometimes cabling) resources as two physically separate display devices.

PROPOSED RESOLUTION: Tiled displays will appear as a single display object in this API.

14) Should the raw EDID data be included in the display information?

PROPOSED RESOLUTION: None. Unclear whether this is a good idea. Provides a path for forward-compatibility as new EDID extensions are introduced, but may be complicated by the outcome of issue 13.

15) Should min and max scaling factor capabilities of overlays be exposed?

PROPOSED RESOLUTION: Yes. This is exposed indirectly by allowing applications to query the min/max position and extent of the source and destination regions from which image contents are fetched by the display engine when using a particular mode and overlay pair.

16) Should devices be able to expose planes that can be moved between displays? If so, how?

PROPOSED RESOLUTION: None.

17) Should there be a way to destroy display modes? If so, does it support destroying "built in" modes?

PROPOSED RESOLUTION: None.

18) What should the lifetime of display and built-in display mode objects be?

PROPOSED RESOLUTION: The lifetime of the instance. These objects can not be destroyed. A future extension may be added to expose a way to destroy these objects and/or support display hotplug.

19) Should persistent mode for smart panels be enabled/disabled at swap chain creation time, or on a per-present basis.

PROPOSED RESOLUTION: On a per-present basis.

C.2.4.7 Examples

Example 1

Enumerating displays, display modes, and planes, and creating a VkSurfaceKHR

```
extern VkBool32 ModeMatchesMyCriteria(const VkDisplayModePropertiesKHR *m);
extern VkInstance instance;
extern VkPhysicalDevice physDevice;
extern VkSurfaceKHR surface;

uint32_t displayCount, planeCount, i, j, k;
VkDisplayPropertiesKHR* pDisplayProps;
VkDisplayPlanePropertiesKHR* pPlaneProps;
```

```
VkDisplayModeKHR myMode = VK_NULL_HANDLE;
VkDisplayKHR myDisplay = VK_NULL_HANDLE;
uint32_t bestPlane = UINT32_MAX;
VkDisplayPlaneAlphaFlagBitsKHR alphaMode = 0;
PFN_vkGetPhysicalDeviceDisplayPropertiesKHR \leftrightarrow
    pfnGetPhysicalDeviceDisplayPropertiesKHR;
PFN_vkGetPhysicalDeviceDisplayPlanePropertiesKHR \leftarrow
    pfnGetPhysicalDeviceDisplayPlanePropertiesKHR;
PFN_vkGetDisplayModePropertiesKHR pfnGetDisplayModePropertiesKHR;
PFN_vkGetDisplayPlaneCapabilitiesKHR pfnGetDisplayPlaneCapabilitiesKHR;
PFN_vkGetDisplayPlaneSupportedDisplaysKHR pfnGetDisplayPlaneSupportedDisplaysKHR;
PFN_vkCreateDisplayPlaneSurfaceKHR pfnCreateDisplayPlaneSurfaceKHR;
pfnGetPhysicalDeviceDisplayPropertiesKHR =
    (PFN_vkGetPhysicalDeviceDisplayPropertiesKHR)
    vkGetInstanceProcAddr(instance, "vkGetPhysicalDeviceDisplayPropertiesKHR");
pfnGetPhysicalDeviceDisplayPlanePropertiesKHR =
    (PFN_vkGetPhysicalDeviceDisplayPlanePropertiesKHR)
     wk \texttt{GetInstanceProcAddr(instance, "} vk \texttt{GetPhysicalDeviceDisplayPlanePropertiesKHR"} \ \leftarrow \\
pfnGetDisplayModePropertiesKHR =
    (PFN_vkGetDisplayModePropertiesKHR)
    vkGetInstanceProcAddr(instance, "vkGetDisplayModePropertiesKHR");
pfnGetDisplayPlaneCapabilitiesKHR =
    (PFN_vkGetDisplayPlaneCapabilitiesKHR)
    vkGetInstanceProcAddr(instance, "vkGetDisplayPlaneCapabilitiesKHR");
pfnGetDisplayPlaneSupportedDisplaysKHR =
    (PFN_vkGetDisplayPlaneSupportedDisplaysKHR)
    vkGetInstanceProcAddr(instance, "vkGetDisplayPlaneSupportedDisplaysKHR");
pfnCreateDisplayPlaneSurfaceKHR =
    (PFN_vkCreateDisplayPlaneSurfaceKHR)
    vkGetInstanceProcAddr(instance, "vkCreateDisplayPlaneSurfaceKHR");
// Get a list of displays on a physical device
displayCount = 0;
pfnGetPhysicalDeviceDisplayPropertiesKHR(physDevice, &displayCount, NULL);
pDisplayProps = (VkDisplayPropertiesKHR*) malloc(sizeof(VkDisplayPropertiesKHR) * \leftarrow
    displayCount);
pfnGetPhysicalDeviceDisplayPropertiesKHR(physDevice, &displayCount, pDisplayProps) ←
// Get a list of display planes on a physical device
planeCount = 0;
pfnGetPhysicalDeviceDisplayPlanePropertiesKHR (physDevice, &planeCount, NULL);
pPlaneProps = (VkDisplayPlanePropertiesKHR*)malloc(sizeof( ↔
    VkDisplayPlanePropertiesKHR) * planeCount);
pfnGetPhysicalDeviceDisplayPlanePropertiesKHR(physDevice, &planeCount, pPlaneProps \leftarrow
   );
// Get the list of pModes each display supports
for (i = 0; i < displayCount; ++i)</pre>
    VkDisplayKHR display = pDisplayProps[i].display;
    VkDisplayModePropertiesKHR* pModes;
    uint32_t modeCount;
```

```
vkGetDisplayModePropertiesKHR(physDevice, display, &modeCount, NULL);
pModes = (VkDisplayModePropertiesKHR*)malloc(sizeof(VkDisplayModePropertiesKHR ↔
   ) * modeCount);
vkGetDisplayModePropertiesKHR(physDevice, display, &modeCount, pModes);
myMode = VK_NULL_HANDLE;
for (j = 0; j < modeCount; ++j)
    const VkDisplayModePropertiesKHR* mode = &pModes[i];
    if (ModeMatchesMyCriteria(mode))
        myMode = mode->displayMode;
        break;
}
free (pModes);
// If there are no usable pModes found then check the next display.
if (myMode == VK_NULL_HANDLE)
    continue;
// Find a plane that matches these criteria, in order of preference:
// -Is compatible with the chosen display + mode.
// -Isn't currently bound to another display.
// -Supports per-pixel alpha, if possible.
for (j = 0; j < planeCount; ++j)
{
    uint32_t supportedCount = 0;
    VkDisplayKHR* pSupportedDisplays;
    VkDisplayPlaneCapabilitiesKHR planeCaps;
    // See if the plane is compatible with the current display.
    pfnGetDisplayPlaneSupportedDisplaysKHR(physDevice, j, &supportedCount, ←
       NULL);
    // Plane doesn't support any displays. This might happen on a card
    // that has a fixed mapping between planes and connectors when no
    // displays are currently attached to this plane's connector, for
    // example.
    if (supportedCount == 0)
        continue;
    pSupportedDisplays = (VkDisplayKHR*) malloc(sizeof(VkDisplayKHR)* \leftrightarrow 
        supportedCount);
    pfnGetDisplayPlaneSupportedDisplaysKHR(physDevice, j, &supportedCount, ↔
        pSupportedDisplays);
    for (k = 0; k < supportedCount; ++k)</pre>
        if (pSupportedDisplays[k] == display) {
            // If no supported plane has yet been found, this is
            // currently the best available plane.
            if (bestPlane == UINT32_MAX)
                bestPlane = j;
            break;
```

```
// If the plane can't be used with the chosen display, keep looking.
        // Each display must have at least one compatible plane.
        if (k == supportedCount)
            continue;
        // If the plane passed the above test and is currently bound to the
        // desired display, or is not in use, it is the best plane found so
        if ((pPlaneProps[j].currentDisplay == VK_NULL_HANDLE) &&
            (pPlaneProps[j].currentDisplay == display))
           bestPlane = j;
        else
            continue;
        pfnGetDisplayPlaneCapabilitiesKHR (physDevice, myMode, j, &planeCaps);
        // Prefer a plane that supports per-pixel alpha.
        if (planeCaps.supportedAlpha & VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR)
            // This is the perfect plane for the given criteria. Use it.
            bestPlane = j;
            alphaMode = VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR;
            break;
    }
}
free(pDisplayProps);
if (myDisplay == VK_NULL_HANDLE || myMode == VK_NULL_HANDLE) {
    // No suitable display + mode could be found. Abort.
   abort();
} else {
    // Success. Create a VkSurfaceKHR object for this plane.
   const VkDisplaySurfaceCreateInfoKHR createInfo =
       VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR, // sType
        NULL,
                                                            // pNext
        Ο,
                                                             // flags
        myMode,
                                                            // displayMode
                                                            // planeIndex
        bestPlane,
                                                            // planeStackIndex
        pPlaneProps[bestPlane].currentStackIndex,
       VK_SURFACE_TRANSFORM_IDENTITY_KHR,
                                                            // transform
        1.0f,
                                                            // globalAlpha
                                                             // alphaMode
        alphaMode,
        . . .
   pfnCreateDisplayPlaneSurfaceKHR(instance, &createInfo, NULL, &surface);
```

C.2.4.8 Version History

• Revision 1, 2015-02-24 (James Jones)

- Initial draft
- Revision 2, 2015-03-12 (Norbert Nopper)
 - Added overlay enumeration for a display.
- Revision 3, 2015-03-17 (Norbert Nopper)
 - Fixed typos and namings as discussed in Bugzilla.
 - Reordered and grouped functions.
 - Added functions to query count of display, mode and overlay.
 - Added native display handle, which is maybe needed on some platforms to create a native Window.
- Revision 4, 2015-03-18 (Norbert Nopper)
 - Removed primary and virtualPostion members (see comment of James Jones in Bugzilla).
 - Added native overlay handle to info structure.
 - Replaced, with; in struct.
- Revision 6, 2015-03-18 (Daniel Rakos)
 - Added WSI extension suffix to all items.
 - Made the whole API more "Vulkanish".
 - Replaced all functions with a single vkGetDisplayInfoKHR function to better match the rest of the API.
 - Made the display, display mode, and overlay objects be first class objects, not subclasses of VkBaseObject as they
 don't support the common functions anyways.
 - Renamed *Info structures to *Properties.
 - Removed overlayIndex field from VkOverlayProperties as there is an implicit index already as a result of moving to a "Vulkanish" API.
 - Displays aren't get through device, but through physical GPU to match the rest of the Vulkan API. Also this is something ISVs explicitly requested.
 - Added issue (6) and (7).
- Revision 7, 2015-03-25 (James Jones)
 - Added an issues section
 - Added rotation and mirroring flags
- Revision 8, 2015-03-25 (James Jones)
 - Combined the duplicate issues sections introduced in last change.
 - Added proposed resolutions to several issues.
- Revision 9, 2015-04-01 (Daniel Rakos)
 - Rebased extension against Vulkan 0.82.0
- Revision 10, 2015-04-01 (James Jones)
 - Added issues (10) and (11).
 - Added more straw-man issue resolutions, and cleaned up the proposed resolution for issue (4).

- Updated the rotation and mirroring enums to have proper bitmask semantics.
- Revision 11, 2015-04-15 (James Jones)
 - Added proposed resolution for issues (1) and (2).
 - Added issues (12), (13), (14), and (15)
 - Removed pNativeHandle field from overlay structure.
 - Fixed small compilation errors in example code.
- Revision 12, 2015-07-29 (James Jones)
 - Rewrote the guts of the extension against the latest WSI swapchain specifications and the latest Vulkan API.
 - Address overlay planes by their index rather than an object handle and refer to them as "planes" rather than
 "overlays" to make it slightly clearer that even a display with no "overlays" still has at least one base "plane" that
 images can be displayed on.
 - Updated most of the issues.
 - Added an "extension type" section to the specification header.
 - Re-used the VK_EXT_KHR_surface surface transform enumerations rather than redefining them here.
 - Updated the example code to use the new semantics.
- Revision 13, 2015-08-21 (Ian Elliott)
 - Renamed this extension and all of its enumerations, types, functions, etc. This makes it compliant with the proposed standard for Vulkan extensions.
 - Switched from "revision" to "version", including use of the VK MAKE VERSION macro in the header file.
- Revision 14, 2015-09-01 (James Jones)
 - Restore single-field revision number.
- Revision 15, 2015-09-08 (James Jones)
 - Added alpha flags enum.
 - Added premultiplied alpha support.
- Revision 16, 2015-09-08 (James Jones)
 - Added description section to the spec.
 - Added issues 16 18.
- Revision 17, 2015-10-02 (James Jones)
 - Planes are now a property of the entire device rather than individual displays. This allows planes to be moved between multiple displays on devices that support it.
 - Added a function to create a VkSurfaceKHR object describing a display plane and mode to align with the new per-platform surface creation conventions.
 - Removed detailed mode timing data. It was agreed that the mode extents and refresh rate are sufficient for current use cases. Other information could be added back2 in as an extension if it is needed in the future.
 - Added support for smart/persistent/buffered display devices.
- Revision 18, 2015-10-26 (Ian Elliott)

- Renamed from VK_EXT_KHR_display to VK_KHR_display.
- Revision 19, 2015-11-02 (James Jones)
 - Updated example code to match revision 17 changes.
- Revision 20, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to creation functions.
- Revision 21, 2015-11-10 (Jesse Hall)
 - Added VK_DISPLAY_PLANE_ALPHA_OPAQUE_BIT_KHR, and use VkDisplayPlaneAlphaFlagBitsKHR for VkDisplayPlanePropertiesKHR::alphaMode instead of VkDisplayPlaneAlphaFlagsKHR, since it only represents one mode.
 - Added reserved flags bitmask to VkDisplayPlanePropertiesKHR.
 - Use VkSurfaceTransformFlagBitsKHR instead of obsolete VkSurfaceTransformKHR.
 - Renamed vkGetDisplayPlaneSupportedDisplaysKHR parameters for clarity.
- Revision 22, 2015-12-18 (James Jones)
 - Added missing "planeIndex" parameter to vkGetDisplayPlaneSupportedDisplaysKHR()

C.2.5 VK_KHR_display_swapchain

Name String

VK_KHR_display_swapchain

Extension Type

Device extension

Registered Extension Number

4

Status

Draft.

Last Modified Date

11/10/2015

Revision

9

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_swapchain and VK_KHR_display

Contributors

· James Jones, NVIDIA

- · Jeff Vigil, Qualcomm
- · Jesse Hall, Google

Contacts

• James Jones (jajones at nvidia.com)

This extension provides an API to create a swapchain directly on a device's display without any underlying window system.

C.2.5.1 New Object Types

None

C.2.5.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_DISPLAY_PRESENT_INFO_KHR
- Extending VkResult:
 - VK_ERROR_INCOMPATIBLE_DISPLAY_KHR

C.2.5.3 New Enums

None

C.2.5.4 New Structures

• VkDisplayPresentInfoKHR

C.2.5.5 New Functions

• vkCreateSharedSwapchainsKHR

C.2.5.6 Issues

1) Should swapchains sharing images each hold a reference to the images, or should it be up to the application to destroy the swapchains and images in an order that avoids the need for reference counting?

```
PROPOSED RESOLUTION: Take a reference. The lifetime of presentable images is already complex enough.
```

2) Should the srcRect/dstRect parameters be specified as part of the present command, or at swapchain creation time?

PROPOSED RESOLUTION: As part of the presentation command. This allows moving and scaling the image on the screen without the need to respecify the mode or create a new swapchain presentable images.

3) Should srcRect/dstRect be specified as rects, or separate offset/extent values?

PROPOSED RESOLUTION: As rects. Specifying them separately might make it easier for hardware to expose support for one but not the other, but in such cases applications must just take care to obey the reported capabilities and not use non-zero offsets or extents that require scaling, as appropriate.

4) How can applications create multiple swapchains that use the same images?

```
RESOLUTION: By calling vkCreateSharedSwapchainsKHR().
```

An earlier resolution used vkCreateSwapchainKHR(), chaining multiple VkSwapchainCreateInfoKHR structures through pNext. In order to allow each swapchain to also allow other extension structs, a level of indirection was used: VkSwapchainCreateInfoKHR::pNext pointed to a different structure, which had both an sType/pNext for additional extensions, and also had a pointer to the next VkSwapchainCreateInfoKHR structure. The number of swapchains to be created could only be found by walking this linked list of alternating structures, and the pSwapchains out parameter was reinterpreted to be an array of VkSwapchainKHR handles.

Another option considered was a method to specify a "shared" swapchain when creating a new swapchain, such that groups of swapchains using the same images could be built up one at a time. This was deemed unusable because drivers need to know all of the displays an image will be used on when determining which internal formats and layouts to use for that image.

C.2.5.7 Examples

Example 1

Create a swapchain on a display mode and plane

```
// See VK_KHR_display for an example of how to pick a display,
// display mode, plane, and how to create a VkSurfaceKHR for
// that plane.
extern VkPhysicalDevice physDevice;
extern VkDisplayModePropertiesKHR myModeProps;
extern VkSurfaceKHR mySurface;
extern VkDevice device;

uint32_t queueCount, i;
uint32_t presentQueueFamilyIndex = UINT32_MAX;
VkBool32 supportsPresent;

// Find a queue family that supports presenting to this surface
uint32_t familyCount;
vkGetPhysicalDeviceQueueFamilyProperties(physDevice, &familyCount, NULL);

for (i = 0; i < familyCount; ++i)
{</pre>
```

```
vkGetPhysicalDeviceSurfaceSupportKHR(physDevice, i, mySurface, & \leftrightarrow
        supportsPresent);
    if (supportsPresent) {
        // Found a queue family that supports present. See
        // VK_KHR_surface for an example of how to find a queue that
        // supports both presentation and graphics
        presentQueueFamilyIndex = i;
        break;
    }
}
// Figure out which formats and present modes this surface supports.
uint32_t formatCount;
vkGetPhysicalDeviceSurfaceFormatsKHR(physDevice, mySurface, &formatCount, NULL);
VkSurfaceFormatKHR* formats = (VkSurfaceFormatKHR*) malloc(sizeof( ←
   VkSurfaceFormatKHR) * formatCount);
vkGetPhysicalDeviceSurfaceFormatsKHR(physDevice, mySurface, &formatCount, formats) \leftrightarrow
const VkSwapchainCreateInfoKHR createInfo =
    VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR,
                                                     // sType
    NULL,
                                                     // pNext
                                                     // flags
    Ο,
    mySurface,
                                                     // surface
                                                    // minImageCount
    formats[0].format,
                                                    // imageFormat
                                                    // imageColorSpace
    formats[0].colorSpace,
    myModeProps.parameters.visibleRegion,
                                                    // imageExtent
                                                    // imageArrayLayers
                                                    // imageUsage
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,
    VK_SHARING_MODE_EXCLUSIVE,
                                                    // imageSharingMode
                                                    // queueFamilyIndexCount
                                                    // pQueueFamilyIndices
    NULL,
    VK_SURFACE_TRANSFORM_IDENTITY_BIT_KHR,
                                                    // preTransform
    VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR,
                                                    // compositeAlpha
                                                    // presentMode
    VK_PRESENT_MODE_FIFO_KHR,
    VK_TRUE,
                                                    // clipped
    VK_NULL_HANDLE
                                                    // oldSwapchain
};
VkSwapchainKHR swapchain;
// This is equivalent to vkCreateSwapchainKHR.
vkCreateSharedSwapchainsKHR(device, 1, &createInfo, NULL, &swapchain);
```

C.2.5.8 Version History

- Revision 1, 2015-07-29 (James Jones)
 - Initial draft
- Revision 2, 2015-08-21 (Ian Elliott)

- Renamed this extension and all of its enumerations, types, functions, etc. This makes it compliant with the proposed standard for Vulkan extensions.
- Switched from "revision" to "version", including use of the VK MAKE VERSION macro in the header file.
- Revision 3, 2015-09-01 (James Jones)
 - Restore single-field revision number.
- Revision 4, 2015-09-08 (James Jones)
 - Allow creating multiple swap chains that share the same images using a single call to vkCreateSwapChainKHR().
- Revision 5, 2015-09-10 (Alon Or-bach)
 - Removed underscores from SWAP_CHAIN in two enums.
- Revision 6, 2015-10-02 (James Jones)
 - Added support for smart panels/buffered displays.
- Revision 7, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_display_swapchain to VK_KHR_display_swapchain.
- Revision 8, 2015-11-03 (Daniel Rakos)
 - Updated sample code based on the changes to VK_KHR_swapchain.
- Revision 9, 2015-11-10 (Jesse Hall)
 - Replaced VkDisplaySwapchainCreateInfoKHR with vkCreateSharedSwapchainsKHR, changing resolution of issue #4.

C.2.6 VK_KHR_android_surface

Name String

VK_KHR_android_surface

Extension Type

Instance extension

Registered Extension Number

9

Last Modified Date

2016-14-01

Revision

6

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0 of the Vulkan API.

• This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- Ian Elliott, LunarG
- Courtney Goeltzenleuchter, LunarG
- · Jesse Hall, Google
- · James Jones, NVIDIA
- · Antoine Labour, Google
- Jon Leech, Khronos
- · David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- · Ray Smith, ARM
- · Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG

Contacts

• Jesse Hall, Google

The VK_KHR_android_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to an ANativeWindow, Android's native surface type. The ANativeWindow represents the producer endpoint of any buffer queue, regardless of consumer endpoint. Common consumer endpoints for ANativeWindows are the system window compositor, video encoders, and application-specific compositors importing the images through a SurfaceTexture.

C.2.6.1 New Object Types

None

C.2.6.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_ANDROID_SURFACE_CREATE_INFO_KHR

C.2.6.3 New Enums

None

C.2.6.4 New Structures

• VkAndroidSurfaceCreateInfoKHR

C.2.6.5 New Functions

• vkCreateAndroidSurfaceKHR

C.2.6.6 Issues

1) Does Android need a way to query for compatibility between a particular physical device (and queue family?) and a specific Android display?

RESOLVED: No. Currently on Android, any GPU is expected to be able to present to the system compositor, and all queue families must support the necessary image layout transitions and synchronization operations.

C.2.6.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft.
- Revision 2, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_android_surface to VK_KHR_android_surface.
- Revision 3, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to surface creation function.
- Revision 4, 2015-11-10 (Jesse Hall)
 - Removed VK_ERROR_INVALID_ANDROID_WINDOW_KHR.
- Revision 5, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.
- Revision 6, 2016-01-14 (James Jones)
 - Moved VK_ERROR_NATIVE_WINDOW_IN_USE_KHR from the VK_KHR_android_surface to the VK_KHR_ surface extension.

C.2.7 VK KHR mir surface

Name String

VK_KHR_mir_surface

Extension Type

Instance extension

Registered Extension Number

8

Last Modified Date

10/28/2015

Revision

4

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- · Ian Elliott, LunarG
- · Courtney Goeltzenleuchter, LunarG
- · Jesse Hall, Google
- James Jones, NVIDIA
- Antoine Labour, Google
- Jon Leech, Khronos
- · David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- · Ray Smith, ARM
- · Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG

Contacts

- · Jesse Hall, Google
- · Ian Elliott, LunarG

The VK_KHR_mir_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to a Mir surface, as well as a query to determine support for rendering to the windows desktop.

C.2.7.1 New Object Types

None

C.2.7.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_MIR_SURFACE_CREATE_INFO_KHR

C.2.7.3 New Enums

None

C.2.7.4 New Structures

• VkMirSurfaceCreateInfoKHR

C.2.7.5 New Functions

- vkCreateMirSurfaceKHR
- vkGetPhysicalDeviceMirPresentationSupportKHR

C.2.7.6 Issues

1) Does Mir need a way to query for compatibility between a particular physical device (and queue family?) and a specific Mir connection, screen, window, etc.?

RESOLVED: Yes, vkGetPhysicalDeviceMirPresentationSupportKHR() was added to address this.

C.2.7.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft, based on the previous contents of VK_EXT_KHR_swapchain (later renamed VK_EXT_KHR_surface).
- Revision 2, 2015-10-02 (James Jones)
 - Added vkGetPhysicalDeviceMirPresentationSupportKHR() to resolve issue #1.
- Revision 3, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_mir_surface to VK_KHR_mir_surface.
- Revision 4, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.

C.2.8 VK_KHR_wayland_surface

Name String

VK_KHR_wayland_surface

Extension Type

Instance extension

Registered Extension Number

Last Modified Date

11/28/2015

Revision

5

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- · Ian Elliott, LunarG
- Courtney Goeltzenleuchter, LunarG
- · Jesse Hall, Google
- James Jones, NVIDIA
- Antoine Labour, Google
- Jon Leech, Khronos
- · David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- Graham Sellers, AMD
- · Ray Smith, ARM
- Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG

Contacts

- · Jesse Hall, Google
- Ian Elliott, LunarG

The VK_KHR_wayland_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to a Wayland wl_surface, as well as a query to determine support for rendering to the windows desktop.

C.2.8.1 New Object Types

None

C.2.8.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_WAYLAND_SURFACE_CREATE_INFO_KHR

C.2.8.3 New Enums

None

C.2.8.4 New Structures

• VkWaylandSurfaceCreateInfoKHR

C.2.8.5 New Functions

- vkCreateWaylandSurfaceKHR
- $\bullet \ \, \text{vkGetPhysicalDeviceWaylandPresentationSupportKHR}$

C.2.8.6 Issues

1) Does Wayland need a way to query for compatibility between a particular physical device and a specific Wayland display? This would be a more general query than vkGetPhysicalDeviceSurfaceSupportKHR: if the Wayland- specific query returned true for a (VkPhysicalDevice, struct wl_display*) pair, then the physical device could be assumed to support presentation to any VkSurface for surfaces on the display.

RESOLVED: Yes. vkGetPhysicalDeviceWaylandPresentationSupportKHR() was added to address this issue.

C.2.8.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft, based on the previous contents of VK_EXT_KHR_swapchain (later renamed VK_EXT_KHR_surface).
- Revision 2, 2015-10-02 (James Jones)
 - Added vkGetPhysicalDeviceWaylandPresentationSupportKHR() to resolve issue #1.
 - Adjusted wording of issue #1 to match the agreed-upon solution.
 - Renamed "window" parameters to "surface" to match Wayland conventions.
- Revision 3, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_wayland_surface to VK_KHR_wayland_surface.
- Revision 4, 2015-11-03 (Daniel Rakos)
- Added allocation callbacks to vkCreateWaylandSurfaceKHR.
- Revision 5, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.

C.2.9 VK_KHR_win32_surface

Name String

VK_KHR_win32_surface

Extension Type

Instance extension

Registered Extension Number

10

Last Modified Date

11/28/2015

Revision

5

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- Ian Elliott, LunarG
- Courtney Goeltzenleuchter, LunarG
- Jesse Hall, Google
- James Jones, NVIDIA
- Antoine Labour, Google
- Jon Leech, Khronos
- David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- Ray Smith, ARM
- · Jeff Vigil, Qualcomm
- Chia-I Wu, LunarG

Contacts

- Jesse Hall, Google
- · Ian Elliott, LunarG

The VK_KHR_win32_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to a Win32 HWND, as well as a query to determine support for rendering to the windows desktop.

C.2.9.1 New Object Types

None

C.2.9.2 New Enum Constants

- Extending VkStructureType:
 - VK STRUCTURE TYPE WIN32 SURFACE CREATE INFO KHR

C.2.9.3 New Enums

None

C.2.9.4 New Structures

• VkWin32SurfaceCreateInfoKHR

C.2.9.5 New Functions

- vkCreateWin32SurfaceKHR
- vkGetPhysicalDeviceWin32PresentationSupportKHR

C.2.9.6 Issues

1) Does Win32 need a way to query for compatibility between a particular physical device and a specific screen? Compatibility between a physical device and a window generally only depends on what screen the window is on. However, there isn't an obvious way to identify a screen without already having a window on the screen.

RESOLVED: No. While it may be useful, there isn't a clear way to do this on Win32. However, a method was added to query support for presenting to the windows desktop as a whole.

C.2.9.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft, based on the previous contents of VK_EXT_KHR_swapchain (later renamed VK_EXT_KHR_surface).
- Revision 2, 2015-10-02 (James Jones)
 - Added presentation support query for win32 desktops.
- Revision 3, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_win32_surface to VK_KHR_win32_surface.
- Revision 4, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to vkCreateWin32SurfaceKHR.
- Revision 5, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.

C.2.10 VK_KHR_xcb_surface

Name String

VK_KHR_xcb_surface

Extension Type

Instance extension

Registered Extension Number

6

Last Modified Date

11/28/2015

Revision

6

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- Ian Elliott, LunarG
- Courtney Goeltzenleuchter, LunarG
- Jesse Hall, Google
- James Jones, NVIDIA
- · Antoine Labour, Google
- Jon Leech, Khronos
- · David Mao, AMD
- Norbert Nopper, Freescale
- · Alon Or-bach, Samsung
- · Daniel Rakos, AMD
- Graham Sellers, AMD
- Ray Smith, ARM
- · Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG

Contacts

- Jesse Hall, Google
- · Ian Elliott, LunarG

The VK_KHR_xcb_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to an X11 window, using the XCB client-side library, as well as a query to determine support for rendering via XCB.

C.2.10.1 New Object Types

None

C.2.10.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_XCB_SURFACE_CREATE_INFO_KHR

C.2.10.3 New Enums

None

C.2.10.4 New Structures

• VkXcbSurfaceCreateInfoKHR

C.2.10.5 New Functions

- vkCreateXcbSurfaceKHR
- vkGetPhysicalDeviceXcbPresentationSupportKHR

C.2.10.6 Issues

1) Does XCB need a way to query for compatibility between a particular physical device and a specific screen? This would be a more general query than vkGetPhysicalDeviceSurfaceSupportKHR: If it returned true, then the physical device could be assumed to support presentation to any window on that screen.

RESOLVED: Yes, this is needed for toolkits that want to create a VkDevice before creating a window. To ensure the query is reliable, it must be made against a particular X visual rather than the screen in general.

C.2.10.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft, based on the previous contents of VK_EXT_KHR_swapchain (later renamed VK_EXT_KHR_surface).
- Revision 2, 2015-10-02 (James Jones)
 - Added presentation support query for an (xcb_connection_t*, xcb_visualid_t) pair.
 - Removed "root" parameter from CreateXcbSurfaceKHR(), as it is redundant when a window on the same screen is specified as well.
 - Adjusted wording of issue #1 and added agreed upon resolution.
- Revision 3, 2015-10-14 (Ian Elliott)

- Removed "root" parameter from CreateXcbSurfaceKHR() in one more place.
- Revision 4, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_xcb_surface to VK_KHR_xcb_surface.
- Revision 5, 2015-10-23 (Daniel Rakos)
 - Added allocation callbacks to vkCreateXcbSurfaceKHR.
- Revision 6, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.

C.2.11 VK_KHR_xlib_surface

Name String

VK_KHR_xlib_surface

Extension Type

Instance extension

Registered Extension Number

5

Last Modified Date

11/28/2015

Revision

6

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- This extension requires VK_KHR_surface.

Contributors

- · Patrick Doane, Blizzard
- · Jason Ekstrand, Intel
- Ian Elliott, LunarG
- Courtney Goeltzenleuchter, LunarG
- · Jesse Hall, Google
- James Jones, NVIDIA
- Antoine Labour, Google
- Jon Leech, Khronos
- · David Mao, AMD
- Norbert Nopper, Freescale
- Alon Or-bach, Samsung

- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- · Ray Smith, ARM
- Jeff Vigil, Qualcomm
- · Chia-I Wu, LunarG

Contacts

- Jesse Hall, Google
- · Ian Elliott, LunarG

The VK_KHR_xlib_surface extension is an instance extension. It provides a mechanism to create a VkSurfaceKHR object (defined by the VK_KHR_surface extension) that refers to an X11 window, using the Xlib client-side library, as well as a query to determine support for rendering via Xlib.

C.2.11.1 New Object Types

None

C.2.11.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_XLIB_SURFACE_CREATE_INFO_KHR

C.2.11.3 New Enums

None

C.2.11.4 New Structures

• VkXlibSurfaceCreateInfoKHR

C.2.11.5 New Functions

- vkCreateXlibSurfaceKHR
- vkGetPhysicalDeviceXlibPresentationSupportKHR

C.2.11.6 Issues

1) Does X11 need a way to query for compatibility between a particular physical device and a specific screen? This would be a more general query than vkGetPhysicalDeviceSurfaceSupportKHR: If it returned true, then the physical device could be assumed to support presentation to any window on that screen.

RESOLVED: Yes, this is needed for toolkits that want to create a VkDevice before creating a window. To ensure the query is reliable, it must be made against a particular X visual rather than the screen in general.

C.2.11.7 Version History

- Revision 1, 2015-09-23 (Jesse Hall)
 - Initial draft, based on the previous contents of VK_EXT_KHR_swapchain (later renamed VK_EXT_KHR_surface).
- Revision 2, 2015-10-02 (James Jones)
 - Added presentation support query for (Display*, VisualID) pair.
 - Removed "root" parameter from CreateXlibSurfaceKHR(), as it is redundant when a window on the same screen is specified as well.
 - Added appropriate X errors.
 - Adjusted wording of issue #1 and added agreed upon resolution.
- Revision 3, 2015-10-14 (Ian Elliott)
 - Renamed this extension from VK_EXT_KHR_x11_surface to VK_EXT_KHR_xlib_surface.
- Revision 4, 2015-10-26 (Ian Elliott)
 - Renamed from VK_EXT_KHR_xlib_surface to VK_KHR_xlib_surface.
- Revision 5, 2015-11-03 (Daniel Rakos)
 - Added allocation callbacks to vkCreateXlibSurfaceKHR.
- Revision 6, 2015-11-28 (Daniel Rakos)
 - Updated the surface create function to take a pCreateInfo structure.

C.3 VK_EXT_debug_marker

Name String

VK_EXT_debug_marker

Extension Type

Device extension

Registered Extension Number

23

Last Modified Date

2016-04-23

Revision

3

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0.11 of the Vulkan API.

Contributors

- · Baldur Karlsson
- Dan Ginsburg, Valve
- · Jon Ashburn, LunarG
- Kyle Spagnoli, NVIDIA

Contacts

· Baldur Karlsson

The VK_EXT_debug_marker extension is a device extension. It introduces concepts of object naming and tagging, for better tracking of Vulkan objects, as well as additional commands for recording annotations of named sections of a workload to aid organisation and offline analysis in external tools.

C.3.1 New Object Types

None

C.3.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_NAME_INFO_EXT
 - VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_TAG_INFO_EXT
 - VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT

C.3.3 New Enums

None

C.3.4 New Structures

- VkDebugMarkerObjectNameInfoEXT
- VkDebugMarkerObjectTagInfoEXT
- VkDebugMarkerMarkerInfoEXT

C.3.5 New Functions

- vkDebugMarkerSetObjectTagEXT
- vkDebugMarkerSetObjectNameEXT
- vkCmdDebugMarkerBeginEXT
- vkCmdDebugMarkerEndEXT
- vkCmdDebugMarkerInsertEXT

C.3.6 Examples

Example 1

Associate a name with an image, for easier debugging in external tools or with validation layers that can print a friendly name when referring to objects in error messages.

```
extern VkDevice device;
extern VkImage image;
// Must call extension functions through a function pointer:
{\tt PFN\_vkDebugMarkerSetObjectNameEXT} \ \ {\tt pfnDebugMarkerSetObjectNameEXT} \ = \ ( \ \hookleftarrow \ )
    PFN_vkDebugMarkerSetObjectNameEXT)vkGetDeviceProcAddr(device, " <math>\leftarrow
    vkDebugMarkerSetObjectNameEXT");
// Set a name on the image
const VkDeviceCreateInfo imageNameInfo =
    VK STRUCTURE TYPE DEBUG MARKER OBJECT NAME INFO EXT, // sType
                                                       // pNext
    NULL,
    VK_DEBUG_REPORT_OBJECT_TYPE_IMAGE_EXT,
                                                      // objectType
                                                       // object
    (uint64_t)image,
    "Brick Diffuse Texture",
                                                       // pObjectName
};
pfnDebugMarkerSetObjectNameEXT(device, &imageNameInfo);
// A subsequent error might print:
// Image 'Brick Diffuse Texture' (0xc0dec0dedeadbeef) is used in a
// command buffer with no memory bound to it.
```

Example 2

Annotating regions of a workload with naming information so that offline analysis tools can display a more usable visualisation of the commands submitted.

```
extern VkDevice device;
extern VkCommandBuffer commandBuffer;
// Must call extension functions through a function pointer:
PFN_vkCmdDebugMarkerBeginEXT pfnCmdDebugMarkerBeginEXT = ( \leftrightarrow
            {\tt PFN\_vkCmdDebugMarkerBeginEXT)} \ {\tt vkGetDeviceProcAddr(device, "} \leftarrow
             vkCmdDebugMarkerBeginEXT");
 PFN\_vkCmdDebugMarkerEndEXT \ pfnCmdDebugMarkerEndEXT \ = \ (PFN\_vkCmdDebugMarkerEndEXT) \ \leftrightarrow \ (PFN\_vkCmdDebugMarkerEndEXT) \ \leftrightarrow \ (PFN\_vkCmdDebugMarkerEndEXT) \ (PFN\_vkCmdDebugMarkerEn
             vkGetDeviceProcAddr(device, "vkCmdDebugMarkerEndEXT");
PFN_vkCmdDebugMarkerInsertEXT pfnCmdDebugMarkerInsertEXT = ( ←
             PFN vkCmdDebugMarkerInsertEXT)vkGetDeviceProcAddr(device, " ←
             vkCmdDebugMarkerInsertEXT");
// Describe the area being rendered
const VkDebugMarkerMarkerInfoEXT houseMarker =
{
              VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT, // sType
             NULL,
                                                                                                                                                                                              // pNext
              "Brick House",
                                                                                                                                                                                              // pMarkerName
                                                                                                                                                                                            // color
               { 1.0f, 0.0f, 0.0f, 1.0f },
};
```

```
// Start an annotated group of calls under the 'Brick House' name
pfnCmdDebugMarkerBeginEXT(commandBuffer, &houseMarker);
    // A mutable structure for each part being rendered
    VkDebugMarkerMarkerInfoEXT housePartMarker =
        VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT, // sType
        NULT.
                                                        // pMarkerName
        { 0.0f, 0.0f, 0.0f, 0.0f },
                                                        // color
    };
    // Set the name and insert the marker
    housePartMarker.pMarkerName = "Walls";
    pfnCmdDebugMarkerInsertEXT(commandBuffer, &housePartMarker);
    // Insert the drawcall for the walls
    vkCmdDrawIndexed(commandBuffer, 1000, 1, 0, 0, 0);
    // Insert a recursive region for two sets of windows
    housePartMarker.pMarkerName = "Windows";
    pfnCmdDebugMarkerBeginEXT(commandBuffer, &housePartMarker);
    {
        vkCmdDrawIndexed(commandBuffer, 75, 6, 1000, 0, 0);
        vkCmdDrawIndexed(commandBuffer, 100, 2, 1450, 0, 0);
    pfnCmdDebugMarkerEndEXT(commandBuffer);
    housePartMarker.pMarkerName = "Front Door";
    pfnCmdDebugMarkerInsertEXT(commandBuffer, &housePartMarker);
    vkCmdDrawIndexed(commandBuffer, 350, 1, 1650, 0, 0);
    housePartMarker.pMarkerName = "Roof";
    pfnCmdDebugMarkerInsertEXT(commandBuffer, &housePartMarker);
    vkCmdDrawIndexed(commandBuffer, 500, 1, 2000, 0, 0);
// End the house annotation started above
pfnCmdDebugMarkerEndEXT(commandBuffer);
```

C.3.7 Issues

1) Should the tag or name for an object be specified using the pNext parameter in the object's CreateInfo structure?

RESOLVED: No. While this fits with other Vulkan patterns and would allow more type safety and future proofing against future objects, it has notable downsides. In particular passing the name at CreateInfo time does not allow renaming, prevents late binding of naming information, and does not allow naming of implicitly created objects such as queues and swap chain images.

2) Should the command annotation functions vkCmdDebugMarkerBeginEXT() and vkCmdDebugMarkerEndEXT() support the ability to specify a color?

RESOLVED: Yes. The functions have been expanded to take an optional color which can be used at will by implementations consuming the command buffer annotations in their visualisation.

3) Should the functions added in this extension accept an extensible structure as their parameter for a more flexible API, as opposed to direct function parameters? If so, which functions?

RESOLVED: Yes. All functions have been modified to take a structure type with extensible pNext pointer, to allow future extensions to add additional annotation information in the same commands.

C.3.8 Version History

- Revision 1, 2015-02-24 (Baldur Karlsson)
 - Initial draft, based on LunarG marker spec
- Revision 2, 2015-02-26 (Baldur Karlsson)
 - Renamed Dbg to DebugMarker in function names
 - Allow markers in secondary command buffers under certain circumstances
 - Minor language tweaks and edits
- Revision 3, 2015-04-23 (Baldur Karlsson)
 - Reorganise spec layout to closer match desired organisation
 - Added optional color to markers (both regions and inserted labels)
 - Changed functions to take extensible structs instead of direct function parameters

C.4 VK_EXT_debug_report

Due to the nature of the Vulkan interface, there is very little error information available to the developer/application. By enabling optional validation layers and using the Debug Report extension a developer has much more detailed feedback on the application's use of Vulkan.

This extension adds two entry points (vkCreateDebugReportCallbackEXT,

vkDestroyDebugReportCallbackEXT) and an extension structure that together define a way for layers and the implementation to call back to the application for events of interest to the application. An application enables this extension by including "VK_EXT_debug_report" in the list of ppEnabledExtensionNames at vkCreateInstance. Layers indicating support will return VK_EXT_debug_report in the extension list queried via vkEnumerateInstanceExtensionProperties.

To register a callback, an application uses vkCreateDebugReportCallbackEXT.

- instance the instance the callback will be logged on.
- pCreateInfo points to a VkDebugReportCallbackCreateInfoEXT structure which defines the conditions under which this callback will be called.
- pCallback is a pointer to record the VkDebugReportCallbackEXT object created.

Valid Usage

- instance must be a valid VkInstance handle
- pCreateInfo must be a pointer to a valid VkDebuqReportCallbackCreateInfoEXT structure
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- pCallback must be a pointer to a VkDebugReportCallbackEXT handle

Return Codes

Success

• VK SUCCESS

Failure

• VK_ERROR_OUT_OF_HOST_MEMORY

The definition of VkDebugReportCallbackCreateInfoEXT is:

- *sType* is the type of this structure.
- pNext is NULL or a pointer to an extension-specific structure.
- flags indicate which event(s) will cause this callback to be called. Flags are interpreted as bitmasks and multiple may be set. Bits which can be set include:

```
typedef enum VkDebugReportFlagBitsEXT {
    VK_DEBUG_REPORT_INFORMATION_BIT_EXT = 0x00000001,
    VK_DEBUG_REPORT_WARNING_BIT_EXT = 0x00000002,
    VK_DEBUG_REPORT_PERFORMANCE_WARNING_BIT_EXT = 0x00000004,
    VK_DEBUG_REPORT_ERROR_BIT_EXT = 0x00000008,
    VK_DEBUG_REPORT_DEBUG_BIT_EXT = 0x00000010,
} VkDebugReportFlagBitsEXT;
```

- VK_DEBUG_REPORT_ERROR_BIT_EXT indicates an error that may cause undefined results, including an application crash.
- VK_DEBUG_REPORT_WARNING_BIT_EXT indicates an unexpected use. E.g. Not destroying objects prior to
 destroying the containing object or potential inconsistencies between descriptor set layout and the layout in the
 corresponding shader, etc.
- VK_DEBUG_REPORT_PERFORMANCE_WARNING_BIT_EXT indicates a potentially non-optimal use of Vulkan. E.g. using **vkCmdClearImage** when a RenderPass load_op would have worked.
- VK_DEBUG_REPORT_INFORMATION_BIT_EXT indicates an informational message such as resource details that may be handy when debugging an application.
- VK_DEBUG_REPORT_DEBUG_BIT_EXT indicates diagnostic information from the loader and layers.
- pfnCallback is the application callback function to call.
- pUserData is user data to be passed to the callback.

For each VkDebugReportCallbackEXT that is created the flags determine when that function is called. A callback will be made for issues that match any bit set in its flags. The callback will come directly from the component that detected the event, unless some other layer intercepts the calls for its own purposes (filter them in different way, log to system error log, etc.) An application may receive multiple callbacks if multiple VkDebugReportCallbackEXT objects were created. A callback will always be executed in the same thread as the originating Vulkan call. A callback may be called from multiple threads simultaneously (if the application is making Vulkan calls from multiple threads).

Valid Usage

- stype must be VK_STRUCTURE_TYPE_DEBUG_REPORT_CALLBACK_CREATE_INFO_EXT
- pNext must be NULL
- flags must be a valid combination of VkDebugReportFlagBitsEXT values
- flags must not be 0

The prototype for the callback function implemented by the app is:

```
const char*
const char*
void*

pLayerPrefix,
pMessage,
pUserData);
```

- flags indicates the VkDebuqReportFlagBitsEXT that triggered this callback.
- *objType* the type of object being used / created at the time the event was triggered. Object types are defined by VkDebugReportObjectTypeEXT.
- object gives the object where the issue was detected. object may be VK_NULL_OBJECT if there is no object associated with the event.
- location is a component (layer, driver, loader) defined value that indicates the *location* of the trigger. This is an optional value.
- messageCode a layer defined value indicating what test triggered this callback.
- pLayerPrefix is the abbreviation of the component making the callback.
- pMessage is a null-terminated string detailing the trigger conditions.
- pUserData is the user data given when the DebugReportCallback was created.

The callback returns a VkBool32 that indicates to the calling layer if the Vulkan call should be aborted or not. Applications should always return VK_FALSE so that they see the same behavior with and without validation layers enabled.

If the application returns VK_TRUE from its callback and the Vulkan call being aborted returns a VkResult, the layer will return VK_ERROR_VALIDATION_FAILED_EXT.



Note

The primary expected use of VK_ERROR_VALIDATION_FAILED_EXT is for validation layer testing. It is not expected that an application would see this this error code during normal use of the validation layers.

To inject its own messages into the debug stream an application uses vkDebugReportMessageEXT.

```
void vkDebugReportMessageEXT(
   VkInstance
                                                  instance,
   VkDebugReportFlagsEXT
                                                  flags,
   VkDebugReportObjectTypeEXT
                                                  objectType,
   uint64 t
                                                  object,
   size_t
                                                  location,
   int32_t
                                                  messageCode,
    const char*
                                                  pLayerPrefix,
    const char*
                                                  pMessage);
```

- instance the instance the callback will be logged on.
- flags indicates the VkDebugReportFlagBitsEXT that triggered this callback.
- objType is the type of object being used / created at the time the event was triggered. Object types are defined by VkDebugReportObjectTypeEXT.

- object is object where the issue was detected. object may be VK_NULL_OBJECT if there is no object associated with the event.
- *location* is a component (layer, driver, loader) defined value that indicates the *location* of the trigger. This is an optional value.
- messageCode is a layer defined value indicating what test triggered this callback.
- *pLayerPrefix* is the abbreviation of the component making the callback.
- pMessage is a null-terminated string detailing the trigger conditions.

The call will propagate through the layers and cause a callback to the application. The parameters are passed on to the callback in addition to the pUserData value that was defined at the time the callback was registered.

Valid Usage

- instance must be a valid VkInstance handle
- flags must be a valid combination of VkDebugReportFlagBitsEXT values
- flags must not be 0
- objectType must be a valid VkDebugReportObjectTypeEXT value
- pLayerPrefix must be a pointer to a valid
- pMessage must be a pointer to a valid
- instance must be a valid VkInstance handle
- flags must be a combination of one or more of VkDebugReportFlagBitsEXT
- objType must be one of VkDebugReportObjectTypeEXT, VK_DEBUG_REPORT_OBJECT_TYPE_UNKNOWN_EXT if object is NULL
- object may be a Vulkan object
- pLayerPrefix must be a NULL terminated string
- pMsq must be a NULL terminated string

To destroy a VkDebugReportCallbackEXT object an application calls vkDestroyDebugReportCallbackEXT.

• instance the instance where the callback was created.

• callback the VkDebugReportCallbackEXT object to destroy.

Valid Usage

- instance must be a valid VkInstance handle
- callback must be a valid VkDebugReportCallbackEXT handle
- If pAllocator is not NULL, pAllocator must be a pointer to a valid VkAllocationCallbacks structure
- callback must have been created, allocated, or retrieved from instance
- If VkAllocationCallbacks were provided when *instance* was created, a compatible set of callbacks must be provided here
- If no VkAllocationCallbacks were provided when instance was created, pAllocator must be NULL

Host Synchronization

• Host access to callback must be externally synchronized

Using the VK_EXT_debug_report allows an application to register multiple callbacks with the validation layers. Some callbacks may log the information to a file, others may cause a debug break point or other application defined behavior. An application can register callbacks even when no validation layers are enabled, but they will only be called for loader and, if implemented, driver events.

To capture issues found while creating or destroying an instance an application can link a VkDebugReportCallbackCreateInfoEXT structure to the pNext element of the VkInstanceCreateInfo structure given to vkCreateInstance. This callback is only valid for the duration of the vkCreateInstance and the vkDestroyInstance call. Use vkCreateDebugReportCallbackEXT to create persistent callback objects.

Example uses: Create three callback objects. One will log errors and warnings to the debug console using Windows OutputDebugString. The second will cause the debugger to break at that callback when an error happens and the third will log warnings to stdout.

```
NULT.
                                                                      // pUserData
};
res = vkCreateDebugReportCallbackEXT(instance, &callback1, &cb1);
if (res != VK_SUCCESS)
   /* Do error handling for VK_ERROR_OUT_OF_MEMORY */
callback.flags = VK_DEBUG_REPORT_ERROR_BIT_EXT;
callback.pfnCallback = myDebugBreak;
callback.pUserData = NULL;
res = vkCreateDebugReportCallbackEXT(instance, &callback, &cb2);
if (res != VK_SUCCESS)
   /* Do error handling for VK_ERROR_OUT_OF_MEMORY */
VkDebugReportCallbackCreateInfoEXT callback3 = {
        VK_STRUCTURE_TYPE_DEBUG_REPORT_CALLBACK_CREATE_INFO_EXT,
                                                                     // sType
                                                                      // pNext
                                                                      // flags
        VK_DEBUG_REPORT_WARNING_BIT_EXT,
                                                                      // pfnCallback
        mystdOutLogger,
                                                                      // pUserData
        NULL
res = vkCreateDebugReportCallbackEXT(instance, &callback3, &cb3);
if (res != VK_SUCCESS)
   /* Do error handling for VK_ERROR_OUT_OF_MEMORY */
/\star remove callbacks when cleaning up \star/
vkDestroyDebugReportCallbackEXT(instance, cb1);
vkDestroyDebugReportCallbackEXT(instance, cb2);
vkDestroyDebugReportCallbackEXT(instance, cb3);
```

Note



In the initial release of the VK_EXT_debug_report extension, the token VK_STRUCTURE_TYPE_DEBUG_REPORT_CREATE_INFO_EXT was used. Starting in VK_EXT_DEBUG_REPORT_SPEC_VERS ION 2 of the extension branch, VK_STRUCTURE_TYPE_DEBUG_REPORT_CALLBACK_CREATE_INFO_EXT is used instead for consistency with Vulkan naming rules. The older enum is still available for backwards compatibility.

C.5 VK_AMD_draw_indirect_count

Name String

VK_AMD_draw_indirect_count

Extension Type

Device extension

Registered Extension Number

34

Last Modified Date

2016-08-23

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0 of the Vulkan API.

Contributors

- · Matthaeus G. Chajdas, AMD
- · Derrick Owens, AMD
- · Graham Sellers, AMD
- · Daniel Rakos, AMD
- Dominik Witczak, AMD

Contacts

• Matthaeus G. Chajdas, AMD (matthaeus.chajdas@amd.com)

This extension allows an application to source the number of draw calls for indirect draw calls from a buffer. This enables applications to generate arbitrary amounts of draw commands and execute them without host intervention.

C.5.1 New Functions

- vkCmdDrawIndirectCountAMD
- vkCmdDrawIndexedIndirectCountAMD

C.5.2 Version History

- Revision 2, 2016-08-23 (Dominik Witczak)
 - Minor fixes
- Revision 1, 2016-07-21 (Matthaeus Chajdas)
 - Initial draft

C.6 VK_AMD_gcn_shader

Name String

VK_AMD_gcn_shader

Extension Type

Device extension

Registered Extension Number

26

Last Modified Date

2016-05-30

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0.15 of the Vulkan API.

Contributors

- · Dominik Witczak, AMD
- Daniel Rakos, AMD
- Rex Xu, AMD
- Graham Sellers, AMD

Contacts

- Dominik Witczak, AMD (dominik.witczak@amd.com)
- SPV_AMD_gcn_shader

C.6.1 Version History

- Revision 1, 2016-05-30 (Dominik Witczak)
 - Initial draft

C.7 VK_AMD_rasterization_order

Name String

VK_AMD_rasterization_order

Extension Type

Device extension

Registered Extension Number

19

Last Modified Date

2016-04-25

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0.11 of the Vulkan API.

Contributors

- · Matthaeus G. Chajdas, AMD
- · Jaakko Konttinen, AMD
- Daniel Rakos, AMD
- · Graham Sellers, AMD
- · Dominik Witczak, AMD

Contacts

· Daniel Rakos, AMD

This extension introduces the possibility for the application to control the order of primitive rasterization. In unextended Vulkan, the following stages are guaranteed to execute in *API order*:

- · depth bounds test
- stencil test, stencil op, and stencil write
- · depth test and depth write
- · occlusion queries
- blending, logic op, and color write

This extension enables applications to opt into a relaxed, implementation defined primitive rasterization order that may allow better parallel processing of primitives and thus enabling higher primitive throughput. It is applicable in cases where the primitive rasterization order is known to not affect the output of the rendering or any differences caused by a different rasterization order are not a concern from the point of view of the application's purpose.

A few examples of cases when using the relaxed primitive rasterization order would not have an effect on the final rendering:

- If the primitives rendered are known to not overlap in framebuffer space.
- If depth testing is used with a comparison operator of VK_COMPARE_OP_LESS, VK_COMPARE_OP_LESS_OR_ EQUAL, VK_COMPARE_OP_GREATER, or VK_COMPARE_OP_GREATER_OR_EQUAL, and the primitives rendered are known to not overlap in clip space.
- If depth testing is not used and blending is enabled for all attachments with a commutative blend operator.

C.7.1 New Object Types

None

C.7.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_RASTERIZATION_ORDER_AMD

C.7.3 New Enums

• VkRasterizationOrderAMD

C.7.4 New Structures

• VkPipelineRasterizationStateRasterizationOrderAMD

C.7.5 New Functions

None

C.7.6 Issues

1) How is this extension useful to application developers?

RESOLVED: Allows them to increase primitive throughput for cases when strict API order rasterization is not important due to the nature of the content, the configuration used, or the requirements towards the output of the rendering.

2) How does this extension interact with content optimizations aiming to reduce overdraw by appropriately ordering the input primitives?

RESOLVED: While the relaxed rasterization order might somewhat limit the effectiveness of such content optimizations, most of the benefits of it are expected to be retained even when the relaxed rasterization order is used so applications are still recommended to apply these optimizations even if they intend to use the extension.

3) Are there any guarantees about the primitive rasterization order when using the new relaxed mode?

RESOLVED: No. The rasterization order is completely implementation dependent in this case but in practice it is expected to partially still follow the order of incoming primitives.

4) Does the new relaxed rasterization order have any adverse effect on repeatability and other invariance rules of the API?

RESOLVED: Yes, in the sense that it extends the list of exceptions when the repeatability requirement does not apply.

C.7.7 Examples

None

C.7.8 Issues

None

C.7.9 Version History

- Revision 1, 2016-04-25 (Daniel Rakos)
 - Initial draft.

C.8 VK_AMD_shader_explicit_vertex_parameter

Name String

VK_AMD_shader_explicit_vertex_parameter

Extension Type

Device extension

Registered Extension Number

22

Last Modified Date

2016-05-10

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0.11 of the Vulkan API.

Contributors

- Matthaeus G. Chajdas, AMD
- Qun Lin, AMD
- · Daniel Rakos, AMD
- · Graham Sellers, AMD
- Rex Xu, AMD

Contacts

- Qun Lin, AMD (quentin.lin@amd.com)
- SPV_AMD_shader_explicit_vertex_parameter

C.8.1 Version History

- Revision 1, 2016-05-10 (Daniel Rakos)
 - Initial draft

C.9 VK_AMD_shader_trinary_minmax

Name String

VK_AMD_shader_trinary_minmax

Extension Type

Device extension

Registered Extension Number

21

Last Modified Date

2016-05-10

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0.11 of the Vulkan API.

Contributors

- Matthaeus G. Chajdas, AMD
- Qun Lin, AMD
- Daniel Rakos, AMD
- Graham Sellers, AMD
- Rex Xu, AMD

Contacts

- Qun Lin, AMD (quentin.lin@amd.com)
- SPV_AMD_shader_trinary_minmax

C.9.1 Version History

- Revision 1, 2016-05-10 (Daniel Rakos)
 - Initial draft

C.10 VK_AMD_negative_viewport_height

Name String

VK_AMD_negative_viewport_height

Extension Type

Device extension

Registered Extension Number

36

Last Modified Date

2016-09-02

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0 of the Vulkan API.

Contributors

- Matthaeus G. Chajdas, AMD
- Graham Sellers, AMD
- Baldur Karlsson

Contacts

• Matthaeus G. Chajdas, AMD (matthaeus.chajdas@amd.com)

This extension allows an application to specify a negative viewport height. The result is that the viewport transformation will flip along the y-axis.

- Revision 1, 2016-09-02 (Matthaeus Chajdas)
 - Initial draft

C.11 VK_IMG_filter_cubic

Name String

VK_IMG_filter_cubic

Extension Type

Device extension

Registered Extension Number

16

Status

Final

Last Modified Date

2016-02-23

Revision

1

Dependencies

• This extension is written against version 1.0 of the Vulkan API.

Contributors

• Tobias Hector, Imagination Technologies

Contacts

• Tobias Hector (tobias.hector@imgtec.com)

VK_IMG_filter_cubic adds an additional, high quality cubic filtering mode to Vulkan, using a Catmull-Rom bicubic filter. Performing this kind of filtering can be done in a shader by using 16 samples and a number of instructions, but this can be inefficient. The cubic filter mode exposes an optimized high quality texture sampling using fixed texture sampling hardware.

C.11.1 New Enum Constants

```
Extending VkFilter:VK_FILTER_CUBIC_IMG
```

C.11.2 Example

Creating a sampler with the new filter for both magnification and minification

```
VkSamplerCreateInfo createInfo =
{
    VK_STRUCTURE_TYPE_SAMPLER_CREATE_INFO // sType
    // Other members set to application-desired values
};

createInfo.magFilter = VK_FILTER_CUBIC_IMG;
createInfo.minFilter = VK_FILTER_CUBIC_IMG;

VkSampler sampler;
VkResult result = vkCreateSampler(
    device,
    &createInfo,
    &sampler);
```

C.11.3 Version History

- Revision 1, 2016-02-23 (Tobias Hector)
 - Initial version

C.12 VK_NV_dedicated_allocation

Name String

VK_NV_dedicated_allocation

Extension Type

Device extension

Registered Extension Number

27

Status

Draft.

Last Modified Date

2016-05-31

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0. of the Vulkan API.

Contributors

Jeff Bolz, NVIDIA

Contacts

• Jeff Bolz (jbolz at nvidia.com)

This extension allows device memory to be allocated for a particular buffer or image resource, which on some devices can significantly improve the performance of that resource. Normal device memory allocations must support memory aliasing and sparse binding, which could interfere with optimizations like framebuffer compression or efficient page table usage. This is important for render targets and very large resources, but need not (and probably should not) be used for smaller resources that can benefit from suballocation.

This extension adds a few small structures to resource creation and memory allocation: a new structure that flags whether am image/buffer will have a dedicated allocation, and a structure indicating the image or buffer that an allocation will be bound to.

C.12.1 New Object Types

None.

C.12.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_IMAGE_CREATE_INFO_NV
 - VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_BUFFER_CREATE_INFO_NV
 - VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_MEMORY_ALLOCATE_INFO_NV

C.12.3 New Enums

None.

C.12.4 New Structures

- VkDedicatedAllocationImageCreateInfoNV
- VkDedicatedAllocationBufferCreateInfoNV
- VkDedicatedAllocationMemoryAllocateInfoNV

C.12.5 New Functions

None.

C.12.6 Issues

None.

C.12.7 Examples

```
// Create an image with
// VkDedicatedAllocationImageCreateInfoNV::dedicatedAllocation
// set to VK_TRUE
VkDedicatedAllocationImageCreateInfoNV dedicatedImageInfo =
    VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_IMAGE_CREATE_INFO_NV,
       sType
   NULL,
       pNext
                                                                             // ←
   VK_TRUE,
       dedicatedAllocation
};
VkImageCreateInfo imageCreateInfo =
   VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO, // sType
    &dedicatedImageInfo
                                           // pNext
    // Other members set as usual
};
VkImage image;
VkResult result = vkCreateImage(
   device,
    &imageCreateInfo,
   NULL,
                               // pAllocator
    &image);
VkMemoryRequirements memoryRequirements;
vkGetImageMemoryRequirements(
    device,
    image,
    &memoryRequirements);
// Allocate memory with VkDedicatedAllocationMemoryAllocateInfoNV::image
```

```
// pointing to the image we are allocating the memory for
VkDedicatedAllocationMemoryAllocateInfoNV dedicatedInfo =
    VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_MEMORY_ALLOCATE_INFO_NV,
                                                                                 // ←
                                                                                 // ←
    NULL,
        pNext
    image,
        image
                                                                                 // ←
   VK_NULL_HANDLE,
       buffer
};
VkMemoryAllocateInfo memoryAllocateInfo =
   VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO,
                                                             // sType
                                                             // pNext
    &dedicatedInfo,
                                                             // allocationSize
   memoryRequirements.size,
   FindMemoryTypeIndex(memoryRequirements.memoryTypeBits), // memoryTypeIndex
};
VkDeviceMemory memory;
vkAllocateMemory(
   device,
   &memoryAllocateInfo,
                               // pAllocator
   NULL,
   &memory);
// Bind the image to the memory
vkBindImageMemory(
   device,
    image,
   memory,
    0);
```

C.12.8 Version History

- Revision 1, 2016-05-31 (Jeff Bolz)
 - Internal revisions

C.13 VK_NV_glsl_shader

```
Name String
```

VK_NV_glsl_shader

Extension Type

Device extension

Registered Extension Number

13

Status

Draft.

Last Modified Date

2016-02-14

Revision

1

IP Status

No known IP claims.

Dependencies

• This extension is written against version 1.0. of the Vulkan API.

Contributors

• Piers Daniell, NVIDIA

Contacts

• Piers Daniell (pdaniell at nvidia.com)

This extension allows GLSL shaders written to the GL_KHR_vulkan_glsl extension specification to be used instead of SPIR-V. The implementation will automatically detect whether the shader is SPIR-V or GLSL and compile it appropriatly.

C.13.1 New Object Types

C.13.2 New Enum Constants

- Extending VkResult:
 - VK_ERROR_INVALID_SHADER_NV
- C.13.3 New Enums
- C.13.4 New Structures
- C.13.5 New Functions
- C.13.6 Issues

C.13.7 Examples

Example 1

Passing in GLSL code

```
char const vss[] =
    "#version 450 core\n"
    "layout(location = 0) in vec2 aVertex; \n"
    "layout(location = 1) in vec4 aColor; \n"
    "out vec4 vColor; \n"
    "void main()\n"
    "{\n"
        vColor = aColor; \n"
        gl_Position = vec4(aVertex, 0, 1); n"
    "}\n"
\label{eq:VkShaderModuleCreateInfo} \mbox{VkShaderInfo} = \{ \quad \hookleftarrow \quad
   VK_STRUCTURE_TYPE_SHADER_MODULE_CREATE_INFO };
vertexShaderInfo.codeSize = sizeof vss;
vertexShaderInfo.pCode = vss;
VkShaderModule vertexShader;
vkCreateShaderModule(device, &vertexShaderInfo, 0, &vertexShader);
```

C.13.8 Version History

- Revision 1, 2016-02-14 (Piers Daniell)
 - Initial draft

C.14 VK_NV_external_memory_capabilities

Name String

VK_NV_external_memory_capabilities

Extension Type

Instance extension

Registered Extension Number

56

Status

Complete

Last Modified Date

2016-08-19

Revision

1

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- Interacts with VK_NV_dedicated_allocation.

Contributors

• James Jones, NVIDIA

Contact

James Jones (jajones at nvidia.com)

Applications may wish to import memory from the Direct 3D API, or export memory to other Vulkan instances. This extension provides a set of capability queries that allow applications determine what types of win32 memory handles an implementation supports for a given set of use cases.

C.14.1 New Object Types

None.

C.14.2 New Enum Constants

None.

C.14.3 New Enums

- VkExternalMemoryHandleTypeFlagBitsNV
- VkExternalMemoryFeatureFlagBitsNV

C.14.4 New Structs

• VkExternalImageFormatPropertiesNV

C.14.5 New Functions

 $\bullet \ \ vkGetPhysicalDeviceExternalImageFormatPropertiesNV\\$

C.14.6 Issues

1) Why do so many external memory capabilities need to be queried on a per-memory-handle-type basis?

RESOLUTION: This is because some handle types are based on OS-native objects that have far more limited capabilities than the very generic Vulkan memory objects. Not all memory handle types can name memory objects that support 3D images, for example. Some handle types can not even support the deferred image and memory binding behavior of Vulkan and require specifying the image when allocating or importing the memory object.

2) Does the VkExternalImageFormatPropertiesNV struct need to include a list of memory type bits that support the given handle type?

RESOLUTION: No. The memory types that do not support the handle types will simply be filtered out of the results returned by vkGetImageMemoryRequirements() when a set of handle types was specified at image creation time.

3) Should the non-opaque handle types be moved to their own extension?

RESOLUTION: Perhaps. However, defining the handle type bits does very little and does not require any platform-specific types on its own, and it is easier to maintain the bitmask values in a single extension for now. Presumably more handle types could be added by separate extensions though, and it would be midly weird to have some platform-specific ones defined in the core spec and some in extensions

C.15 VK_NV_external_memory

Name String

VK_NV_external_memory

Extension Type

Device extension

Registered Extension Number

57

Status

Complete

Last Modified Date

2016-08-19

Revision

1

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- Requires VK_NV_external_memory_capabilities

Contributors

- James Jones, NVIDIA
- Carsten Rohde, NVIDIA

Contact

• James Jones (jajones at nvidia.com)

Applications may wish to export memory to other Vulkan instances or other APIs, or import memory from other Vulkan instances or other APIs to enable Vulkan workloads to be split up across application module, process, or API boundaries. This extension enables applications to create exportable Vulkan memory objects such that the underlying resources can be referenced outside the Vulkan instance that created them.

C.15.1 New Object Types

None.

C.15.2 New Enum Constants

Extending VkStructureType:

- VK_STRUCTURE_TYPE_EXTERNAL_MEMORY_IMAGE_CREATE_INFO_NV
- VK_STRUCTURE_TYPE_EXPORT_MEMORY_ALLOCATE_INFO_NV

C.15.3 New Enums

None.

C.15.4 New Structures

- Extending VkImageCreateInfo:
 - VkExternalMemoryImageCreateInfoNV
- Extending VkMemoryAllocateInfo
 - VkExportMemoryAllocateInfoNV

C.15.5 New Functions

None.

C.15.6 Issues

1) If memory objects are shared between processes and APIs, is this considered aliasing according to the rules outlined in section 11.8, Memory Aliasing?

RESOLUTION: Yes, but strict exceptions to the rules are added to allow some forms of aliasing in these cases. Further, other extensions may build upon these new aliasing rules to define specific support usage within Vulkan for imported native memory objects, or memory objects from other APIs.

2) Are new image layouts or metadata required to specify image layouts and layout transitions compatible with non-Vulkan APIs, or with other instances of the same Vulkan driver?

RESOLUTION: No. Separate instances of the same Vulkan driver running on the same GPU should have identical internal layout semantics, so applictions have the tools they need to ensure views of images are consistent between the two instances. Other APIs will fall into two categories: Those that are Vulkan- compatible (A term to be defined by subsequent interopability extensions), or

Vulkan incompatible. When sharing images with Vulkan incompatible APIs, the Vulkan image must be transitioned to the GENERAL layout before handing it off to the external API.

Note this does not attempt to address cross-device transitions, nor transitions to engines on the same device which are not visible within the Vulkan API. Both of these are beyond the scope of this extension.

C.15.7 Examples

```
// TODO: Write some sample code here.
```

C.15.8 Version History

- Revision 1, 2016-08-19 (James Jones)
 - Initial draft

C.16 VK_NV_external_memory_win32

Name String

VK_NV_external_memory_win32

Extension Type

Device extension

Registered Extension Number

58

Status

Complete

Last Modified Date

2016-08-19

Revision

1

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- Requires VK_NV_external_memory_capabilities
- Requires VK_NV_external_memory

Contributors

- James Jones, NVIDIA
- · Carsten Rohde, NVIDIA

Contact

• James Jones (jajones at nvidia.com)

Applications may wish to export memory to other Vulkan instances or other APIs, or import memory from other Vulkan instances or other APIs to enable Vulkan workloads to be split up across application module, process, or API boundaries. This extension enables win32 applications to export win32 handles from Vulkan memory objects such that the underlying resources can be referenced outside the Vulkan instance that created them, and import win32 handles created in the Direct3D API to Vulkan memory objects.

C.16.1 New Object Types

None.

C.16.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_IMPORT_MEMORY_WIN32_HANDLE_INFO_NV
 - VK_STRUCTURE_TYPE_EXPORT_MEMORY_WIN32_HANDLE_INFO_NV

C.16.3 New Enums

None.

C.16.4 New Structures

- Extending VkMemoryAllocateInfo
 - VkImportMemoryWin32HandleInfoNV
- Extends VkMemoryAllocateInfo
 - VkExportMemoryWin32HandleInfoNV

C.16.5 New Functions

vkGetMemoryWin32HandleNV

C.16.6 Issues

1) If memory objects are shared between processes and APIs, is this considered aliasing according to the rules outlined in section 11.8, Memory Aliasing?

RESOLUTION: Yes, but strict exceptions to the rules are added to allow some forms of aliasing in these cases. Further, other extensions may build upon these new aliasing rules to define specific support usage within Vulkan for imported native memory objects, or memory objects from other APIs.

2) Are new image layouts or metadata required to specify image layouts and layout transitions compatible with non-Vulkan APIs, or with other instances of the same Vulkan driver?

RESOLUTION: No. Separate instances of the same Vulkan driver running on the same GPU should have identical internal layout semantics, so applictions have the tools they need to ensure views of images are consistent between the two instances. Other APIs will fall into two categories: Those that are Vulkan-compatible (A term to be defined by subsequent interopability extensions), or Vulkan incompatible. When sharing images with Vulkan incompatible APIs, the Vulkan image must be transitioned to the GENERAL layout before handing it off to the external API.

Note this does not attempt to address cross-device transitions, nor transitions to engines on the same device which are not visible within the Vulkan API. Both of these are beyond the scope of this extension.

3) Do applications need to call CloseHandle() on the values returned from VkGetMemoryWin32HandleKHR() when handleType is VK_EXTERNAL_SEMAPHORE_HANDLE_TYPE_OPAQUE_WIN32_BIT_KHR?

RESOLUTION: Yes, unless it is passed back in to another driver instance to import the object. A successful get call transfers ownership of the handle to the application, while an import transfers ownership to the associated driver. Destroying the memory object will not destroy the handle or the handle's reference to the underlying memory resource.

C.16.7 Examples

```
extern VkDevice device;
VkPhysicalDeviceMemoryProperties memoryProperties;
VkExternalImageFormatPropertiesNV properties;
VkExternalMemoryImageCreateInfoNV externalMemoryImageCreateInfo;
VkDedicatedAllocationImageCreateInfoNV dedicatedImageCreateInfo;
VkImageCreateInfo imageCreateInfo;
VkImage image;
VkMemoryRequirements imageMemoryRequirements;
uint32_t numMemoryTypes;
uint32_t memoryType;
VkExportMemoryAllocateInfoNV exportMemoryAllocateInfo;
VkDedicatedAllocationMemoryAllocateInfoNV dedicatedAllocationInfo;
VkMemoryAllocateInfo memoryAllocateInfo;
VkDeviceMemory memory;
VkResult result;
HANDLE memoryHnd;
// Figure out how many memory types the device supports
vkGetPhysicalDeviceMemoryProperties(physicalDevice,
                                    &memoryProperties);
numMemoryTypes = memoryProperties.memoryTypeCount;
// Check the external handle type capabilities for the chosen format
// Exportable 2D image support with at least 1 mip level, 1 array
// layer, and VK_SAMPLE_COUNT_1_BIT using optimal tiling and supporting
// texturing and color rendering is required.
result = vkGetPhysicalDeviceExternalImageFormatPropertiesNV(
   physicalDevice,
   format,
   VK_IMAGE_TYPE_2D,
   VK_IMAGE_TILING_OPTIMAL,
   VK_IMAGE_USAGE_SAMPLED_BIT |
   VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,
    handleType,
    &properties);
if ((result != VK_SUCCESS) ||
    ! (properties.externalMemoryFeatures &
      VK_EXTERNAL_MEMORY_FEATURE_EXPORTABLE_BIT_NV)) {
    abort();
// Set up the external memory image creation info
memset (&externalMemoryImageCreateInfo,
       0, sizeof(externalMemoryImageCreateInfo));
externalMemoryImageCreateInfo.sType =
    VK_STRUCTURE_TYPE_EXTERNAL_MEMORY_IMAGE_CREATE_INFO_NV;
externalMemoryImageCreateInfo.handleTypes = handleType;
if (properties.externalMemoryFeatures &
    VK_EXTERNAL_MEMORY_FEATURE_DEDICATED_ONLY_BIT_NV) {
    memset(&dedicatedImageCreateInfo, 0, sizeof(dedicatedImageCreateInfo));
    dedicatedImageCreateInfo.sType =
        VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_IMAGE_CREATE_INFO_NV;
    dedicatedImageCreateInfo.dedicatedAllocation = VK_TRUE;
    externalMemoryImageCreateInfo.pNext = &dedicatedImageCreateInfo;
```

```
// Set up the core image creation info
memset(&imageCreateInfo, 0, sizeof(imageCreateInfo));
imageCreateInfo.sType = VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO;
imageCreateInfo.pNext = &externalMemoryImageCreateInfo;
imageCreateInfo.format = format;
imageCreateInfo.extent.width = 64;
imageCreateInfo.extent.height = 64;
imageCreateInfo.extent.depth = 1;
imageCreateInfo.mipLevels = 1;
imageCreateInfo.arrayLayers = 1;
imageCreateInfo.samples = VK_SAMPLE_COUNT_1_BIT;
imageCreateInfo.tiling = VK_IMAGE_TILING_OPTIMAL;
imageCreateInfo.usage = VK_IMAGE_USAGE_SAMPLED_BIT |
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT;
imageCreateInfo.sharingMode = VK_SHARING_MODE_EXCLUSIVE;
imageCreateInfo.initialLayout = VK_IMAGE_LAYOUT_UNDEFINED;
vkCreateImage(device, &imageCreateInfo, NULL, &image);
vkGetImageMemoryRequirements(device,
                             &imageMemoryRequirements);
// For simplicity, just pick the first compatible memory type.
for (memoryType = 0; memoryType < numMemoryTypes; memoryType++) {</pre>
    if ((1 << memoryType) & imageMemoryRequirements.memoryTypeBits) {</pre>
        break;
// At least one memory type must be supported given the prior external
// handle capability check.
assert (memoryType < numMemoryTypes);</pre>
// Allocate the external memory object.
memset(&exportMemoryAllocateInfo, 0, sizeof(exportMemoryAllocateInfo));
exportMemoryAllocateInfo.sType =
   VK_STRUCTURE_TYPE_EXPORT_MEMORY_ALLOCATE_INFO_NV;
exportMemoryAllocateInfo.handleTypes = handleType;
if (properties.externalMemoryFeatures &
   VK_EXTERNAL_MEMORY_FEATURE_DEDICATED_ONLY_BIT_NV) {
   memset(&dedicatedAllocationInfo, 0, sizeof(dedicatedAllocationInfo));
    dedicatedAllocationInfo.sType =
        VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_MEMORY_ALLOCATE_INFO_NV;
    dedicatedAllocationInfo.image = image;
    exportMemoryAllocateInfo.pNext = &dedicatedAllocationInfo;
memset(&memoryAllocateInfo, 0, sizeof(memoryAllocateInfo));
memoryAllocateInfo.sType = VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO;
memoryAllocateInfo.pNext = &exportMemoryAllocateInfo;
memoryAllocateInfo.allocationSize = imageMemoryRequirements.size;
memoryAllocateInfo.memoryTypeIndex = memoryType;
vkAllocateMemory(device, &memoryAllocateInfo, NULL, &memory);
if (!(properties.externalMemoryFeatures &
```

```
VK_EXTERNAL_MEMORY_FEATURE_DEDICATED_ONLY_BIT_NV)) {
   vkBindImageMemory(device, image, memory, 0);
}

// Get the external memory opaque FD handle
  vkGetMemoryWin32HandleNV(device, memory, &memoryHnd);
```

C.16.8 Version History

- Revision 1, 2016-08-11 (James Jones)
 - Initial draft

C.17 VK_NV_win32_keyed_mutex

Name

VK_NV_win32_keyed_mutex

Extension Type

Device extension

Registered Extension Number

59

Status

Complete

Last Modified Date

2016-08-19

Revision

2

IP Status

No known IP claims.

Dependencies

- This extension is written against version 1.0 of the Vulkan API.
- Requires VK_NV_external_memory_capabilities
- Requires VK_NV_external_memory_win32

Contributors

- · James Jones, NVIDIA
- Carsten Rohde, NVIDIA

Contact

• Carsten Rohde (crohde at nvidia.com)

Applications that wish to import Direct3D 11 memory objects into the Vulkan API may wish to use the native keyed mutex mechanism to synchronize access to the memory between Vulkan and Direct3D. This extension provides a way for an application to access the keyed mutex associated with an imported Vulkan memory object when submitting command buffers to a queue.

C.17.1 New Object Types

None.

C.17.2 New Enum Constants

- Extending VkStructureType:
 - VK_STRUCTURE_TYPE_WIN32_KEYED_MUTEX_ACQUIRE_RELEASE_INFO_NV

C.17.3 New Enums

None.

C.17.4 New Structures

- Extending VkSubmitInfo:
 - VkWin32KeyedMutexAcquireReleaseInfoNV

C.17.5 New Functions

None.

C.17.6 Issues

None.

C.17.7 Examples

```
uint32_t numMemoryTypes;
uint32_t memoryType;
VkImportMemoryWin32HandleInfoNV importMemoryInfo;
VkMemoryAllocateInfo memoryAllocateInfo;
VkDeviceMemory mem;
VkResult result;
// Figure out how many memory types the device supports
vkGetPhysicalDeviceMemoryProperties(physicalDevice,
                                    &memoryProperties);
numMemoryTypes = memoryProperties.memoryTypeCount;
// Check the external handle type capabilities for the chosen format
// Importable 2D image support with at least 1 mip level, 1 array
// layer, and VK_SAMPLE_COUNT_1_BIT using optimal tiling and supporting
// texturing and color rendering is required.
result = vkGetPhysicalDeviceExternalImageFormatPropertiesNV(
   physicalDevice,
   format,
   VK_IMAGE_TYPE_2D,
   VK_IMAGE_TILING_OPTIMAL,
   VK_IMAGE_USAGE_SAMPLED_BIT |
   VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,
   0,
   handleType,
    &properties);
if ((result != VK_SUCCESS) ||
    ! (properties.externalMemoryFeatures &
     VK_EXTERNAL_MEMORY_FEATURE_IMPORTABLE_BIT_NV)) {
   abort();
// Set up the external memory image creation info
memset(&externalMemoryImageCreateInfo,
      0, sizeof(externalMemoryImageCreateInfo));
externalMemoryImageCreateInfo.sType =
   VK_STRUCTURE_TYPE_EXTERNAL_MEMORY_IMAGE_CREATE_INFO_NV;
externalMemoryImageCreateInfo.handleTypes = handleType;
// Set up the core image creation info
memset(&imageCreateInfo, 0, sizeof(imageCreateInfo));
imageCreateInfo.sType = VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO;
imageCreateInfo.pNext = &externalMemoryImageCreateInfo;
imageCreateInfo.format = format;
imageCreateInfo.extent.width = 64;
imageCreateInfo.extent.height = 64;
imageCreateInfo.extent.depth = 1;
imageCreateInfo.mipLevels = 1;
imageCreateInfo.arrayLayers = 1;
imageCreateInfo.samples = VK_SAMPLE_COUNT_1_BIT;
imageCreateInfo.tiling = VK_IMAGE_TILING_OPTIMAL;
imageCreateInfo.usage = VK_IMAGE_USAGE_SAMPLED_BIT |
   VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT;
imageCreateInfo.sharingMode = VK_SHARING_MODE_EXCLUSIVE;
imageCreateInfo.initialLayout = VK_IMAGE_LAYOUT_UNDEFINED;
vkCreateImage(device, &imageCreateInfo, NULL, &image);
```

```
vkGetImageMemoryRequirements(device,
                              &imageMemoryRequirements);
// For simplicity, just pick the first compatible memory type.
for (memoryType = 0; memoryType < numMemoryTypes; memoryType++) {</pre>
    if ((1 << memoryType) & imageMemoryRequirements.memoryTypeBits) {</pre>
        break;
// At least one memory type must be supported given the prior external
// handle capability check.
assert (memoryType < numMemoryTypes);</pre>
// Allocate the external memory object.
memset(&exportMemoryAllocateInfo, 0, sizeof(exportMemoryAllocateInfo));
exportMemoryAllocateInfo.sType =
    VK_STRUCTURE_TYPE_EXPORT_MEMORY_ALLOCATE_INFO_NV;
importMemoryInfo.handleTypes = handleType;
importMemoryInfo.handle = sharedNtHandle;
memset(&memoryAllocateInfo, 0, sizeof(memoryAllocateInfo));
memoryAllocateInfo.sType = VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO;
memoryAllocateInfo.pNext = &exportMemoryAllocateInfo;
memoryAllocateInfo.allocationSize = imageMemoryRequirements.size;
memoryAllocateInfo.memoryTypeIndex = memoryType;
vkAllocateMemory(device, &memoryAllocateInfo, NULL, &mem);
vkBindImageMemory(device, image, mem, 0);
const uint64_t acquireKey = 1;
const uint32_t timeout = INFINITE;
const uint64_t releaseKey = 2;
VkWin32KeyedMutexAcquireReleaseInfoNV keyedMutex =
    { VK STRUCTURE TYPE WIN32 KEYED MUTEX ACQUIRE RELEASE INFO NV };
keyedMutex.acquireCount = 1;
keyedMutex.pAcquireSyncs = &mem;
keyedMutex.pAcquireKeys = &acquireKey;
keyedMutex.pAcquireTimeoutMilliseconds = &timeout;
keyedMutex.releaseCount = 1;
keyedMutex.pReleaseSyncs = &mem;
keyedMutex.pReleaseKeys = &releaseKey;
VkSubmitInfo submit_info = { VK_STRUCTURE_TYPE_SUBMIT_INFO, &keyedMutex };
submit_info.commandBufferCount = 1;
submit_info.pCommandBuffers = &cmd_buf;
vkQueueSubmit(queue, 1, &submit_info, VK_NULL_HANDLE);
```

C.17.8 Version History

- Revision 2, 2016-08-11 (James Jones)
 - Updated sample code based on the NV external memory extensions.
 - Renamed from NVX to NV extension.
 - Added Overview and Description sections.
 - Updated sample code to use the NV external memory extensions.
- Revision 1, 2016-06-14 (Carsten Rohde)
 - Initial draft.

Appendix D

API Boilerplate

This appendix defines Vulkan API features that are infrastructure required for a complete functional description of Vulkan, but do not logically belong elsewhere in the Specification.

D.1 Structure Types

Vulkan structures containing sType members must have a value of sType matching the type of the structure, as described more fully in Valid Usage for Structure Types. Structure types supported by the Vulkan API include:

```
typedef enum VkStructureType {
   VK_STRUCTURE_TYPE_APPLICATION_INFO = 0,
   VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO = 1,
   VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO = 2,
   VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO = 3,
   VK_STRUCTURE_TYPE_SUBMIT_INFO = 4,
   VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO = 5,
   VK_STRUCTURE_TYPE_MAPPED_MEMORY_RANGE = 6,
   VK_STRUCTURE_TYPE_BIND_SPARSE_INFO = 7,
   VK_STRUCTURE_TYPE_FENCE_CREATE_INFO = 8,
   VK_STRUCTURE_TYPE_SEMAPHORE_CREATE_INFO = 9,
   VK_STRUCTURE_TYPE_EVENT_CREATE_INFO = 10,
   VK_STRUCTURE_TYPE_QUERY_POOL_CREATE_INFO = 11,
   VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO = 12,
   VK_STRUCTURE_TYPE_BUFFER_VIEW_CREATE_INFO = 13,
   VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO = 14,
   VK_STRUCTURE_TYPE_IMAGE_VIEW_CREATE_INFO = 15,
   VK STRUCTURE TYPE SHADER MODULE CREATE INFO = 16,
   VK_STRUCTURE_TYPE_PIPELINE_CACHE_CREATE_INFO = 17,
   VK_STRUCTURE_TYPE_PIPELINE_SHADER_STAGE_CREATE_INFO = 18,
   VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO = 19,
   VK_STRUCTURE_TYPE_PIPELINE_INPUT_ASSEMBLY_STATE_CREATE_INFO = 20,
   VK_STRUCTURE_TYPE_PIPELINE_TESSELLATION_STATE_CREATE_INFO = 21,
   VK_STRUCTURE TYPE PIPELINE VIEWPORT_STATE_CREATE_INFO = 22,
   VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_CREATE_INFO = 23,
   VK_STRUCTURE_TYPE_PIPELINE_MULTISAMPLE_STATE_CREATE_INFO = 24,
   VK_STRUCTURE_TYPE_PIPELINE_DEPTH_STENCIL_STATE_CREATE_INFO = 25,
   VK_STRUCTURE_TYPE_PIPELINE_COLOR_BLEND_STATE_CREATE_INFO = 26,
   VK_STRUCTURE_TYPE_PIPELINE_DYNAMIC_STATE_CREATE_INFO = 27,
```

```
VK_STRUCTURE_TYPE_GRAPHICS_PIPELINE_CREATE_INFO = 28,
   VK_STRUCTURE_TYPE_COMPUTE_PIPELINE_CREATE_INFO = 29,
   VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO = 30,
   VK_STRUCTURE_TYPE_SAMPLER_CREATE_INFO = 31,
   VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO = 32,
   VK_STRUCTURE_TYPE_DESCRIPTOR_POOL_CREATE_INFO = 33,
   VK_STRUCTURE_TYPE_DESCRIPTOR_SET_ALLOCATE_INFO = 34,
   VK_STRUCTURE_TYPE_WRITE_DESCRIPTOR_SET = 35,
   VK_STRUCTURE_TYPE_COPY_DESCRIPTOR_SET = 36,
   VK_STRUCTURE_TYPE_FRAMEBUFFER_CREATE_INFO = 37,
   VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO = 38,
   VK_STRUCTURE_TYPE_COMMAND_POOL_CREATE_INFO = 39,
   VK_STRUCTURE_TYPE_COMMAND_BUFFER_ALLOCATE_INFO = 40,
   VK_STRUCTURE_TYPE_COMMAND_BUFFER_INHERITANCE_INFO = 41,
   VK_STRUCTURE_TYPE_COMMAND_BUFFER_BEGIN_INFO = 42,
   VK_STRUCTURE_TYPE_RENDER_PASS_BEGIN_INFO = 43,
   VK_STRUCTURE_TYPE_BUFFER_MEMORY_BARRIER = 44,
   VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER = 45,
   VK_STRUCTURE_TYPE_MEMORY_BARRIER = 46,
   VK_STRUCTURE_TYPE_LOADER_INSTANCE_CREATE_INFO = 47,
   VK_STRUCTURE_TYPE_LOADER_DEVICE_CREATE_INFO = 48,
   VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR = 1000001000,
   VK_STRUCTURE_TYPE_PRESENT_INFO_KHR = 1000001001,
   VK_STRUCTURE_TYPE_DISPLAY_MODE_CREATE_INFO_KHR = 1000002000,
   VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR = 1000002001,
   VK_STRUCTURE_TYPE_DISPLAY_PRESENT_INFO_KHR = 1000003000,
   VK_STRUCTURE_TYPE_XLIB_SURFACE_CREATE_INFO_KHR = 1000004000,
   VK_STRUCTURE_TYPE_XCB_SURFACE_CREATE_INFO_KHR = 1000005000,
   VK_STRUCTURE_TYPE_WAYLAND_SURFACE_CREATE_INFO_KHR = 1000006000,
   VK_STRUCTURE_TYPE_MIR_SURFACE_CREATE_INFO_KHR = 1000007000,
   VK_STRUCTURE_TYPE_ANDROID_SURFACE_CREATE_INFO_KHR = 1000008000,
   VK_STRUCTURE_TYPE_WIN32_SURFACE_CREATE_INFO_KHR = 1000009000,
   VK_STRUCTURE_TYPE_DEBUG_REPORT_CALLBACK_CREATE_INFO_EXT = 1000011000,
    {\tt VK\_STRUCTURE\_TYPE\_PIPELINE\_RASTERIZATION\_STATE\_RASTERIZATION\_ORDER\_AMD} \ = \ \ \longleftrightarrow \ \ \\
       1000018000,
   VK STRUCTURE TYPE DEBUG MARKER OBJECT NAME INFO EXT = 1000022000,
   VK_STRUCTURE_TYPE_DEBUG_MARKER_OBJECT_TAG_INFO_EXT = 1000022001,
   VK_STRUCTURE_TYPE_DEBUG_MARKER_MARKER_INFO_EXT = 1000022002,
   VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_IMAGE_CREATE_INFO_NV = 1000026000,
   VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_BUFFER_CREATE_INFO_NV = 1000026001,
   VK_STRUCTURE_TYPE_DEDICATED_ALLOCATION_MEMORY_ALLOCATE_INFO_NV = 1000026002,
   VK_STRUCTURE_TYPE_EXTERNAL_MEMORY_IMAGE_CREATE_INFO_NV = 1000056000,
   VK_STRUCTURE_TYPE_EXPORT_MEMORY_ALLOCATE_INFO_NV = 1000056001,
   VK_STRUCTURE_TYPE_IMPORT_MEMORY_WIN32_HANDLE_INFO_NV = 1000057000,
   VK_STRUCTURE_TYPE_EXPORT_MEMORY_WIN32_HANDLE_INFO_NV = 1000057001,
   VK_STRUCTURE_TYPE_WIN32_KEYED_MUTEX_ACQUIRE_RELEASE_INFO_NV = 1000058000,
} VkStructureType;
```

D.2 Flag Types

Vulkan flag types are all bitmasks aliasing the base type VkFlags and with corresponding bit flag types defining the valid bits for that flag, as described in Valid Usage for Flags. Flag types supported by the Vulkan API include:

```
typedef VkFlags VkAccessFlags;
```

```
typedef VkFlags VkAttachmentDescriptionFlags;
typedef VkFlags VkBufferCreateFlags;
typedef VkFlags VkBufferUsageFlags;
typedef VkFlags VkBufferViewCreateFlags;
typedef VkFlags VkColorComponentFlags;
typedef VkFlags VkCommandBufferResetFlags;
typedef VkFlags VkCommandBufferUsageFlags;
typedef VkFlags VkCommandPoolCreateFlags;
typedef VkFlags VkCommandPoolResetFlags;
typedef VkFlags VkCullModeFlags;
typedef VkFlags VkDependencyFlags;
typedef VkFlags VkDescriptorPoolCreateFlags;
typedef VkFlags VkDescriptorPoolResetFlags;
typedef VkFlags VkDescriptorSetLayoutCreateFlags;
typedef VkFlags VkDeviceCreateFlags;
typedef VkFlags VkDeviceQueueCreateFlags;
typedef VkFlags VkEventCreateFlags;
typedef VkFlags VkFenceCreateFlags;
typedef VkFlags VkFormatFeatureFlags;
typedef VkFlags VkFramebufferCreateFlags;
typedef VkFlags VkImageAspectFlags;
typedef VkFlags VkImageCreateFlags;
typedef VkFlags VkImageUsageFlags;
typedef VkFlags VkImageViewCreateFlags;
```

```
typedef VkFlags VkInstanceCreateFlags;
typedef VkFlags VkMemoryHeapFlags;
typedef VkFlags VkMemoryMapFlags;
typedef VkFlags VkMemoryPropertyFlags;
typedef VkFlags VkPipelineCacheCreateFlags;
typedef VkFlags VkPipelineColorBlendStateCreateFlags;
typedef VkFlags VkPipelineCreateFlags;
typedef VkFlags VkPipelineDepthStencilStateCreateFlags;
typedef VkFlags VkPipelineDynamicStateCreateFlags;
typedef VkFlags VkPipelineInputAssemblyStateCreateFlags;
typedef VkFlags VkPipelineLayoutCreateFlags;
typedef VkFlags VkPipelineMultisampleStateCreateFlags;
typedef VkFlags VkPipelineRasterizationStateCreateFlags;
typedef VkFlags VkPipelineShaderStageCreateFlags;
typedef VkFlags VkPipelineStageFlags;
typedef VkFlags VkPipelineTessellationStateCreateFlags;
typedef VkFlags VkPipelineVertexInputStateCreateFlags;
typedef VkFlags VkPipelineViewportStateCreateFlags;
typedef VkFlags VkQueryControlFlags;
typedef VkFlags VkQueryPipelineStatisticFlags;
typedef VkFlags VkQueryPoolCreateFlags;
typedef VkFlags VkQueryResultFlags;
typedef VkFlags VkQueueFlags;
typedef VkFlags VkRenderPassCreateFlags;
```

```
typedef VkFlags VkSampleCountFlags;

typedef VkFlags VkSamplerCreateFlags;

typedef VkFlags VkSemaphoreCreateFlags;

typedef VkFlags VkShaderModuleCreateFlags;

typedef VkFlags VkShaderStageFlags;

typedef VkFlags VkSparseImageFormatFlags;

typedef VkFlags VkSparseMemoryBindFlags;

typedef VkFlags VkStencilFaceFlags;

typedef VkFlags VkSubpassDescriptionFlags;
```

D.3 Macro Definitions in vulkan.h

D.3.1 Vulkan Version Number Macros

API Version Numbers are packed into integers. These macros manipulate version numbers in useful ways.

VK_VERSION_MAJOR extracts the API major version number from a packed version number:

```
#define VK_VERSION_MAJOR(version) ((uint32_t)(version) >> 22)
```

VK_VERSION_MINOR extracts the API minor version number from a packed version number:

```
#define VK_VERSION_MINOR(version) (((uint32_t) (version) >> 12) & 0x3ff)
```

VK_VERSION_PATCH extracts the API patch version number from a packed version number:

```
#define VK_VERSION_PATCH(version) ((uint32_t)(version) & 0xfff)
```

VK_API_VERSION_1_0 returns the API version number for Vulkan 1.0. The patch version number in this macro will always be zero. The supported patch version for a physical device can be queried with vkGetPhysicalDeviceProperties.

```
// Vulkan 1.0 version number
#define VK_API_VERSION_1_0 VK_MAKE_VERSION(1, 0, 0)
```

VK_API_VERSION is now commented out of vulkan.h and cannot be used.

```
// DEPRECATED: This define has been removed. Specific version defines (e.g. \leftarrow VK_API_VERSION_1_0), or the VK_MAKE_VERSION macro, should be used instead. 
//#define VK_API_VERSION VK_MAKE_VERSION(1, 0, 0)
```

VK MAKE VERSION constructs an API version number.

```
#define VK_MAKE_VERSION(major, minor, patch) \
   (((major) << 22) | ((minor) << 12) | (patch))</pre>
```

- major is the major version number.
- minor is the minor version number.
- patch is the patch version number.

This macro can be used when constructing the VkApplicationInfo::apiVersion parameter passed to vkCreateInstance.

D.3.2 Vulkan Header File Version Number

VK_HEADER_VERSION is the version number of the vulkan.h header. This value is currently kept synchronized with the release number of the Specification. However, it is not guaranteed to remain synchronized, since most Specification updates have no effect on vulkan.h.

```
// Version of this file
#define VK_HEADER_VERSION 27
```

D.3.3 Vulkan Handle macros

VK_DEFINE_HANDLE defines a dispatchable handle type.

```
#define VK_DEFINE_HANDLE(object) typedef struct object##_T* object;
```

• object is the name of the resulting C type.

The only dispatchable handle types are those related to device and instance management, such as VkDevice.

VK_DEFINE_NON_DISPATCHABLE_HANDLE defines a non-dispatchable handle type.

• object is the name of the resulting C type.

Most Vulkan handle types, such as VkBuffer, are non-dispatchable.

Note



The vulkan.h header allows the VK_DEFINE_NON_DISPATCHABLE_HANDLE definition to be overridden by the application. If VK_DEFINE_NON_DISPATCHABLE_HANDLE is already defined when the vulkan.h header is compiled the default definition is skipped. This allows the application to define a binary-compatible custom handle which may provide more type-safety or other features needed by the application. Behavior is undefined if the application defines a non-binary-compatible handle and may result in memory corruption or application termination. Binary compatibility is platform dependent so the application must be careful if it overrides the default VK_DEFINE_NON_DISPATCHABLE_HANDLE definition.

VK_NULL_HANDLE is a reserved value representing a non-valid object handle. It may be passed to and returned from Vulkan commands only when specifically allowed.

#define VK_NULL_HANDLE 0

D.4 Platform-Specific Macro Definitions in vk_platform.h

In addition to the macros described for vulkan.h, platform-specific macros specified and used in the included vk_platform.h file are described in this section. These macros are specifically used to control platform-dependent behavior and their exact definitions are under the control of specific platforms and Vulkan implementations.

D.4.1 Platform-Specific Calling Conventions

VKAPI_ATTR is a macro placed before the return type in Vulkan API function declarations. If not empty, the interpretation of this macro depends on the platform and compiler in use, but normally controls calling conventions for C++11 and GCC/Clang-style compilers.

VKAPI_CALL is a macro placed after the return type in Vulkan API function declarations. If not empty, the interpretation of this macro depends on the platform and compiler in use, but normally controls calling conventions for MSVC-style compilers.

VKAPI_PTR is a macro placed between the (and * in Vulkan API function pointer declarations. If not empty, the interpretation of this macro depends on the platform and compiler in use, and normally controls calling conventions. VKAPI_PTR typically has the same definition as VKAPI_ATTR or VKAPI_CALL, depending on the compiler.

D.4.2 Platform-Specific Header Control

If the $VK_NO_STDINT_H$ macro is defined by the application at compile time, extended integer types required by vulkan.h, such as $uint8_t$, must also be defined by the application. Otherwise, vulkan.h will not compile. If $VK_NO_STDINT_H$ is not defined, the system stdint.h is used to define these types, or there is a fallback path when Microsoft Visual Studio version 2008 and earlier versions are detected at compile time.

D.4.3 Window System-Specific Header Control

When using different window systems with Khronos extensions, header files for those window systems must be included at compile time in order for the corresponding extension definitions to compile. The Vulkan header files cannot determine whether or not an external header is available at compile time, so applications must include macros enabling those headers. If this is not done, the corresponding extension interfaces will not be defined and they will be unusable.

The extensions, required compile-time symbols to enable them, and window systems they correspond to are defined in the following table.

Table D.1: Window System Extensions and Required Compile-Time Symbol Definitions

Extension Name	Required Compile-Time Symbol	Window System Name
VK_KHR_android_surface	VK_USE_PLATFORM_ANDROID_	Android Native
	KHR	
VK_KHR_mir_surface	VK_USE_PLATFORM_MIR_KHR	Mir
VK_KHR_wayland_surface	VK_USE_PLATFORM_WAYLAND_	Wayland
	KHR	
VK_KHR_win32_surface	VK_USE_PLATFORM_WIN32_	Microsoft Windows
	KHR	
VK_KHR_xcb_surface	VK_USE_PLATFORM_XCB_KHR	X Window System Xcb library
VK_KHR_xlib_surface	VK_USE_PLATFORM_XLIB_KHR	X Window System Xlib library

Appendix E

Invariance

The Vulkan specification is not pixel exact. It therefore does not guarantee an exact match between images produced by different Vulkan implementations. However, the specification does specify exact matches, in some cases, for images produced by the same implementation. The purpose of this appendix is to identify and provide justification for those cases that require exact matches.

E.1 Repeatability

The obvious and most fundamental case is repeated issuance of a series of Vulkan commands. For any given Vulkan and framebuffer state vector, and for any Vulkan command, the resulting Vulkan and framebuffer state must be identical whenever the command is executed on that initial Vulkan and framebuffer state. This repeatability requirement does not apply when using shaders containing side effects (image and buffer variable stores and atomic operations), because these memory operations are not guaranteed to be processed in a defined order.

The repeatability requirement does not apply for rendering done using a graphics pipeline that uses VK_RASTERIZATION_ORDER_RELAXED_AMD.

One purpose of repeatability is avoidance of visual artifacts when a double-buffered scene is redrawn. If rendering is not repeatable, swapping between two buffers rendered with the same command sequence may result in visible changes in the image. Such false motion is distracting to the viewer. Another reason for repeatability is testability.

Repeatability, while important, is a weak requirement. Given only repeatability as a requirement, two scenes rendered with one (small) polygon changed in position might differ at every pixel. Such a difference, while within the law of repeatability, is certainly not within its spirit. Additional invariance rules are desirable to ensure useful operation.

E.2 Multi-pass Algorithms

Invariance is necessary for a whole set of useful multi-pass algorithms. Such algorithms render multiple times, each time with a different Vulkan mode vector, to eventually produce a result in the framebuffer. Examples of these algorithms include:

- "Erasing" a primitive from the framebuffer by redrawing it, either in a different color or using the XOR logical operation.
- Using stencil operations to compute capping planes.

E.3 Invariance Rules

For a given instantiation of an Vulkan rendering context:

Rule 1 For any given Vulkan and framebuffer state vector, and for any given Vulkan command, the resulting Vulkan and framebuffer state must be identical each time the command is executed on that initial Vulkan and framebuffer state.

Rule 2 Changes to the following state values have no side effects (the use of any other state value is not affected by the change):

Required:

- Framebuffer contents (all bit planes)
- The color buffers enabled for writing
- Scissor parameters (other than enable)
- Writemasks (color, depth, stencil)
- Clear values (color, depth, stencil)

Strongly suggested:

- Stencil Parameters (other than enable)
- Depth test parameters (other than enable)
- Blend parameters (other than enable)
- Logical operation parameters (other than enable)
- Pixel storage state

Corollary 1 Fragment generation is invariant with respect to the state values listed in Rule 2.

Rule 3 The arithmetic of each per-fragment operation is invariant except with respect to parameters that directly control it.

Corollary 2 *Images rendered into different color buffers sharing the same framebuffer, either simultaneously or separately using the same command sequence, are pixel identical.*

Rule 4 The same vertex or fragment shader will produce the same result when run multiple times with the same input. The wording "the same shader" means a program object that is populated with the same SPIR-V binary, which is used to create pipelines, possibly multiple times, and which program object is then executed using the same Vulkan state vector. Invariance is relaxed for shaders with side effects, such as performing stores or atomics.

Rule 5 All fragment shaders that either conditionally or unconditionally assign FragCoord.z to FragDepth are depth-invariant with respect to each other, for those fragments where the assignment to FragDepth actually is done.

If a sequence of Vulkan commands specifies primitives to be rendered with shaders containing side effects (image and buffer variable stores and atomic operations), invariance rules are relaxed. In particular, rule 1, corollary 2, and rule 4 do not apply in the presence of shader side effects.

The following weaker versions of rules 1 and 4 apply to Vulkan commands involving shader side effects:

Rule 6 For any given Vulkan and framebuffer state vector, and for any given Vulkan command, the contents of any framebuffer state not directly or indirectly affected by results of shader image or buffer variable stores or atomic operations must be identical each time the command is executed on that initial Vulkan and framebuffer state.

Rule 7 The same vertex or fragment shader will produce the same result when run multiple times with the same input as long as:

- shader invocations do not use image atomic operations;
- no framebuffer memory is written to more than once by image stores, unless all such stores write the same value; and
- no shader invocation, or other operation performed to process the sequence of commands, reads memory written to by an image store.

When any sequence of Vulkan commands triggers shader invocations that perform image stores or atomic operations, and subsequent Vulkan commands read the memory written by those shader invocations, these operations must be explicitly synchronized.

E.4 Tessellation Invariance

When using a program containing tessellation evaluation shaders, the fixed-function tessellation primitive generator consumes the input patch specified by an application and emits a new set of primitives. The following invariance rules are intended to provide repeatability guarantees. Additionally, they are intended to allow an application with a carefully crafted tessellation evaluation shader to ensure that the sets of triangles generated for two adjacent patches have identical vertices along shared patch edges, avoiding "cracks" caused by minor differences in the positions of vertices along shared edges.

Rule 1 When processing two patches with identical outer and inner tessellation levels, the tessellation primitive generator will emit an identical set of point, line, or triangle primitives as long as the active program used to process the patch primitives has tessellation evaluation shaders specifying the same tessellation mode, spacing, vertex order, and point mode decorations. Two sets of primitives are considered identical if and only if they contain the same number and type of primitives and the generated tessellation coordinates for the vertex numbered m of the primitive numbered n are identical for all values of m and n.

Rule 2 The set of vertices generated along the outer edge of the subdivided primitive in triangle and quad tessellation, and the tessellation coordinates of each, depends only on the corresponding outer tessellation level and the spacing decorations in the tessellation shaders of the pipeline.

Rule 3 The set of vertices generated when subdividing any outer primitive edge is always symmetric. For triangle tessellation, if the subdivision generates a vertex with tessellation coordinates of the form (0, x, 1-x), (x, 0, 1-x), or (x, 1-x, 0), it will also generate a vertex with coordinates of exactly (0, 1-x, x), (1-x, 0, x), or (1-x, x, 0), respectively. For quad tessellation, if the subdivision generates a vertex with coordinates of (x, 0) or (0, x), it will also generate a vertex with coordinates of exactly (1-x, 0) or (0, 1-x), respectively. For isoline tessellation, if it generates vertices at (0, x) and (1, x) where x is not zero, it will also generate vertices at exactly (0, 1-x) and (1, 1-x), respectively.

Rule 4 The set of vertices generated when subdividing outer edges in triangular and quad tessellation must be independent of the specific edge subdivided, given identical outer tessellation levels and spacing. For example, if vertices at (x, 1 - x, 0) and (1-x, x, 0) are generated when subdividing the w = 0 edge in triangular tessellation, vertices must be generated at (x, 0, 1-x) and (1-x, 0, x) when subdividing an otherwise identical v = 0 edge. For quad tessellation, if vertices at (x, 0) and (1-x, 0) are generated when subdividing the v = 0 edge, vertices must be generated at (0, x) and (0, 1-x) when subdividing an otherwise identical v = 0 edge.

Rule 5 When processing two patches that are identical in all respects enumerated in rule 1 except for vertex order, the set of triangles generated for triangle and quad tessellation must be identical except for vertex and triangle order. For each triangle n1 produced by processing the first patch, there must be a triangle n2 produced when processing the second patch each of whose vertices has the same tessellation coordinates as one of the vertices in n1.

Rule 6 When processing two patches that are identical in all respects enumerated in rule 1 other than matching outer tessellation levels and/or vertex order, the set of interior triangles generated for triangle and quad tessellation must be identical in all respects except for vertex and triangle order. For each interior triangle n1 produced by processing the first patch, there must be a triangle n2 produced when processing the second patch each of whose vertices has the same

tessellation coordinates as one of the vertices in n1. A triangle produced by the tessellator is considered an interior triangle if none of its vertices lie on an outer edge of the subdivided primitive.

Rule 7 For quad and triangle tessellation, the set of triangles connecting an inner and outer edge depends only on the inner and outer tessellation levels corresponding to that edge and the spacing decorations.

Rule 8 The value of all defined components of **TessCoord** will be in the range [0, 1]. Additionally, for any defined component x of **TessCoord**, the results of computing 1.0-x in a tessellation evaluation shader will be exact. If any floating-point values in the range [0, 1] fail to satisfy this property, such values must not be used as tessellation coordinate components.

Appendix F

Credits

Vulkan 1.0 is the result of contributions from many people and companies participating in the Khronos Vulkan Working Group, as well as input from the Vulkan Advisory Panel.

Members of the Working Group, including the company that they represented at the time of their contributions, are listed below. Some specific contributions made by individuals are listed together with their name.

- · Adam Jackson, Red Hat
- Adam Śmigielski, Mobica
- Alex Bourd, Qualcomm Technologies, Inc.
- · Alexander Galazin, ARM
- · Allen Hux, Intel
- Alon Or-bach, Samsung Electronics (WSI technical sub-group chair)
- Andrew Cox, Samsung Electronics
- Andrew Garrard, Samsung Electronics (format wrangler)
- Andrew Poole, Samsung Electronics
- · Andrew Rafter, Samsung Electronics
- Andrew Richards, Codeplay Software Ltd.
- · Andrew Woloszyn, Google
- Antoine Labour, Google
- · Aras Pranckevičius, Unity
- Ashwin Kolhe, NVIDIA
- Ben Bowman, Imagination Technologies
- Benj Lipchak
- Bill Hollings, The Brenwill Workshop

- Bill Licea-Kane, Qualcomm Technologies, Inc.
- Brent E. Insko, Intel
- Brian Ellis, Qualcomm Technologies, Inc.
- · Cass Everitt, Oculus VR
- Cemil Azizoglu, Canonical
- Chad Versace, Intel
- Chang-Hyo Yu, Samsung Electronics
- Chia-I Wu, LunarG
- Chris Frascati, Qualcomm Technologies, Inc.
- Christophe Riccio, Unity
- Cody Northrop, LunarG
- Courtney Goeltzenleuchter, LunarG
- Damien Leone, NVIDIA
- Dan Baker, Oxide Games
- Dan Ginsburg, Valve
- Daniel Johnston, Intel
- Daniel Koch, NVIDIA (Shader Interfaces; Features)
- Daniel Rakos, AMD
- David Airlie, Red Hat
- David Neto, Google
- · David Mao, AMD
- David Yu, Pixar
- Dominik Witczak, AMD
- Frank (LingJun) Chen, Qualcomm Technologies, Inc.
- Fred Liao, Mediatek
- Gabe Dagani, Freescale
- Graeme Leese, Broadcom
- Graham Connor, Imagination Technologies
- Graham Sellers, AMD
- Hwanyong Lee, Kyungpook National University
- Ian Elliott, LunarG
- Ian Romanick, Intel

- James Jones, NVIDIA
- James Hughes, Oculus VR
- Jan Hermes, Continental Corporation
- · Jan-Harald Fredriksen, ARM
- Jason Ekstrand, Intel
- Jeff Bolz, NVIDIA (extensive contributions, exhaustive review and rewrites for technical correctness)
- Jeff Juliano, NVIDIA
- Jeff Vigil, Qualcomm Technologies, Inc.
- · Jens Owen, LunarG
- Jeremy Hayes, LunarG
- · Jesse Barker, ARM
- Jesse Hall, Google
- Johannes van Waveren, Oculus VR
- John Kessenich, Google (SPIR-V and GLSL for Vulkan spec author)
- John McDonald, Valve
- Jon Ashburn, LunarG
- Jon Leech, Independent (XML toolchain, normative language, release wrangler)
- Jonas Gustavsson, Sony Mobile
- Jonathan Hamilton, Imagination Technologies
- Jungwoo Kim, Samsung Electronics
- Kenneth Benzie, Codeplay Software Ltd.
- Kerch Holt, NVIDIA (SPIR-V technical sub-group chair)
- Kristian Kristensen, Intel
- Krzysztof Iwanicki, Samsung Electronics
- Larry Seiler, Intel
- Lutz Latta, Lucasfilm
- Maria Rovatsou, Codeplay Software Ltd.
- · Mark Callow
- · Mark Lobodzinski, LunarG
- Mateusz Przybylski, Intel
- Mathias Heyer, NVIDIA
- Mathias Schott, NVIDIA

- Maxim Lukyanov, Samsung Electronics
- Maurice Ribble, Qualcomm Technologies, Inc.
- Michael Lentine, Google
- Michael Worcester, Imagination Technologies
- Michal Pietrasiuk, Intel
- · Mika Isojarvi, Google
- Mike Stroyan, LunarG
- Minyoung Son, Samsung Electronics
- Mitch Singer, AMD
- Mythri Venugopal, Samsung Electronics
- Naveen Leekha, Google
- Neil Henning, Codeplay Software Ltd.
- Neil Trevett, NVIDIA
- Nick Penwarden, Epic Games
- Niklas Smedberg, Epic Games
- Norbert Nopper, Freescale
- Pat Brown, NVIDIA
- · Patrick Doane, Blizzard Entertainment
- Peter Lohrmann, Valve
- Pierre Boudier, NVIDIA
- Pierre-Loup A. Griffais, Valve
- Piers Daniell, NVIDIA (dynamic state, copy commands, memory types)
- Piotr Bialecki, Intel
- Prabindh Sundareson, Samsung Electronics
- Pyry Haulos, Google (Vulkan conformance test subcommittee chair)
- Ray Smith, ARM
- Rob Stepinski, Transgaming
- Robert J. Simpson, Qualcomm Technologies, Inc.
- Rolando Caloca Olivares, Epic Games
- Roy Ju, Mediatek
- Rufus Hamede, Imagination Technologies
- Sean Ellis, ARM

- Sean Harmer, KDAB
- Shannon Woods, Google
- · Slawomir Cygan, Intel
- · Slawomir Grajewski, Intel
- Stefanus Du Toit, Google
- · Steve Hill, Broadcom
- Steve Viggers, Core Avionics & Industrial Inc.
- Stuart Smith, Imagination Technologies
- Tim Foley, Intel
- Timo Suoranta, AMD
- Timothy Lottes, AMD
- Tobias Hector, Imagination Technologies (validity language and toolchain)
- Tobin Ehlis, LunarG
- Tom Olson, ARM (working group chair)
- Tomasz Kubale, Intel
- Tony Barbour, LunarG
- Wayne Lister, Imagination Technologies
- · Yanjun Zhang, Vivante
- Zhenghong Wang, Mediatek

In addition to the Working Group, the Vulkan Advisory Panel members provided important real-world usage information and advice that helped guide design decisions.

Administrative support to the Working Group was provided by members of Gold Standard Group, including Andrew Riegel, Elizabeth Riegel, Glenn Fredericks, Kathleen Mattson and Michelle Clark. Technical support was provided by James Riordon, webmaster of Khronos.org and OpenGL.org.